Distributed transaction.

T_1: Transfer

```
x = GET(X)
y = GET(Y)
if(x>10){
  PUT(X, x-10)
  PUT(Y, y+10)
}
```

T_2: Audit

```
x = GET(X)
y = GET(Y)
print(x+y)
```

ACID guarantees:

- Atomicity: All parts of txn execute or none (x's balance decreases, y's balance does not increase)
- Consistency: Preserves Invariants. (eg. x's balance >0)
- Durability: Txn's effect are not lost (even if servers restart)

-Isolation-

T_1: Transfer

```
x = GET(X)
y = GET(Y)
PUT(X, x-10)
PUT(Y, y+10)
```

T_2: Audit

```
x = GET(X)
y = GET(Y)
print(x+y)
```

what if T_2 executes here? Prints 190

what if T_1 executes here? Prints 210

- Serializability

x: 90 y: 110 P 200

or

Strict serializability

If T_2 started after T_1 committed, then.

- T_2 Stale Read

Rx 100 Ry100 P200

Wy110

T_2 Rx100 Ry100 Wx90

Serializable but not linearizable.

T_2 Rx90 Ry100 P190

Wy110

T_1 Rx100 Ry100 Wx90

Linearizable but not serializable

- Start with single machine

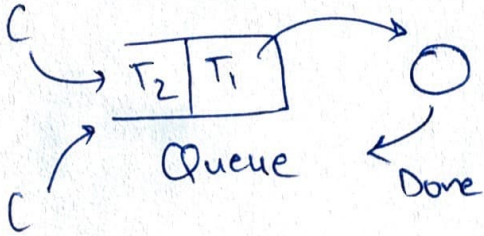

Pull one transaction at a time.

Trivially serializable!

✓ Queue order determines serial order

✗ Bad perf: uses only one CPU. Could have run independent Txn's in parallel

- Multi Cores

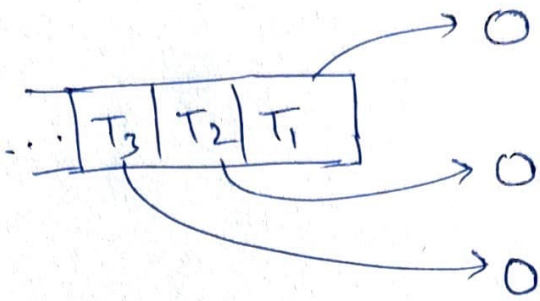

Optimistic
Pessimistic
multi version

Better perf. May not be serializable
(T_2 executes in middle of T_1)
⇒ Need concurrency control mechanisms

- Optimistic concurrency control (Redis)

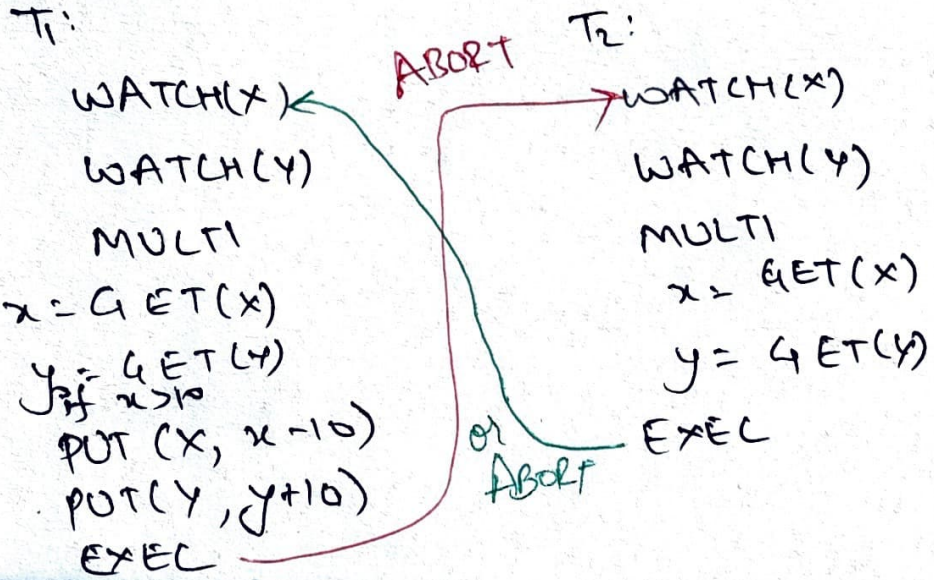

- Pessimistic concurrency control

T_1

```
rlock(X,Y)
x = GET(X)
y = GET(Y)
if x>10 {
wlock(X,Y)
PUT(X, x+10)
PUT(Y, y-10)
}
release(X,Y)
```

T_2

```
rlock(X,Y)
x = GET(X)
y = GET(Y)
release(X,Y)
Print(x+y)
```

- Wound wait deadlock avoidance.

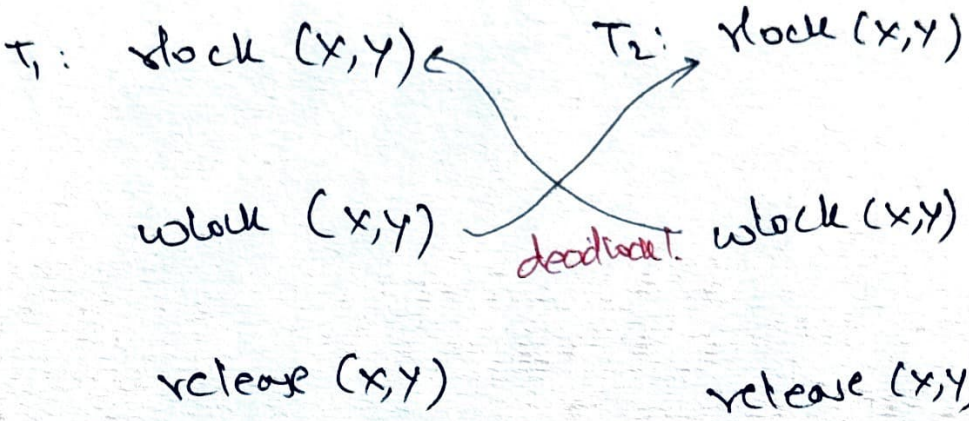

First to wlock, aborts other transaction.

- Optimistic
  - If high conflicts, keep aborting and restarting
  - Need to maintain writes locally that all commit atomically.
- Pessimistic
  - Unnecessary locking if no conflicts
  - Deadlock avoidance is needed

- Challenge: data does not fit on a single machine (eg. X and Y)
- Sharding

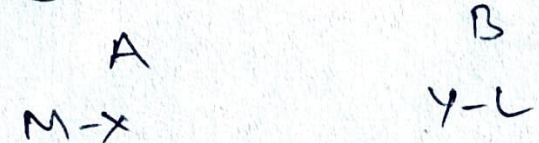

- Might also shard for performance. Transfer from M→N and F→G can happen in parallel on separate machines.

- ⇒ Reads, writes, locks need to happen on separate machines

A (owns x)	B (owns y)
rlock(x)	rlock(y)
x = GET(x)	y = GET(y)
if x > 10	wlock(y)
wlock(x)	PUT(y, y-10)
PUT(x, x-10)	release(y)
release(x)	

- Bad commit protocol

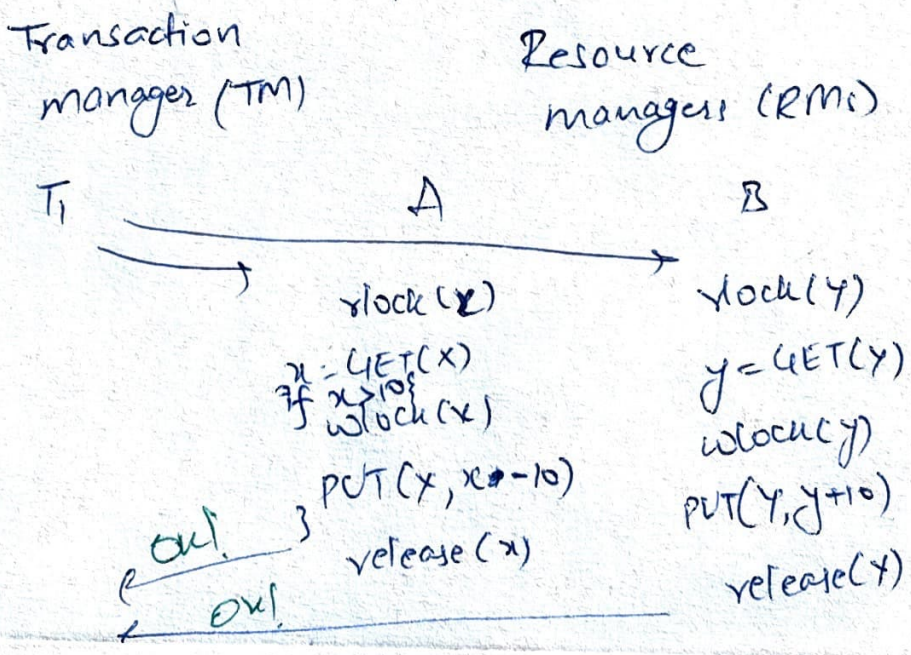

- what can go wrong?
  - Not enough money in x
  - y account does not exist
  - A or B crashes before receiving msgs
  - Network fails
  - TM crashes after sending txn to A but before " " " B

- Safety?

Atomic commits: Everyone commits or everyone aborts.

Keep aborting forever?

- Liveness

- If no failures, A, B can commit, then commit
- If failures, reach a conclusion ASAP

- R/W transactions can be thought of as 2 phases

```
T1: rlock(X,Y)
    x = GET(X)
    assert(x>10)
    y = GET(Y)
    wlock(X,Y)
    PUT(X, x-10)
    PUT(Y, y+10)
    release(X,Y)
```

Prepare phase: Read all values. Take all locks. No writes!

Commit phase: Write and release locks. Write both x,y or none. Must be atomic.

```
PREPARE(x):
    rlock(x)
    x = GET(x)
    assert(x>10)
    wlock(x)
    return x.

COMMIT(x)
    PUT(x, x+10)
    release(x)

ABORT(x)
    release(x)
```

Atomicity

Two phase Commits

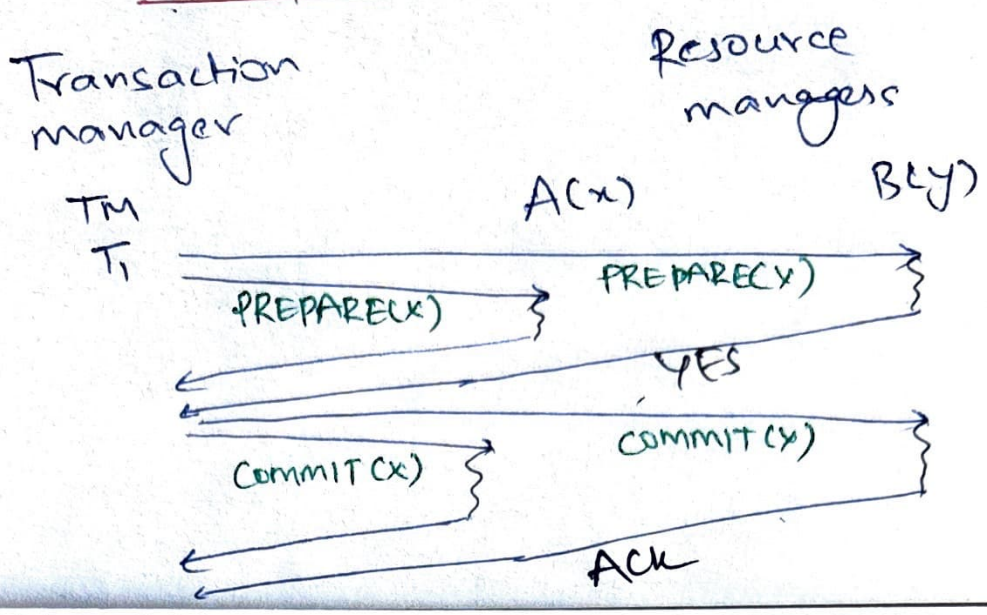

- Why does it give atomic commits?
  - TM can send commit only if it has heard Yes from all RMs.
    - All or nothing
  - Ex: If B cannot lock(y), it replies NO
    - ⇒ TM abort transaction.

- B crashes before sending YES to TM.
  - TM timeouts and unilaterally aborts.
  - or n/w lost YES message

- B crashes after sending YES to TM
  - TM sends commit to A.
  - B restarts
    - must remember it was in middle of Txn. $PREPARE_{T_1}(Y)$ i.e, wlock(Y)
    - TM keeps retrying commit(T_1)
    - Is B guaranteed to get wlock(Y)?

WAL $Prepare_{T_1}(Y) \rightarrow$ Yes

- What if TM restarts before sending prepare to B?
  - Send prepare again.
  - B prepares
  - A should remember it was already prepared and reply YES

- What if TM restarts after prepares?
  - If participant had replied Yes, it is blocked. waiting for commit/abort
  - After restart, TM must commit/abort all pending transactions

- TM log
  - <Txn ID> <details>
  - <commit> <Txn ID>
    or
    <abort> <Txn ID>

- Why is it ok to not log -
  - Sent prepare to A?
    - can just resend prepares
  - received yes from A?
    - can unilaterally abort
  - Sent commit to A?
    - can just resend commit.

- RM/ Participant log
  - <Txn ID> <details> <prepared>
  - " <committed>
    or <aborted>

Safety -

- No commit unless everyone says Yes
- RM cannot back out after saying Yes
  across restarts

Liveness

- Not live if TM crashes forever after prepare (or becomes unreachable)