

○ Distributed transactions

T_1 : Transfer

$x = \text{GET}(x)$

$y = \text{GET}(y)$

$\text{PUT}(x, x-10)$

$\text{PUT}(y, y+10)$

T_2 : Audit

$x = \text{GET}(x)$

$y = \text{GET}(y)$

$\text{print}(x+y)$

● ACID guarantees:

- Atomicity: All parts of txn execute or none (x's balance decreases, y's balance does not increase)
- Consistency: Preserves invariants. (eg. x's balance > 0)
- Durability: Txn's effect are not lost (even if servers restart)

- Isolation-

T_1 : Transfer

$x = \text{GET}(x)$

$y = \text{GET}(y)$

$\text{PUT}(x, x-10)$

$\text{PUT}(y, y+10)$ → what if T_2 executes here?

T_2 : Audit

$x = \text{GET}(x)$

$y = \text{GET}(y)$

$\text{print}(x+y)$

what if T_1 executes here?

prints 210

prints 190

Serializability

T_1 T_2

$x:90$ $y:110$ $P200$

or
 T_2 T_1

strict serializability
if T_2 started after T_1 committed, then.

T_1 T_2

T_2

State Read

$Rx100$ $Ry100$ $P200$

T_2 $Rx100$ $Ry100$ $wx90$

$wy110$

Serializable but not linearizable.

T_2


$Rx90$ $Ry100$ $P190$

T_1 $Rx100$ $Ry100$ $wx90$

$wy110$

Linearizable but not serializable

Chain Replication?



Bad perf: one at a time transaction

or
 T_2 T_1

Challenge: data does not fit on a single machine (eg. x and y)

Sharding

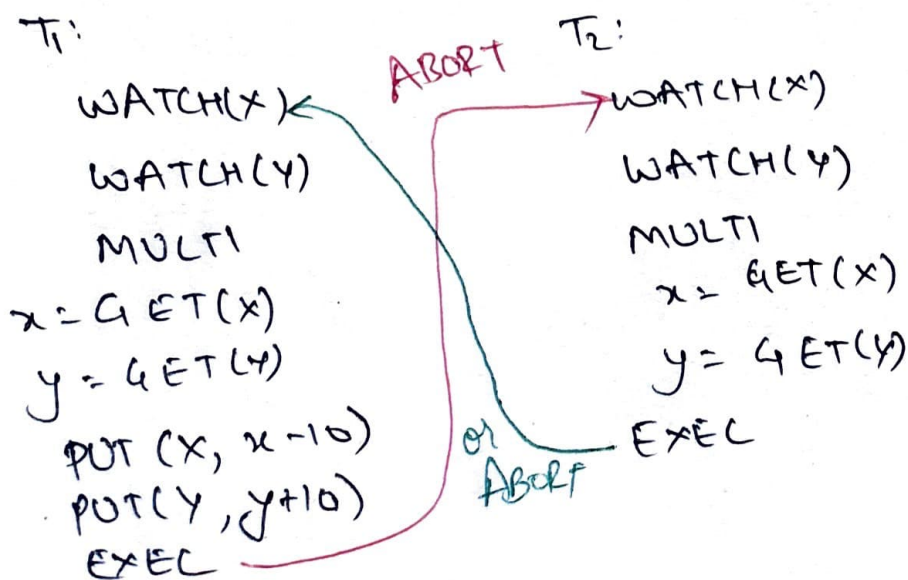
$A_1 \rightarrow A_2 \rightarrow A_3$ M-X	$B_1 \rightarrow B_2 \rightarrow B_3$ Y-L
--	--

- Might also shard for performance.
Transfer from M \rightarrow N, F \rightarrow G can happen
in parallel on separate machines.

Isolation \rightarrow Concurrency control
(optimistic, pessimistic, multi version)

Atomicity \rightarrow Atomic commit
(2 phase commits)

Optimistic concurrency control (Redis)



● Pessimistic concurrency control

T_1

```

lock (X, Y)
x = GET(X)
y = GET(Y)
unlock (X, Y)
PUT (x, x+10)
PUT (y, y-10)
release (X, Y)
    
```

T_2

```

lock (X, Y)
x = GET(X)
y = GET(Y)
release (X, Y)
Print (x+y)
    
```

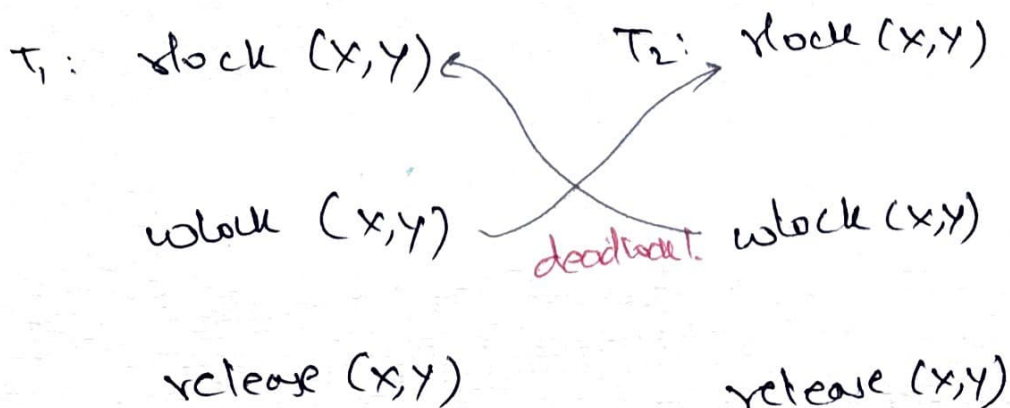
● Optimistic

- If high conflicts, keep aborting and restarting

Pessimistic

- unnecessary locking if no conflicts
- Deadlock avoidance is needed

● Would wait deadlock avoidance.

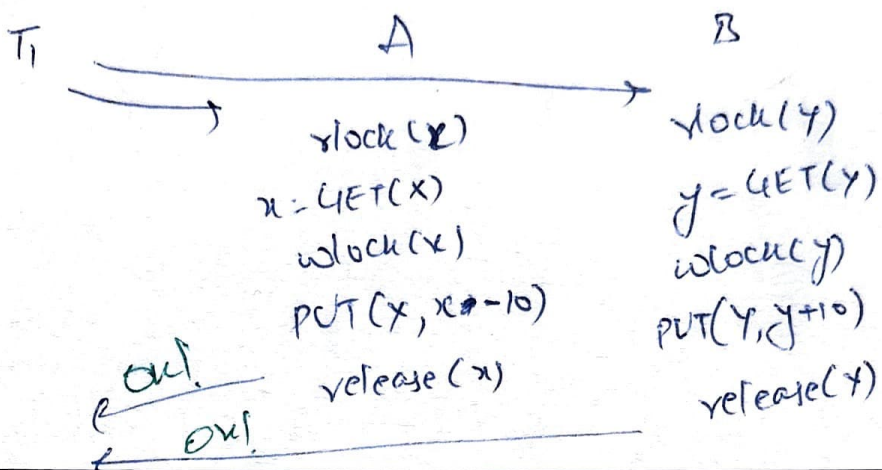


First to unlock, aborts other transaction.

Bad commit protocol

Transaction manager (TM)

Resource managers (RMi)



what can go wrong?

- Not enough money in x
- y account does not exist
- A or B crashes before receiving msg
- Network fails
- TM crashes after sending txn to A but before " " " B

• Safety?

Atomic commits: Everyone commits
or everyone aborts.

Keep aborting forever?

• Liveness

- If no failures, A, B can commit, then commit
- If failures, reach a conclusion ASAP

- R/W transactions can be thought of as 2 phases

T_i : $\text{rlock}(x, y)$
 $x = \text{get}(x)$
 $\text{assert}(x > 10)$
 $y = \text{GET}(y)$
 $\text{wlock}(x, y)$
 $\text{PUT}(x, x + 10)$
 $\text{PUT}(y, y + 10)$
 $\text{release}(x, y)$

Prepare phase: Read all values.
 Take all locks.
 No writes!

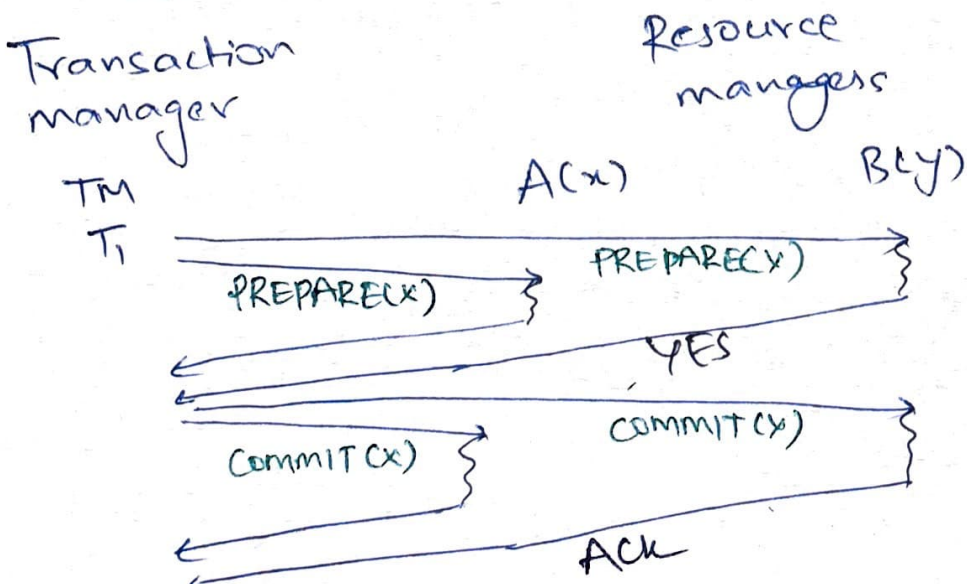
Commit phase: Write and release locks.
 write both x, y or none
 must be atomic.

PREPARE(x):
 $\text{rlock}(x)$
 $x = \text{GET}(x)$
 $\text{assert}(x > 10)$
 $\text{wlock}(x)$
 return x .

COMMIT(x)
 $\text{PUT}(x, x + 10)$
 $\text{release}(x)$

ABORT(x)
 $\text{release}(x)$

Atomicity Two phase Commits



● Why ~~is~~ does it give atomic commits?

- TM can send commit only if it has heard Yes from all RMs.

All or nothing

- Ex: If B cannot wlock(y), it replies NO
⇒ TM abort transaction.

● B crashes before sending YES to TM.

- TM timeouts and unilaterally aborts.
- or n/w lost YES message

● B crashes after sending YES to TM

- TM sends commit to A.
- B restarts
 - must remember it was in middle of Txn. $PREPARE_{T_i}(y)$ i.e., wlock(y)
 - TM keeps retrying commit(T_i)
 - Is B guaranteed to get wlock(y)?

WAL $Prepare_{T_i}(y) \rightarrow \text{Yes}$

What if TM restarts before sending prepare to B?

- Send prepare again.
- B prepares
- A should remember it was already prepared and reply YES

What if TM restarts after prepares?

- If participant had replied Yes, it is blocked. waiting for commit/abort
- After restart, TM must commit/abort all pending transactions

TM log

- $\langle \text{Txn ID} \rangle \langle \text{details} \rangle$
- $\langle \text{commit} \rangle \langle \text{Txn ID} \rangle$
or
 $\langle \text{abort} \rangle \langle \text{Txn ID} \rangle$

why is it ok to not log -

- sent prepare to A?
 - can just resend prepares
- received yes from A?
 - can unilaterally abort
- sent commit to A?
 - can just resend commit.

RM/
Participant log

- <Txn ID> <details> <prepared>
- " <committed>
or <aborted>

Safety -

- No commit unless everyone says yes
- ^{RM} Cannot back out after saying yes
across restarts

Liveness

- Not live if TM crashes forever after
prepare (or becomes unreachable)