
High Throughput Replication with Integrated Membership Management

Pedro Fouto, Nuno Preguiça, João Leitão

Presented By: Aman, Niranjan, Hrishabh

Adapted from authors USENIX ATC 2022 Slides

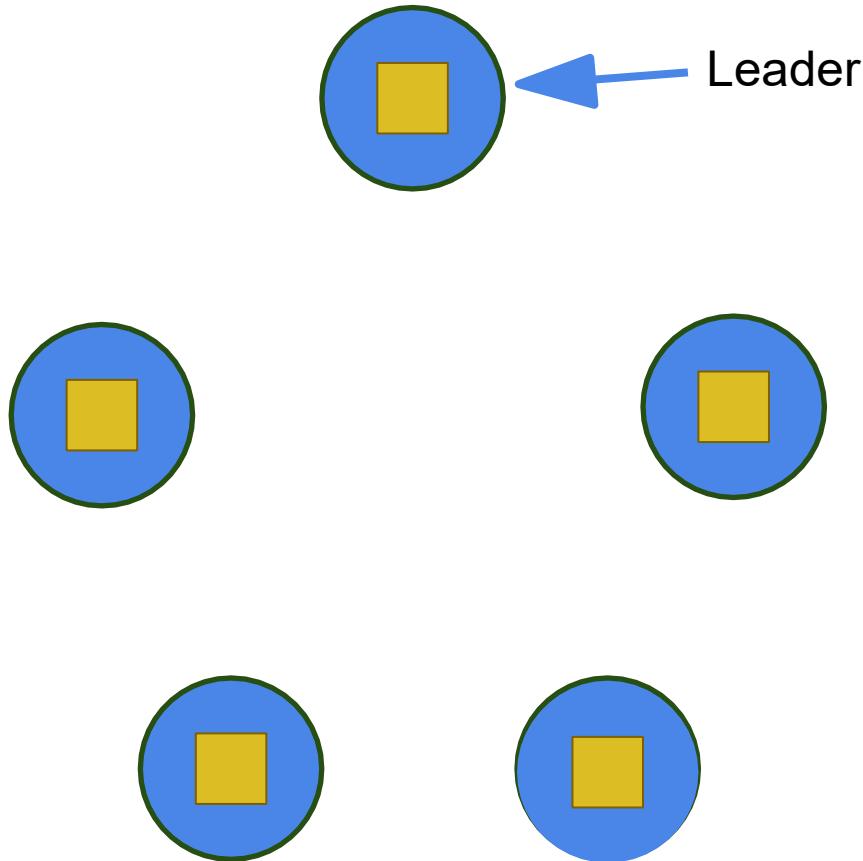
Motivation: Consensus and SMR

- Building blocks of numerous practical replication systems
- Their performance is critical
- Many alternatives have been designed
- Two very relevant ones:
 - (Multi-)Paxos
 - Chain Replication

Motivation: Consensus and SMR

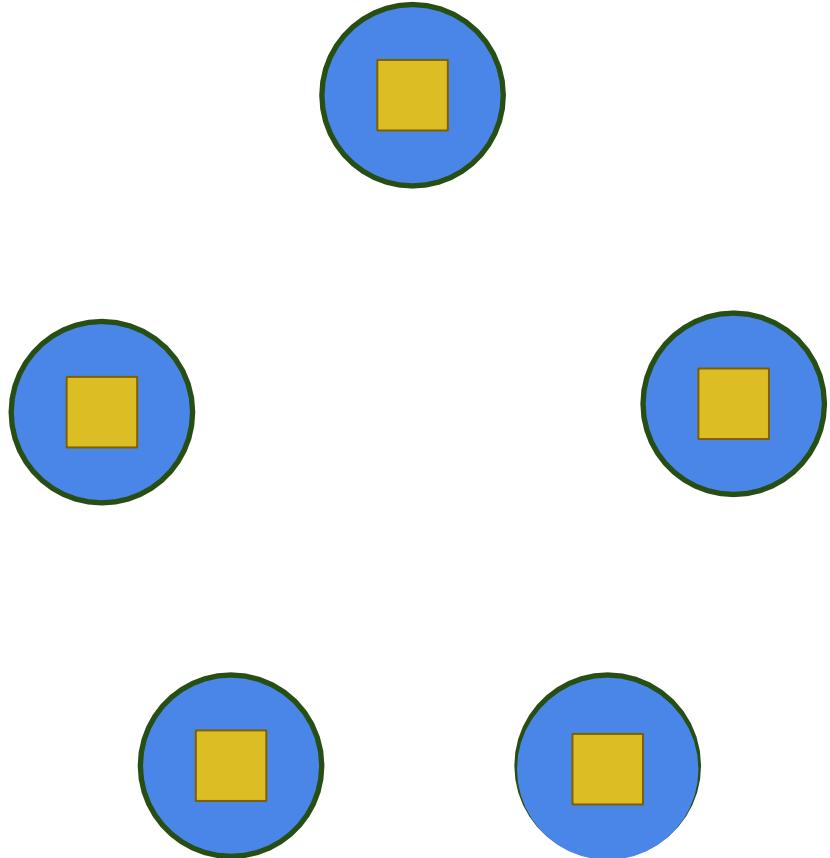
- Building blocks of numerous practical replication systems
- Their performance is critical
- Many alternatives have been designed
- Two very relevant ones:
 - **(Multi-)Paxos**
 - Chain Replication

Multi-Paxos

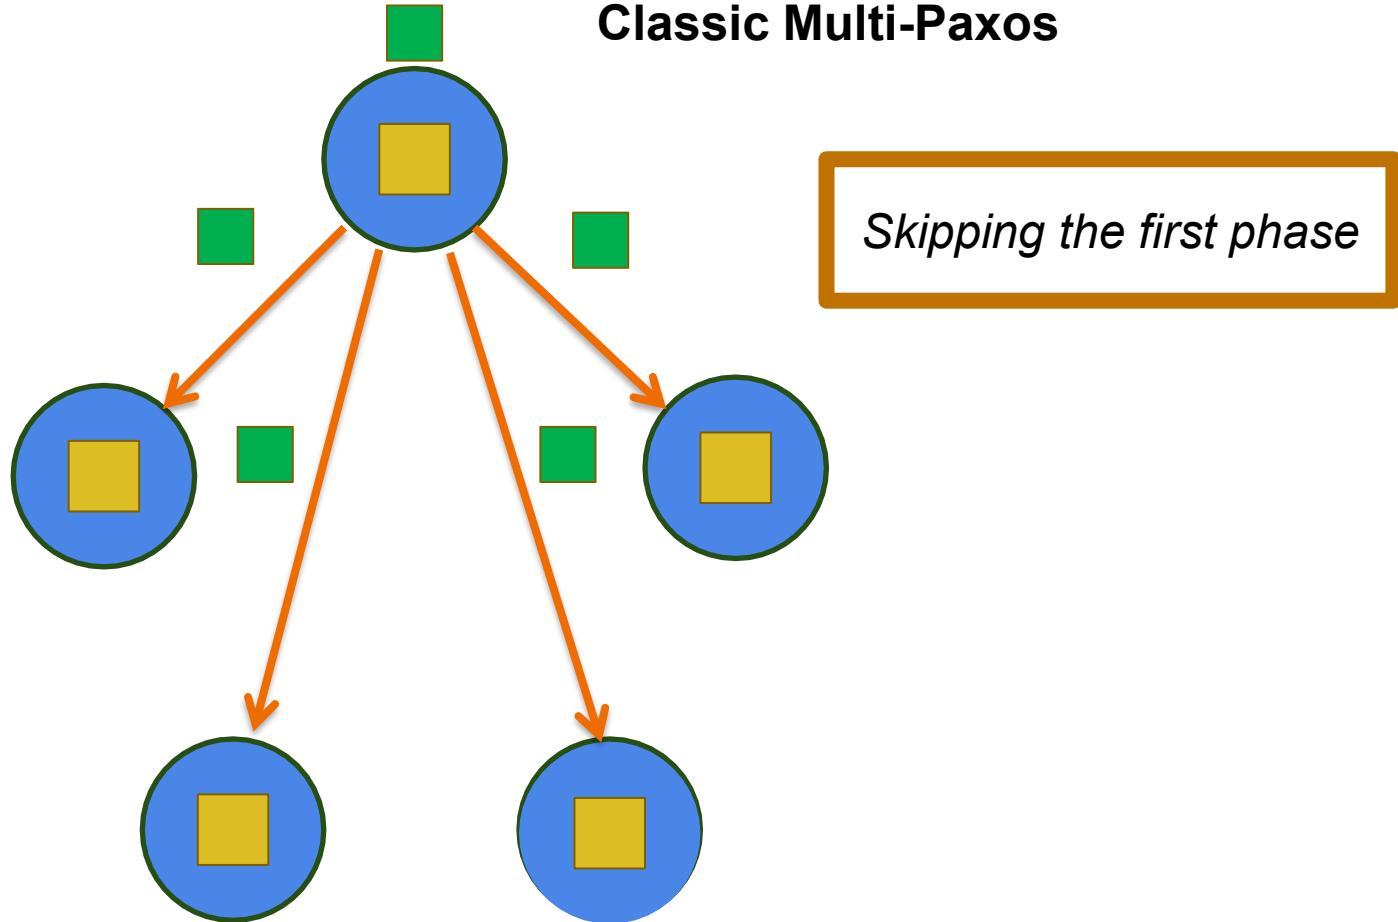


Multi-Paxos

Classic Multi-Paxos

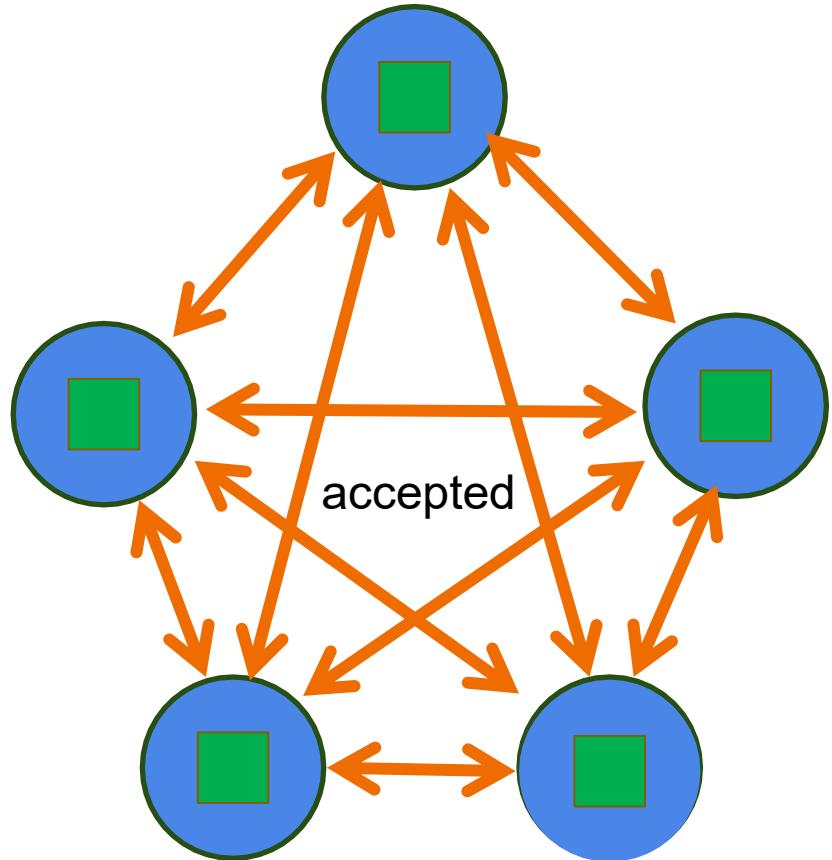


Multi-Paxos

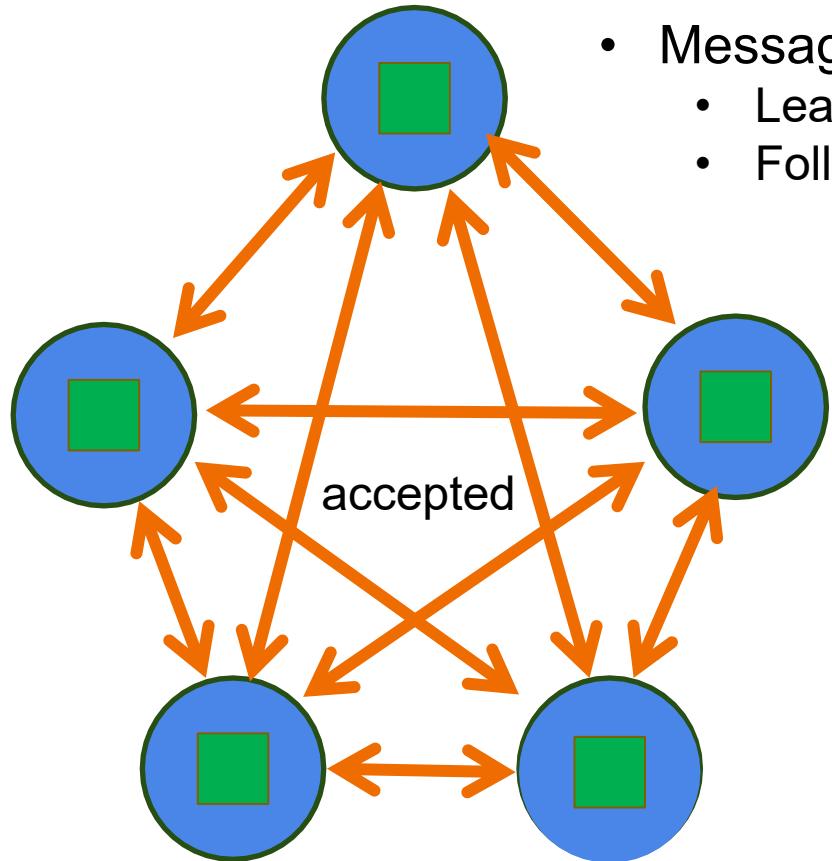


Multi-Paxos

Classic Multi-Paxos



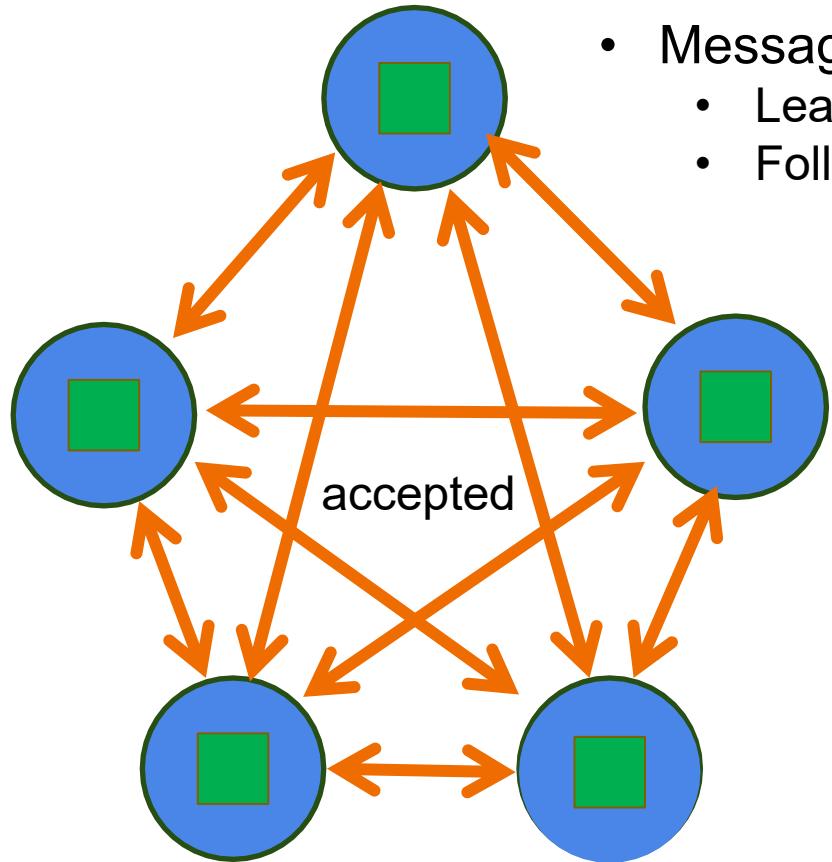
Multi-Paxos



Classic Multi-Paxos

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$

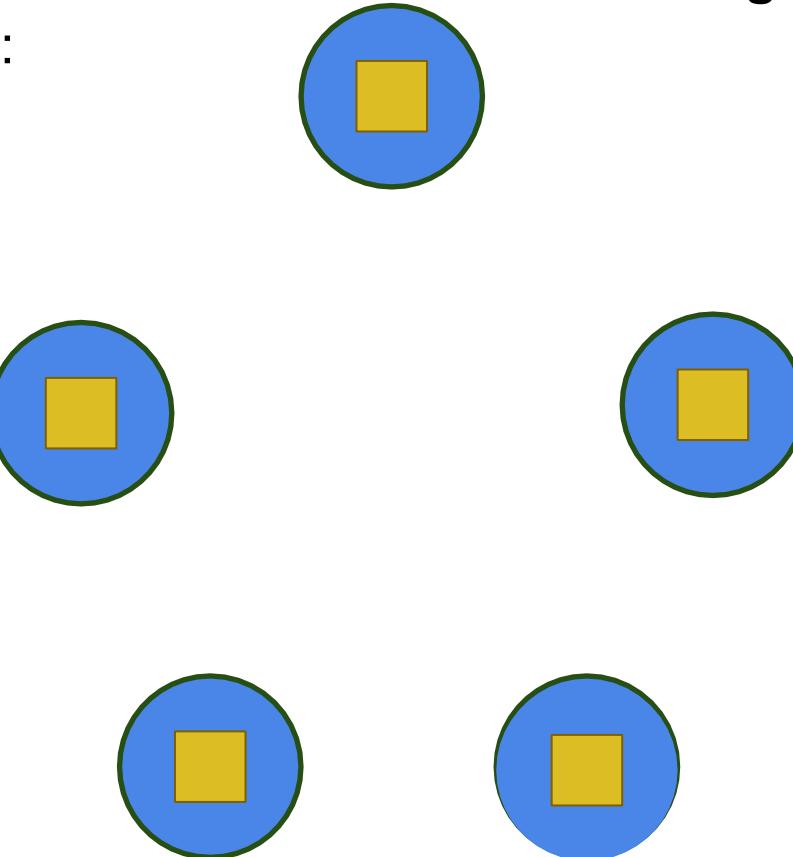
Multi-Paxos



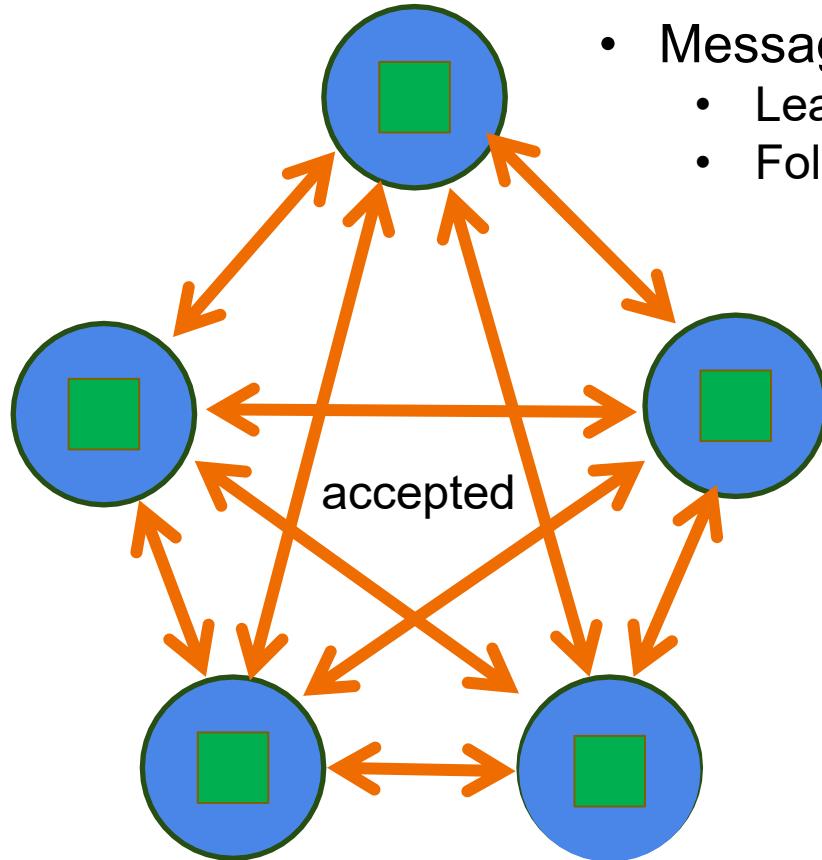
Classic Multi-Paxos

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$

Distinguished Learner

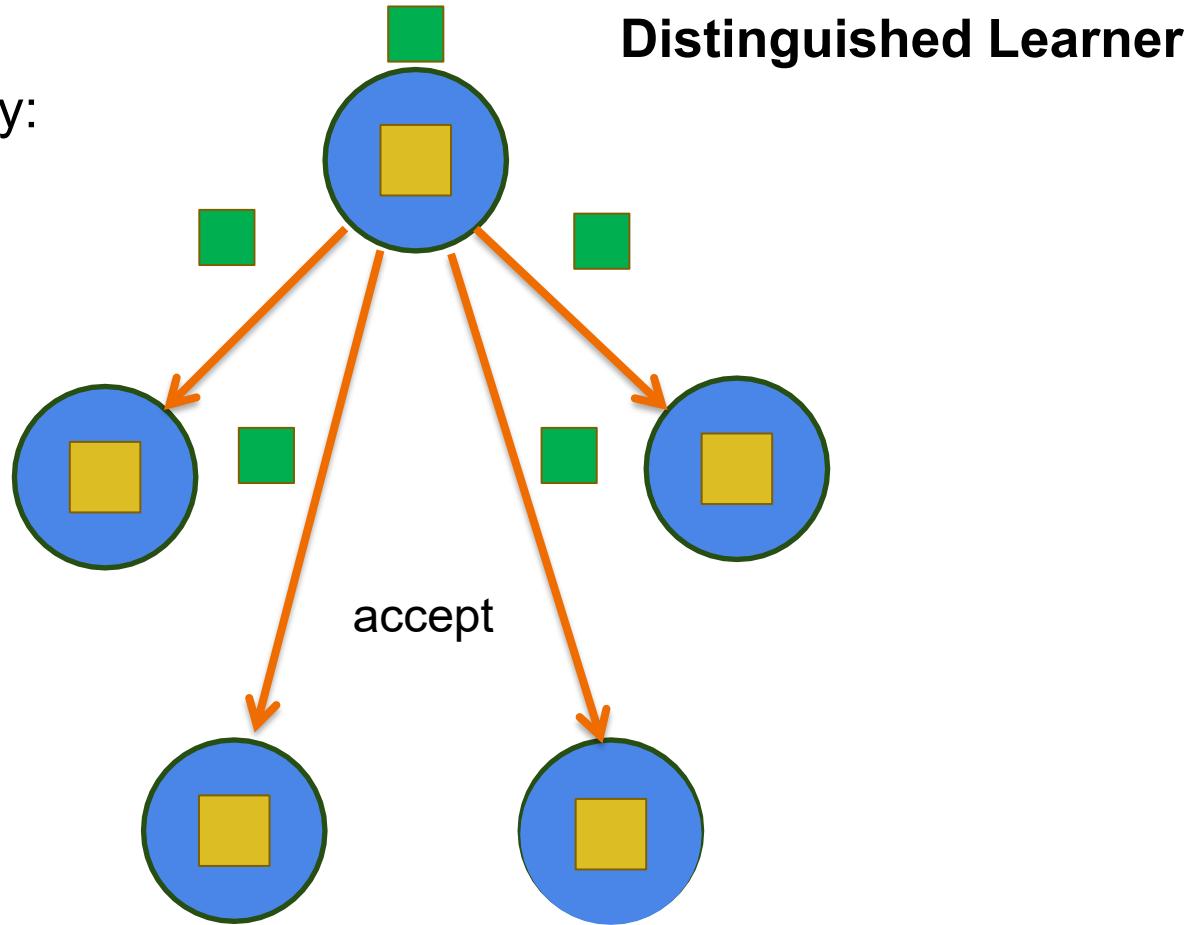


Multi-Paxos



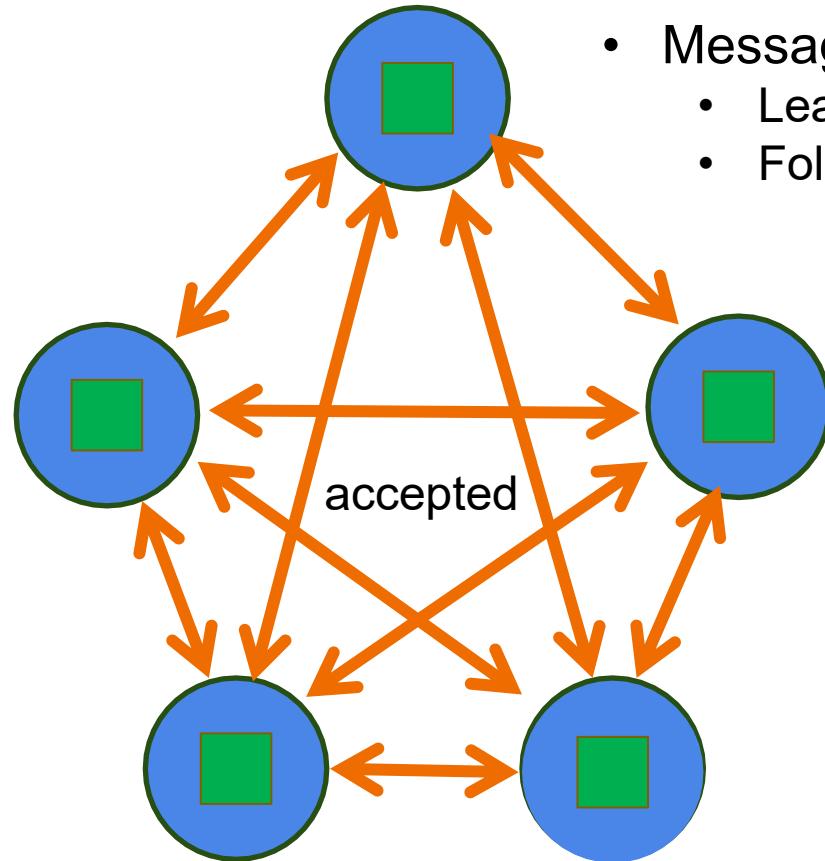
Classic Multi-Paxos

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$



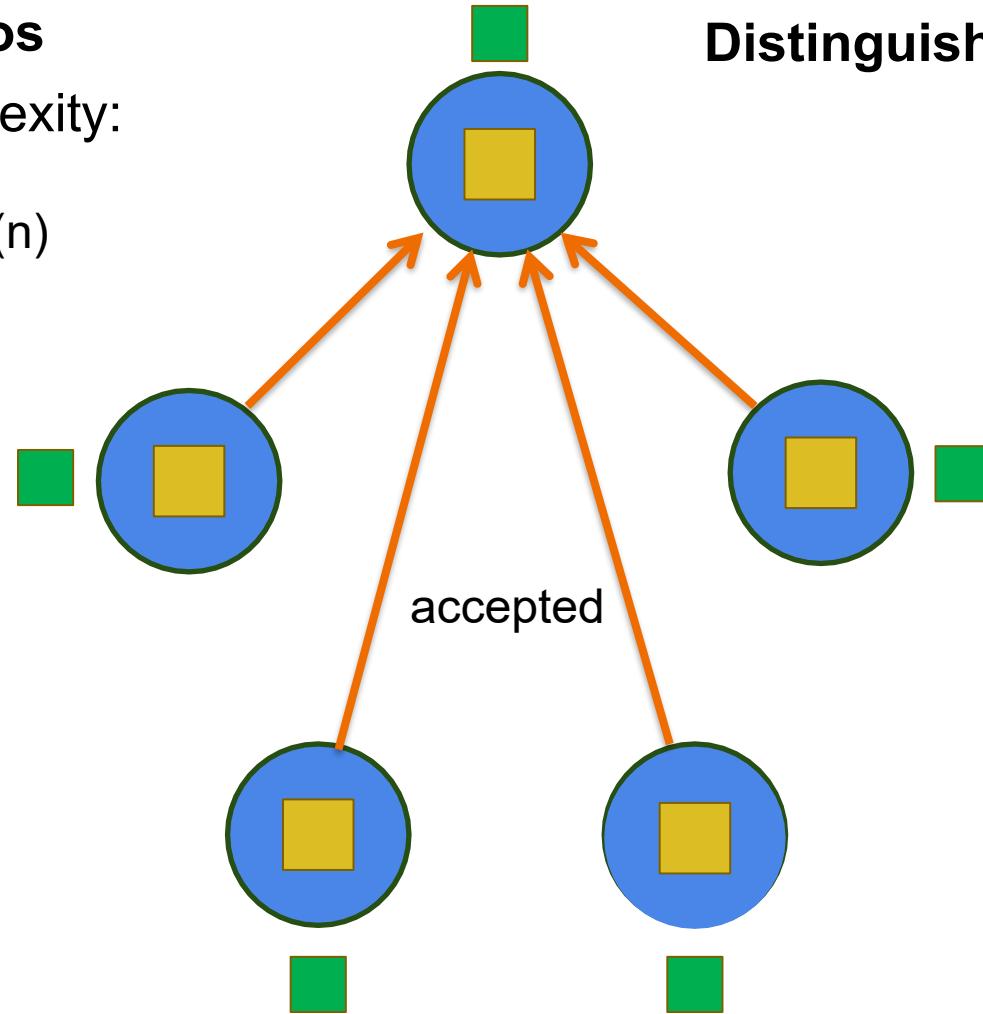
Distinguished Learner

Multi-Paxos



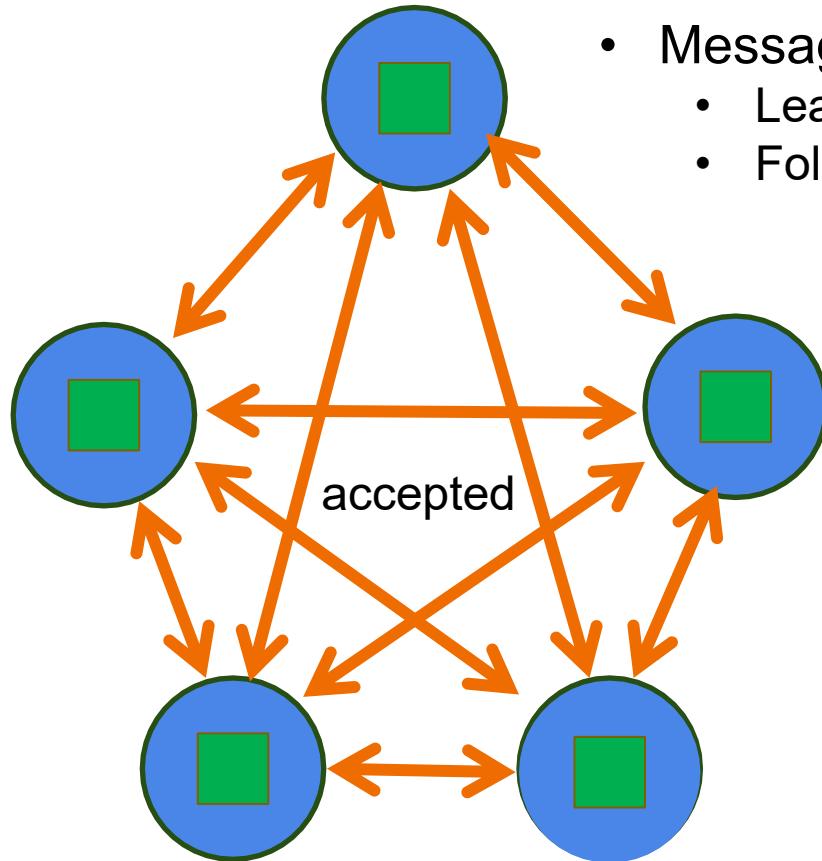
Classic Multi-Paxos

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$



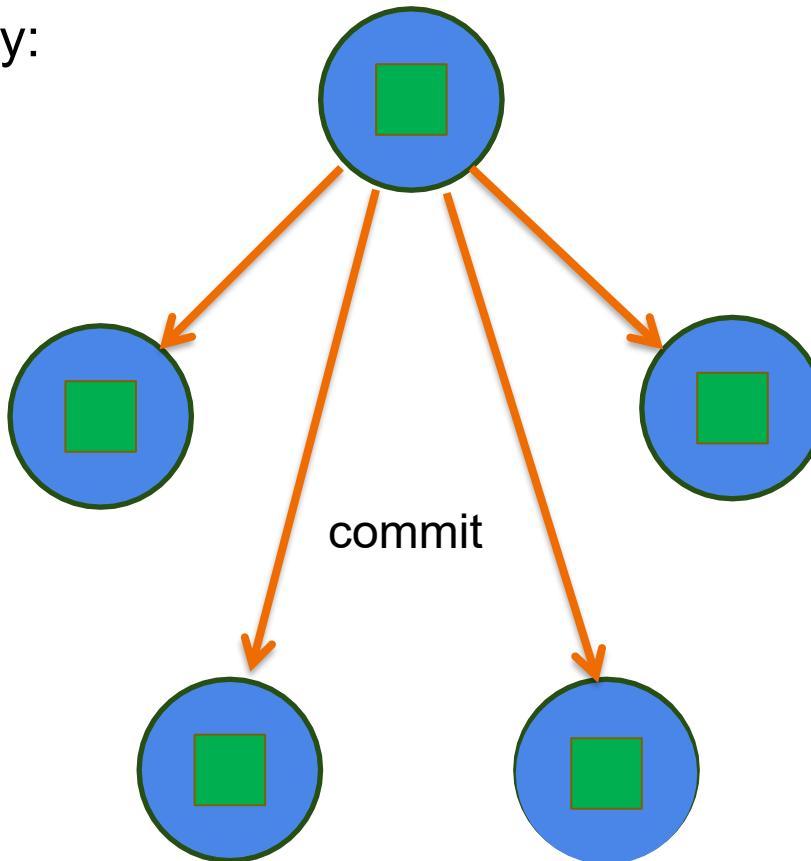
Distinguished Learner

Multi-Paxos



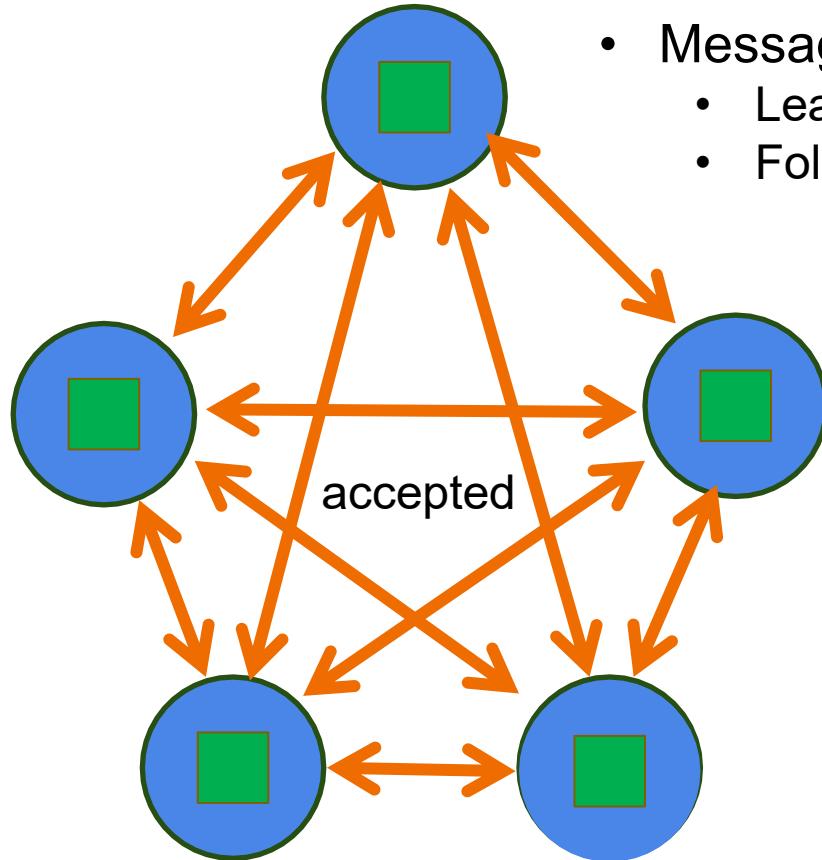
Classic Multi-Paxos

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$



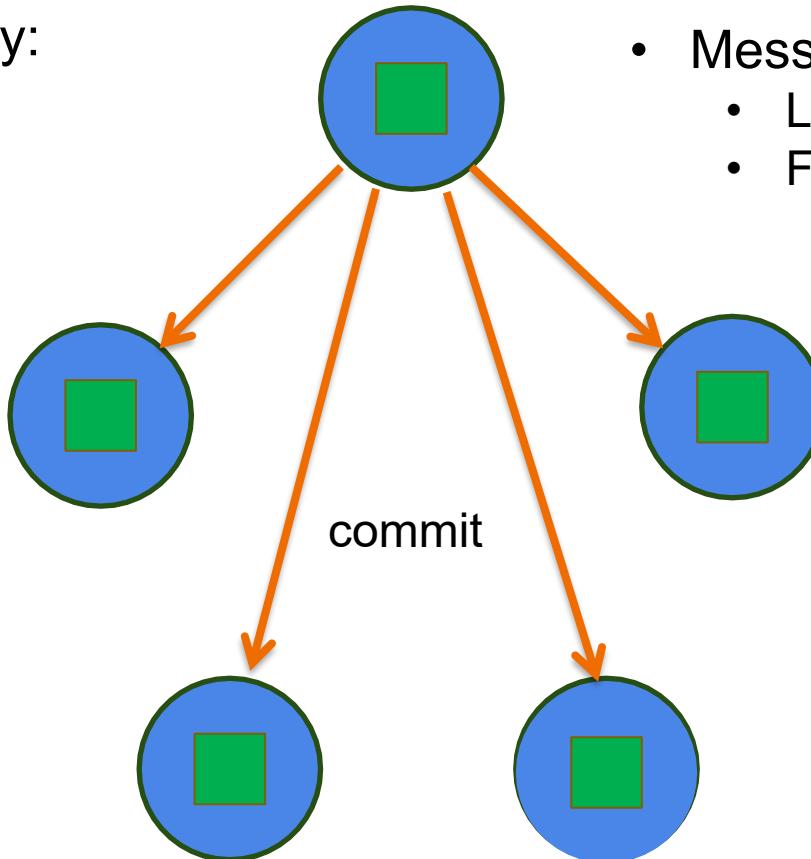
Distinguished Learner

Multi-Paxos



Classic Multi-Paxos

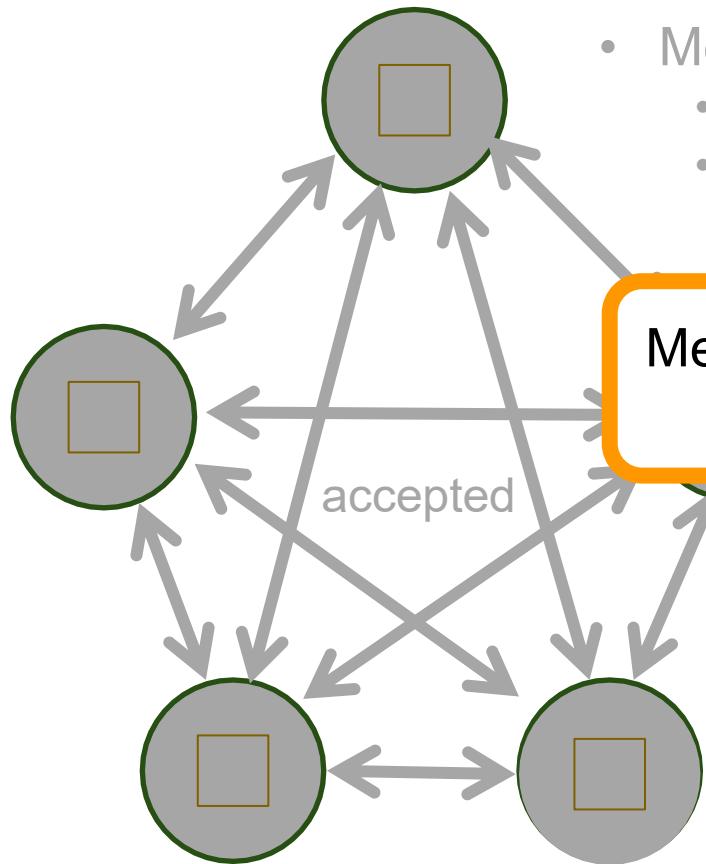
- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$



Distinguished Learner

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(1)$

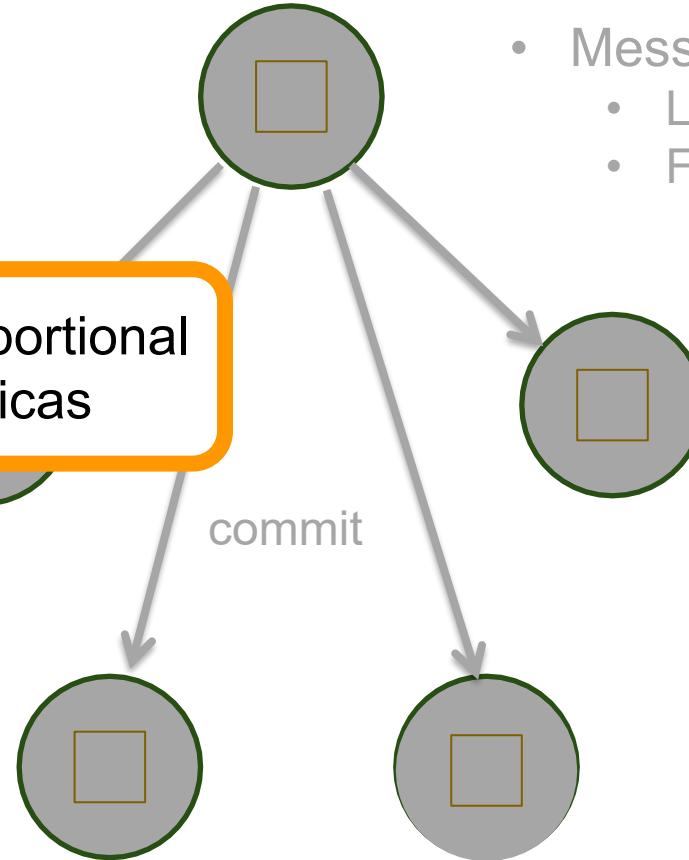
Multi-Paxos



Classic Multi-Paxos

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(n)$

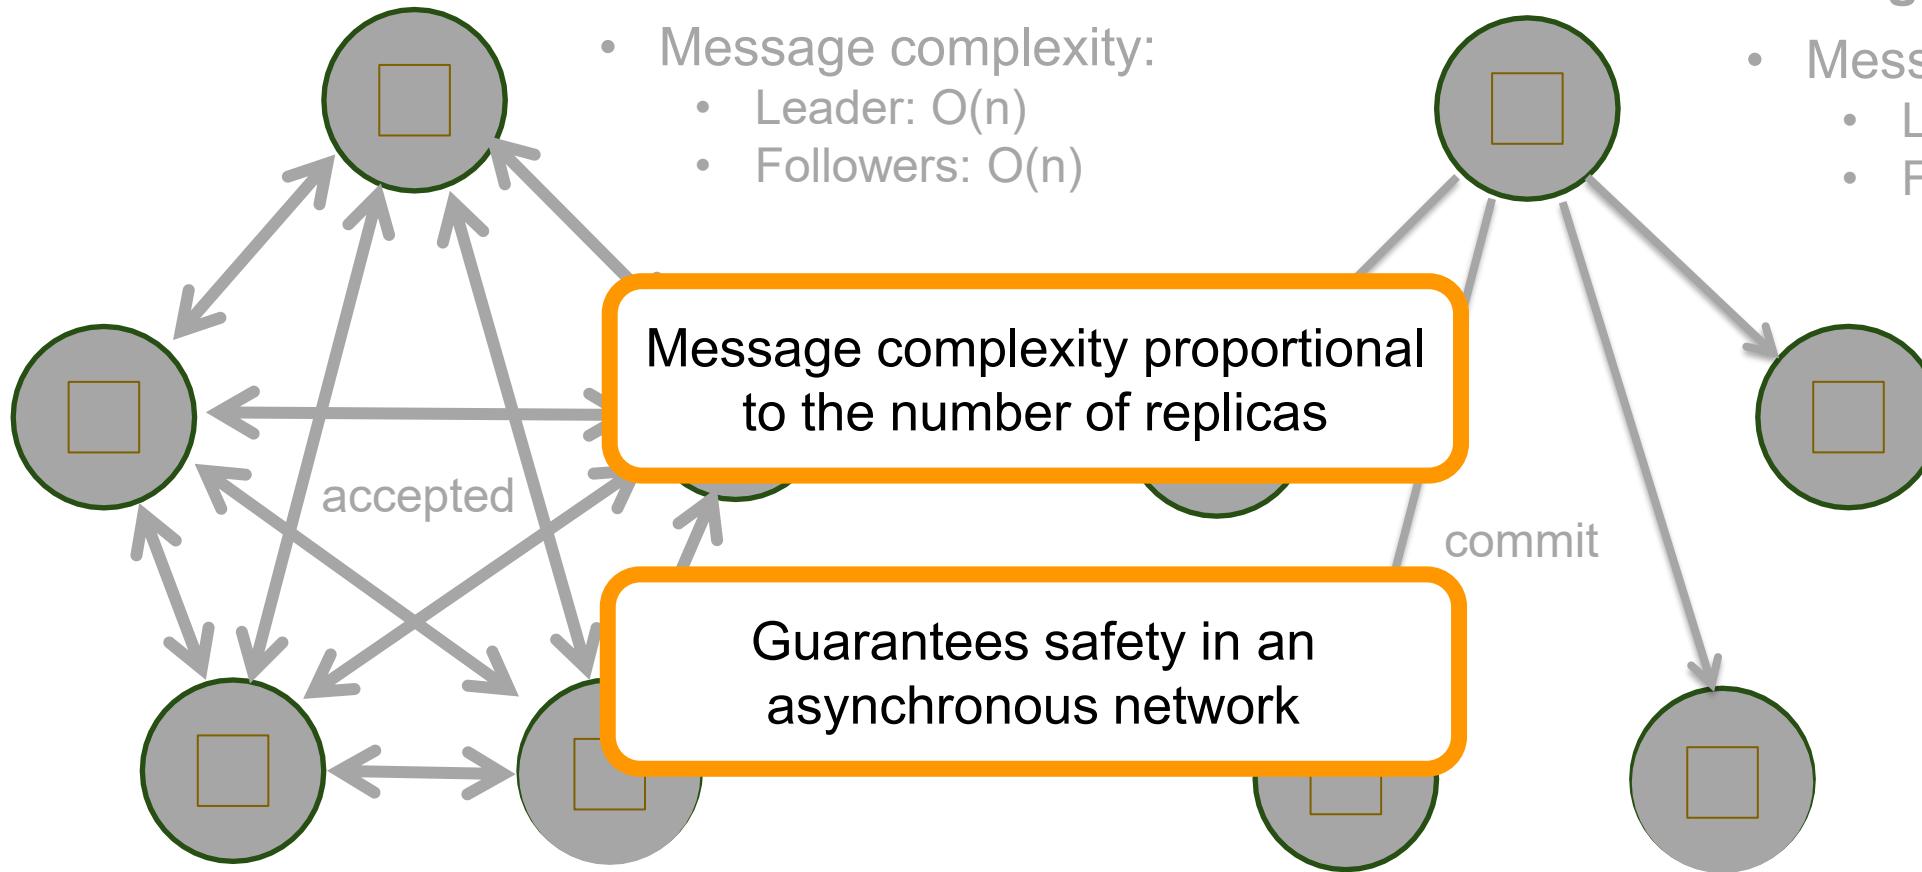
Message complexity proportional
to the number of replicas



Distinguished Learner

- Message complexity:
 - Leader: $O(n)$
 - Followers: $O(1)$

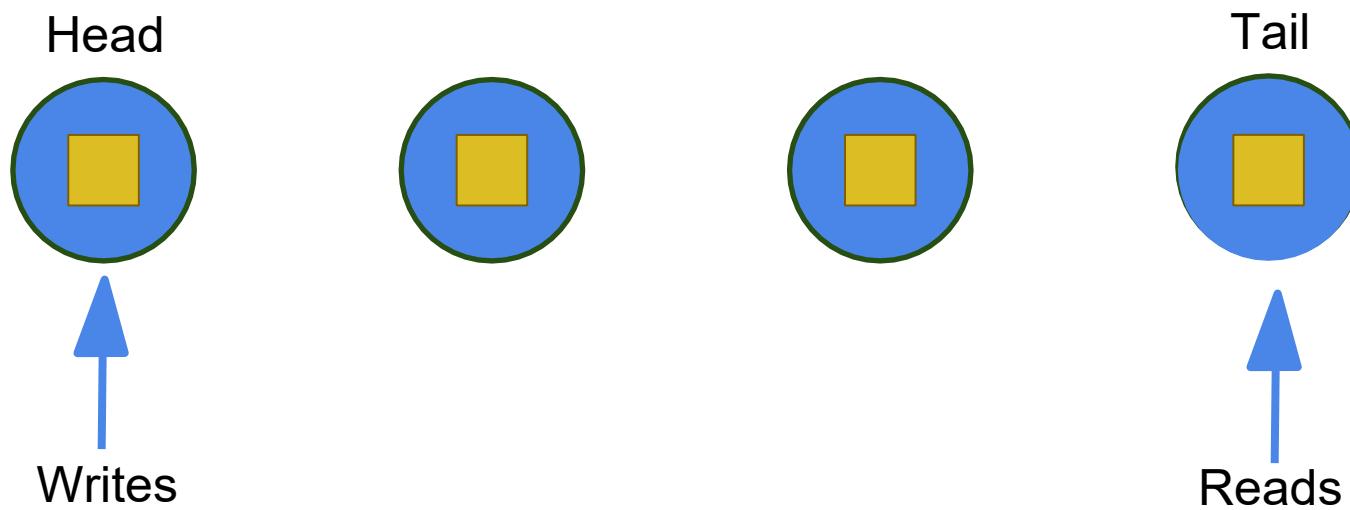
Multi-Paxos



Motivation: Consensus and SMR

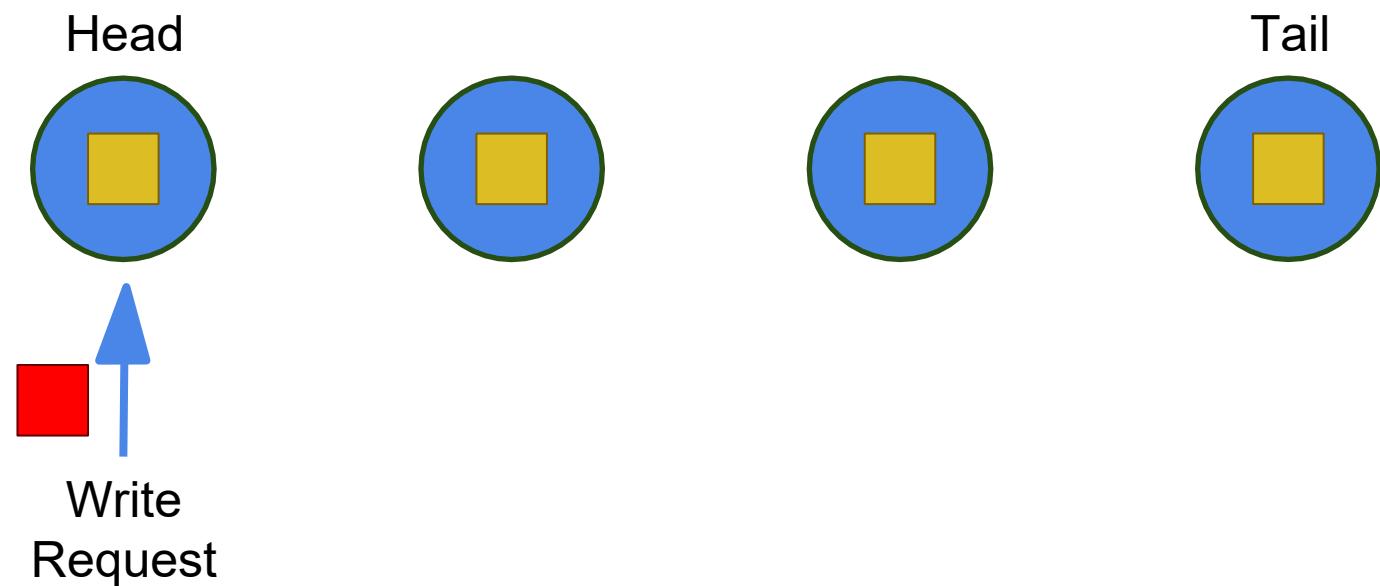
- Building blocks of numerous practical replication systems
- Their performance is critical
- Many alternatives have been designed
- Two very relevant ones:
 - (Multi-)Paxos
 - **Chain Replication**

Chain Replication (CR)

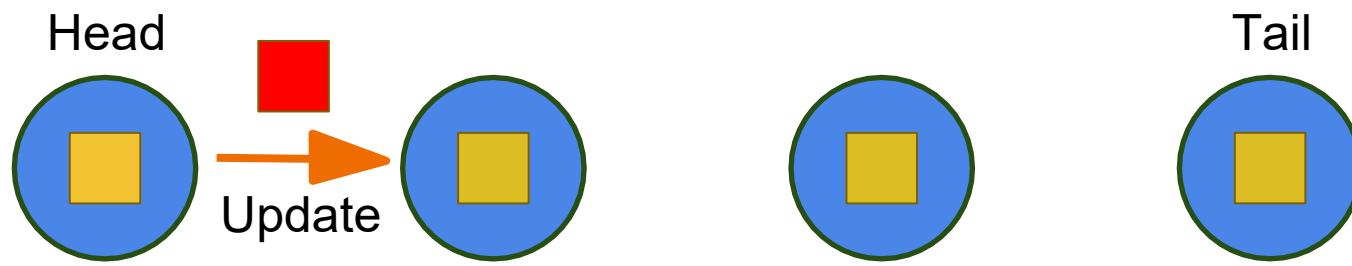


Any object.
Eg: Key-Value Store: $X = 1$

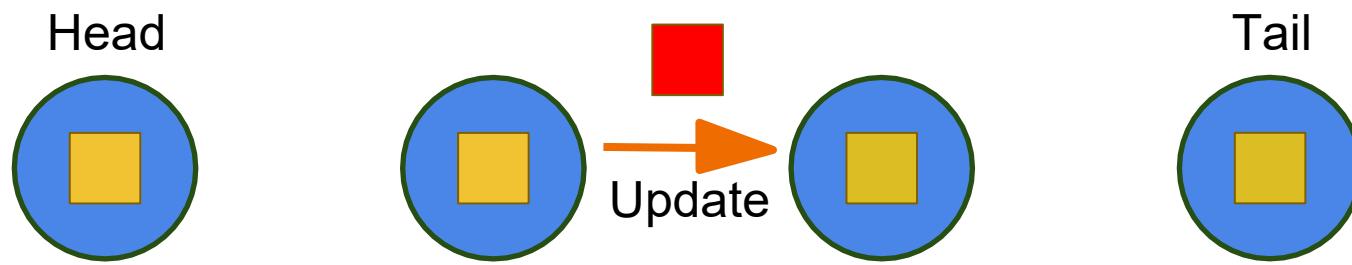
Chain Replication (CR)



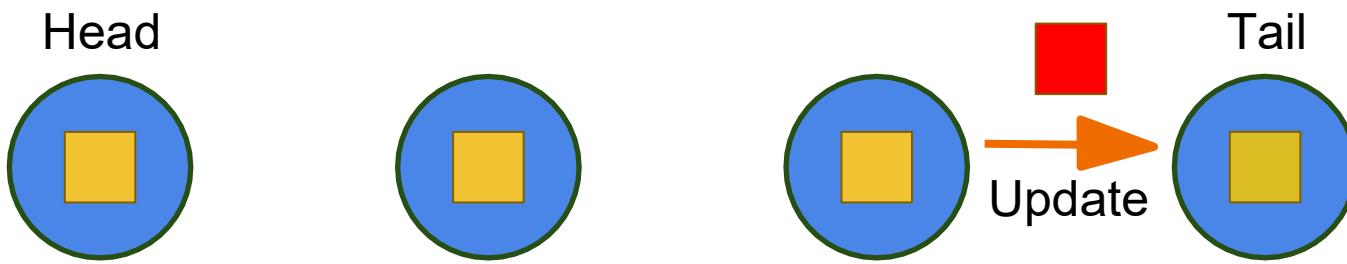
Chain Replication (CR)



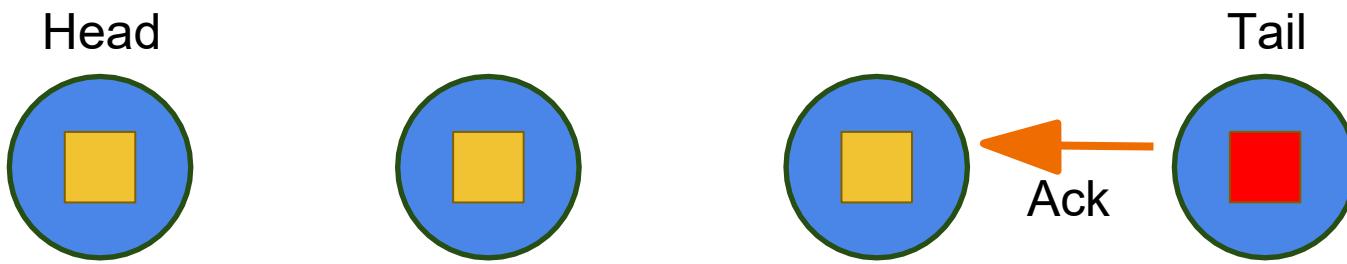
Chain Replication (CR)



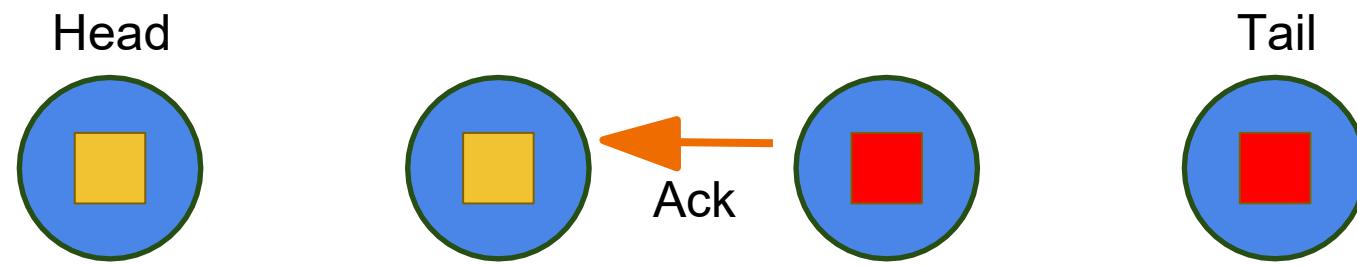
Chain Replication (CR)



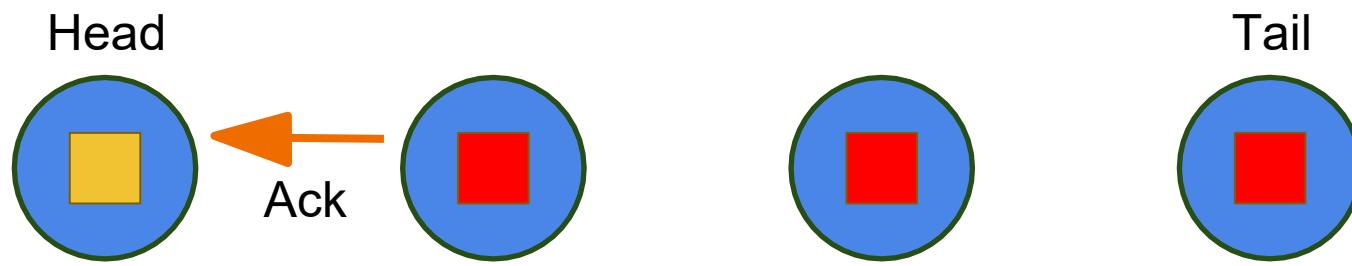
Chain Replication (CR)



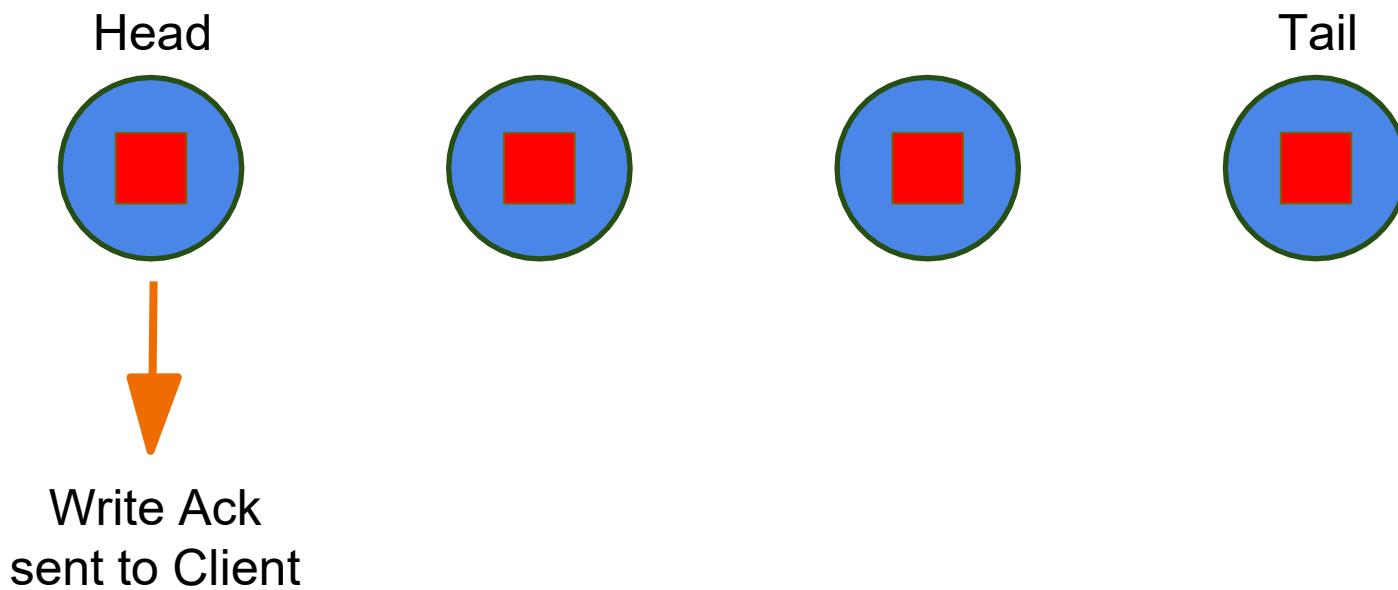
Chain Replication (CR)



Chain Replication (CR)

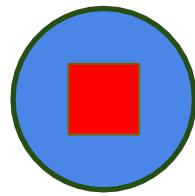
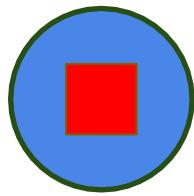
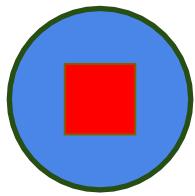


Chain Replication (CR)

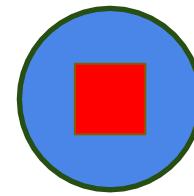


Chain Replication (CR)

Head



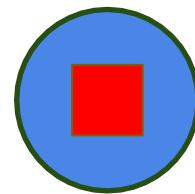
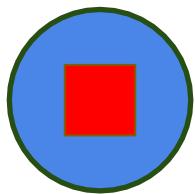
Tail



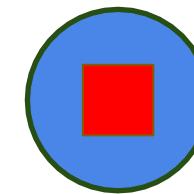
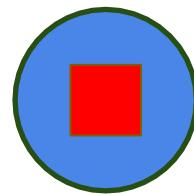
- Message complexity:
 - All replicas: $O(1)$

Chain Replication (CR)

Head



Tail



- Message complexity:
 - All replicas: $O(1)$

However, it has its limitations...

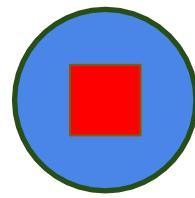
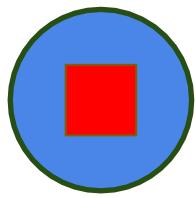
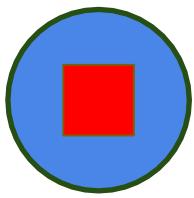
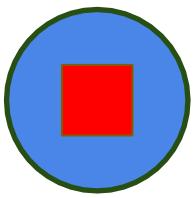
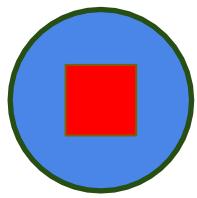
Motivation: Consensus and SMR

- Building blocks of numerous practical replication systems
- Their performance is critical
- Many alternatives have been designed
- Two very relevant ones:
 - (Multi-)Paxos
 - Chain Replication
- Two aspects are often overlooked:
 - **Performant linearizable reads**
 - Membership management and reconfiguration

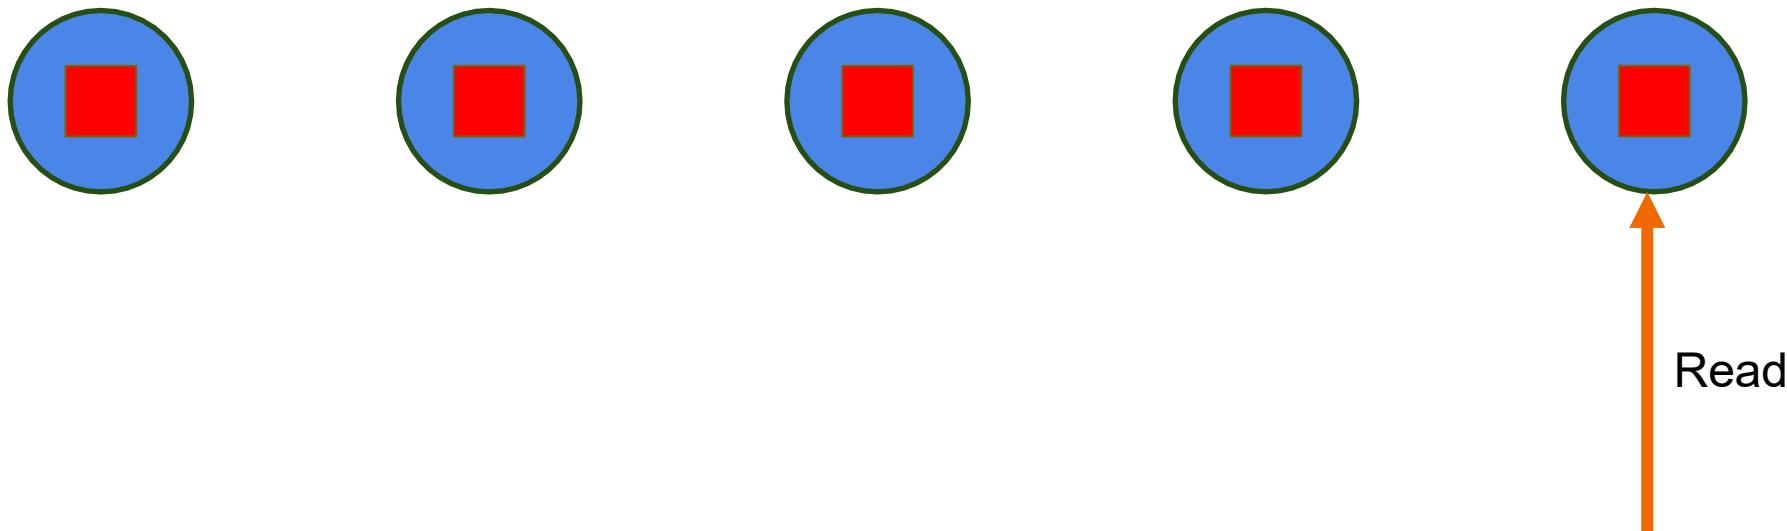
Motivation: Linearizable Reads

- Some existing solutions assume a synchronous model (e.g. Chain Replication)
- Dealing with asynchrony is complicated. One can:
 - Relax consistency (e.g. ZooKeeper)
 - Add extra (costly) steps to write operations
 - Synchronize with other replicas when reading

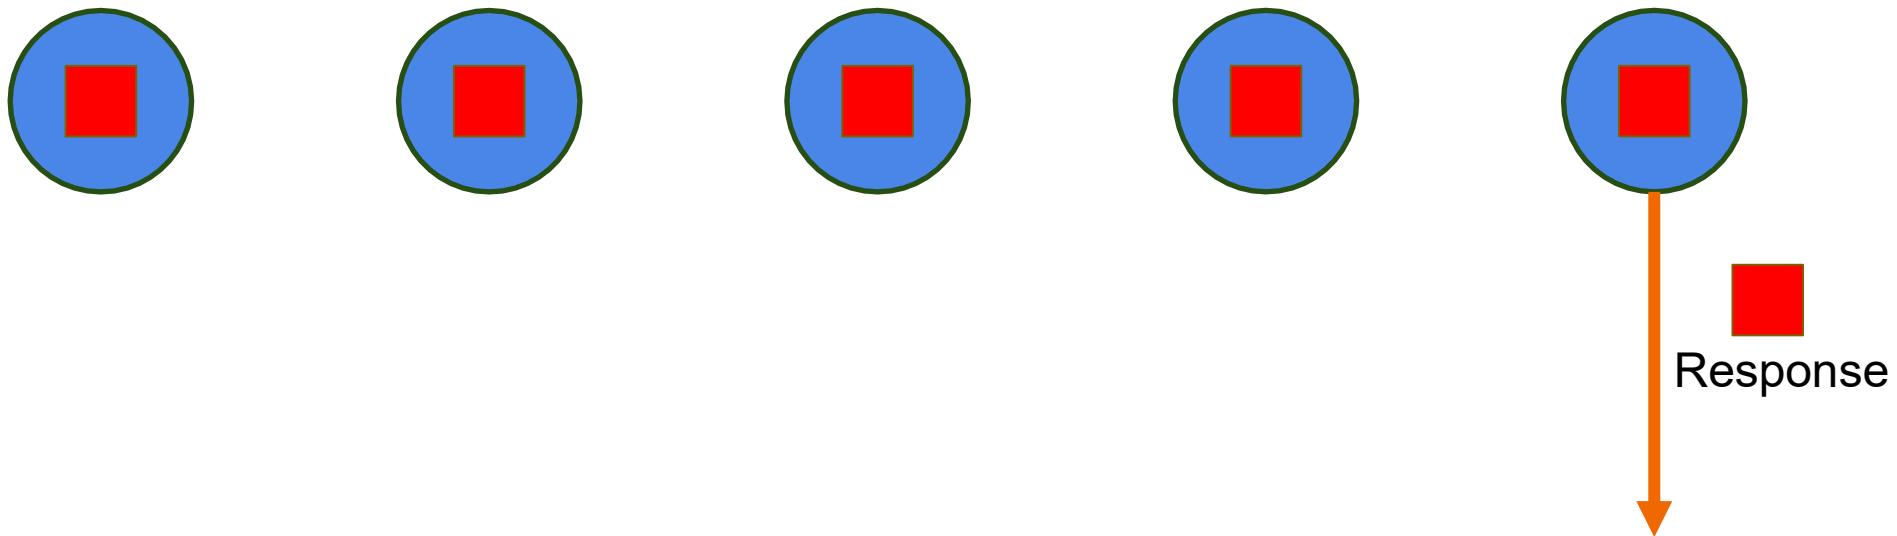
Motivation: Linearizable Reads



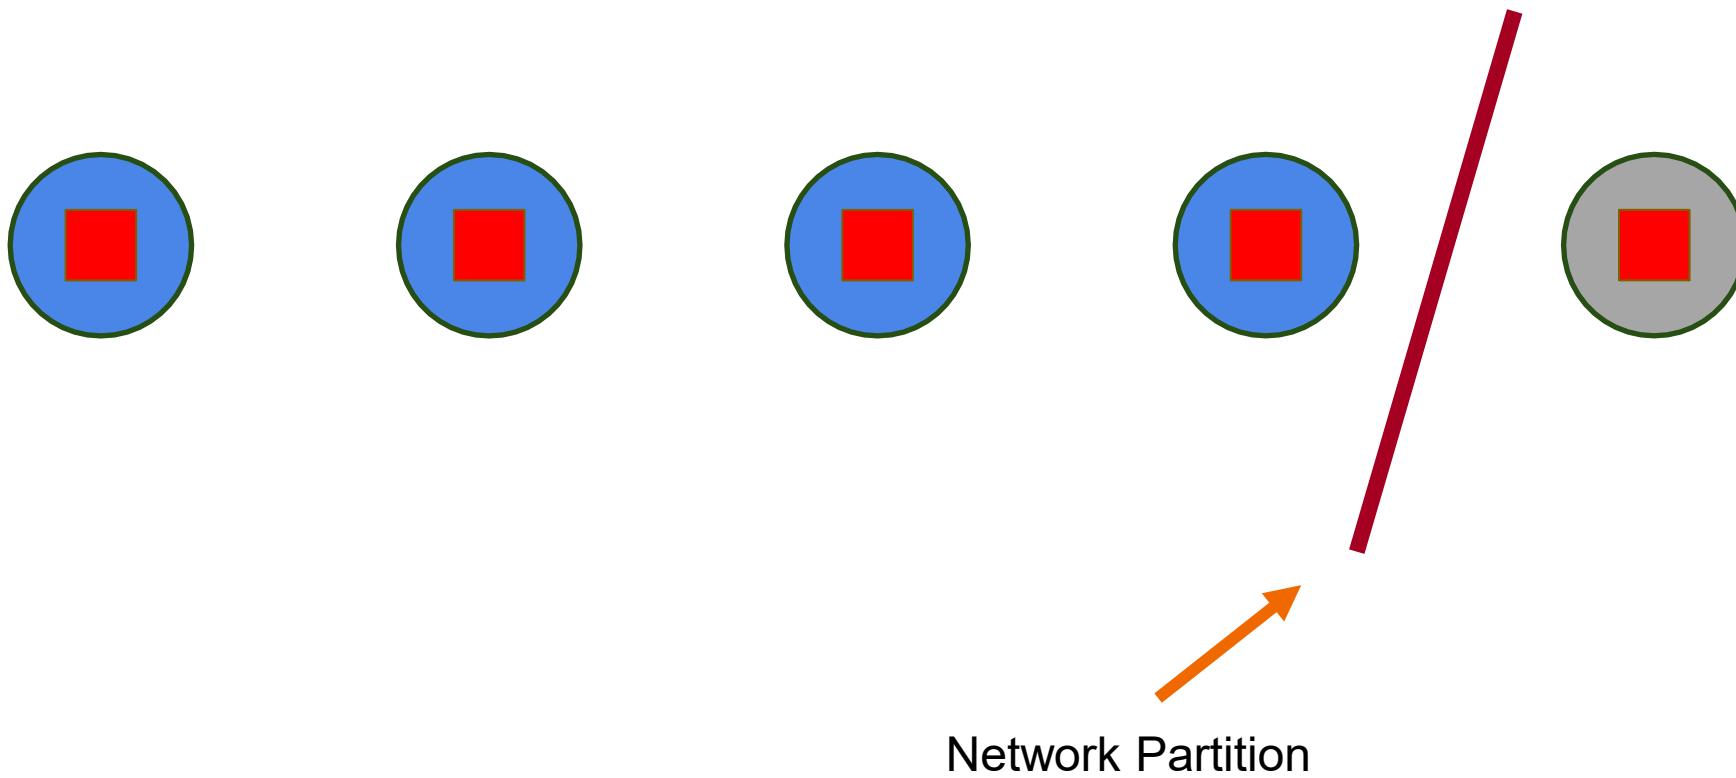
Motivation: Linearizable Reads



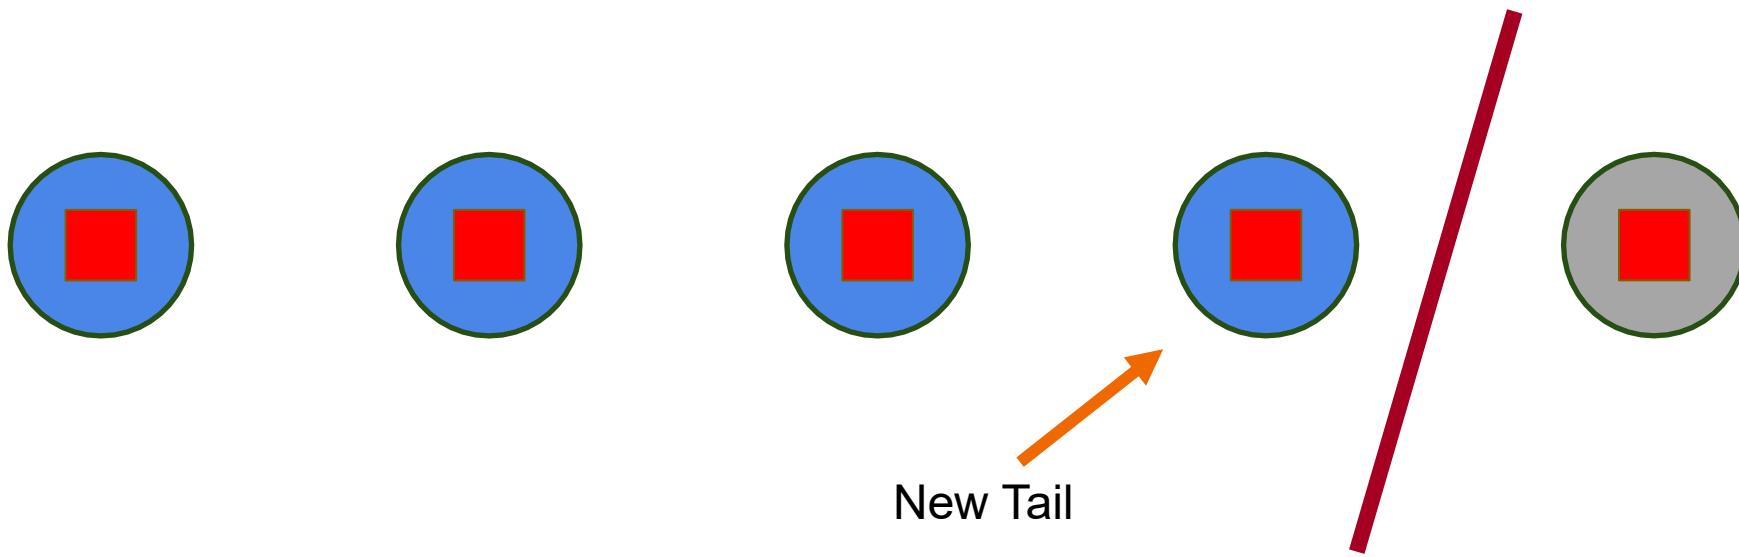
Motivation: Linearizable Reads



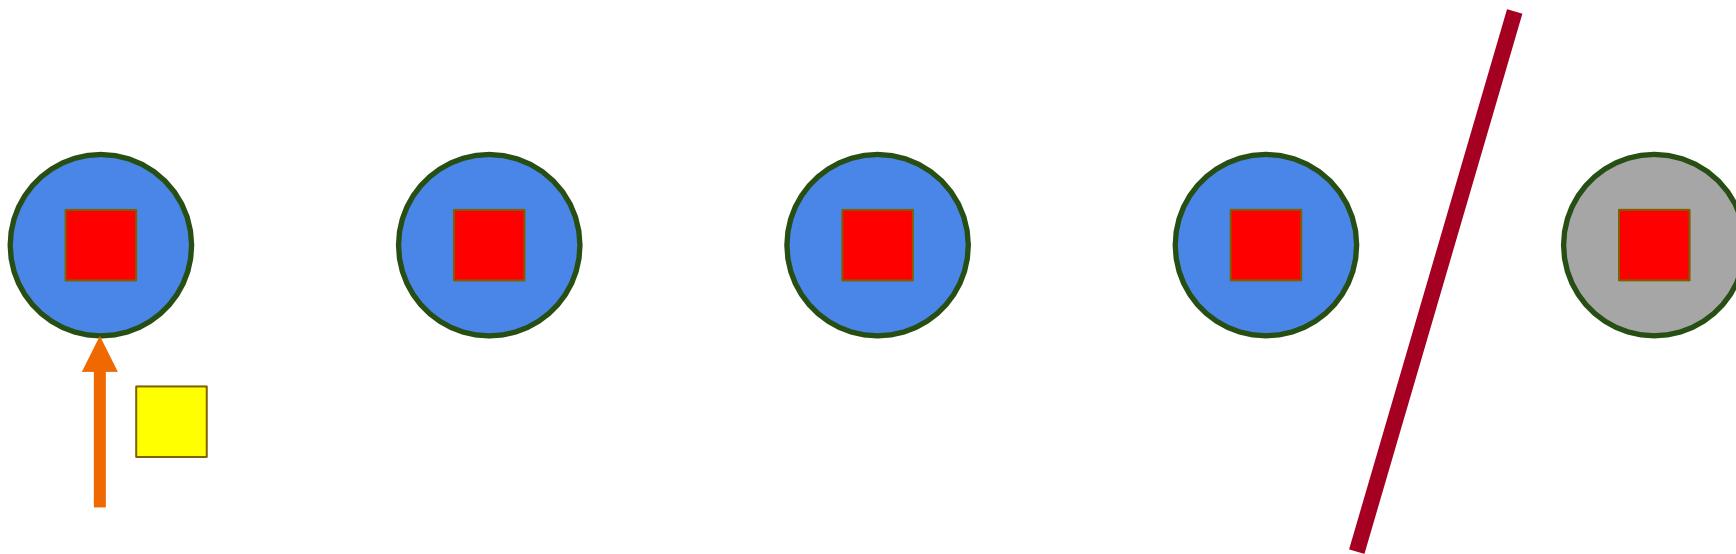
Motivation: Linearizable Reads



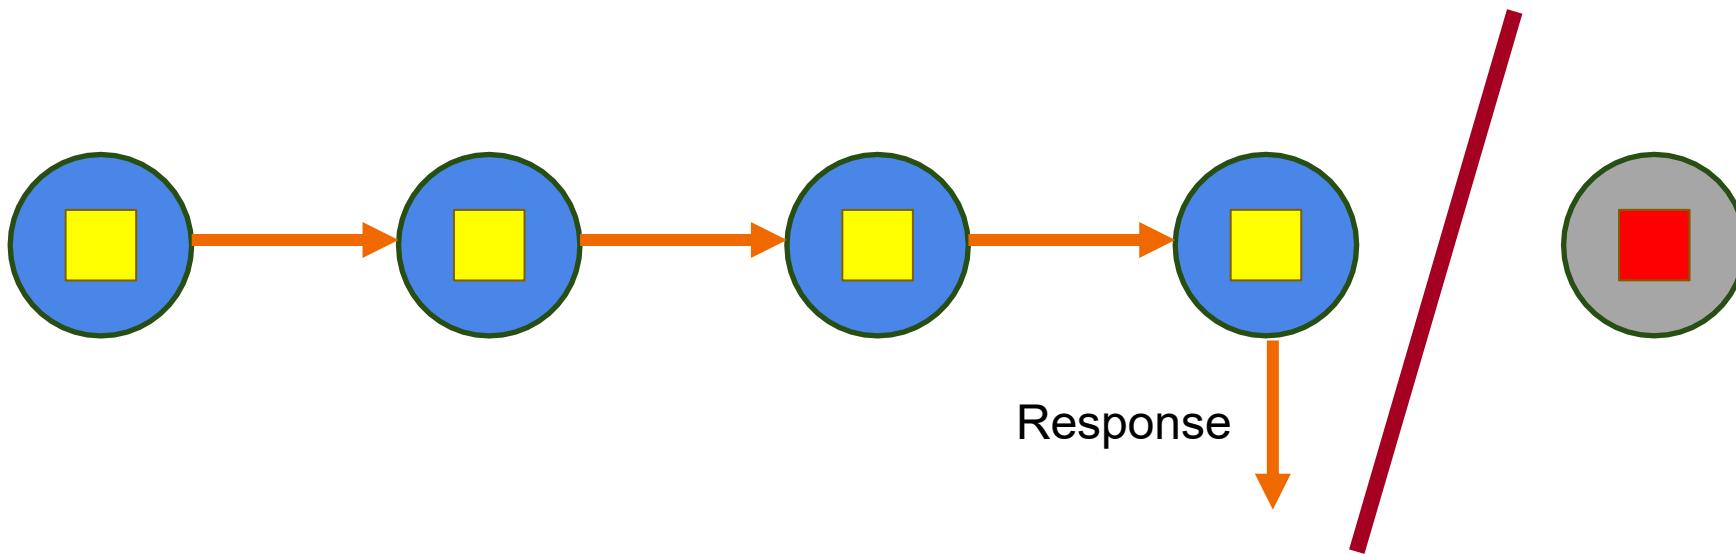
Motivation: Linearizable Reads



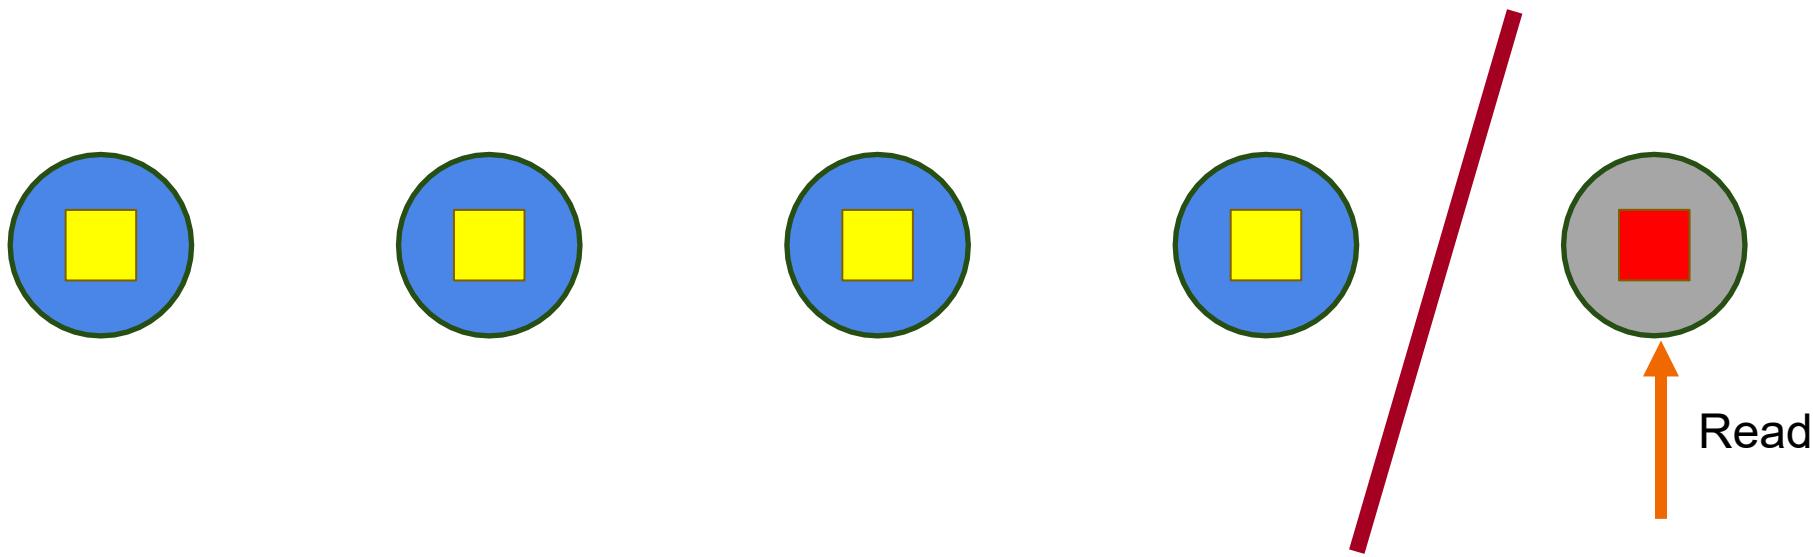
Motivation: Linearizable Reads



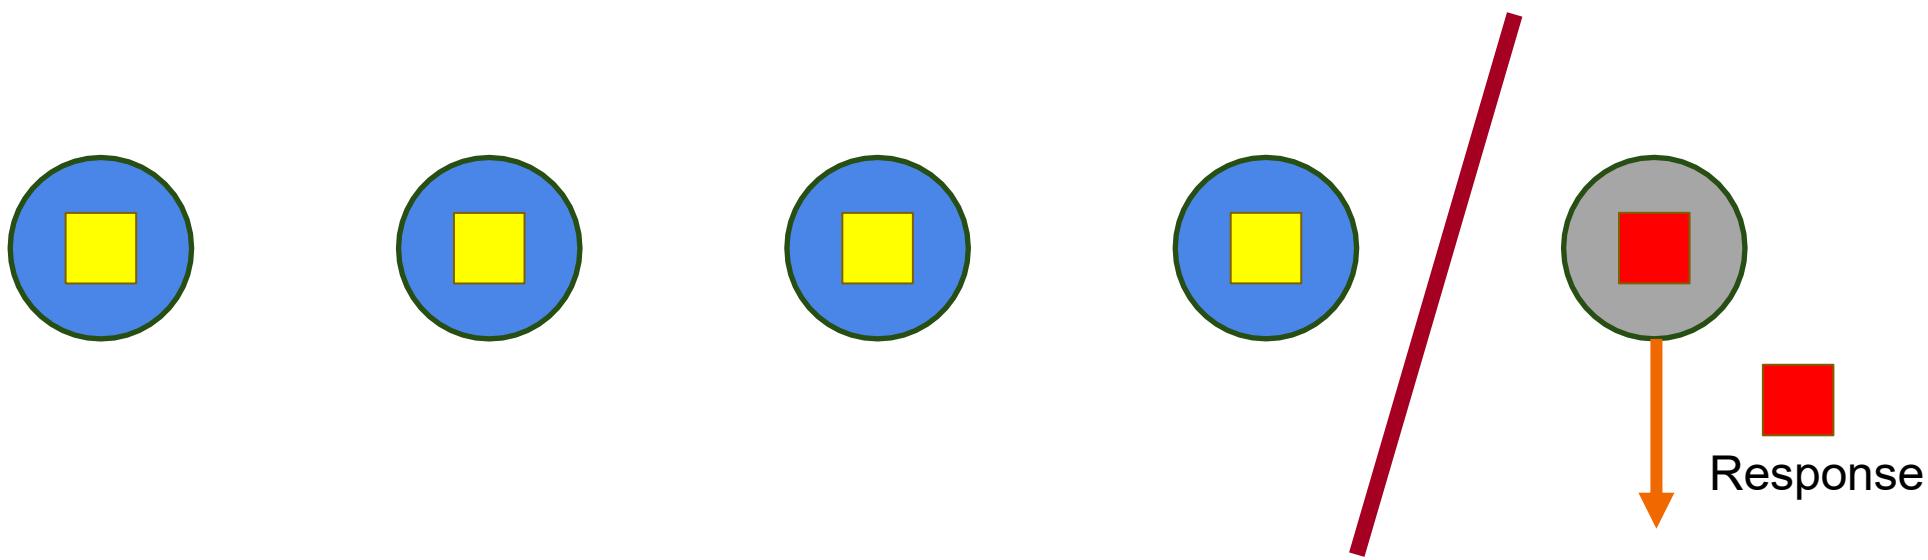
Motivation: Linearizable Reads



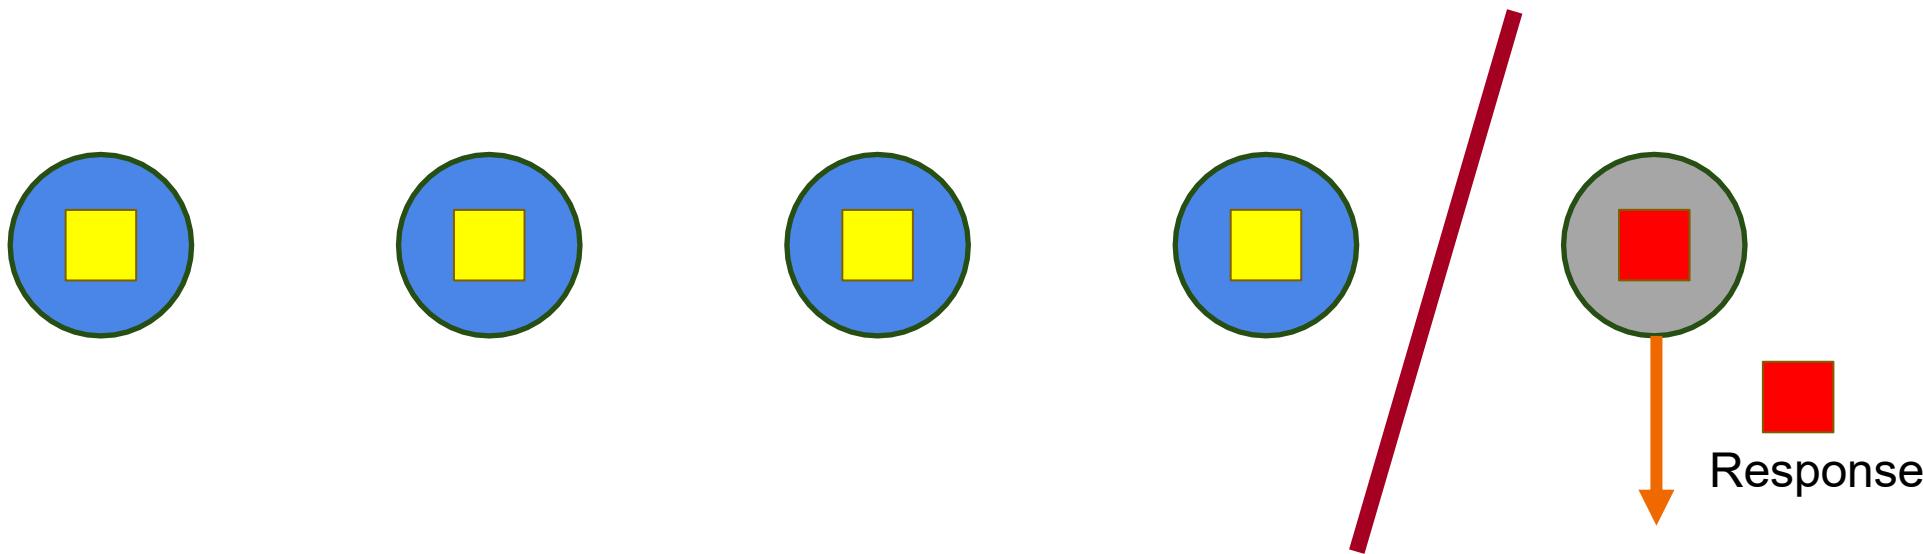
Motivation: Linearizable Reads



Motivation: Linearizable Reads



Motivation: Linearizable Reads



Breaks linearizability in an asynchronous network model

Motivation: Linearizable Reads

- Some existing solutions assume a synchronous model (e.g. Chain Replication)
- Dealing with asynchrony is complicated. One can:
 - **Relax consistency** (e.g. ZooKeeper)
 - Add **extra (costly) steps** to write operations
 - **Synchronize** with other replicas when reading

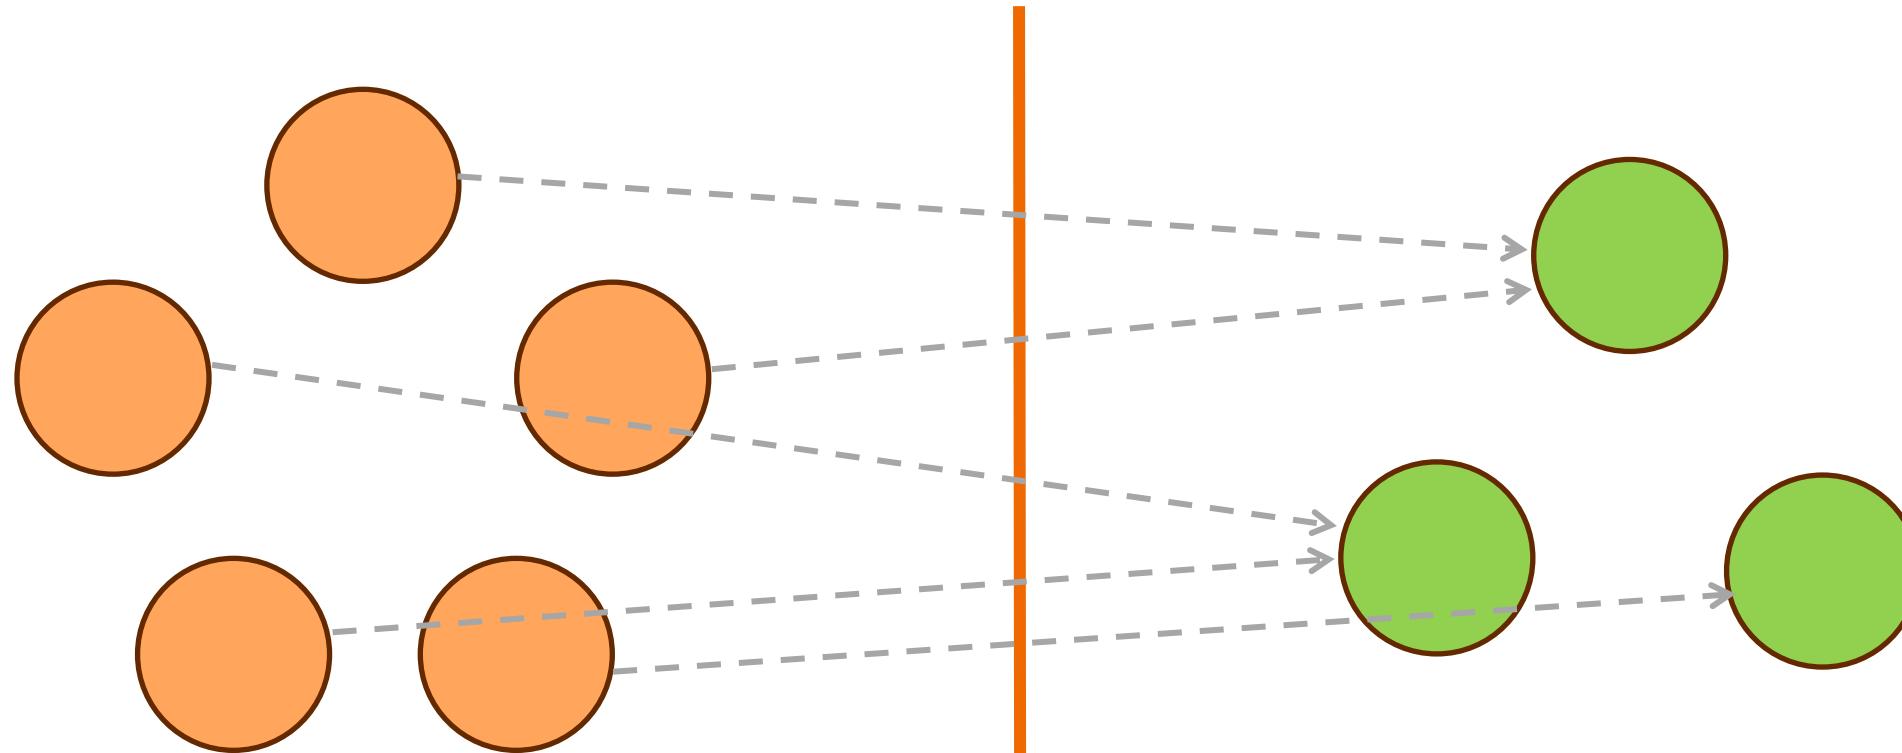
Motivation: Consensus and SMR

- Building blocks of numerous practical replication systems
- Their performance is critical
- Many alternatives have been designed
- Two very relevant ones:
 - (Multi-)Paxos
 - Chain Replication
- Two aspects are often overlooked:
 - Performant linearizable reads
 - **Membership management and reconfiguration**

Motivation : Membership Management

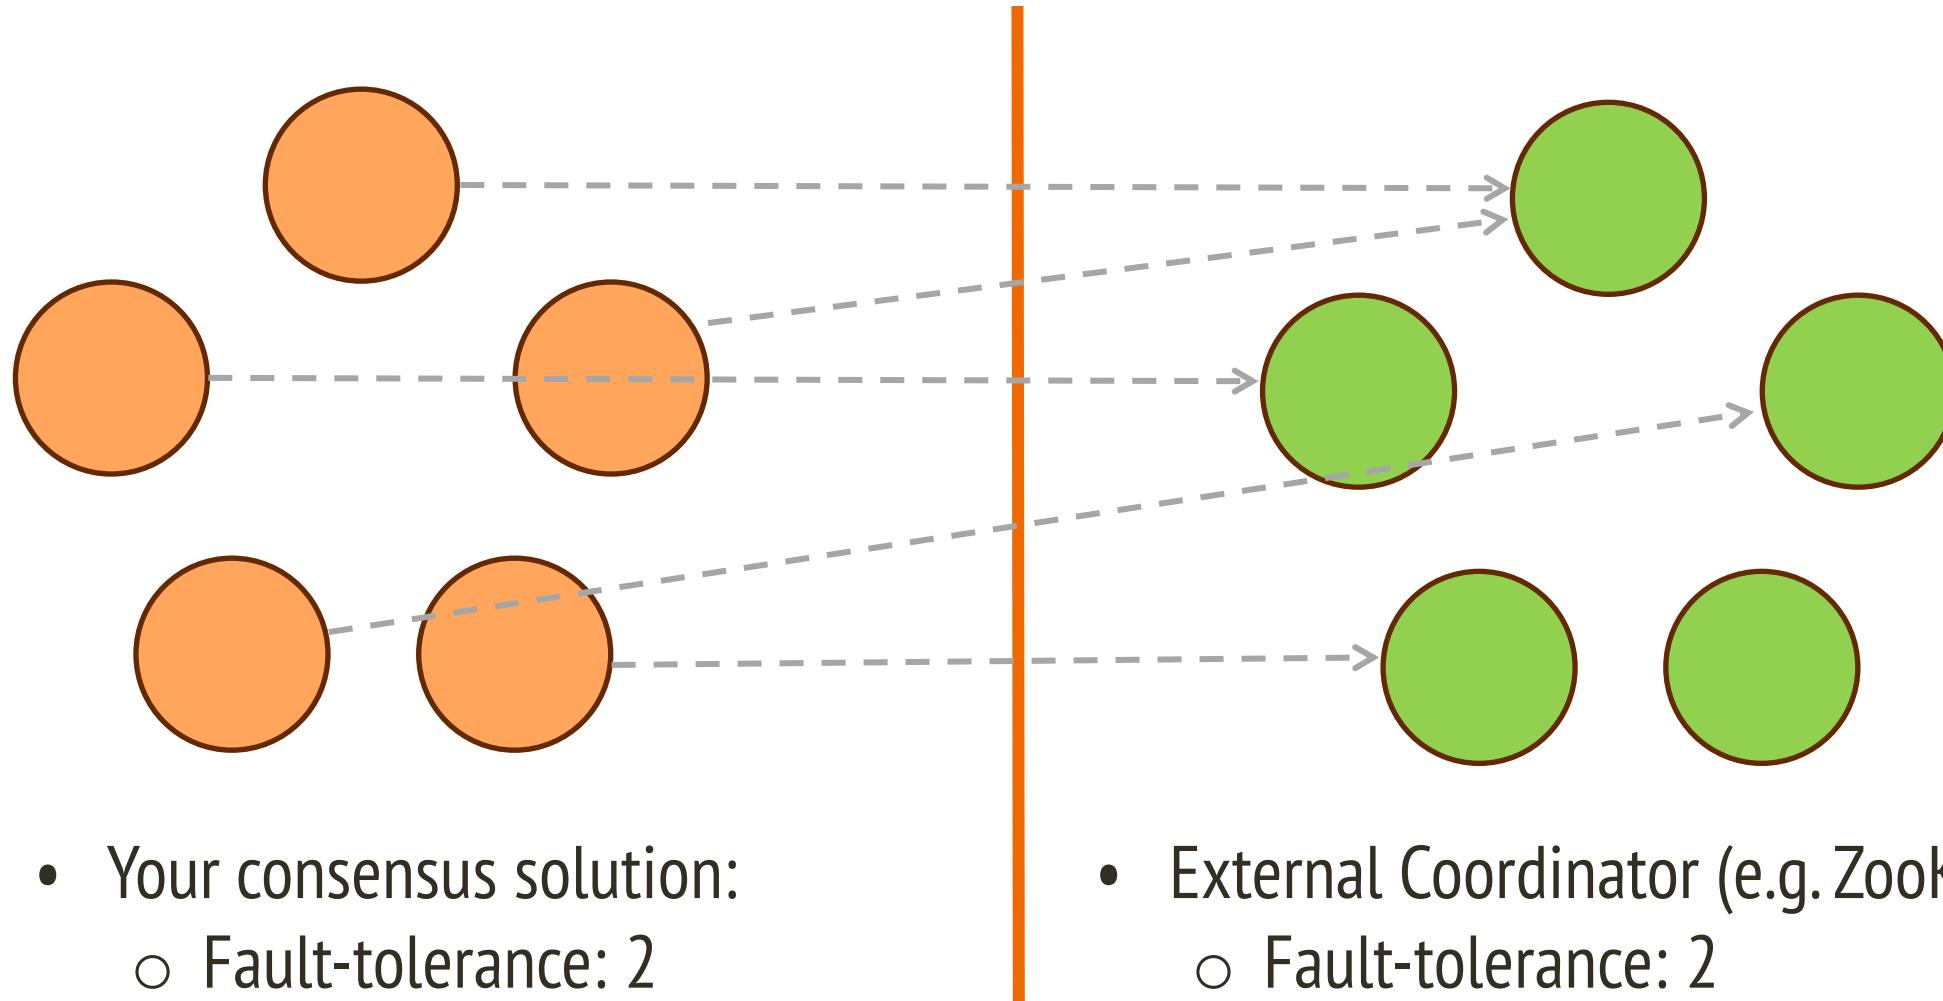
- Most consensus solutions overlook membership management
- Often assume that using an external coordinator service (e.g. ZooKeeper) is trivial and the best solution
- This is not the case:
 - Fault-tolerance becomes complex
 - Complex (and redundant) integration with consensus
 - More vulnerable to partial network partitions

Motivation : Membership Management

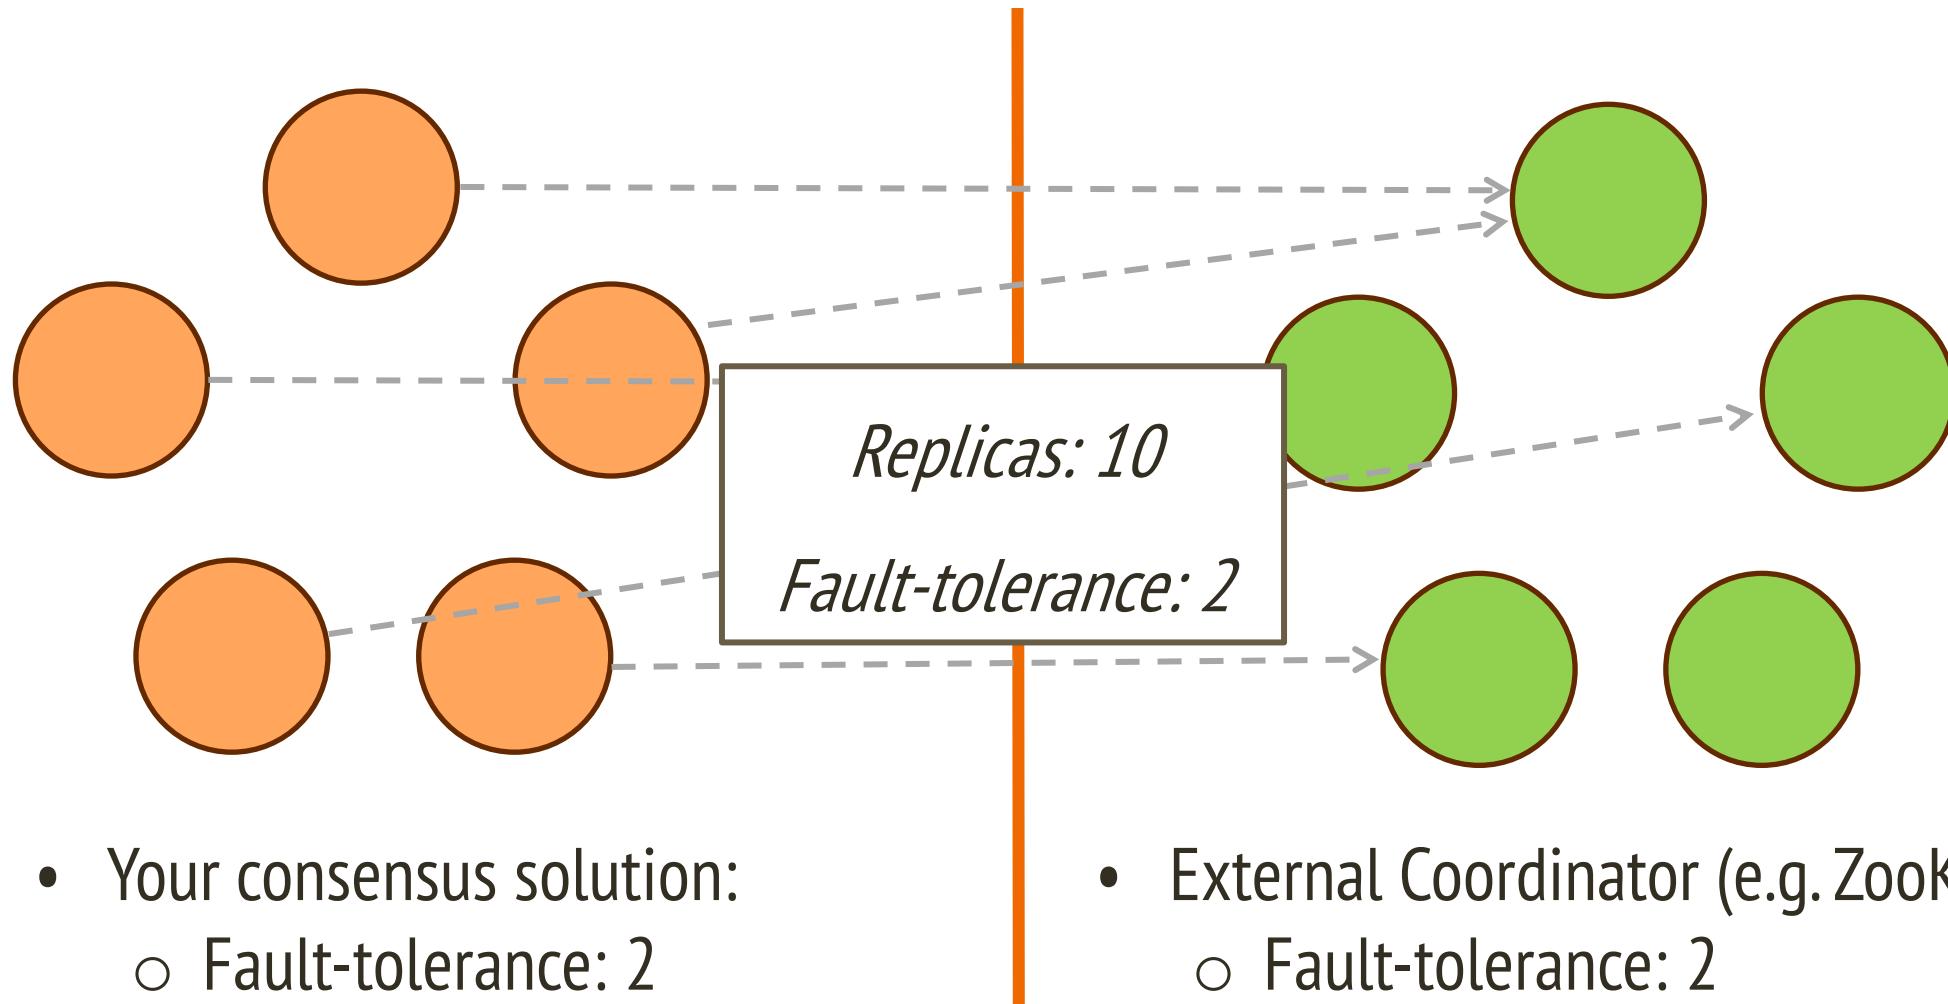


- Your consensus solution:
 - Fault-tolerance: 2
- External Coordinator (e.g. ZooKeeper)
 - Fault-tolerance: 1

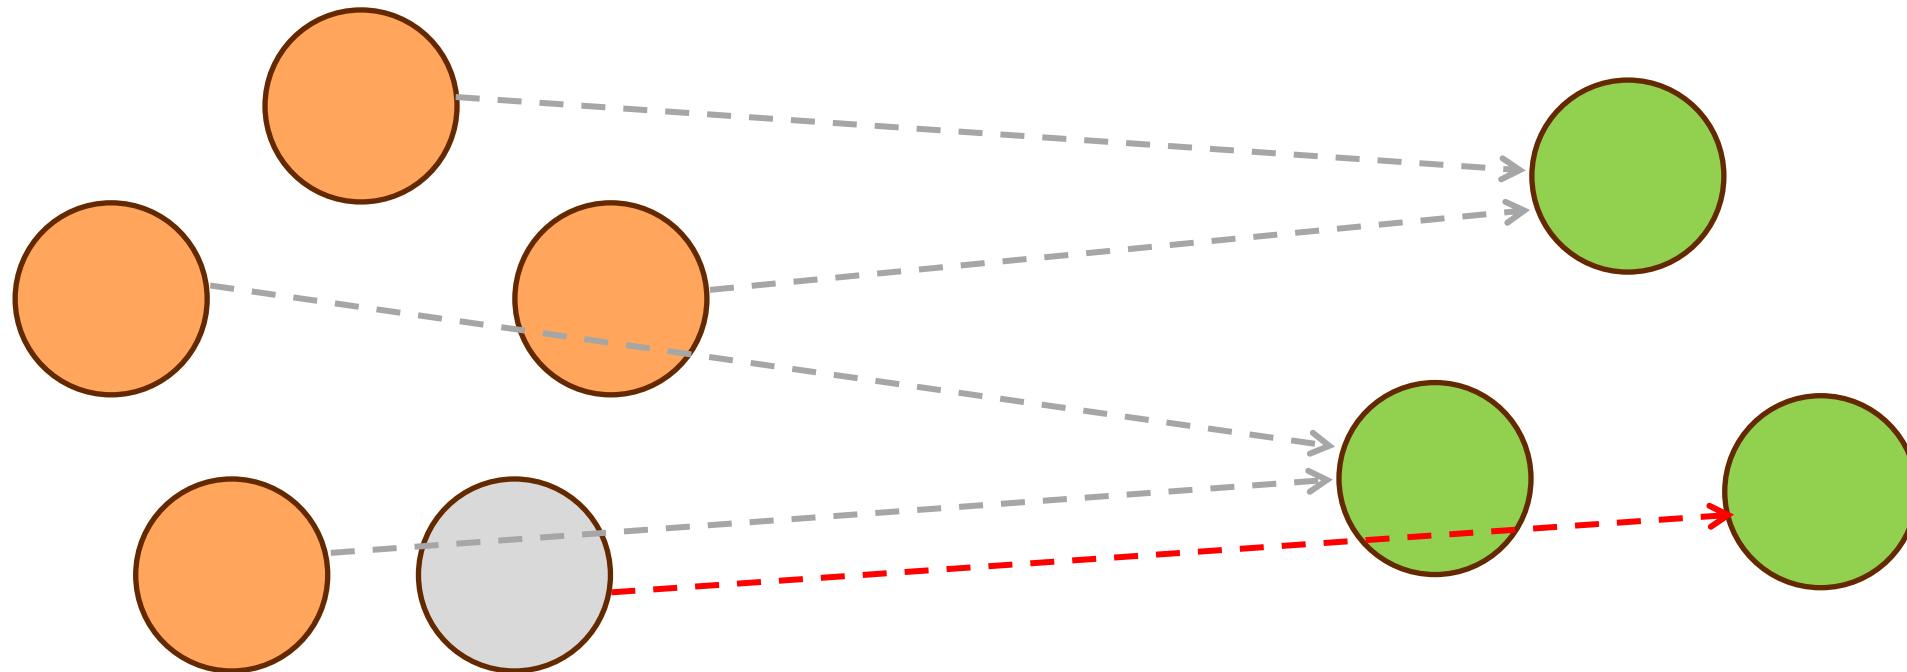
Motivation : Membership Management



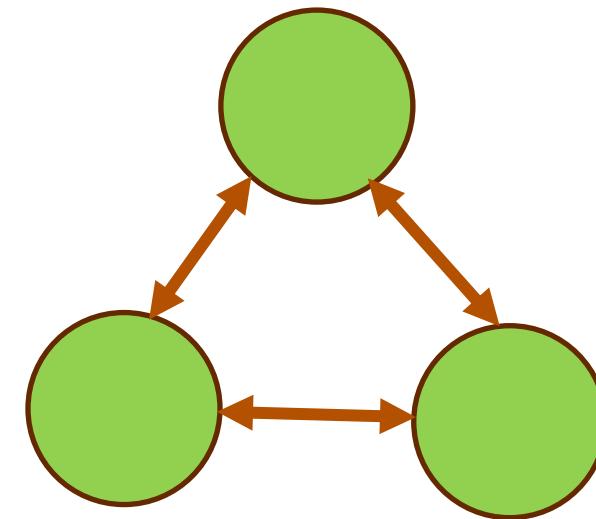
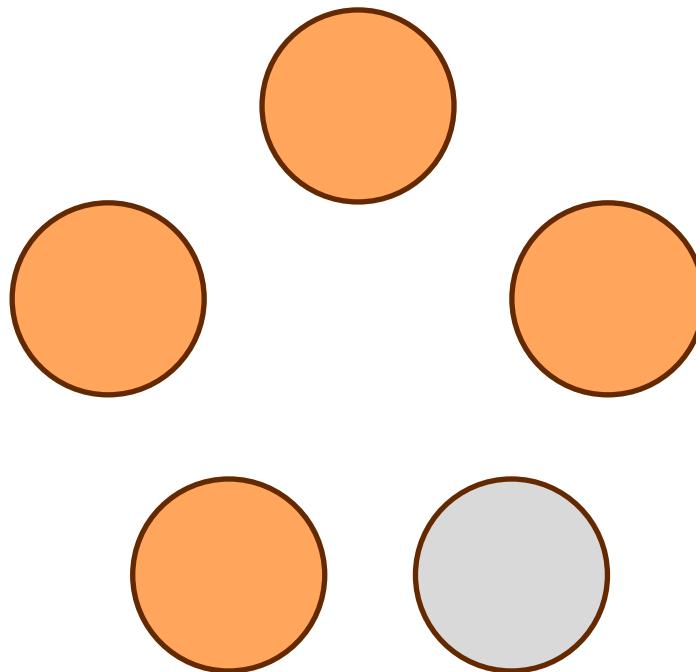
Motivation : Membership Management



Motivation : Membership Management

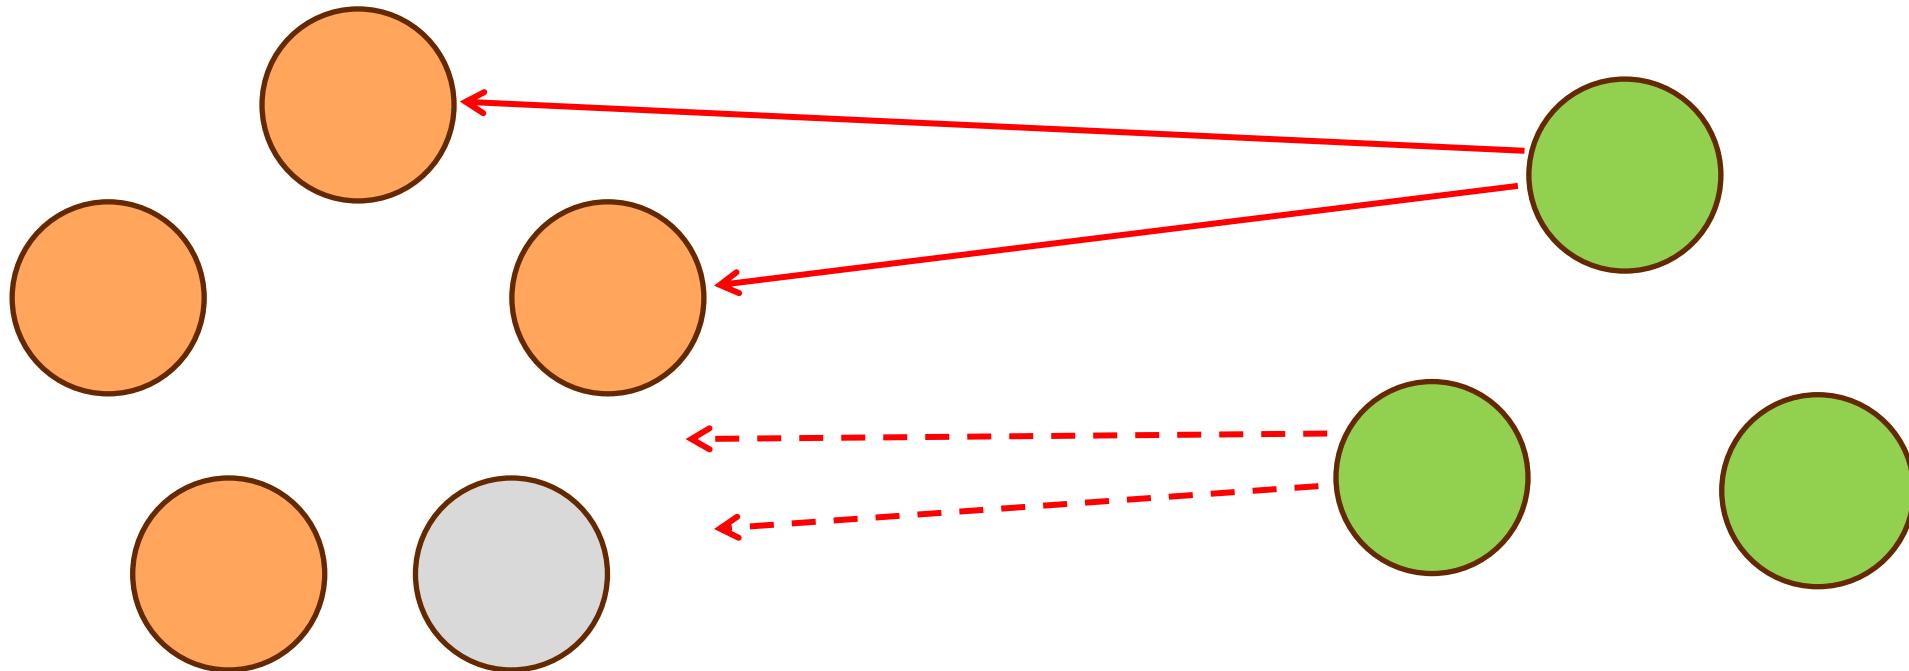


Motivation : Membership Management



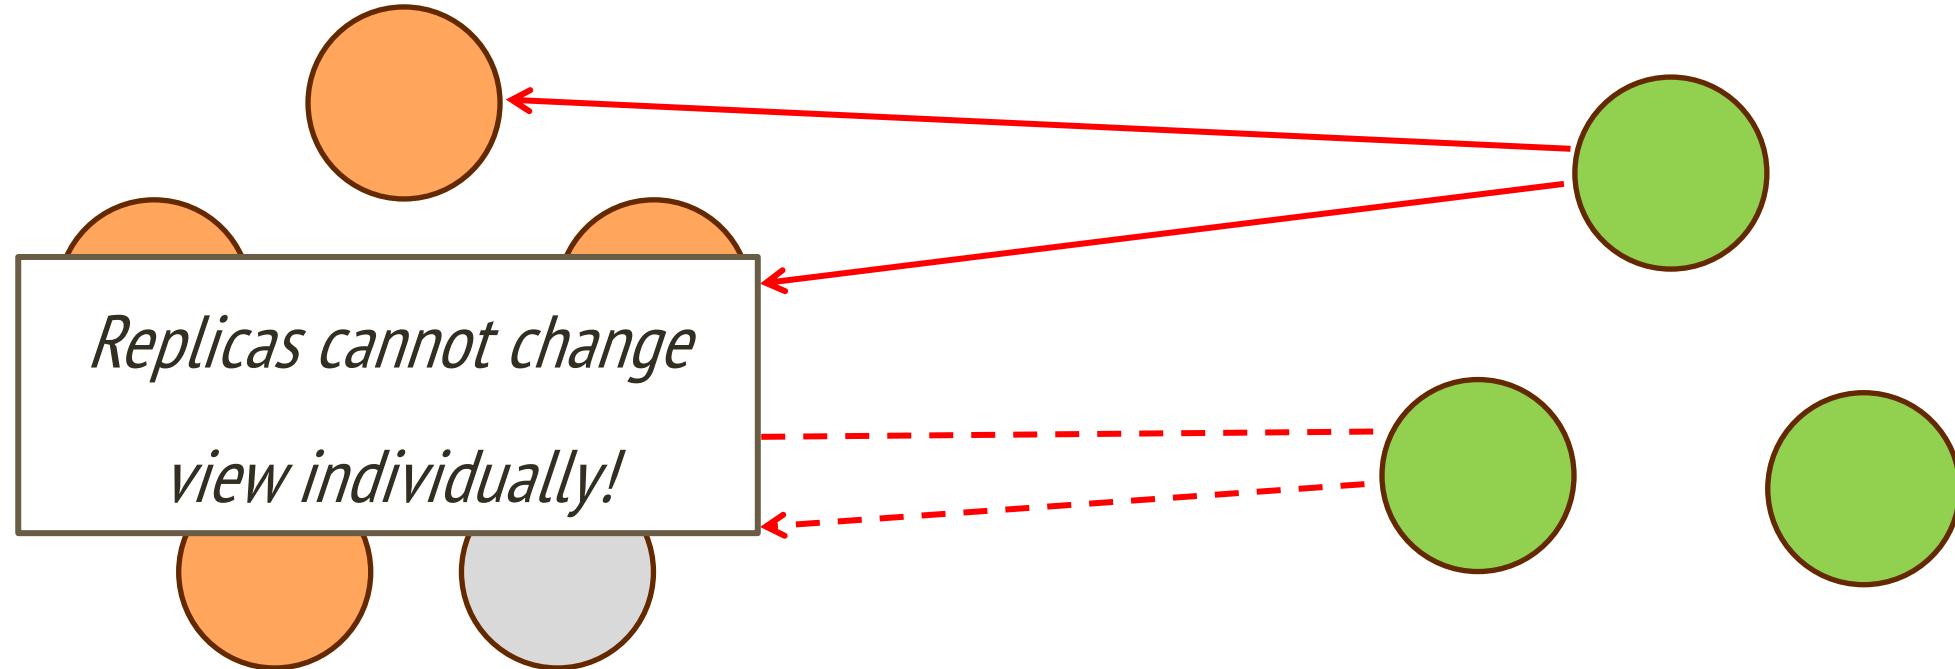
*Consensus round to
decide removal*

Motivation : Membership Management

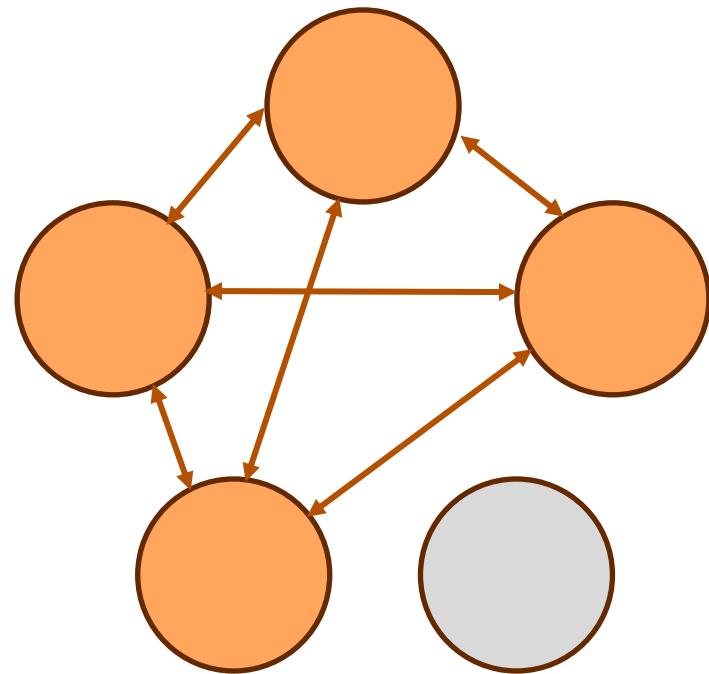


Asynchronous design propagation to consensus replicas

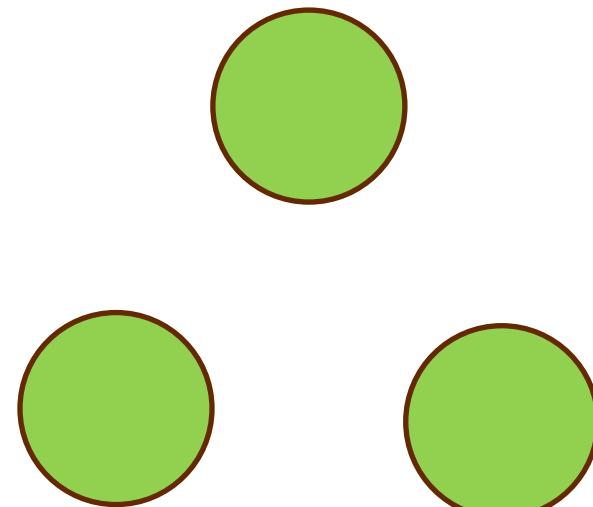
Motivation : Membership Management



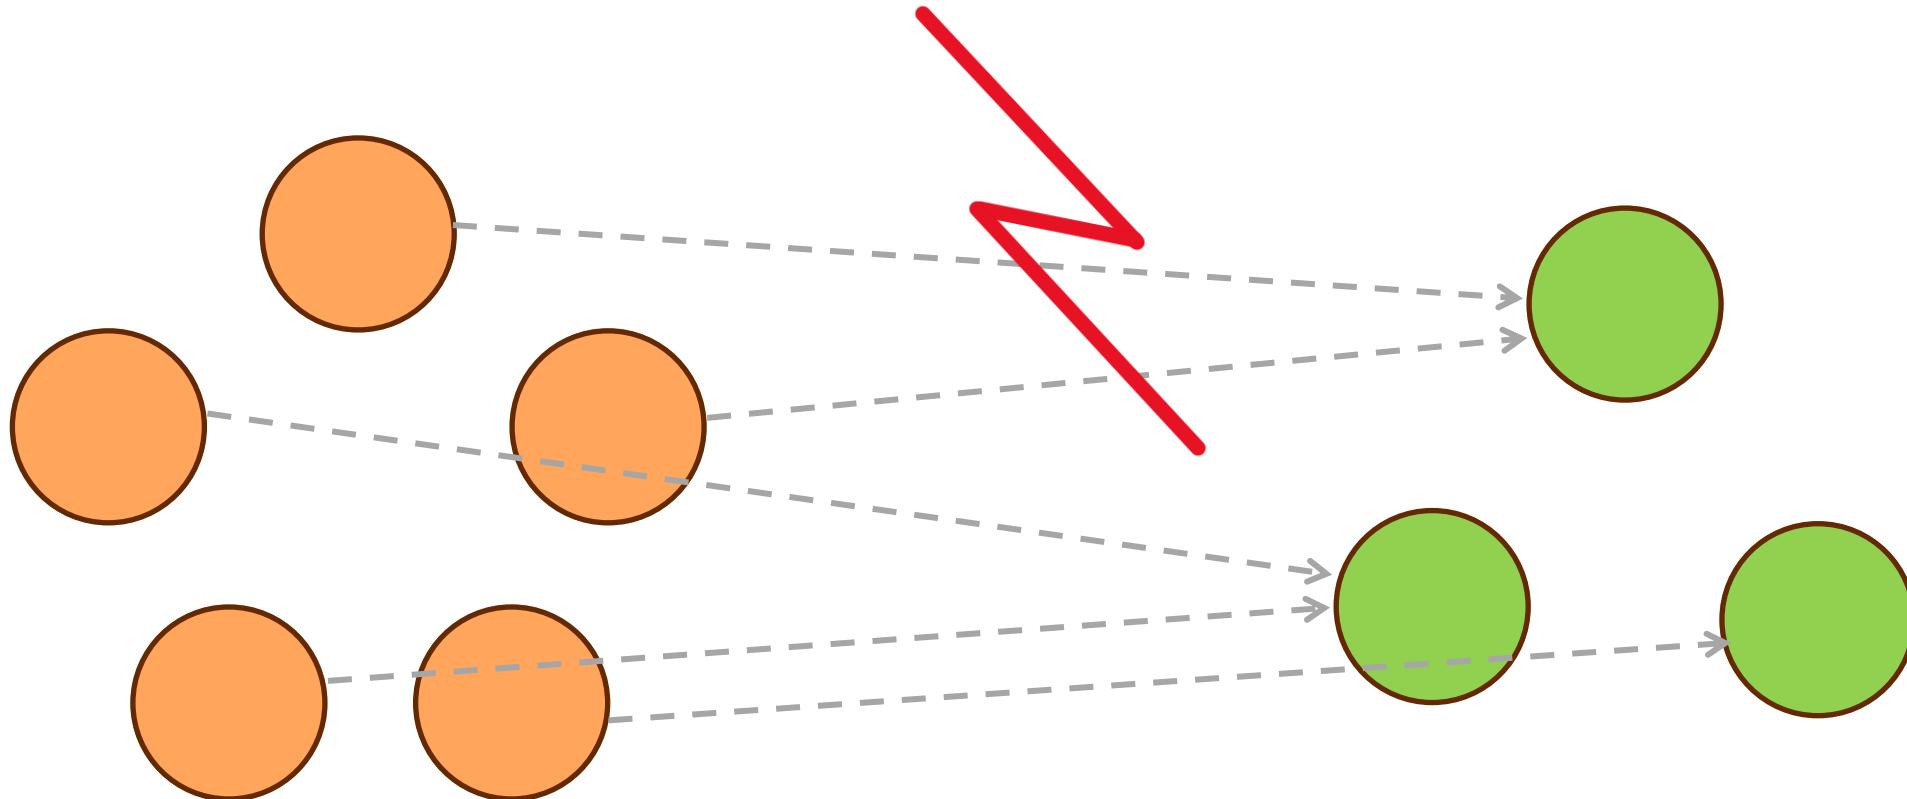
Motivation : Membership Management



*Another redundant consensus
round is required*

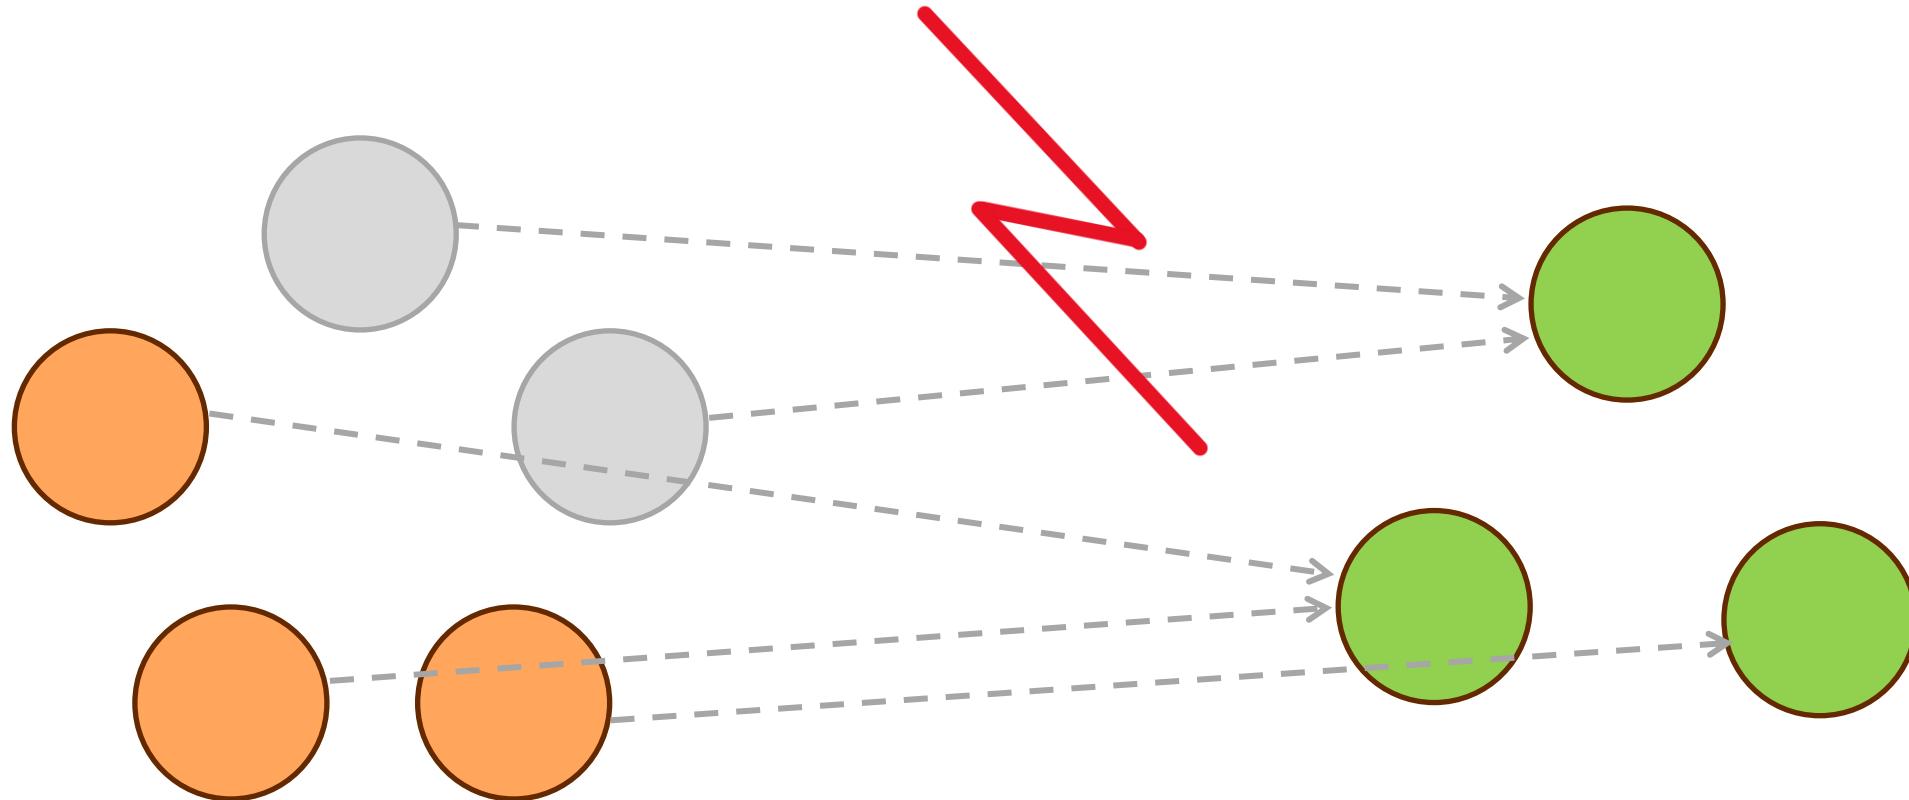


Motivation : Membership Management



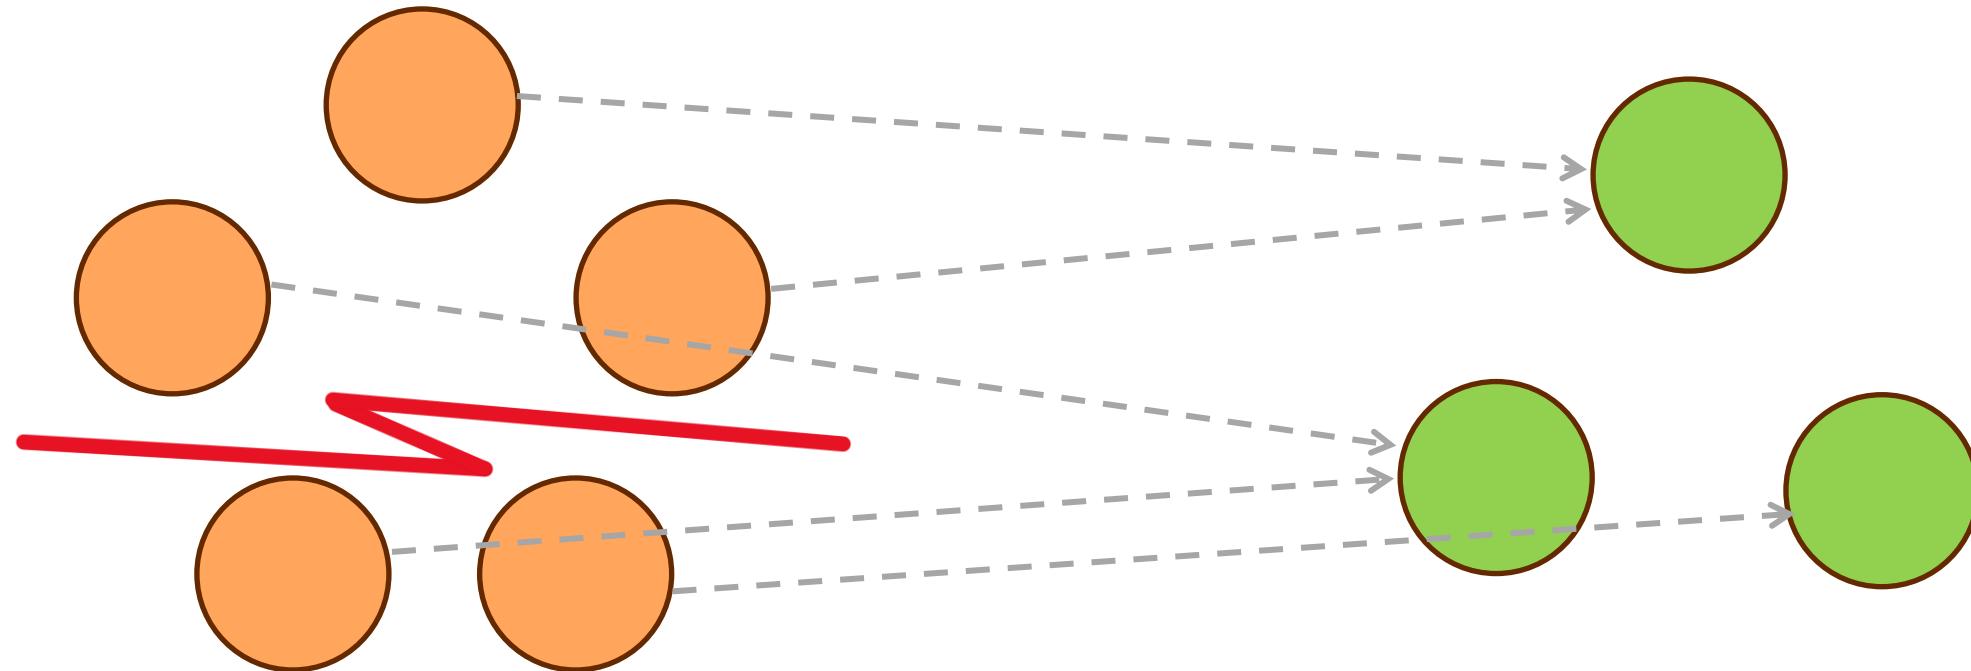
*Partition between coordinator and
consensus replicas*

Motivation : Membership Management



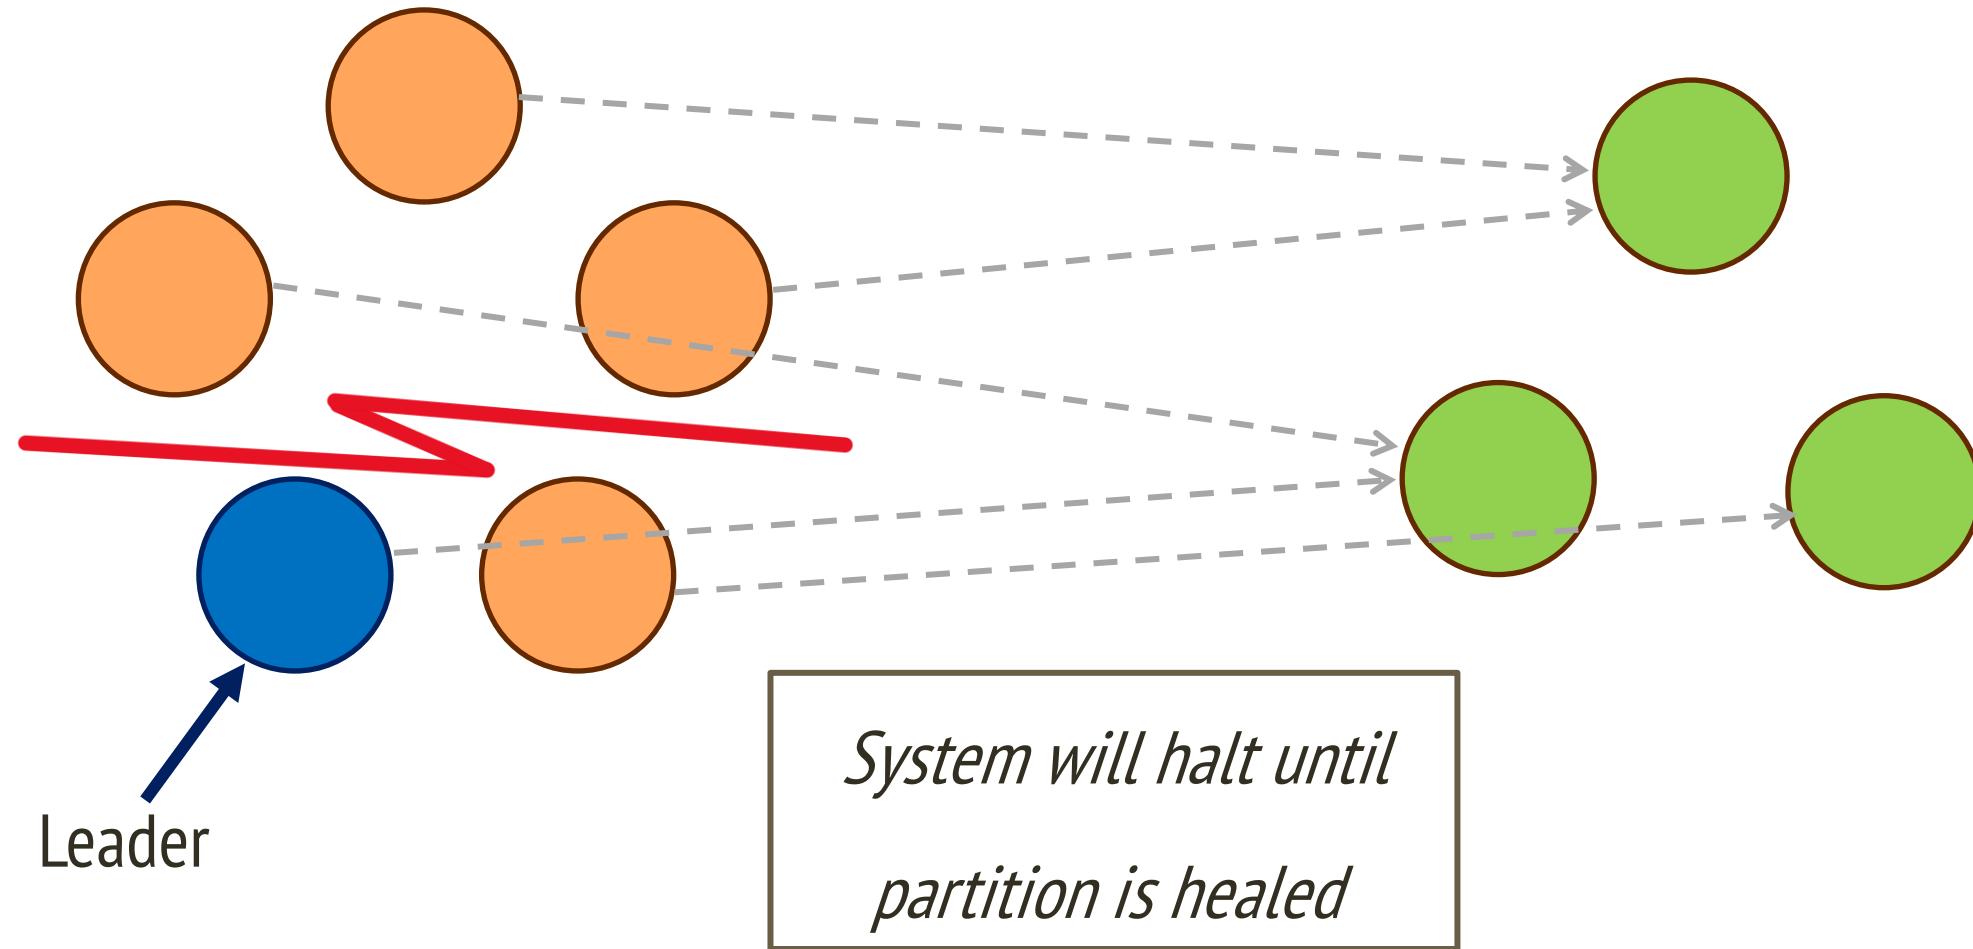
Correct replicas will be removed

Motivation : Membership Management



Partition between consensus replicas

Motivation : Membership Management



Proposal : ChainPaxos

Novel consensus algorithm:

- Combining the best properties of Multi-Paxos and Chain Replication
 - Correction in an asynchronous network
 - Constant message complexity
- Going beyond existing solutions:
 - Maximizing throughput of both read and write operations
 - Providing local linearizable reads in any replica
 - Integrated reconfiguration and fault-tolerance

State of ChainPaxos Nodes

1

chain : *array of nodes*

self : *node*

c_{nextok} : *node*

c_{sleader} : *node*

marked : *set of node*

2

n_{p_{leader}} : *int*

inst : *map int × PaxosInst*

3

submitted : *set of requests*

pending : *set of requests*

4

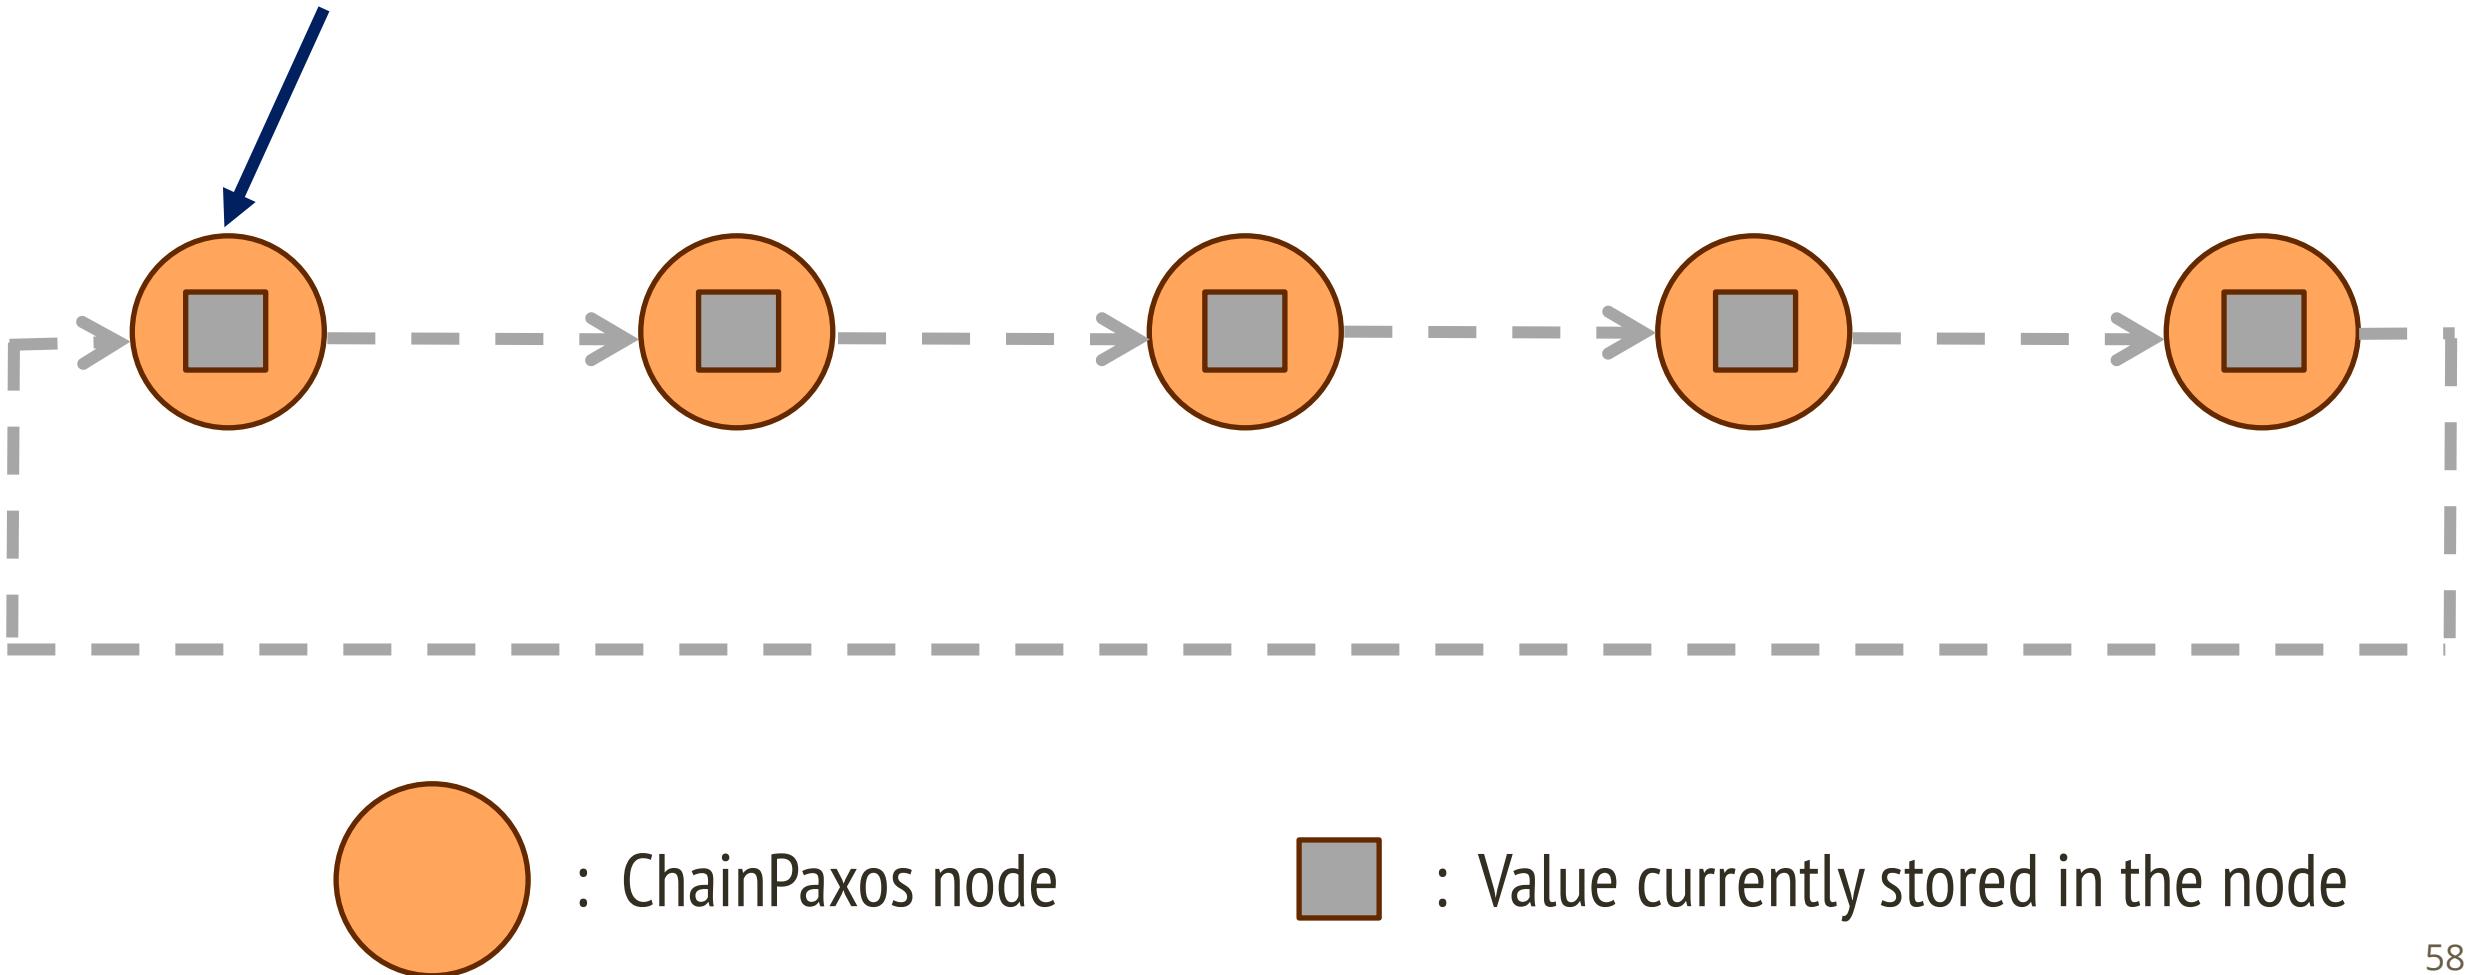
max_{ack} : *int*

max_{acpt} : *int*

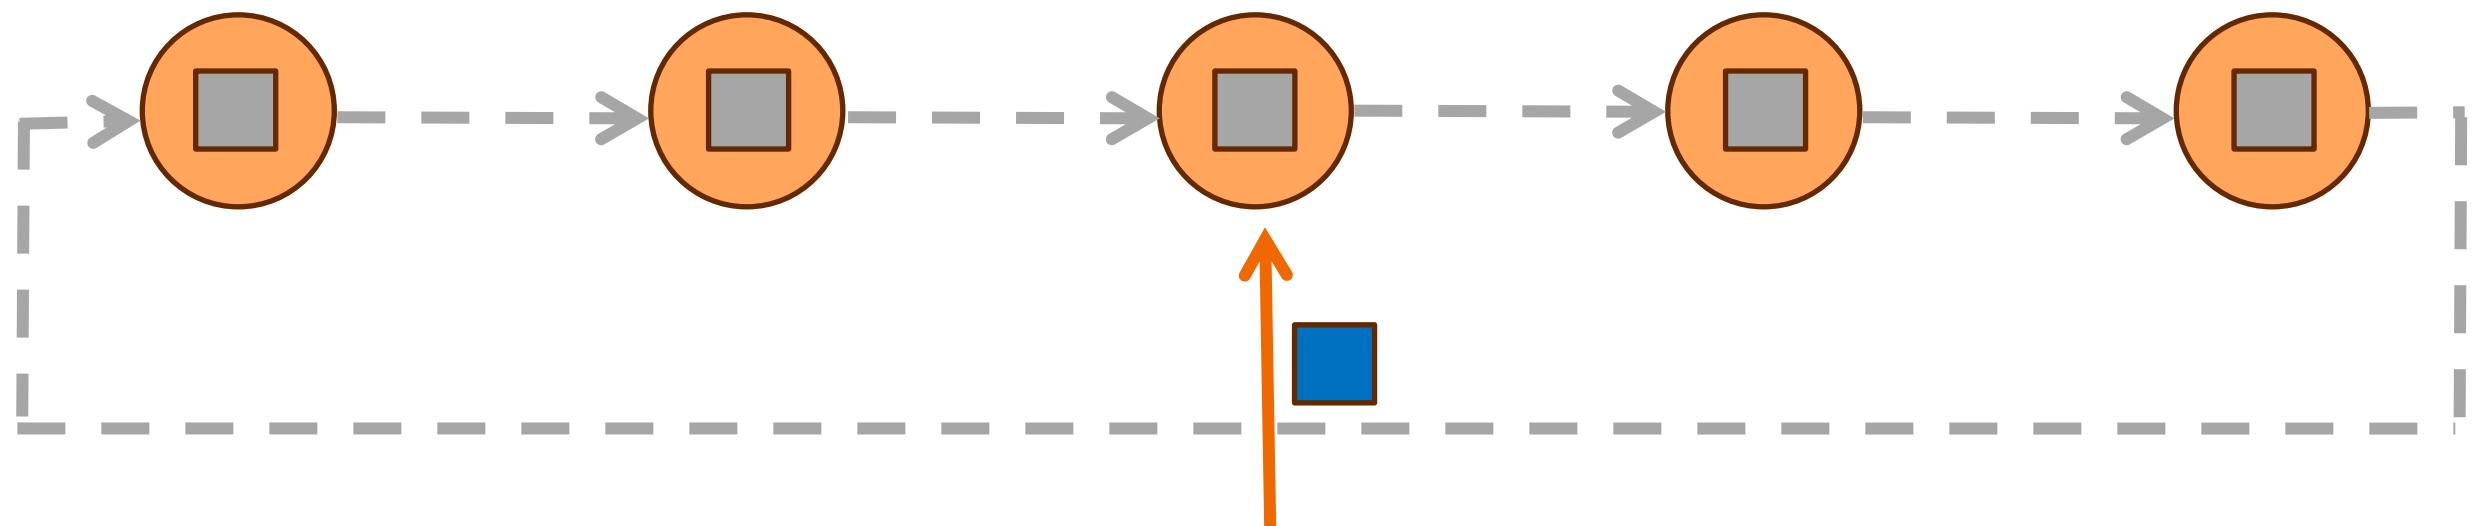
amLeader : *bool*

ChainPaxos : Write Path

Leader (regular Multi-Paxos election)

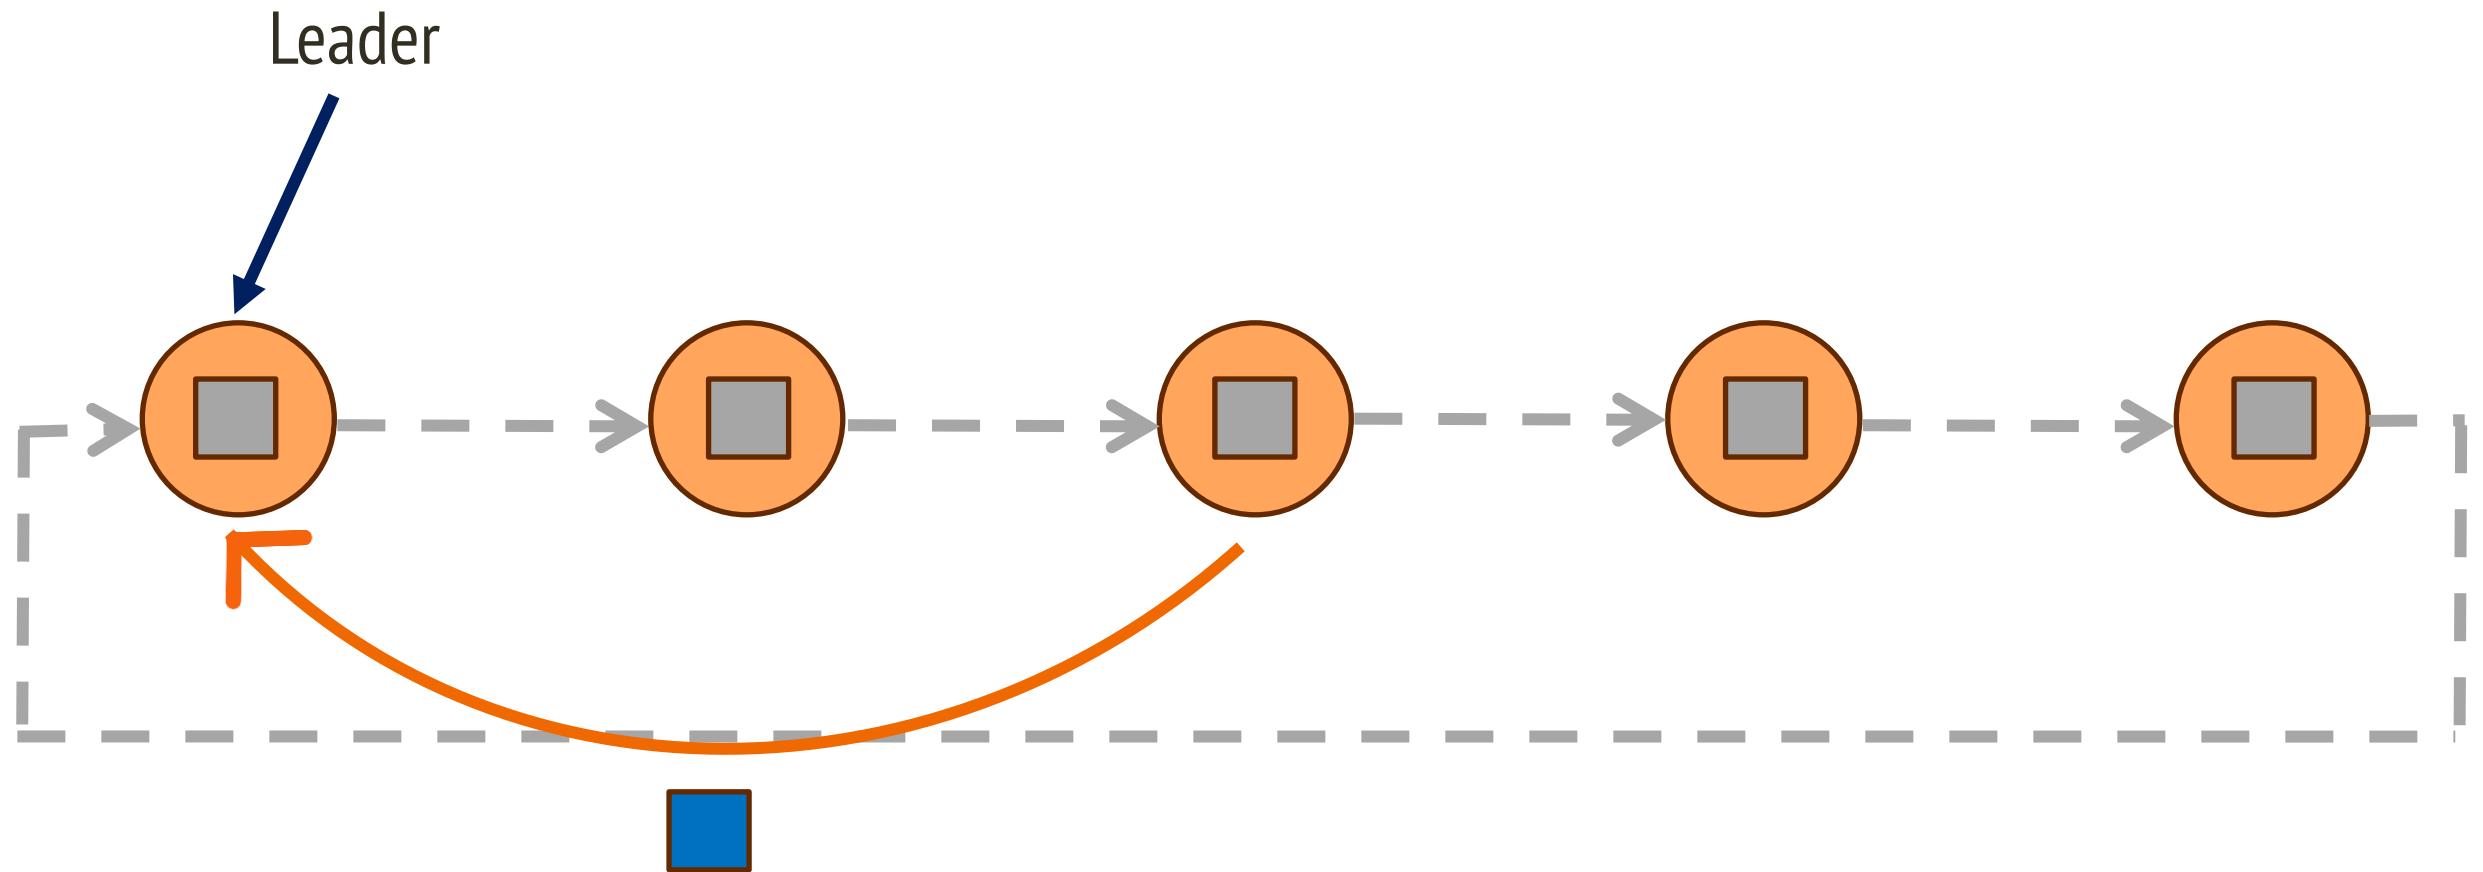


ChainPaxos : Write Path

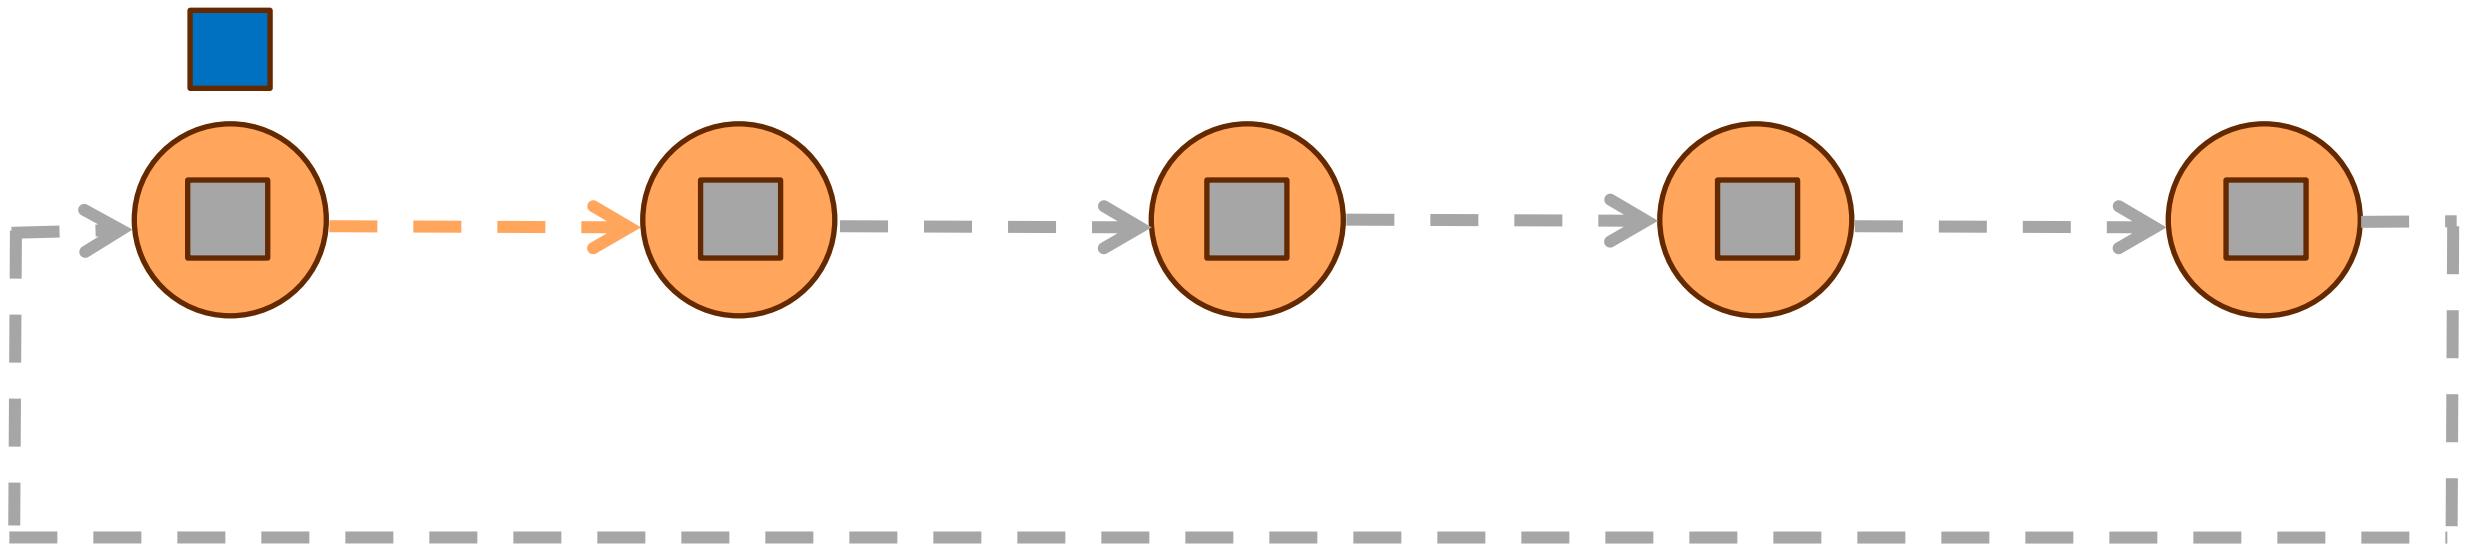


: New value to be written

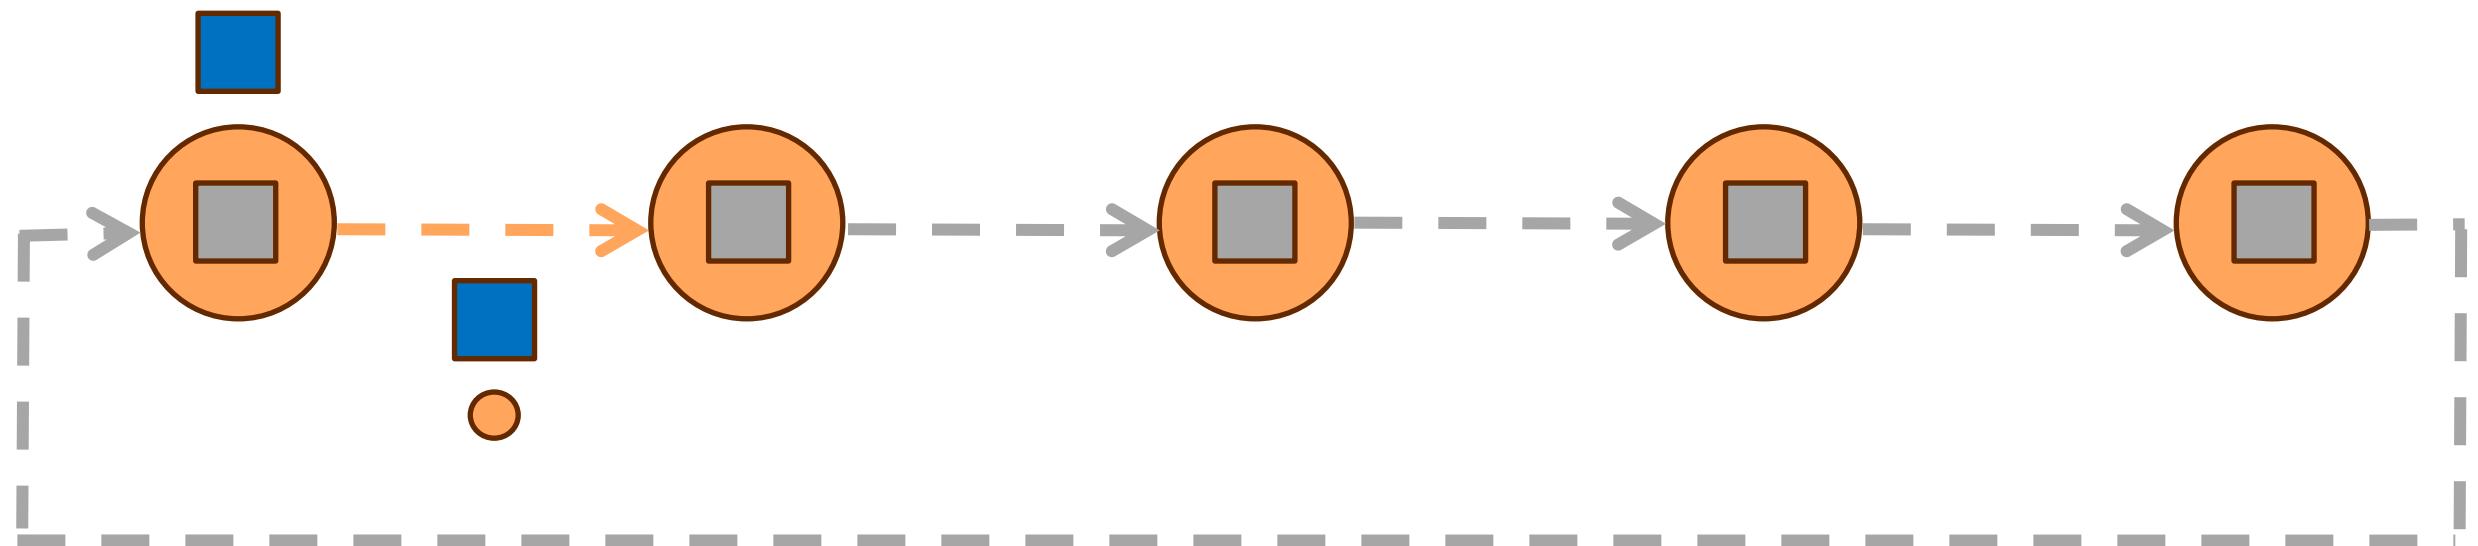
ChainPaxos : Write Path



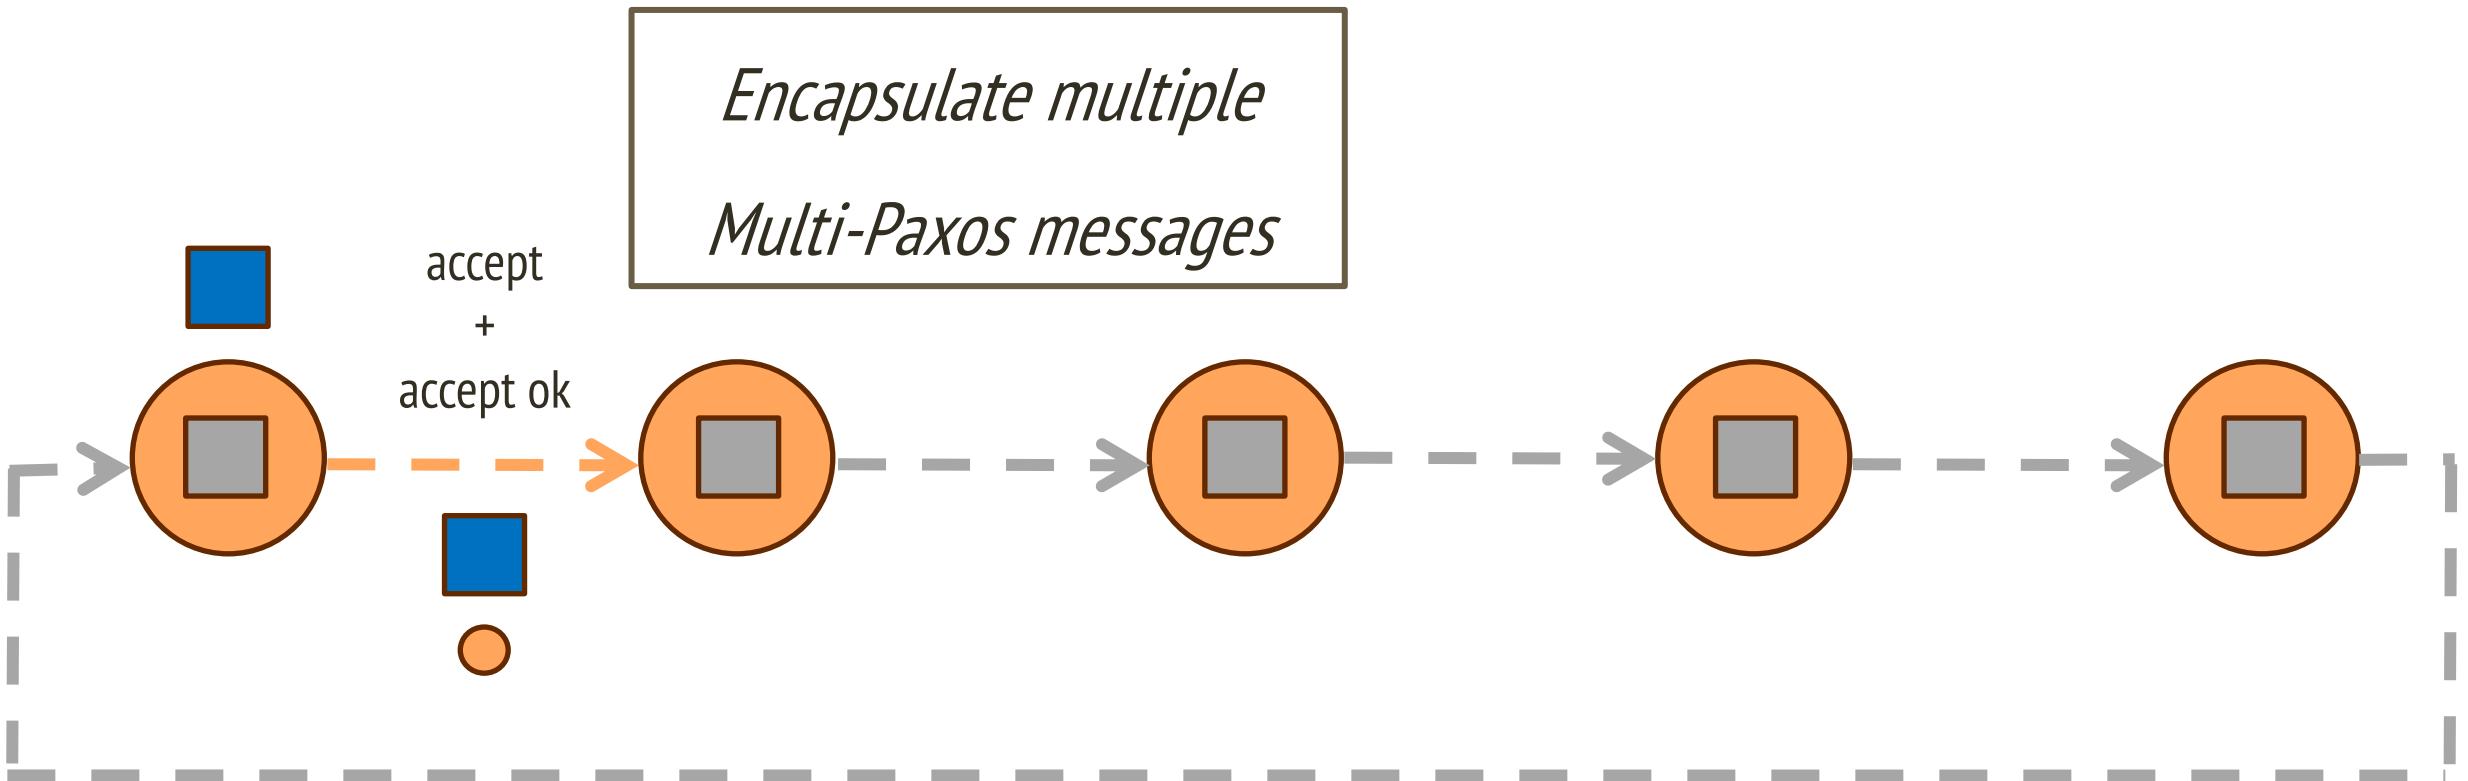
ChainPaxos : Write Path



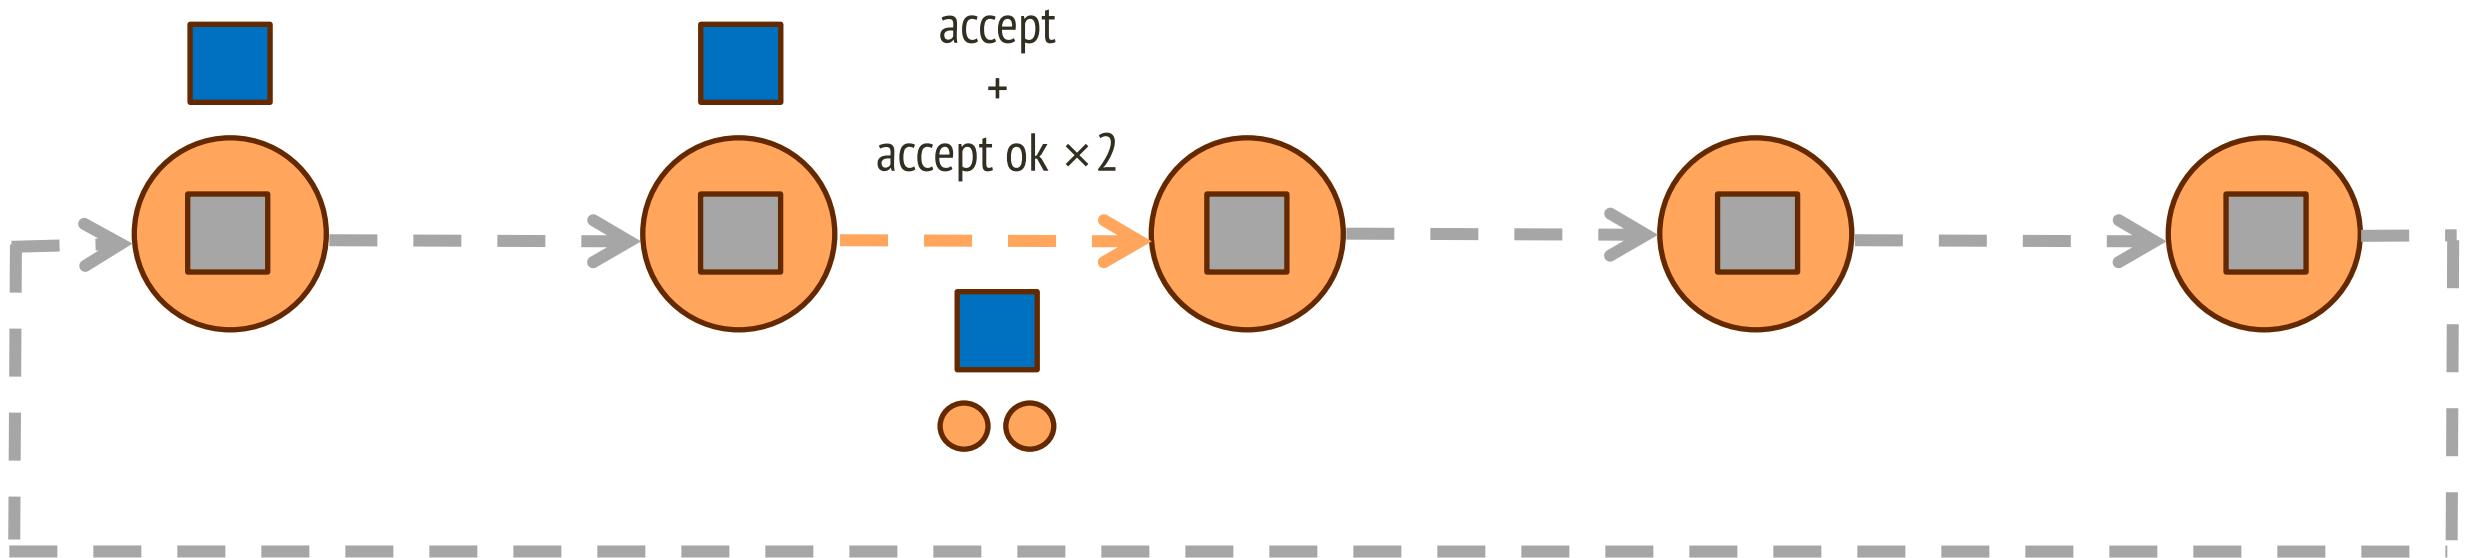
ChainPaxos : Write Path



ChainPaxos : Write Path

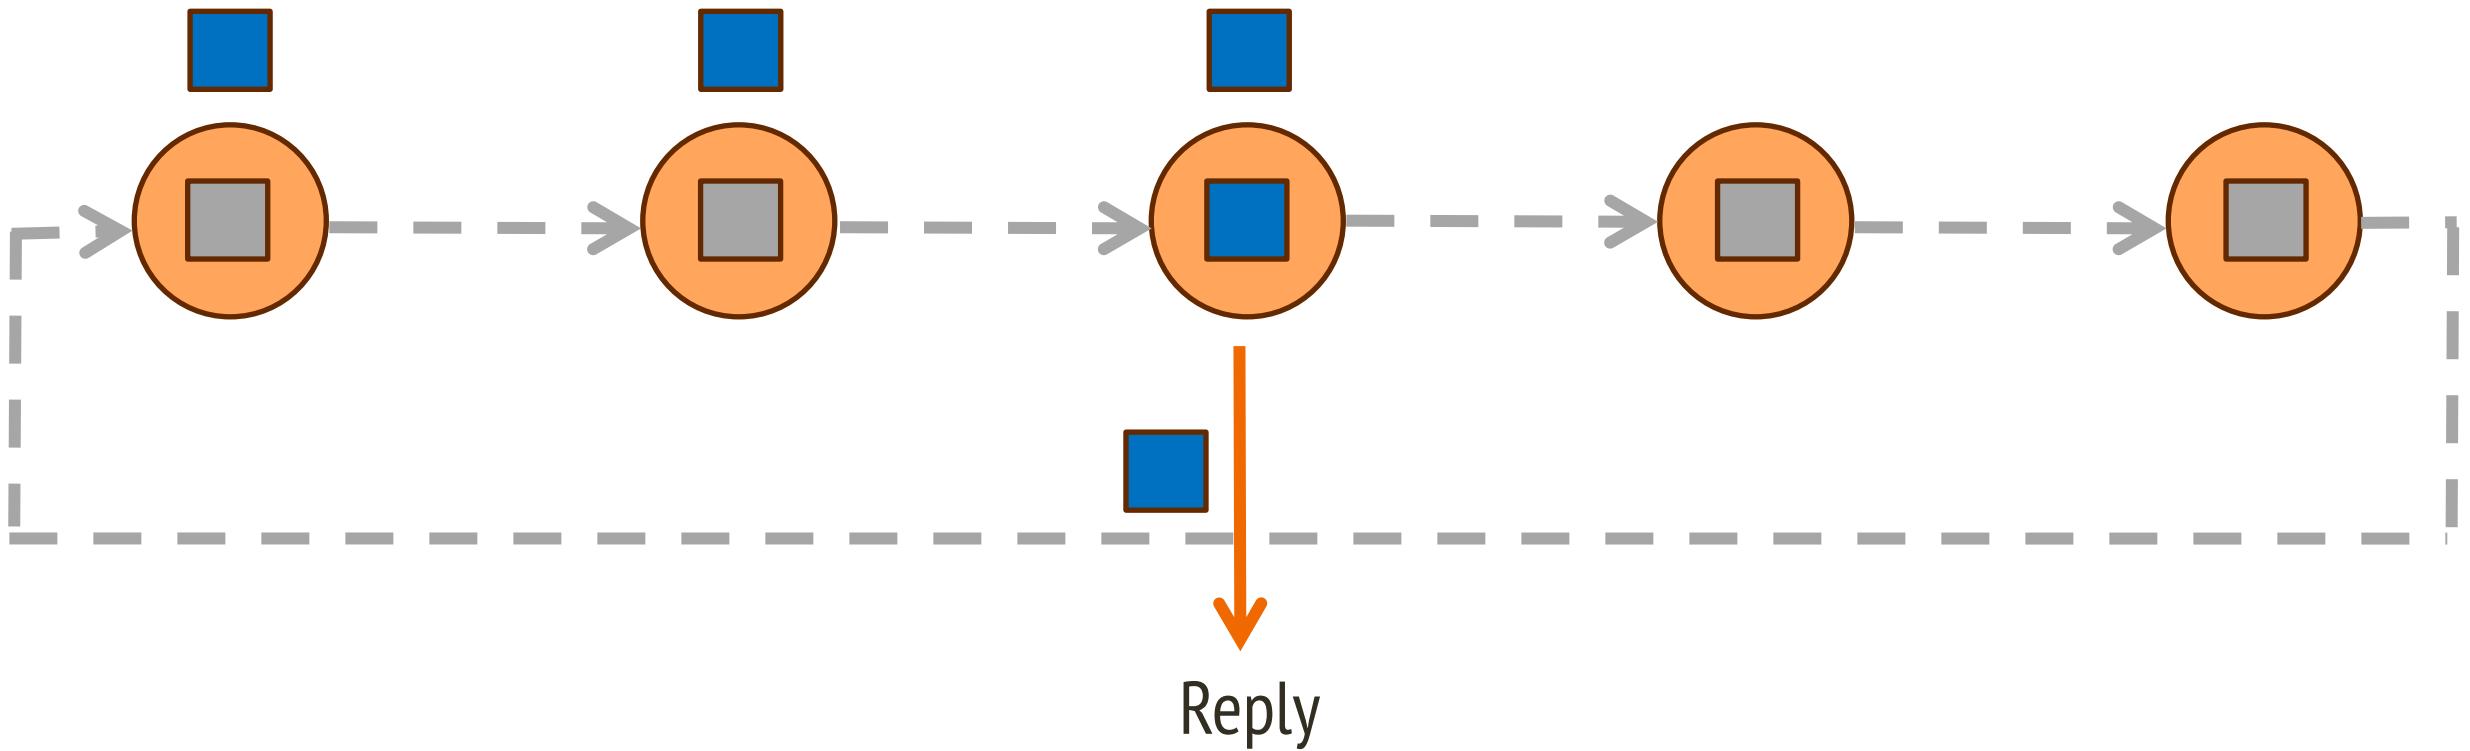


ChainPaxos : Write Path

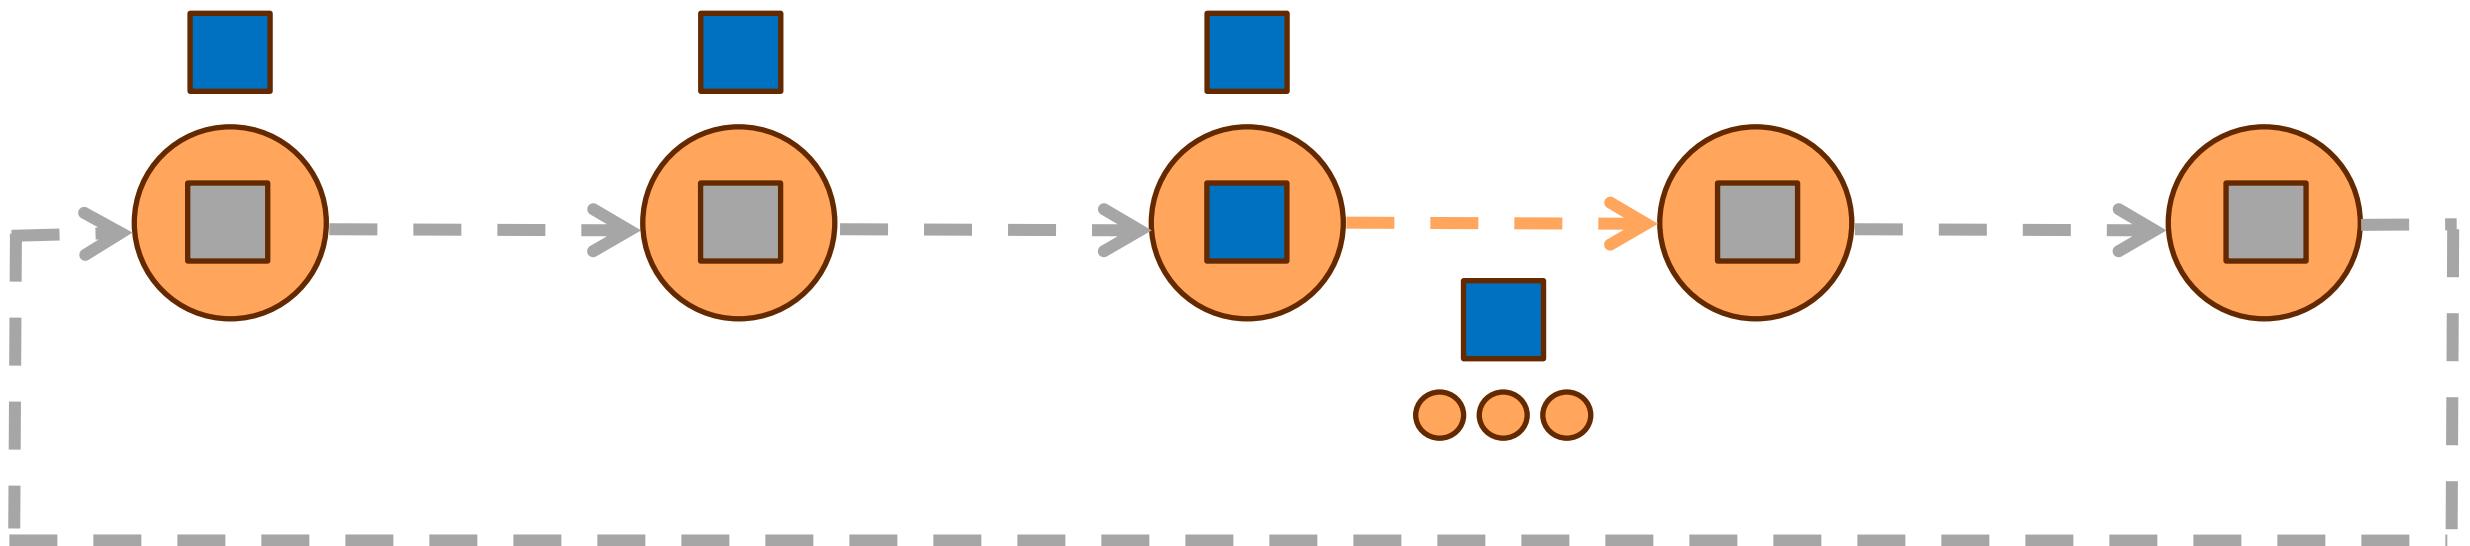


ChainPaxos : Write Path

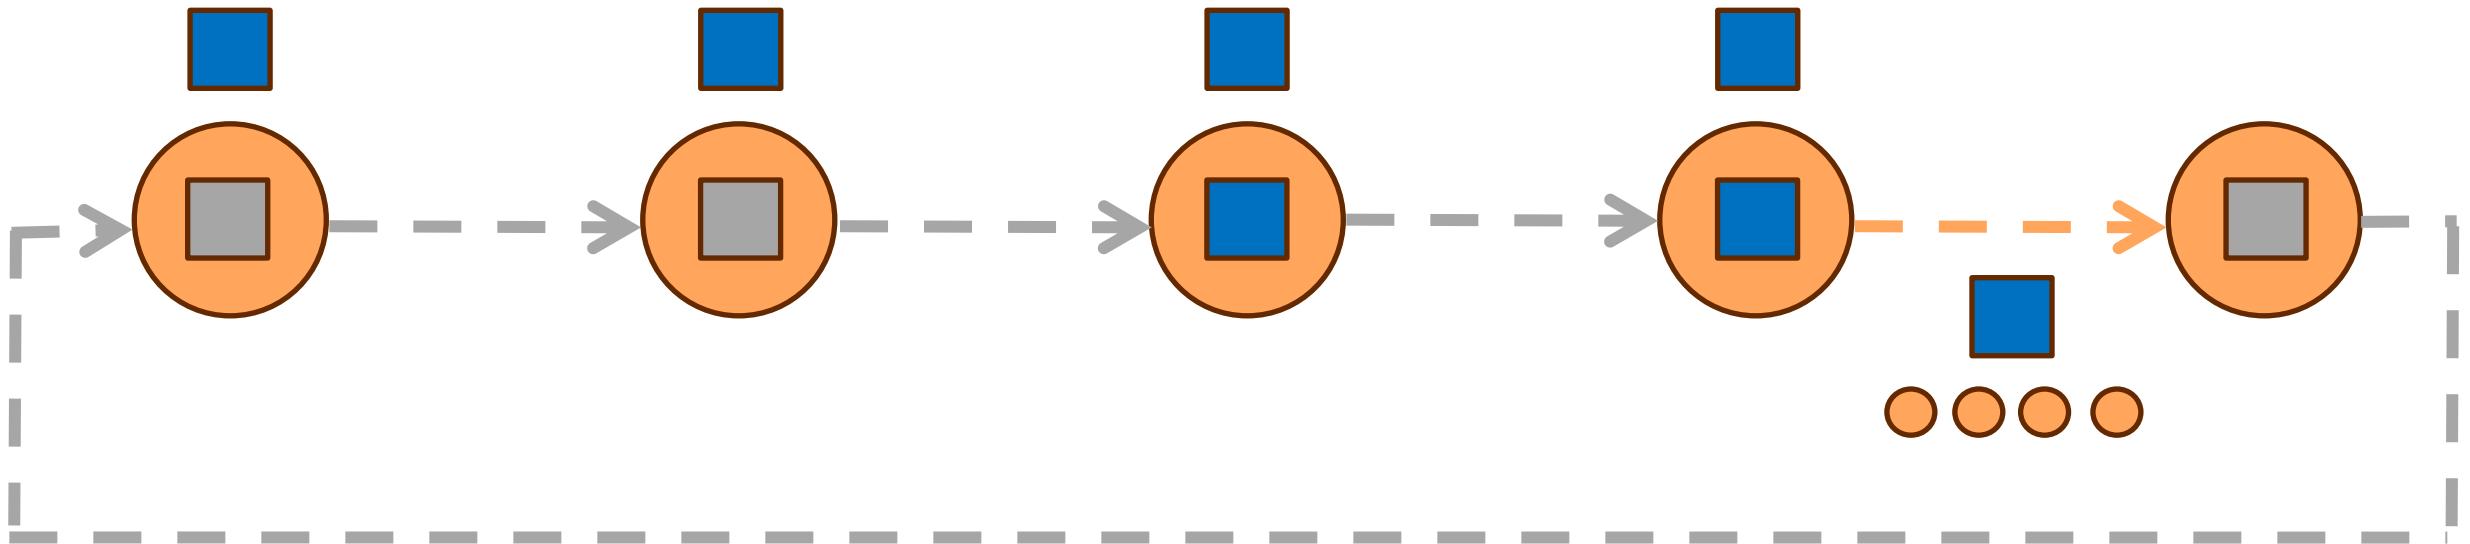
Operation is "locked-in"



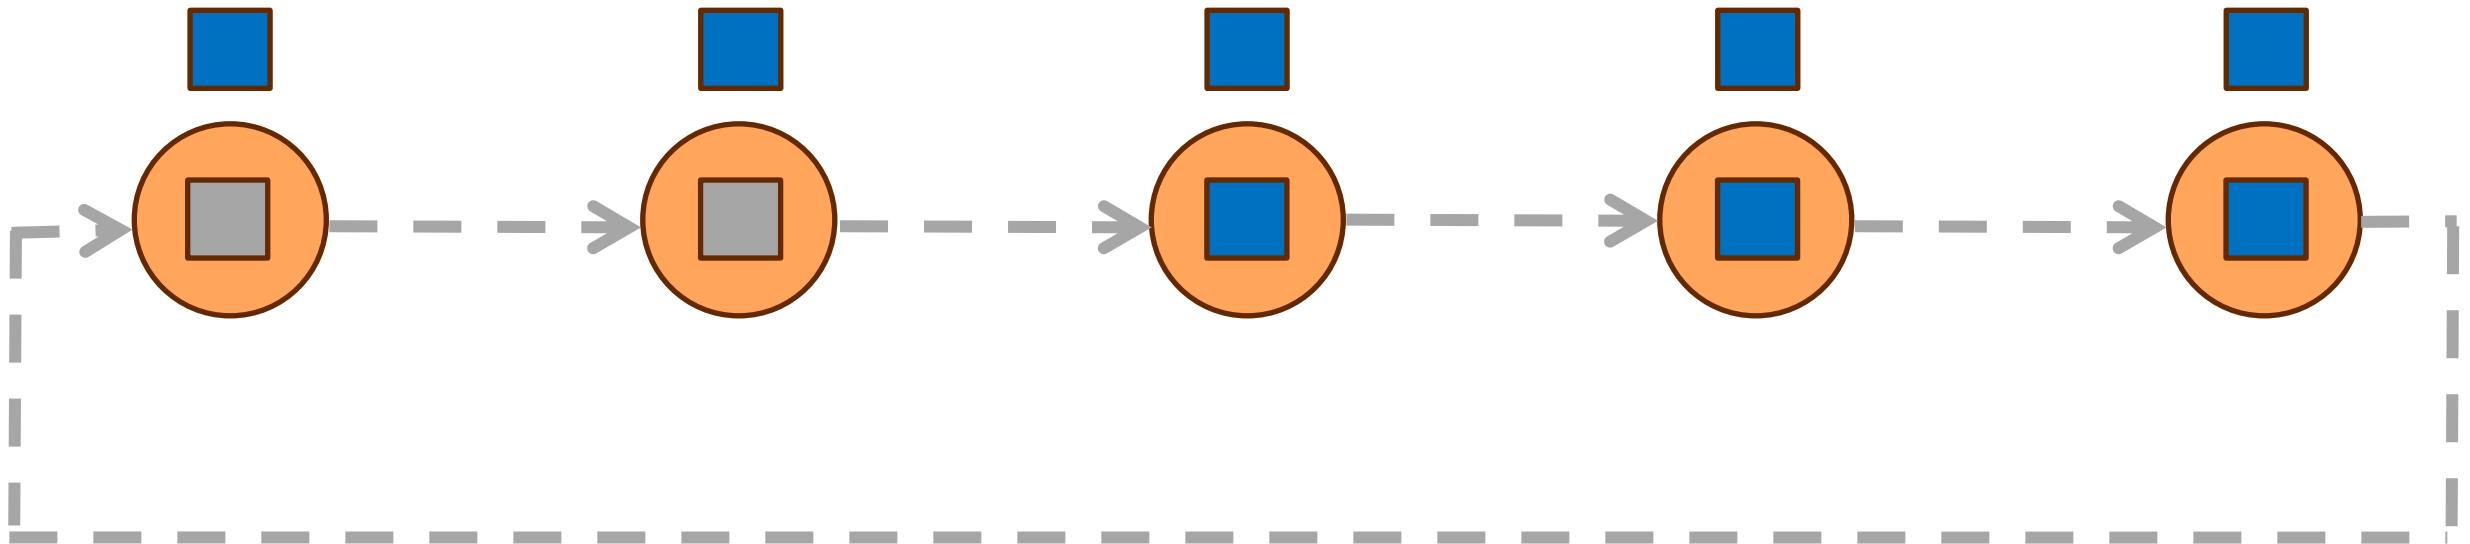
ChainPaxos : Write Path



ChainPaxos : Write Path

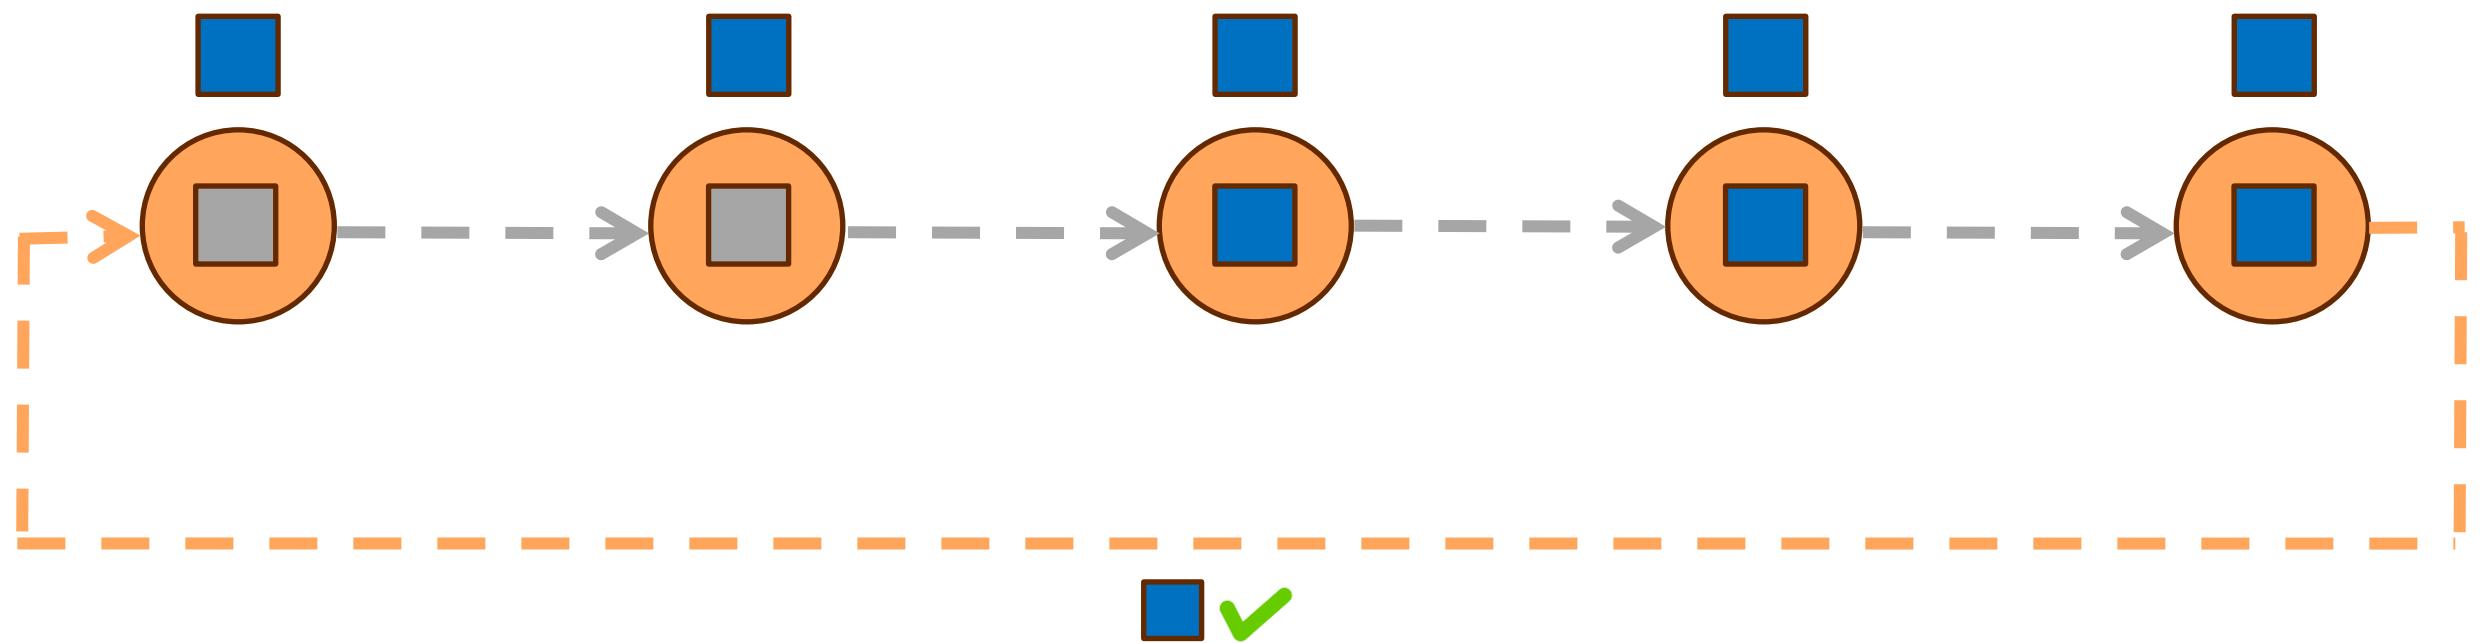


ChainPaxos : Write Path

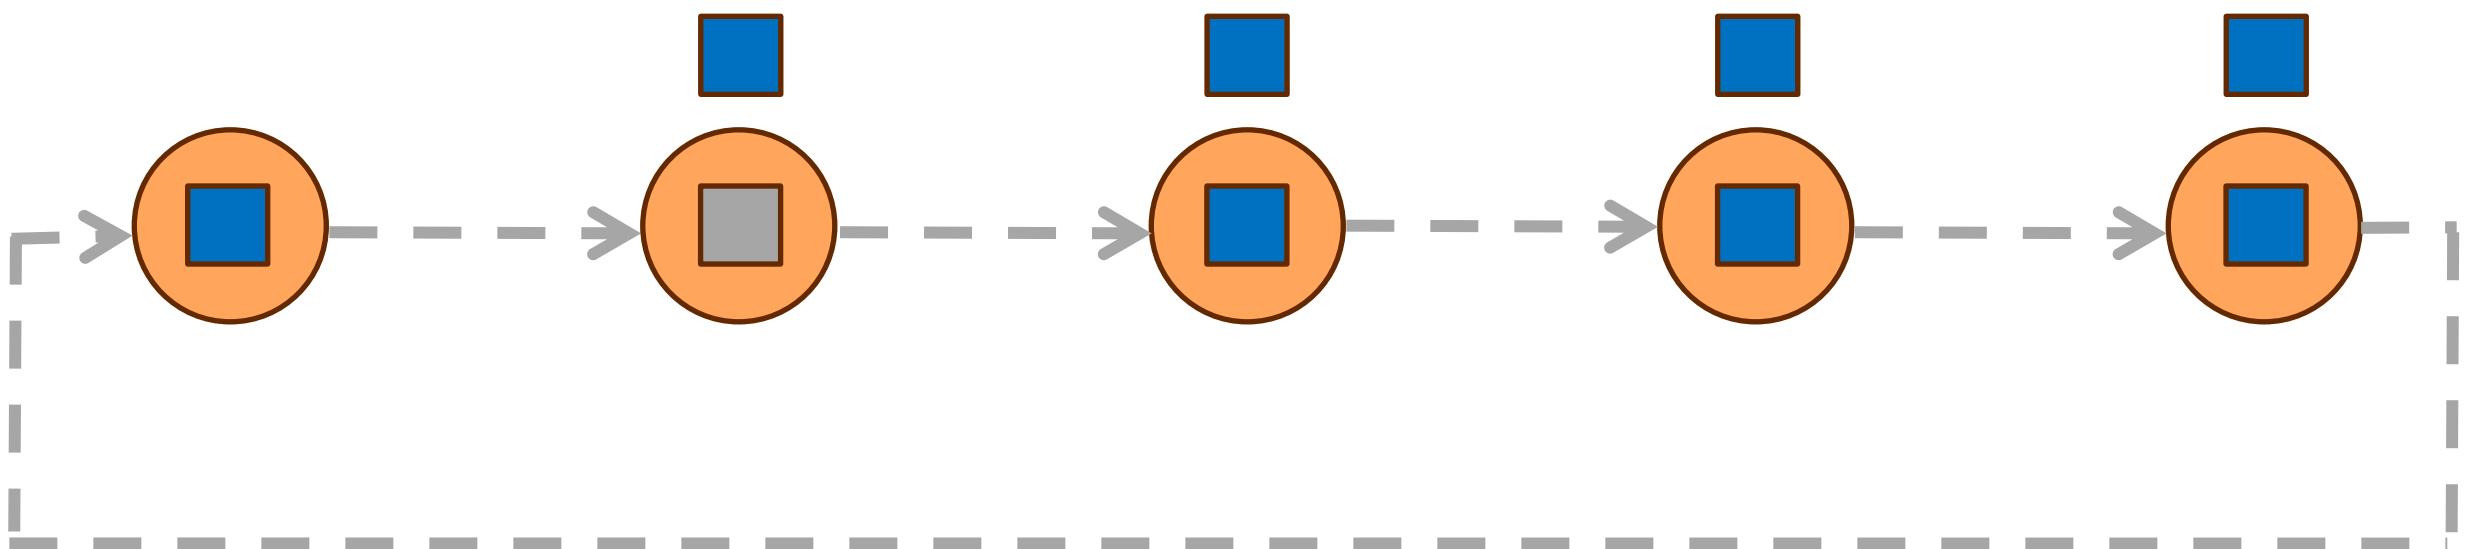


*Need to garbage collect +
execute on the first replicas*

ChainPaxos : Write Path

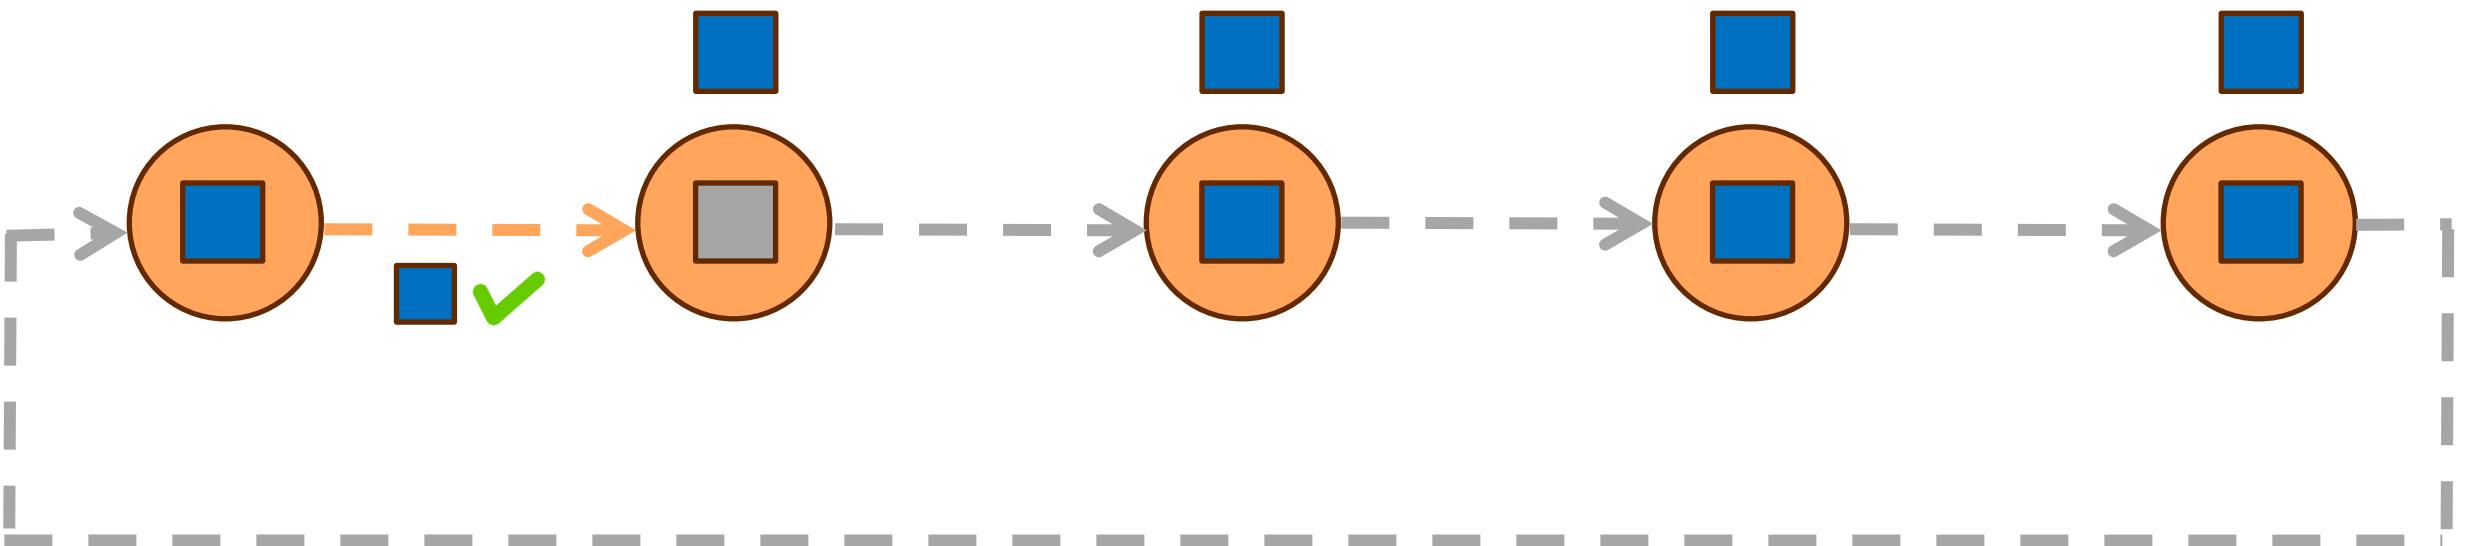


ChainPaxos : Write Path

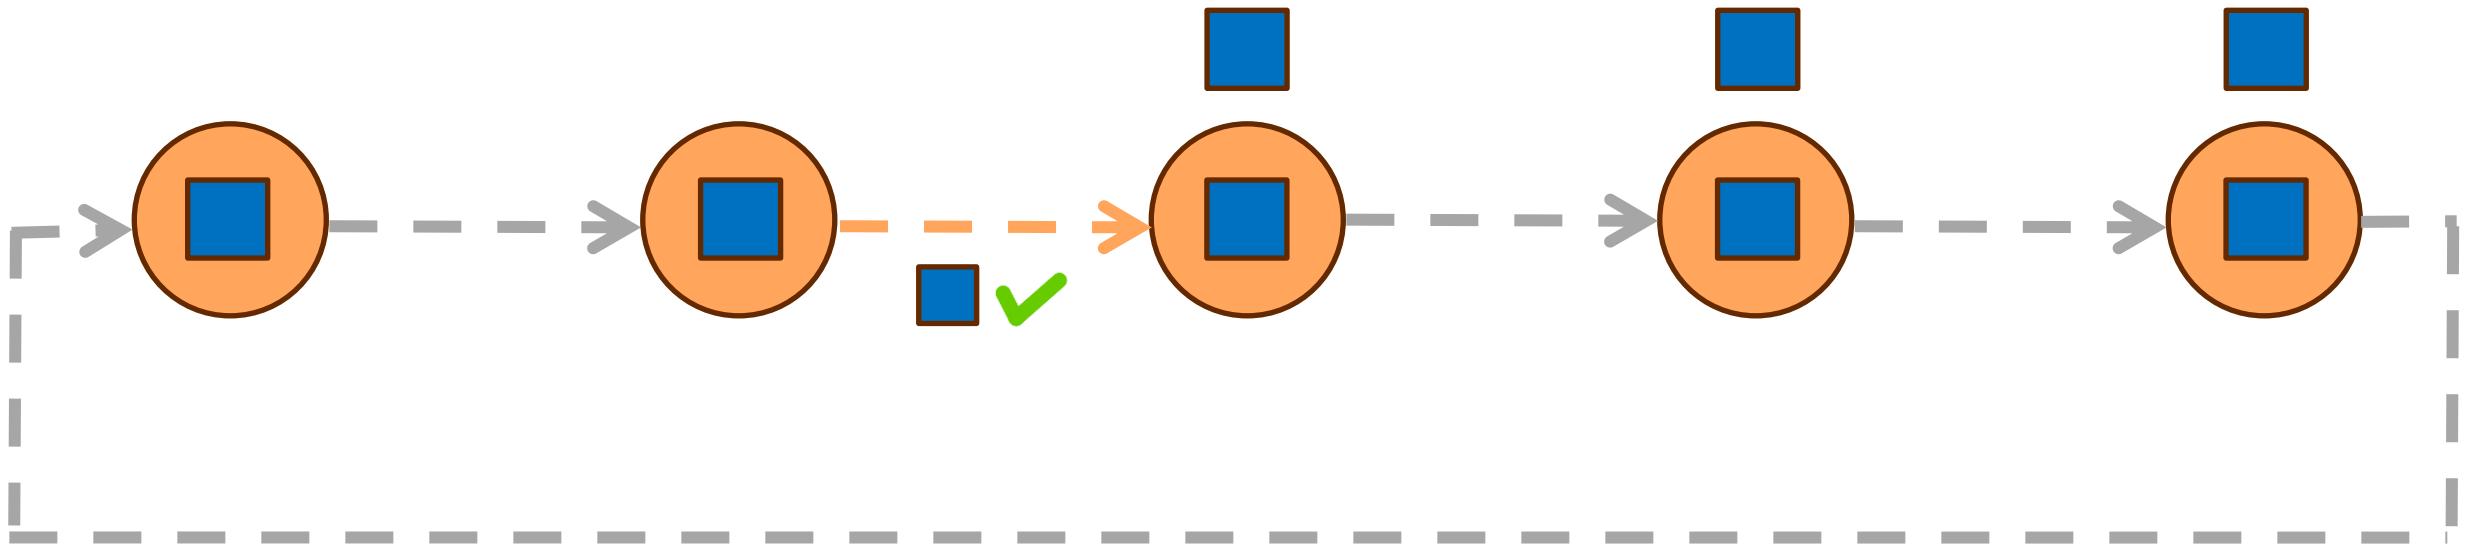


ChainPaxos : Write Path

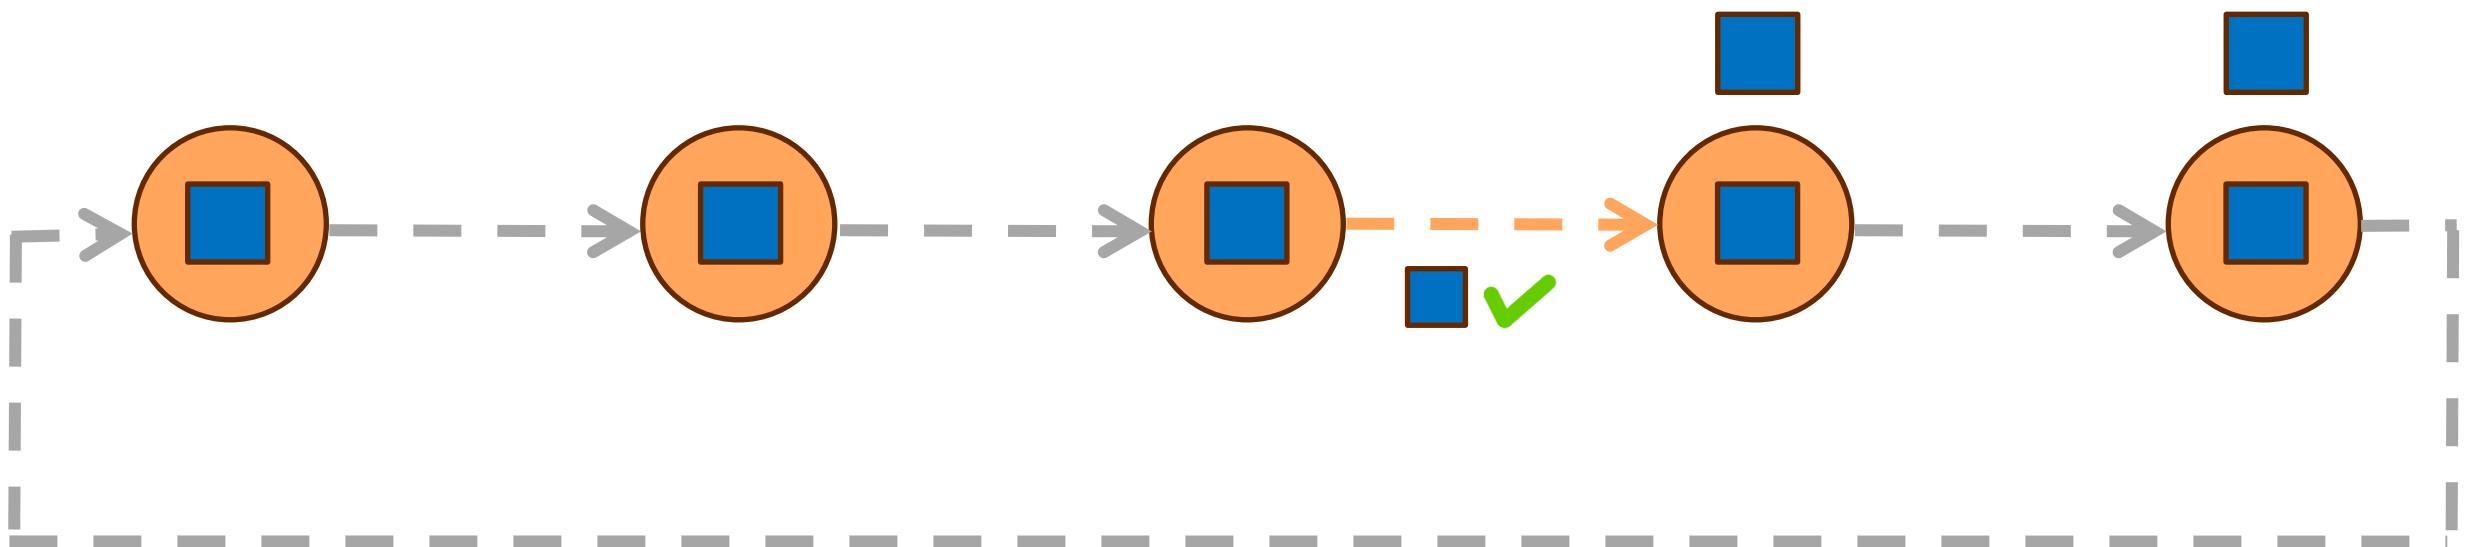
*Ack is "piggybacked" along
with next write**



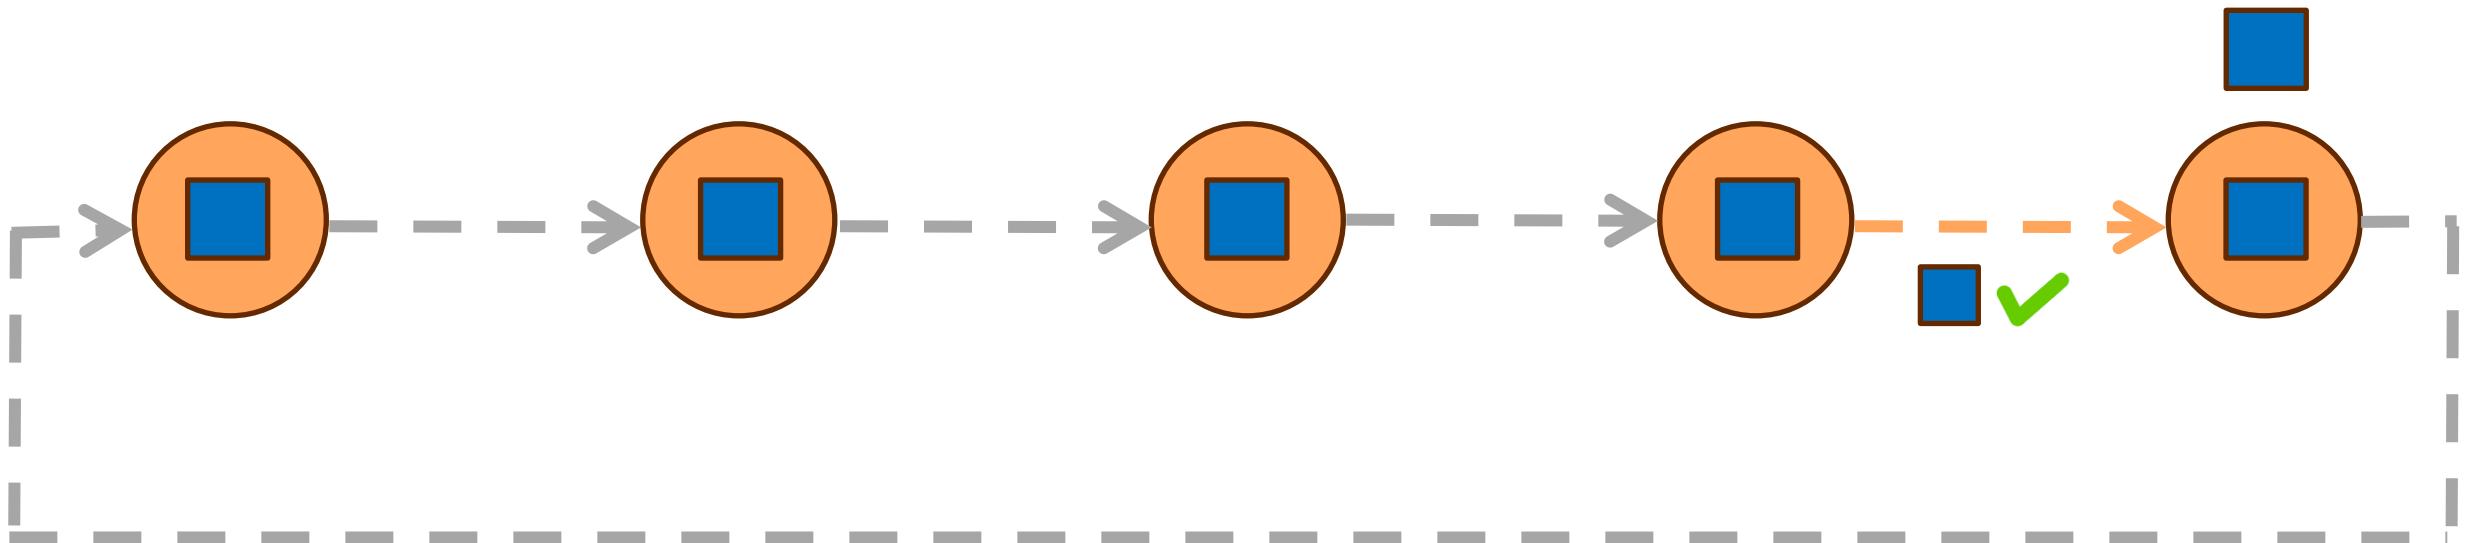
ChainPaxos : Write Path



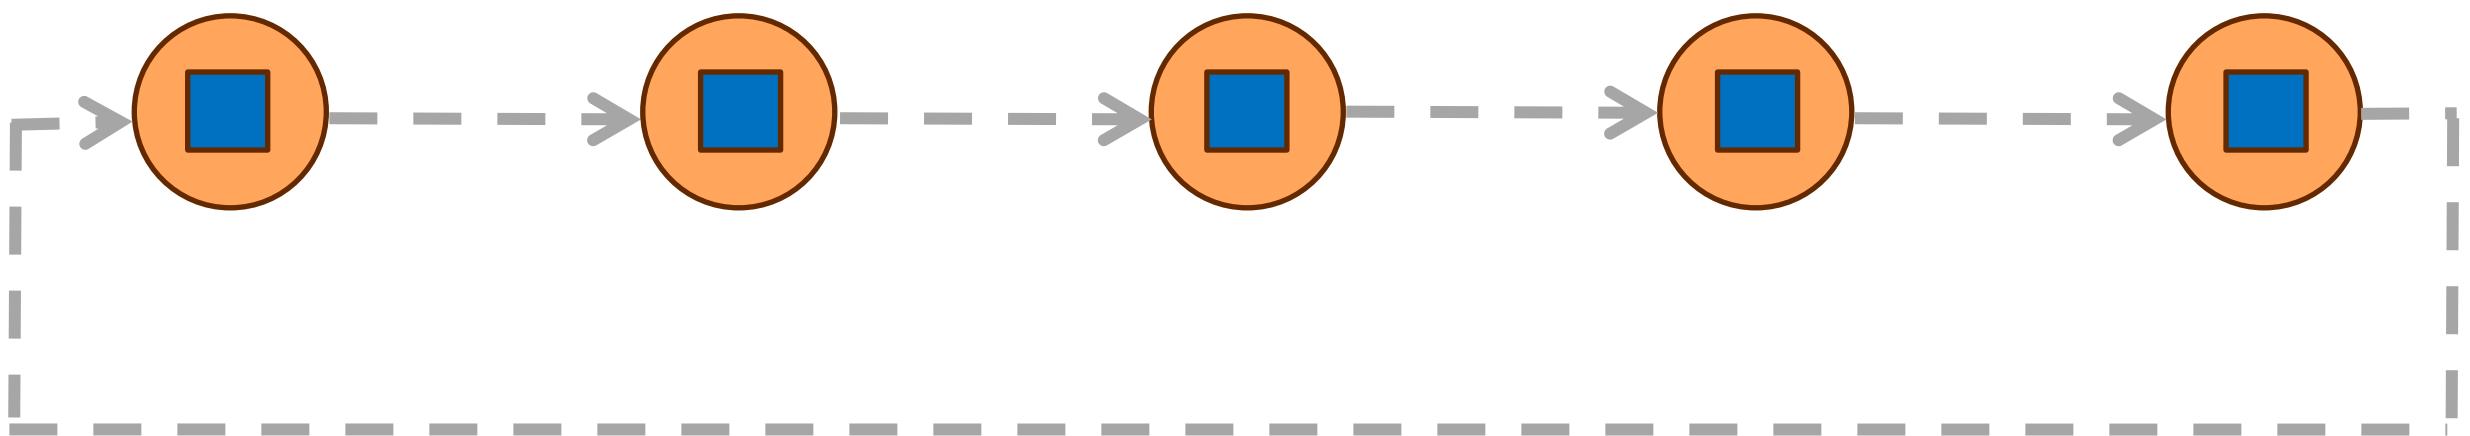
ChainPaxos : Write Path



ChainPaxos : Write Path



ChainPaxos : Write Path



Local Linearizable Reads

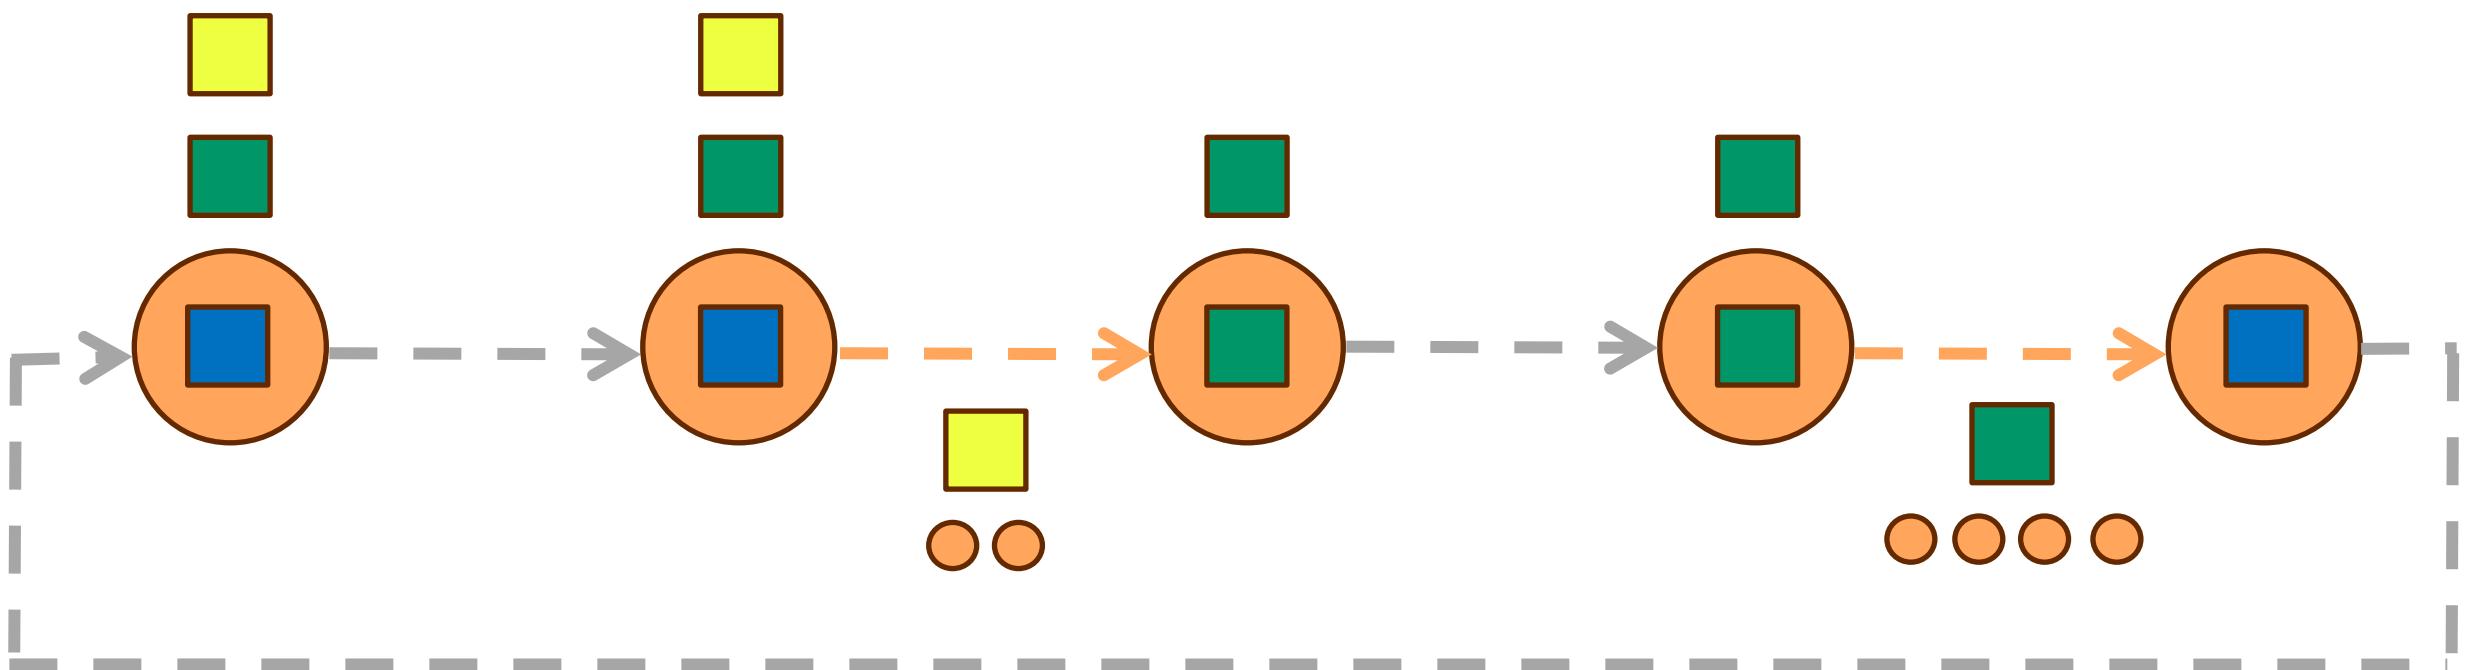
Requirements for Linearizability:

- The result of a read must contain **all writes that completed** before it started
- The result of a read must contain the result of **all reads that completed** before it started

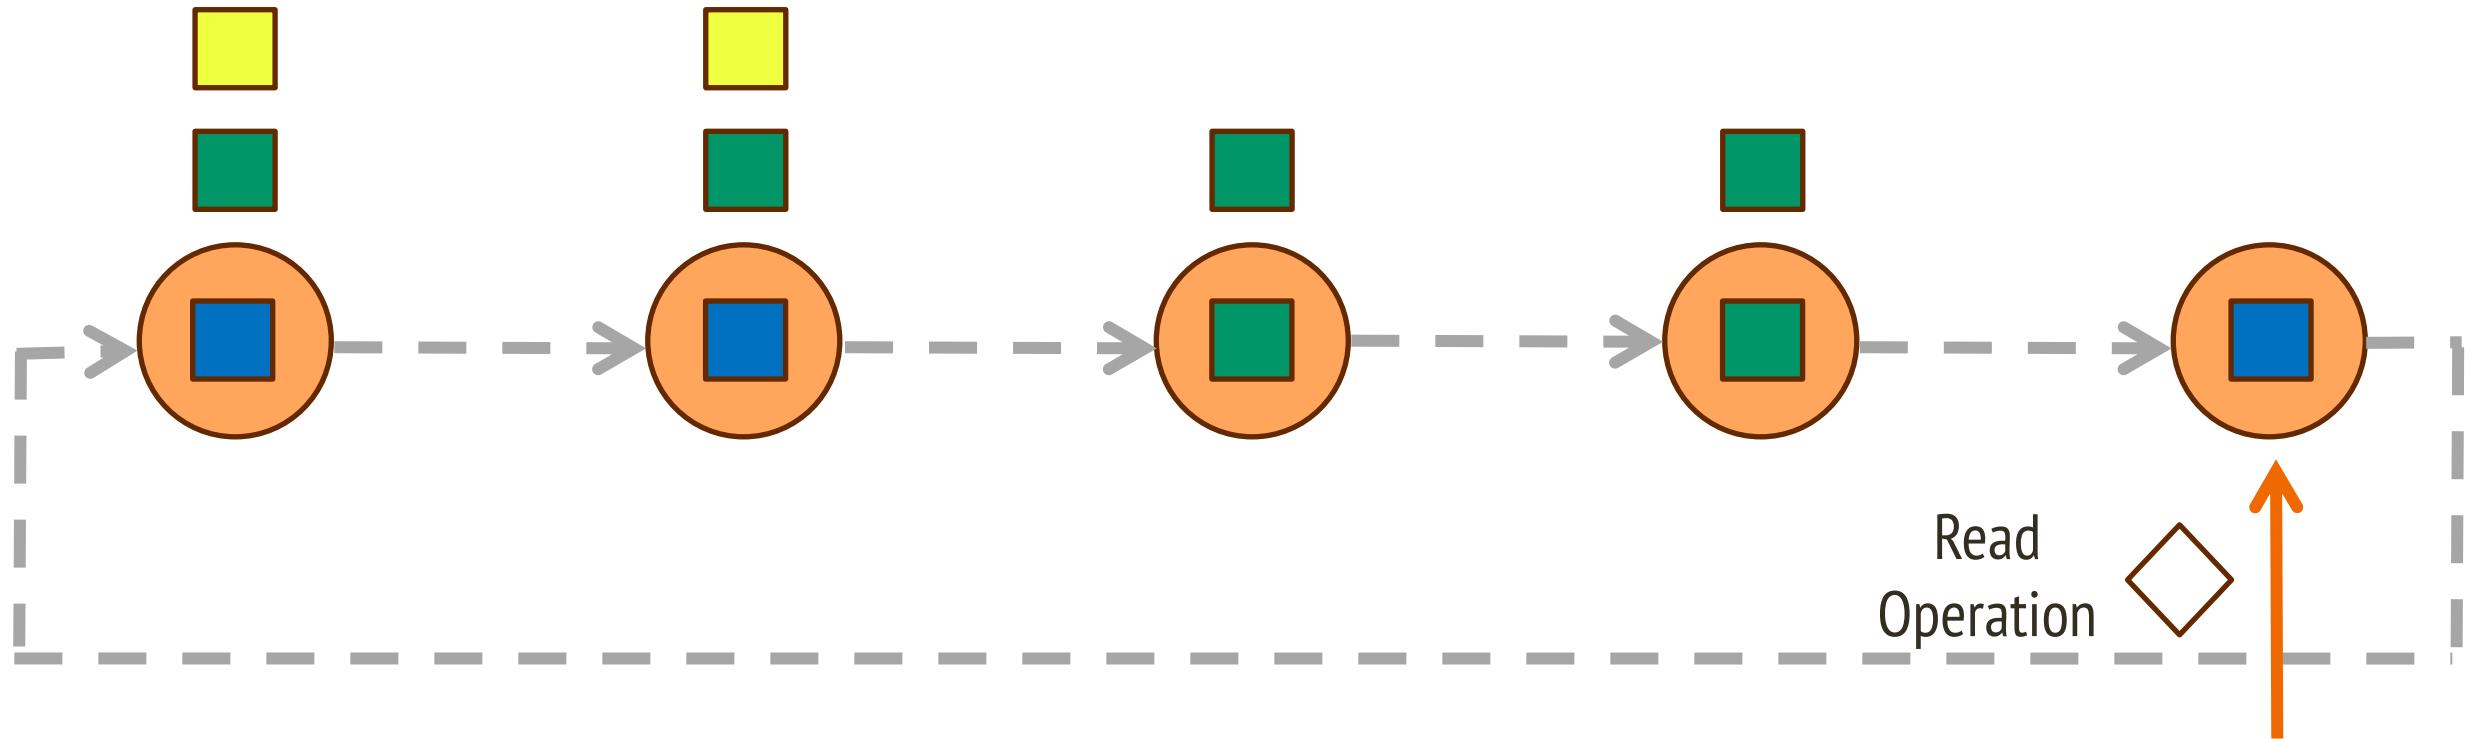
Challenge:

- Read from **any replica**
- **No extra communication steps**

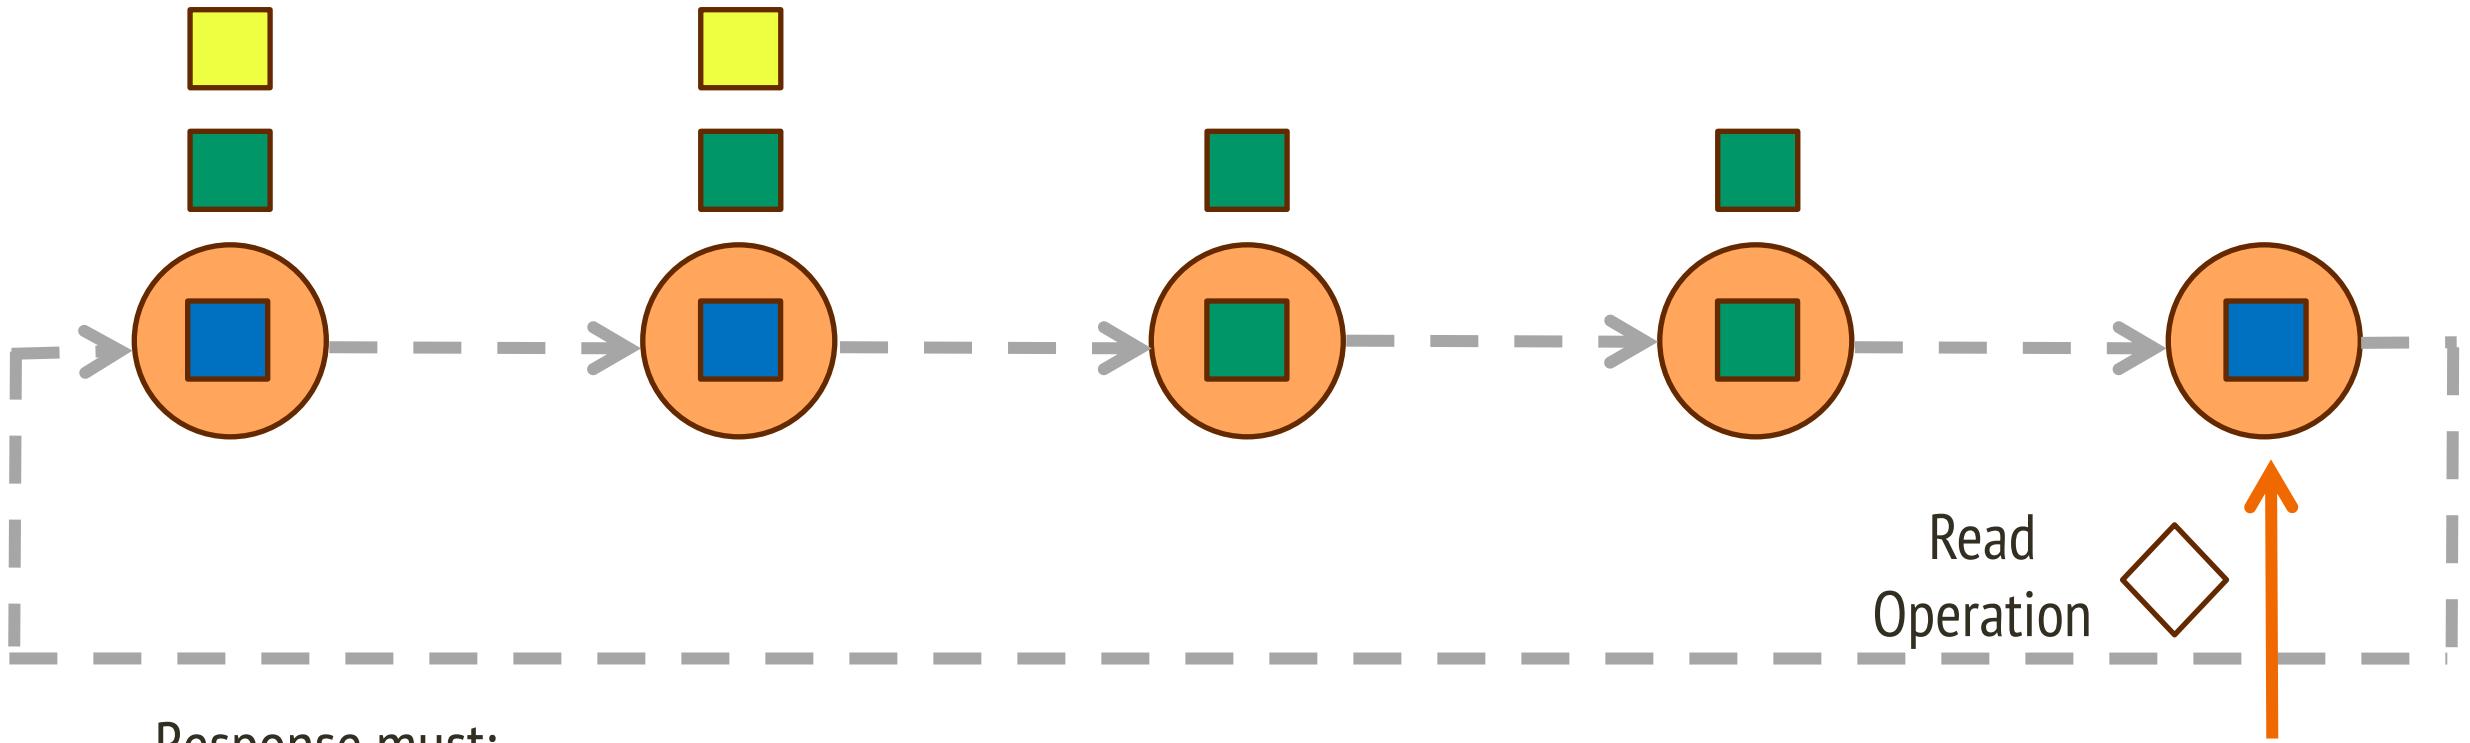
Local Linearizable Reads



Local Linearizable Reads



Local Linearizable Reads

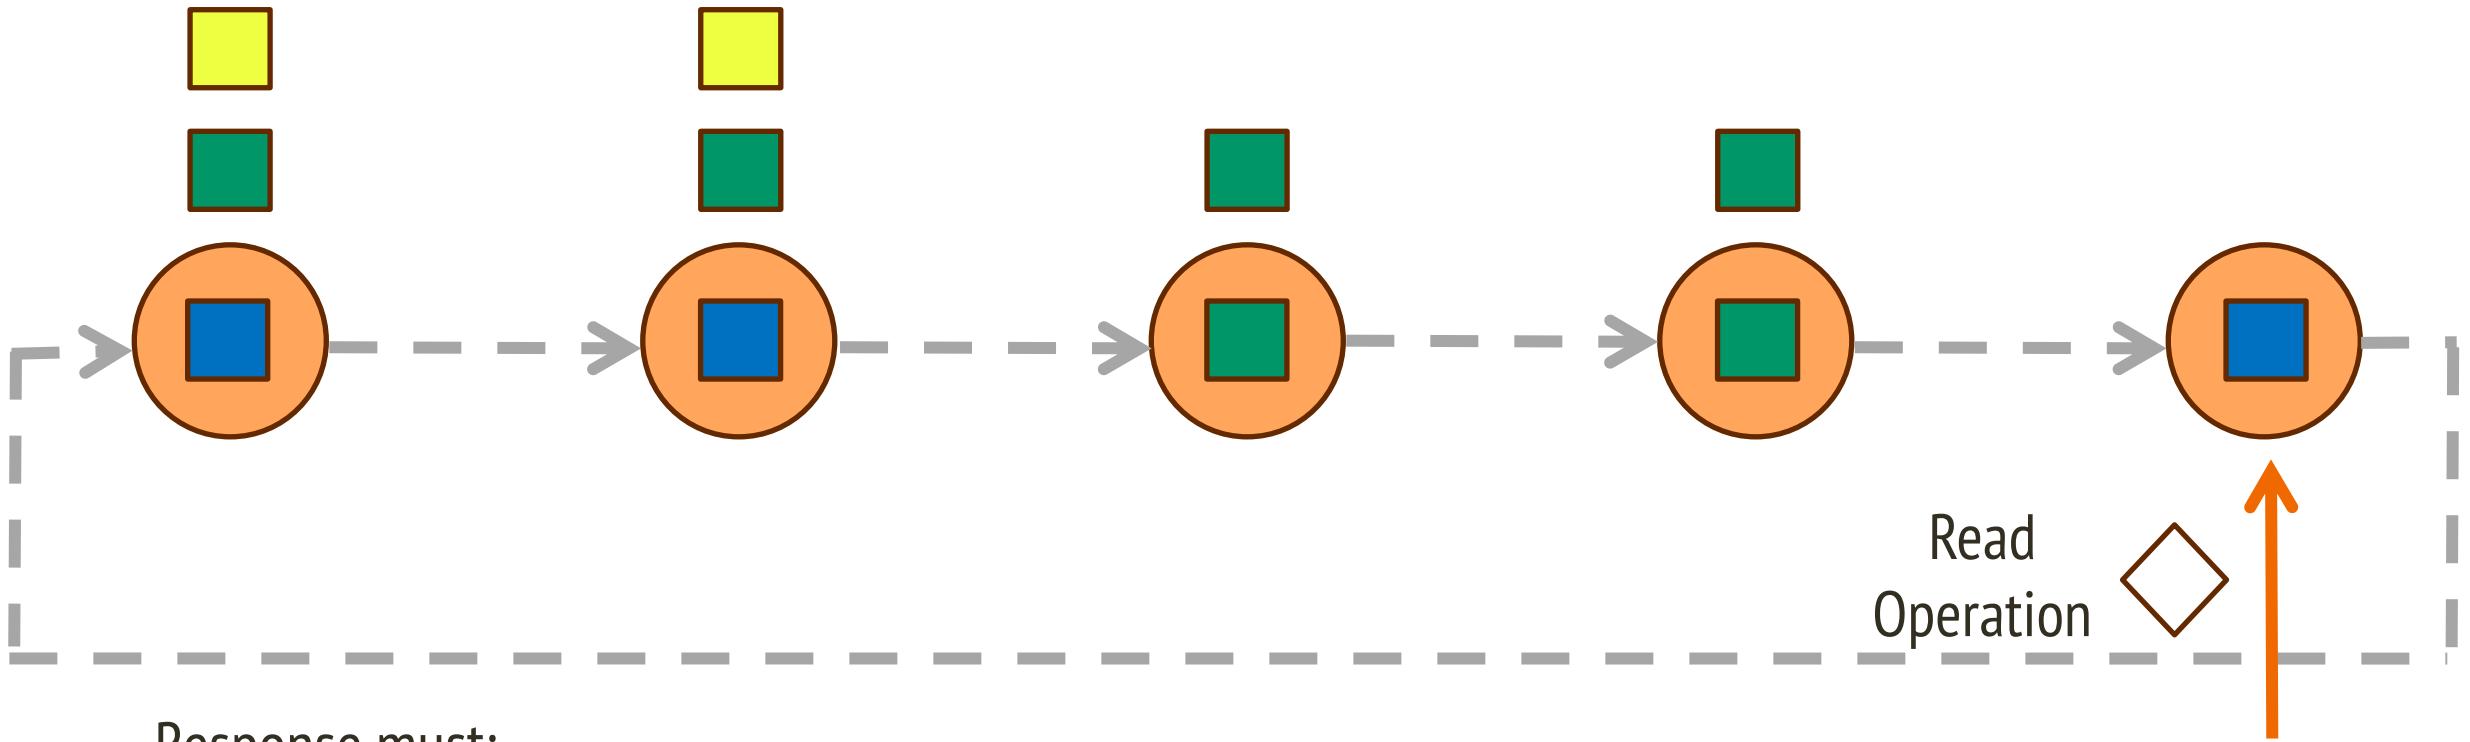


Response must:

- Contain all completed writes
- Contain all completed reads

Read
Operation

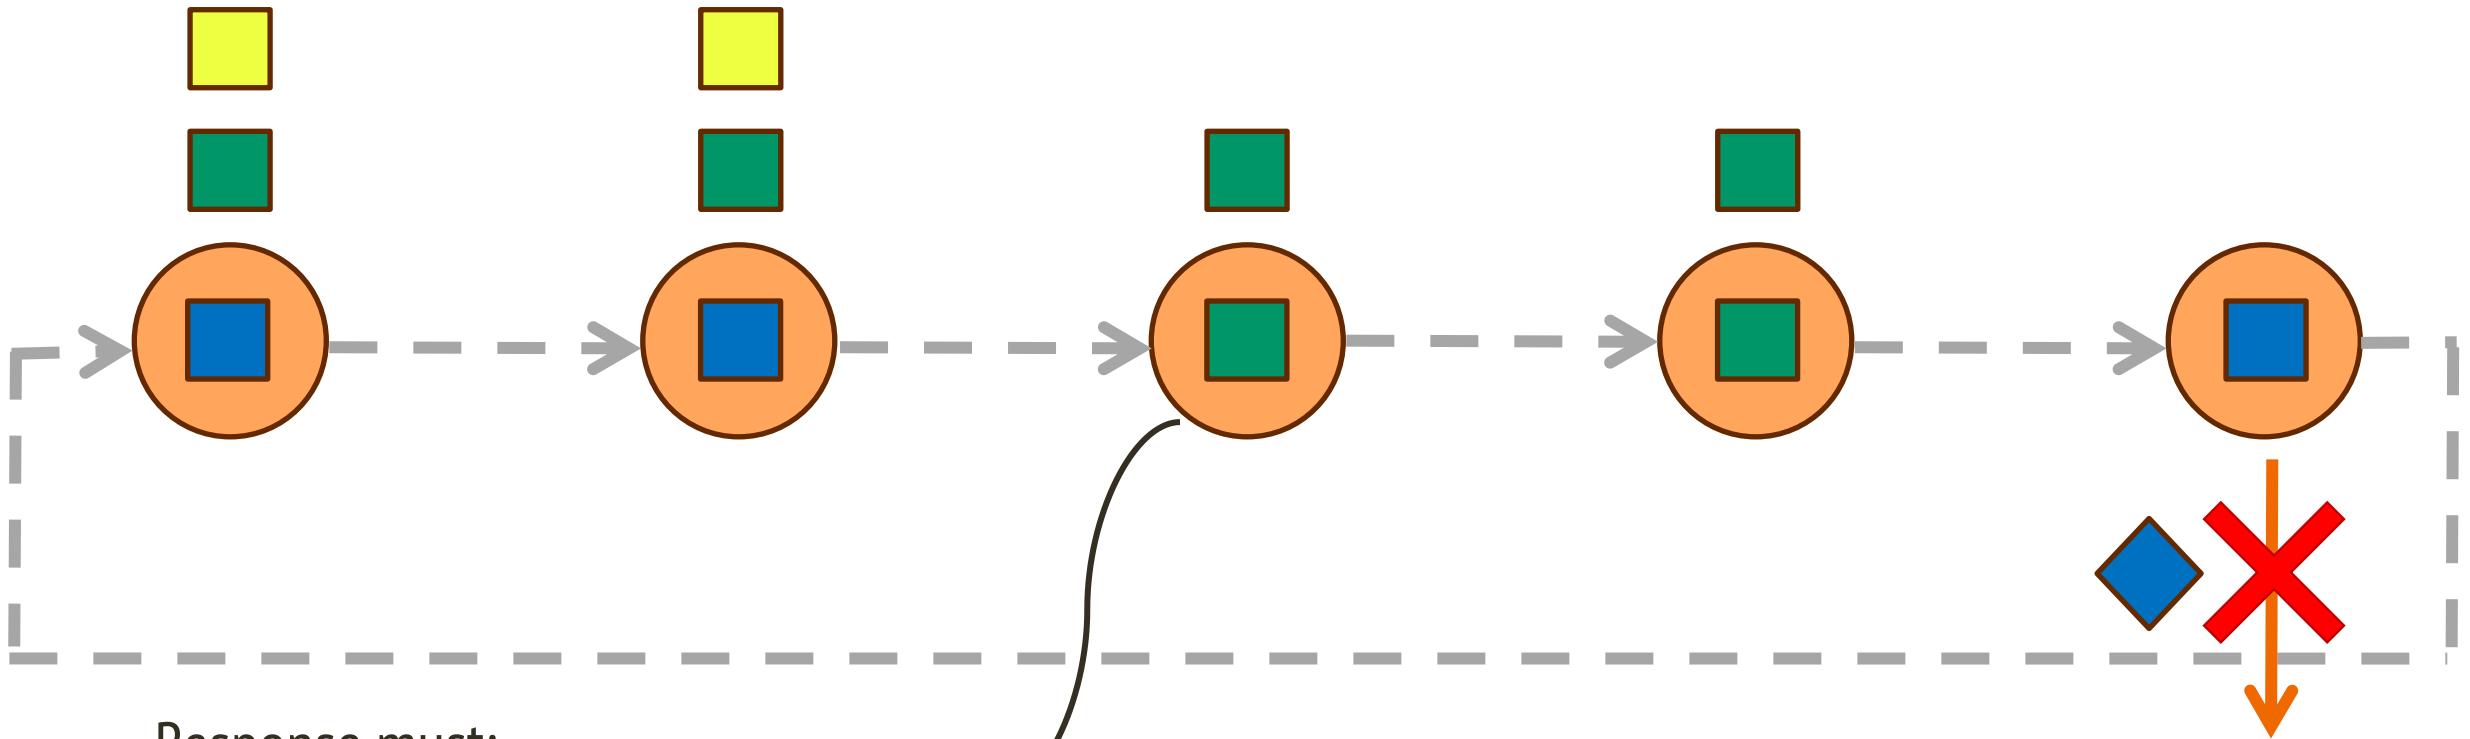
Local Linearizable Reads



Response must:

- Contain all completed writes
- Contain all completed reads

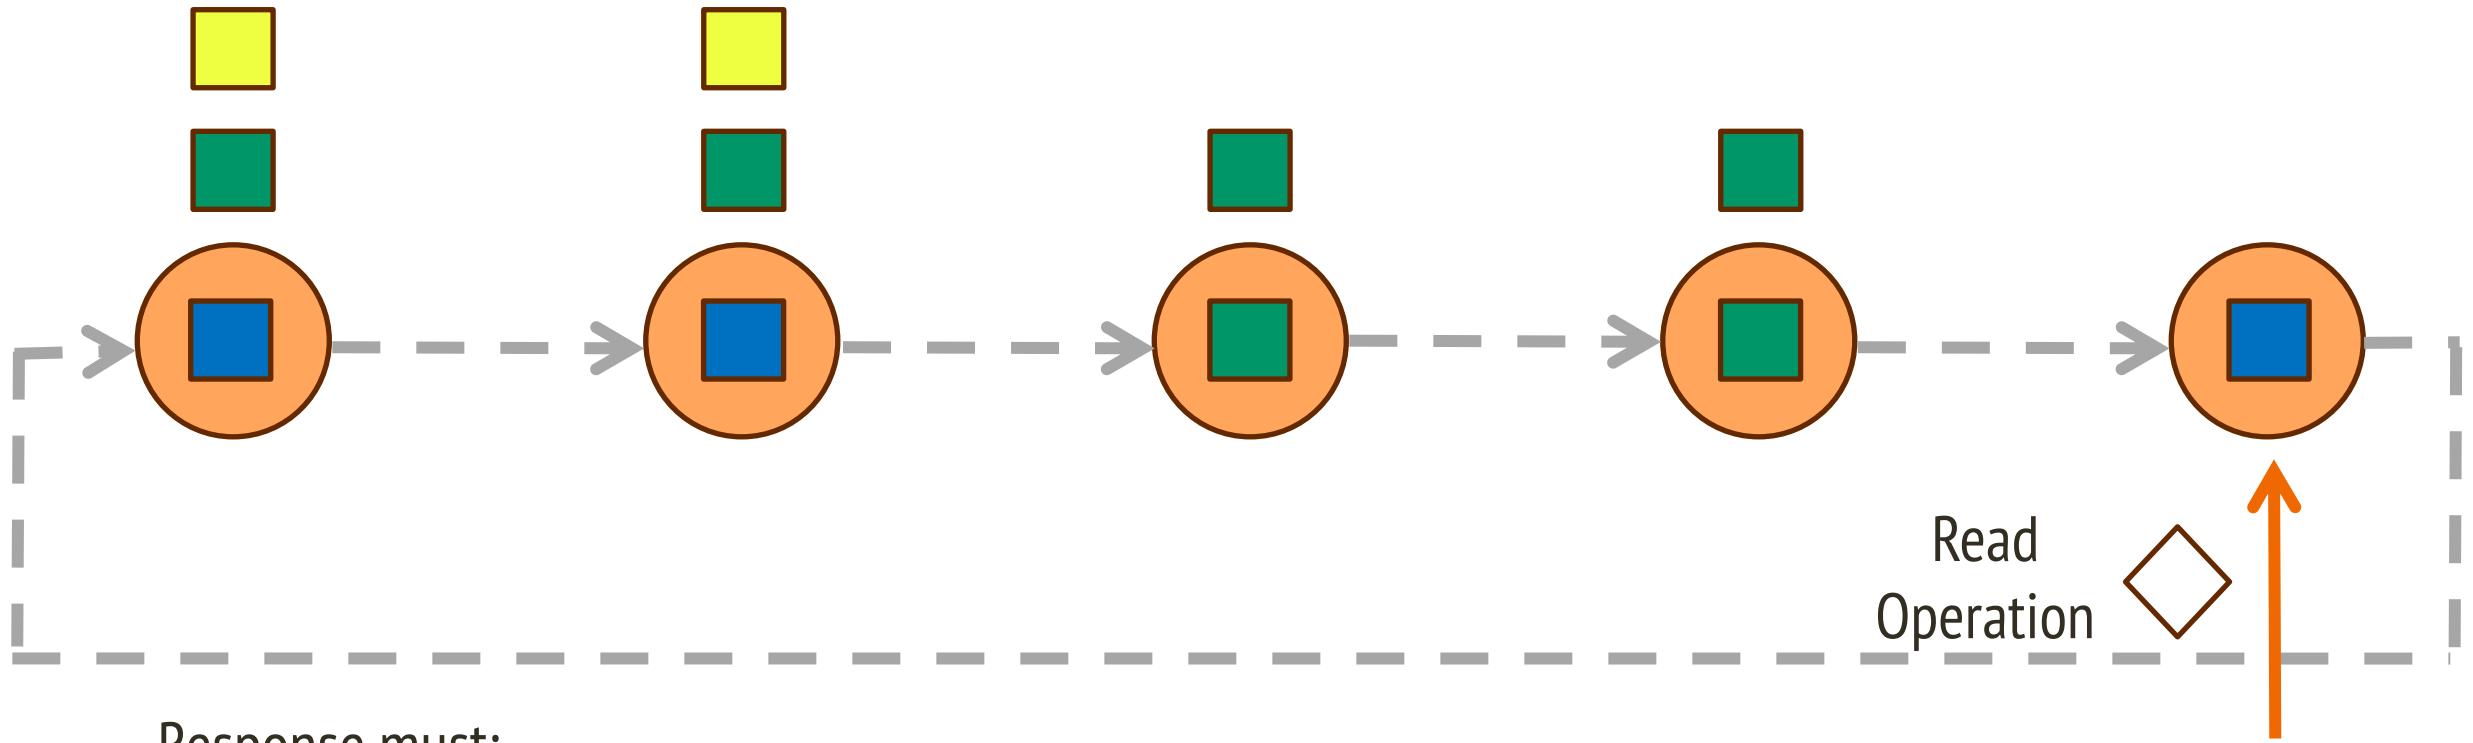
Local Linearizable Reads



Response must:

- Contain all completed writes ()
- Contain all completed reads

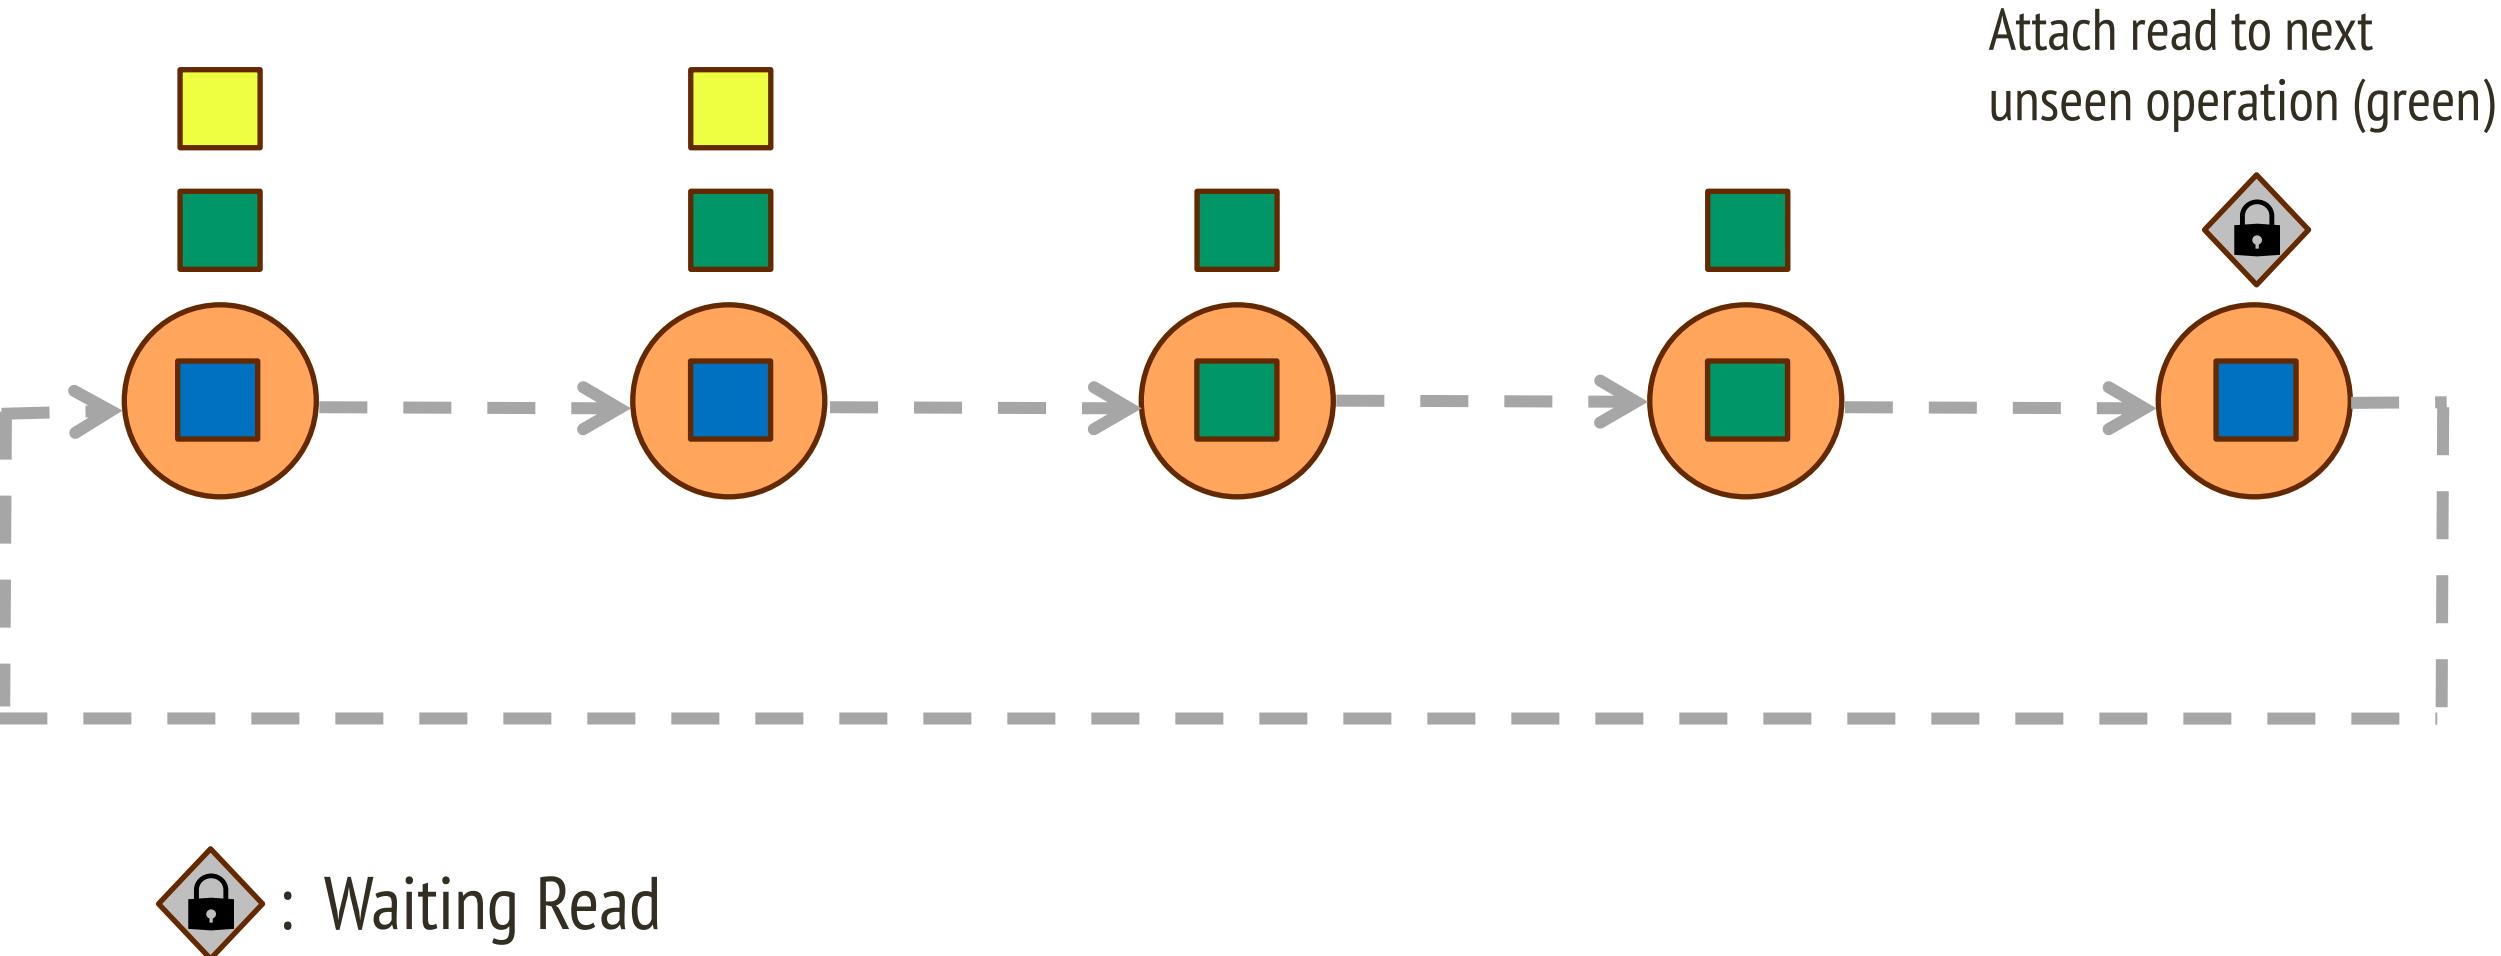
Local Linearizable Reads



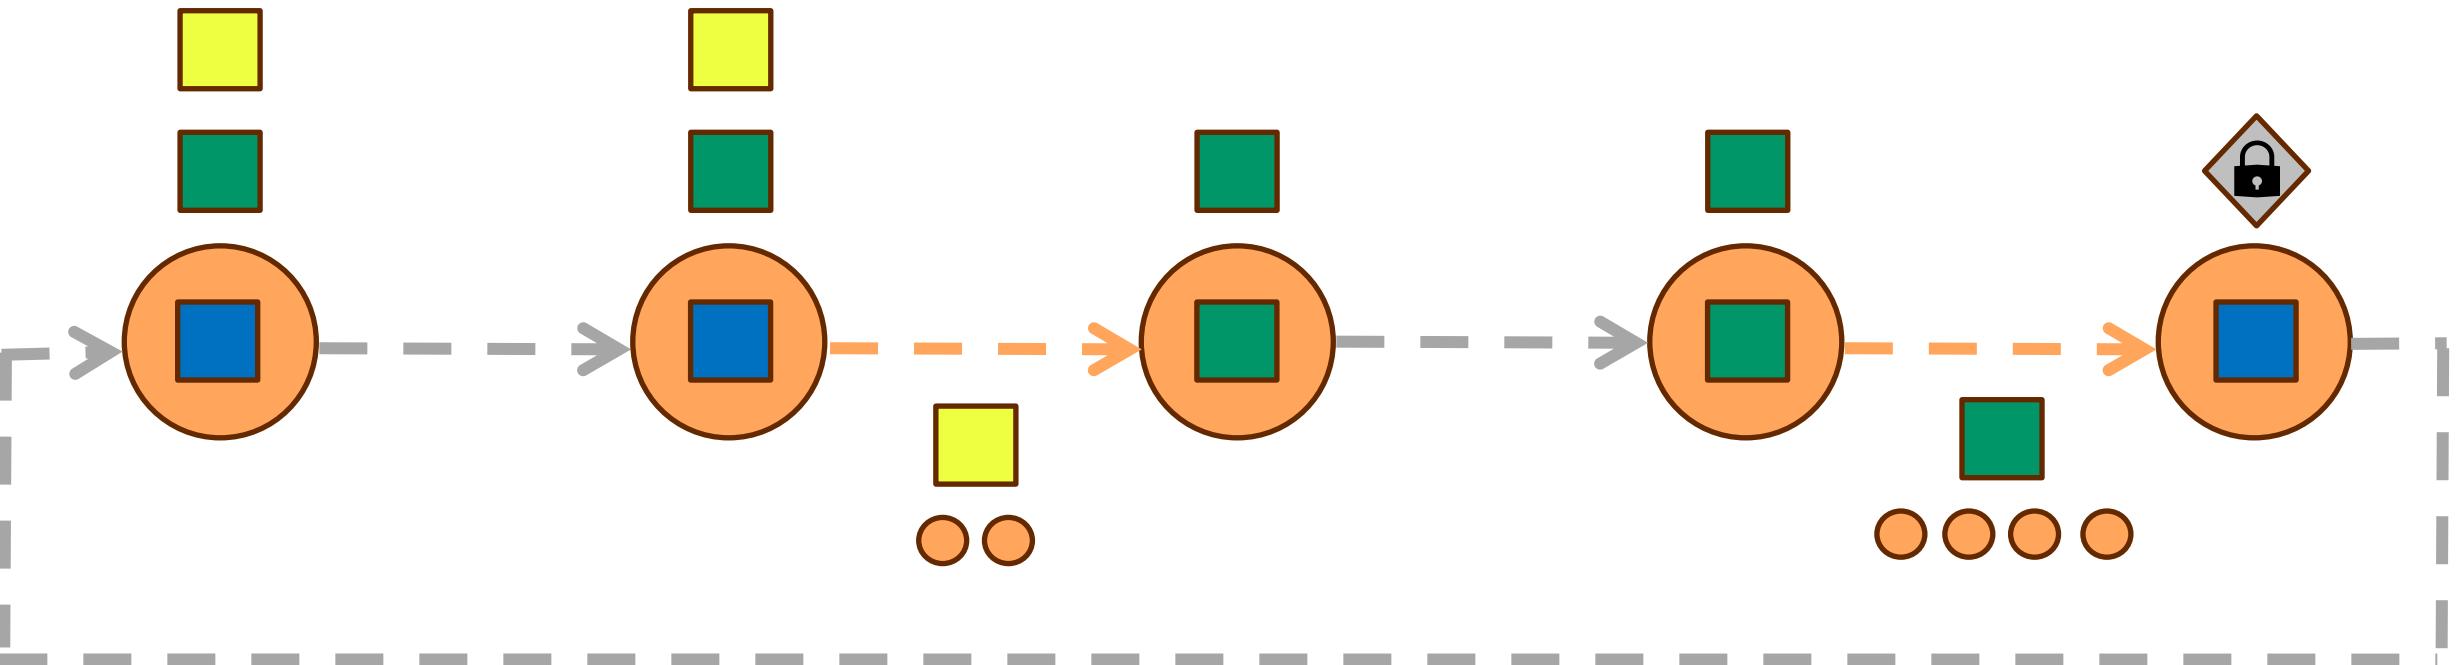
Response must:

- Contain all completed writes (█)
- Contain all completed reads

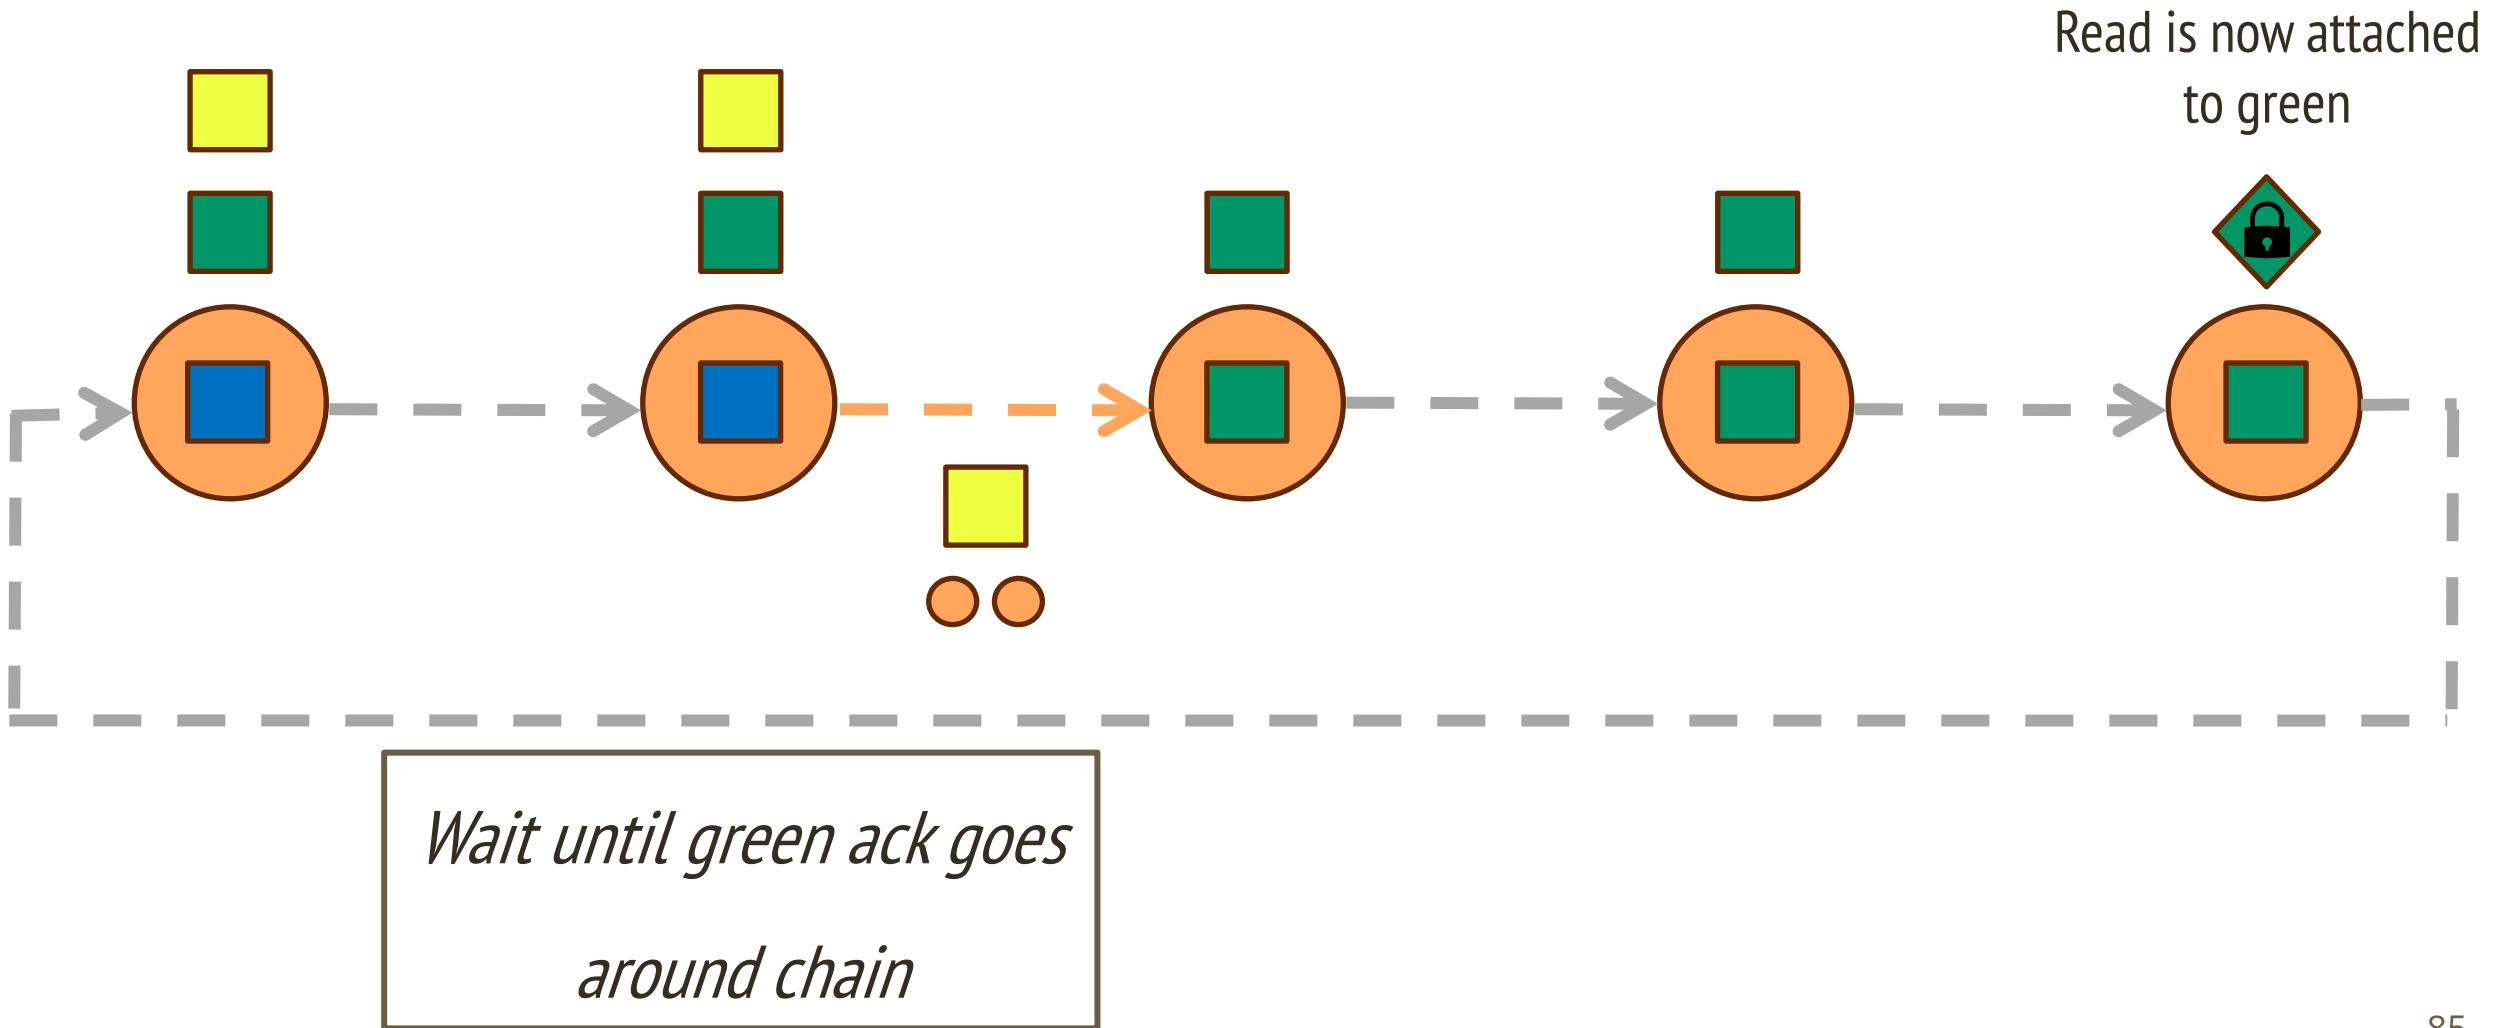
Local Linearizable Reads



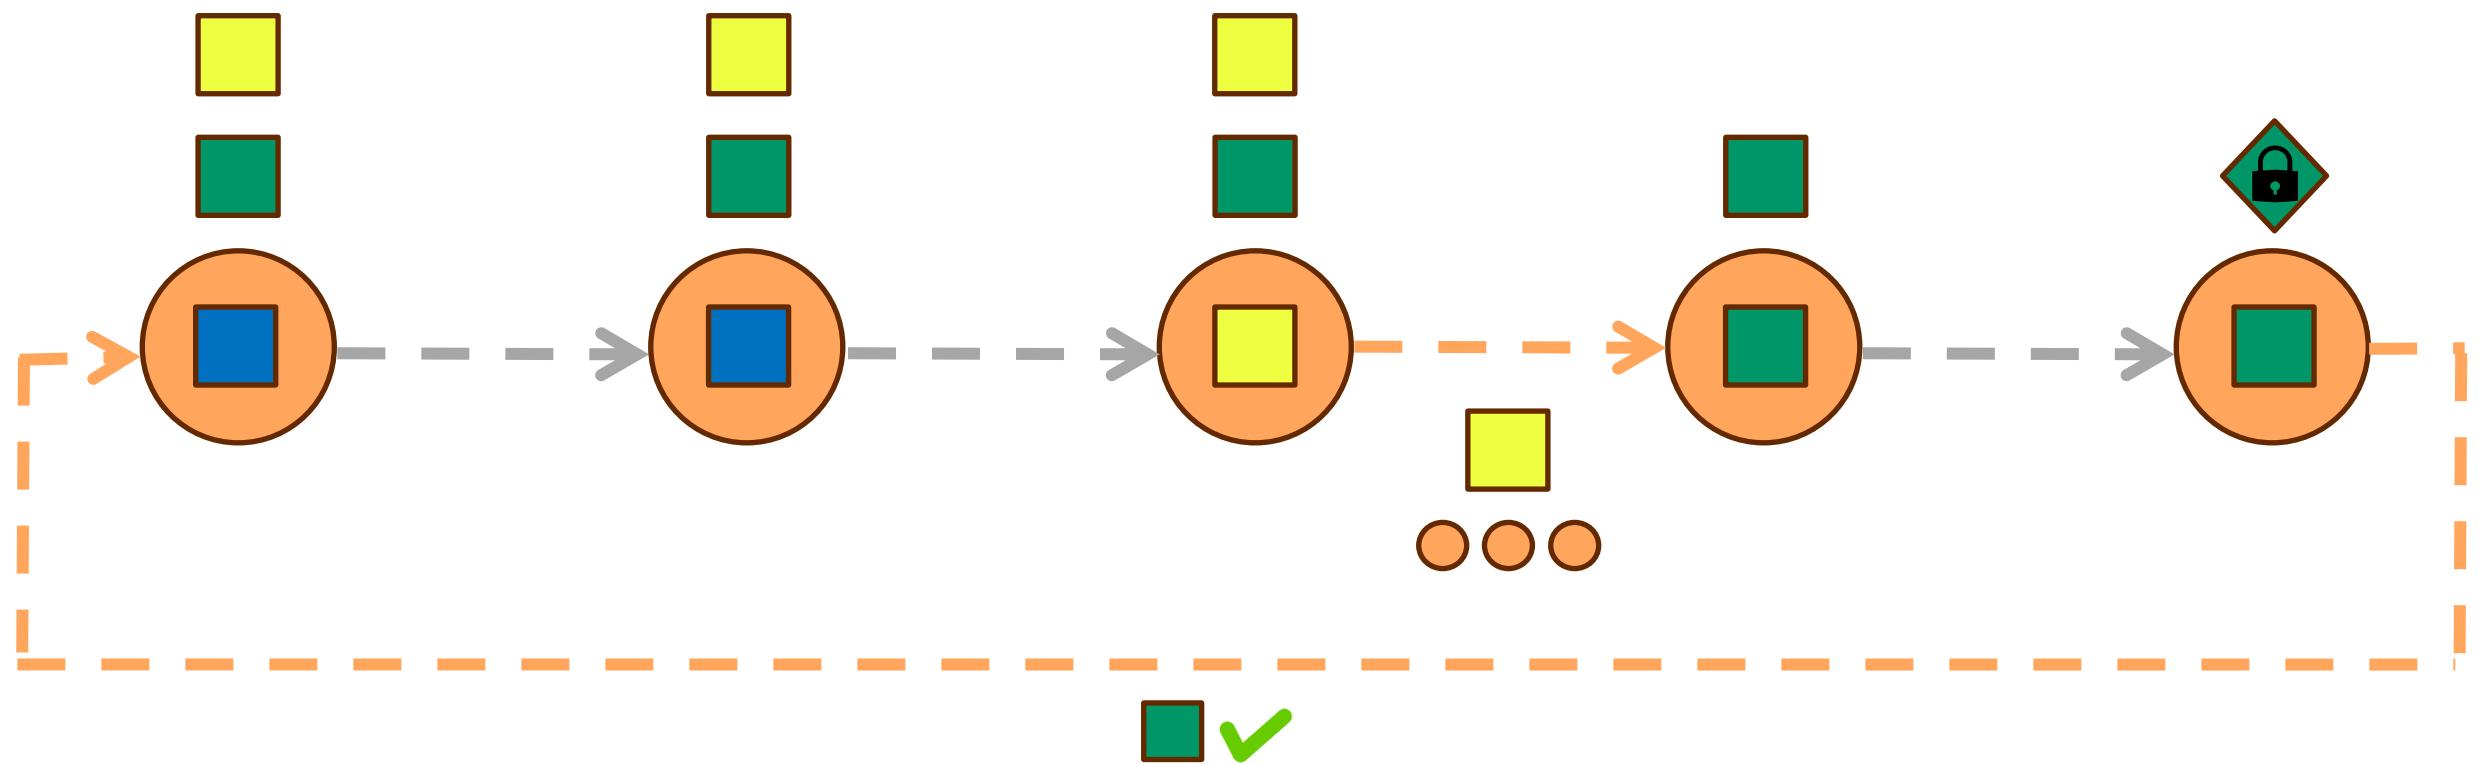
Local Linearizable Reads



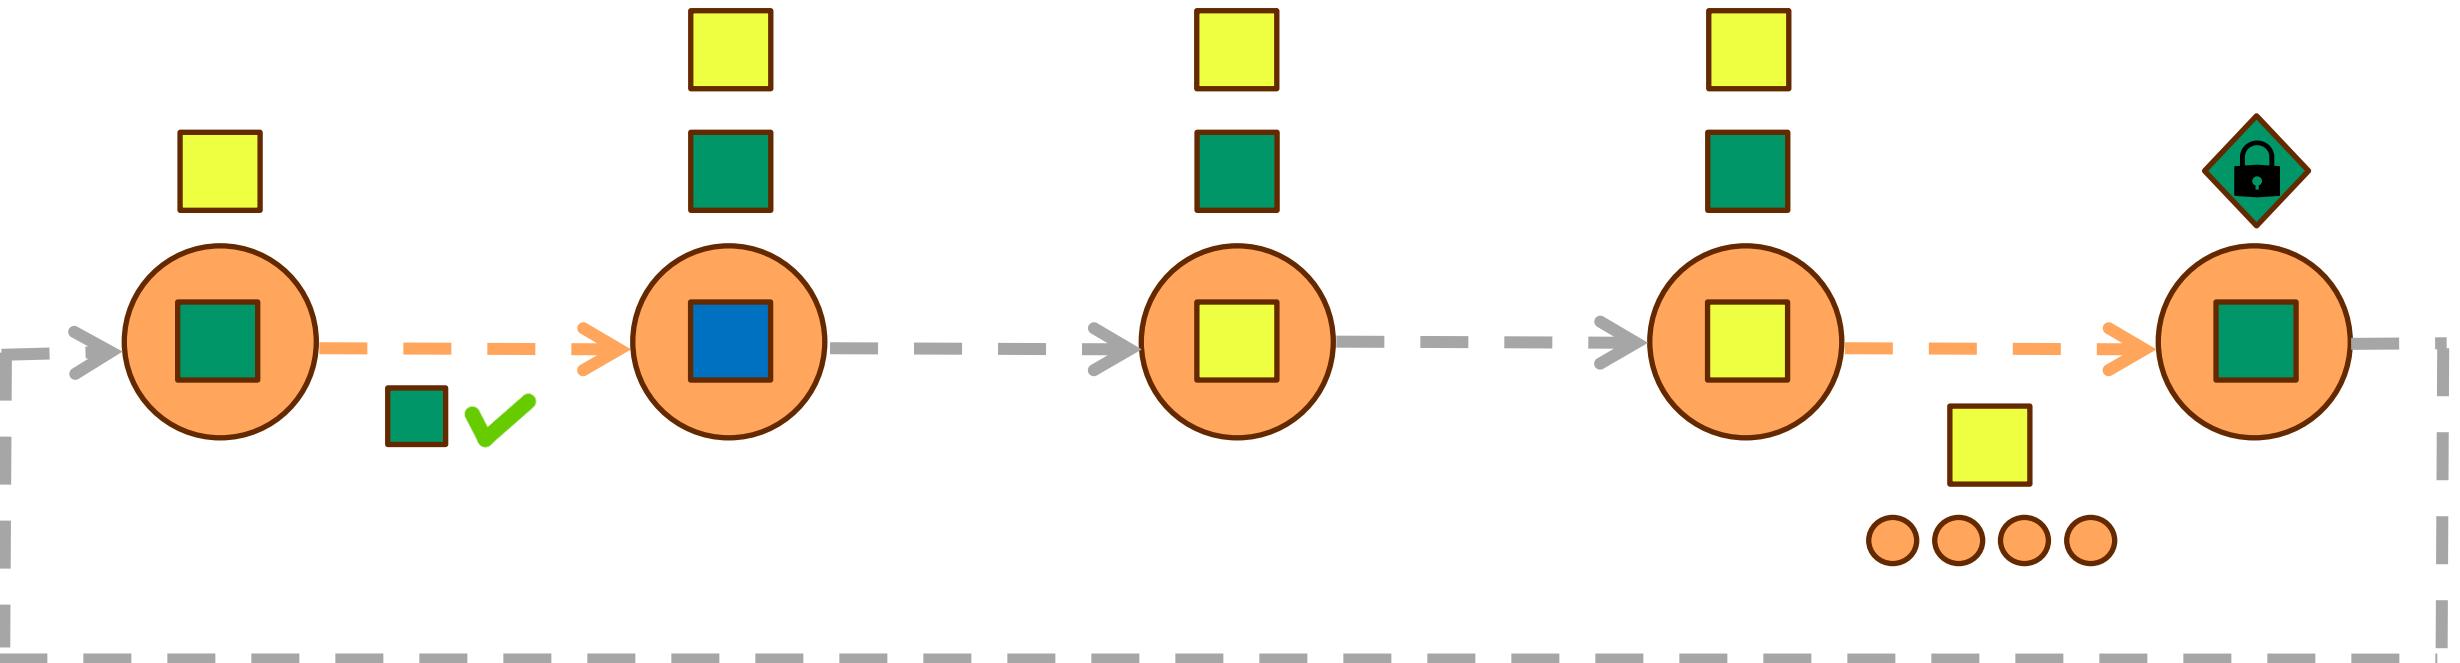
Local Linearizable Reads



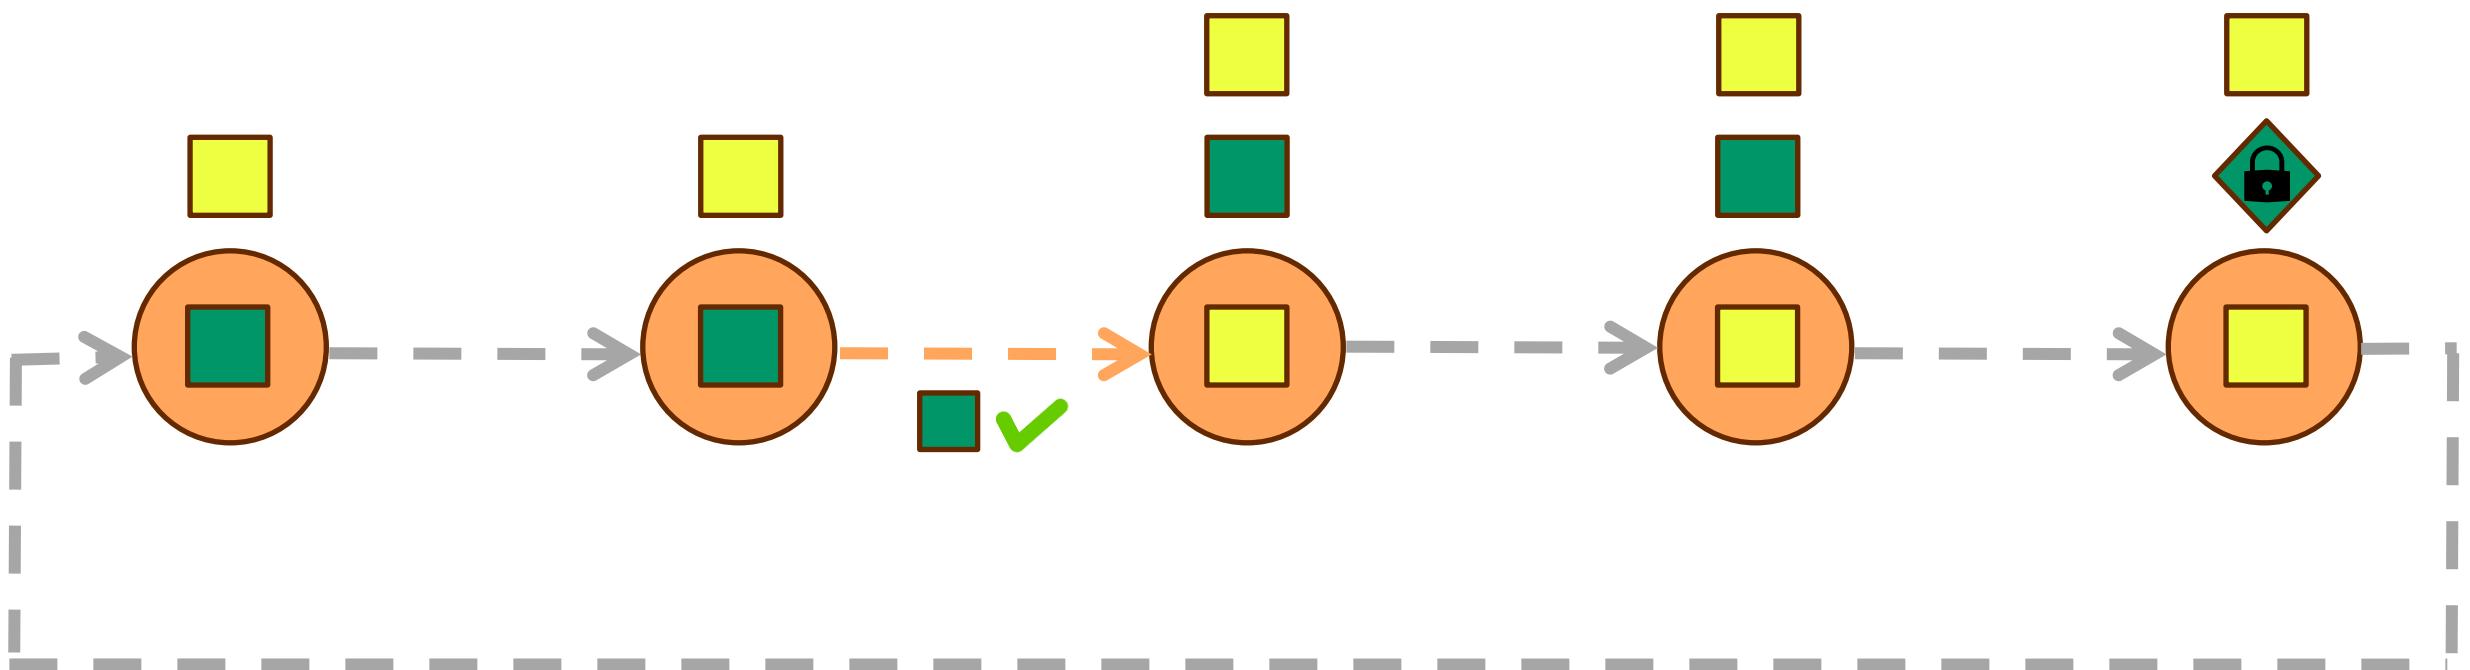
Local Linearizable Reads



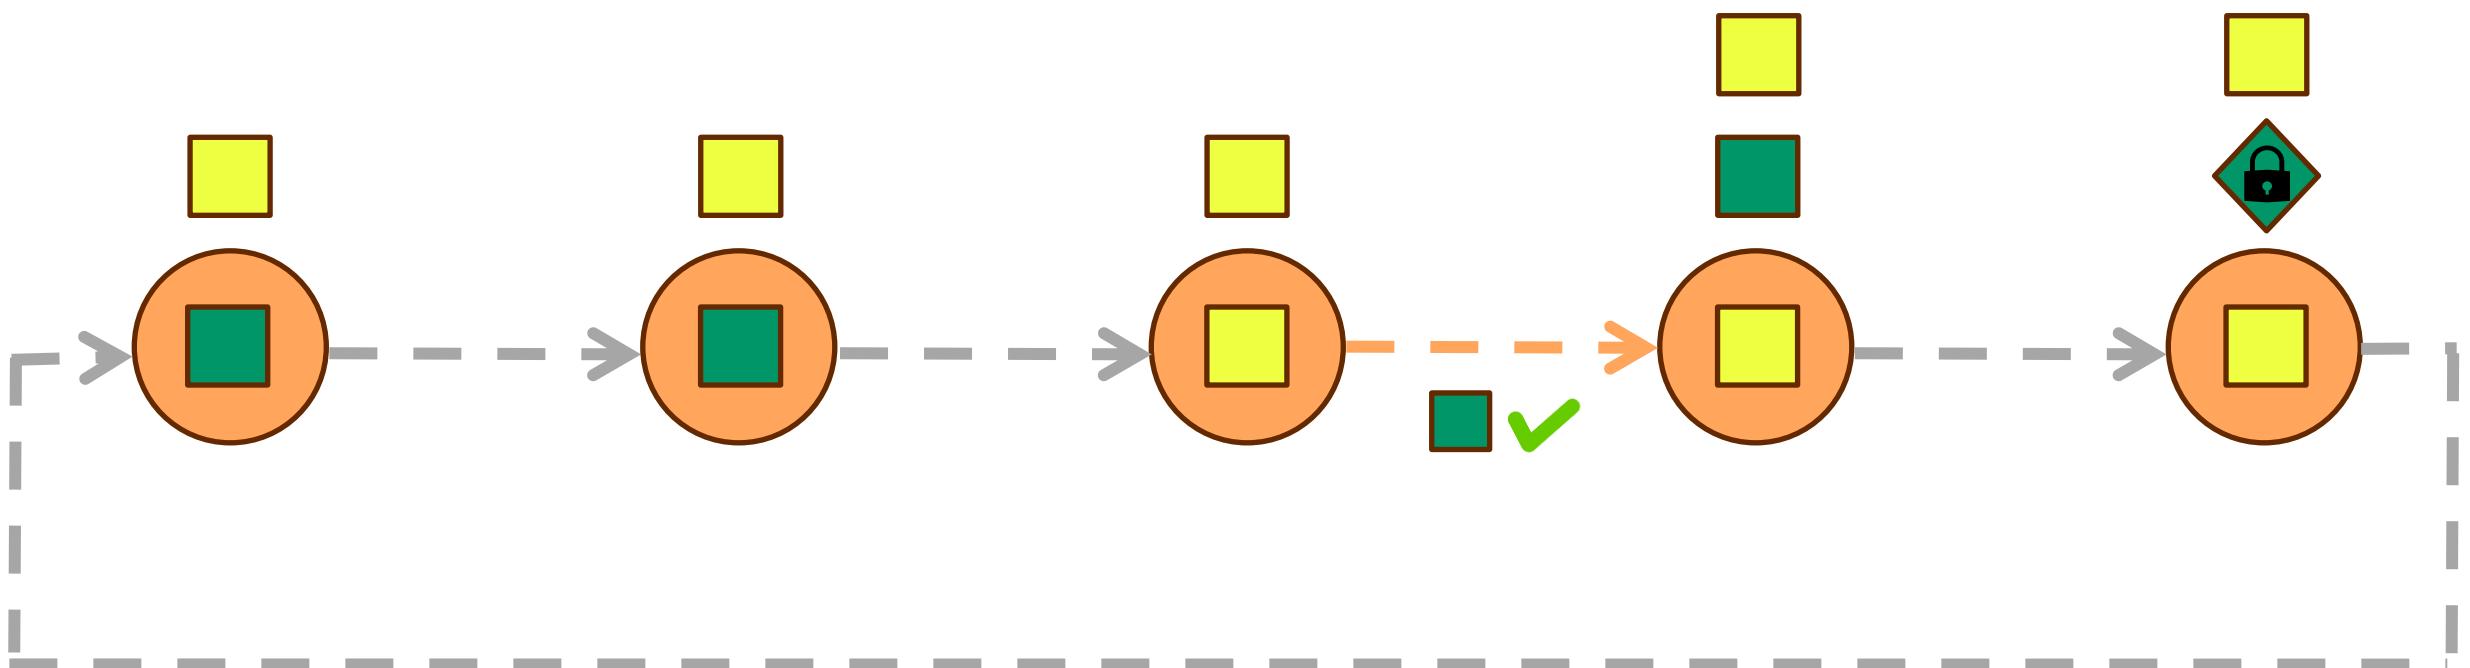
Local Linearizable Reads



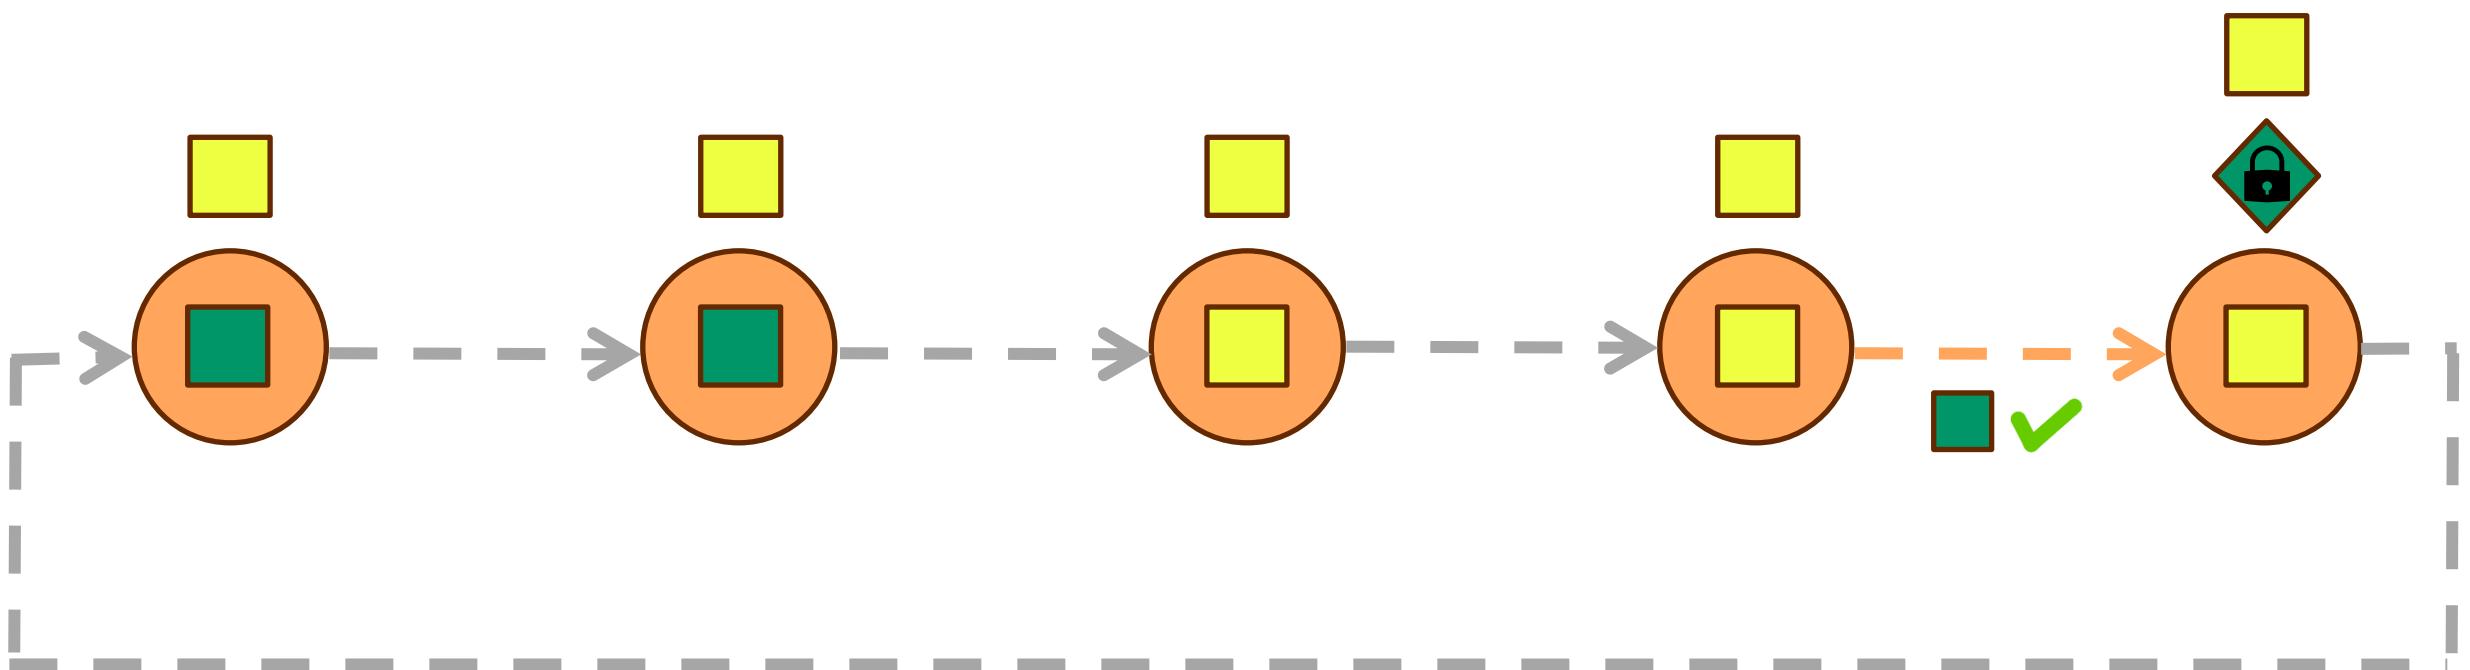
Local Linearizable Reads



Local Linearizable Reads

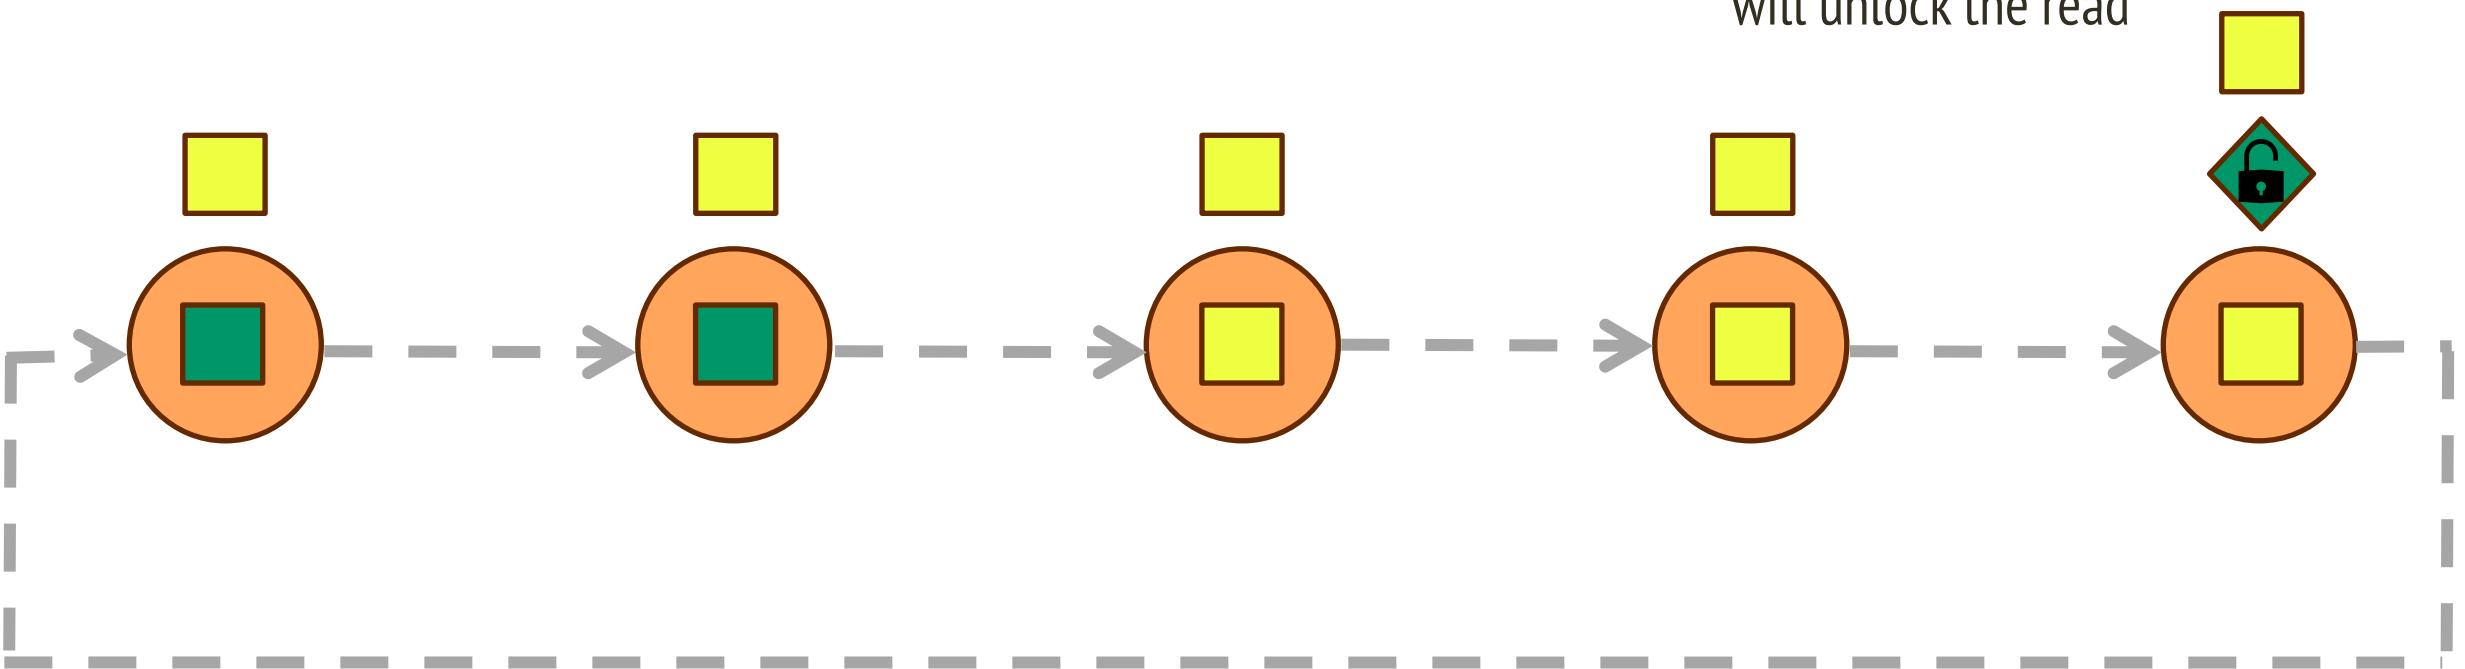


Local Linearizable Reads

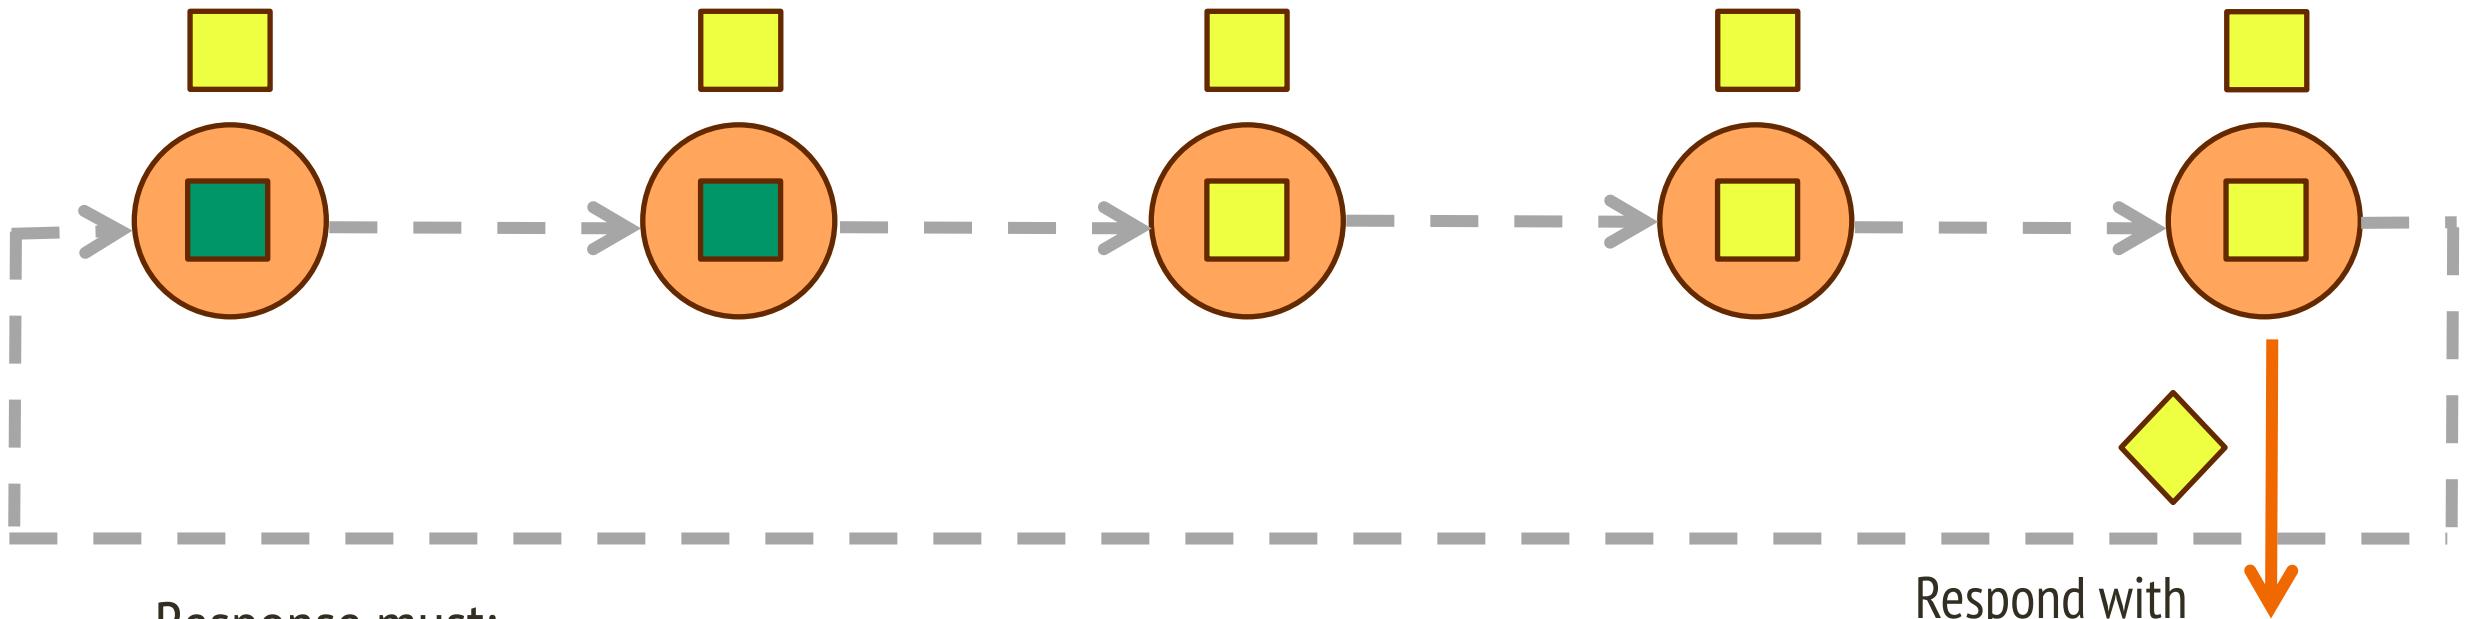


Local Linearizable Reads

Acknowledge of green
will unlock the read



Local Linearizable Reads



Response must:

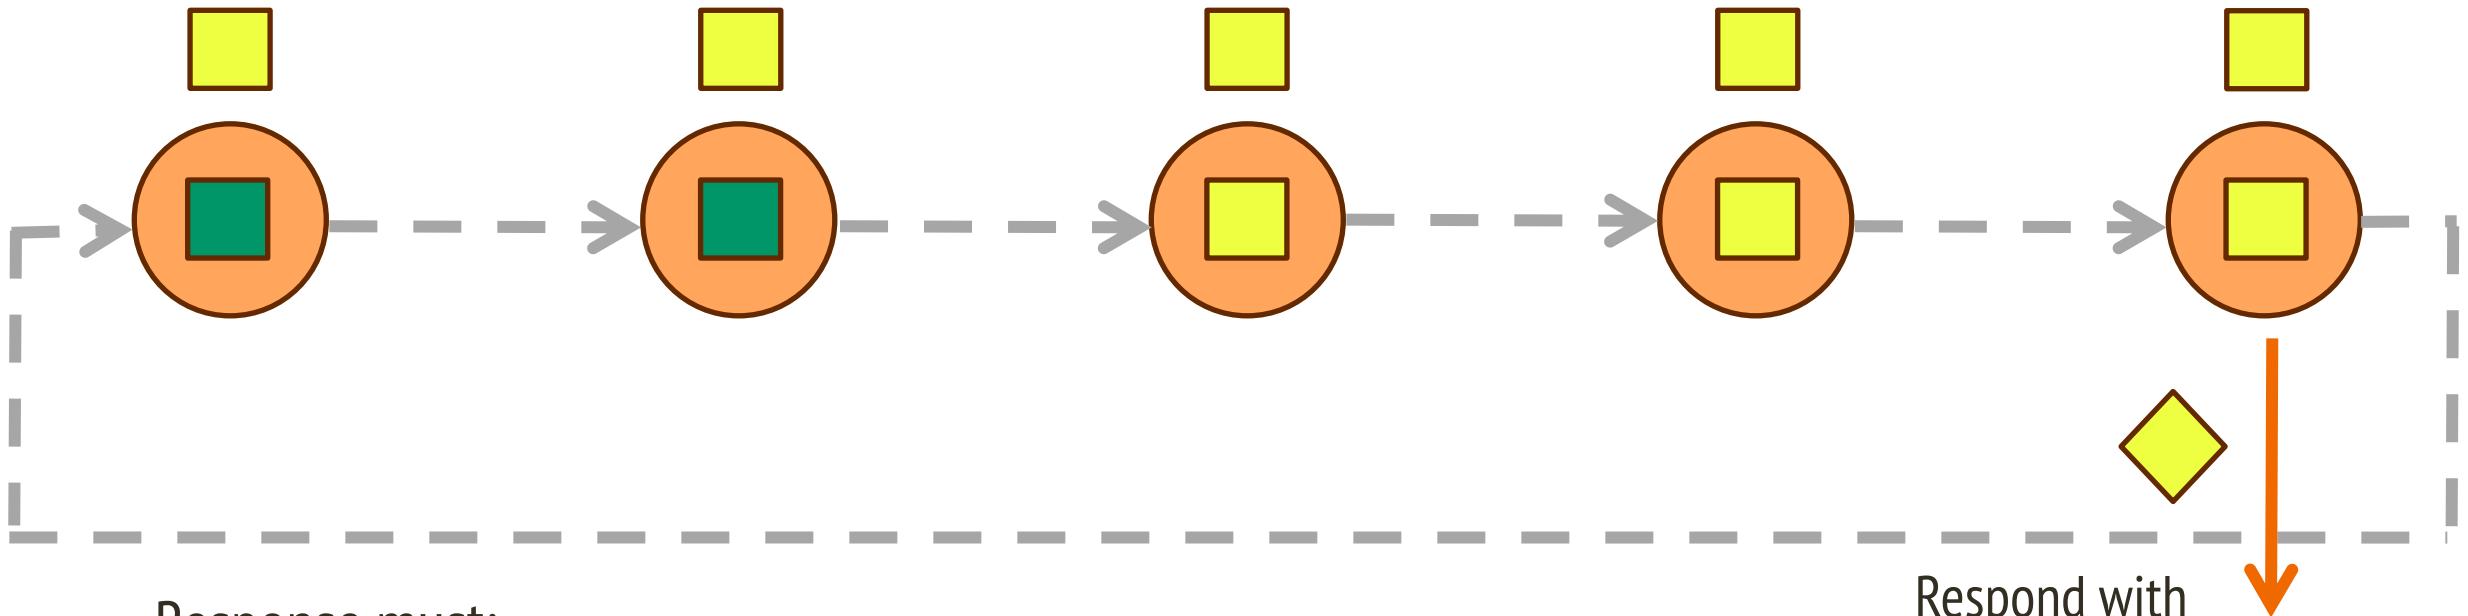
- Contain all completed writes (■)
- Contain all completed reads

Respond with
current state

Local Linearizable Reads

Stronger than required (█)

Best we can do without extra communication

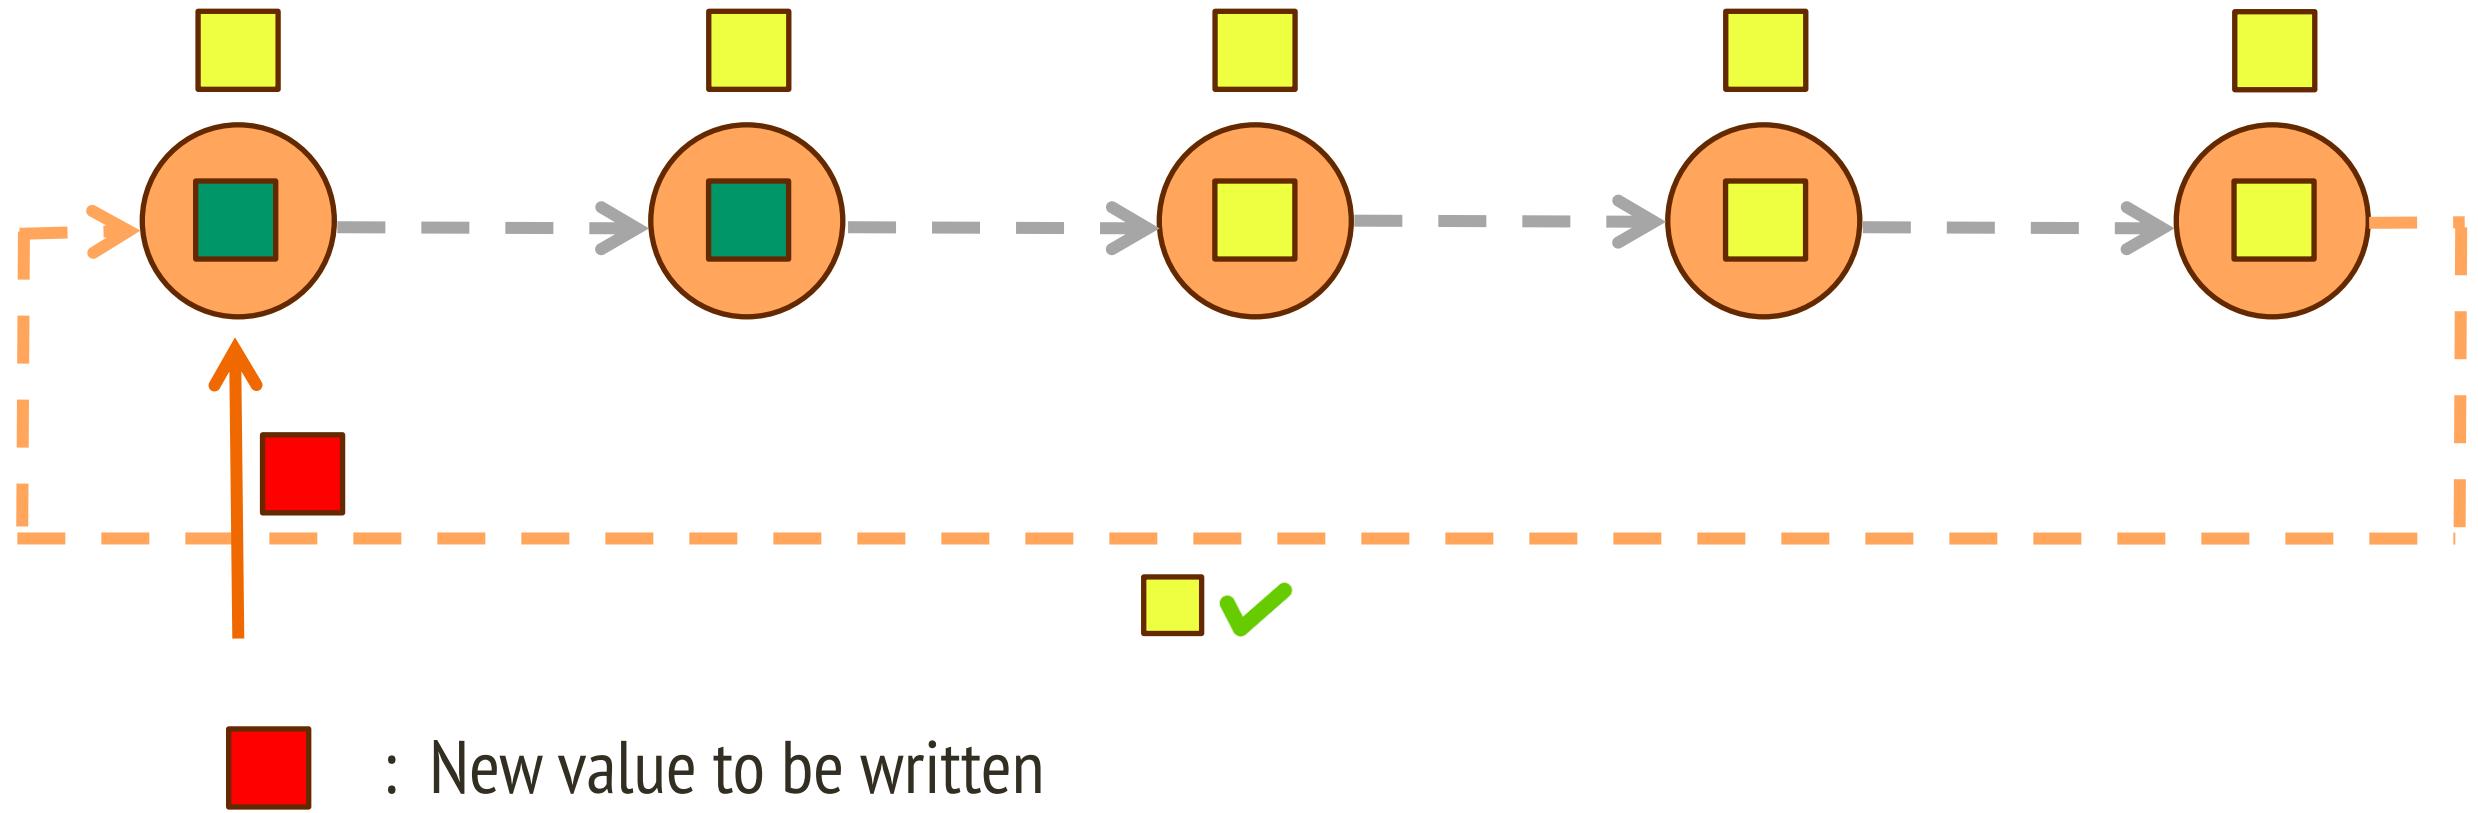


Response must:

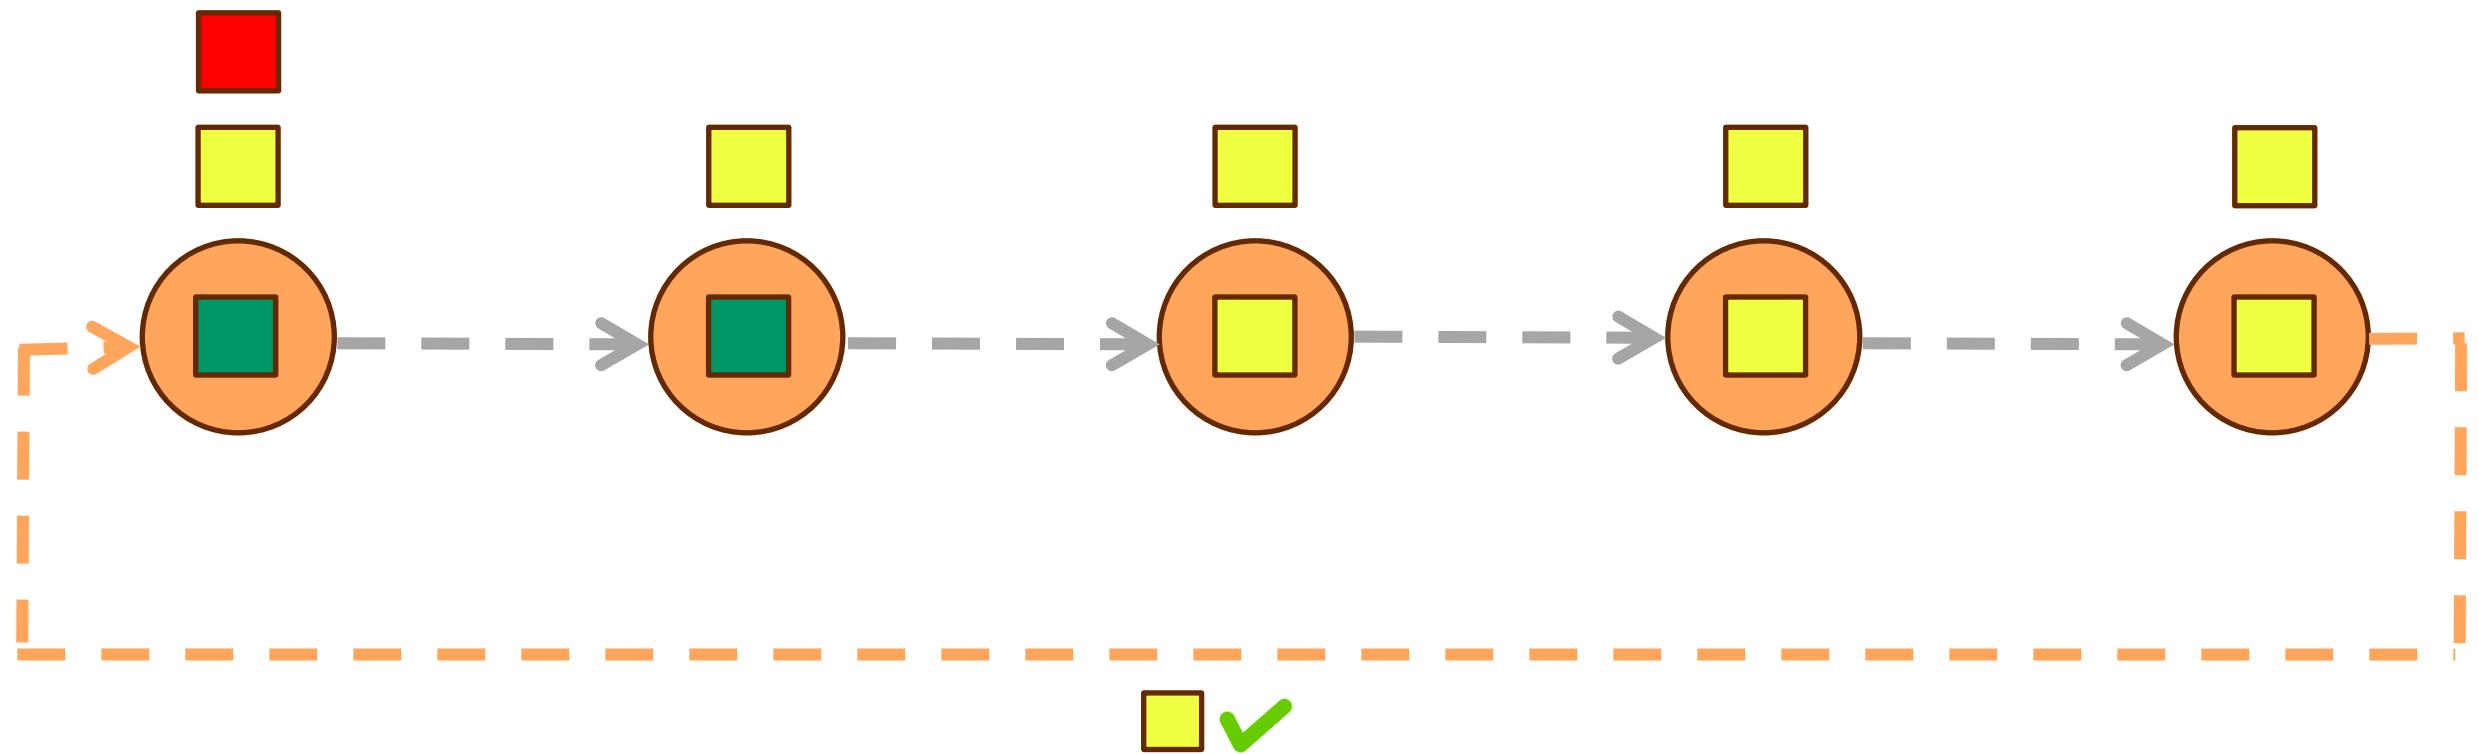
- Contain all completed writes (█)
- Contain all completed reads

Respond with
current state

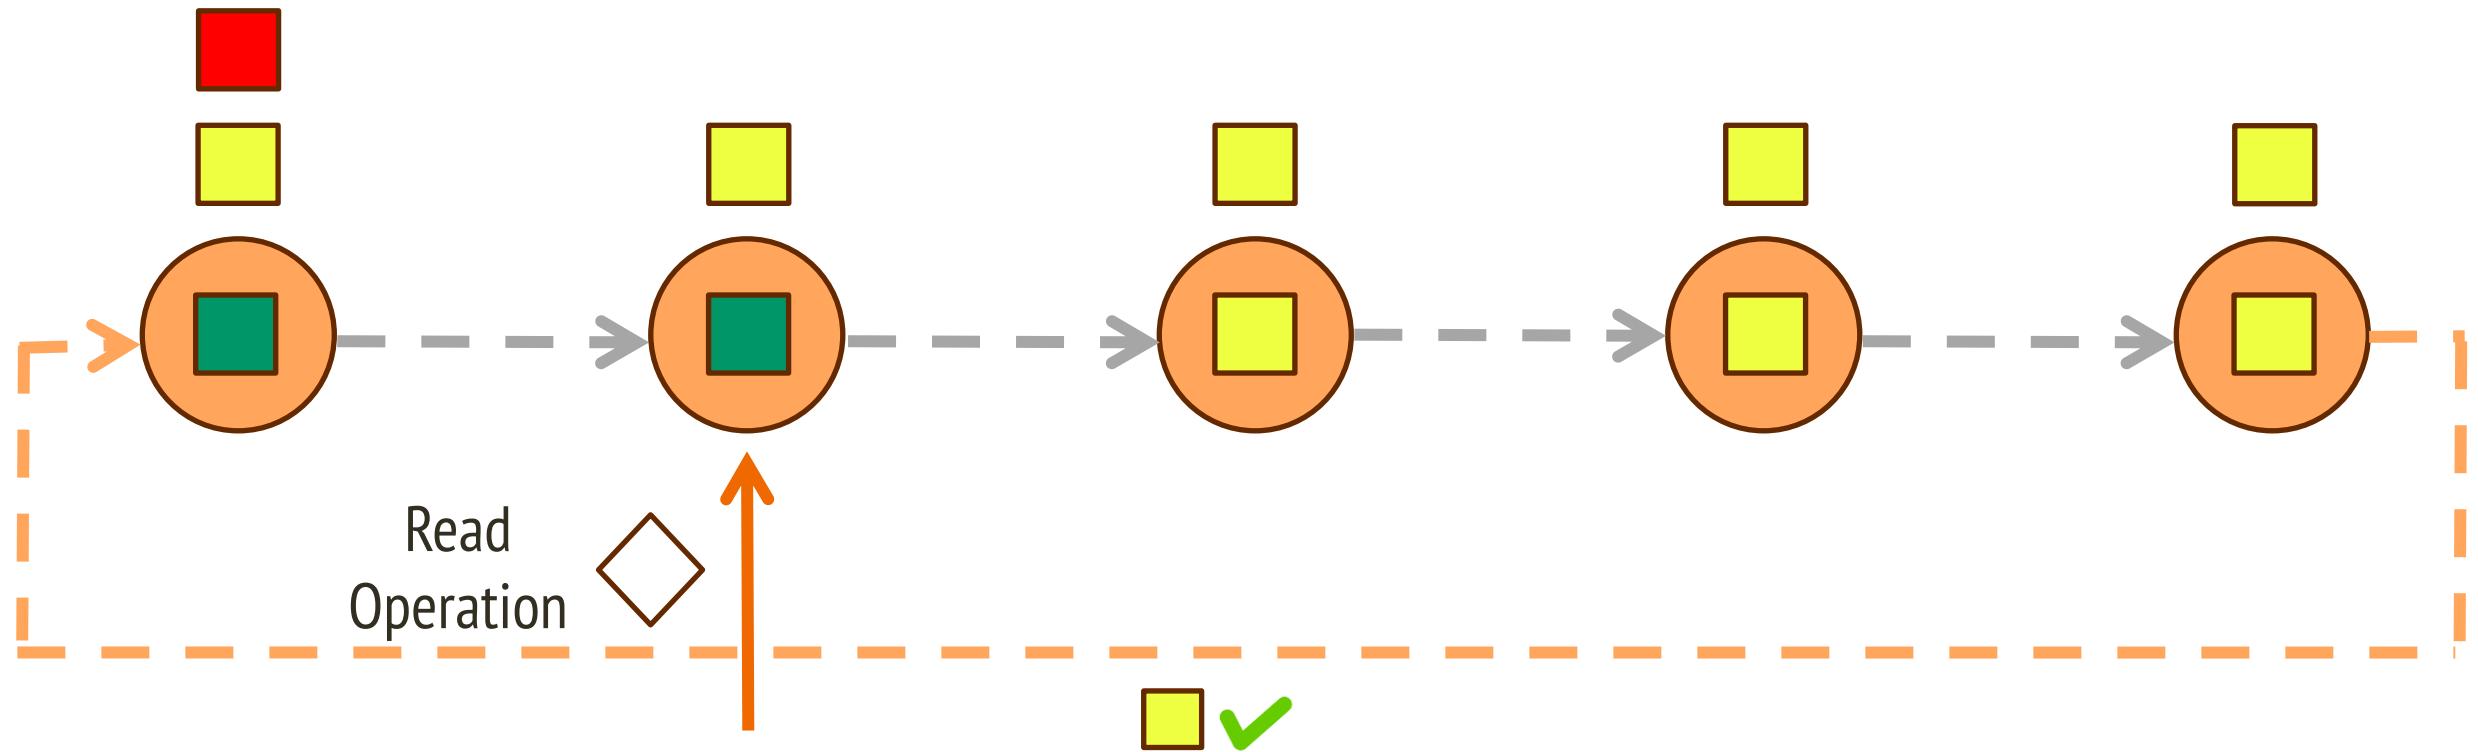
Local Linearizable Reads



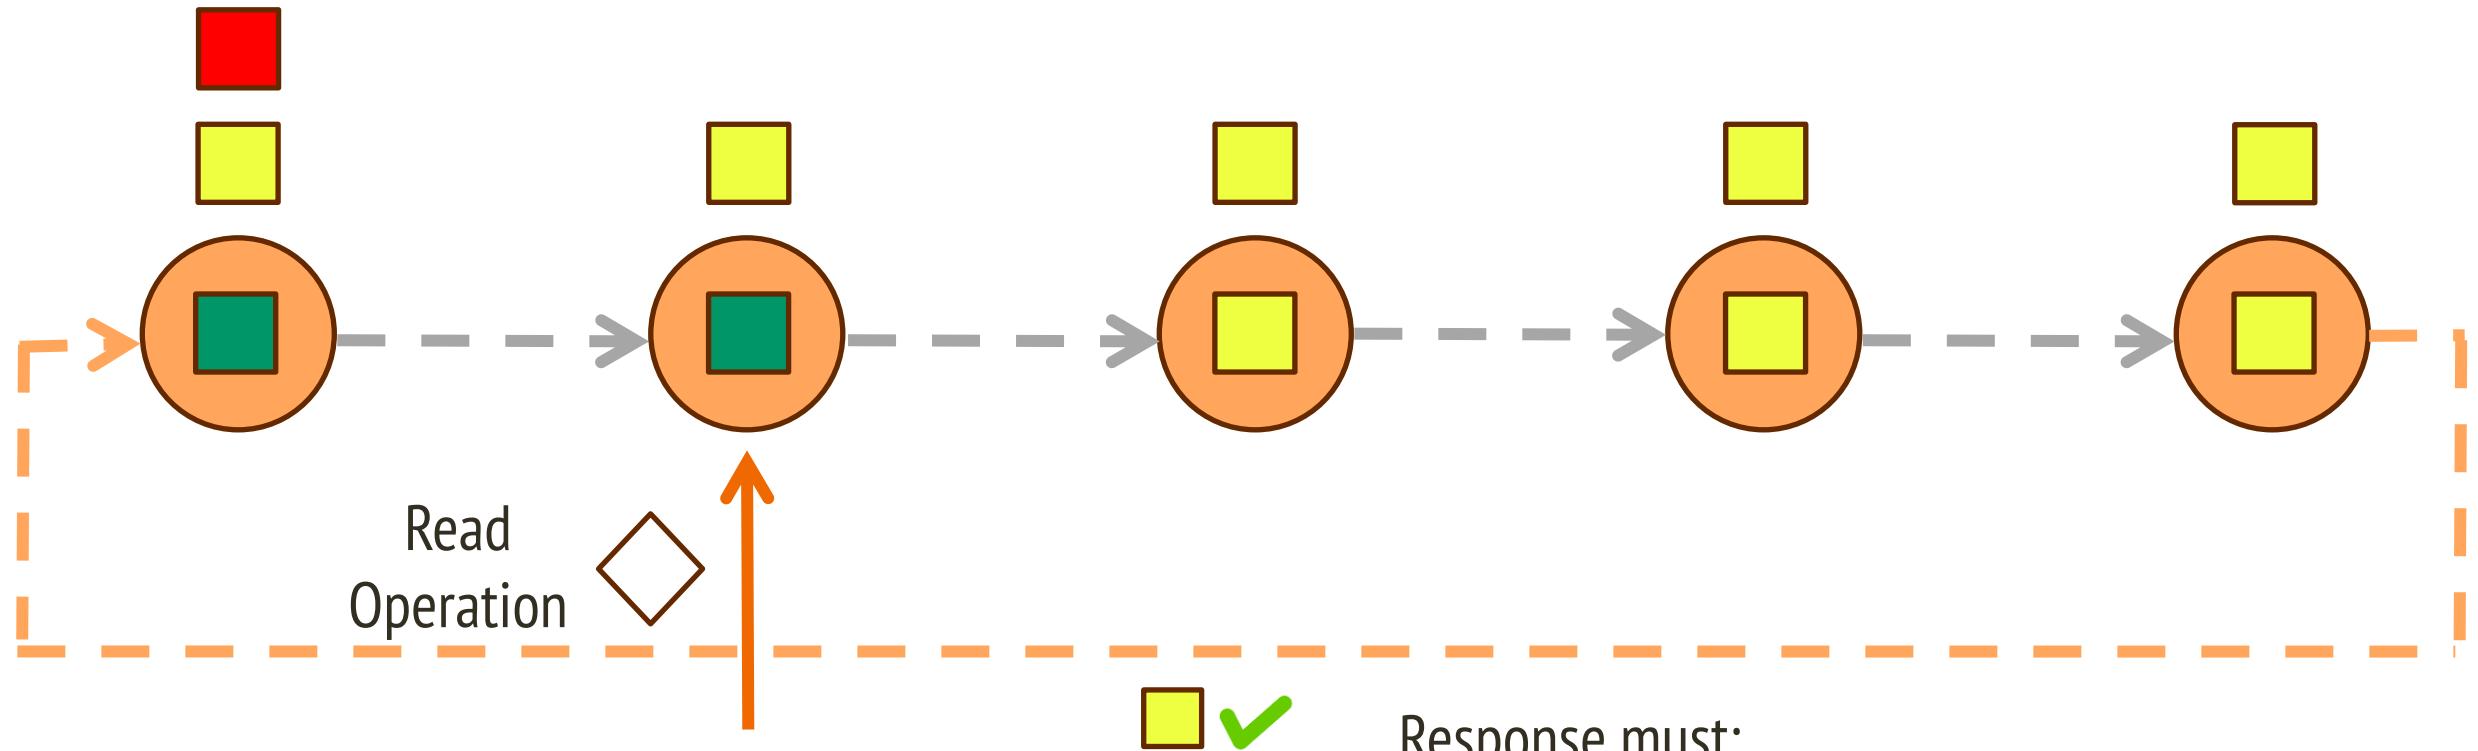
Local Linearizable Reads



Local Linearizable Reads



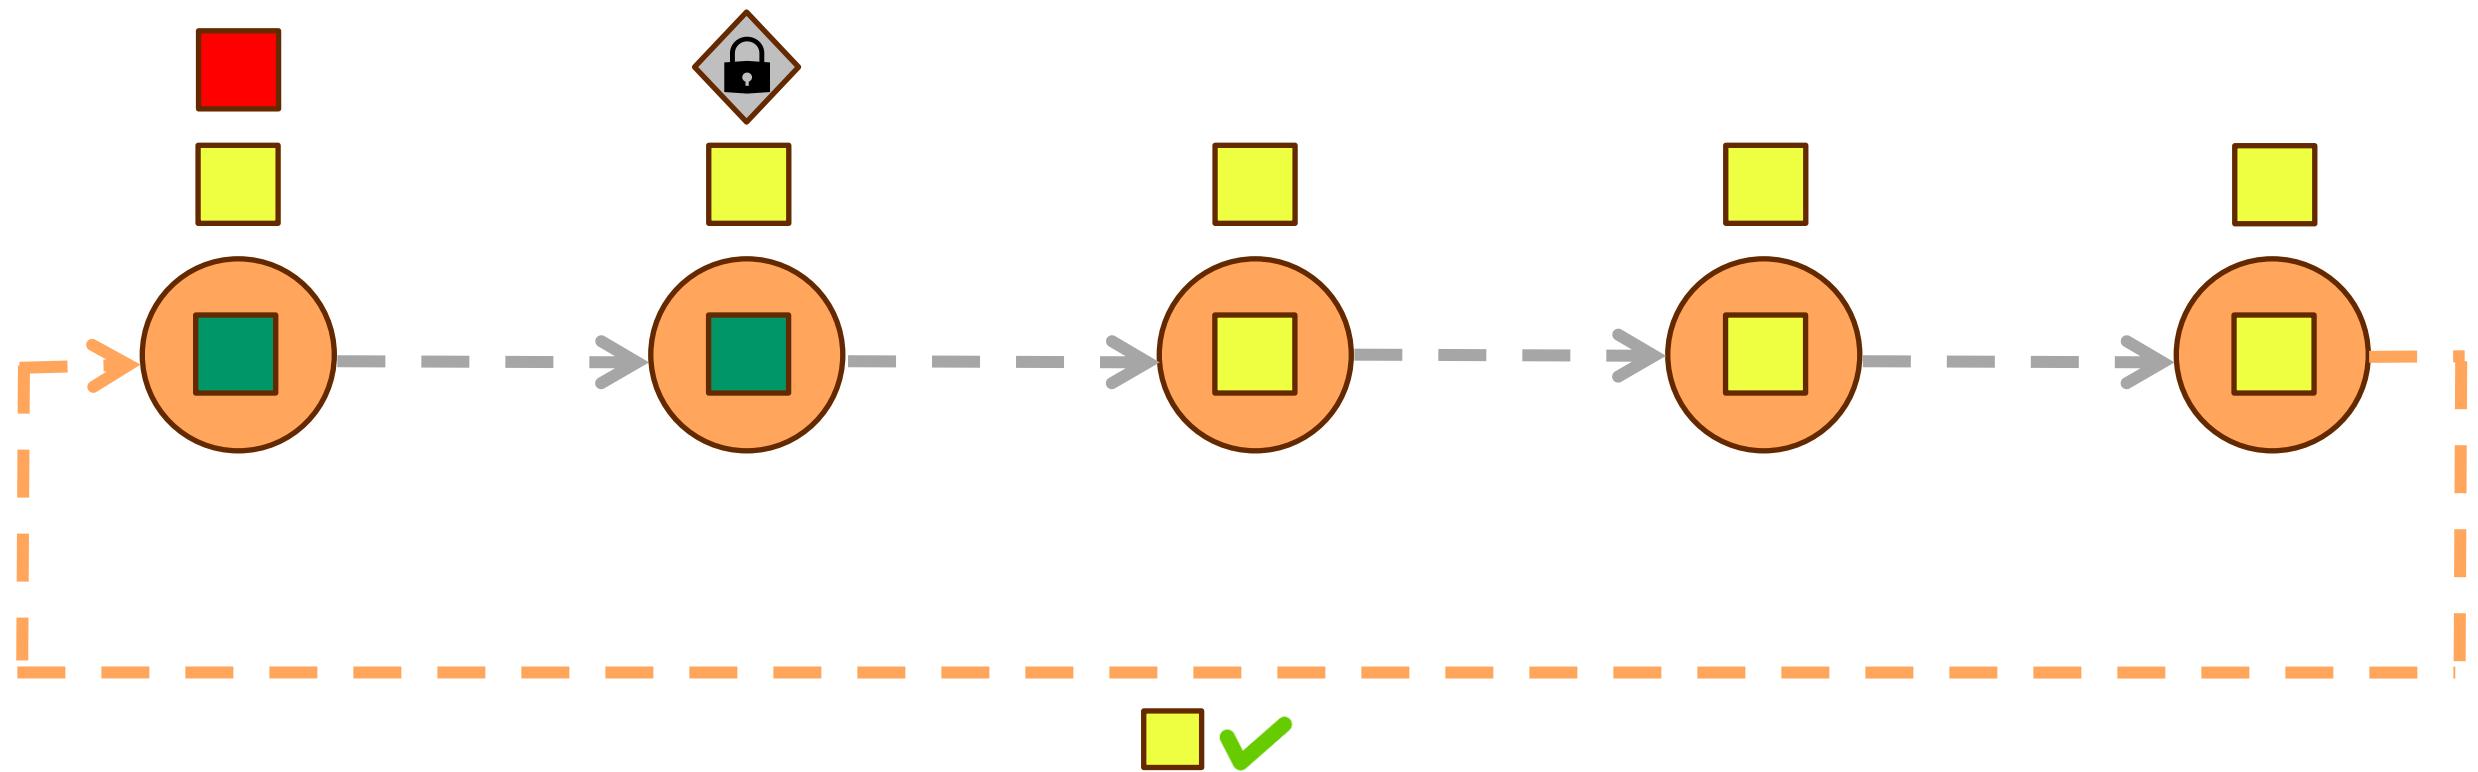
Local Linearizable Reads



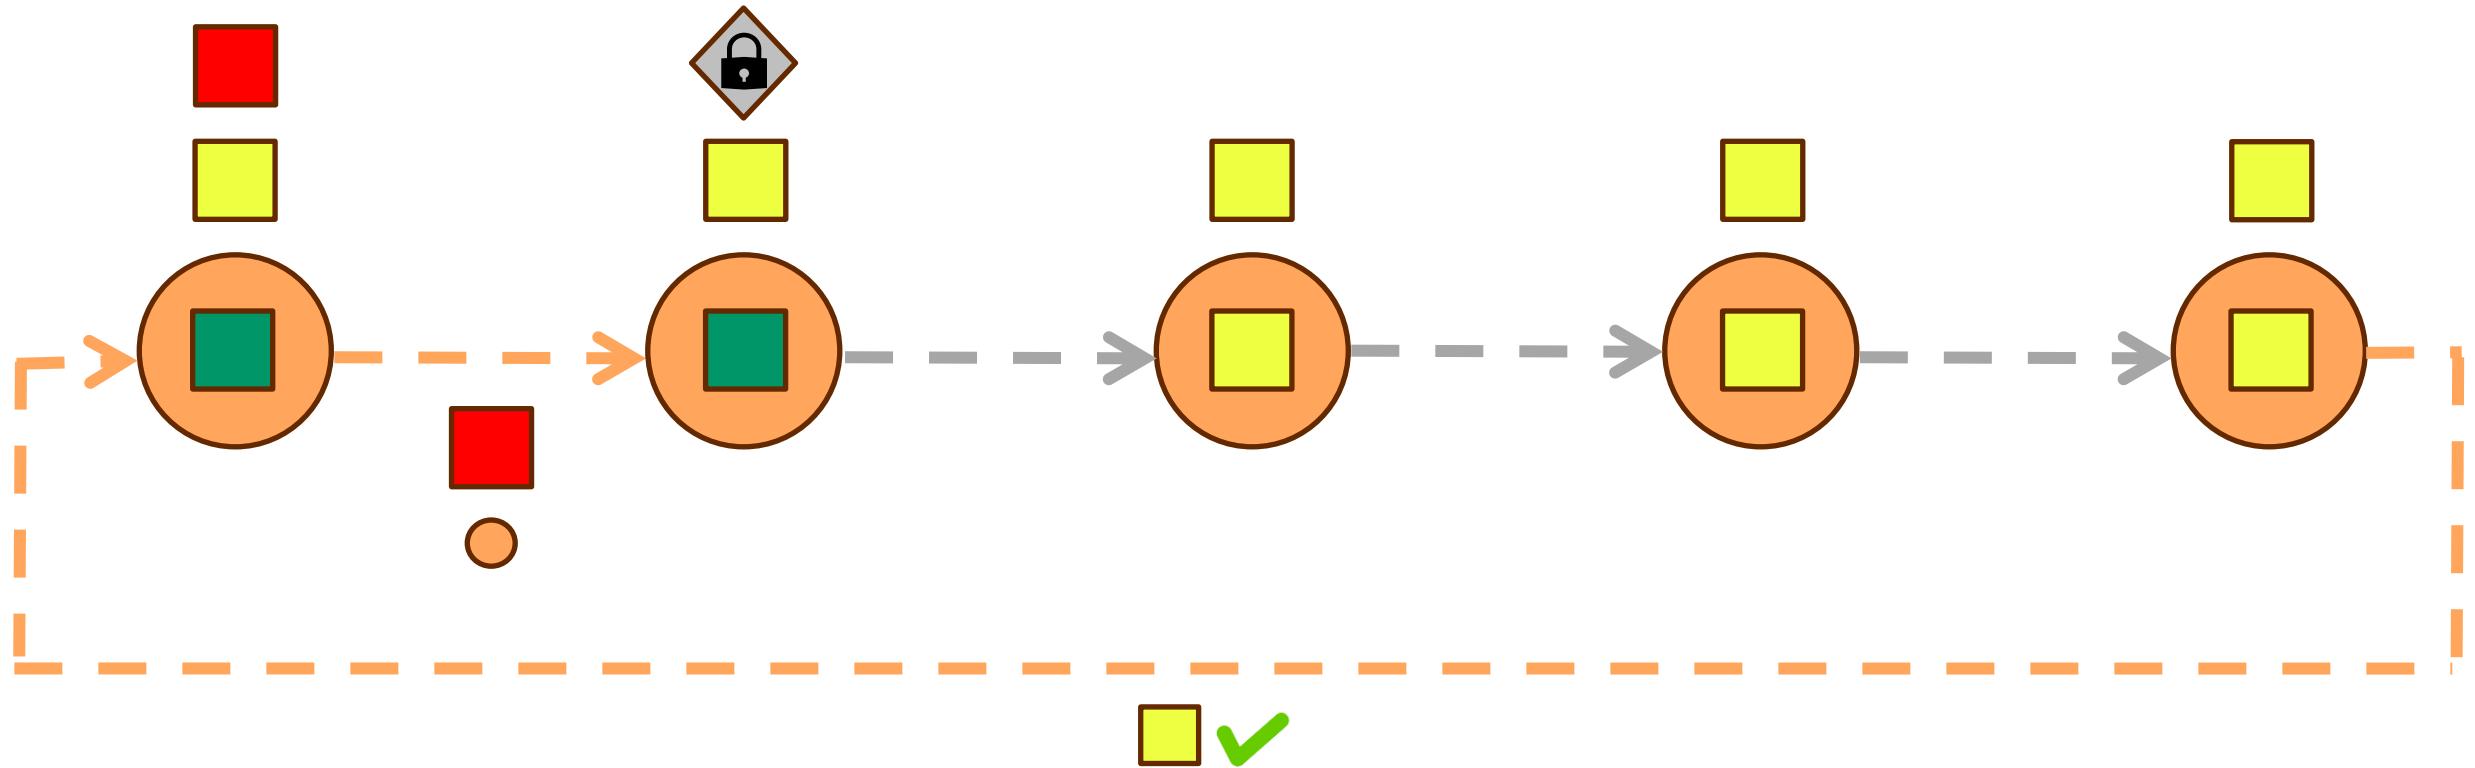
Response must:

- Contain all completed writes
- Contain all completed reads (■)

Local Linearizable Reads

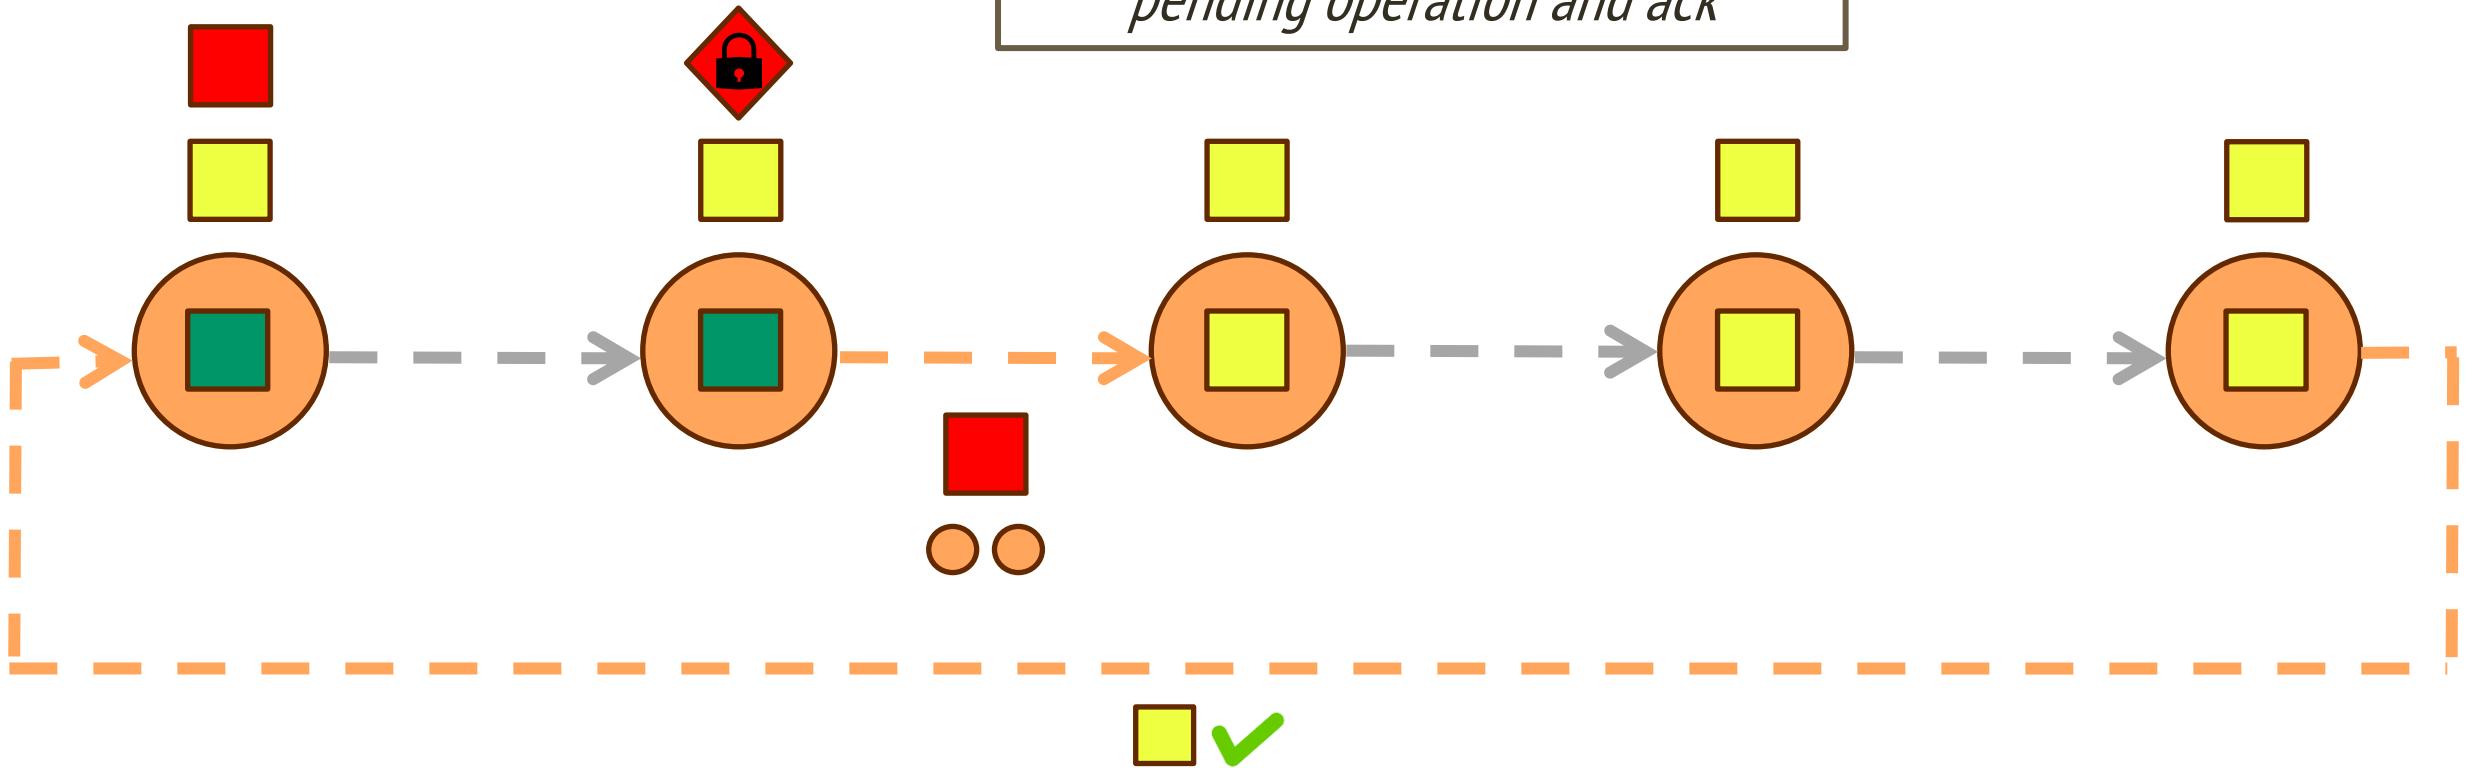


Local Linearizable Reads

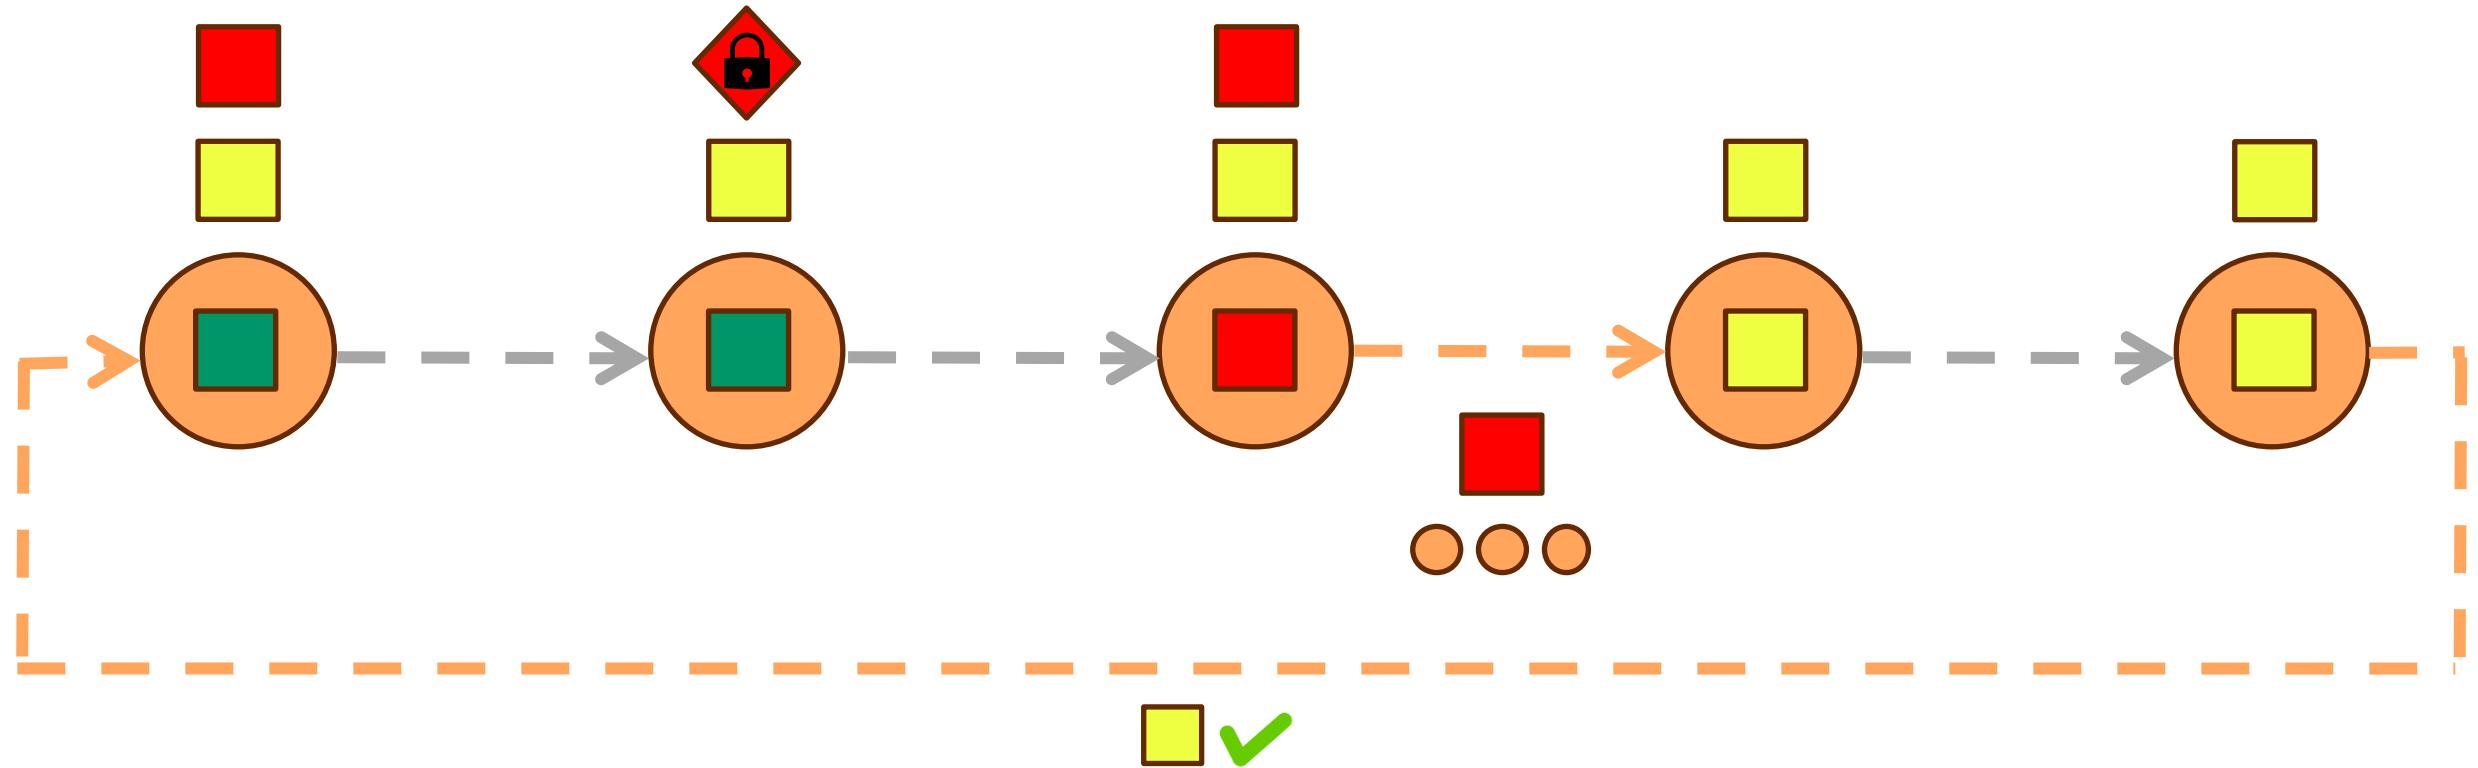


Local Linearizable Reads

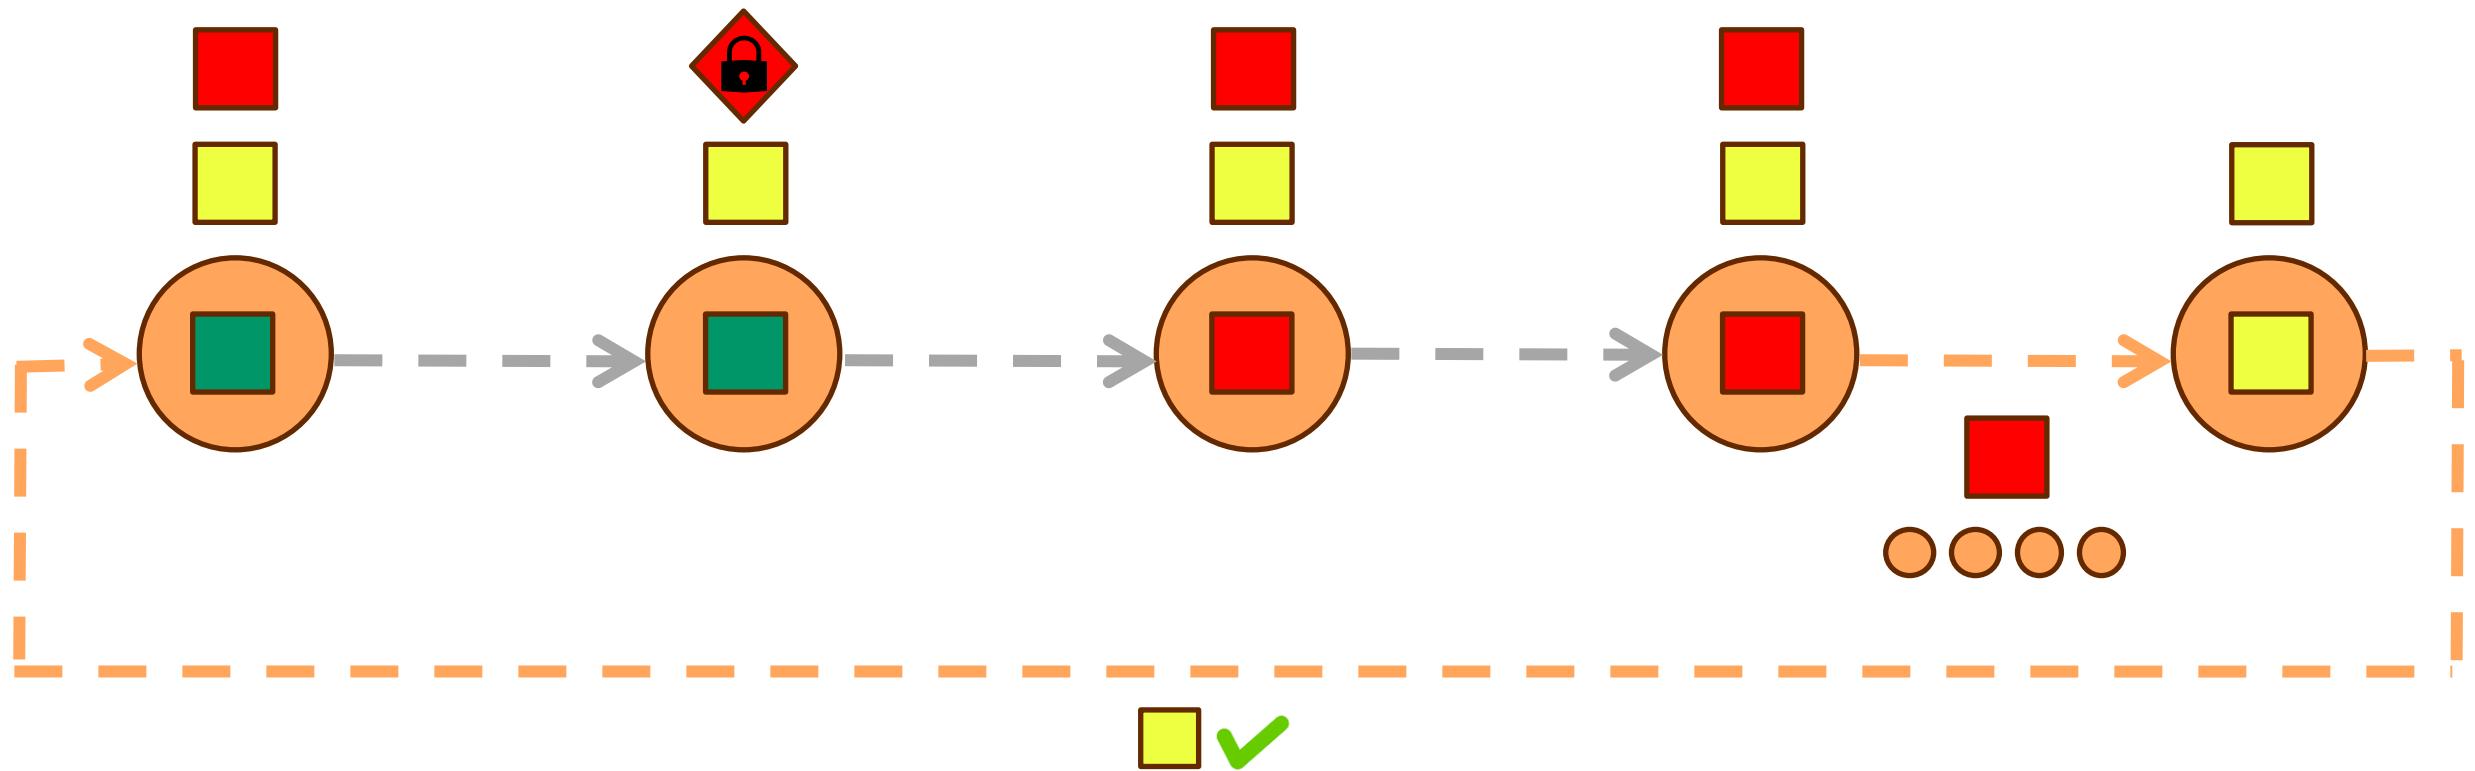
Red operation will "push" every pending operation and ack



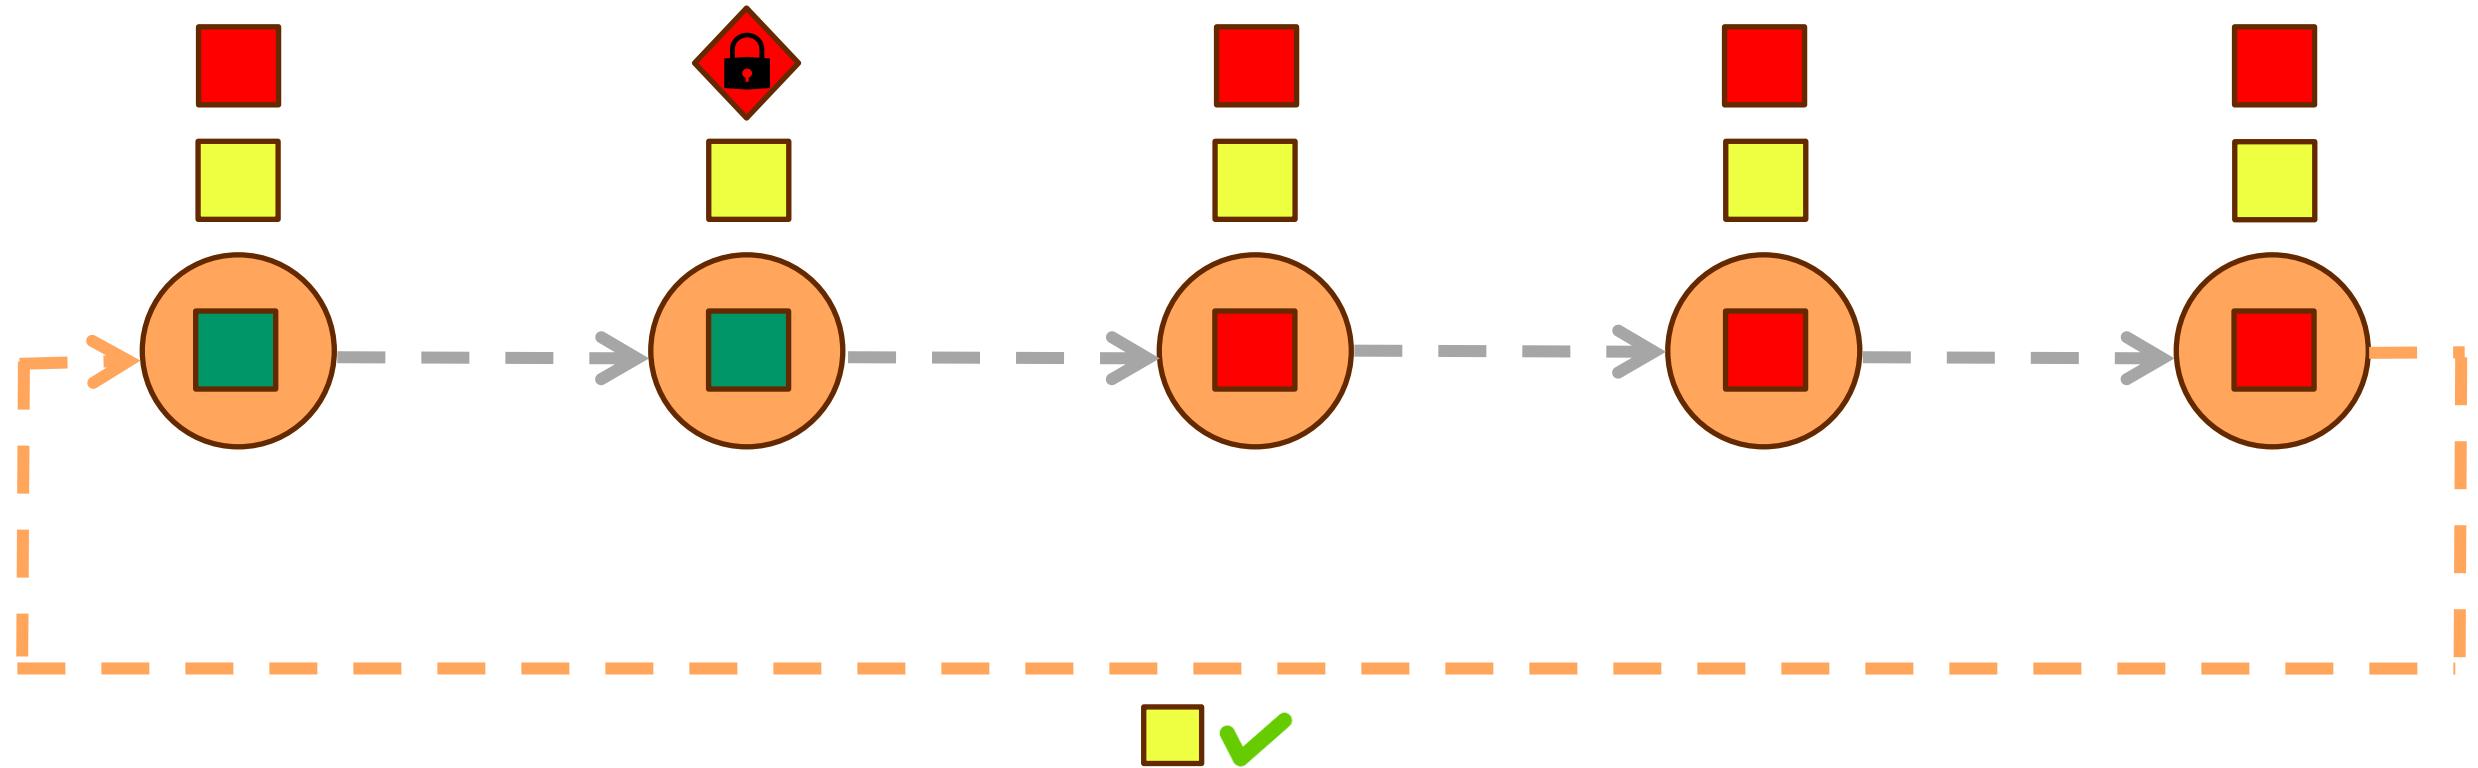
Local Linearizable Reads



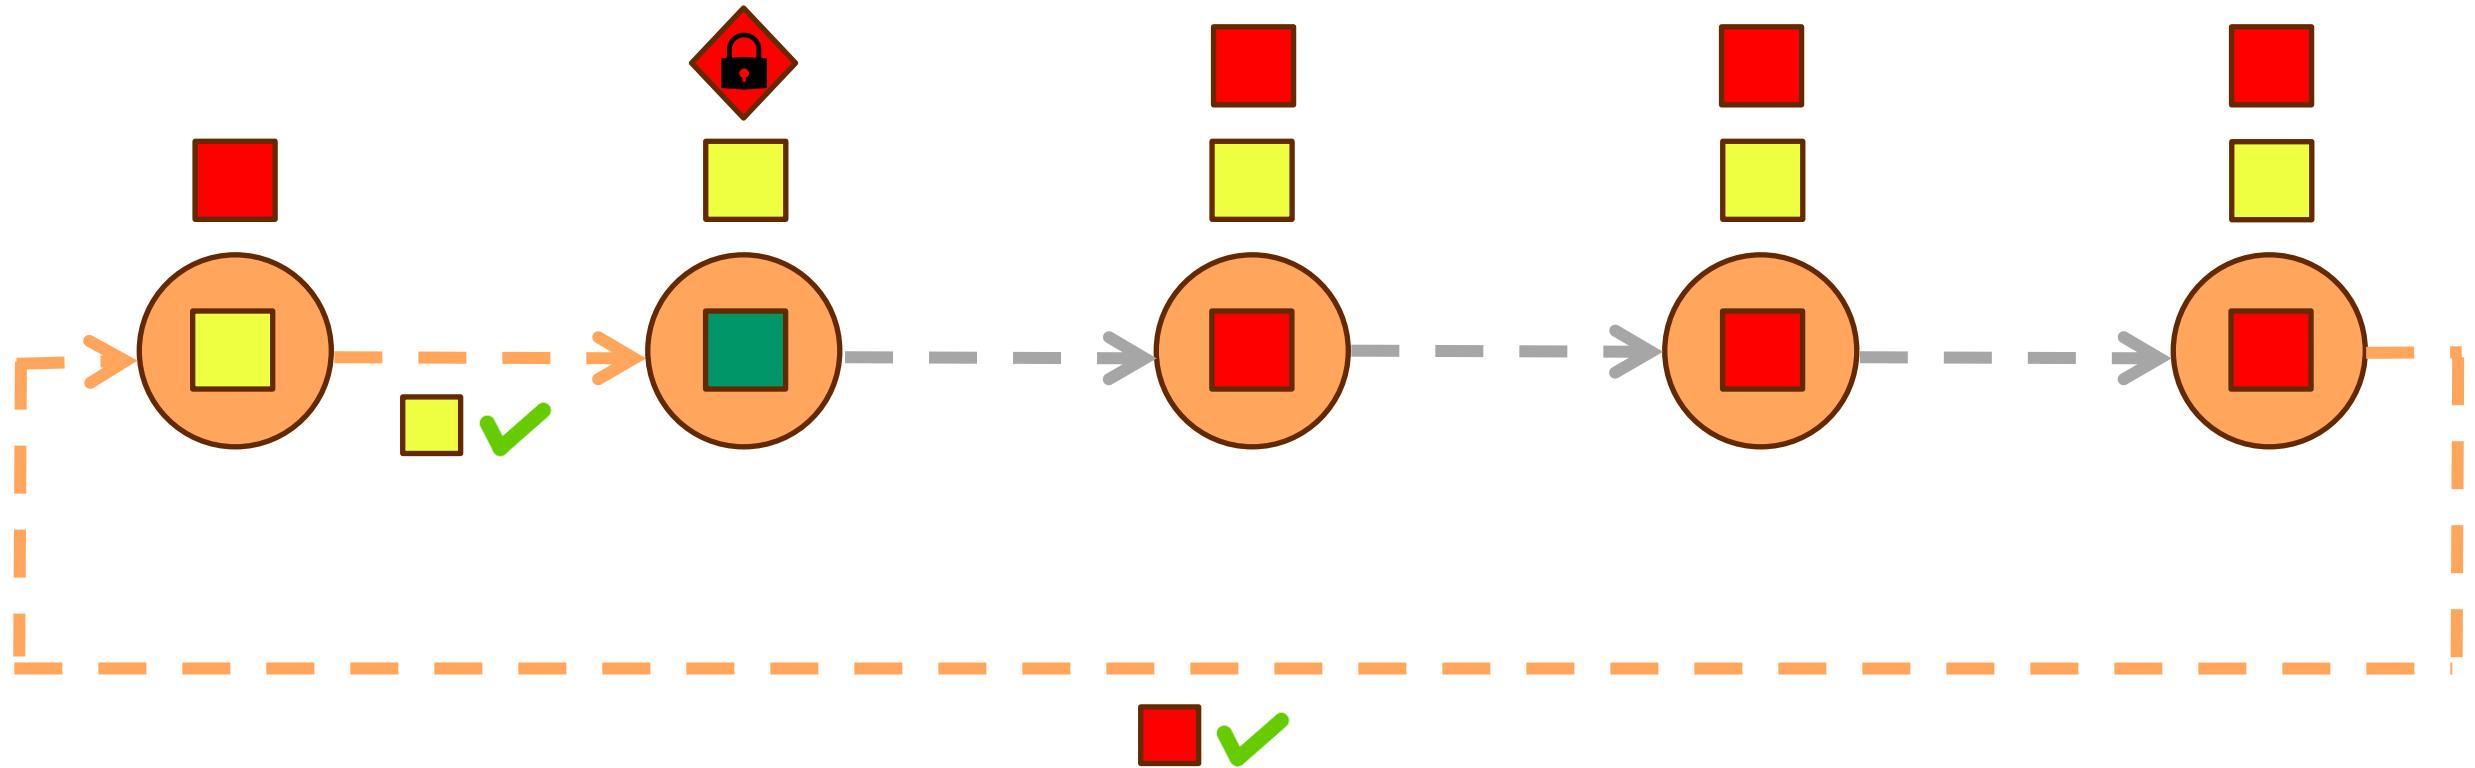
Local Linearizable Reads



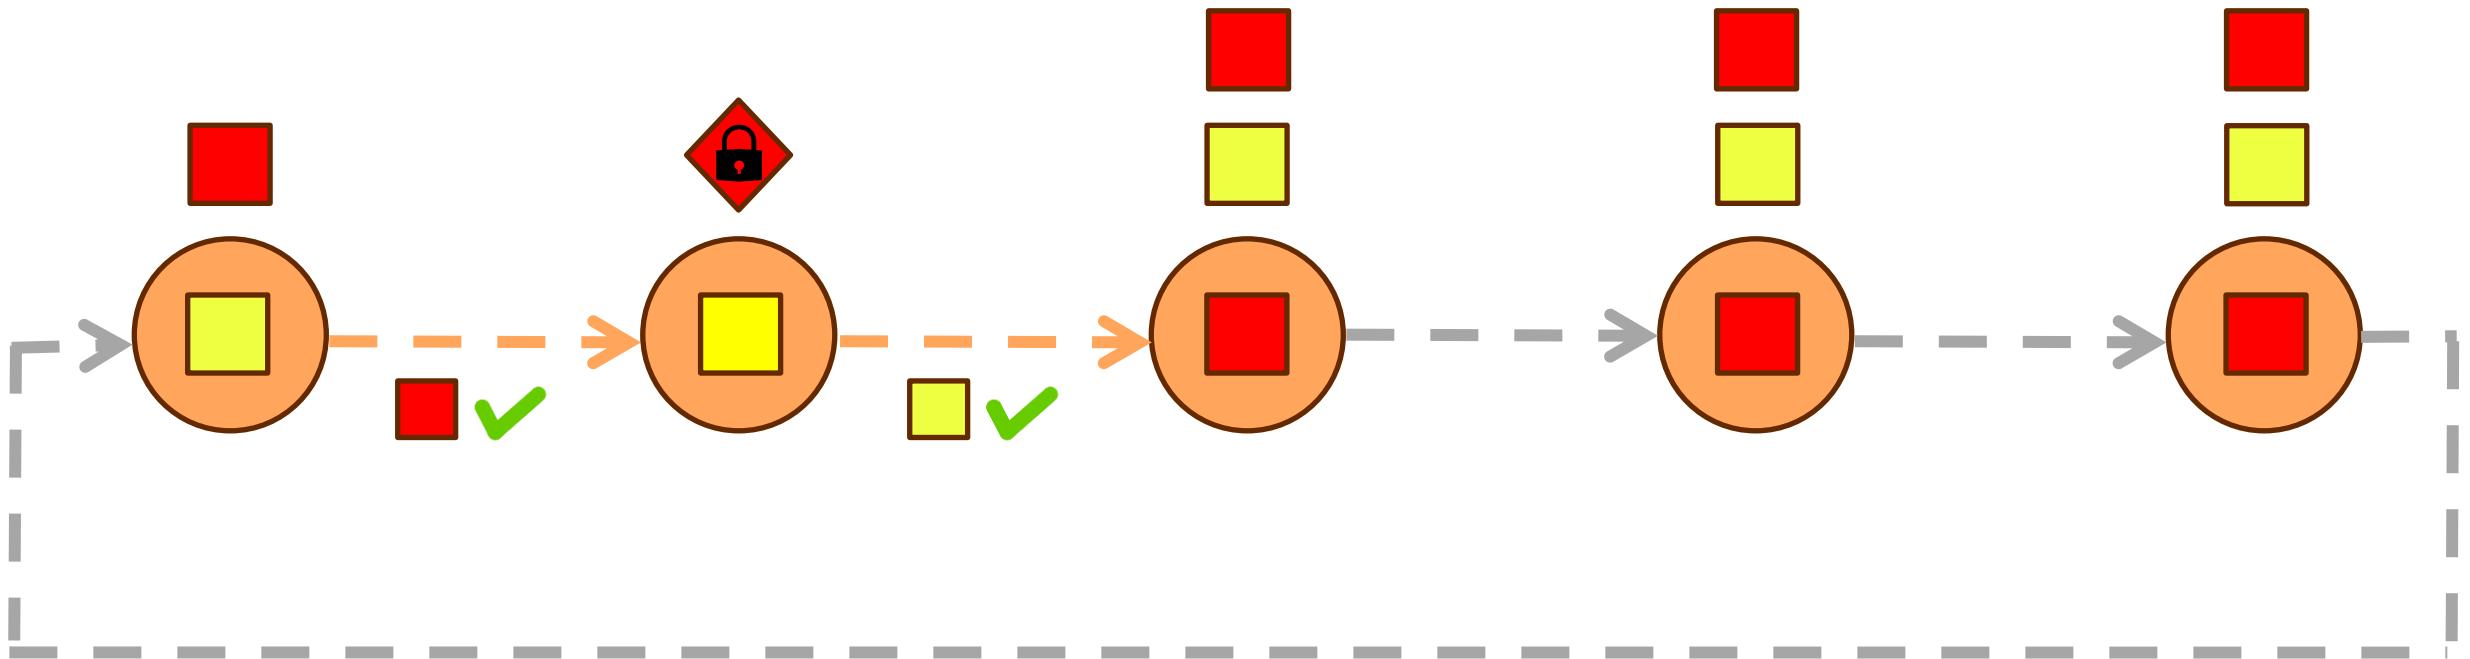
Local Linearizable Reads



Local Linearizable Reads



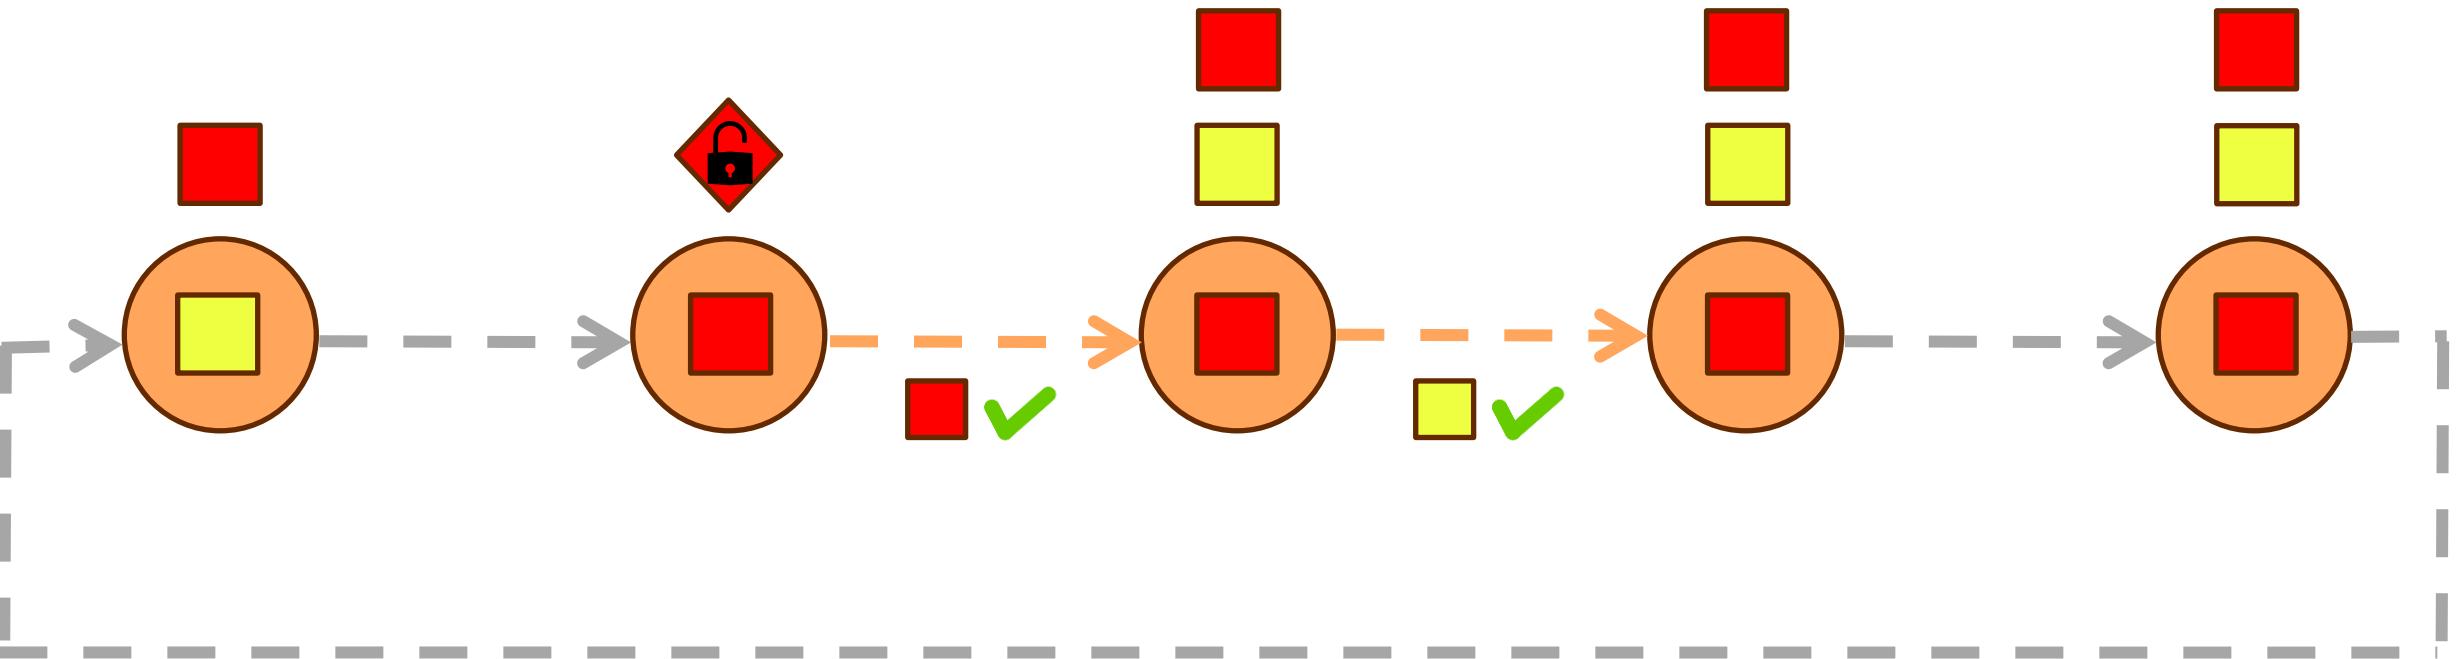
Local Linearizable Reads



Response must:

- Contain all completed writes
- Contain all completed reads (Yellow)

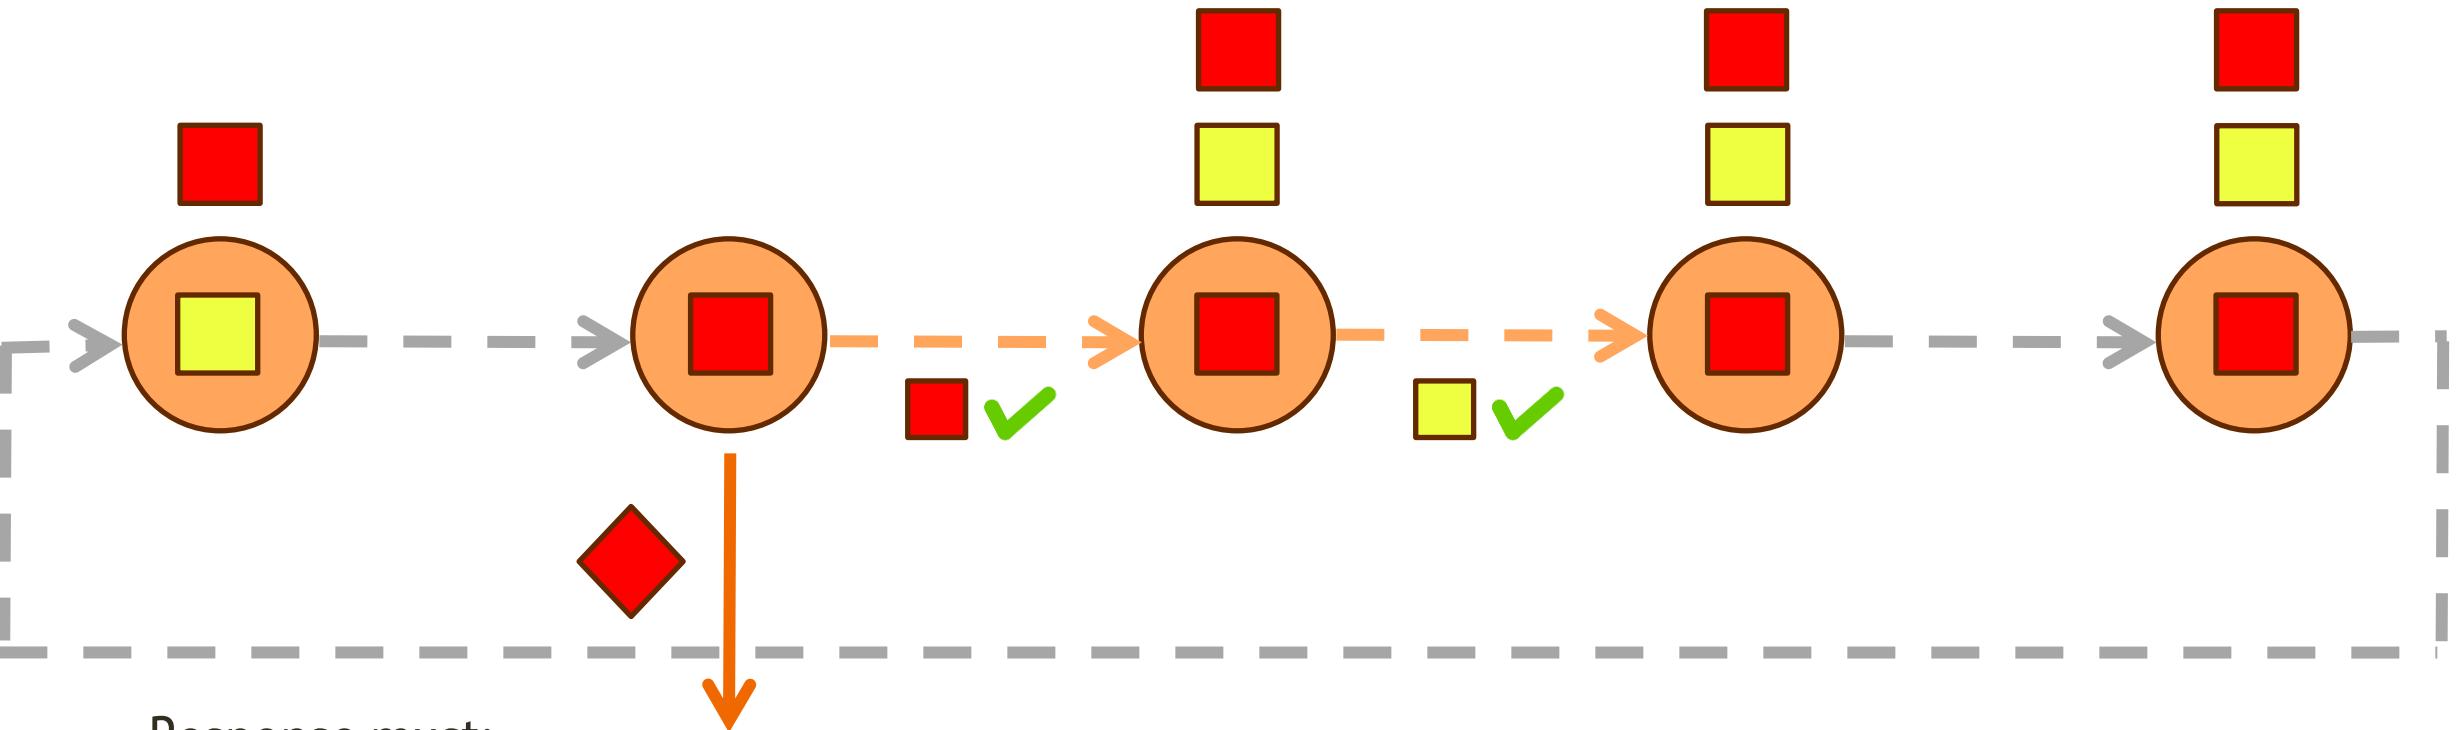
Local Linearizable Reads



Response must:

- Contain all completed writes
- Contain all completed reads ()

Local Linearizable Reads



Response must:

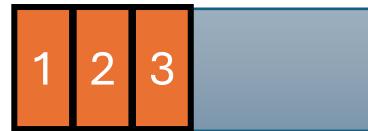
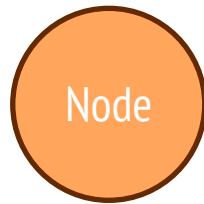
- Contain all completed writes
- Contain all completed reads (Yellow)

Local Linearizable Reads: Summary

- Local read operations do not break linearizability
 - Ensures that all previously **completed reads and writes** are visible
- **No additional communication** steps are required
 - More **conservative** than required, but **unavoidable without coordination**
- **Periodic No-Op** messages ensure that reads are replied to during low write workloads

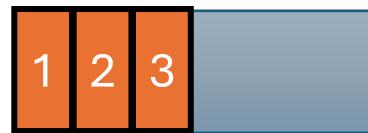
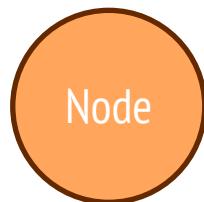
In Depth: The Log Representation

*Log contains the ordered
list of all operations seen*



In Depth: The Log Representation

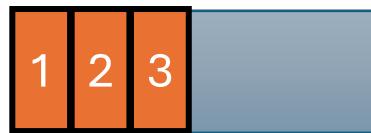
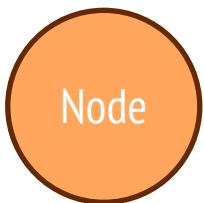
*Log contains the ordered
list of all operations seen*



*Each Node/Replica has its
own view of the log*

In Depth: The Log Representation

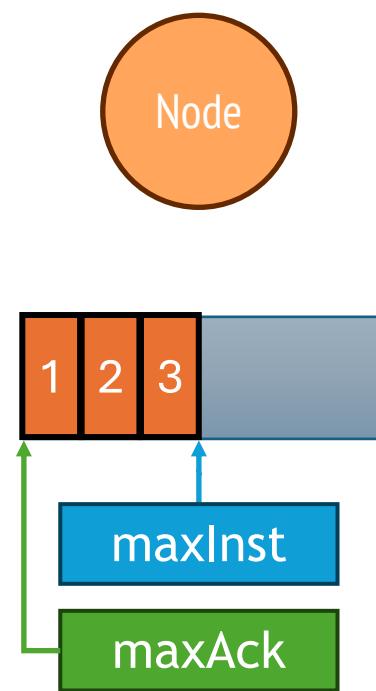
Log contains the ordered list of all operations seen



Each Node/Replica has its own view of the log

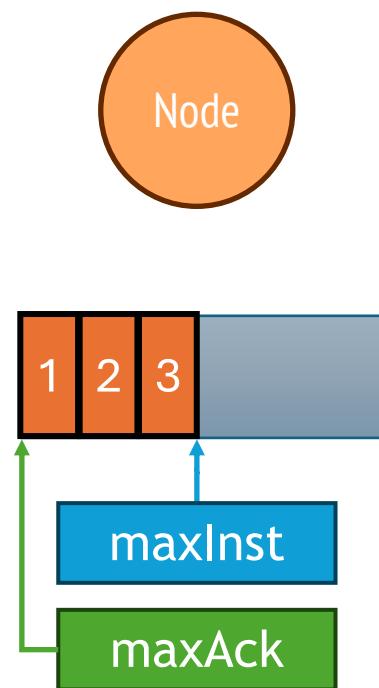
Each Instance writes to a new log index

In Depth: The Log Representation



In Depth: The Log Representation

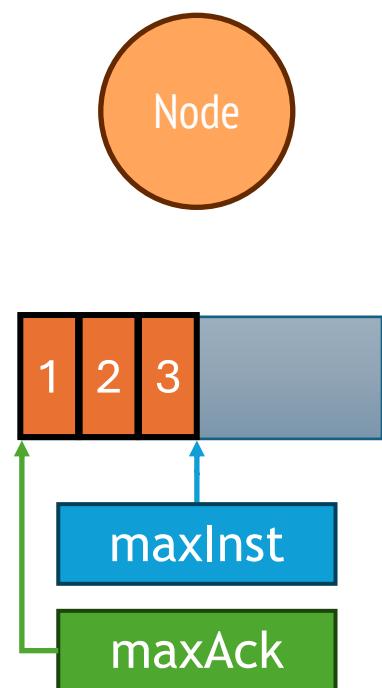
*maxInst keeps track of
replica's max seen instance*



In Depth: The Log Representation

*maxInst keeps track of
replica's max seen instance*

*maxInst is essentially the
length of the replica's log*

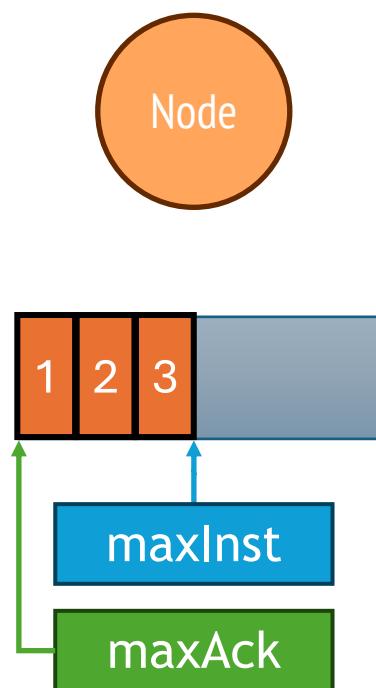


In Depth: The Log Representation

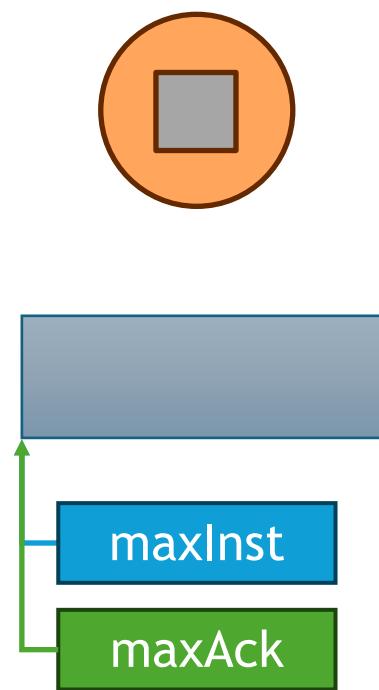
*maxInst keeps track of
replica's max seen instance*

*maxInst is essentially the
length of the replica's log*

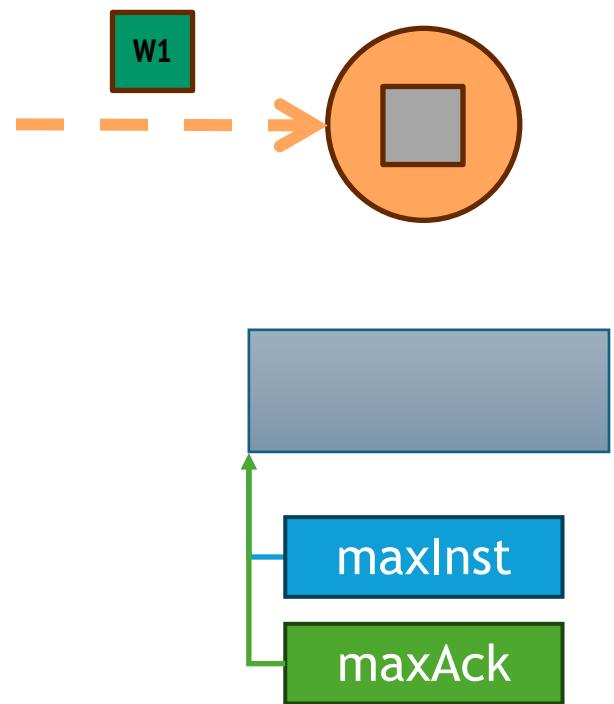
*maxAck keeps track of replica's
max seen acknowledgment*



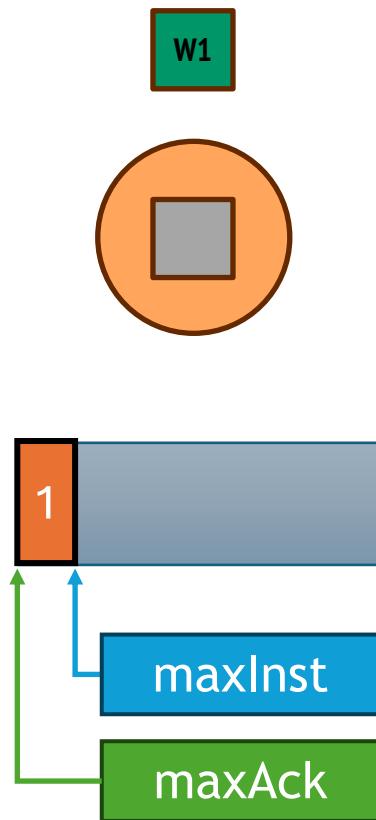
In Depth: The Log Representation



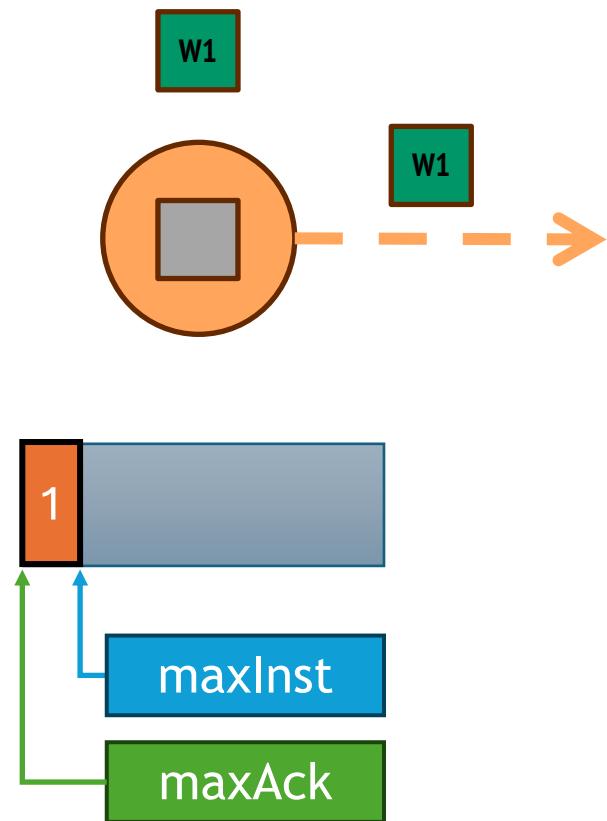
In Depth: The Log Representation



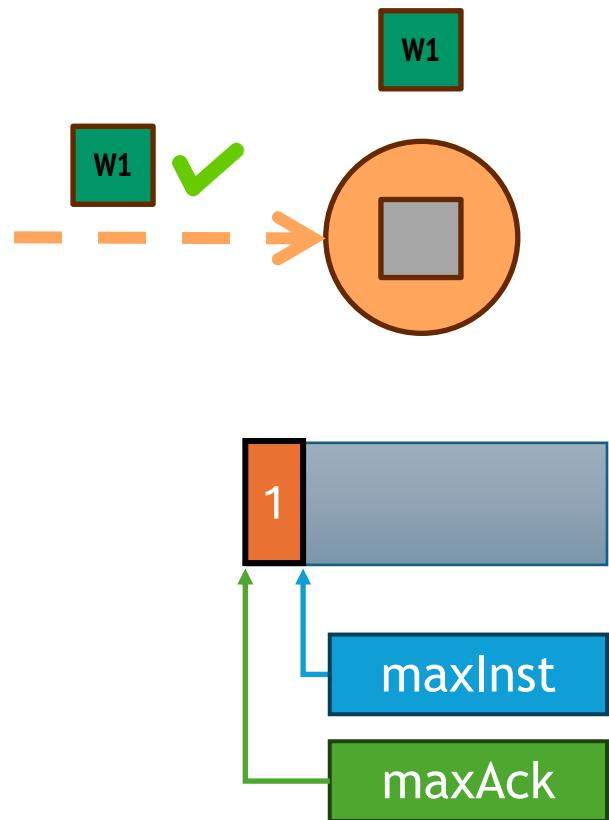
In Depth: The Log Representation



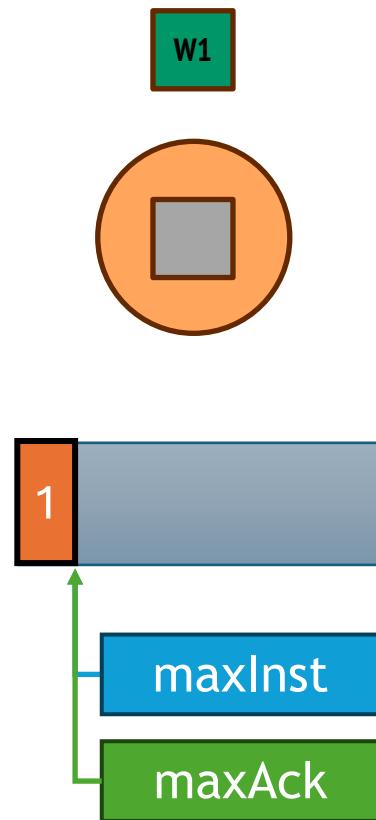
In Depth: The Log Representation



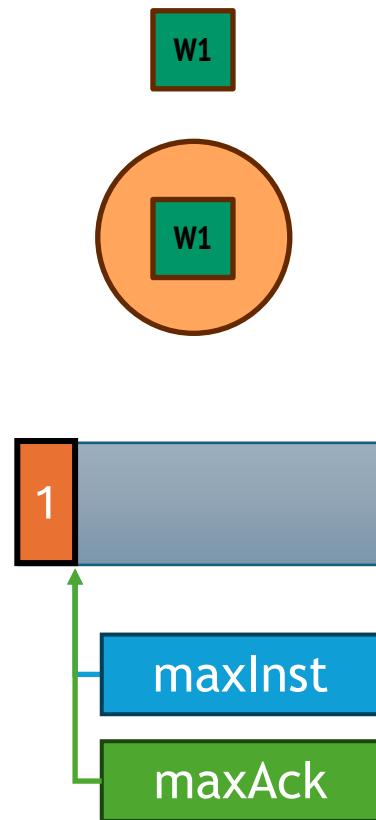
In Depth: The Log Representation



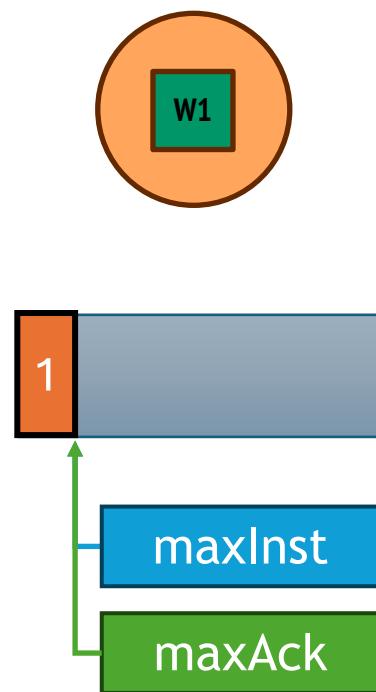
In Depth: The Log Representation



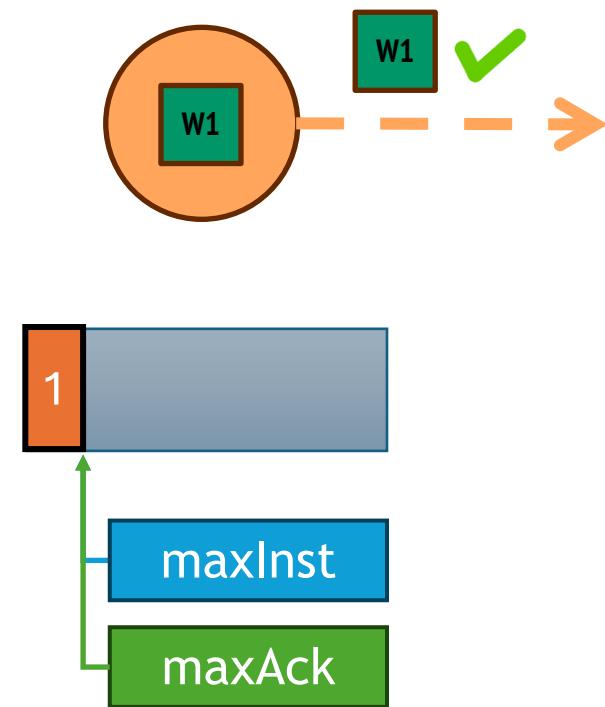
In Depth: The Log Representation



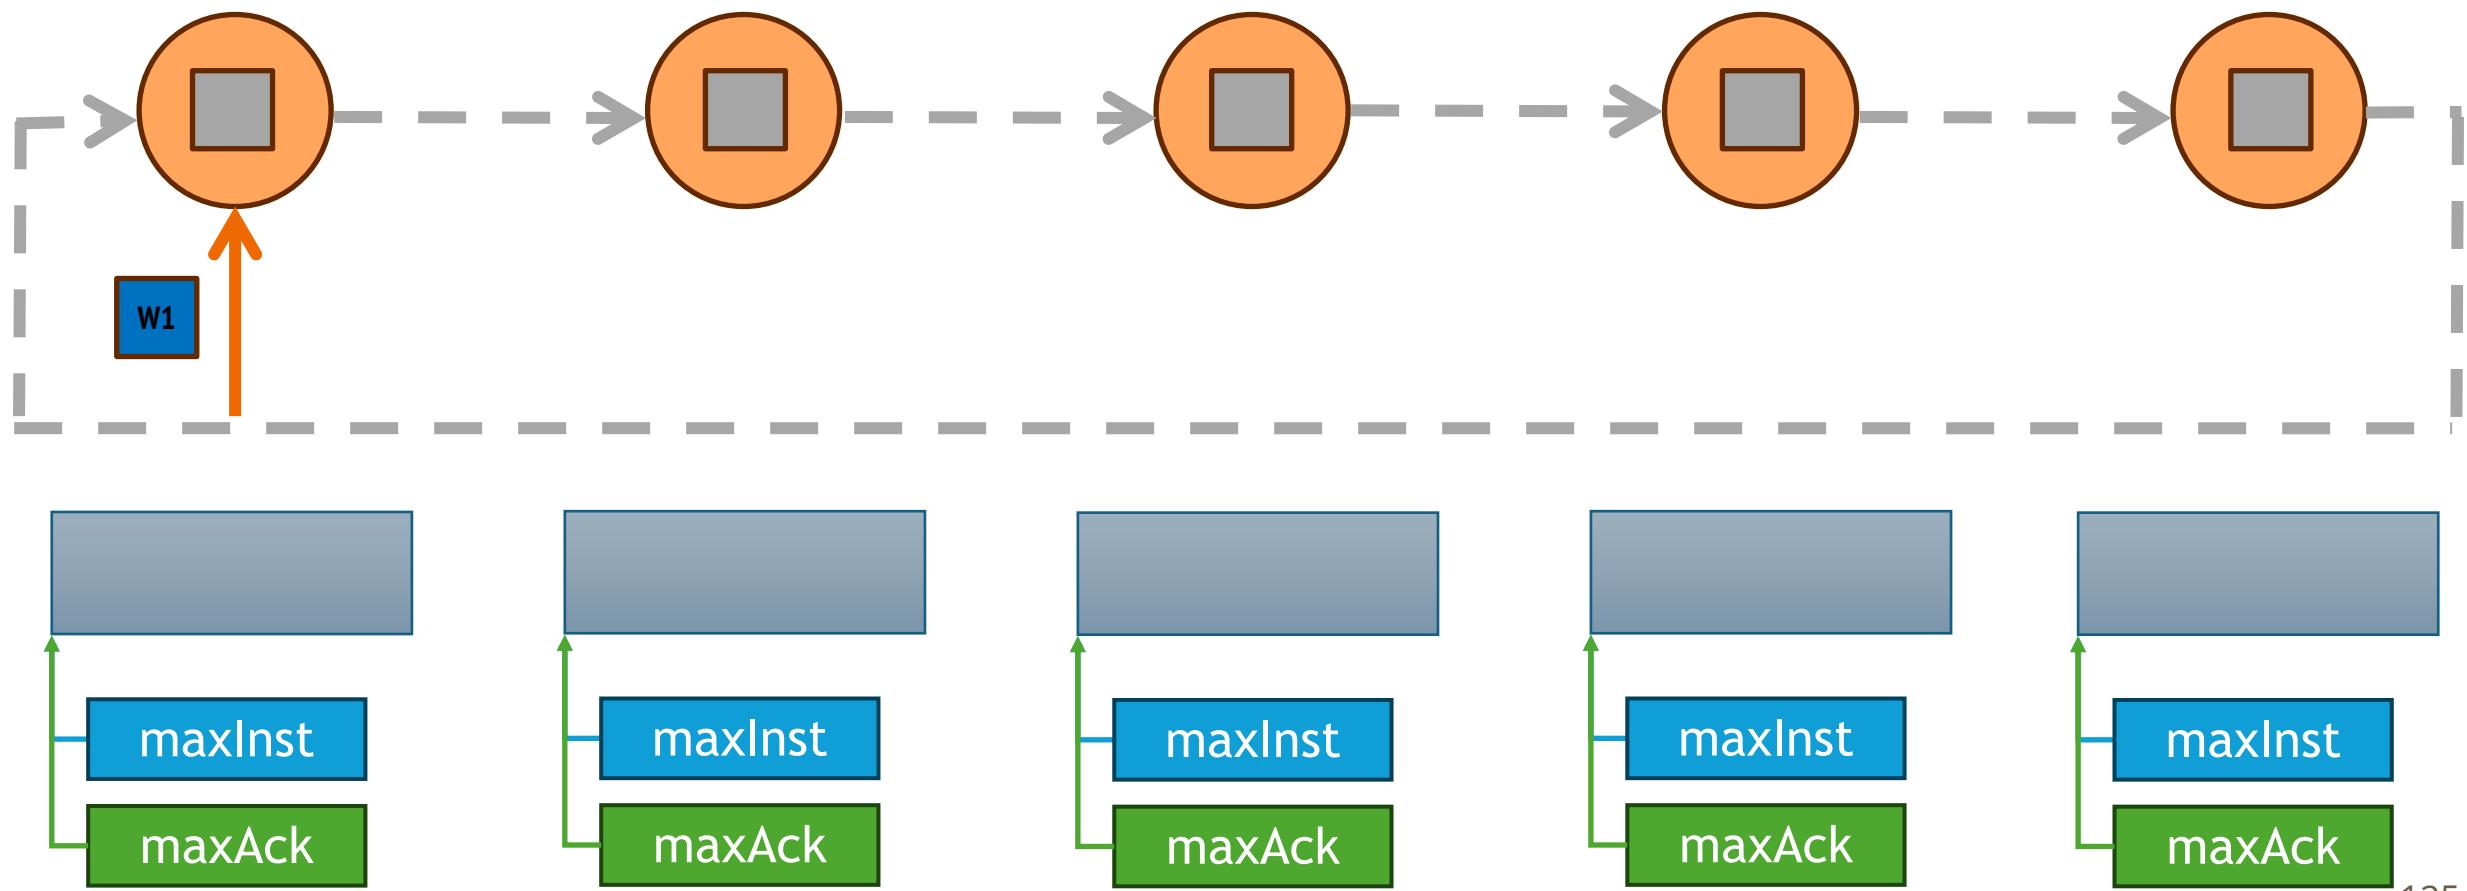
In Depth: The Log Representation



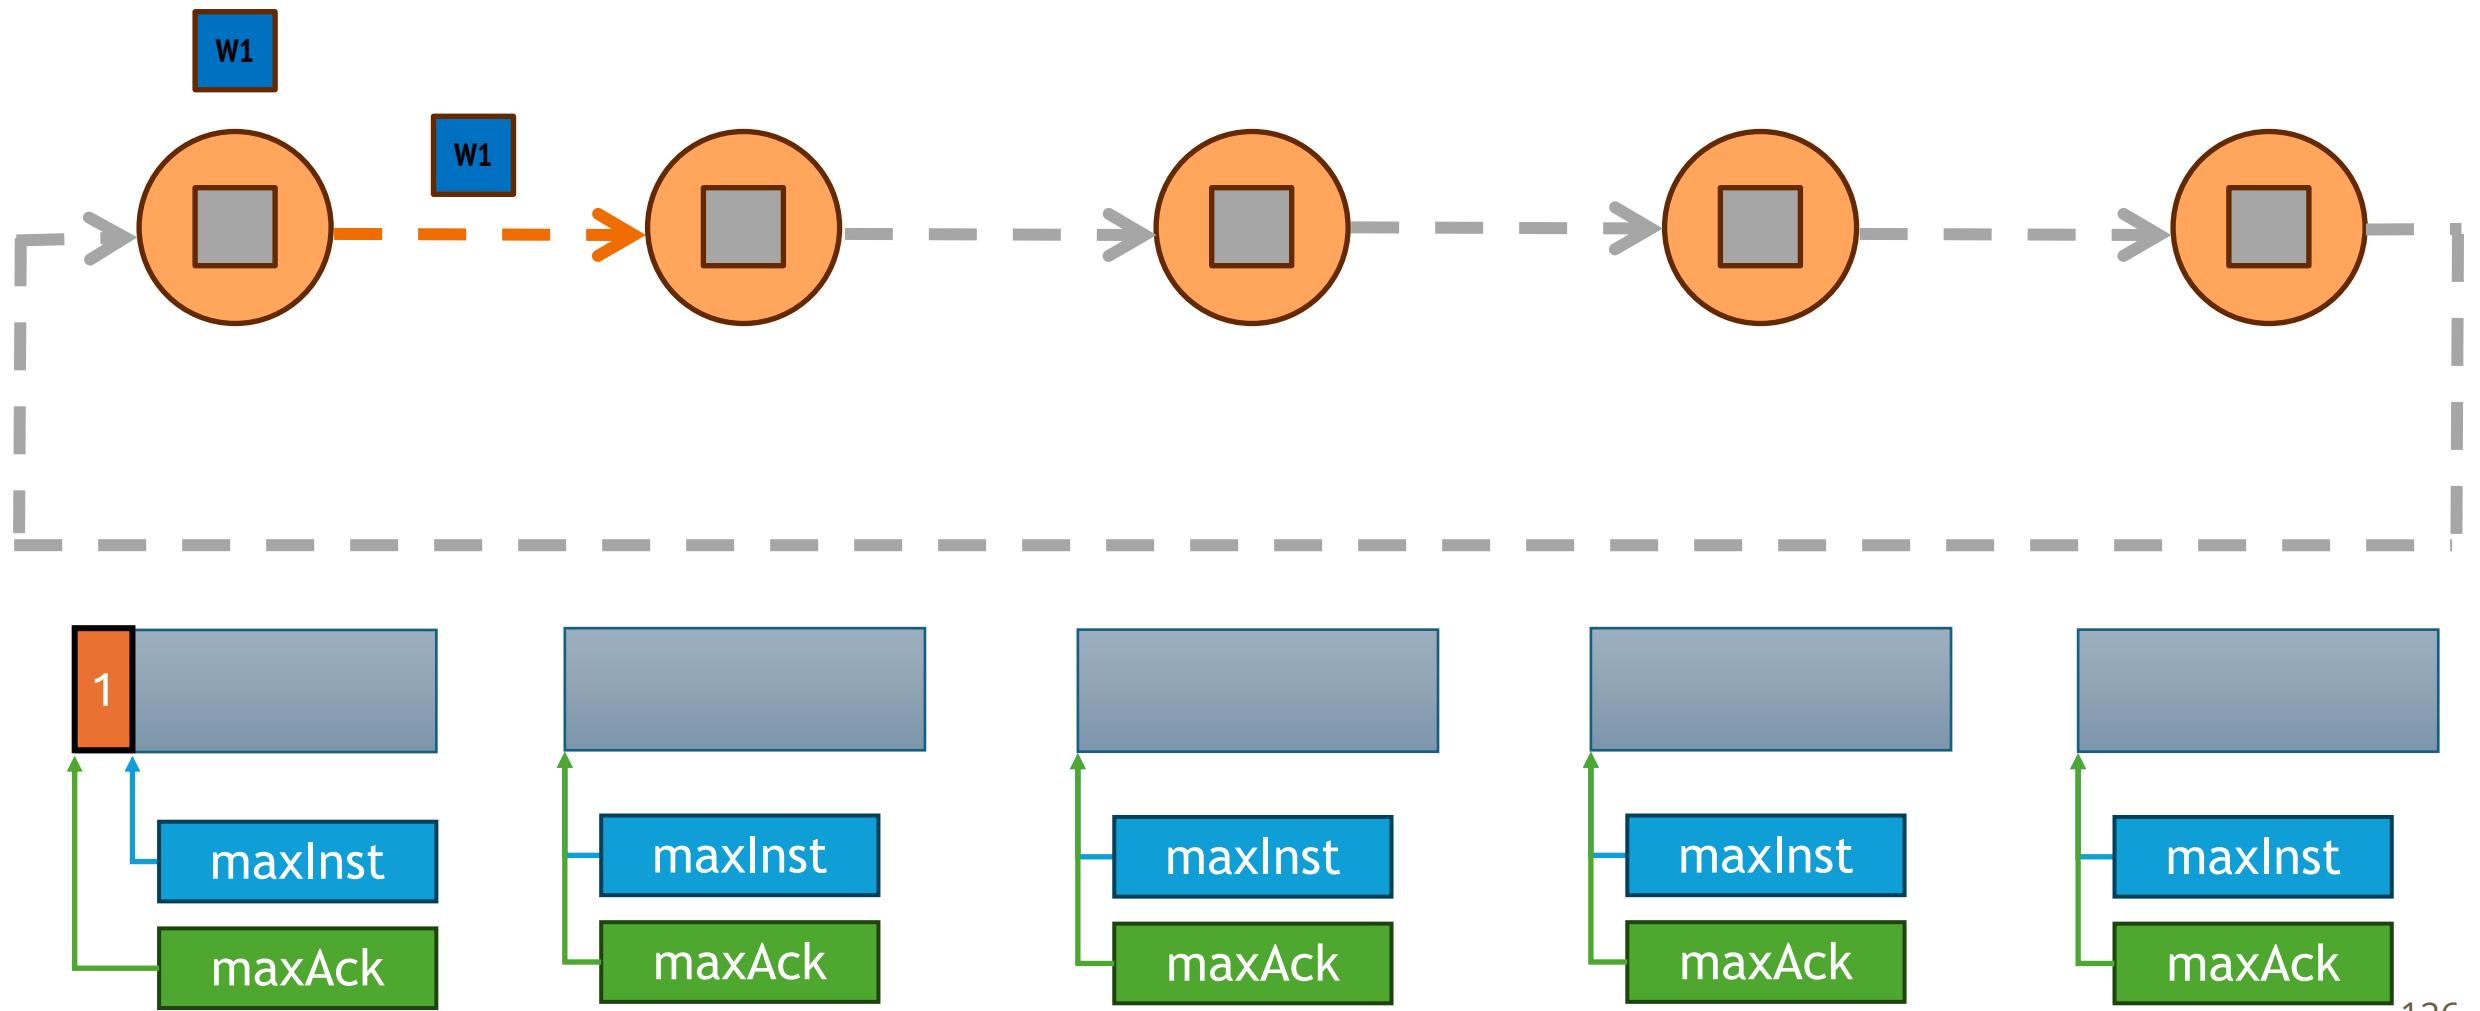
In Depth: The Log Representation



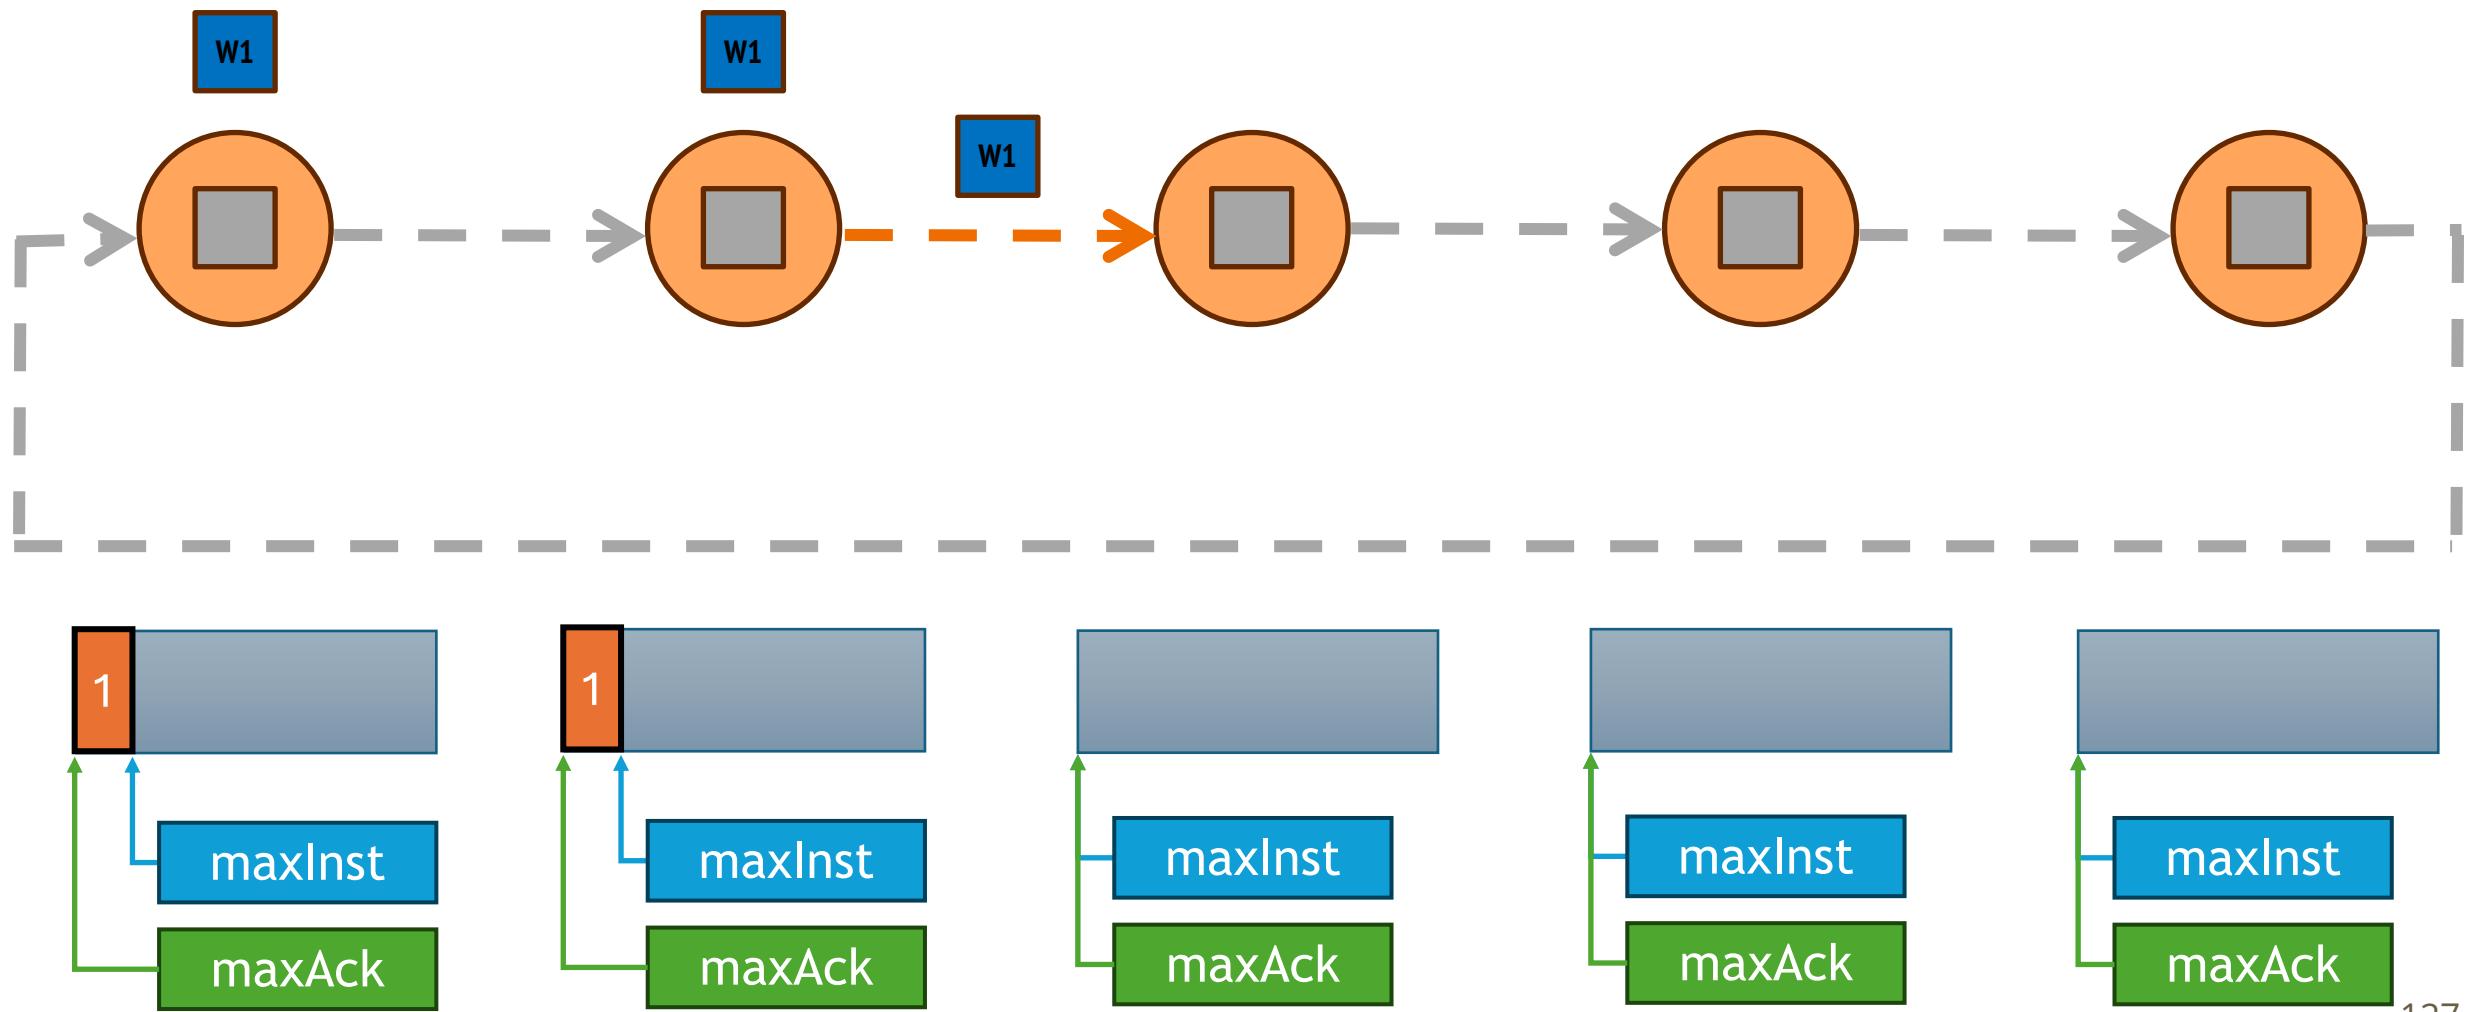
In Depth: Writes



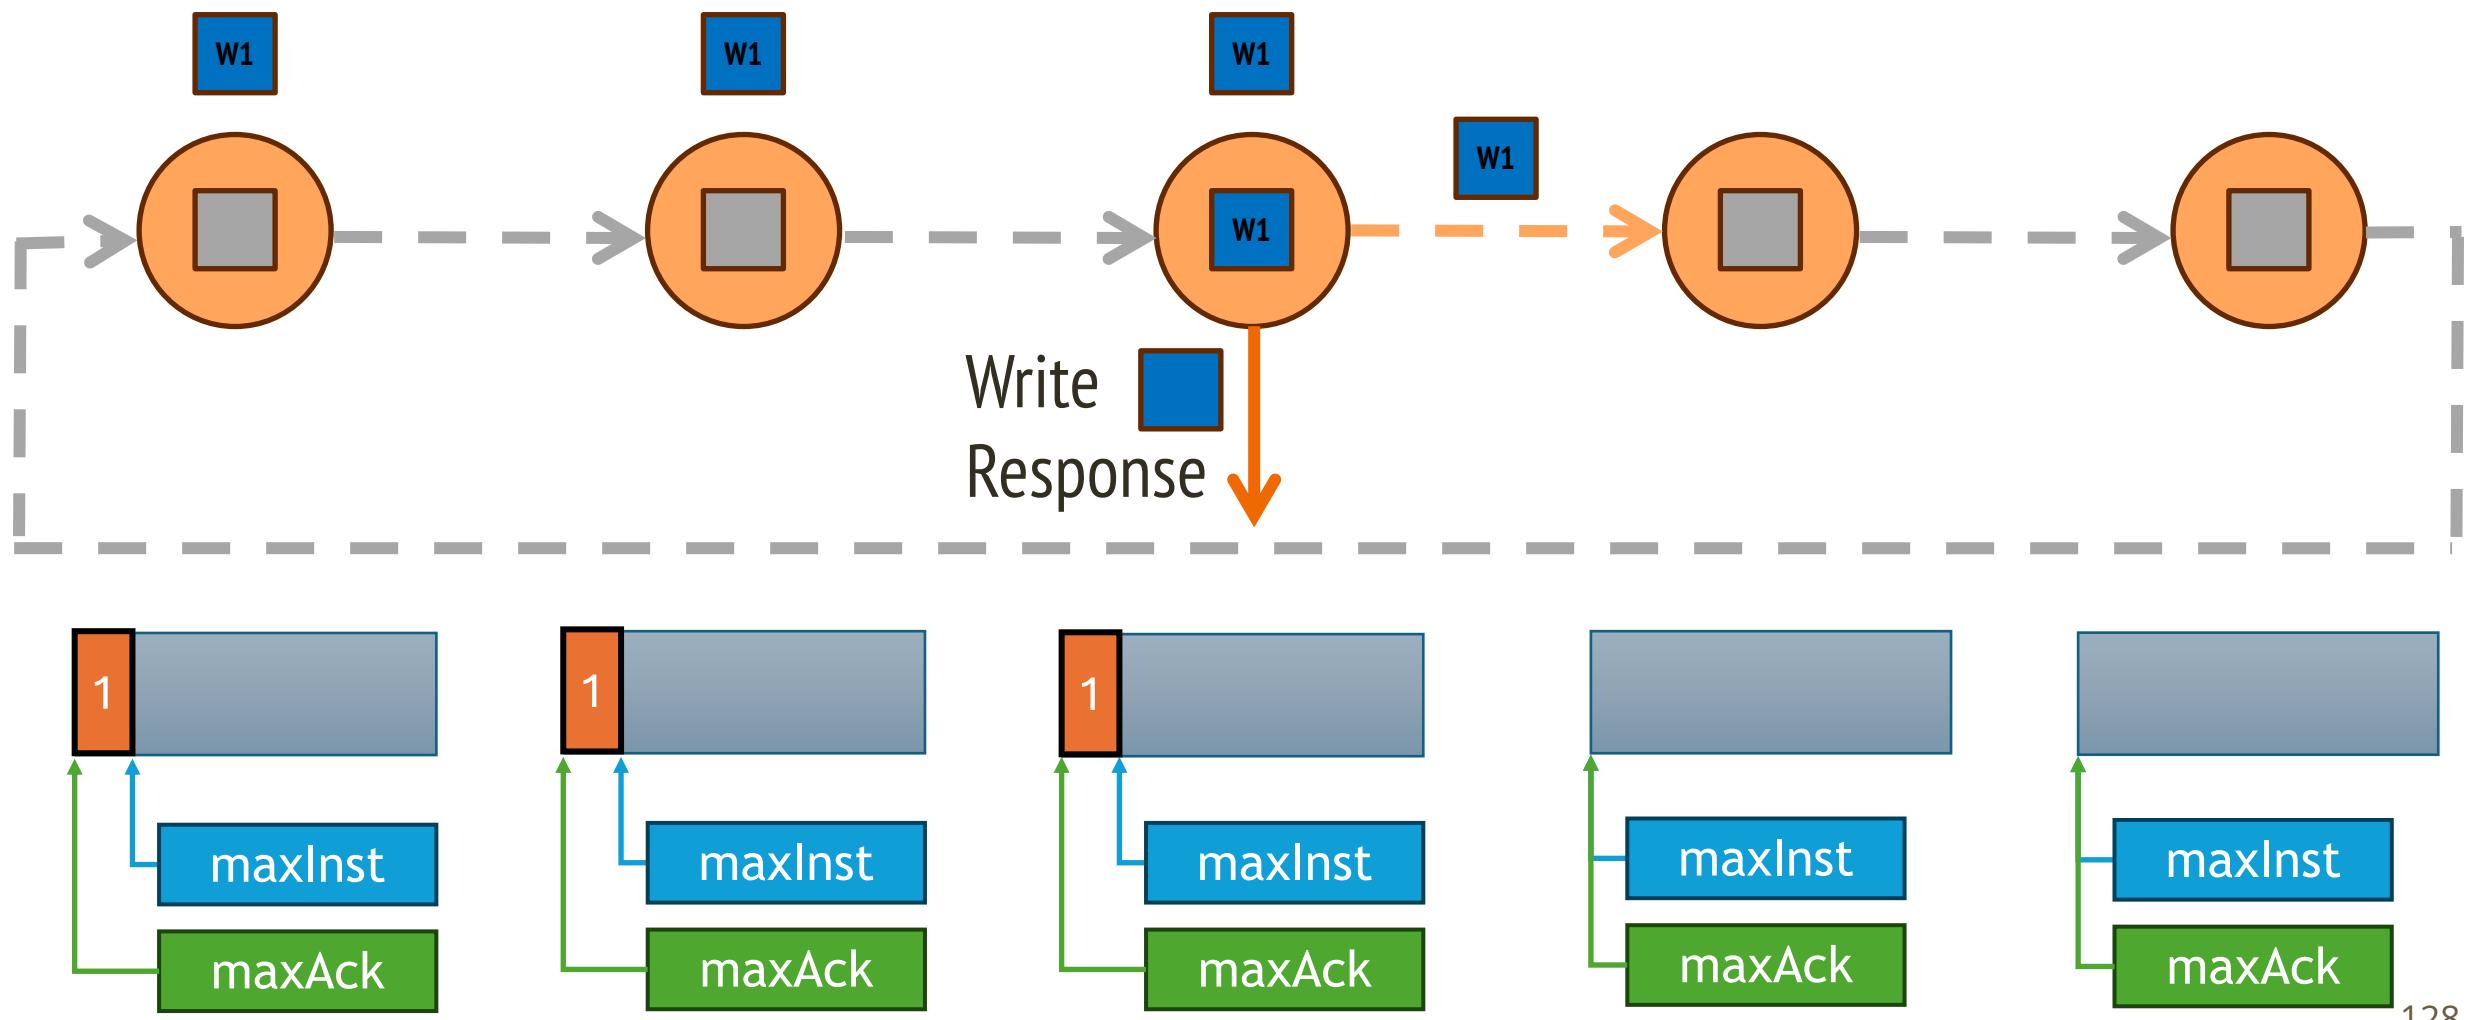
In Depth: Writes



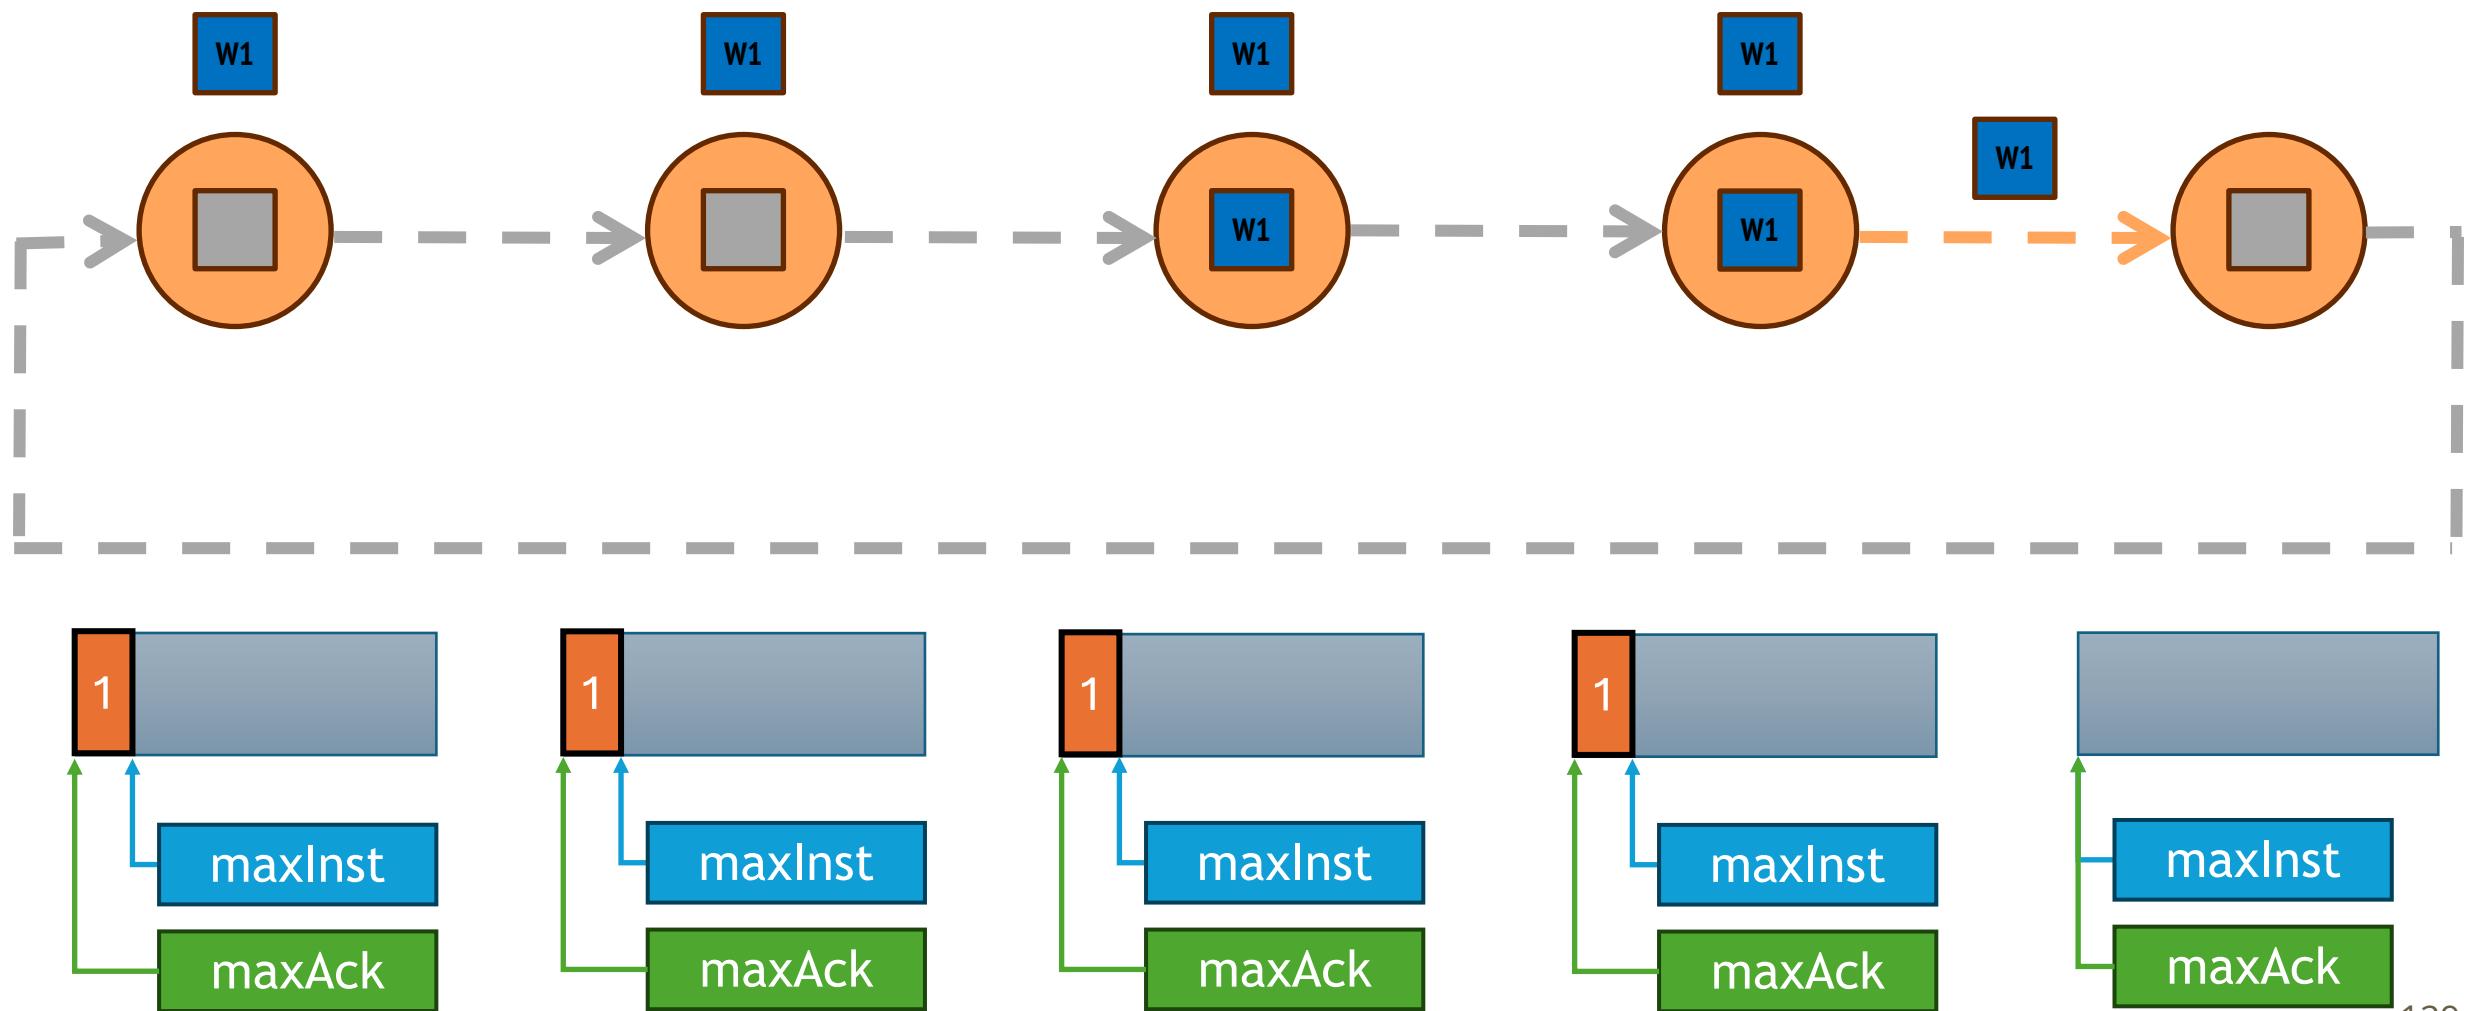
In Depth: Writes



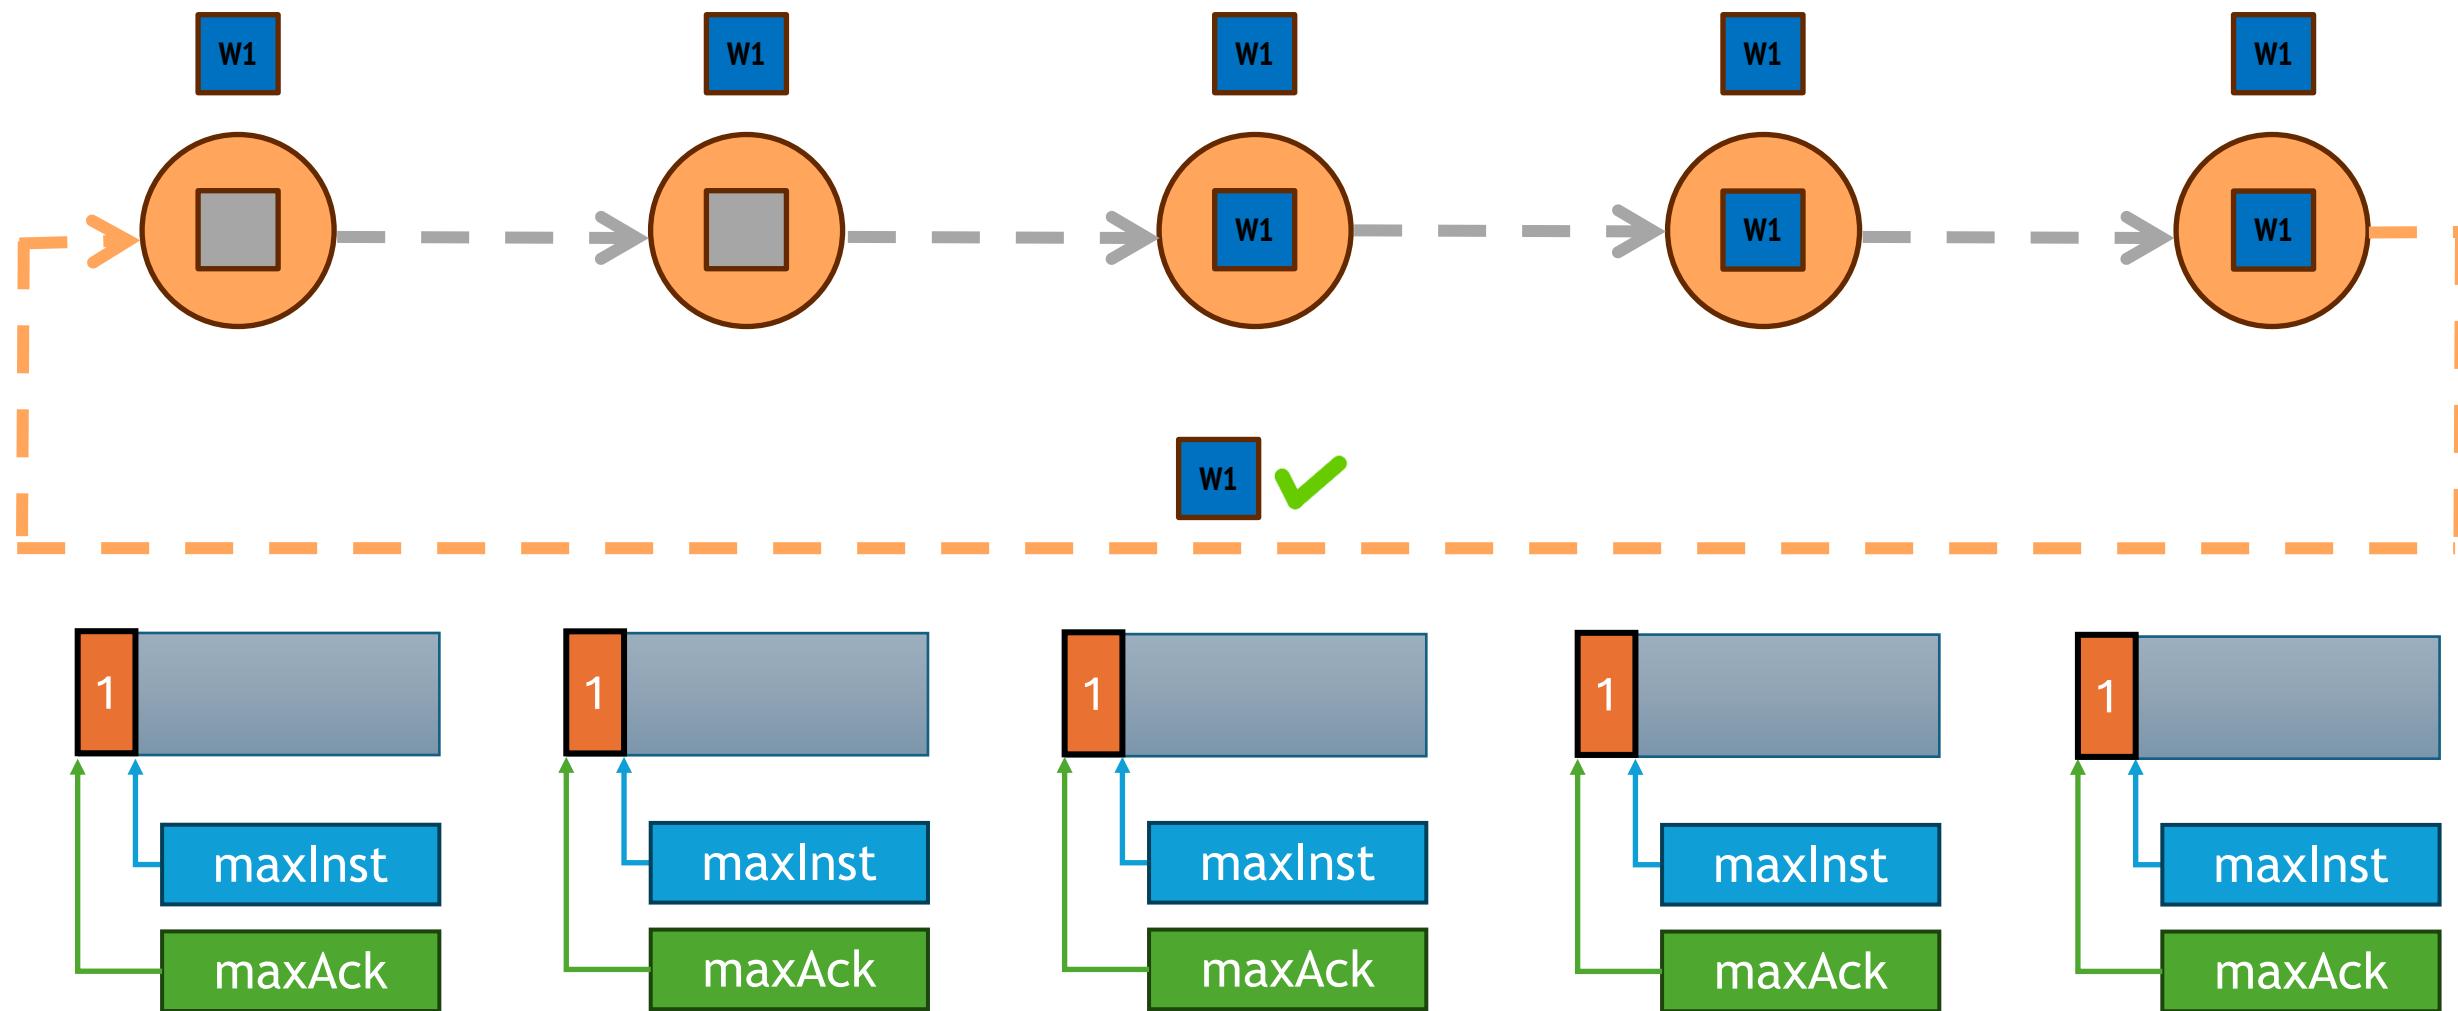
In Depth: Writes



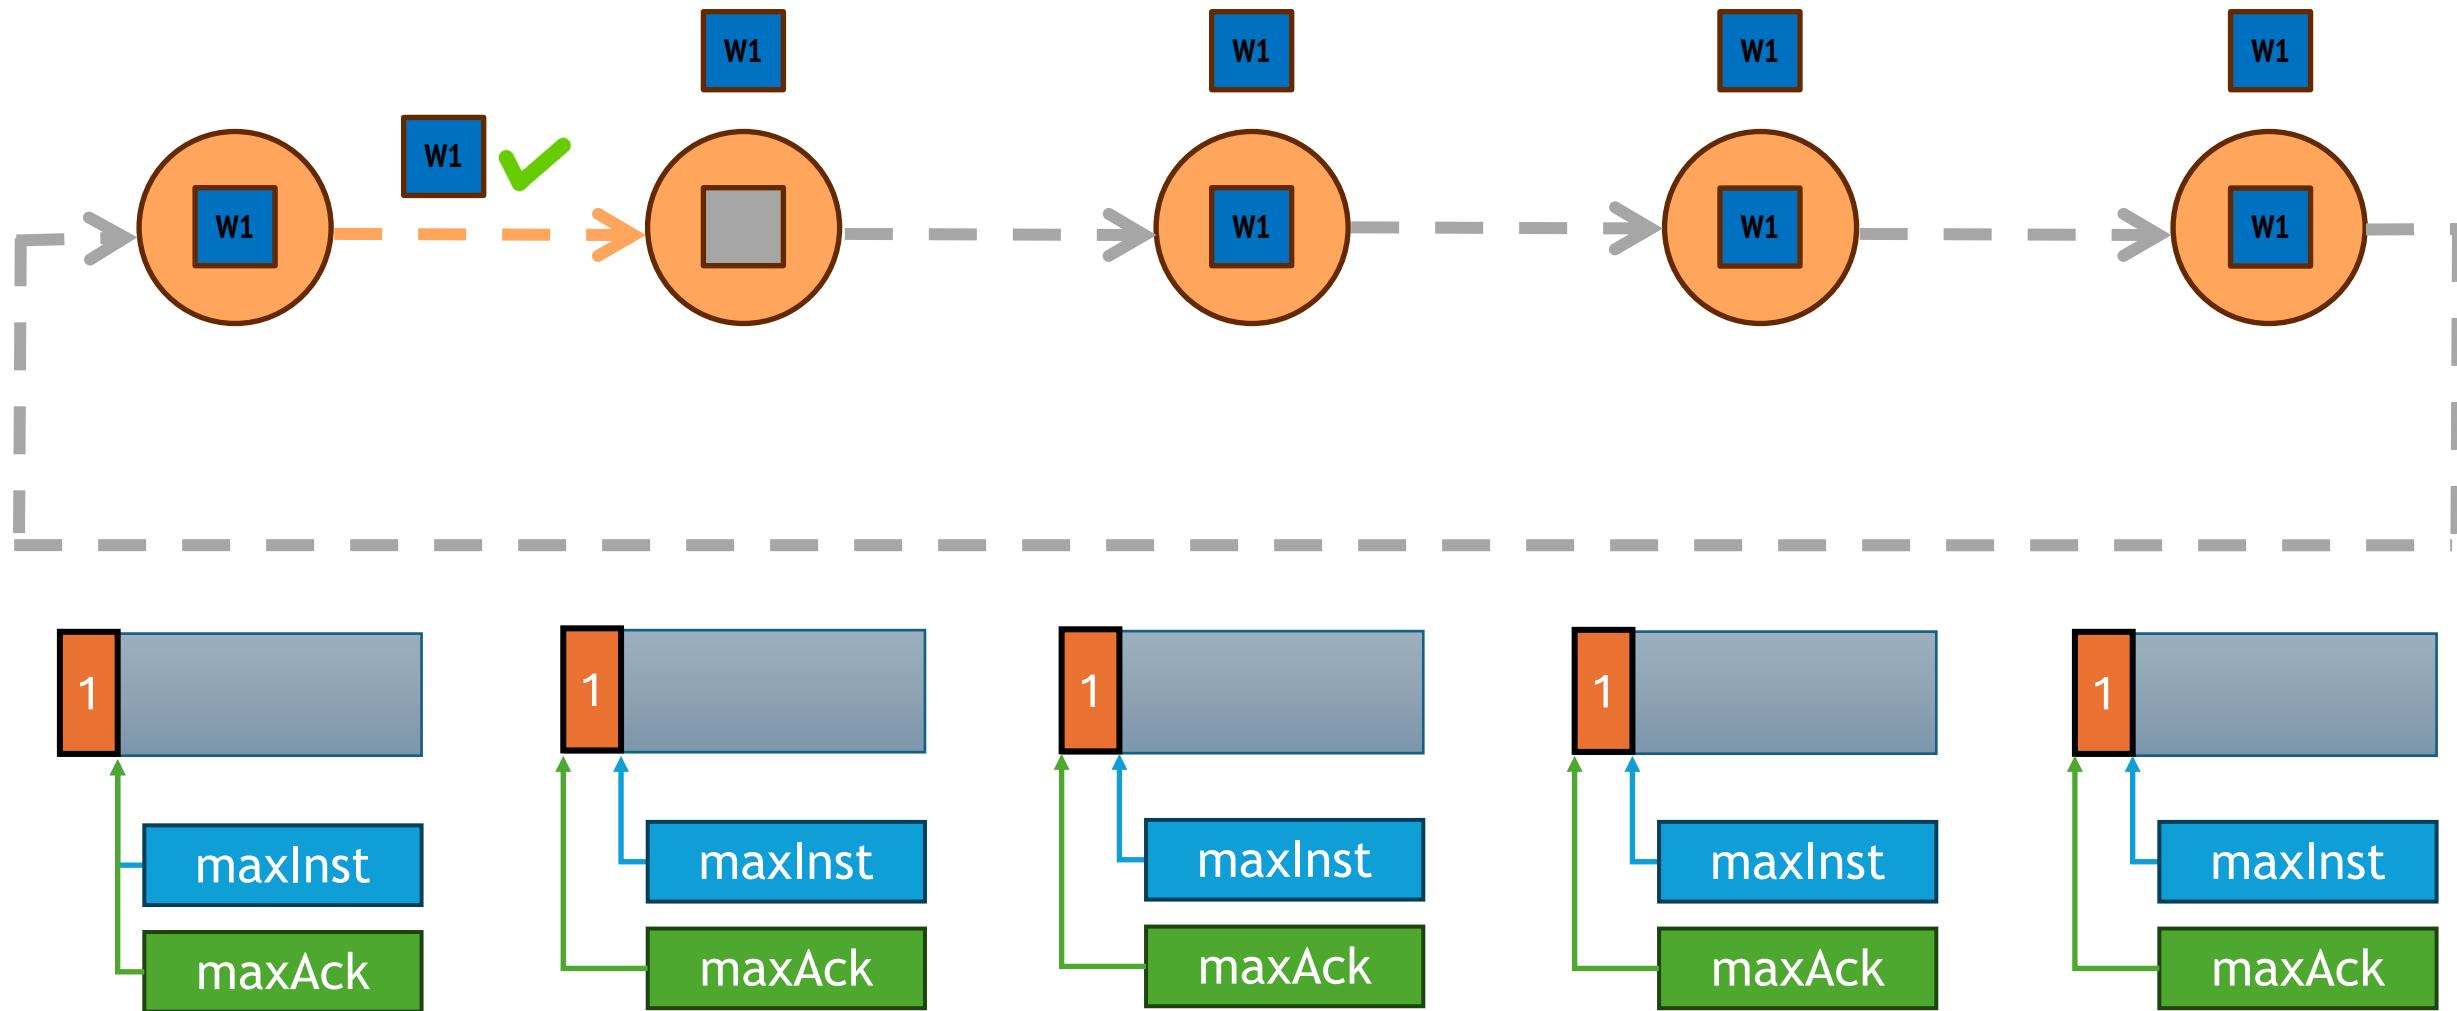
In Depth: Writes



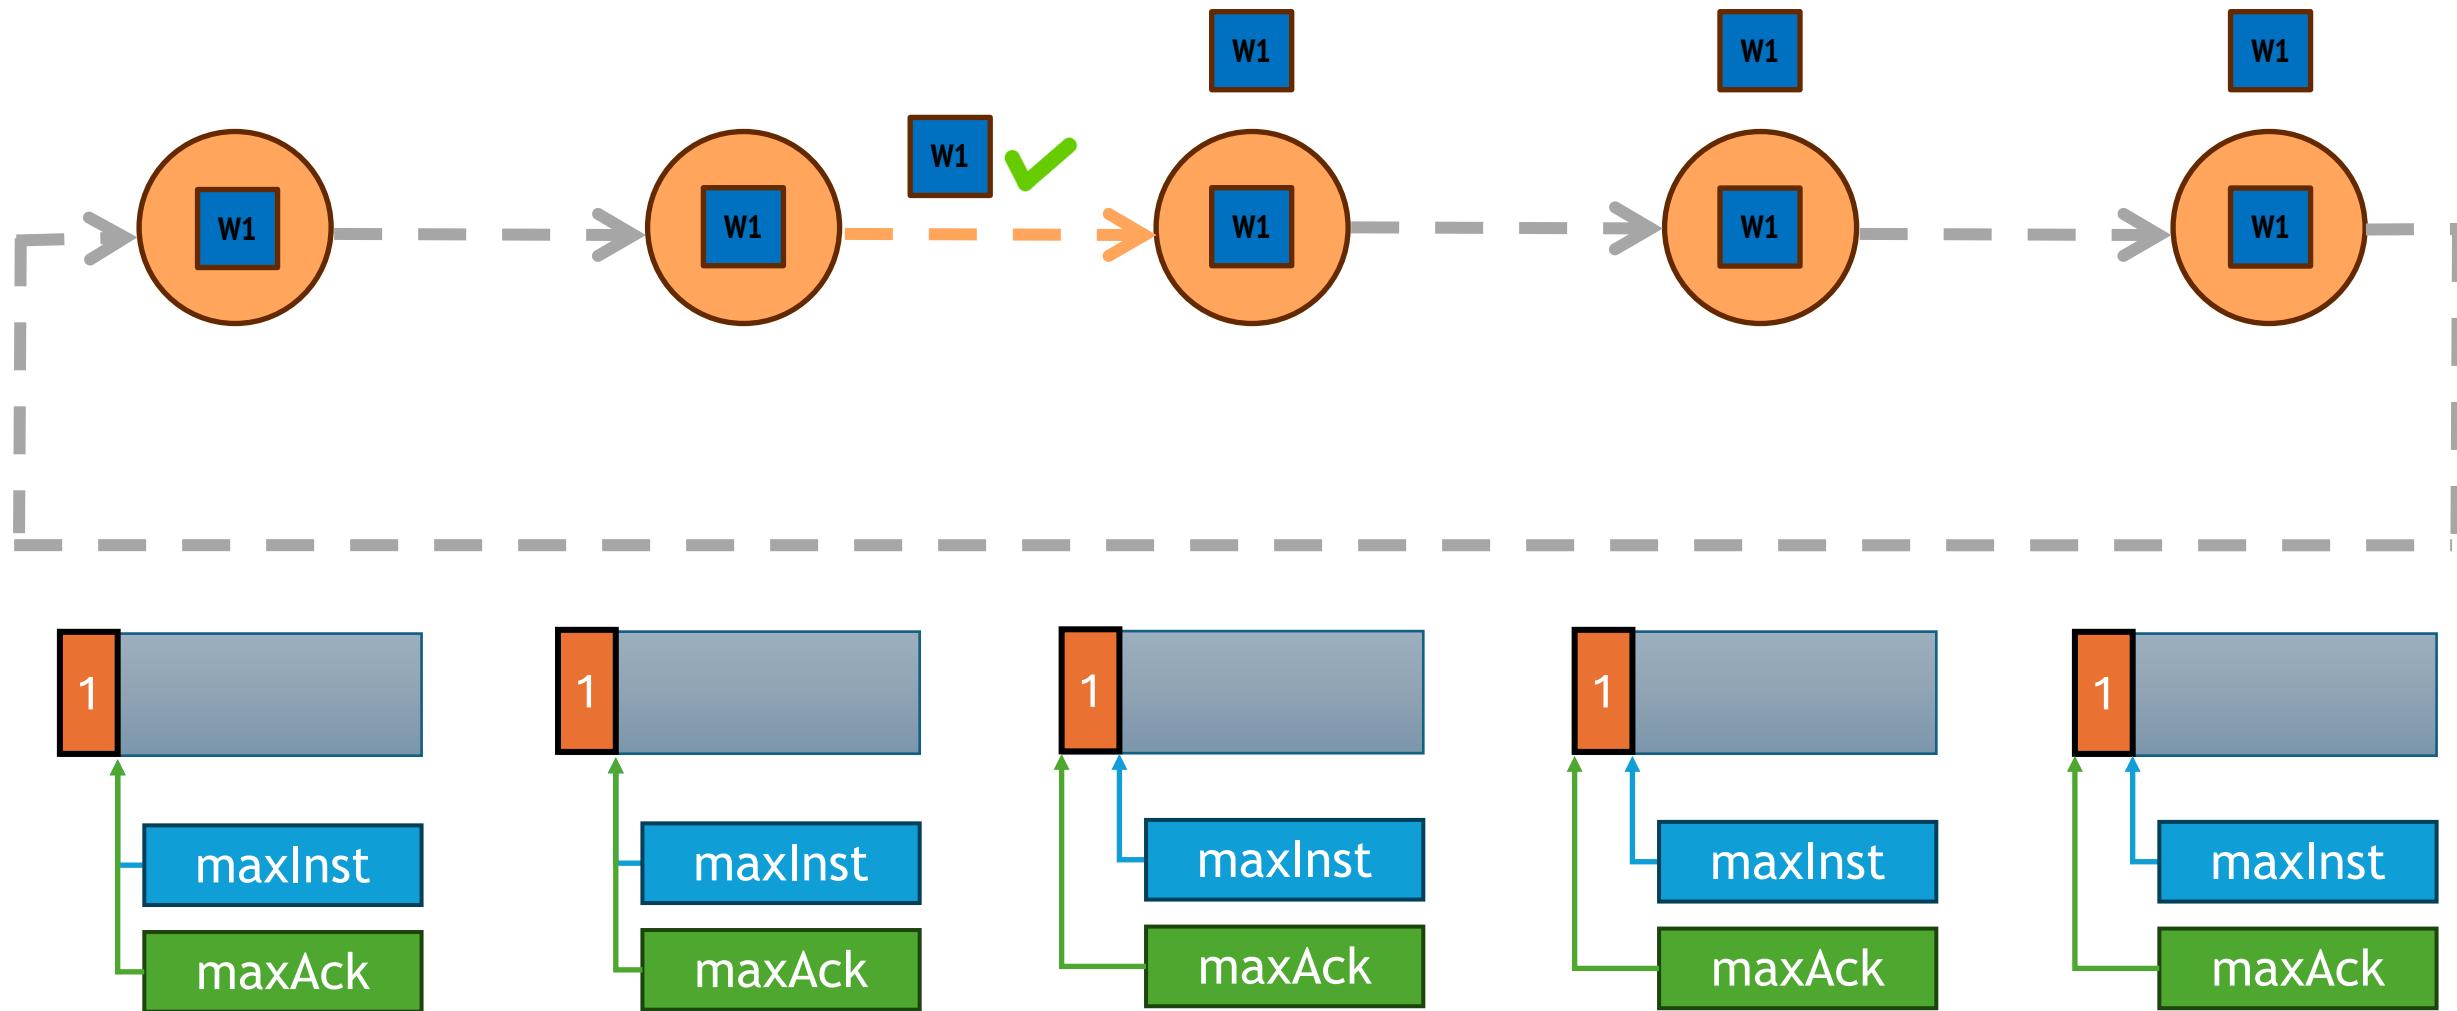
In Depth: Writes



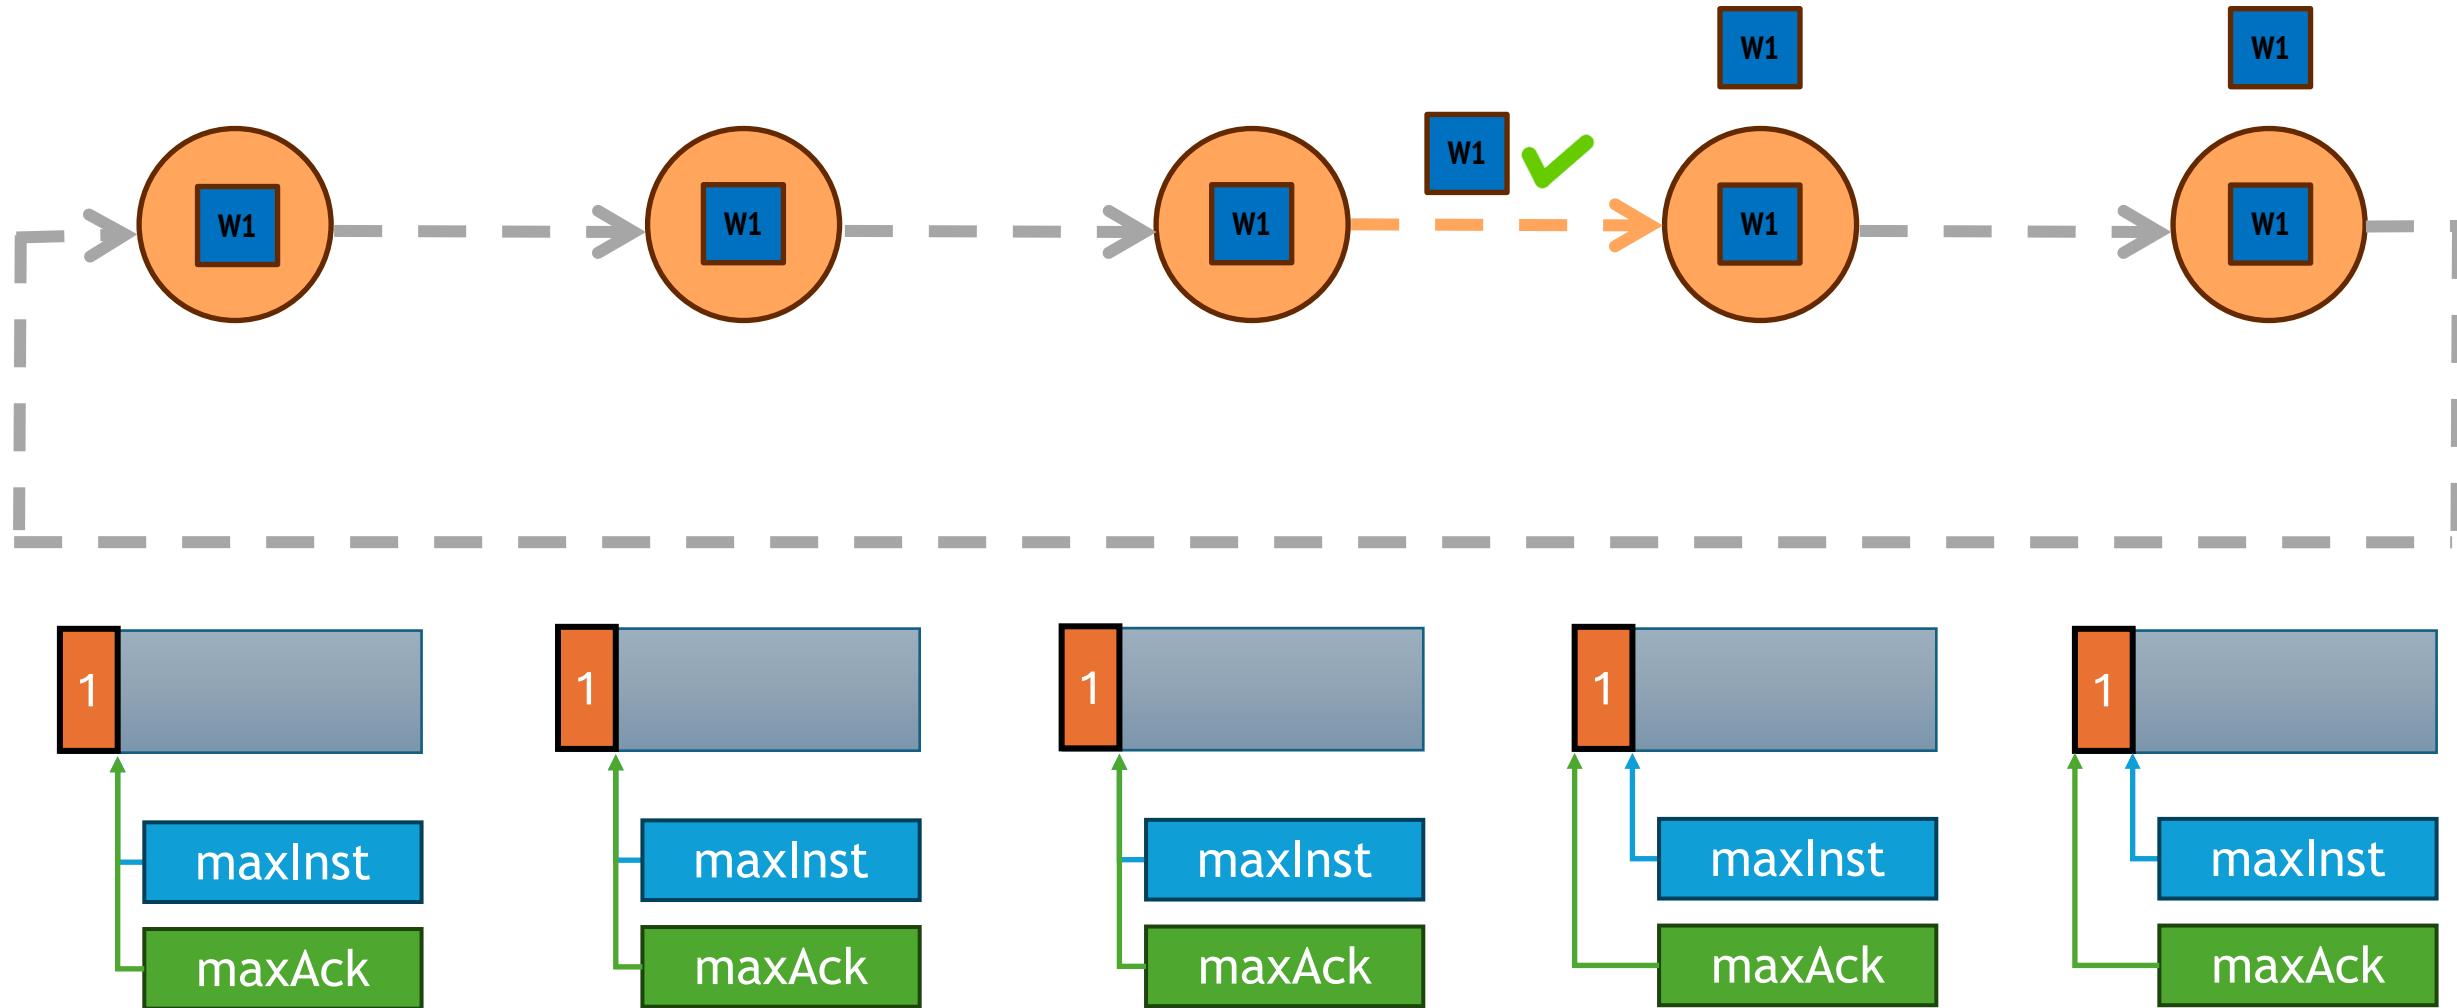
In Depth: Writes



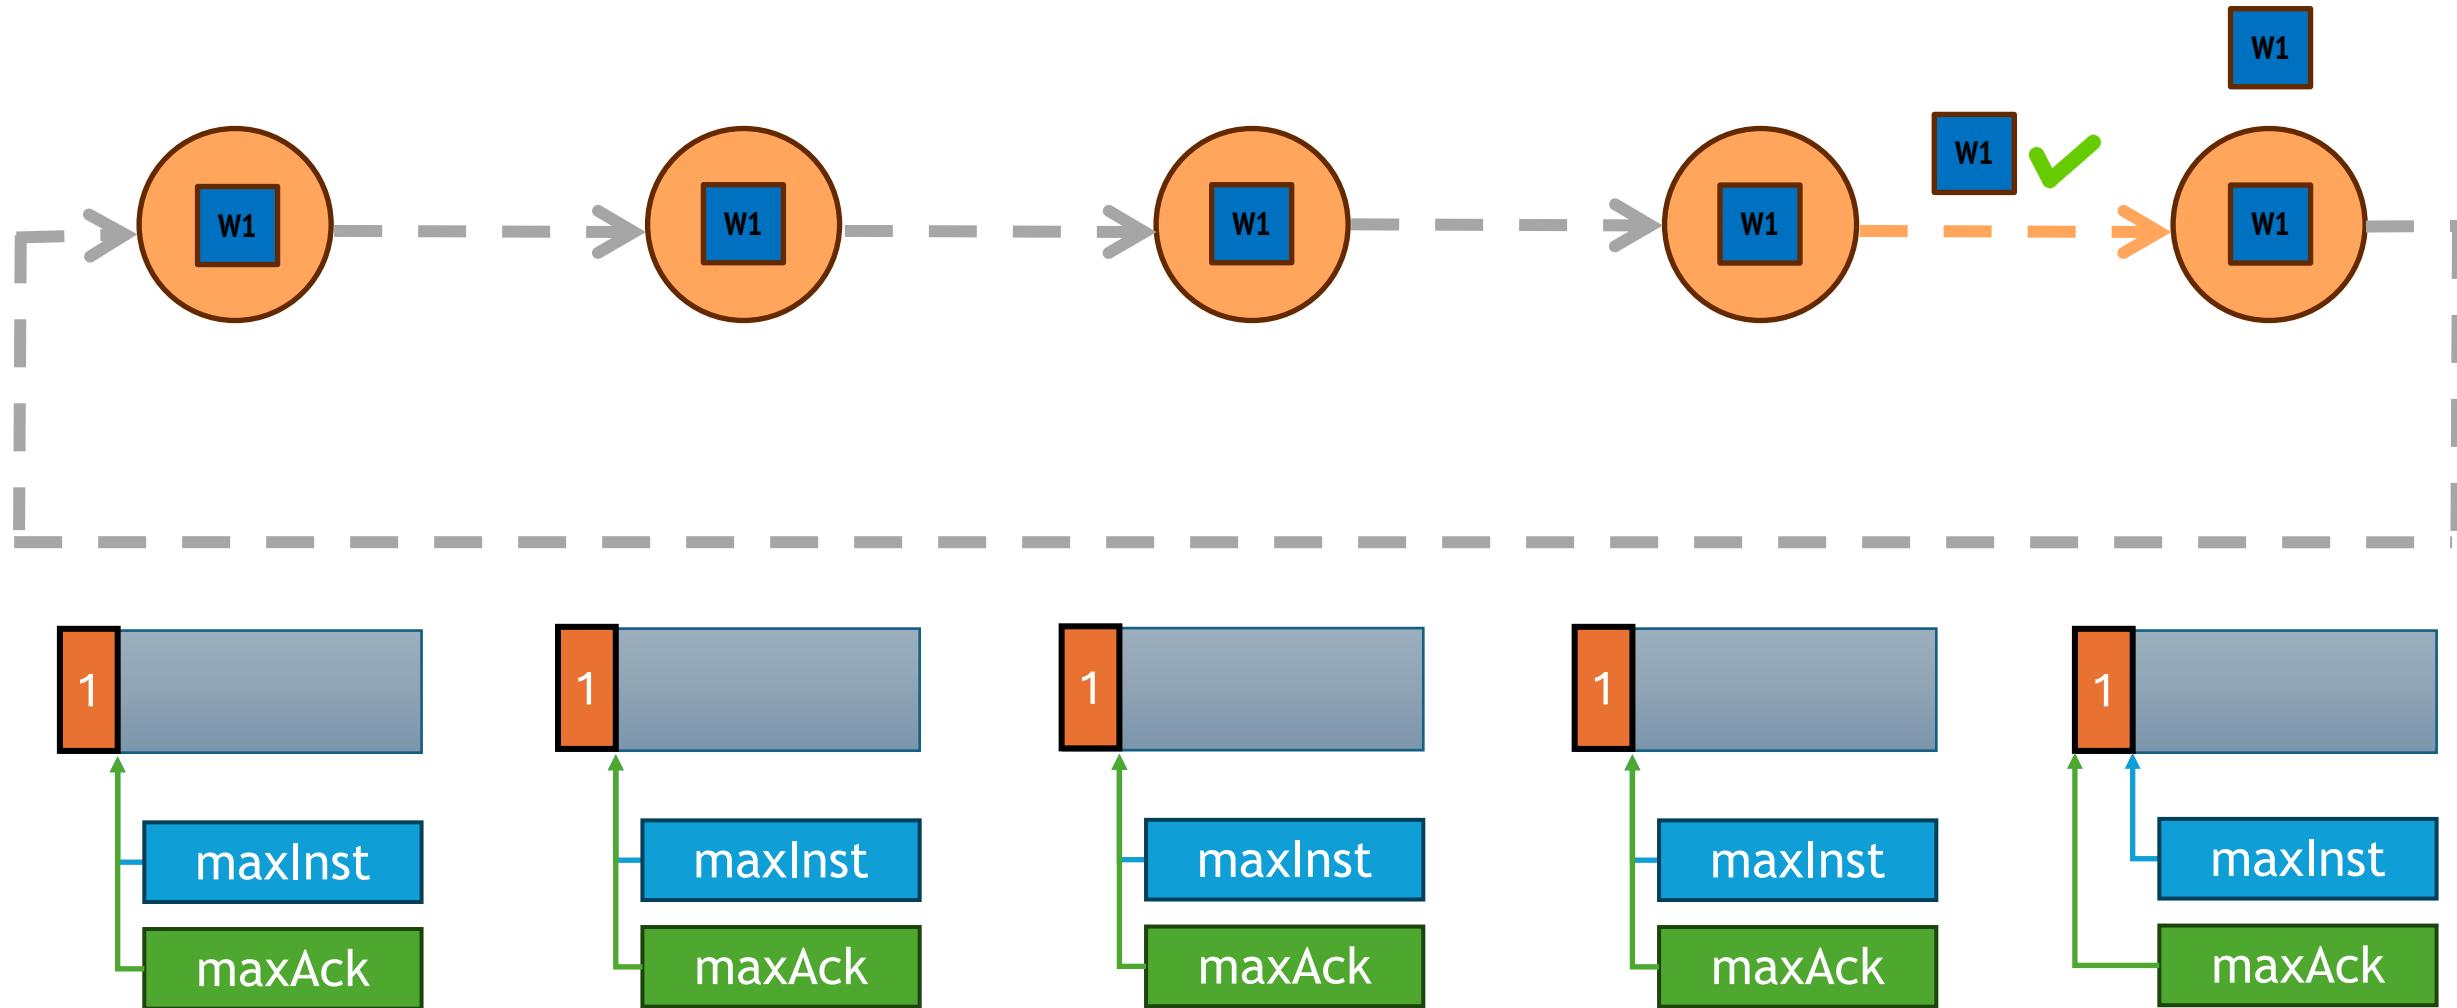
In Depth: Writes



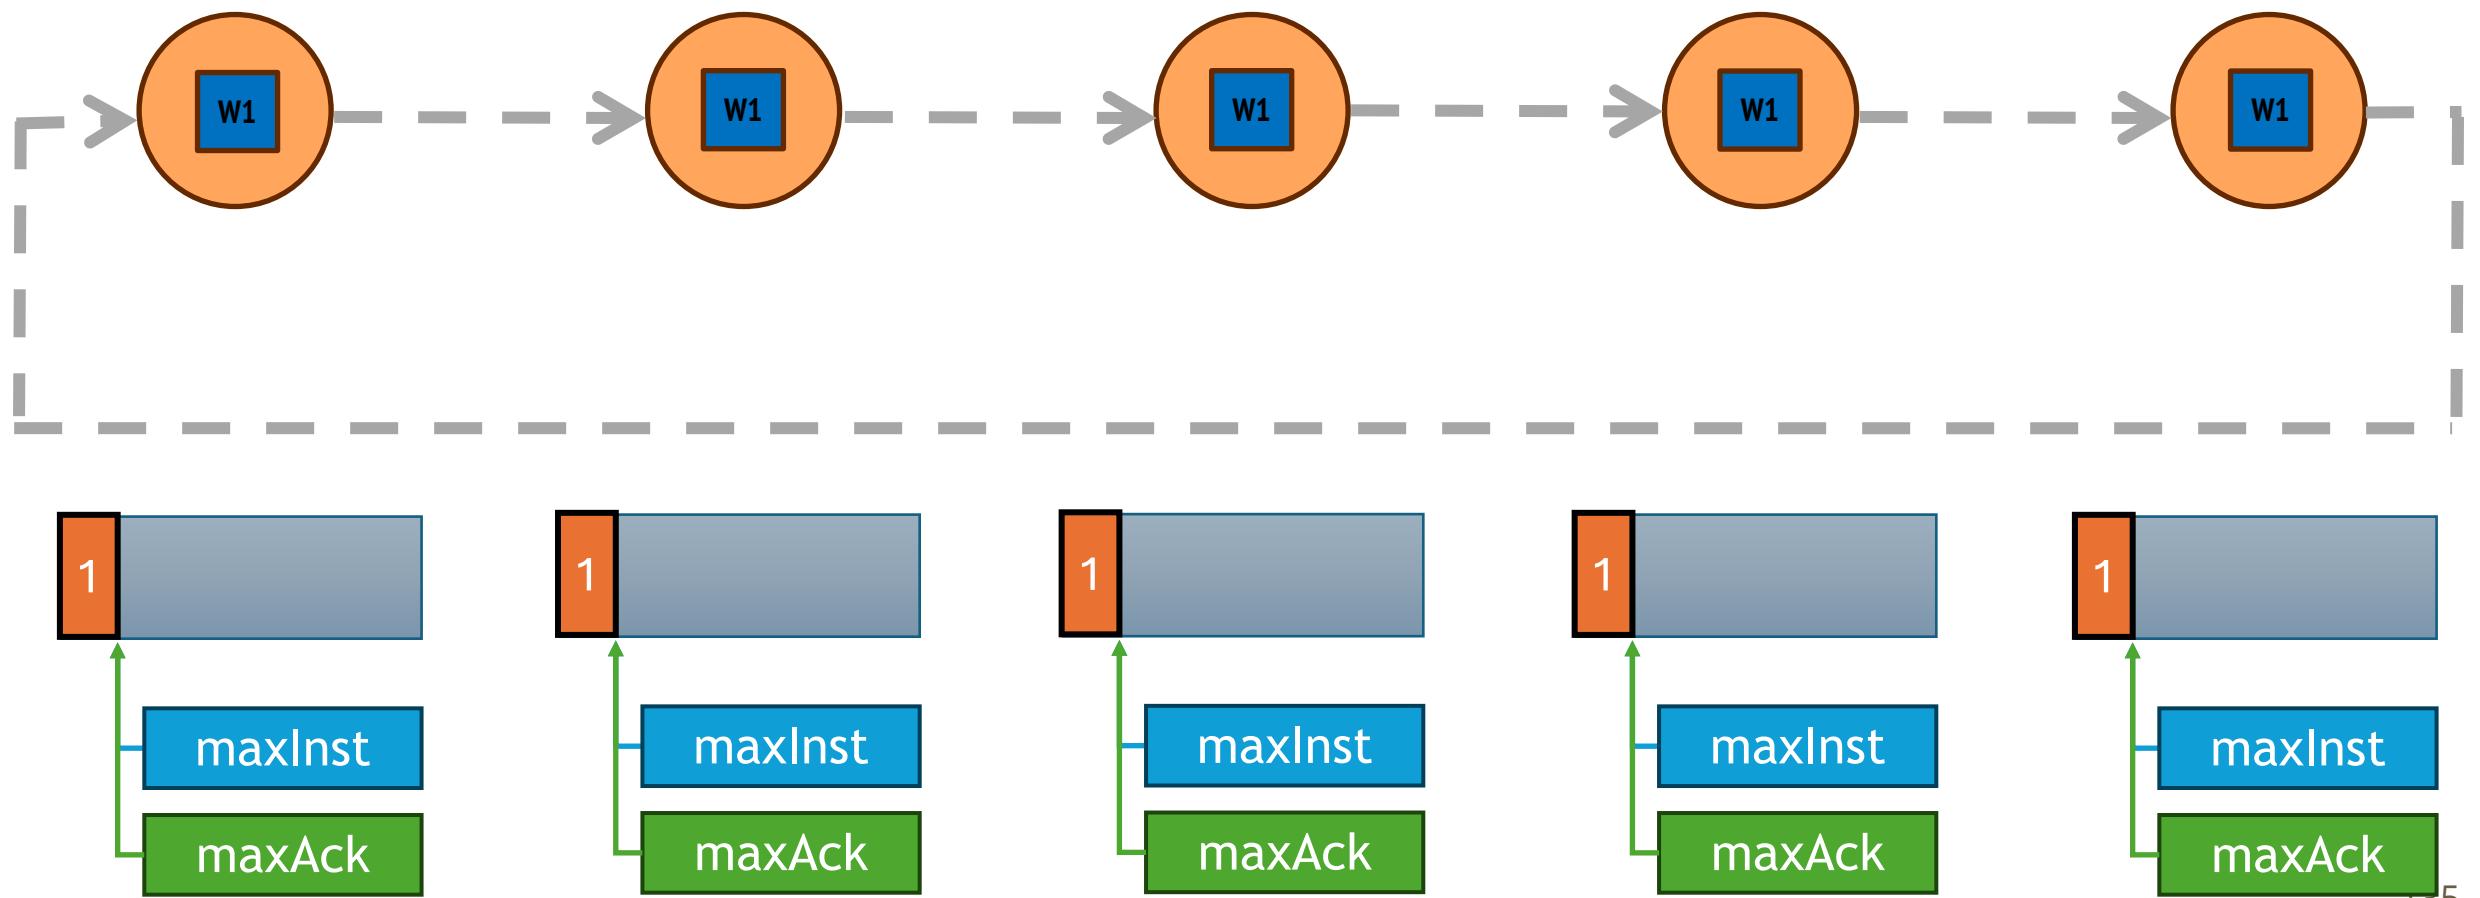
In Depth: Writes



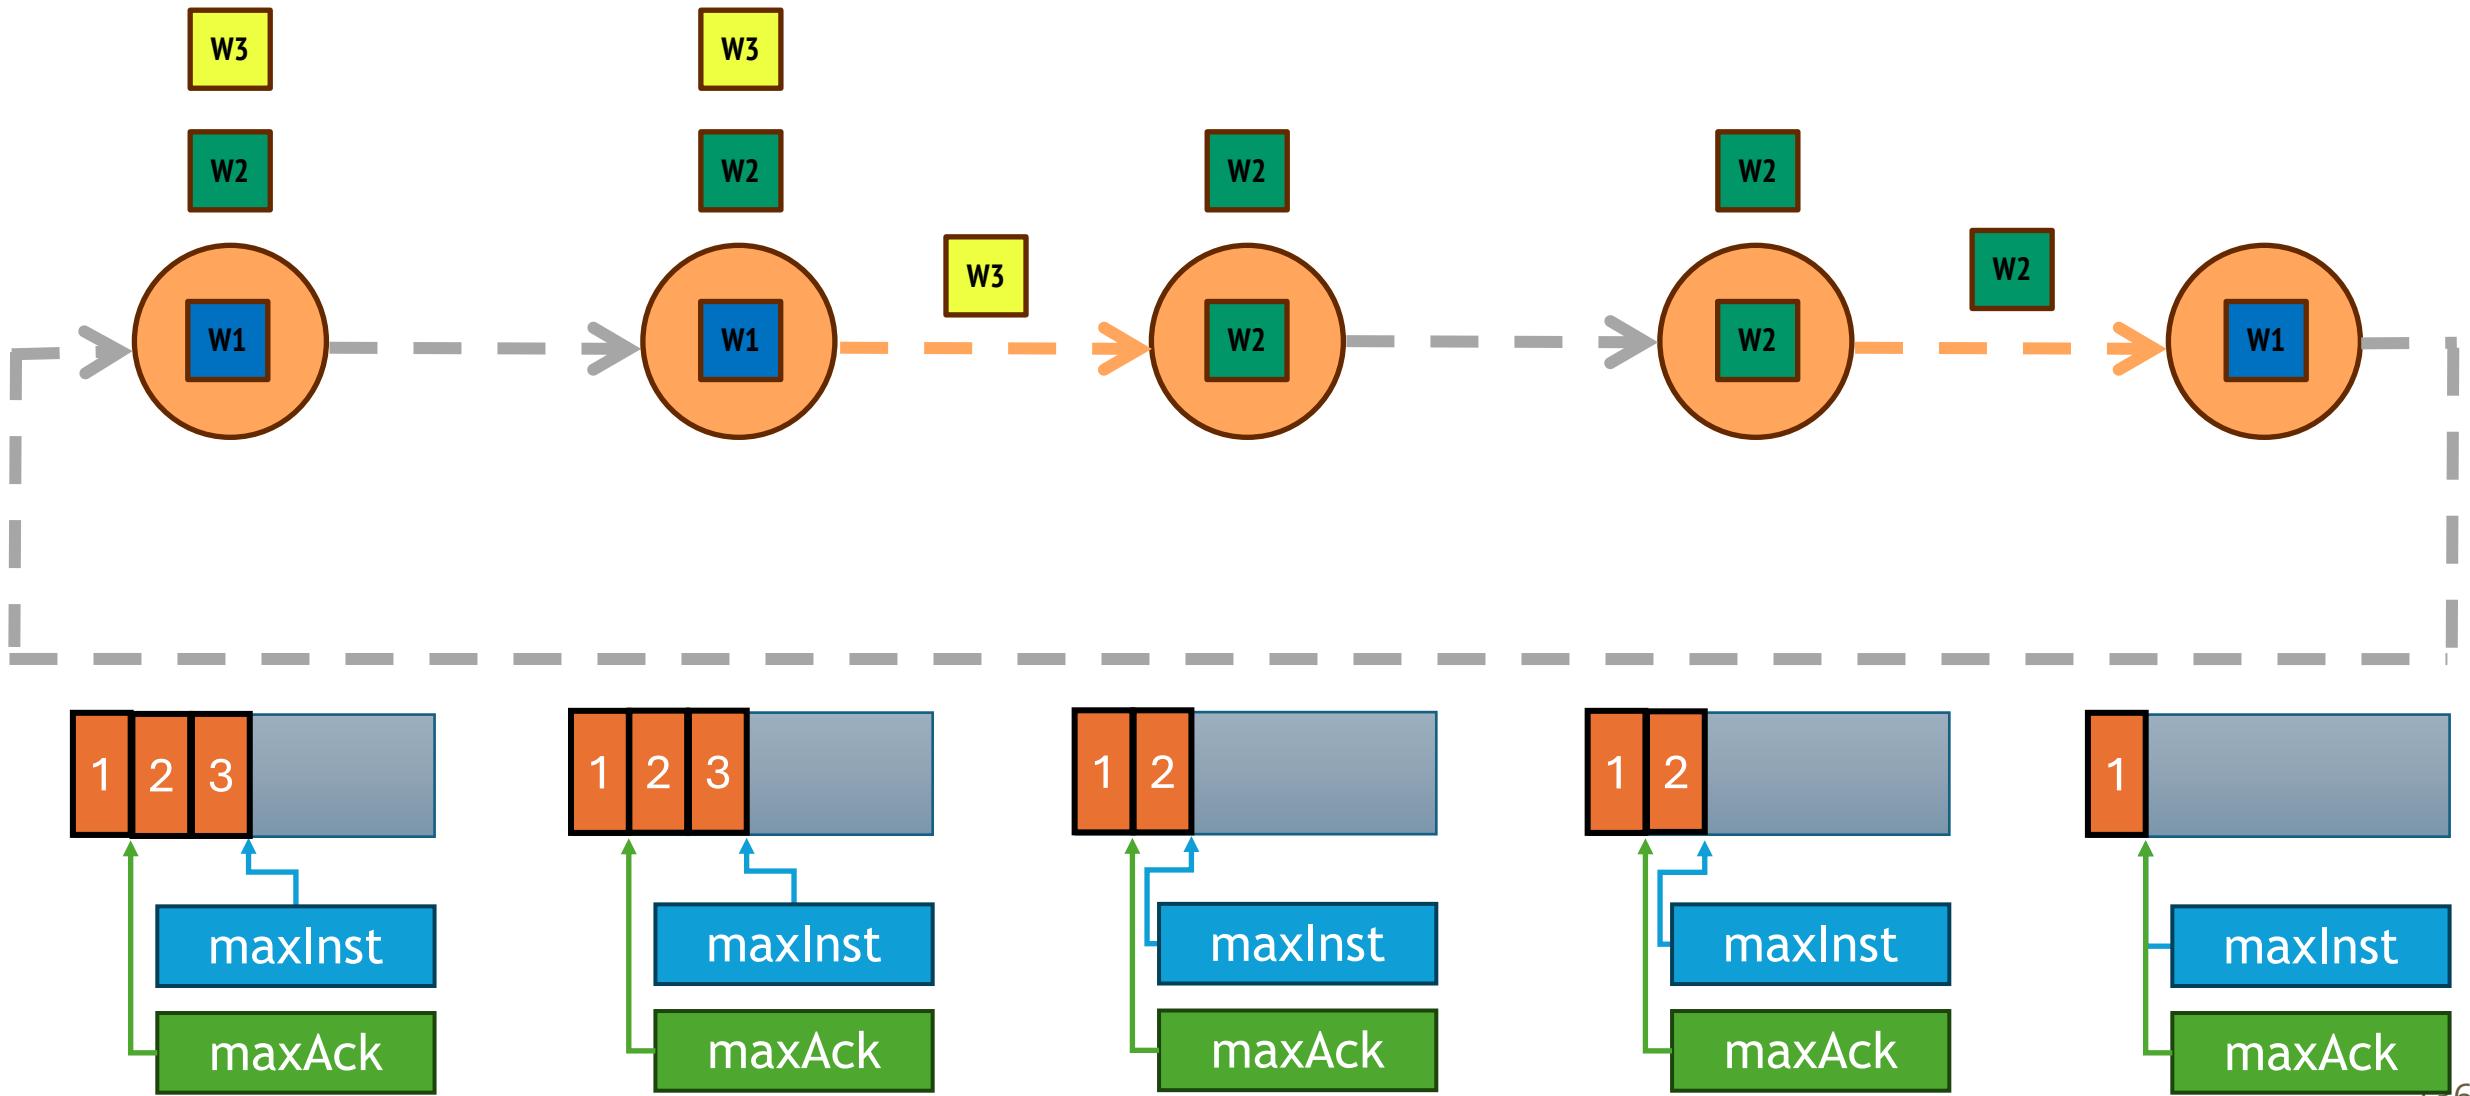
In Depth: Writes



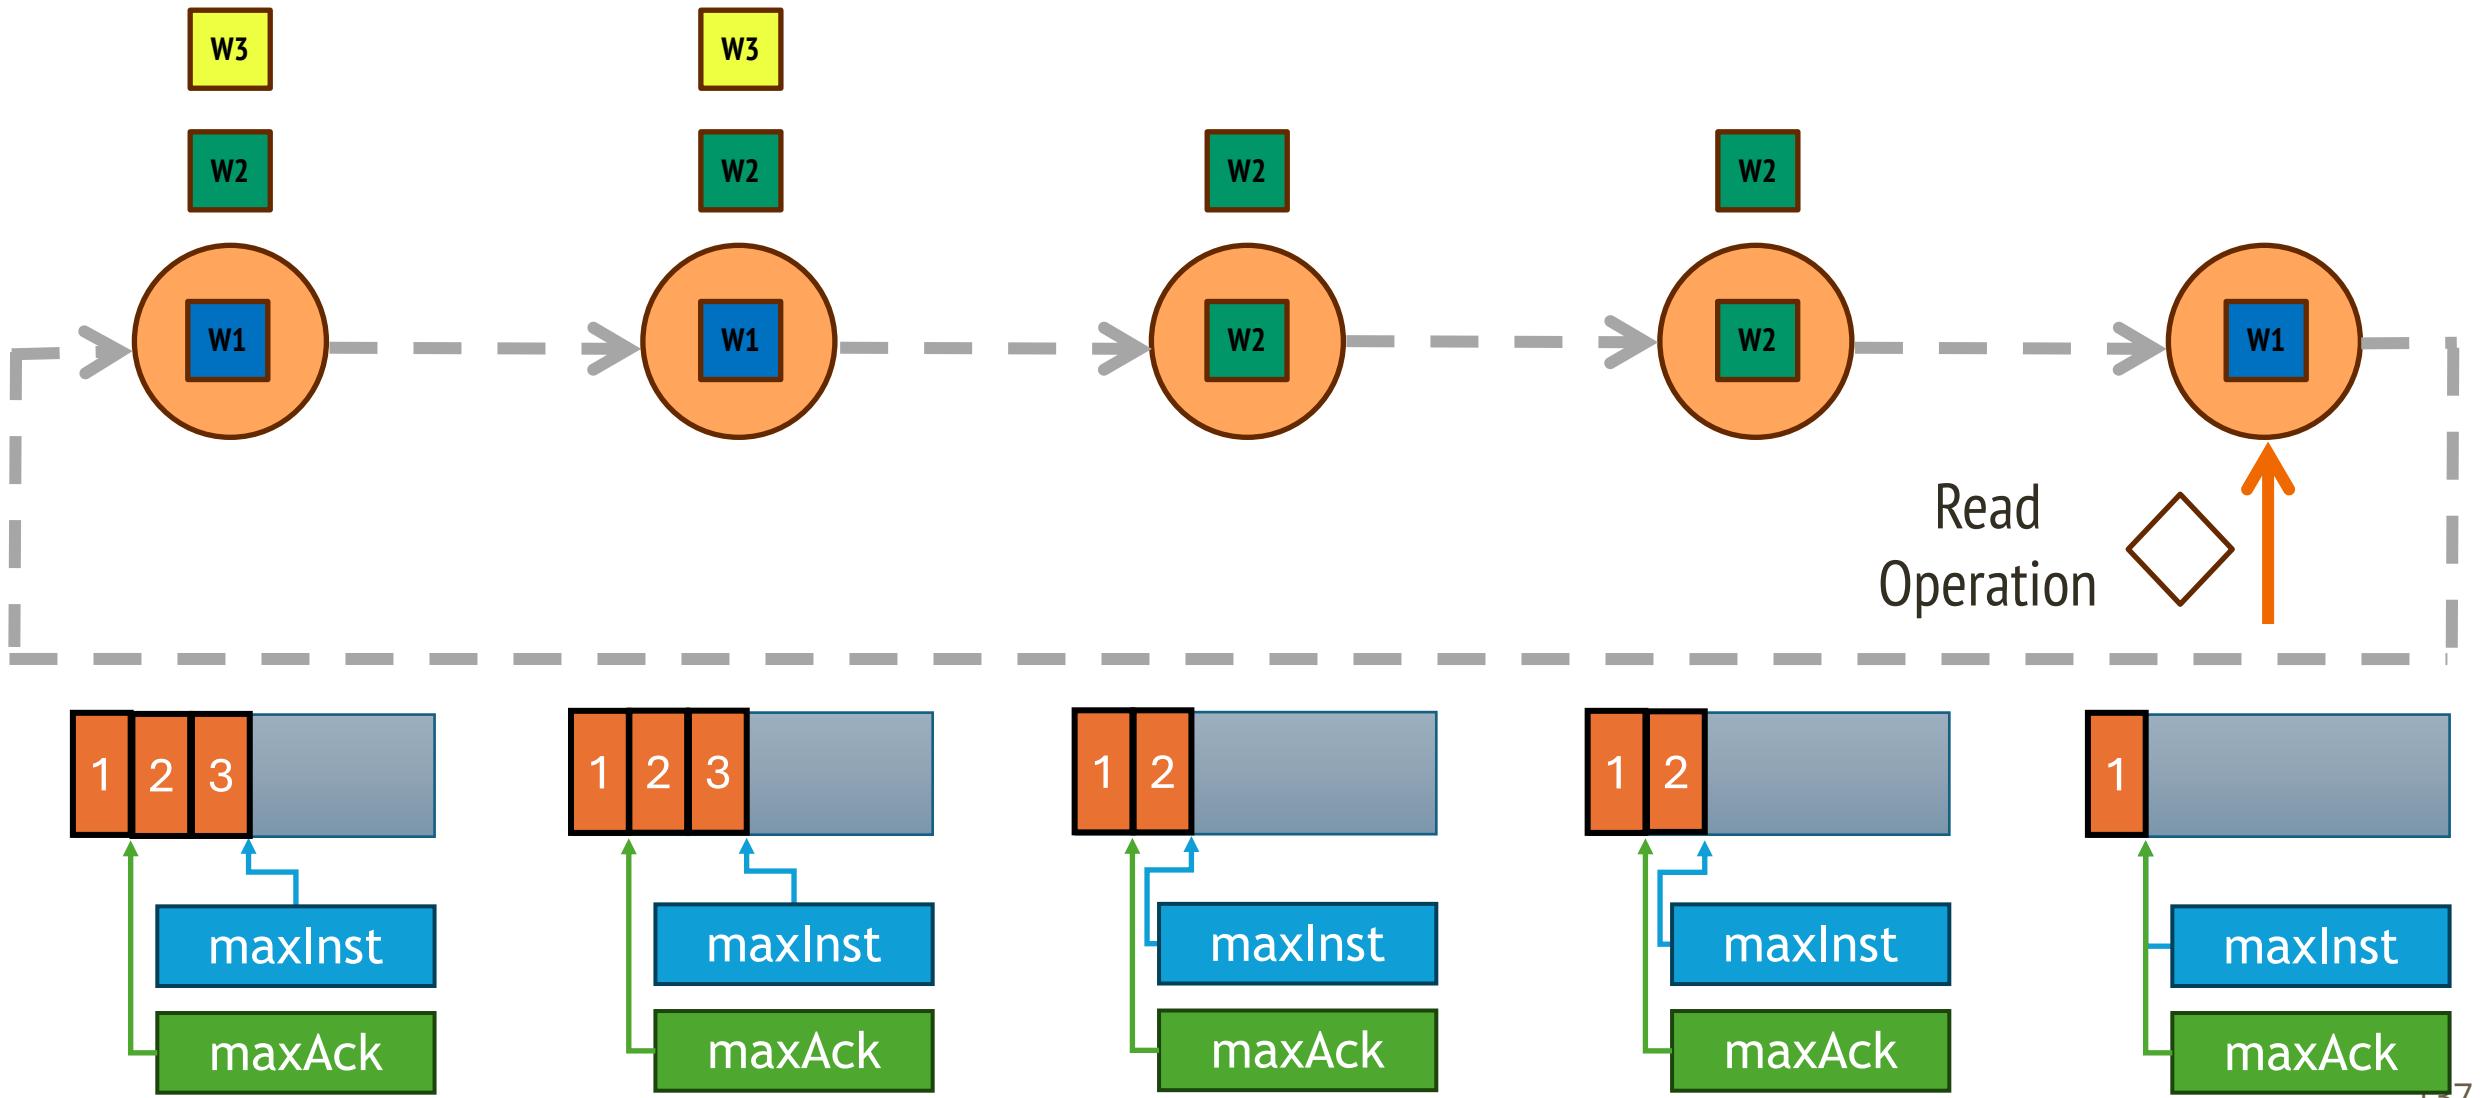
In Depth: Writes



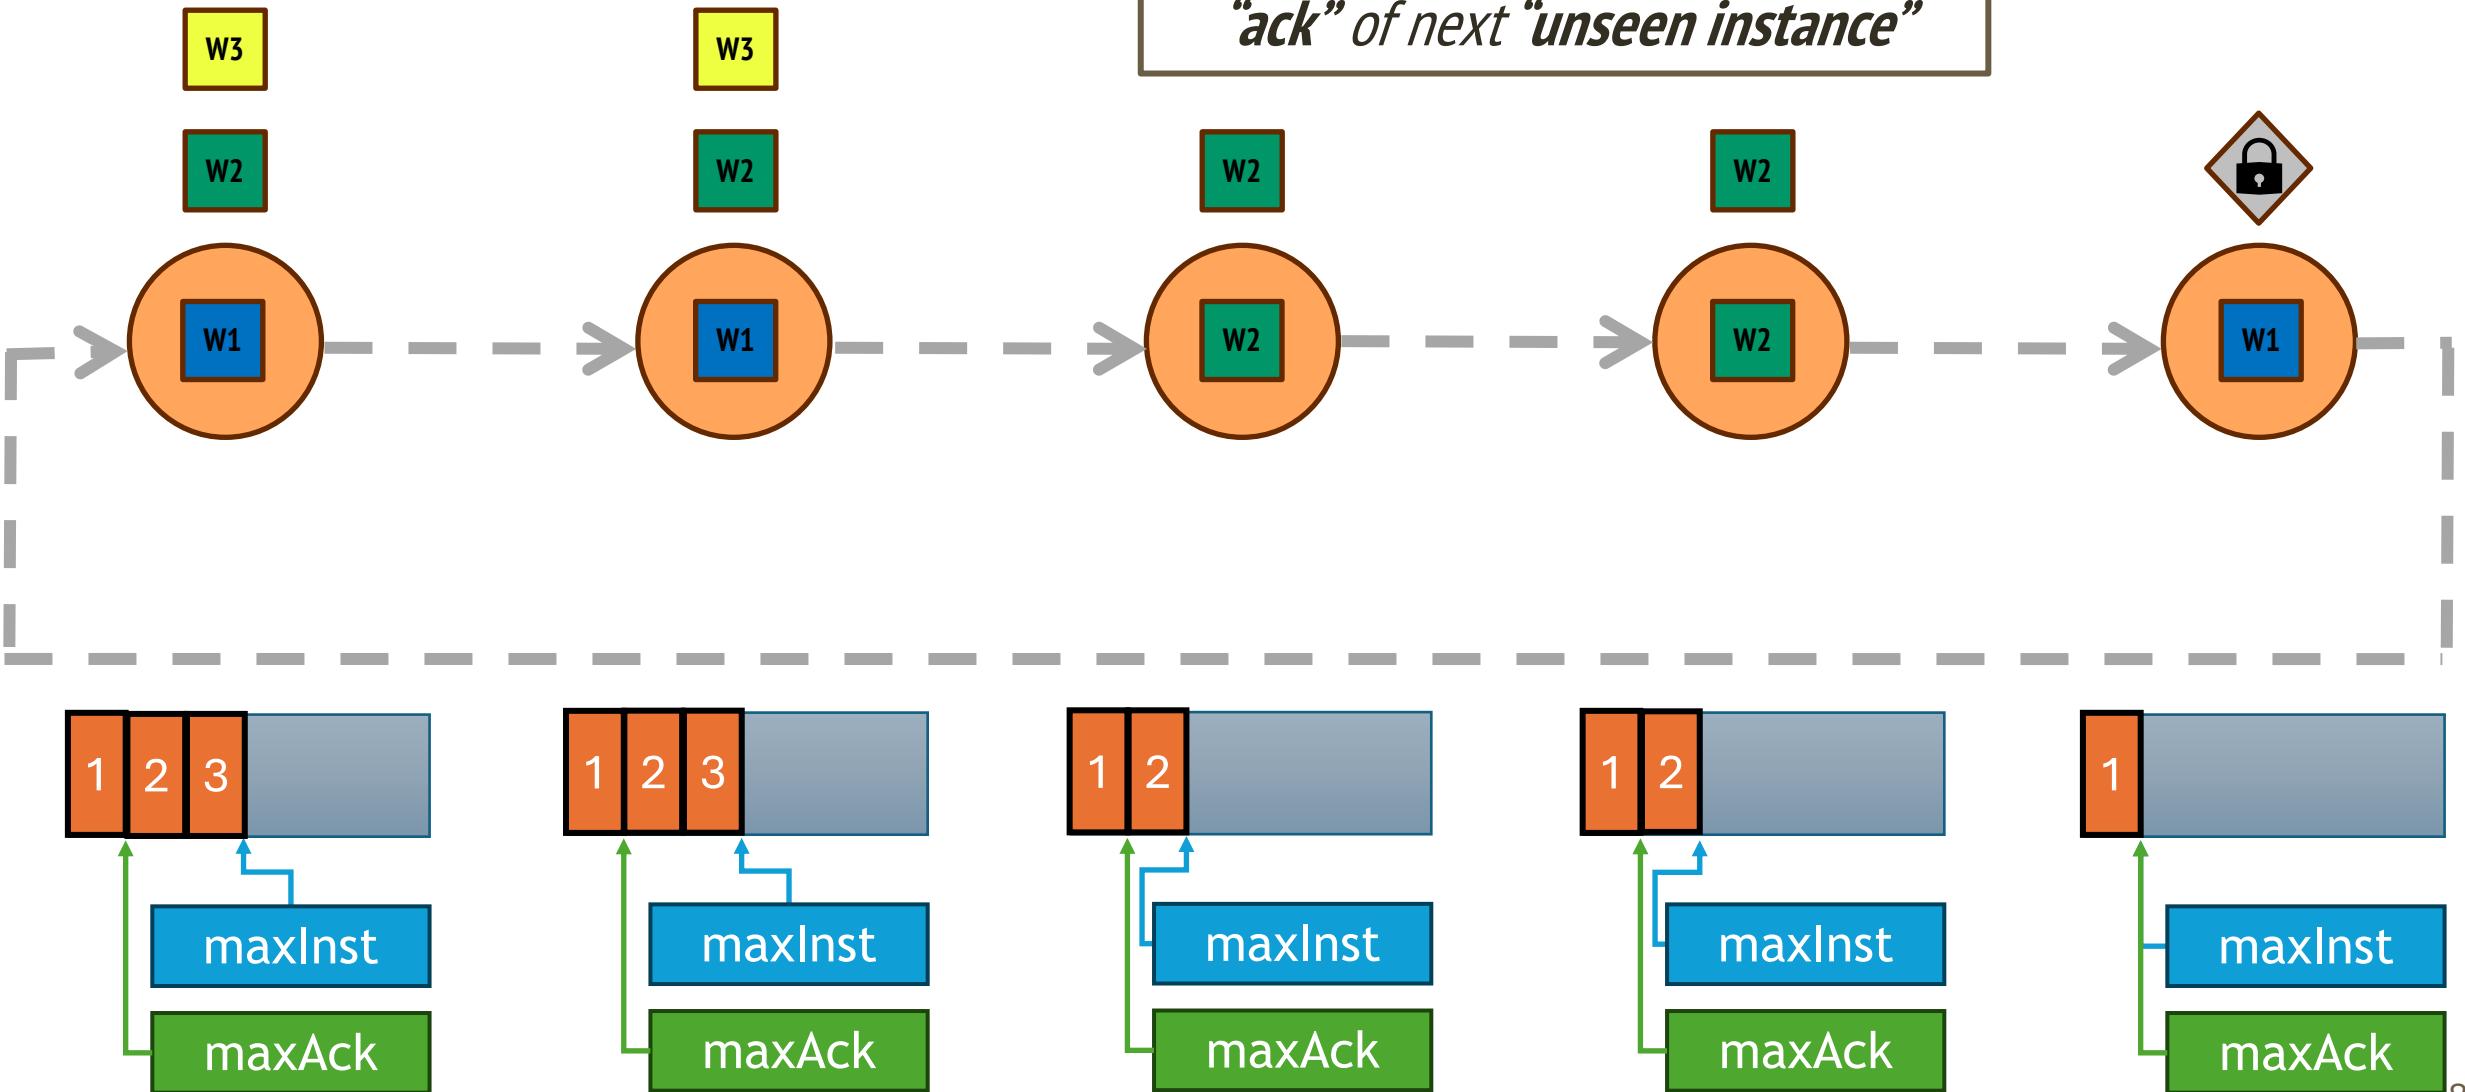
In Depth: Reads



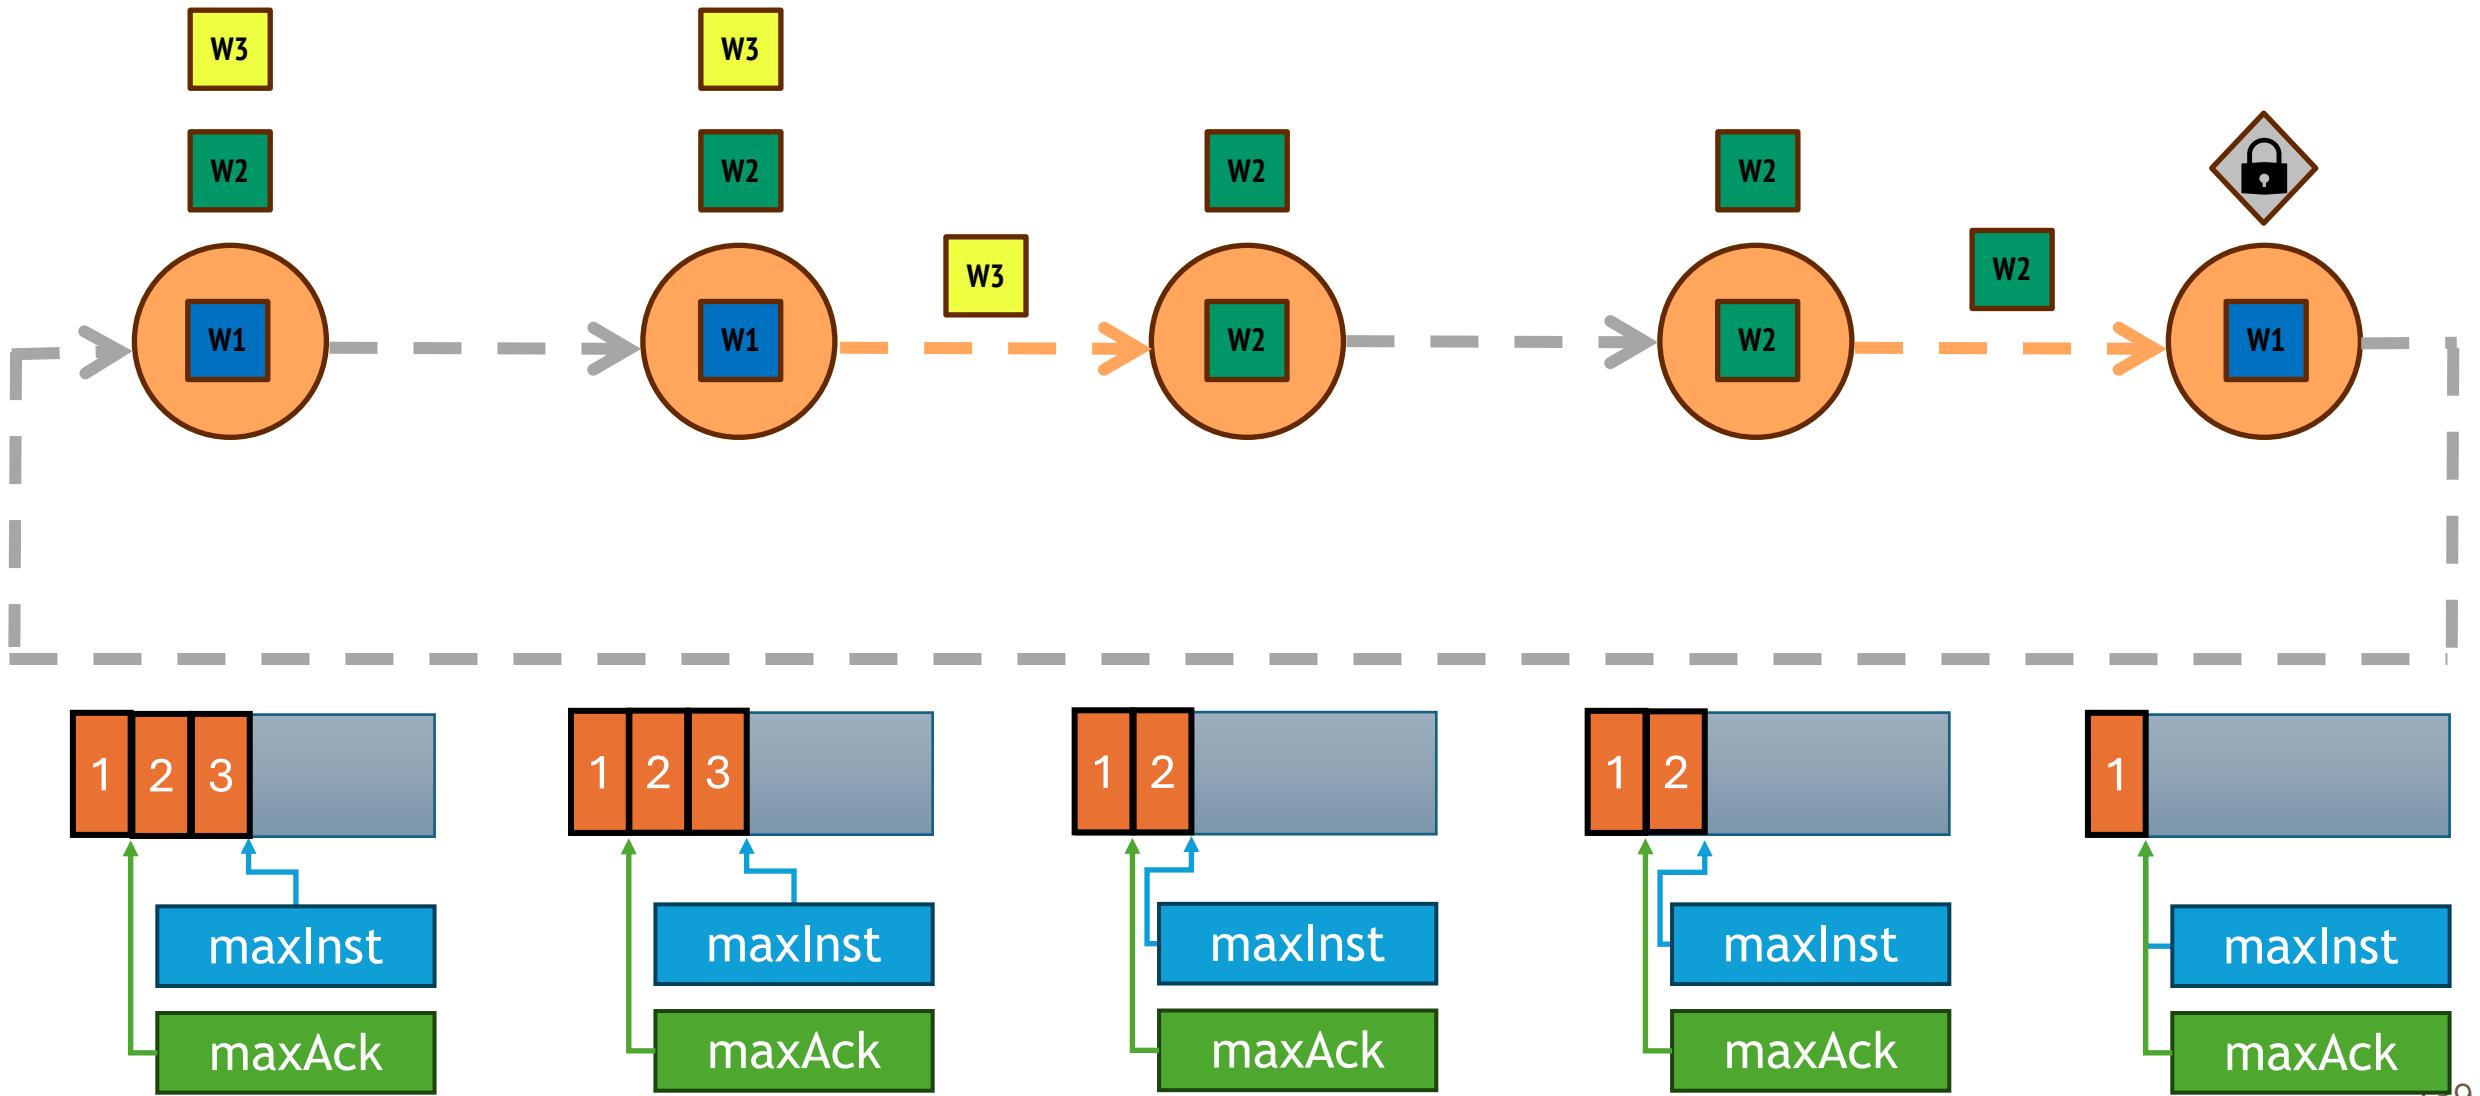
In Depth: Reads



In Depth: Reads

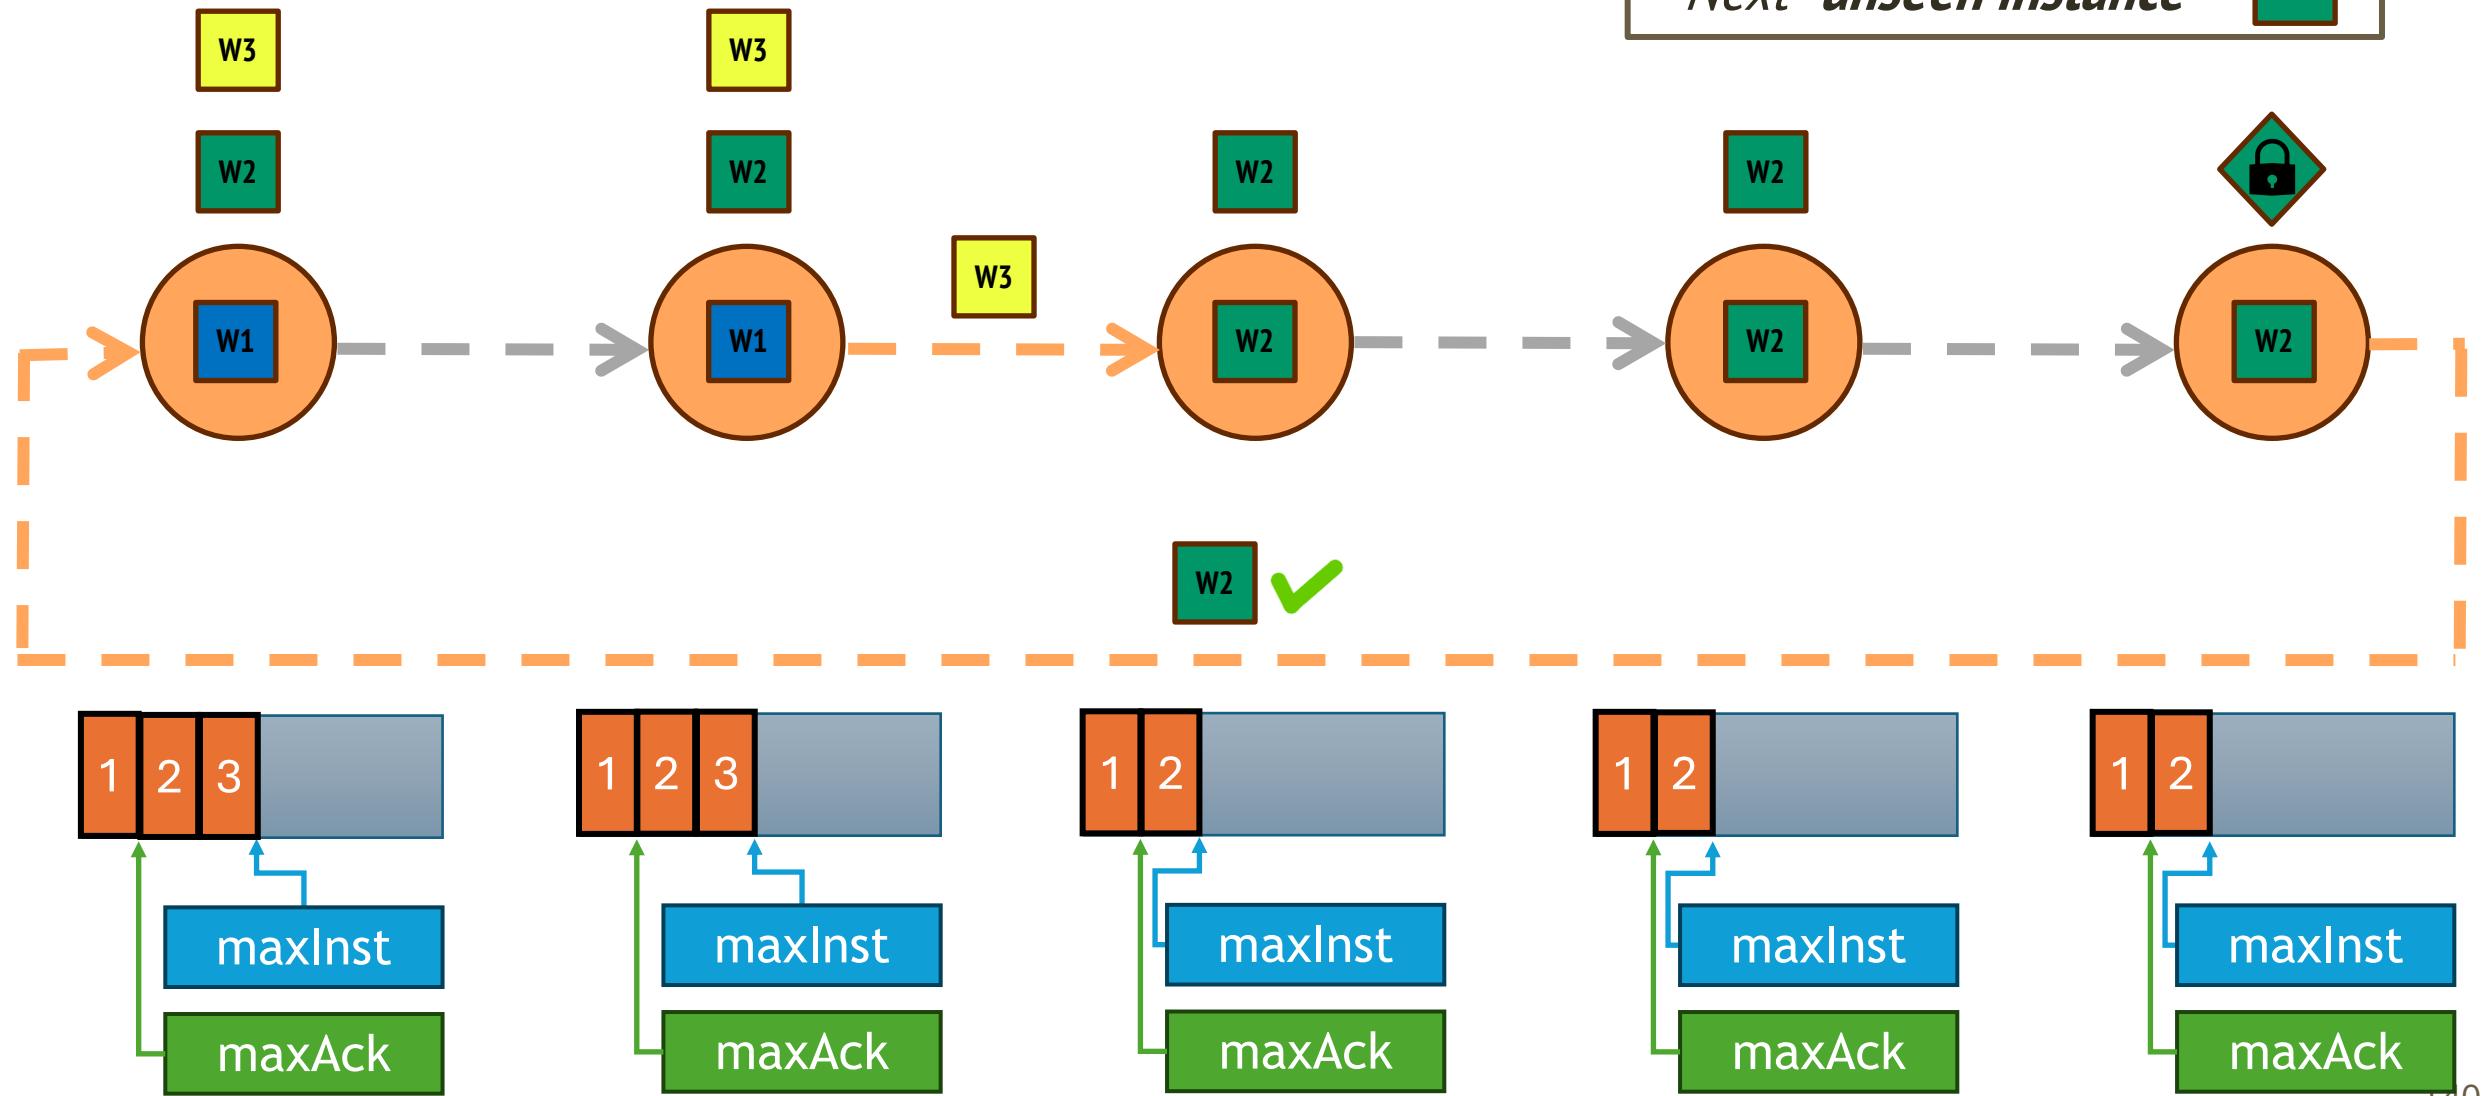


In Depth: Reads

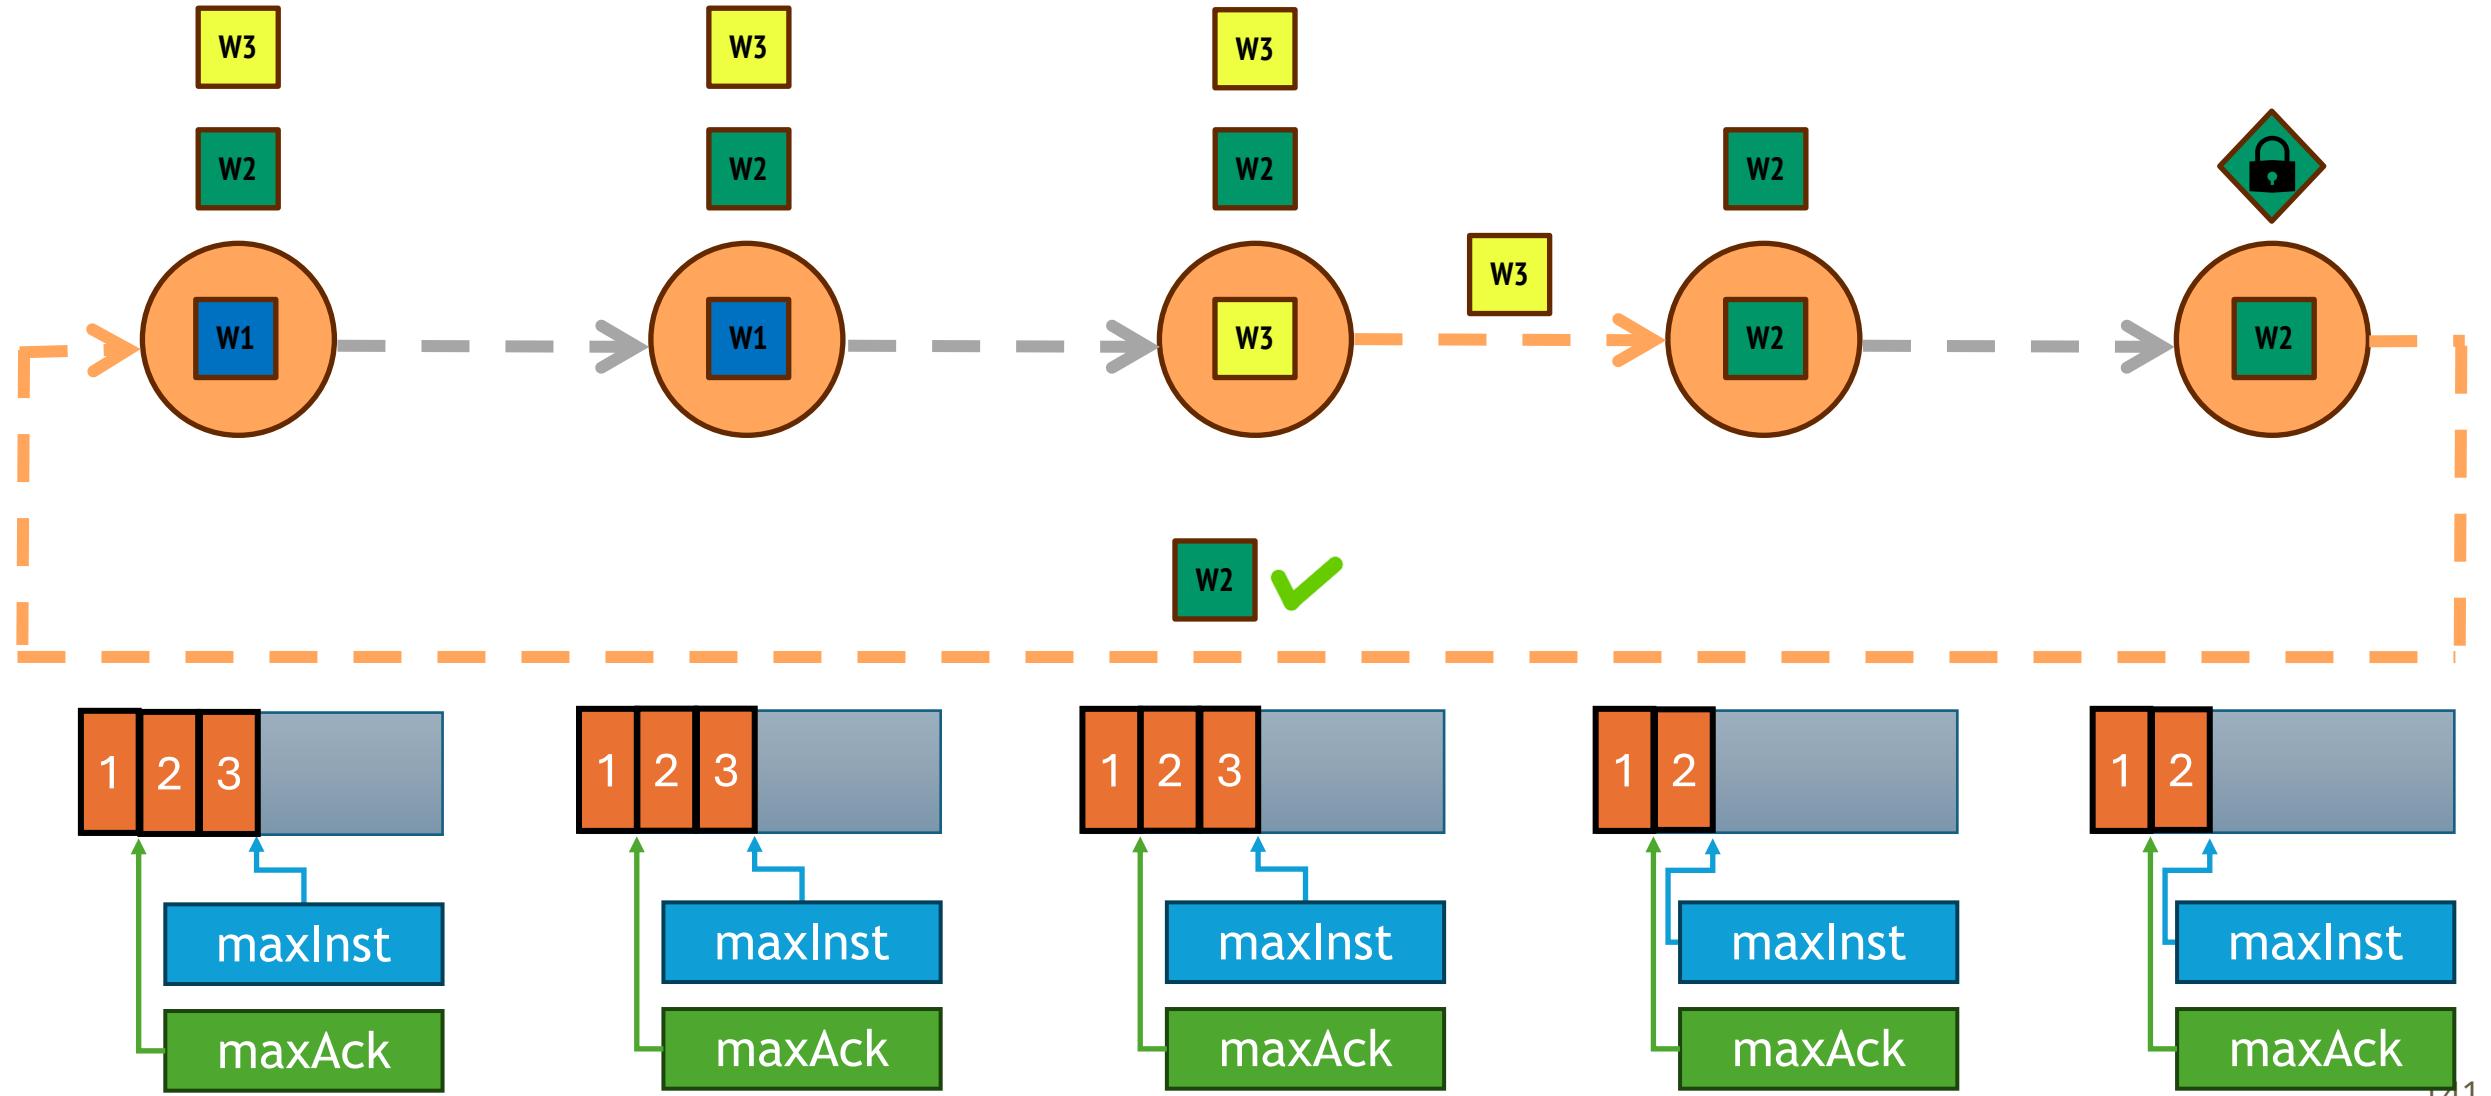


In Depth: Reads

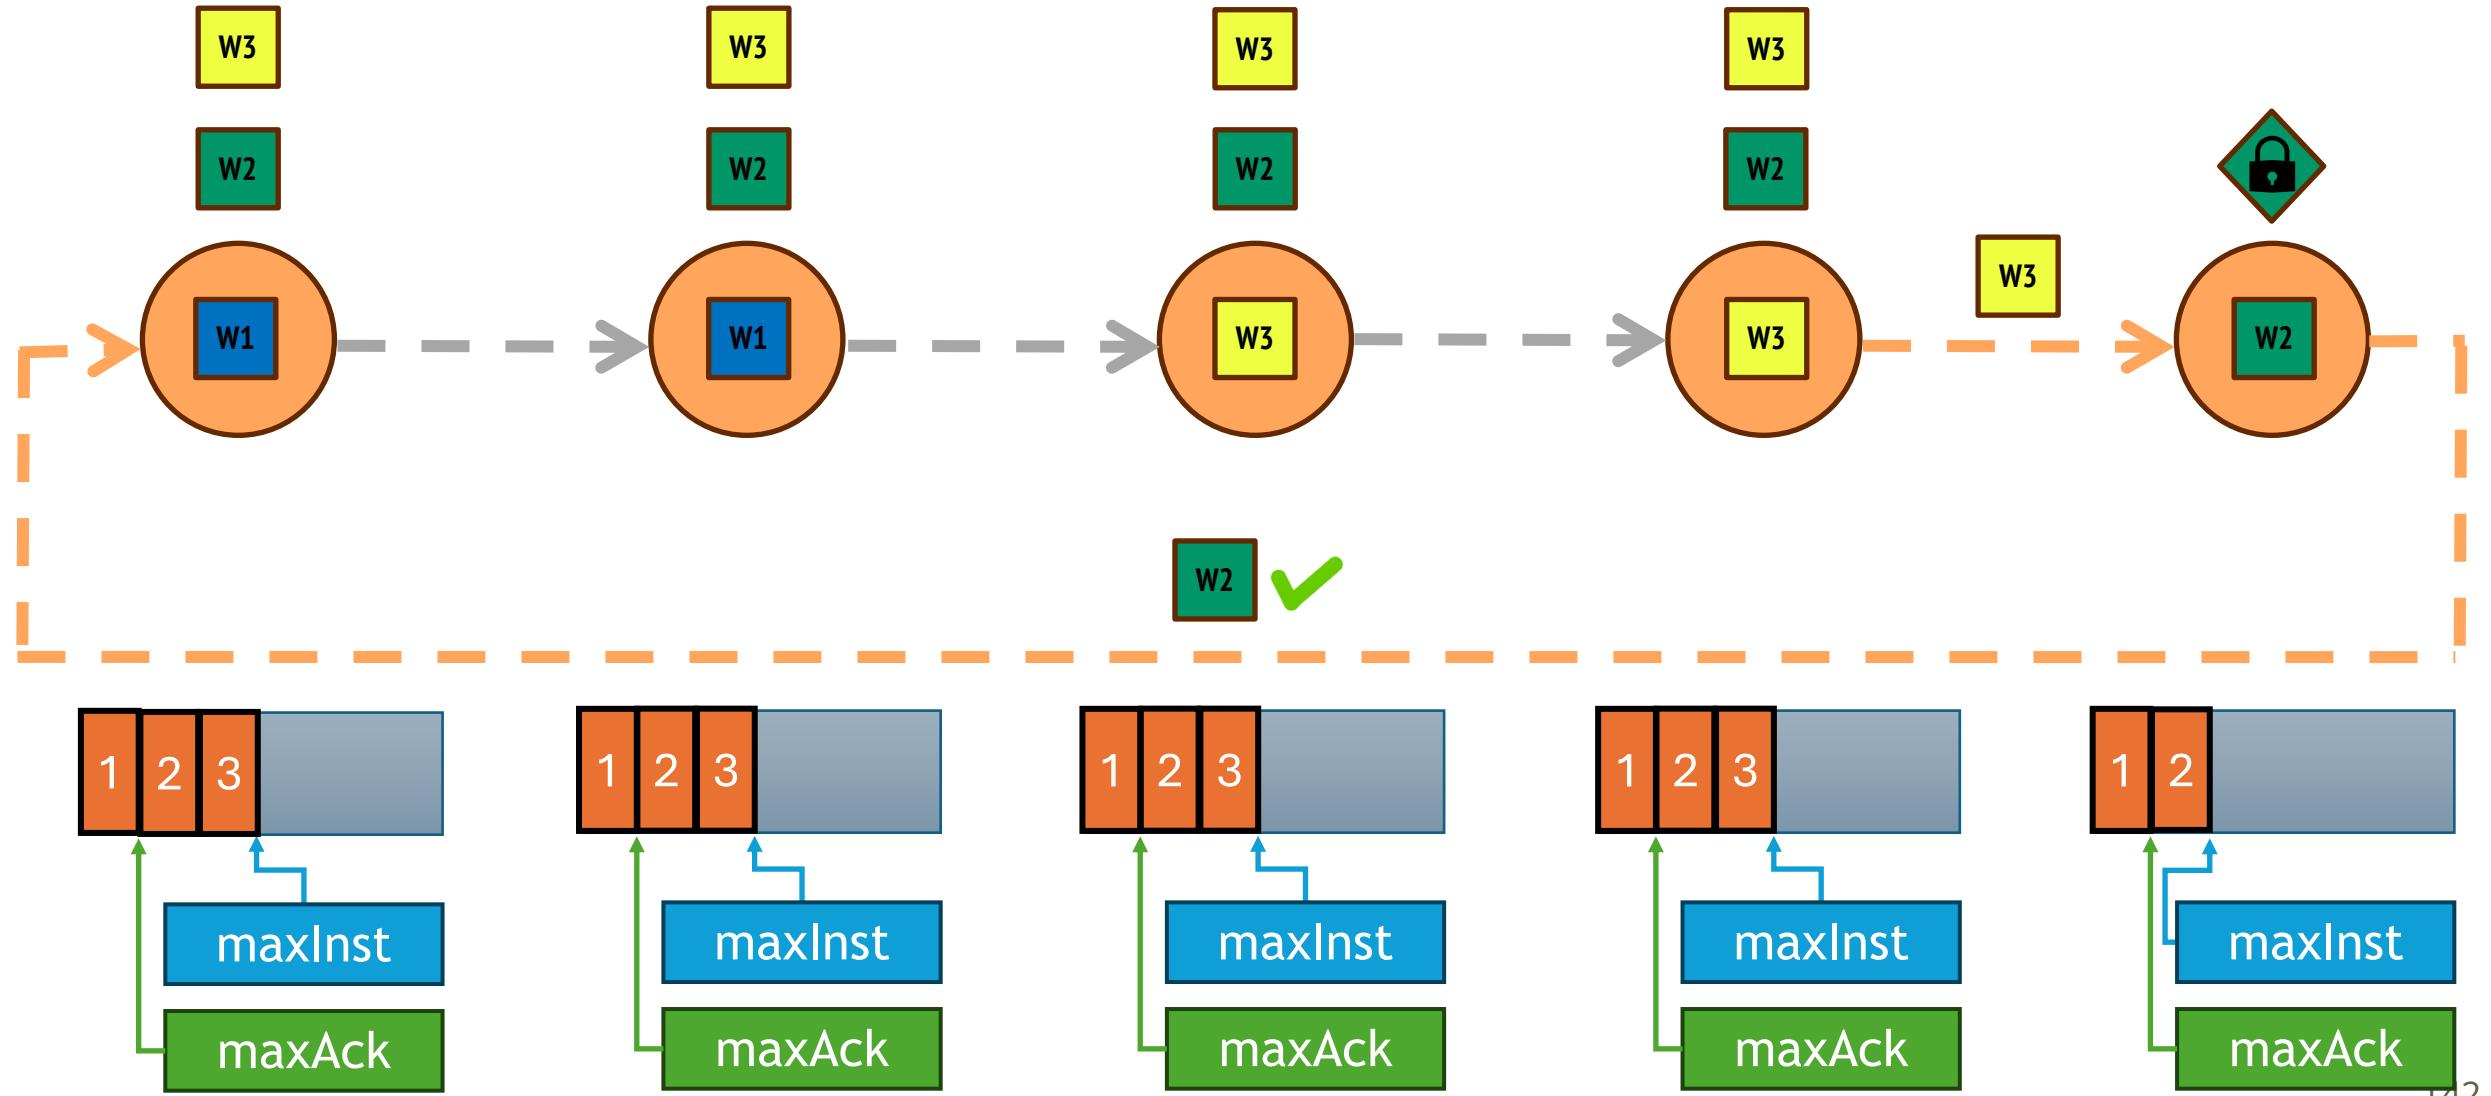
Next “*unseen instance*” - **W2**



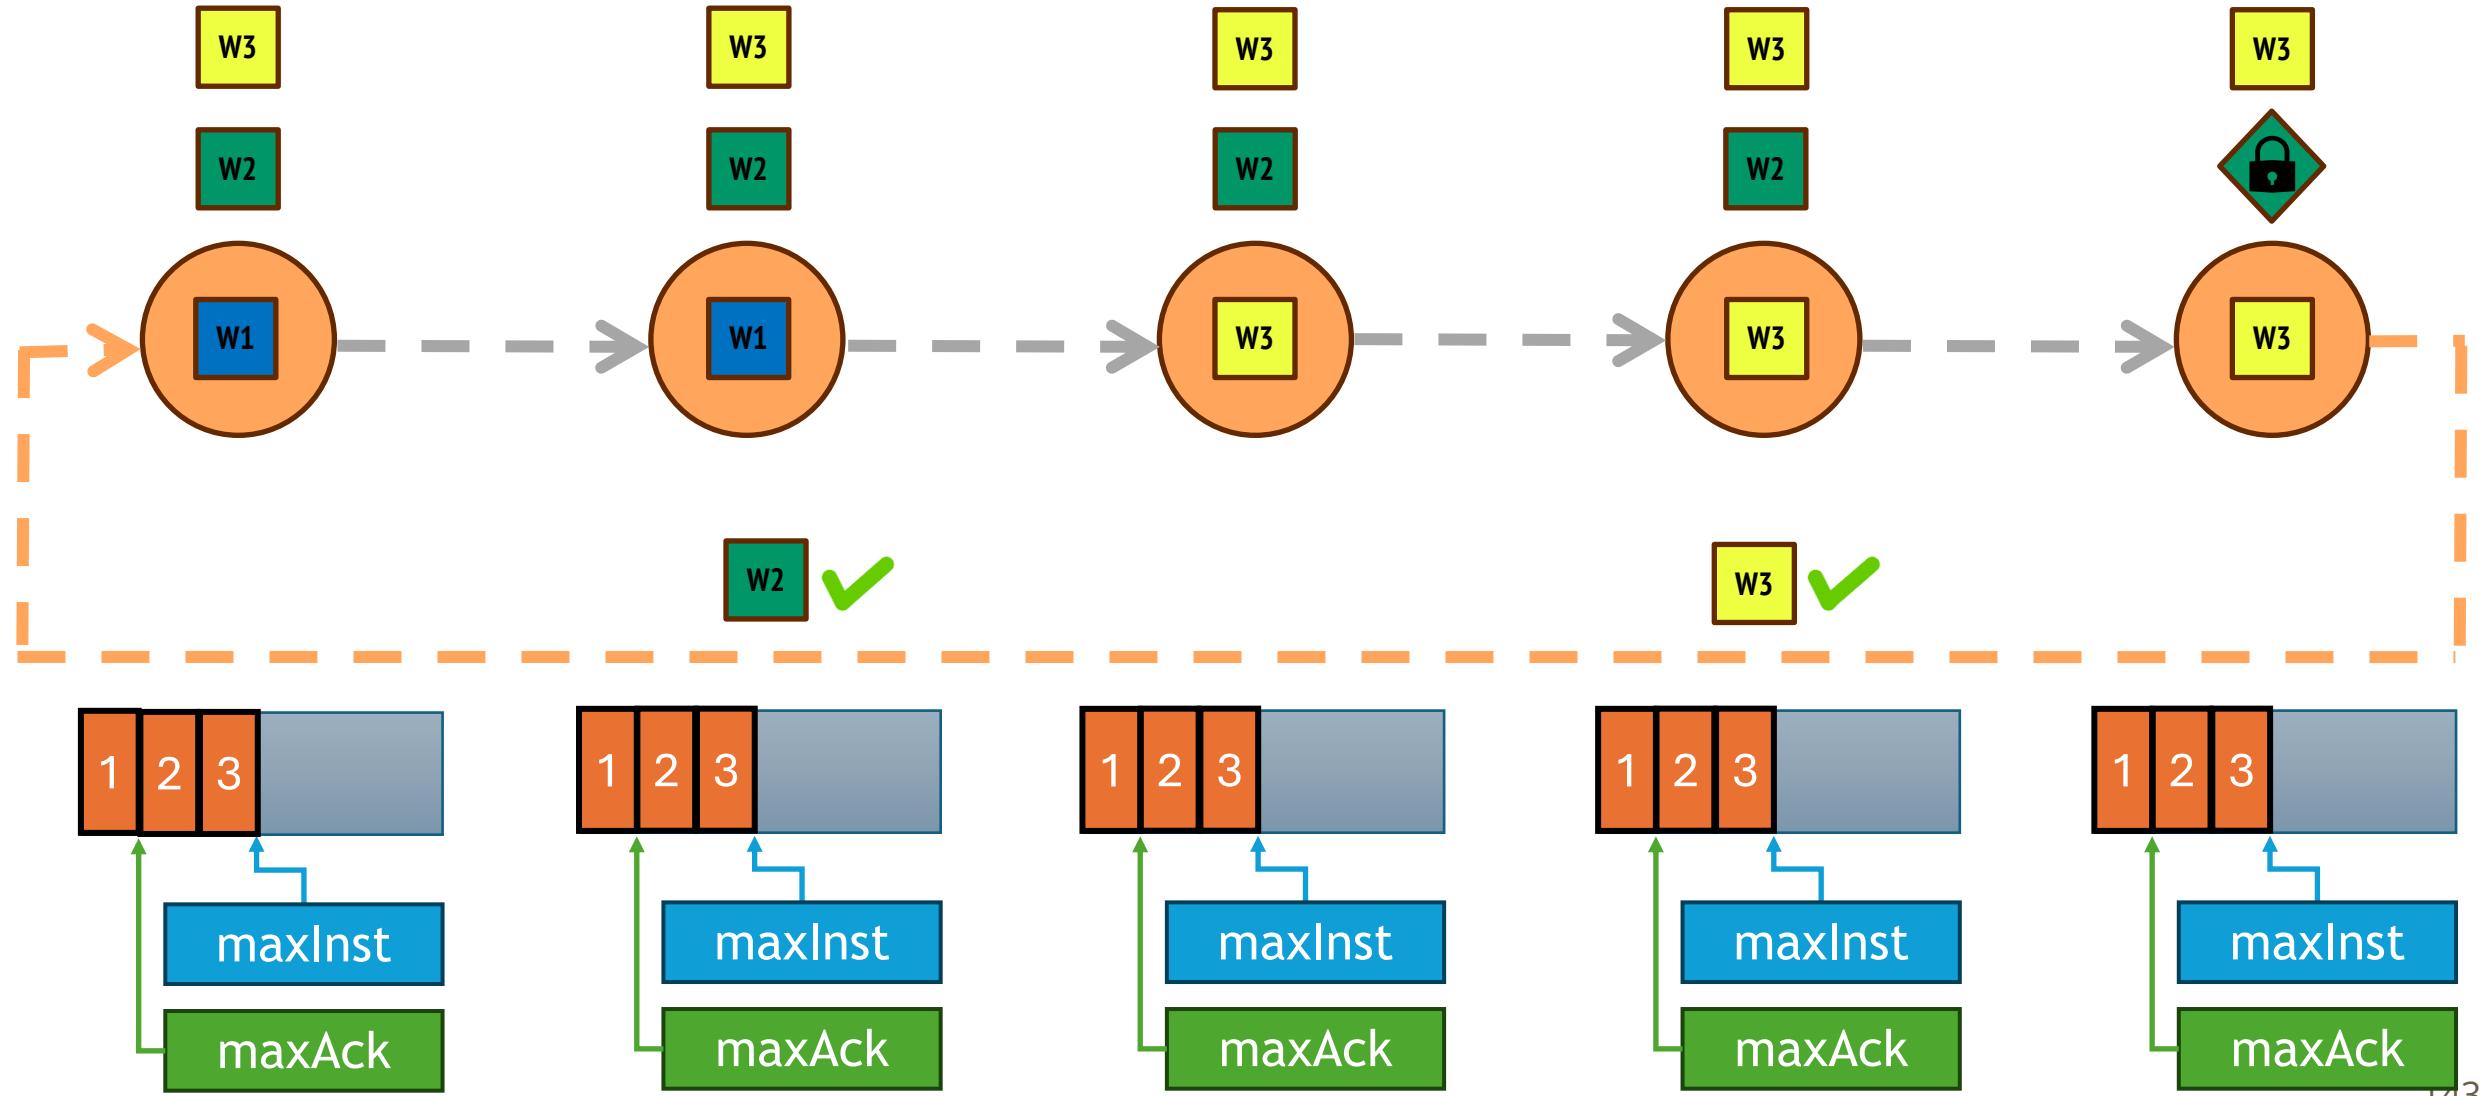
In Depth: Reads



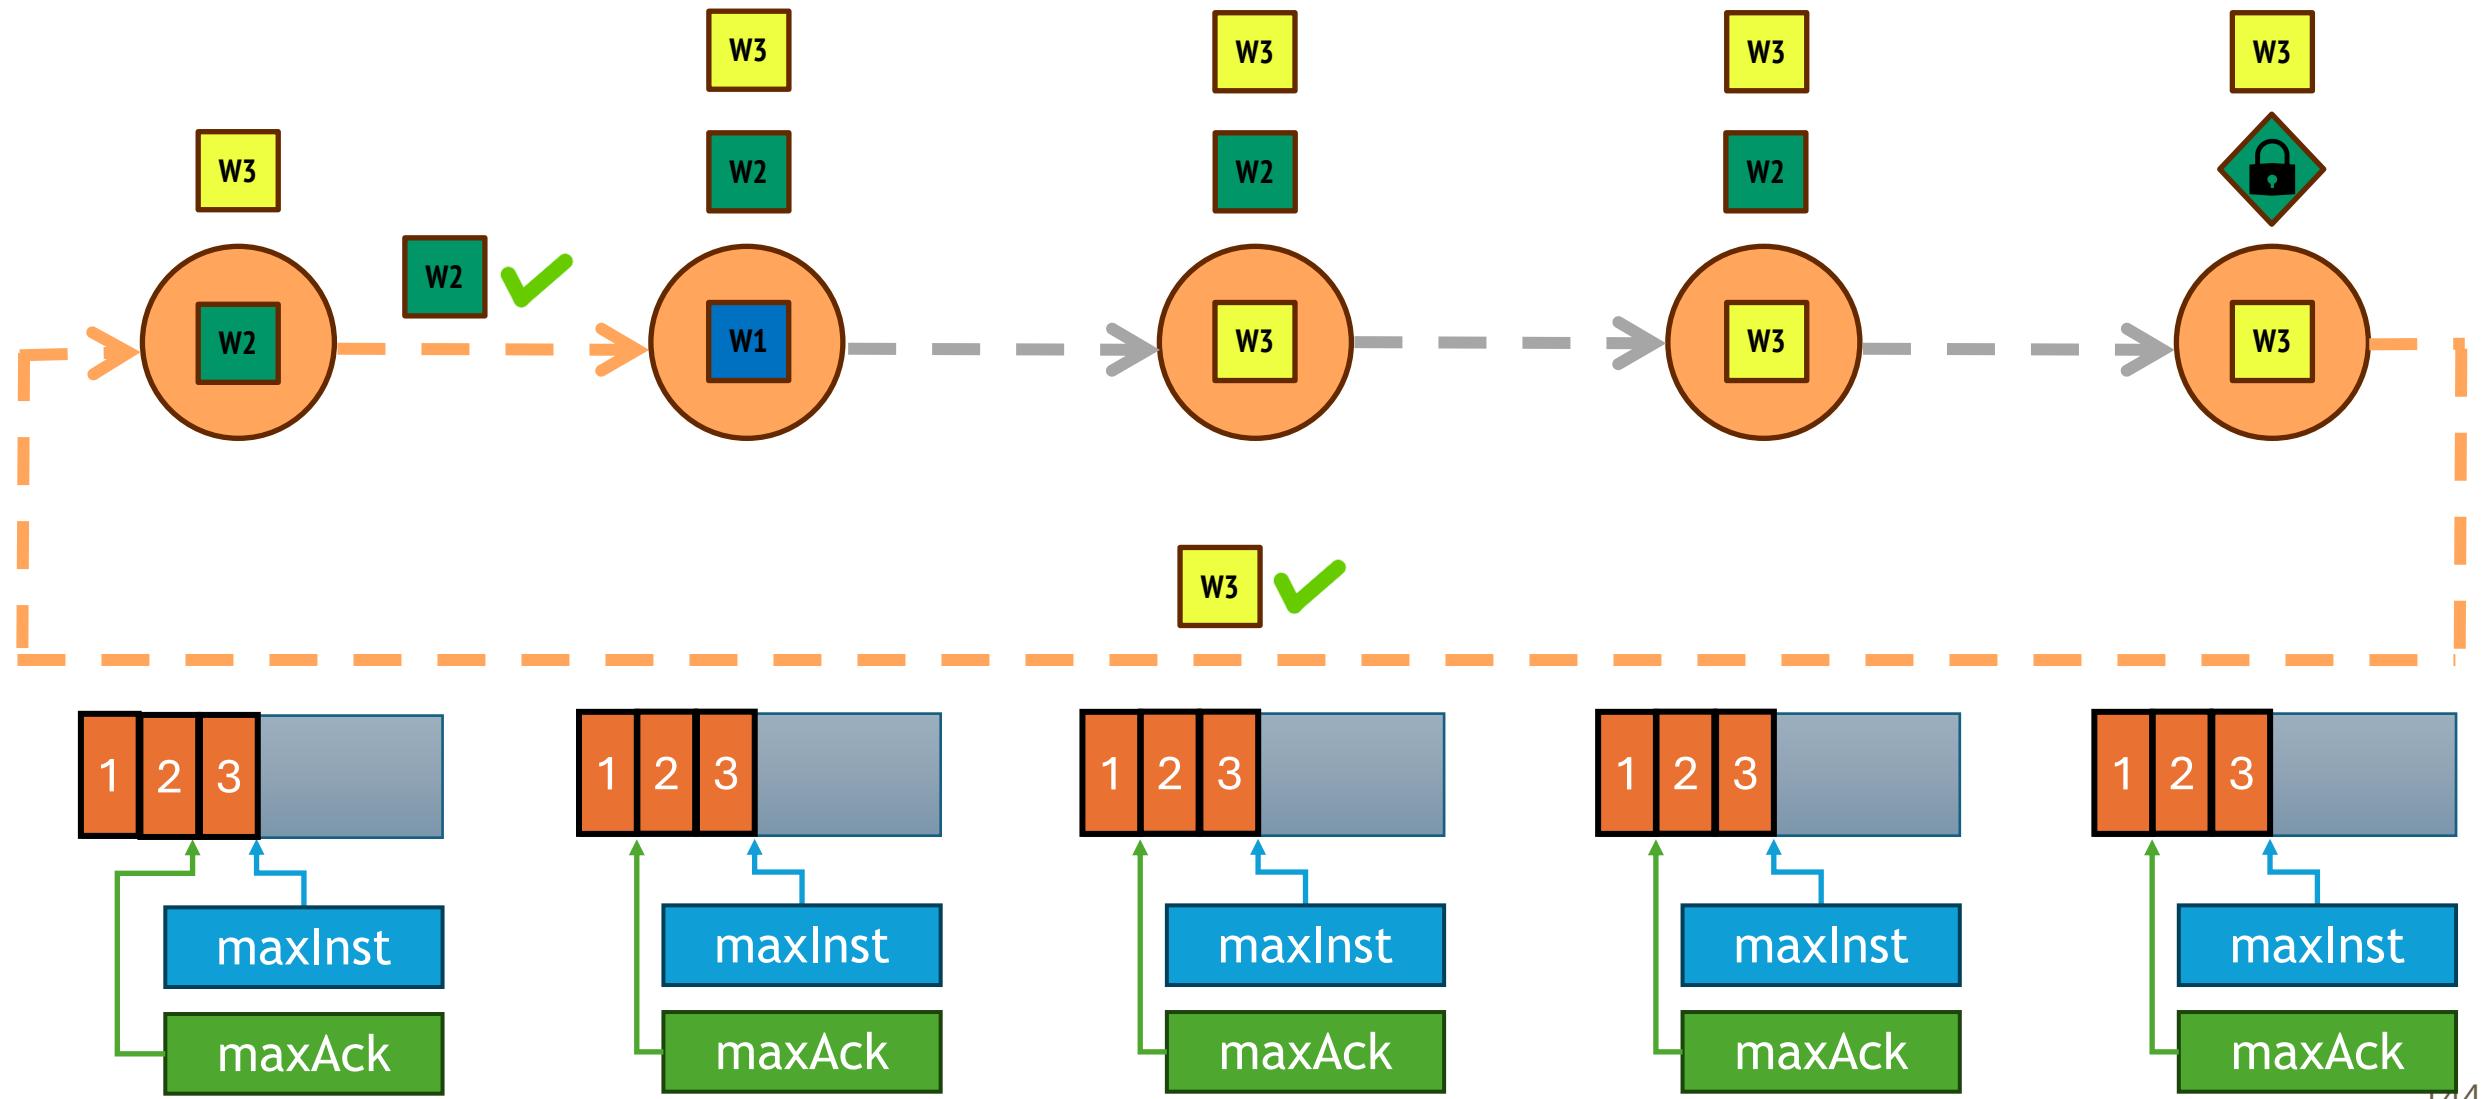
In Depth: Reads



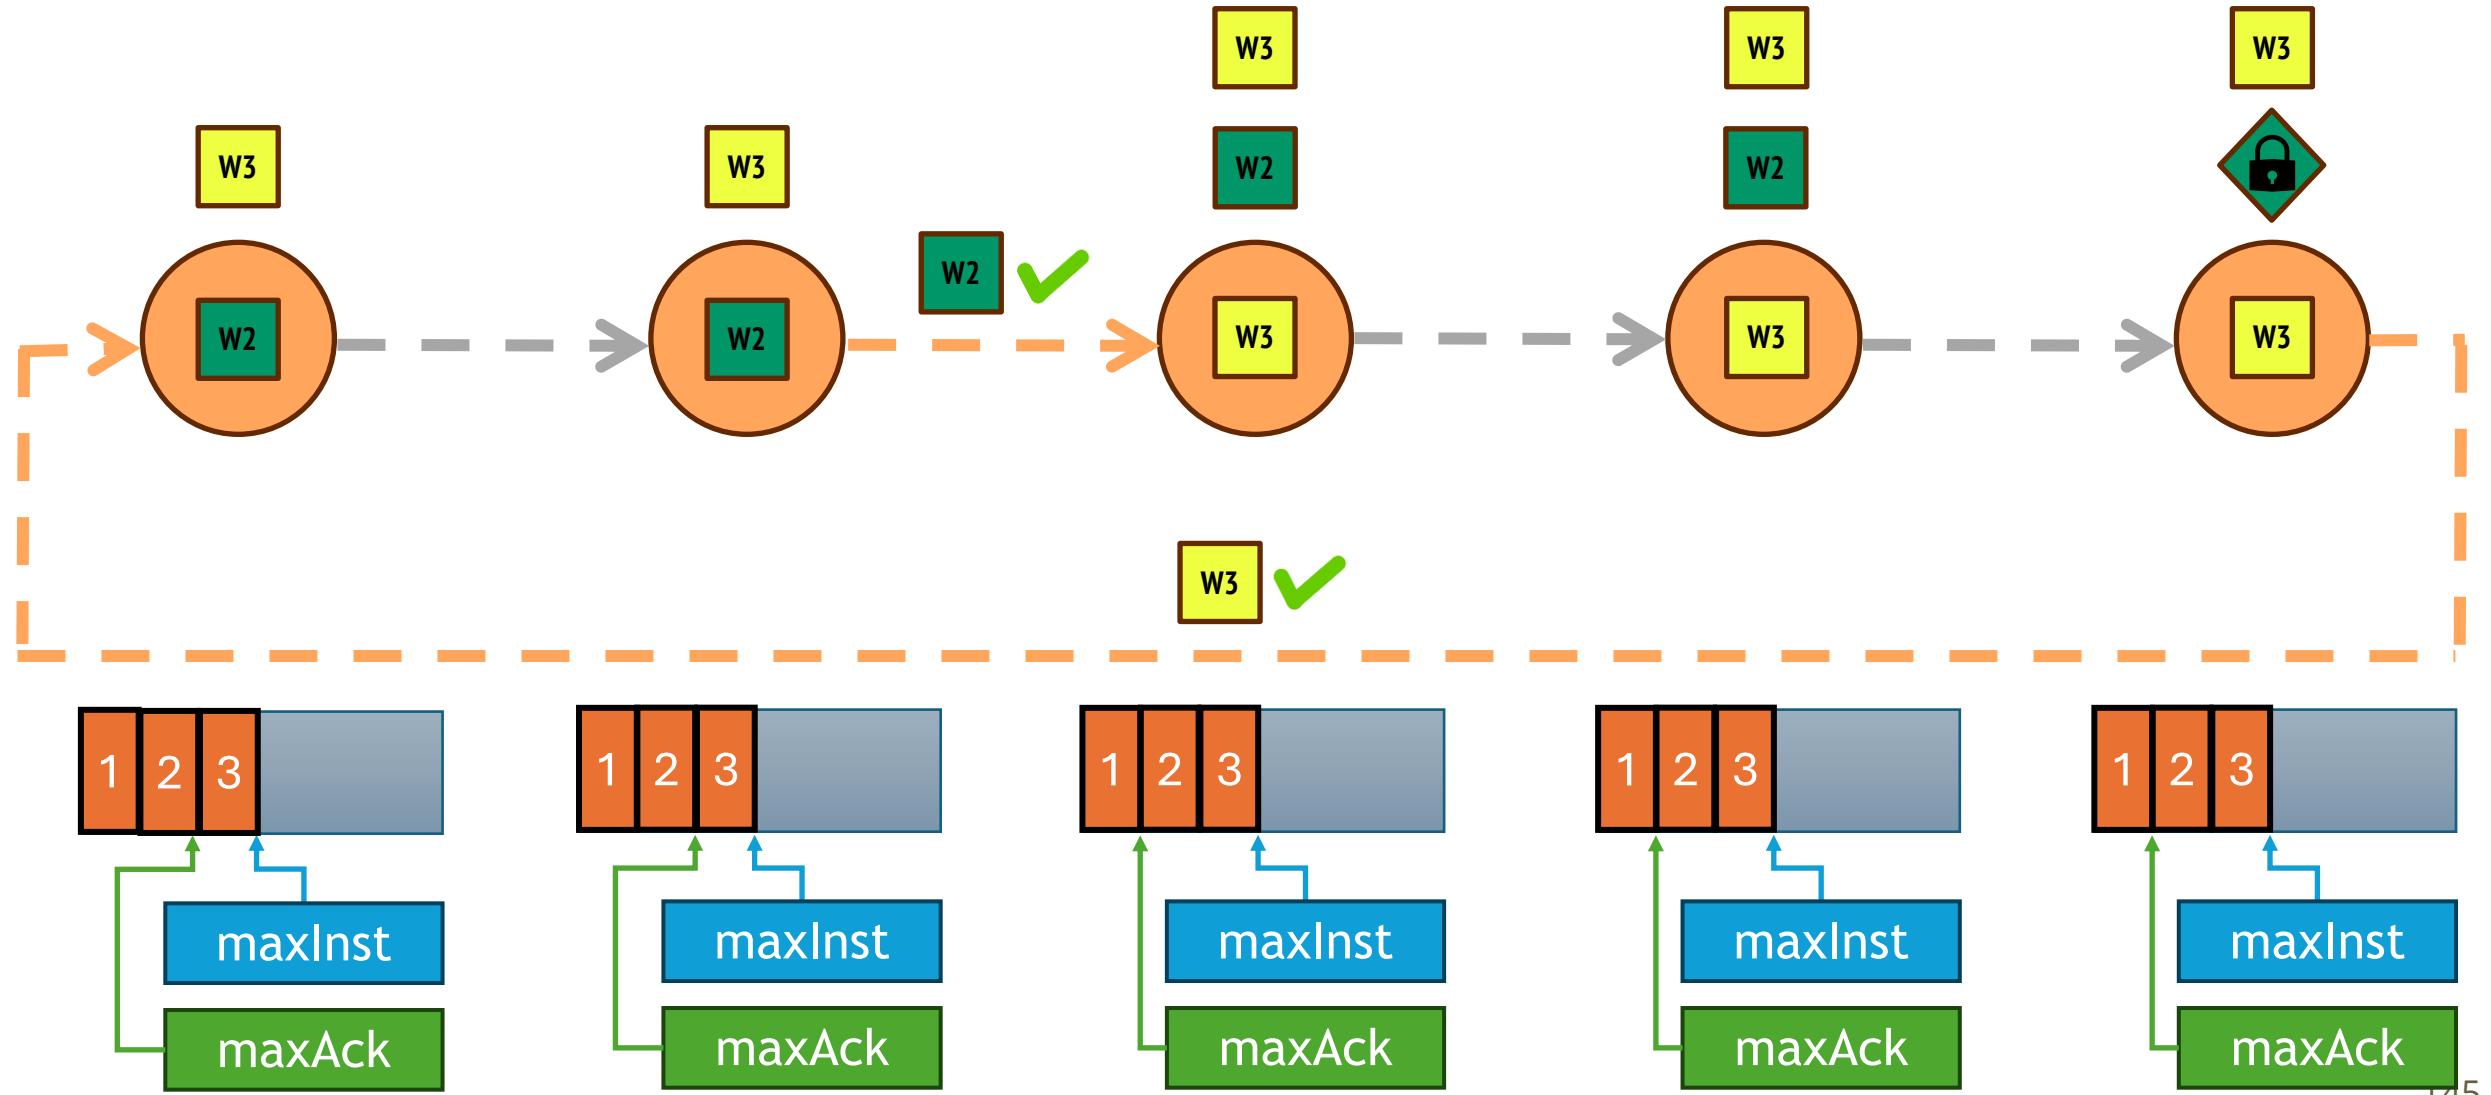
In Depth: Reads



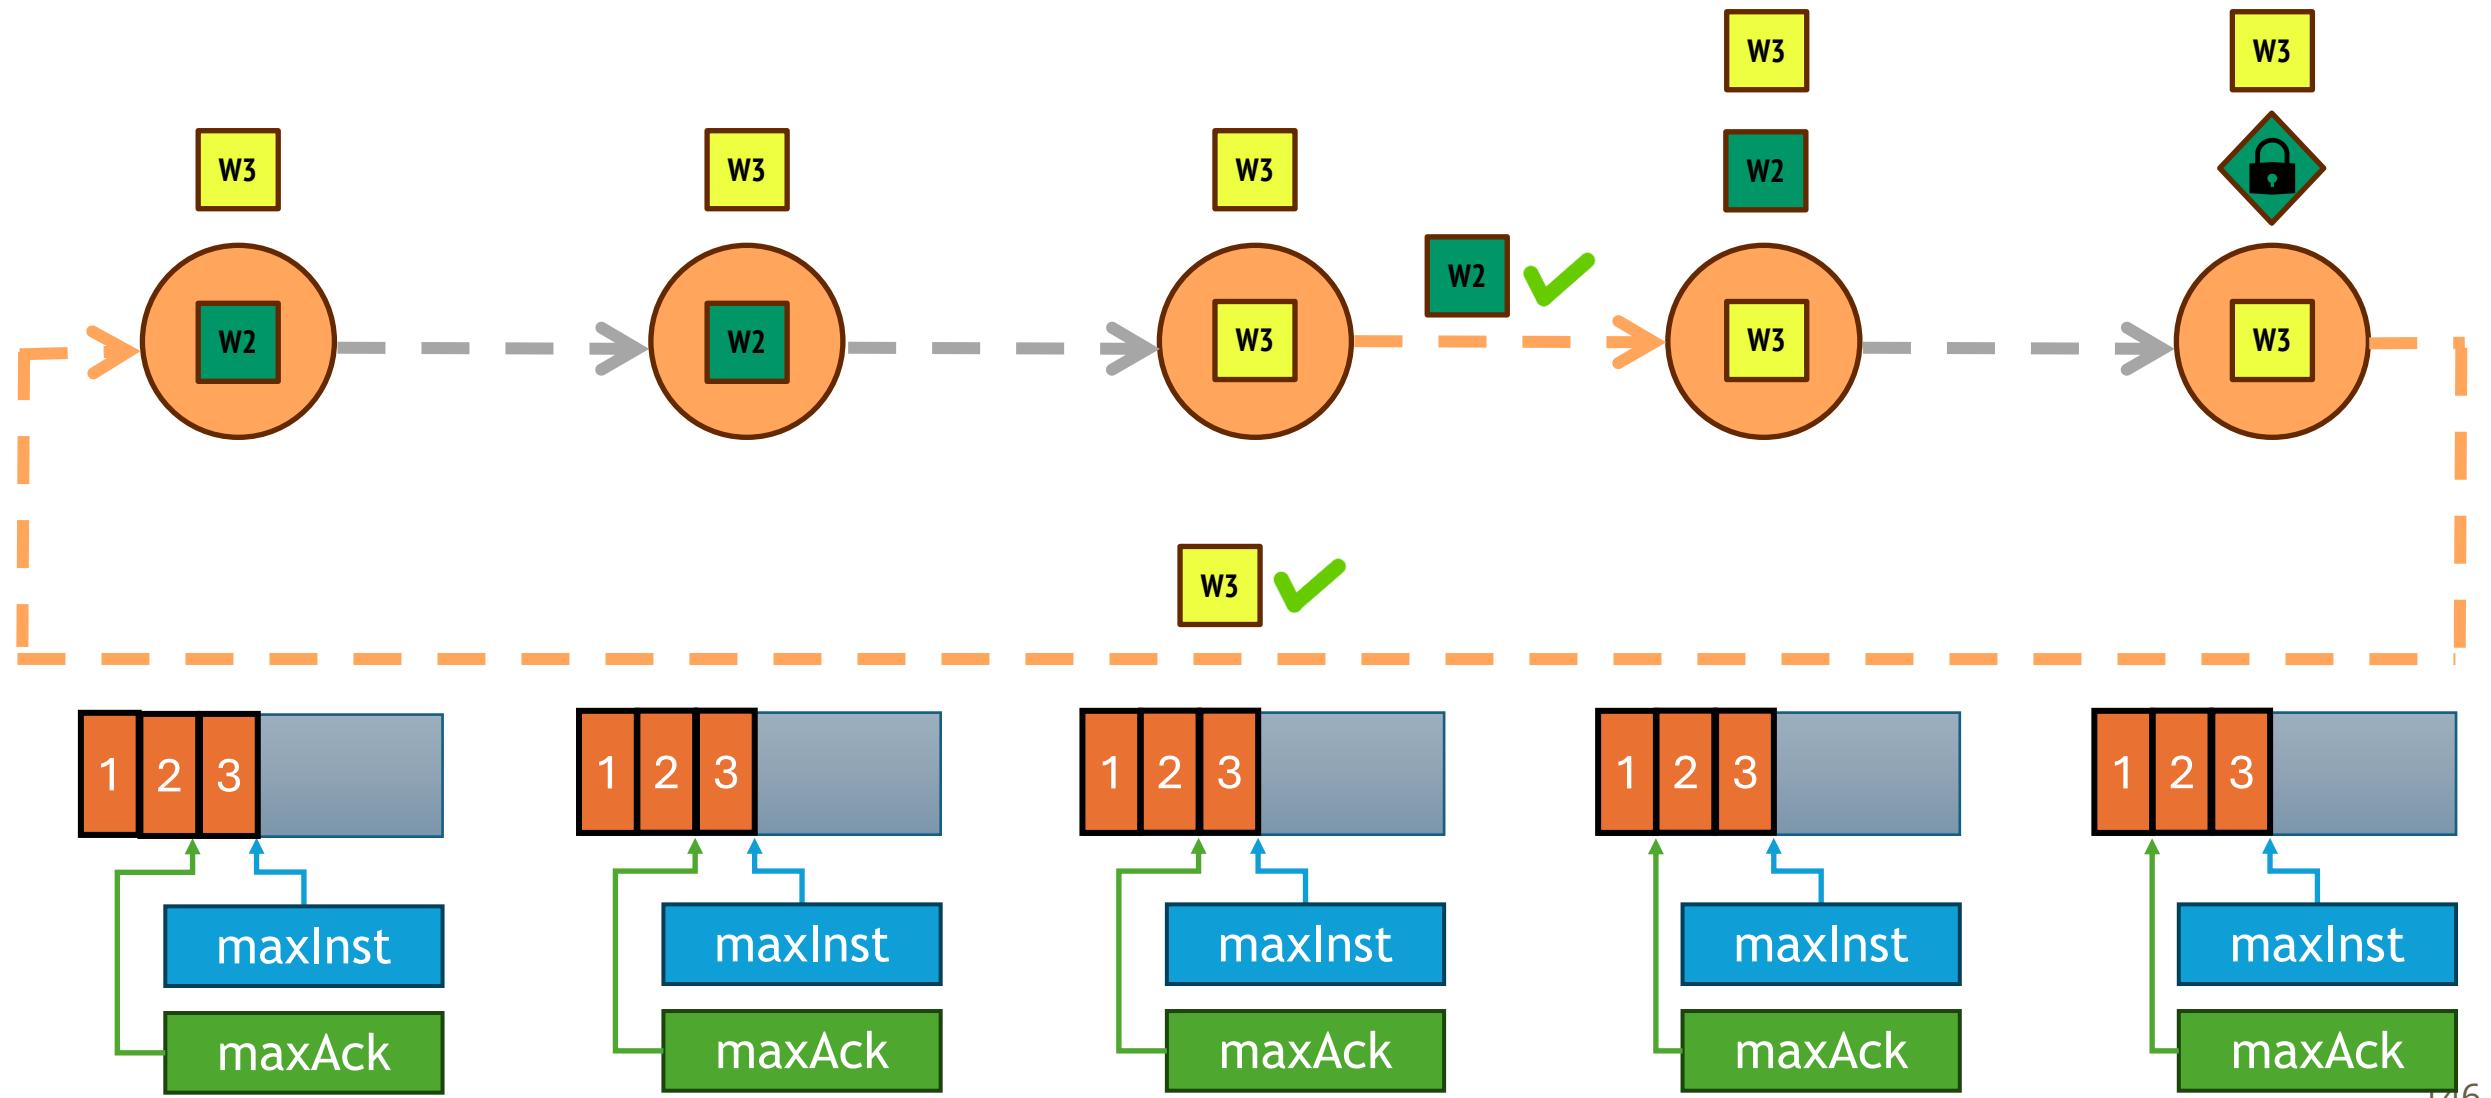
In Depth: Reads



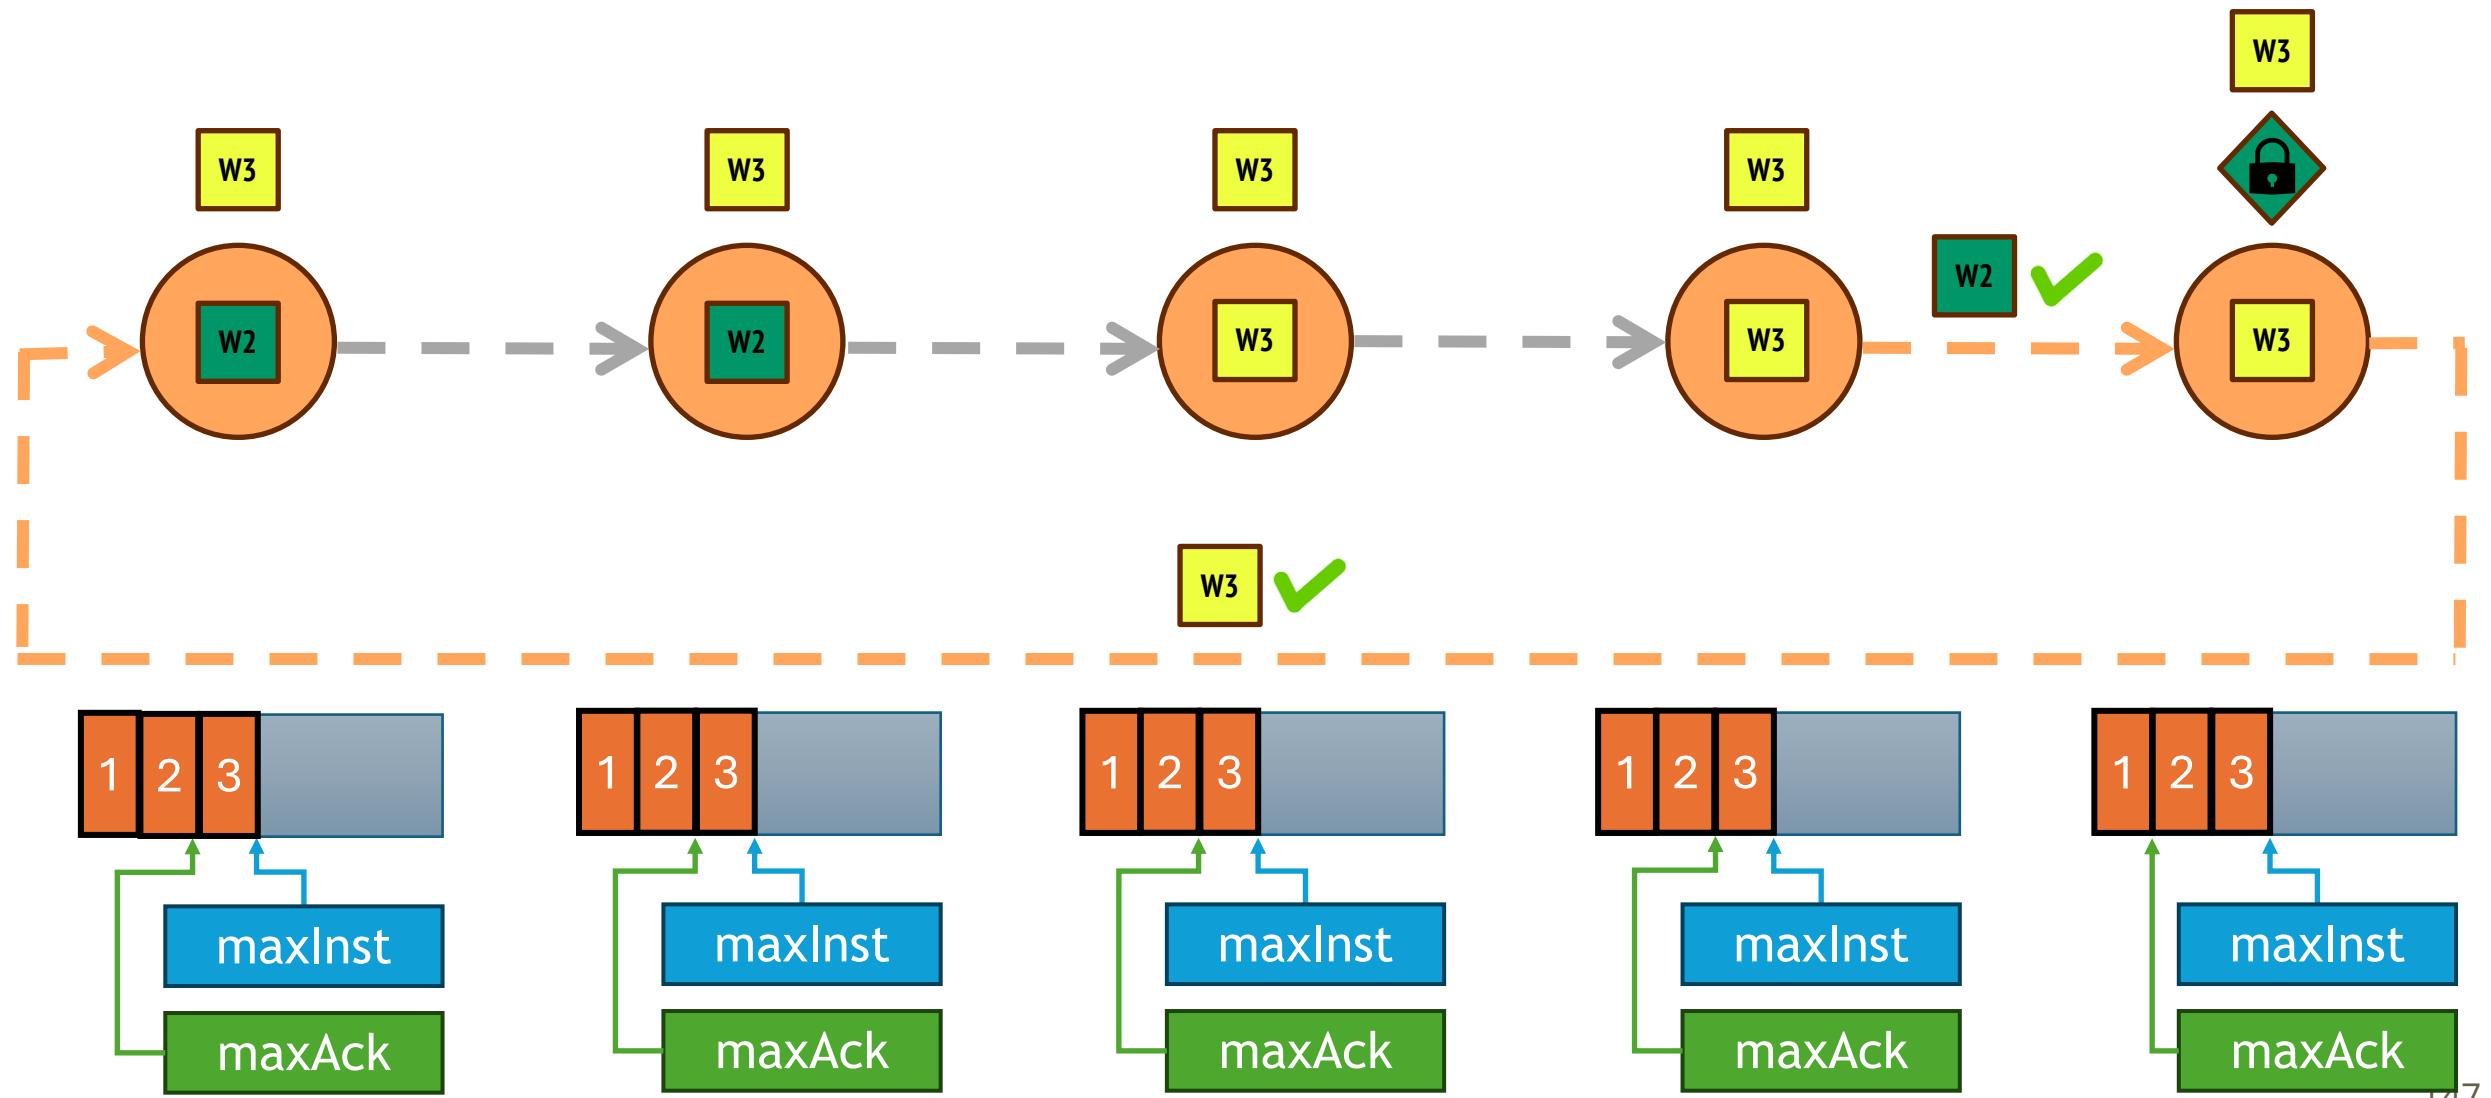
In Depth: Reads



In Depth: Reads

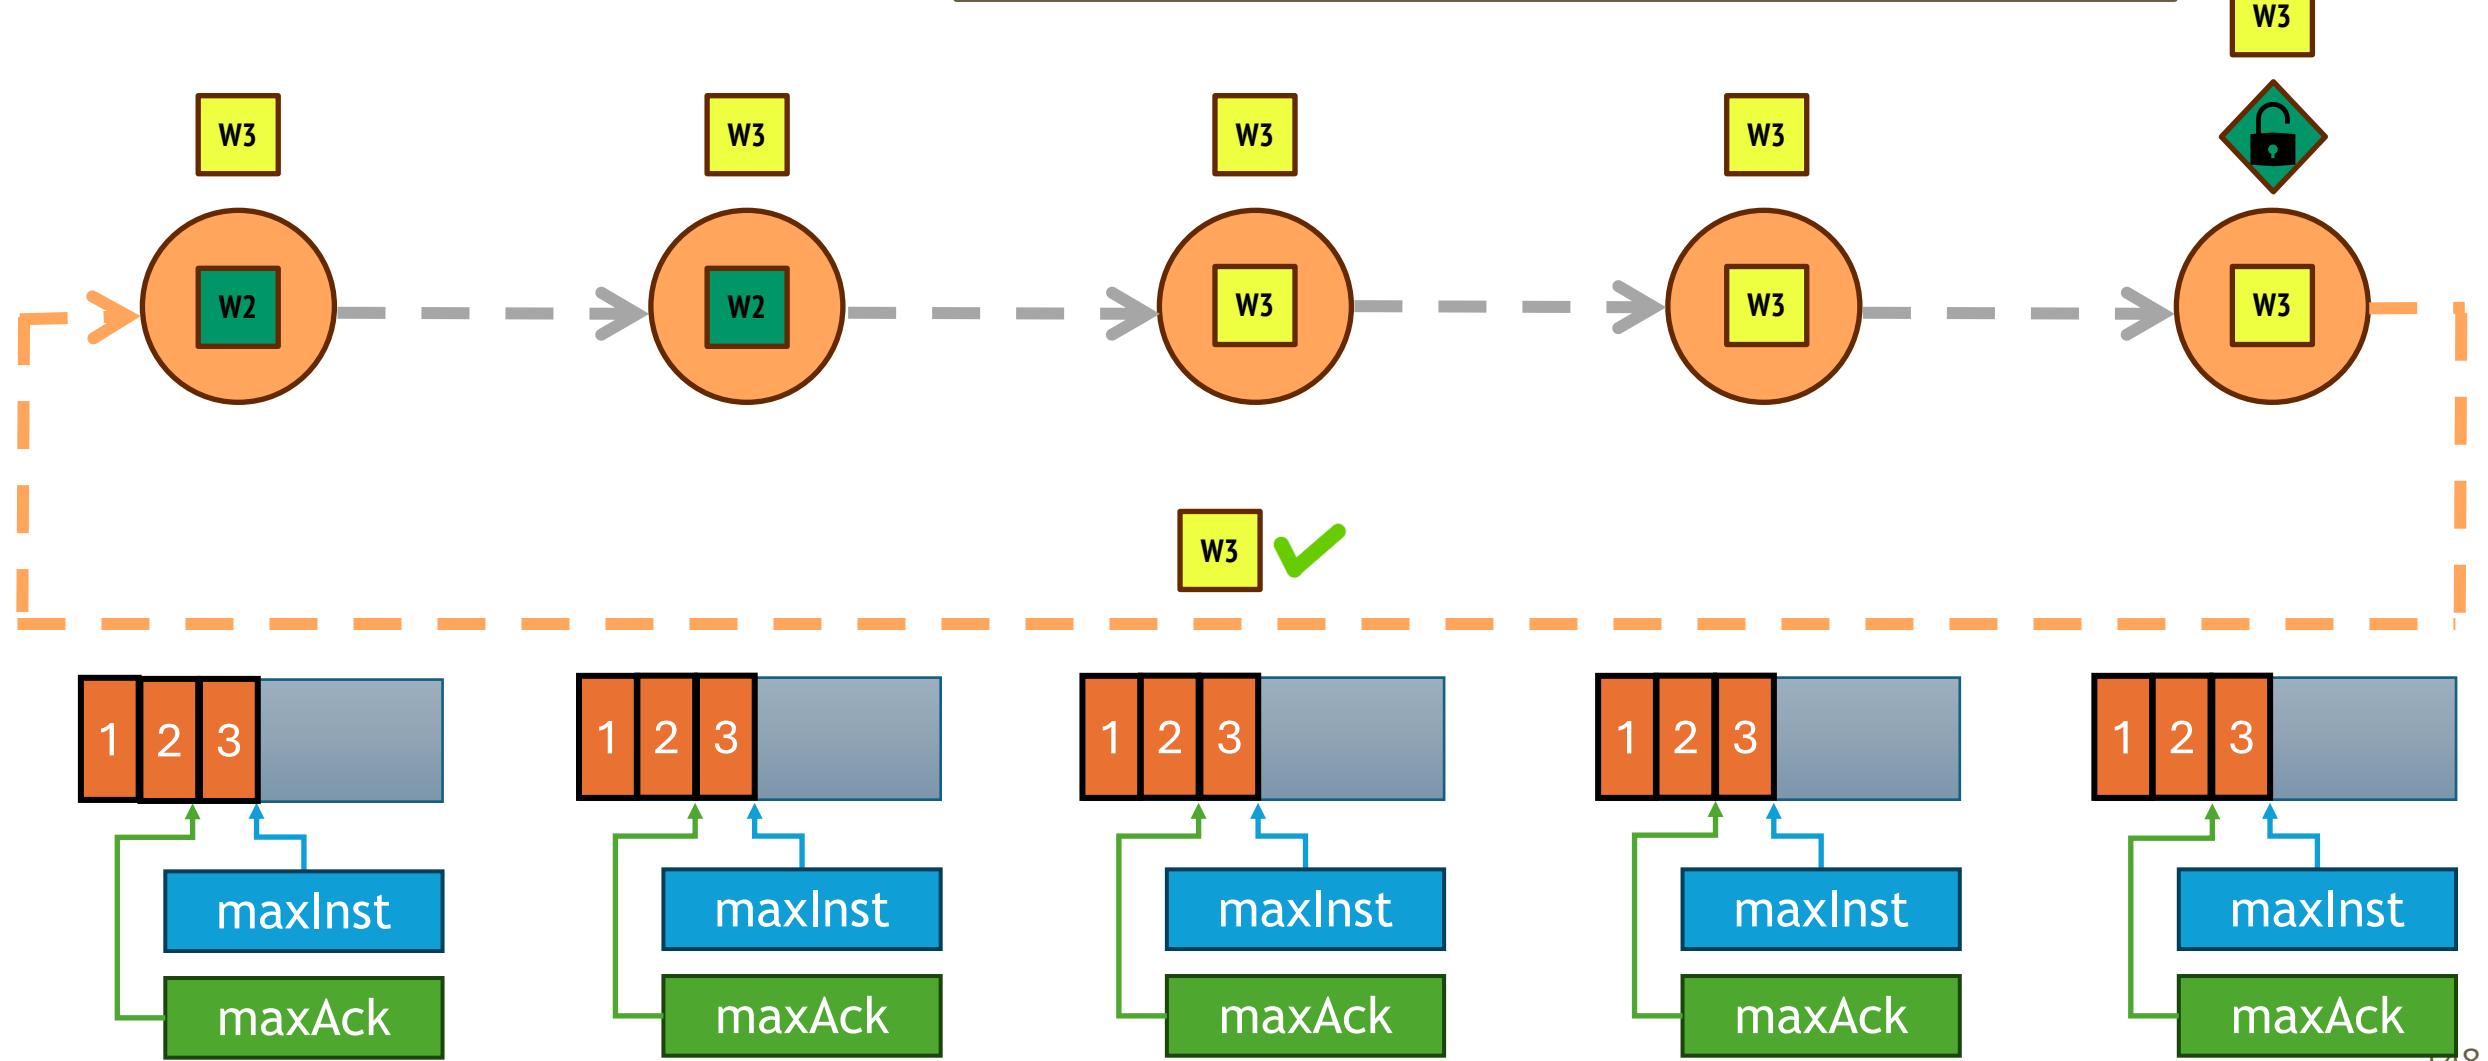


In Depth: Reads

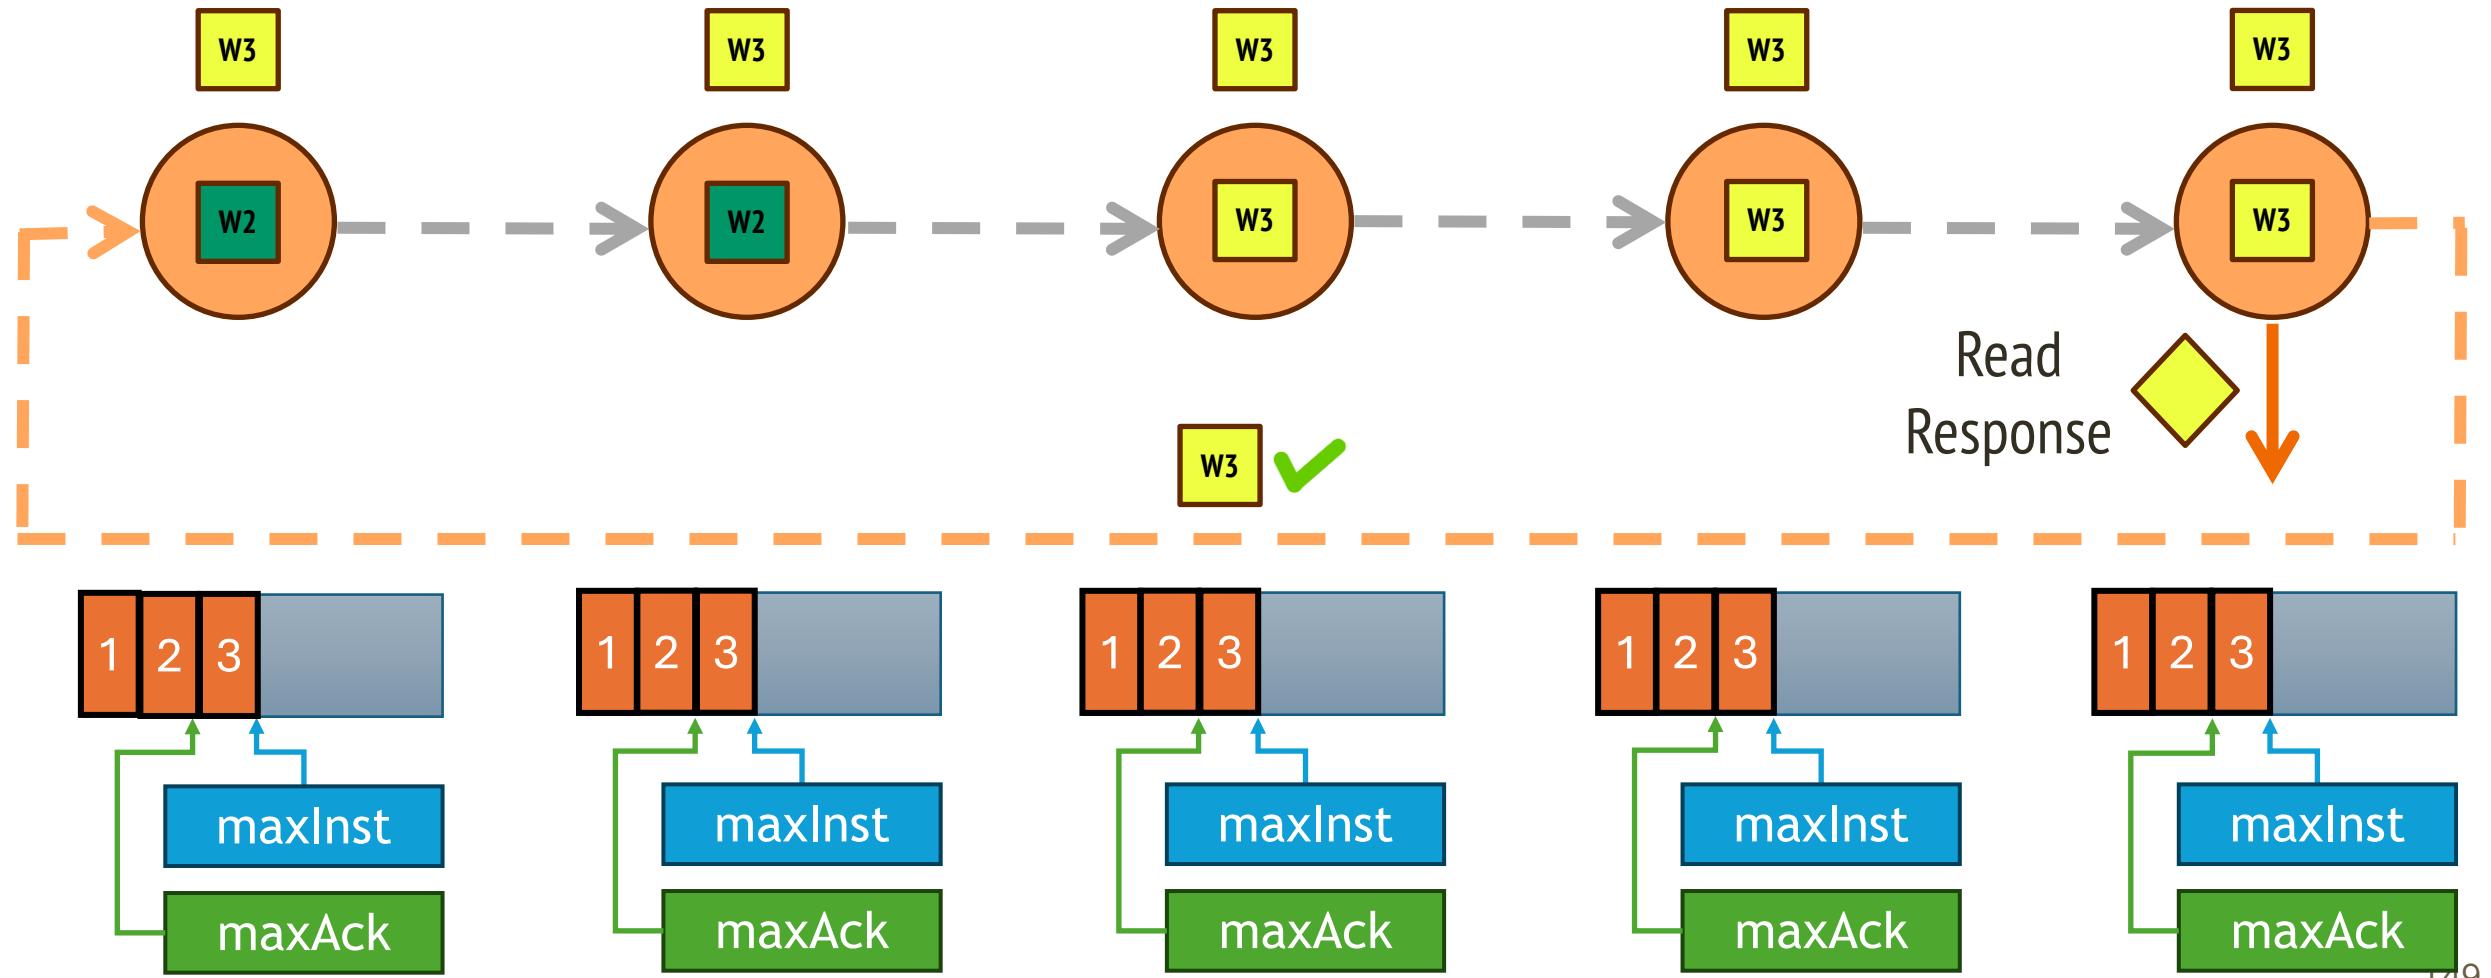


In Depth: Reads

"ack" of next "unseen instance" - W2 ✓



In Depth: Reads



In Depth: Reads

*Reads wait for
“ack” of next “unseen instance”*

In Depth: Reads

Why acks of next unseen instance?

In Depth: Reads

Why acks of next unseen instance?

Why not immediately reply?

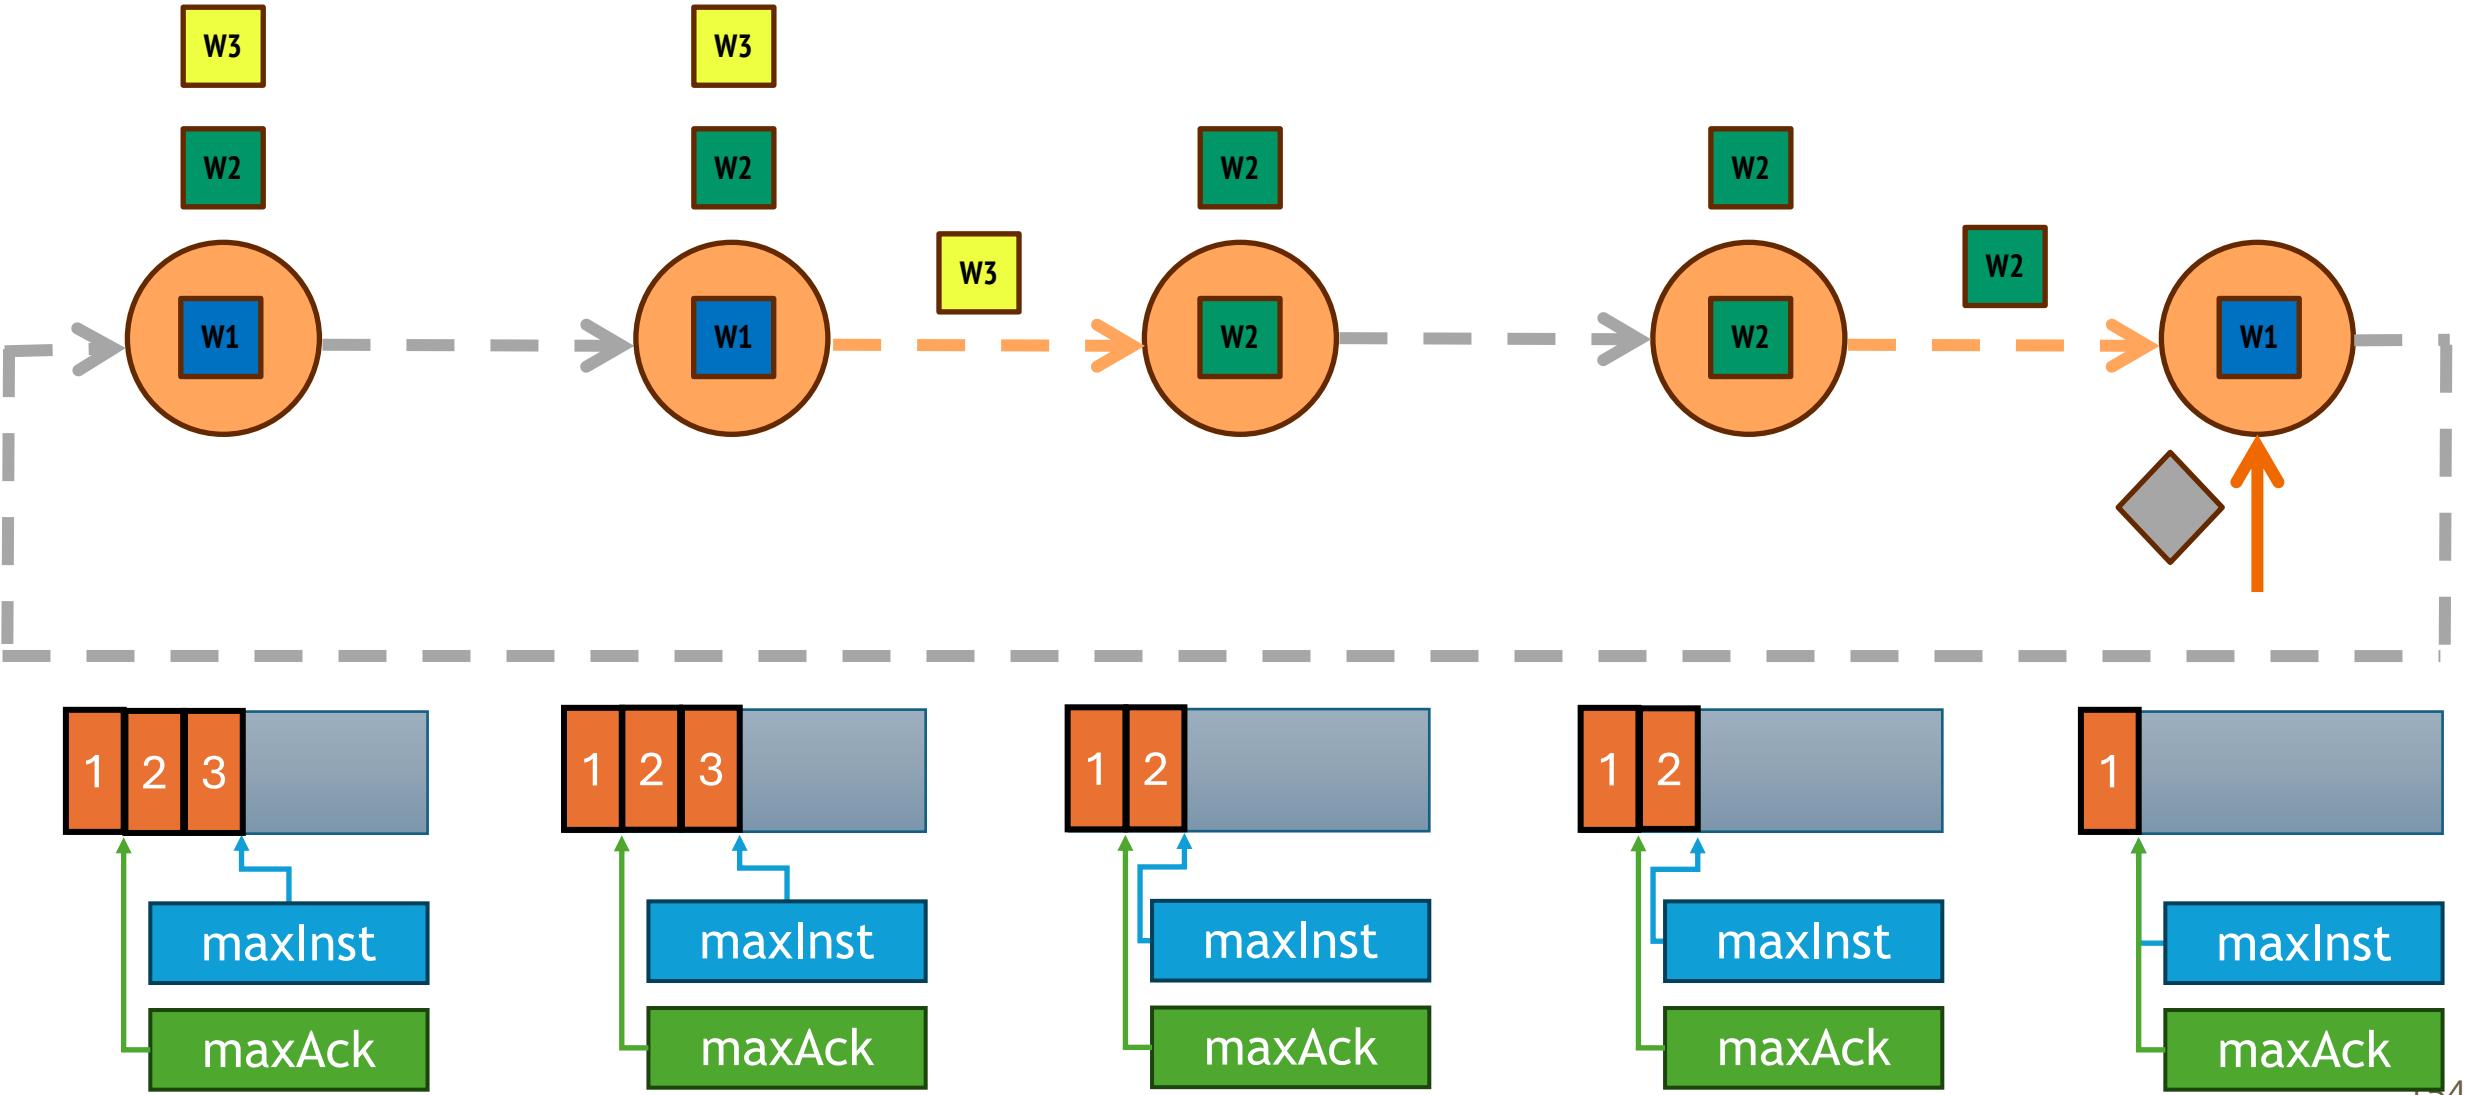
In Depth: Reads

Why acks of next unseen instance?

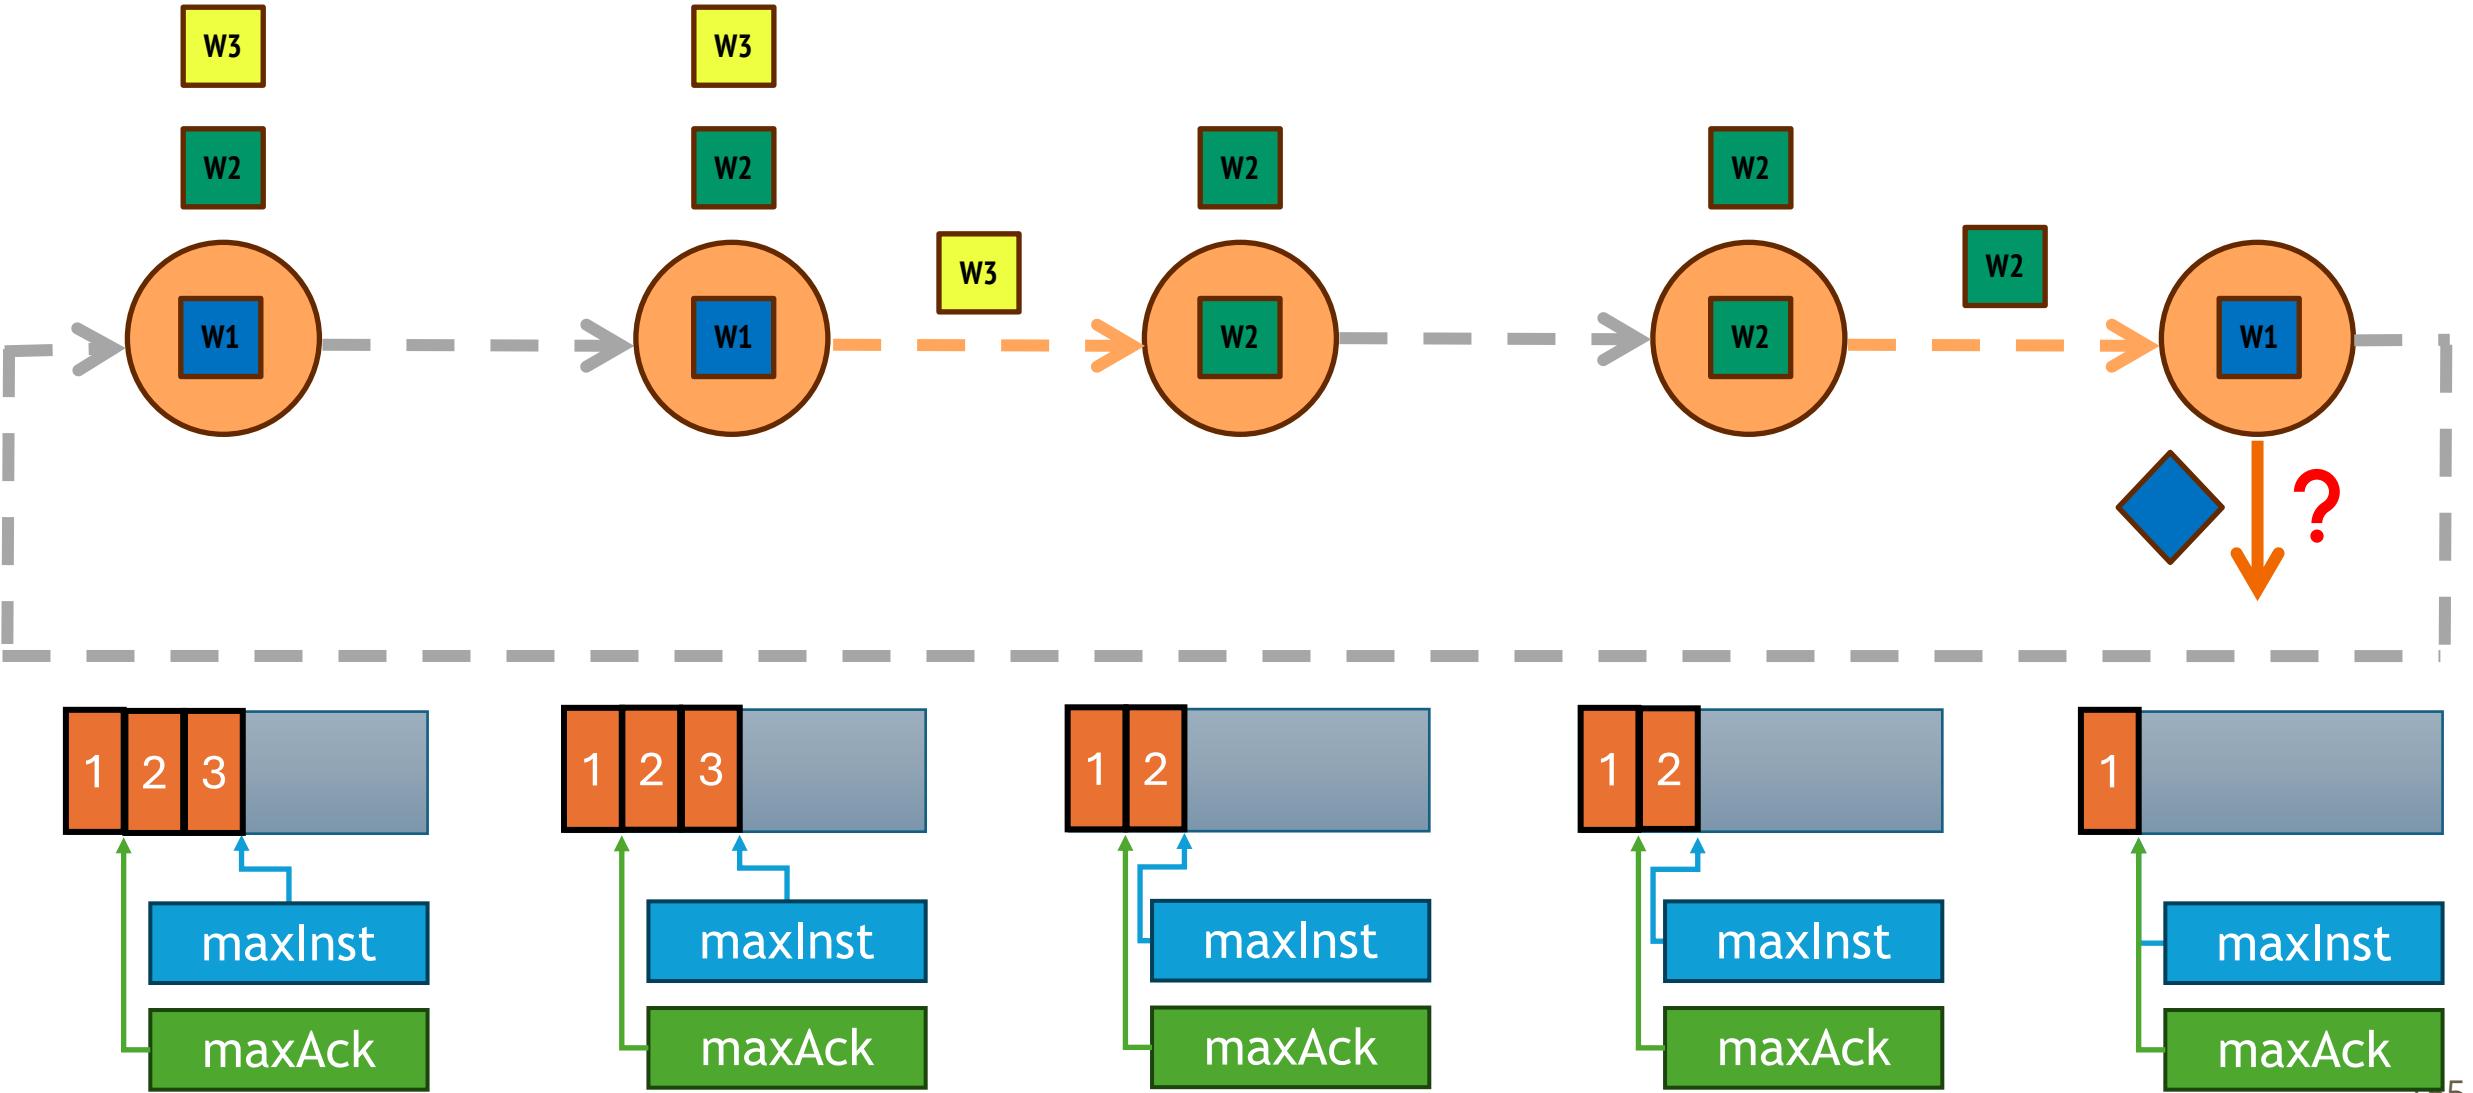
Why not immediately reply?

Replicas cannot be sure whether other replicas have replied to new writes

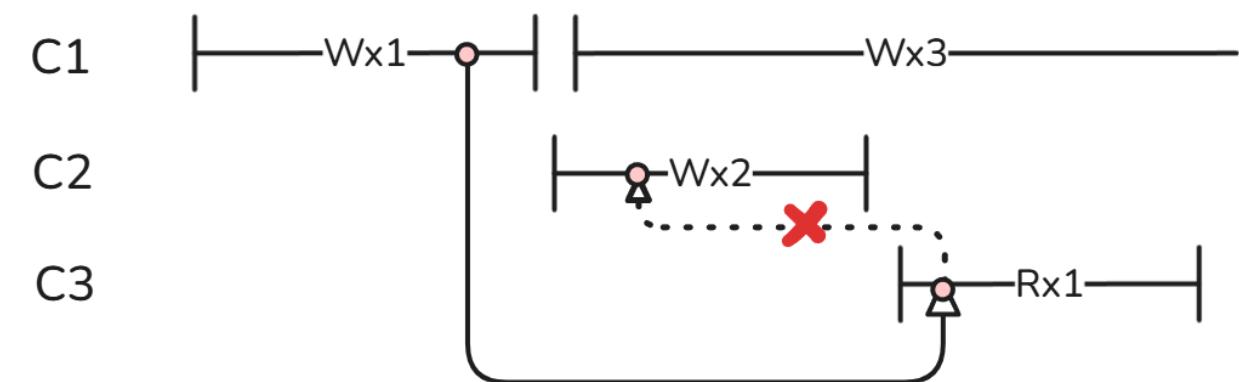
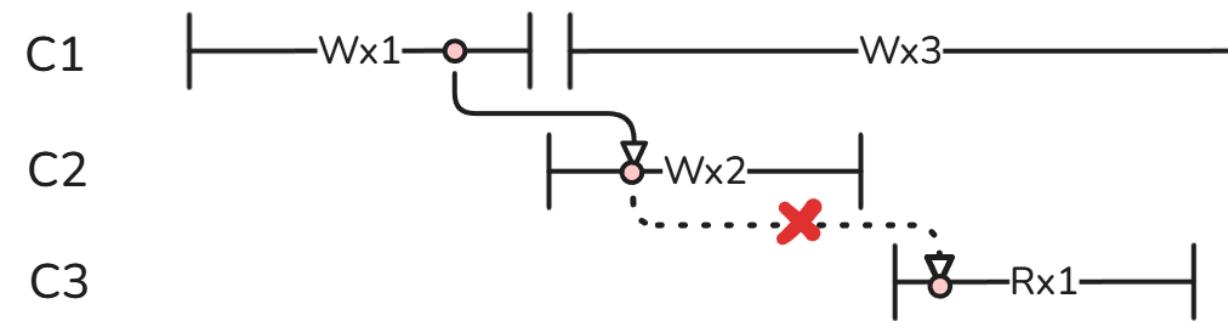
In Depth: Why not immediately reply?



In Depth: Why not immediately reply?



In Depth: Why not immediately reply?



Not Linearizable



In Depth: Reads

Why acks of next unseen instance?

~~*Why not immediately reply?*~~

In Depth: Reads

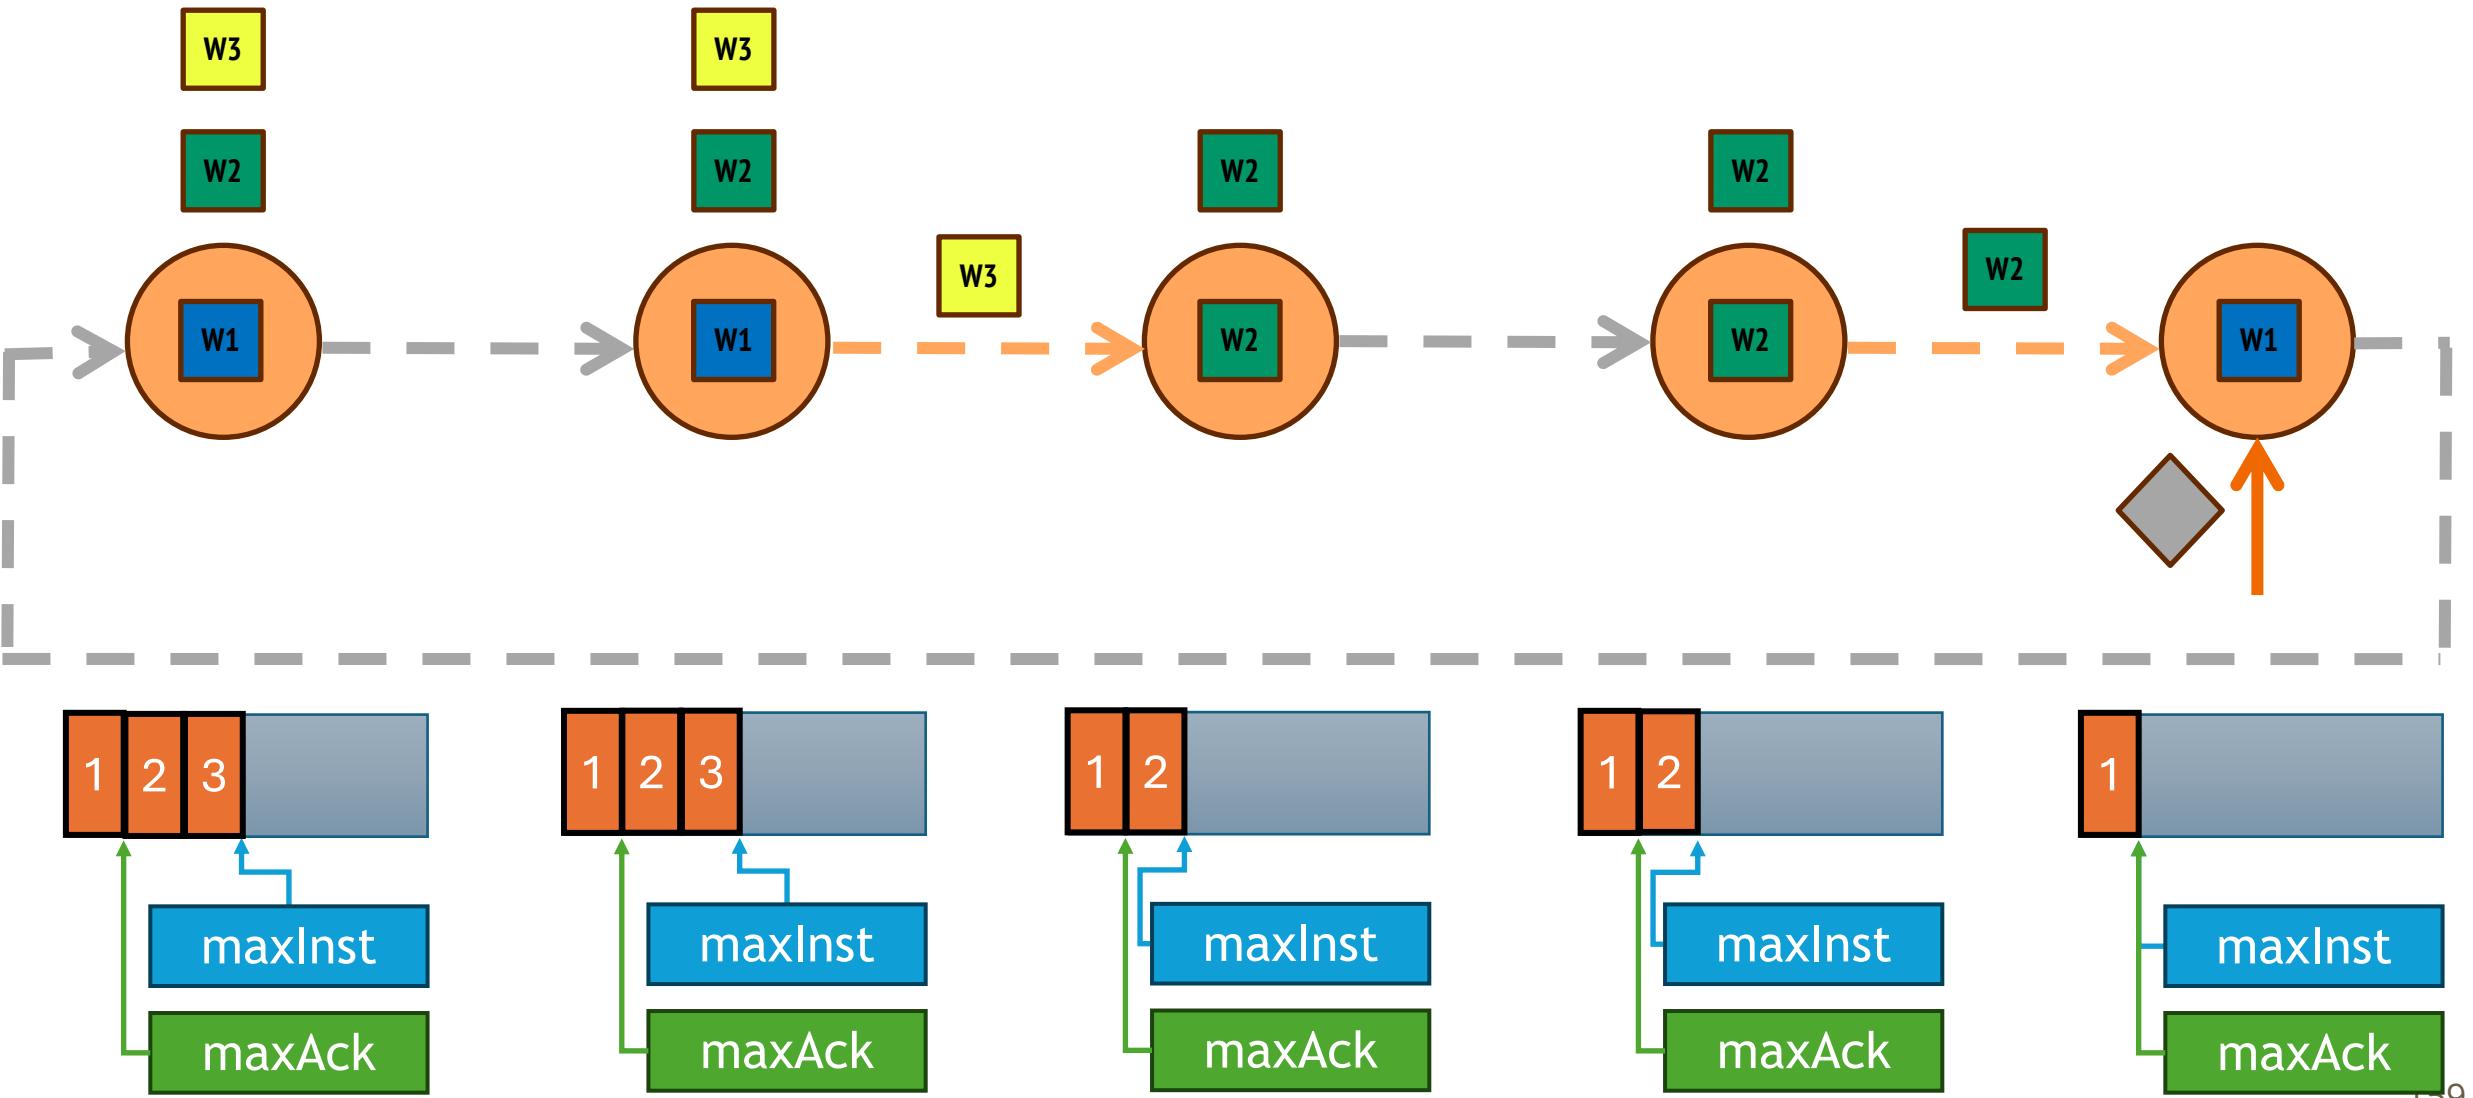
Why acks of next unseen instance?

~~*Why not immediately reply?*~~

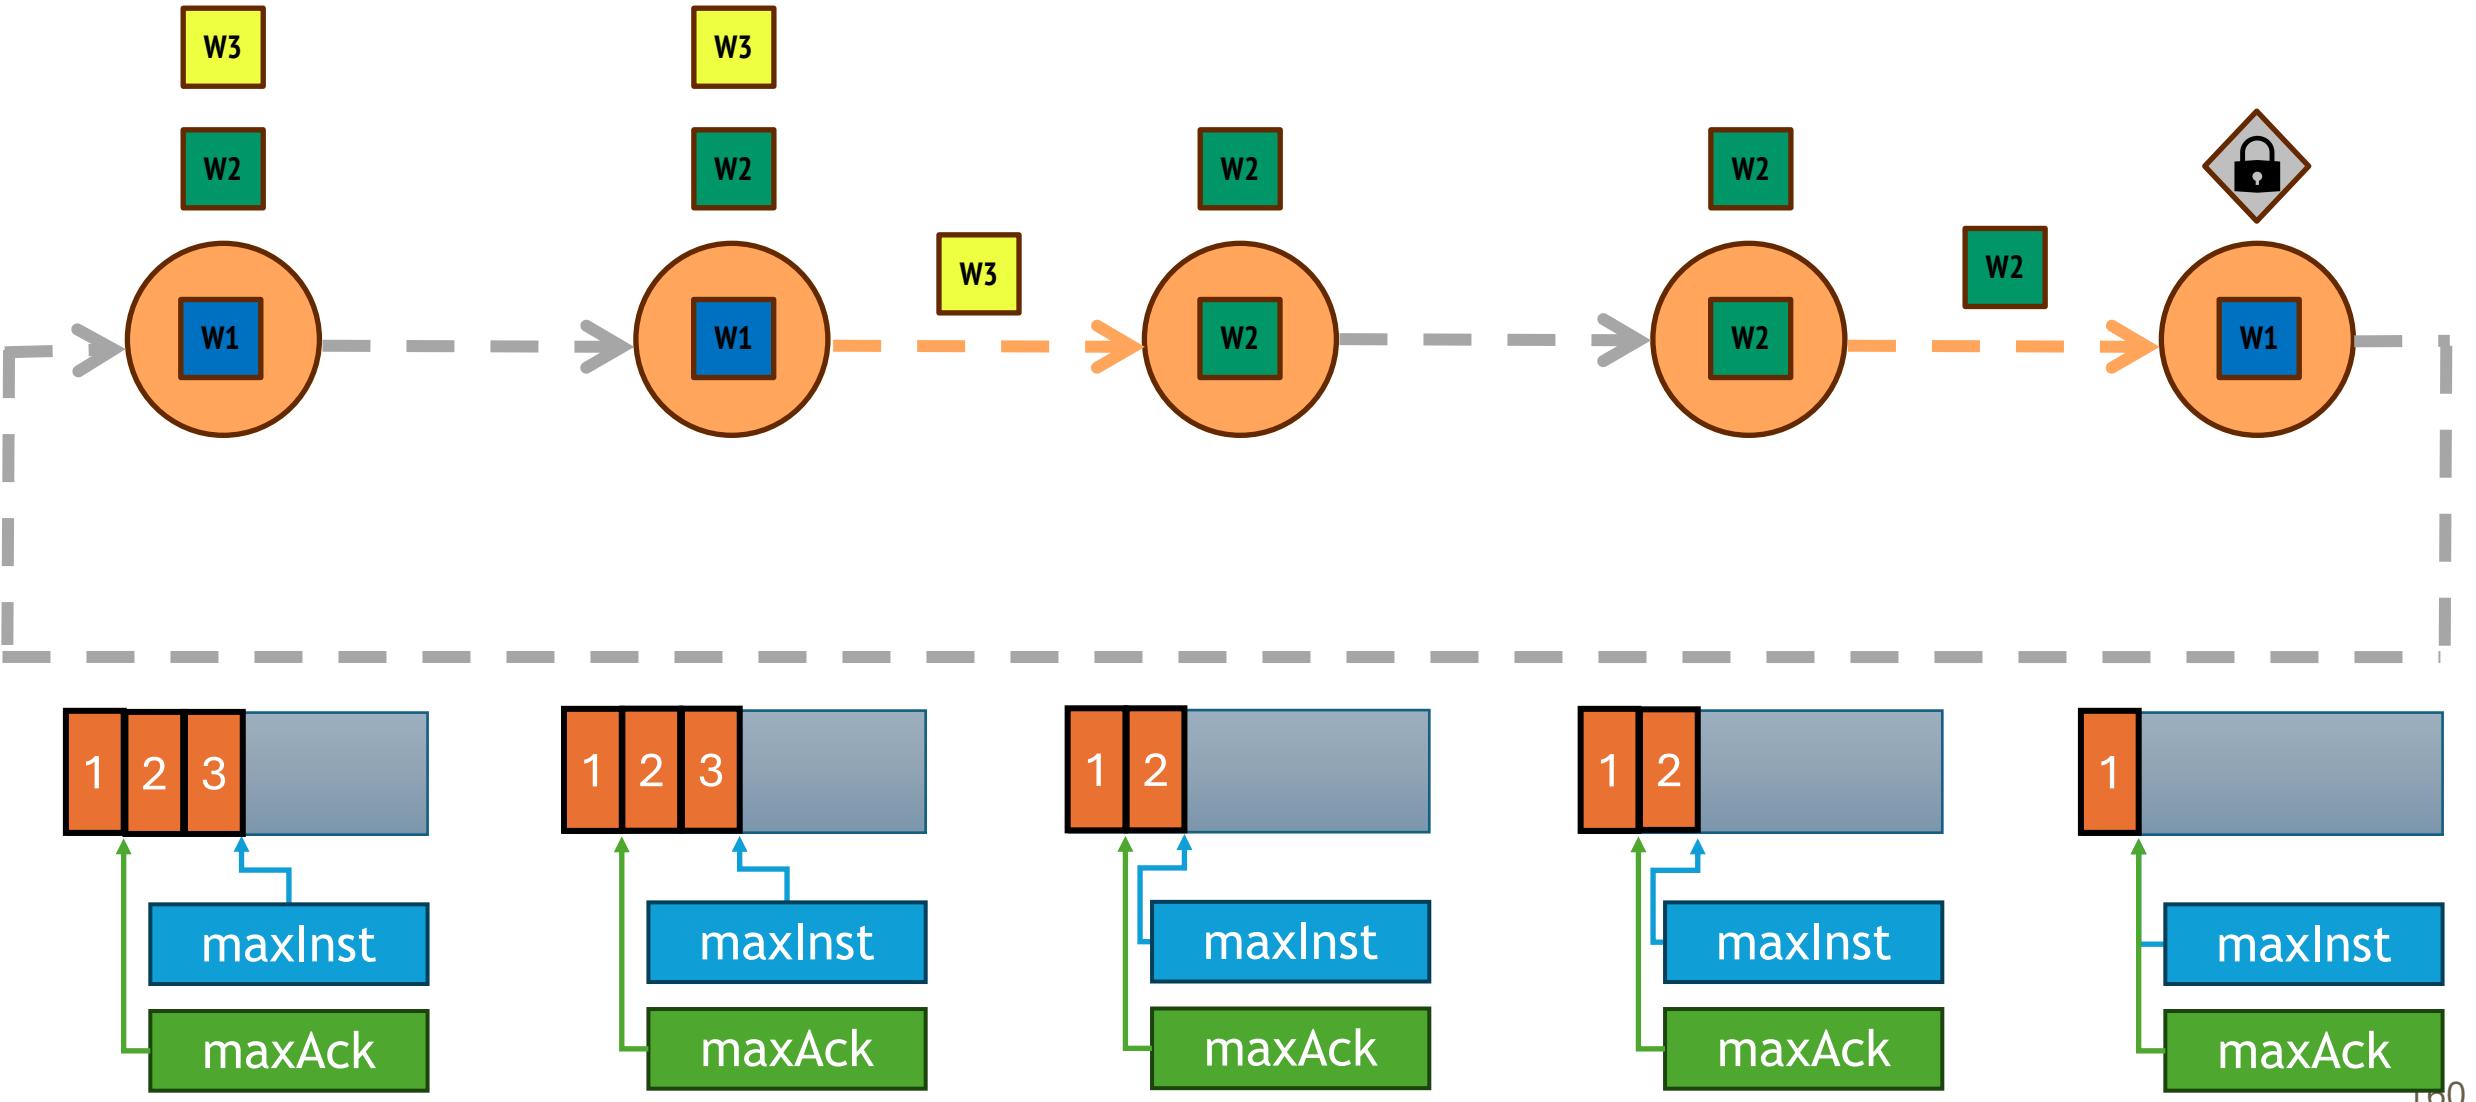
Why not next unseen instance?

Again, replicas cannot be sure whether other replicas have replied to new writes

In Depth: Why not next unseen instance?

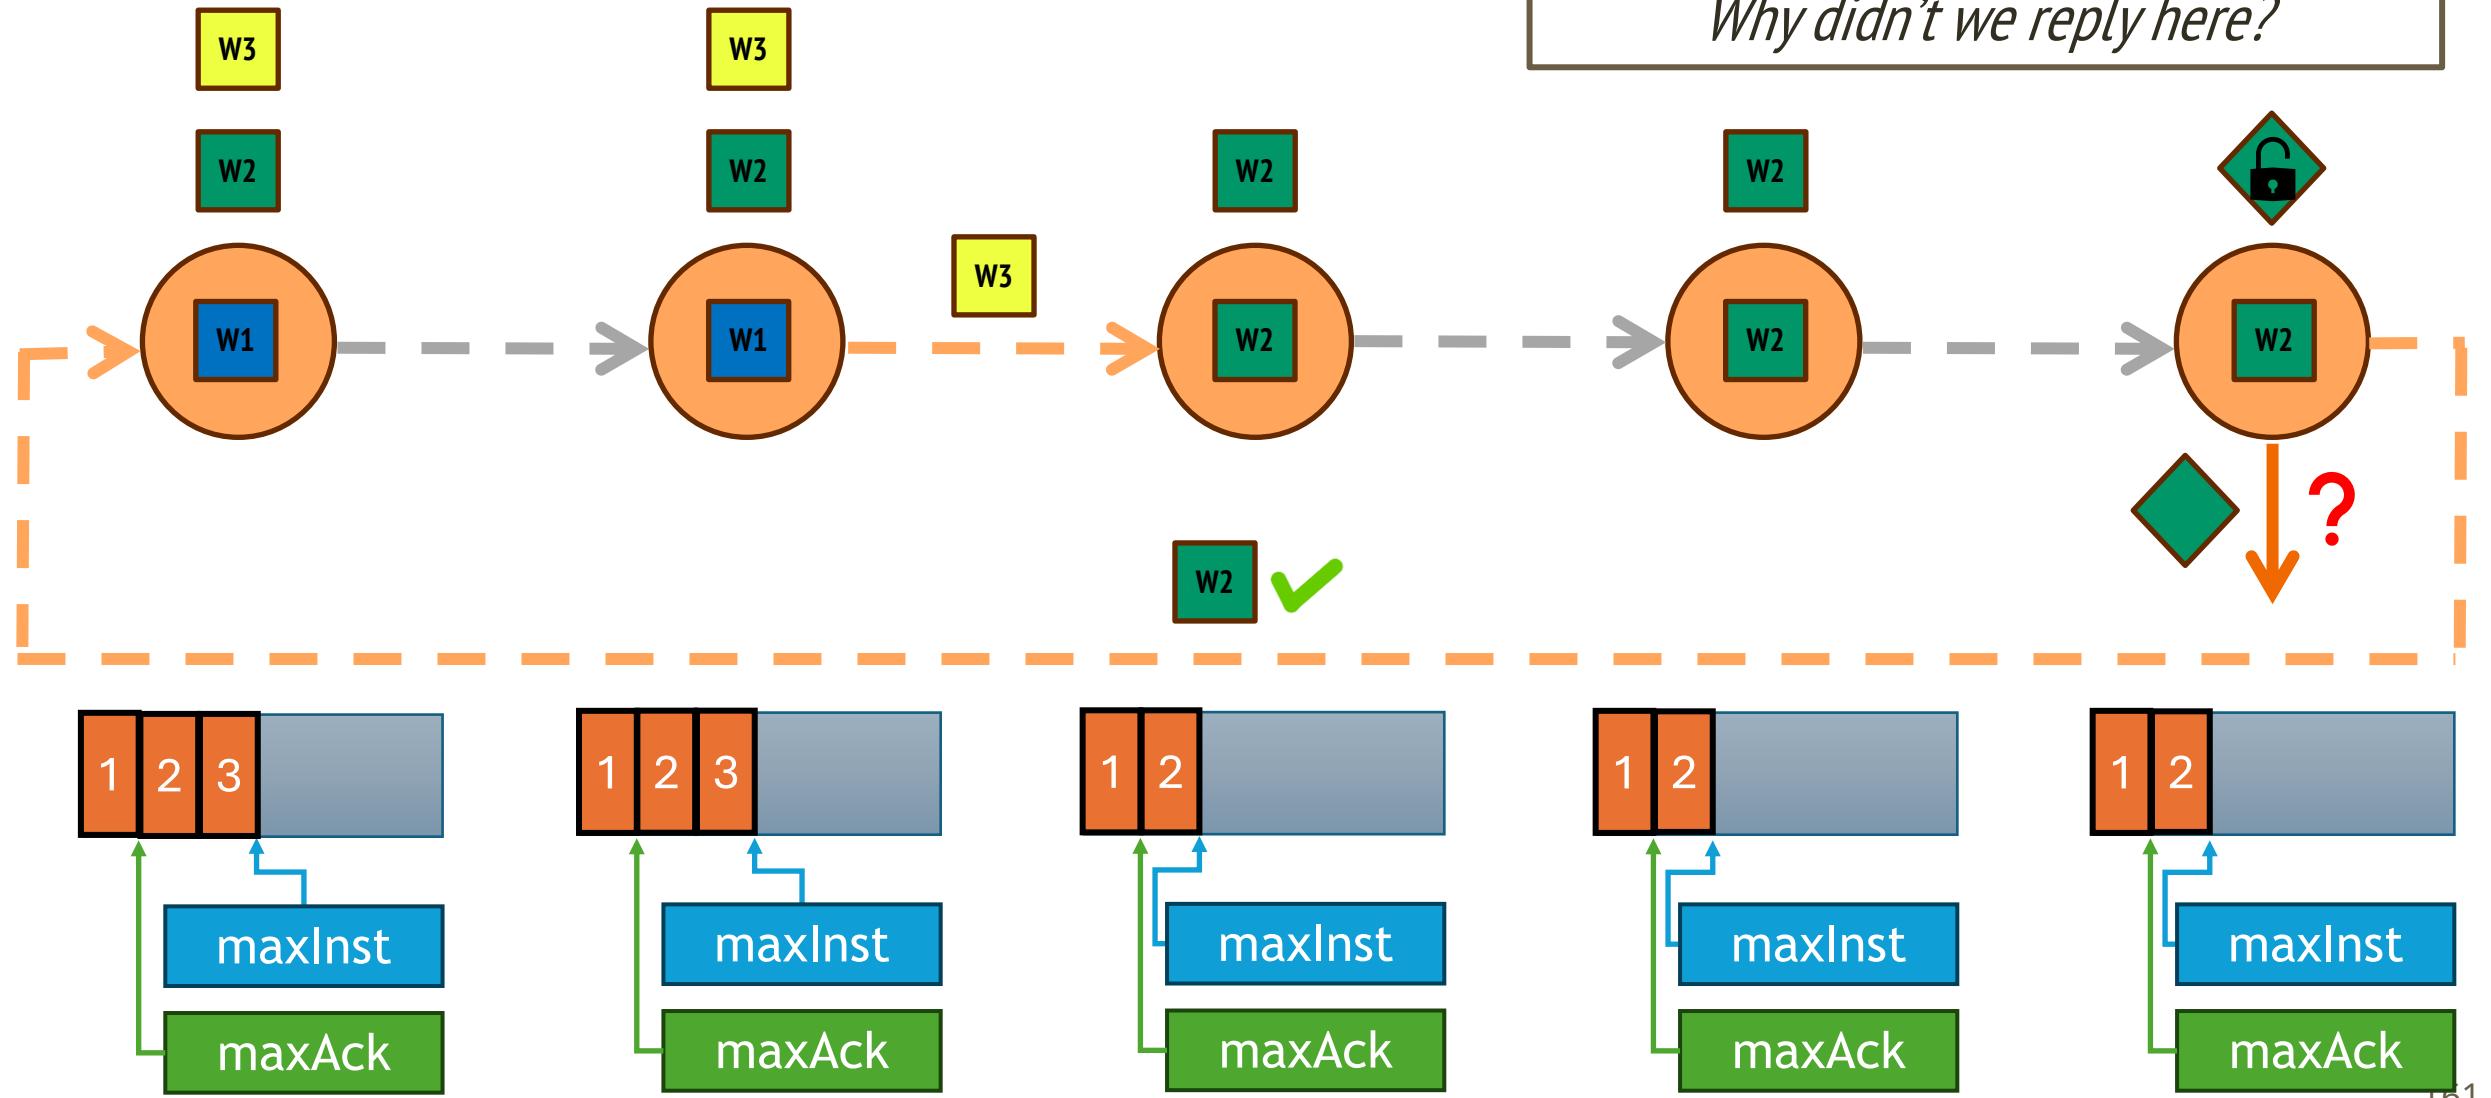


In Depth: Why not next unseen instance?



In Depth: Why not next unseen instance?

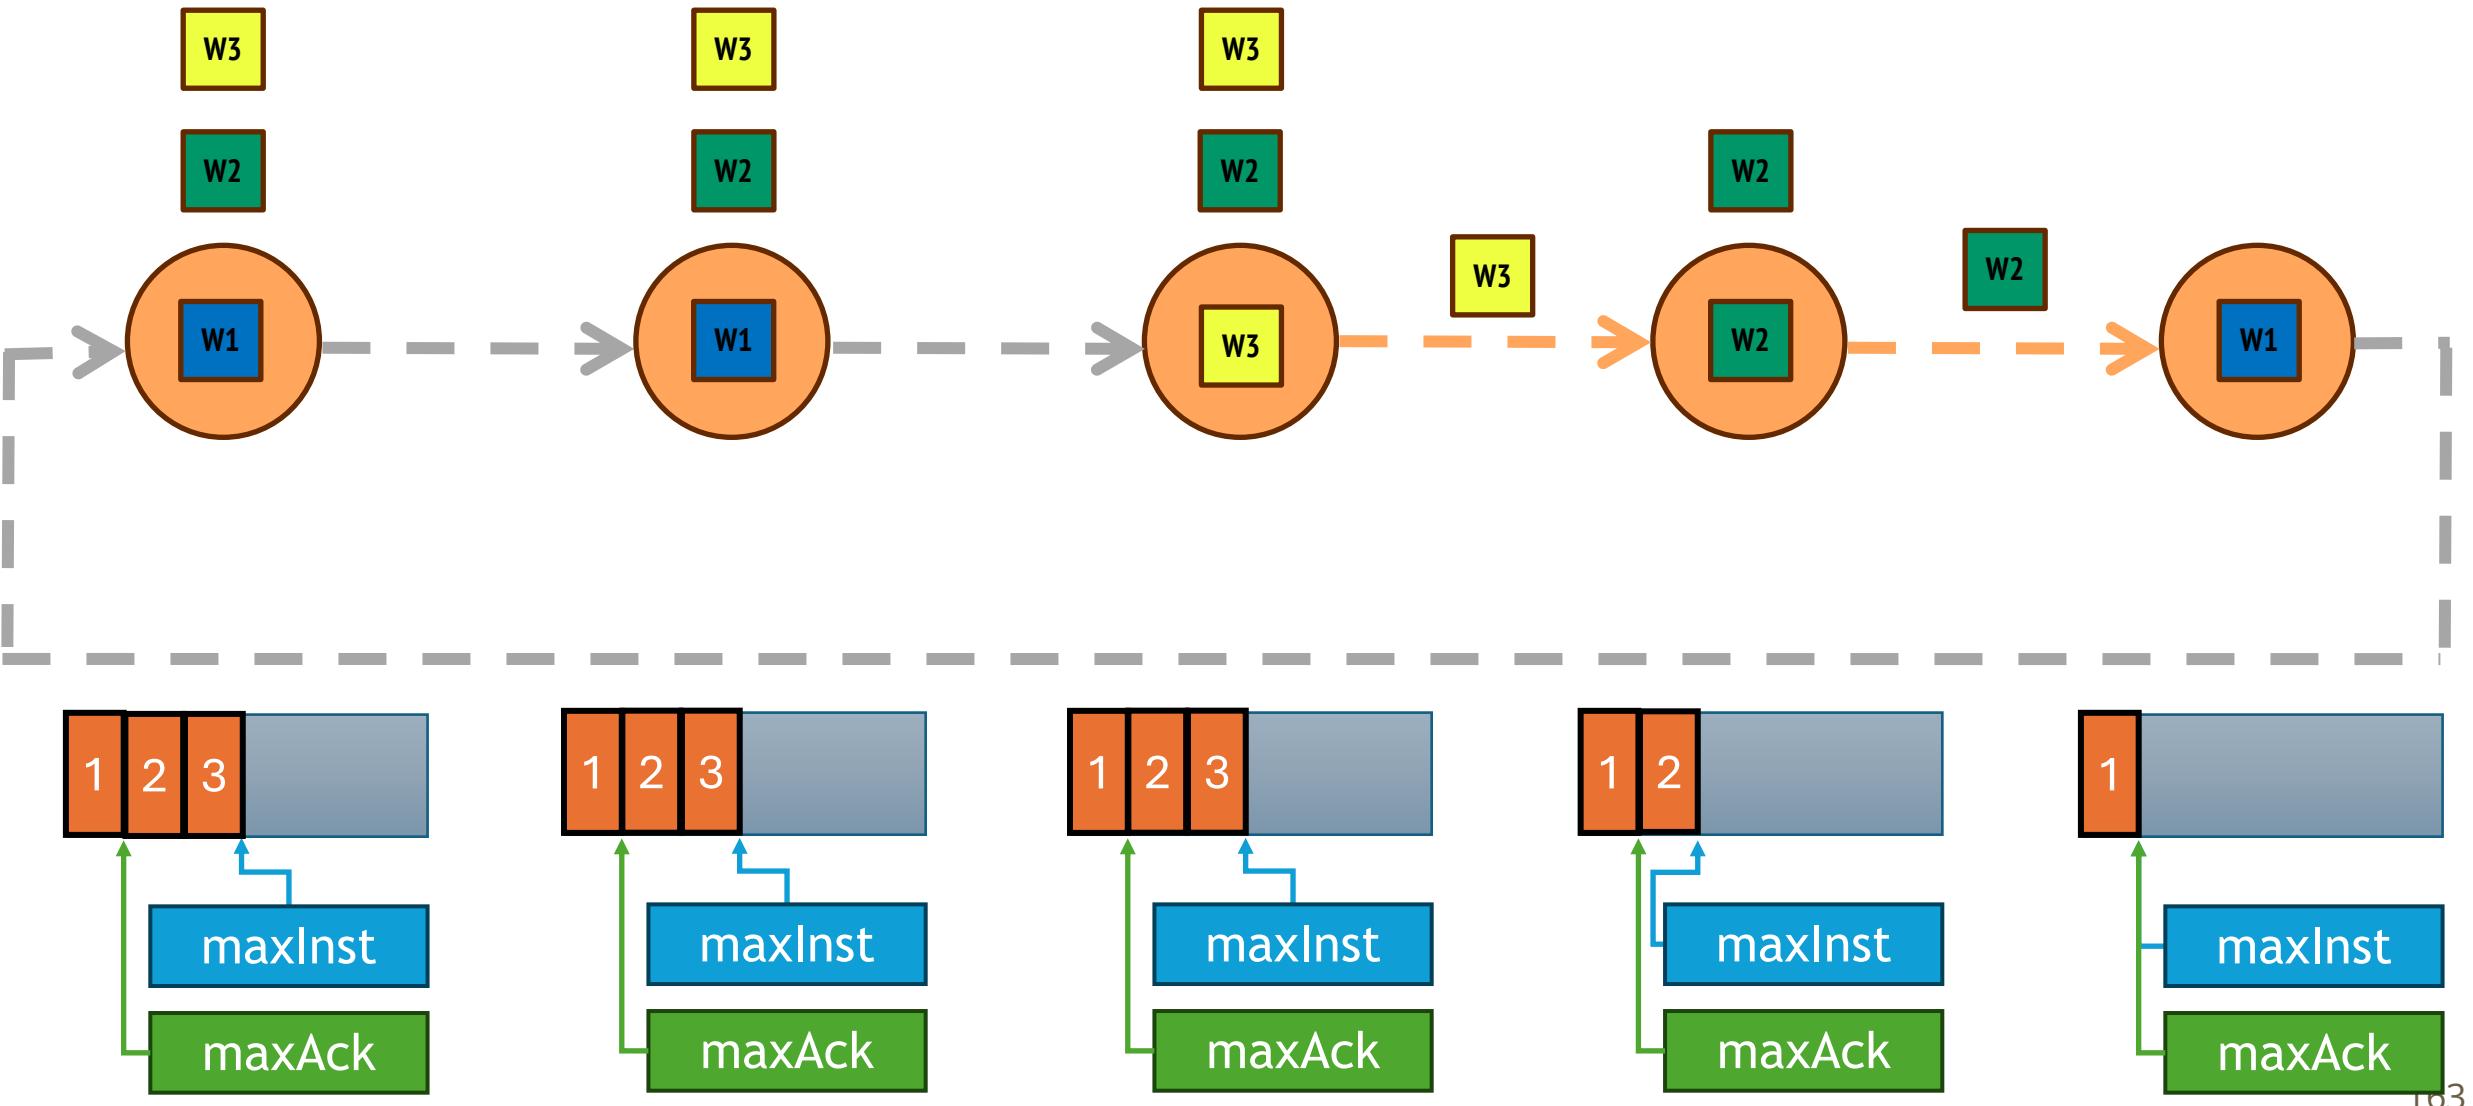
Why didn't we reply here?



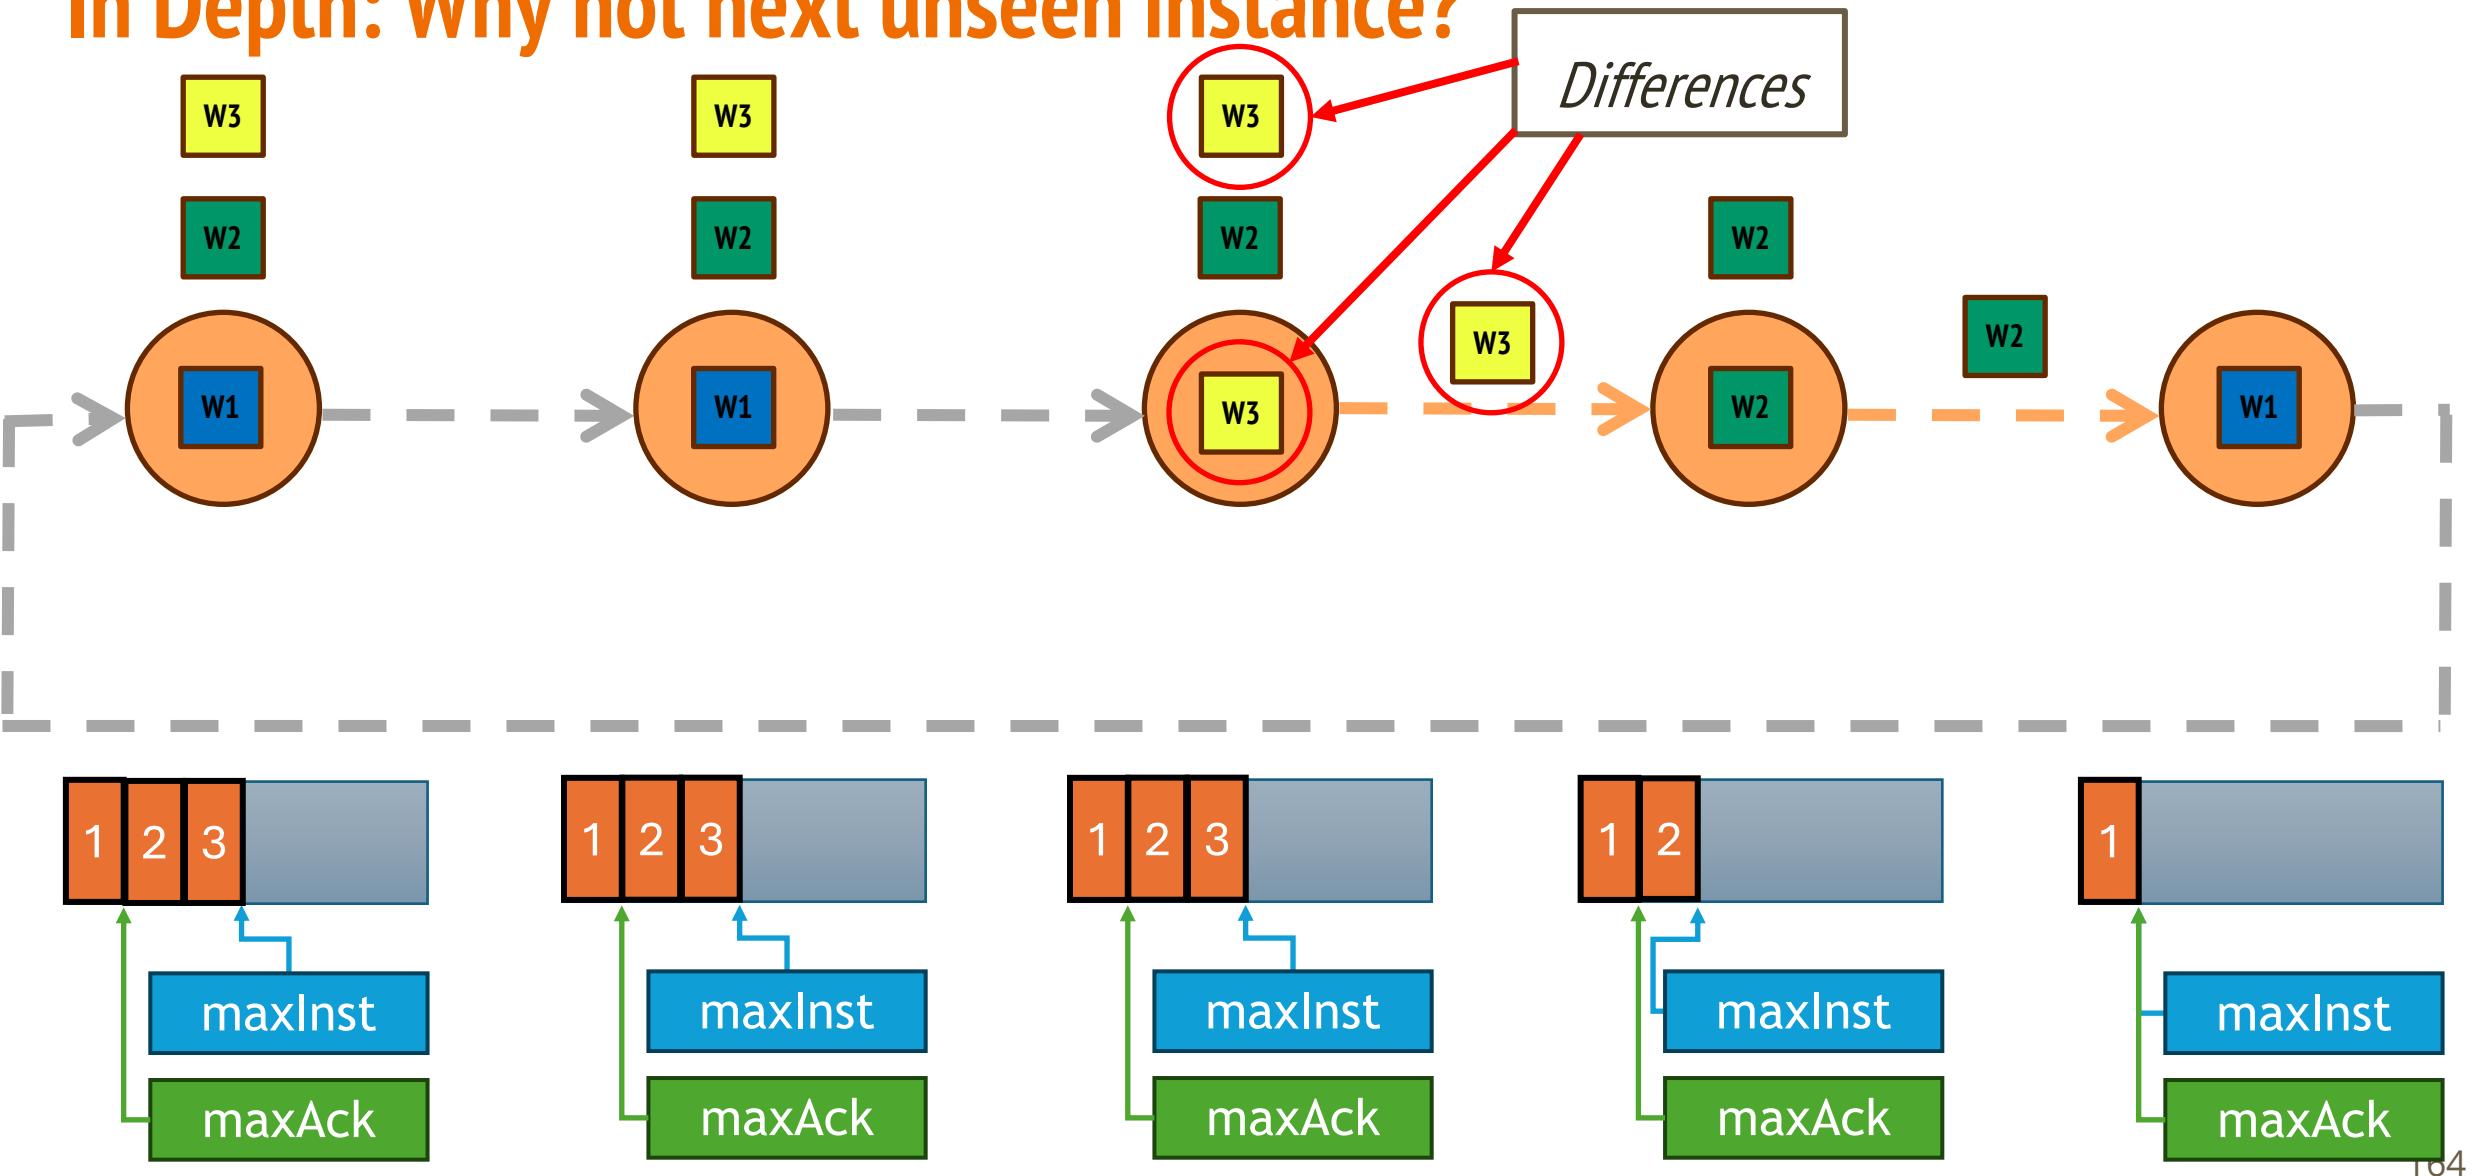
In Depth: Why not next unseen instance?

*Consider a slightly different
situation of the system ...*

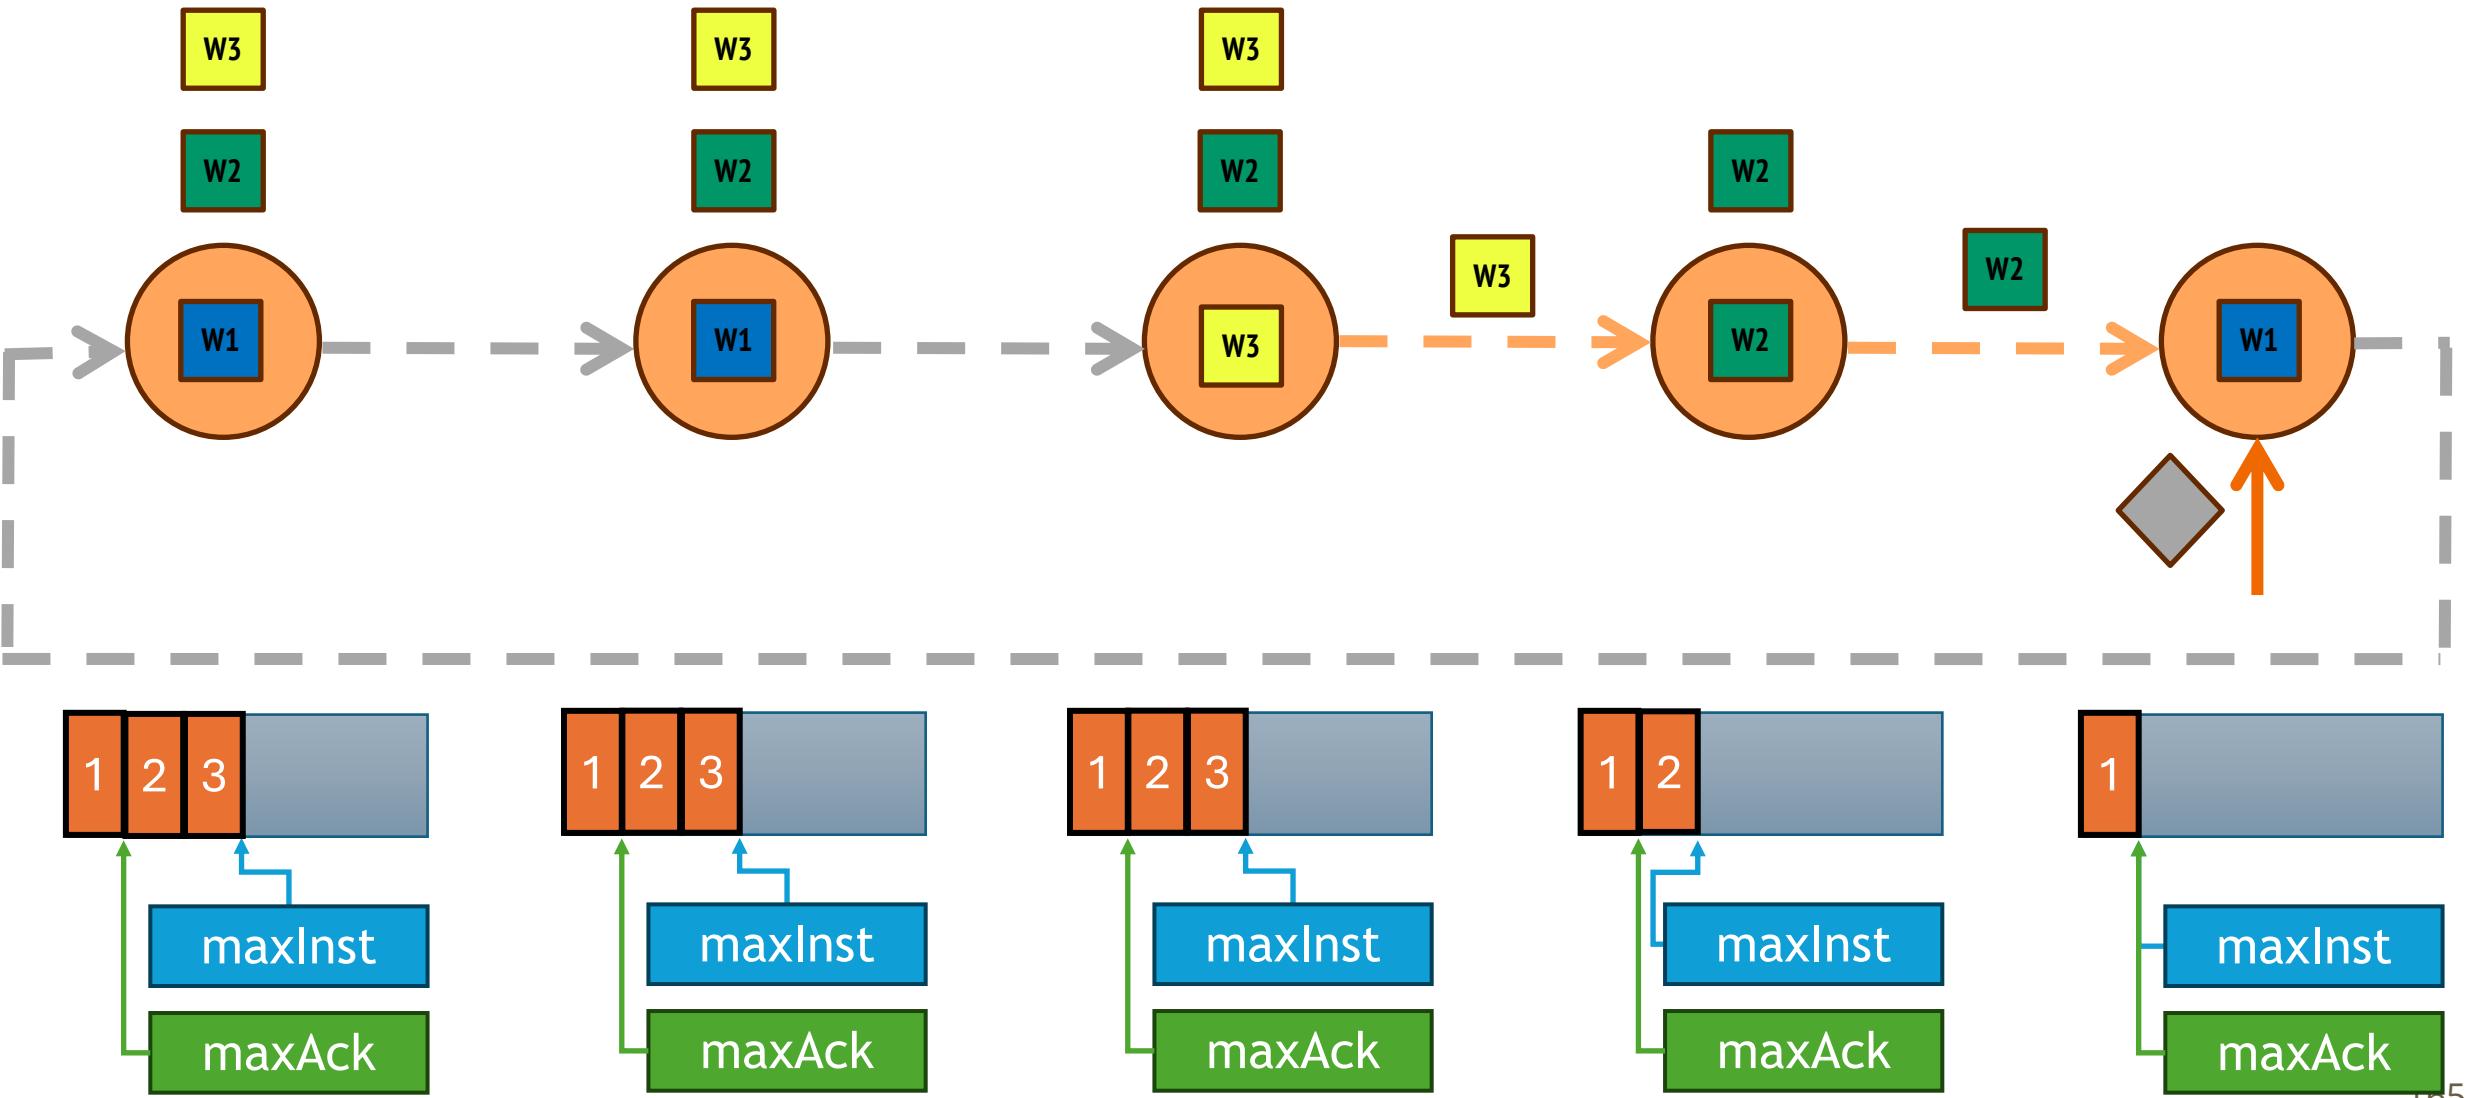
In Depth: Why not next unseen instance?



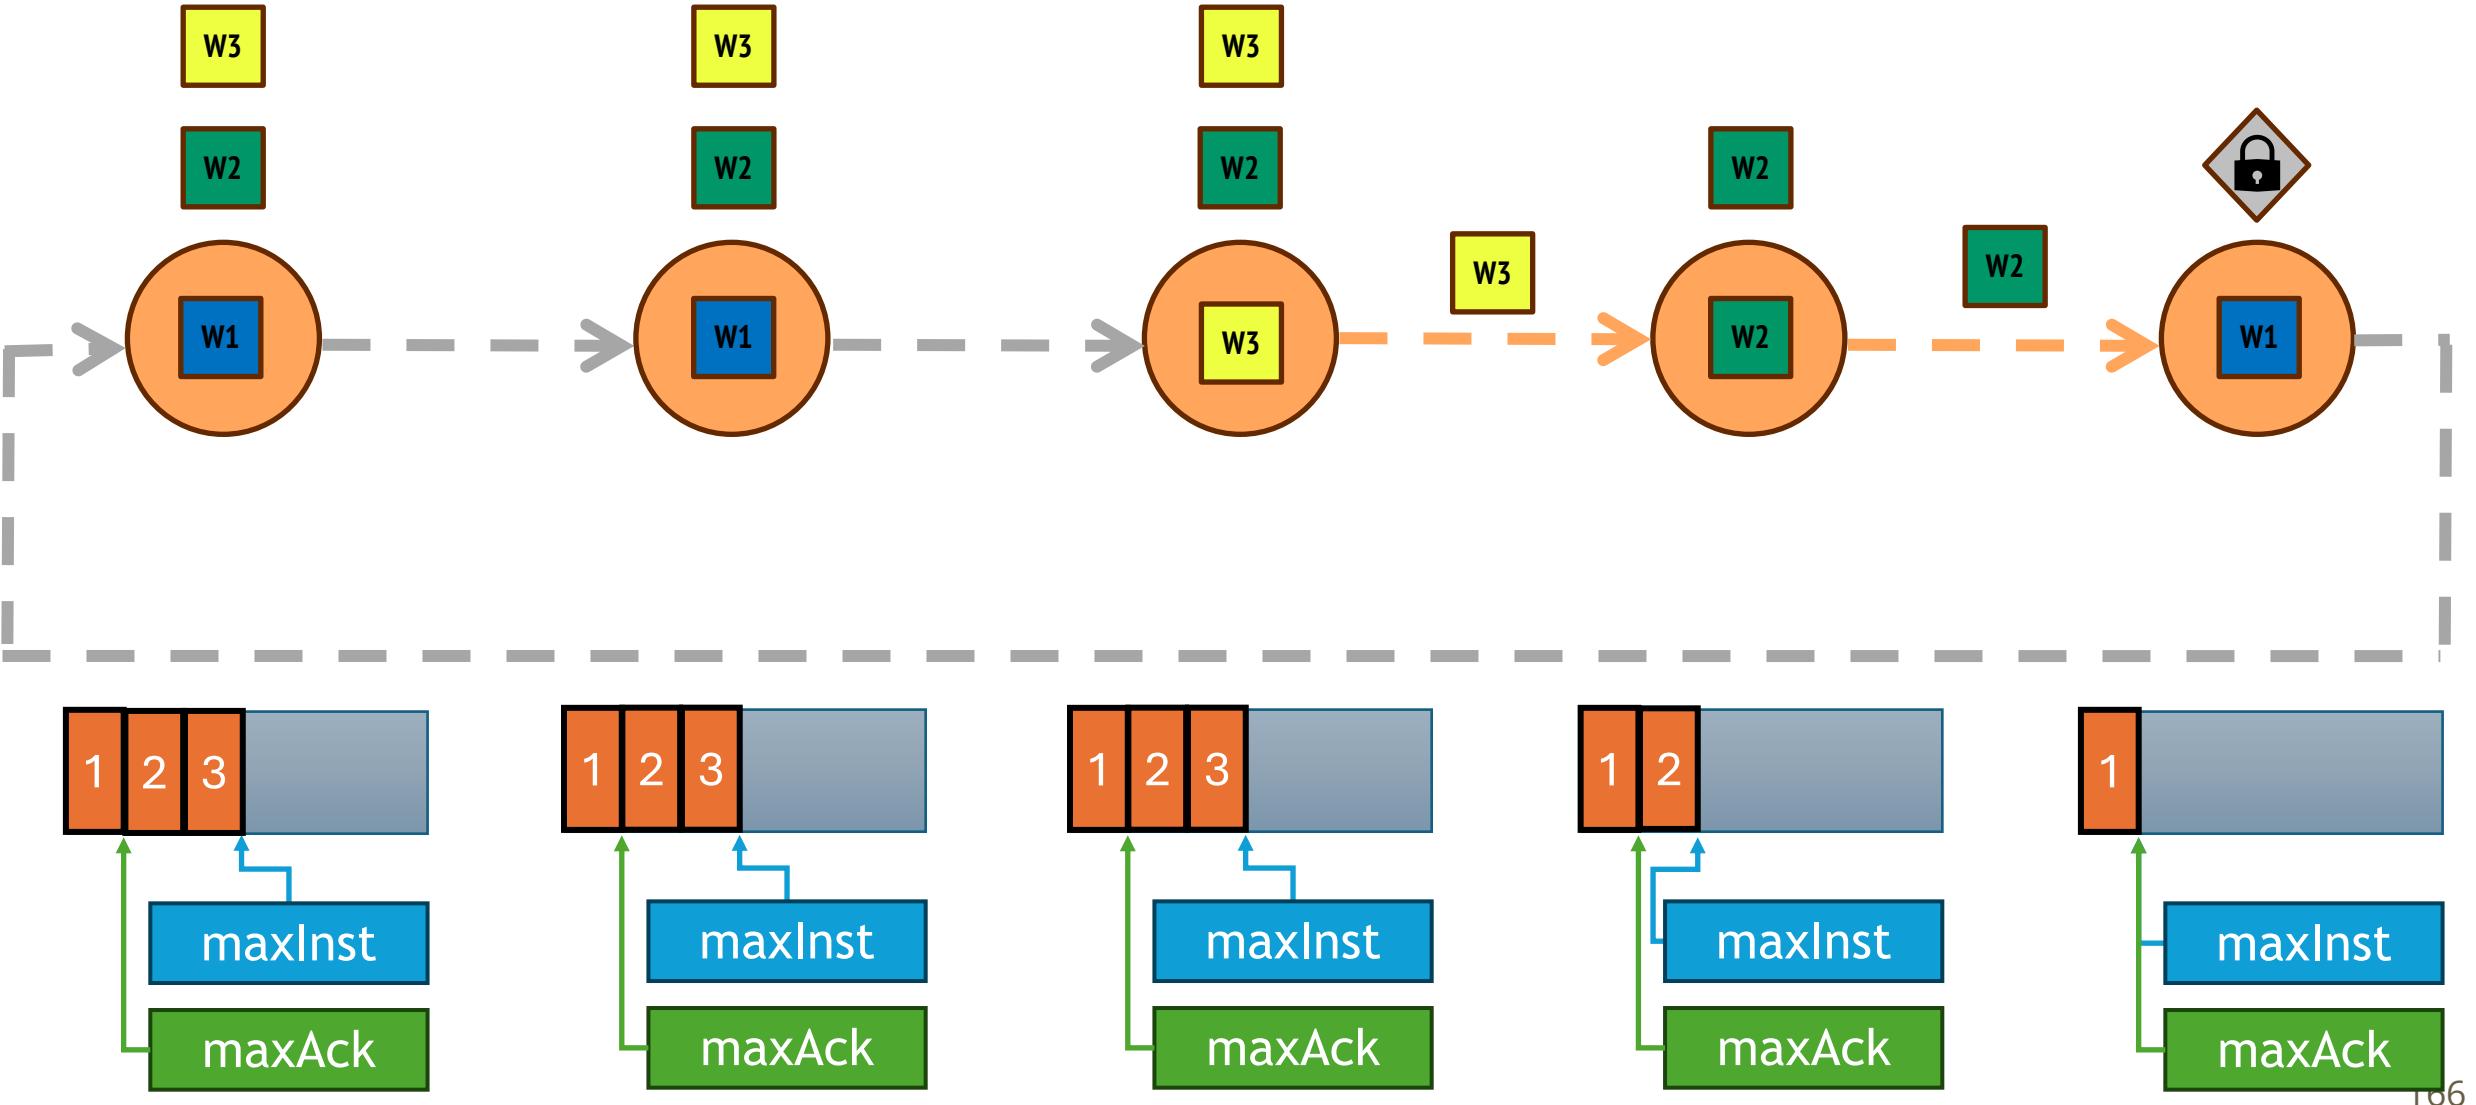
In Depth: Why not next unseen instance?



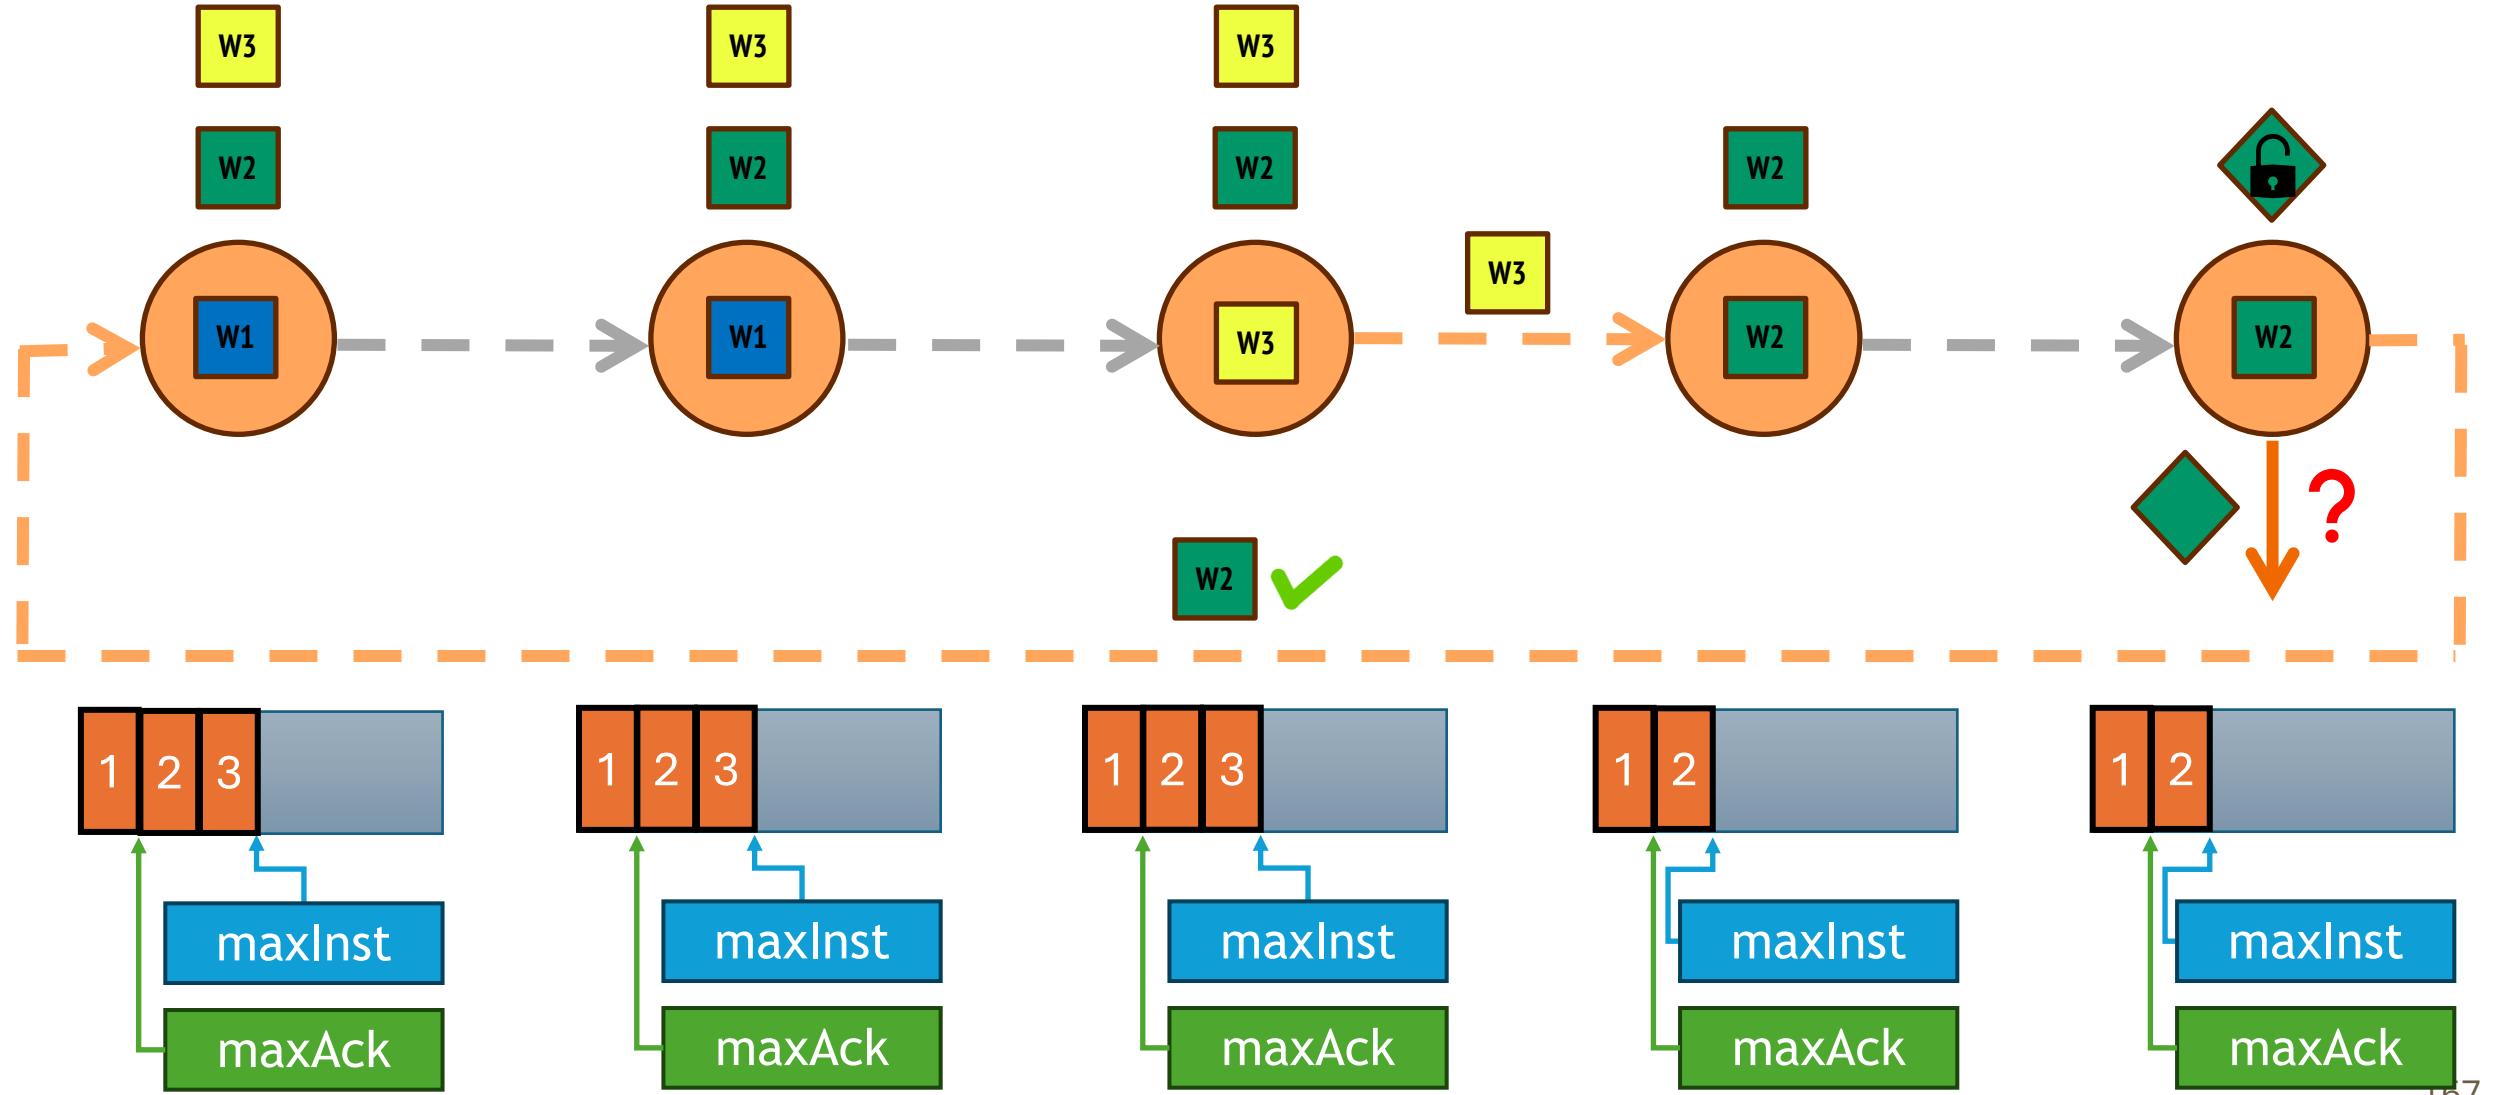
In Depth: Why not next unseen instance?



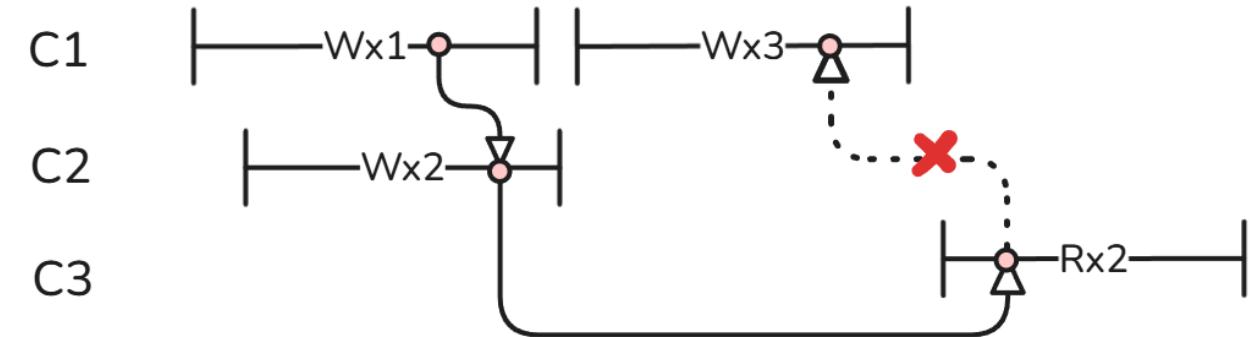
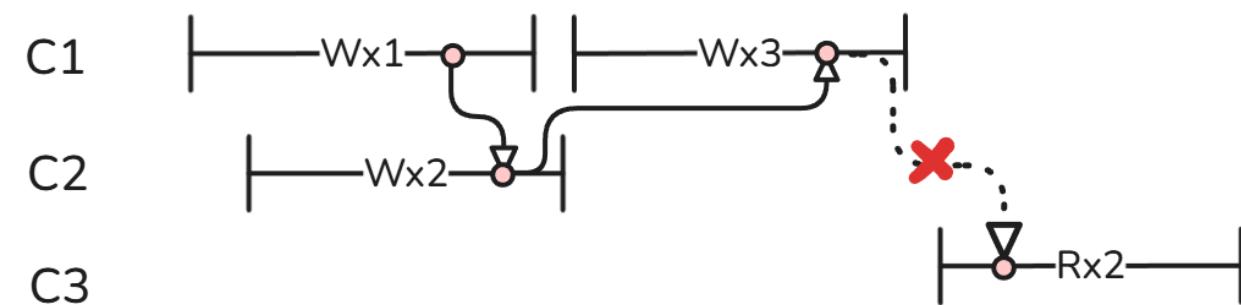
In Depth: Why not next unseen instance?



In Depth: Why not next unseen instance?



In Depth: Why not next unseen instance?



Not Linearizable



In Depth: Reads

Why acks of next unseen instance?

~~*Why not immediately reply?*~~

~~*Why not next unseen instance?*~~

In Depth: Reads

Why acks of next unseen instance?

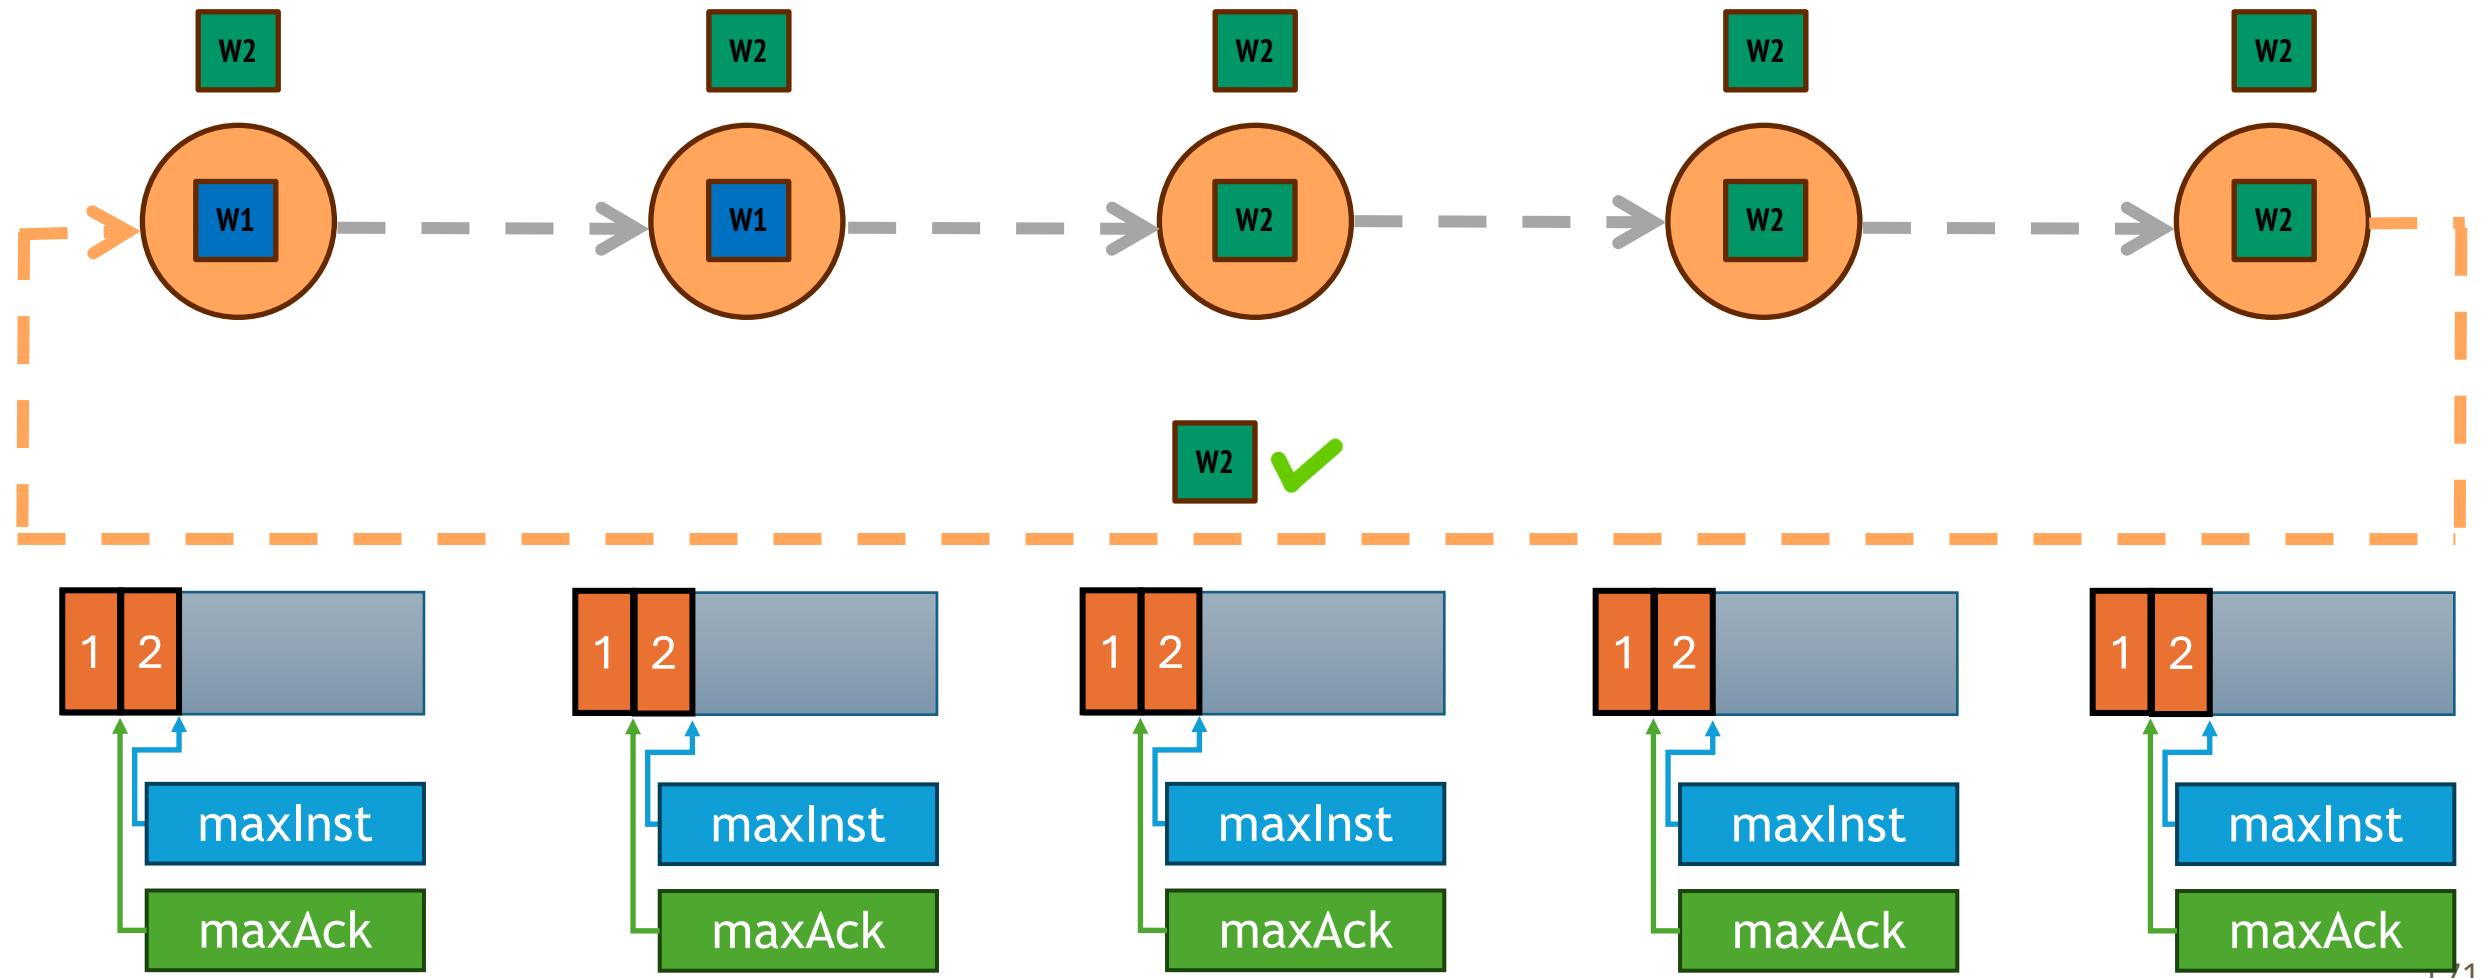
~~*Why not immediately reply?*~~

~~*Why not next unseen instance?*~~

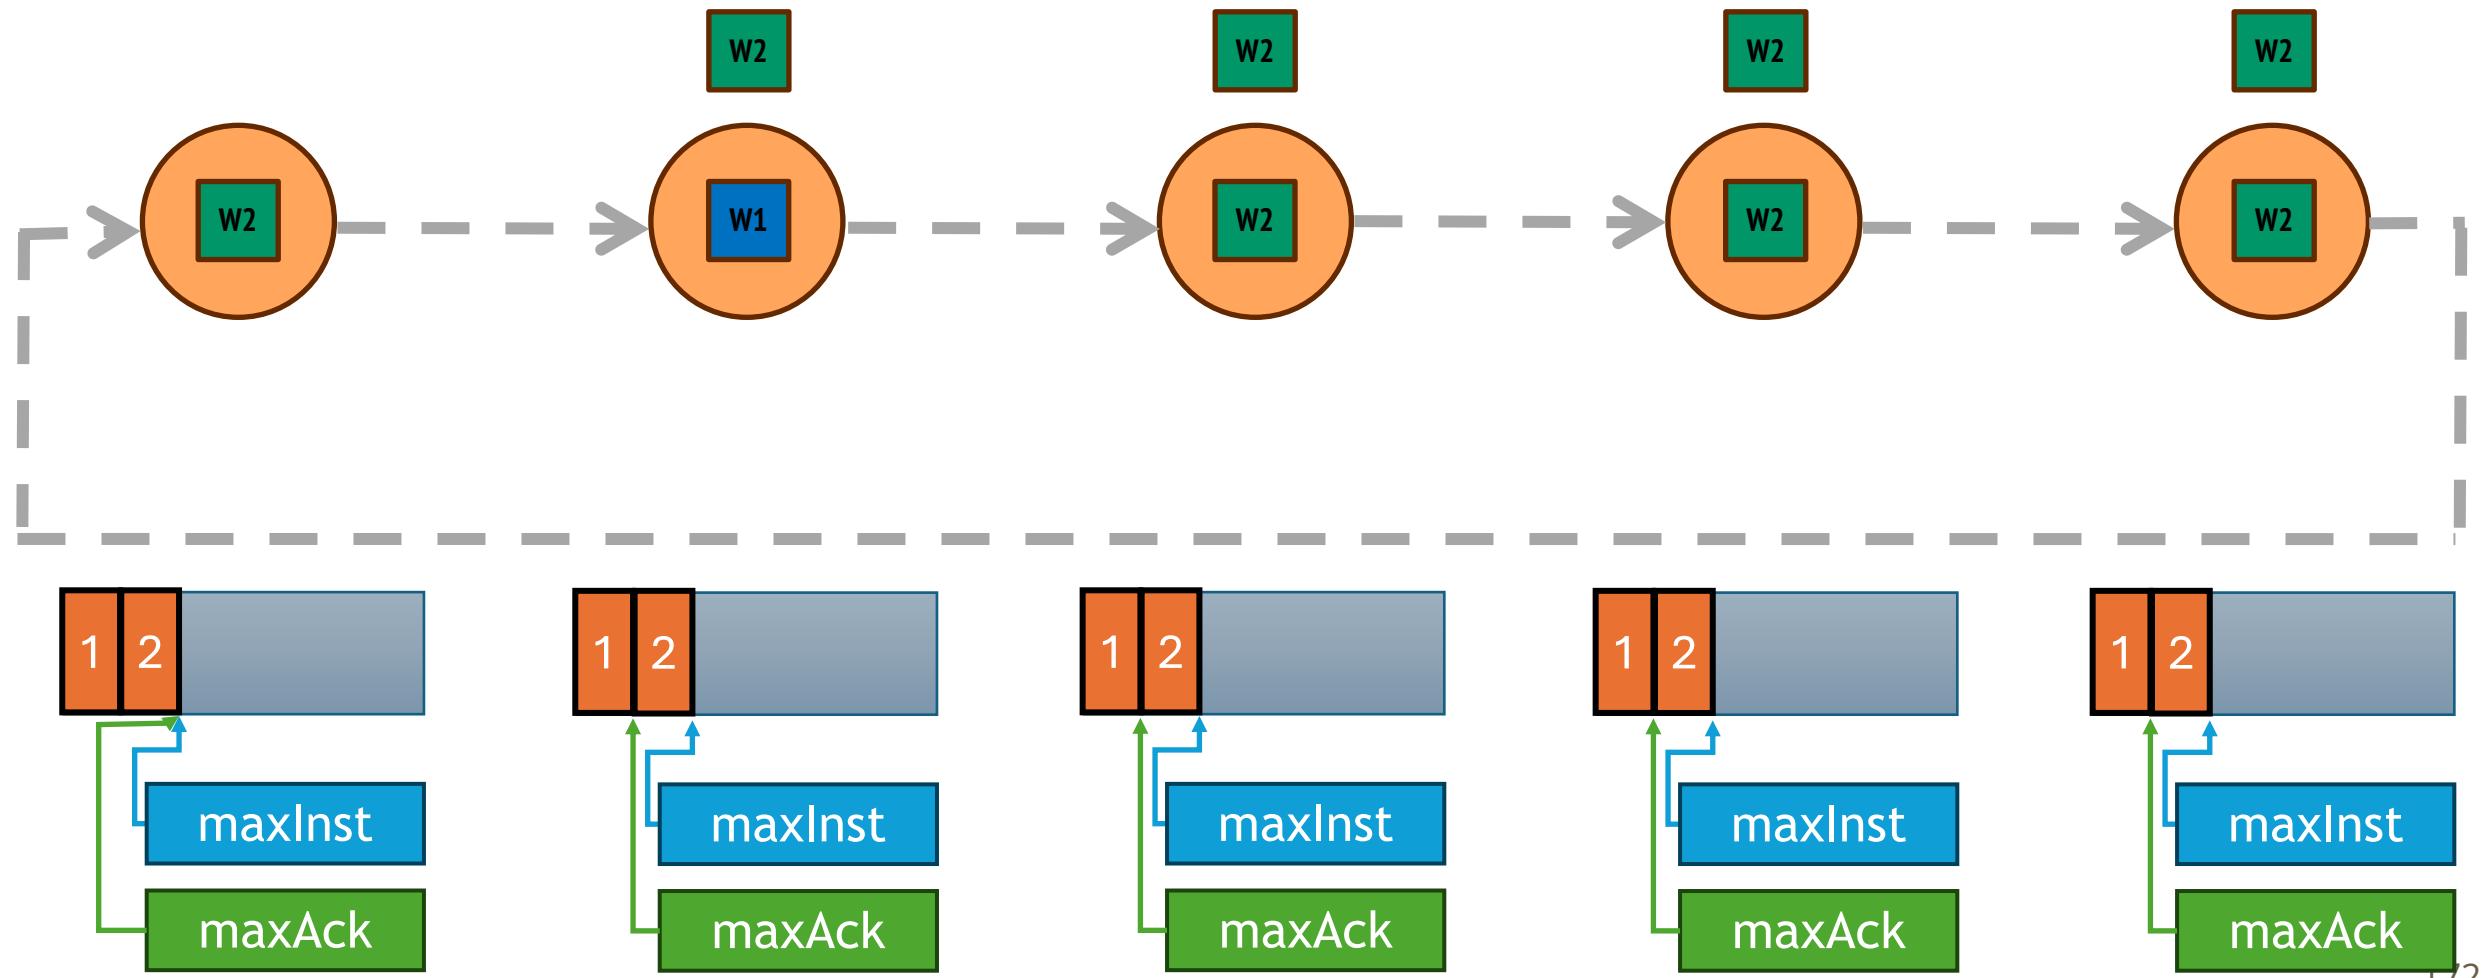
Why not ack of max seen instance?

Slightly more complicated to show, but it still leads to non-linearizable history

In Depth: Why not ack of max seen instance?

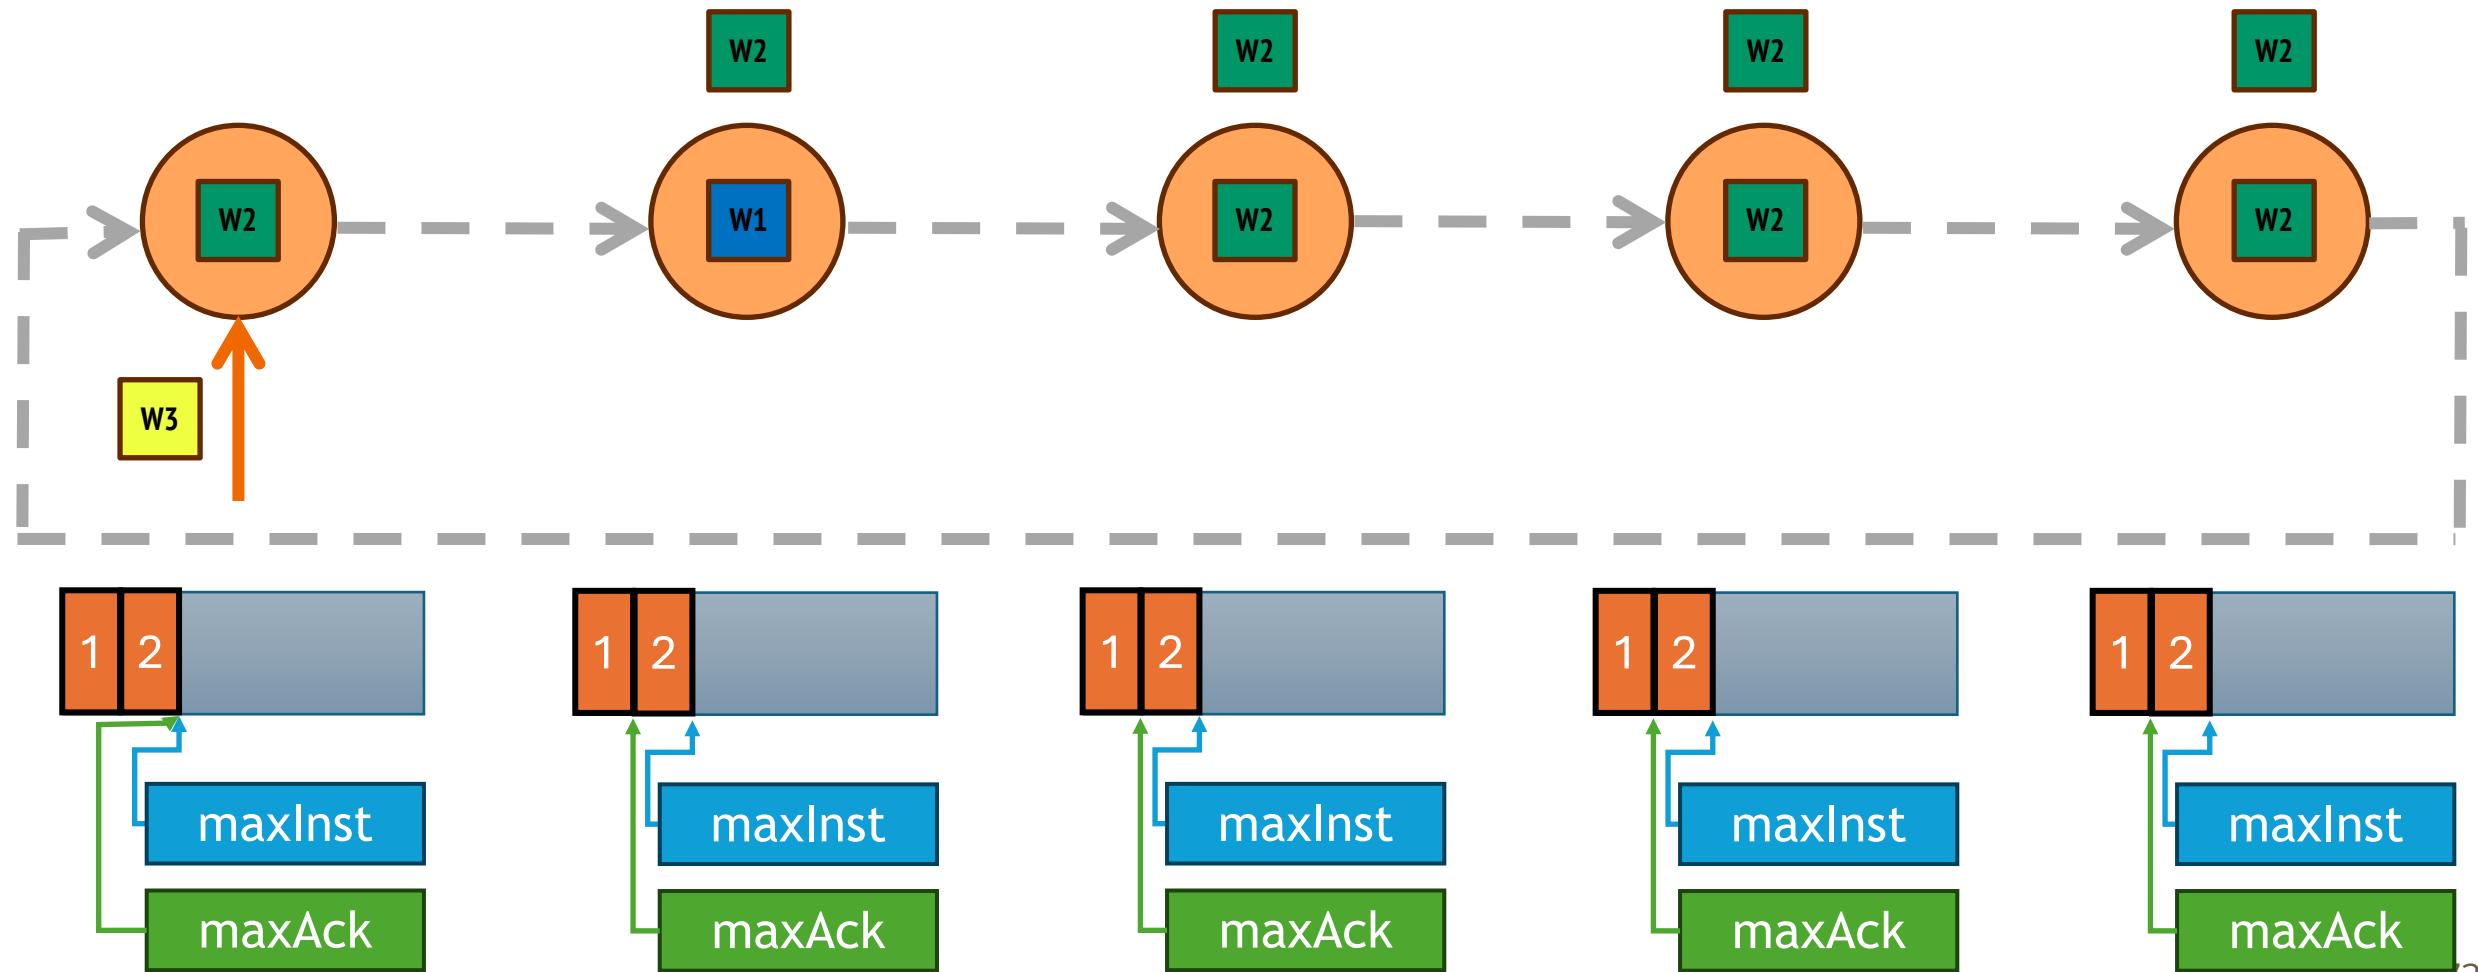


In Depth: Why not ack of max seen instance?



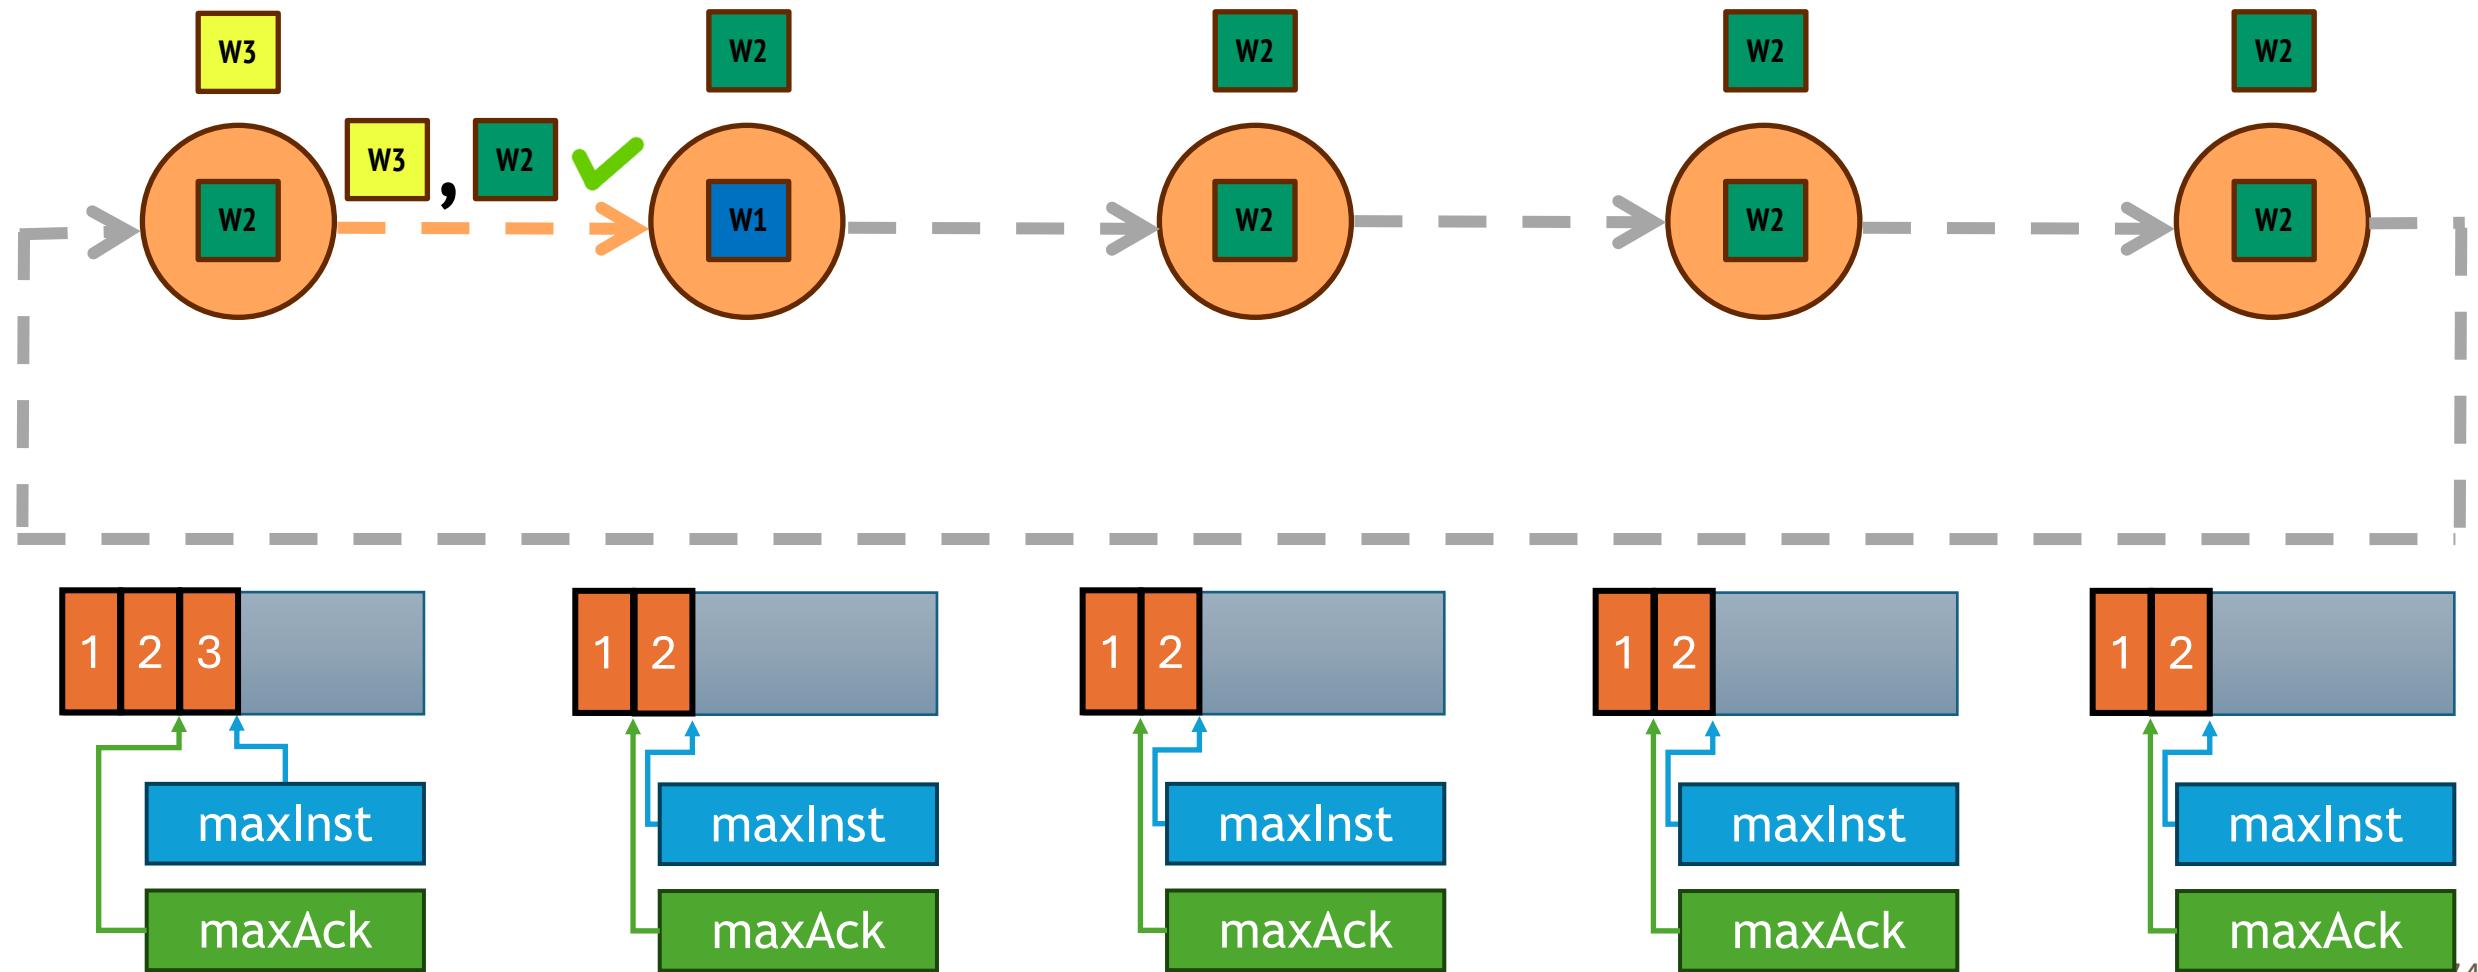
In Depth: Why not ack of max seen instance?

Ack is "piggybacked" along with next write

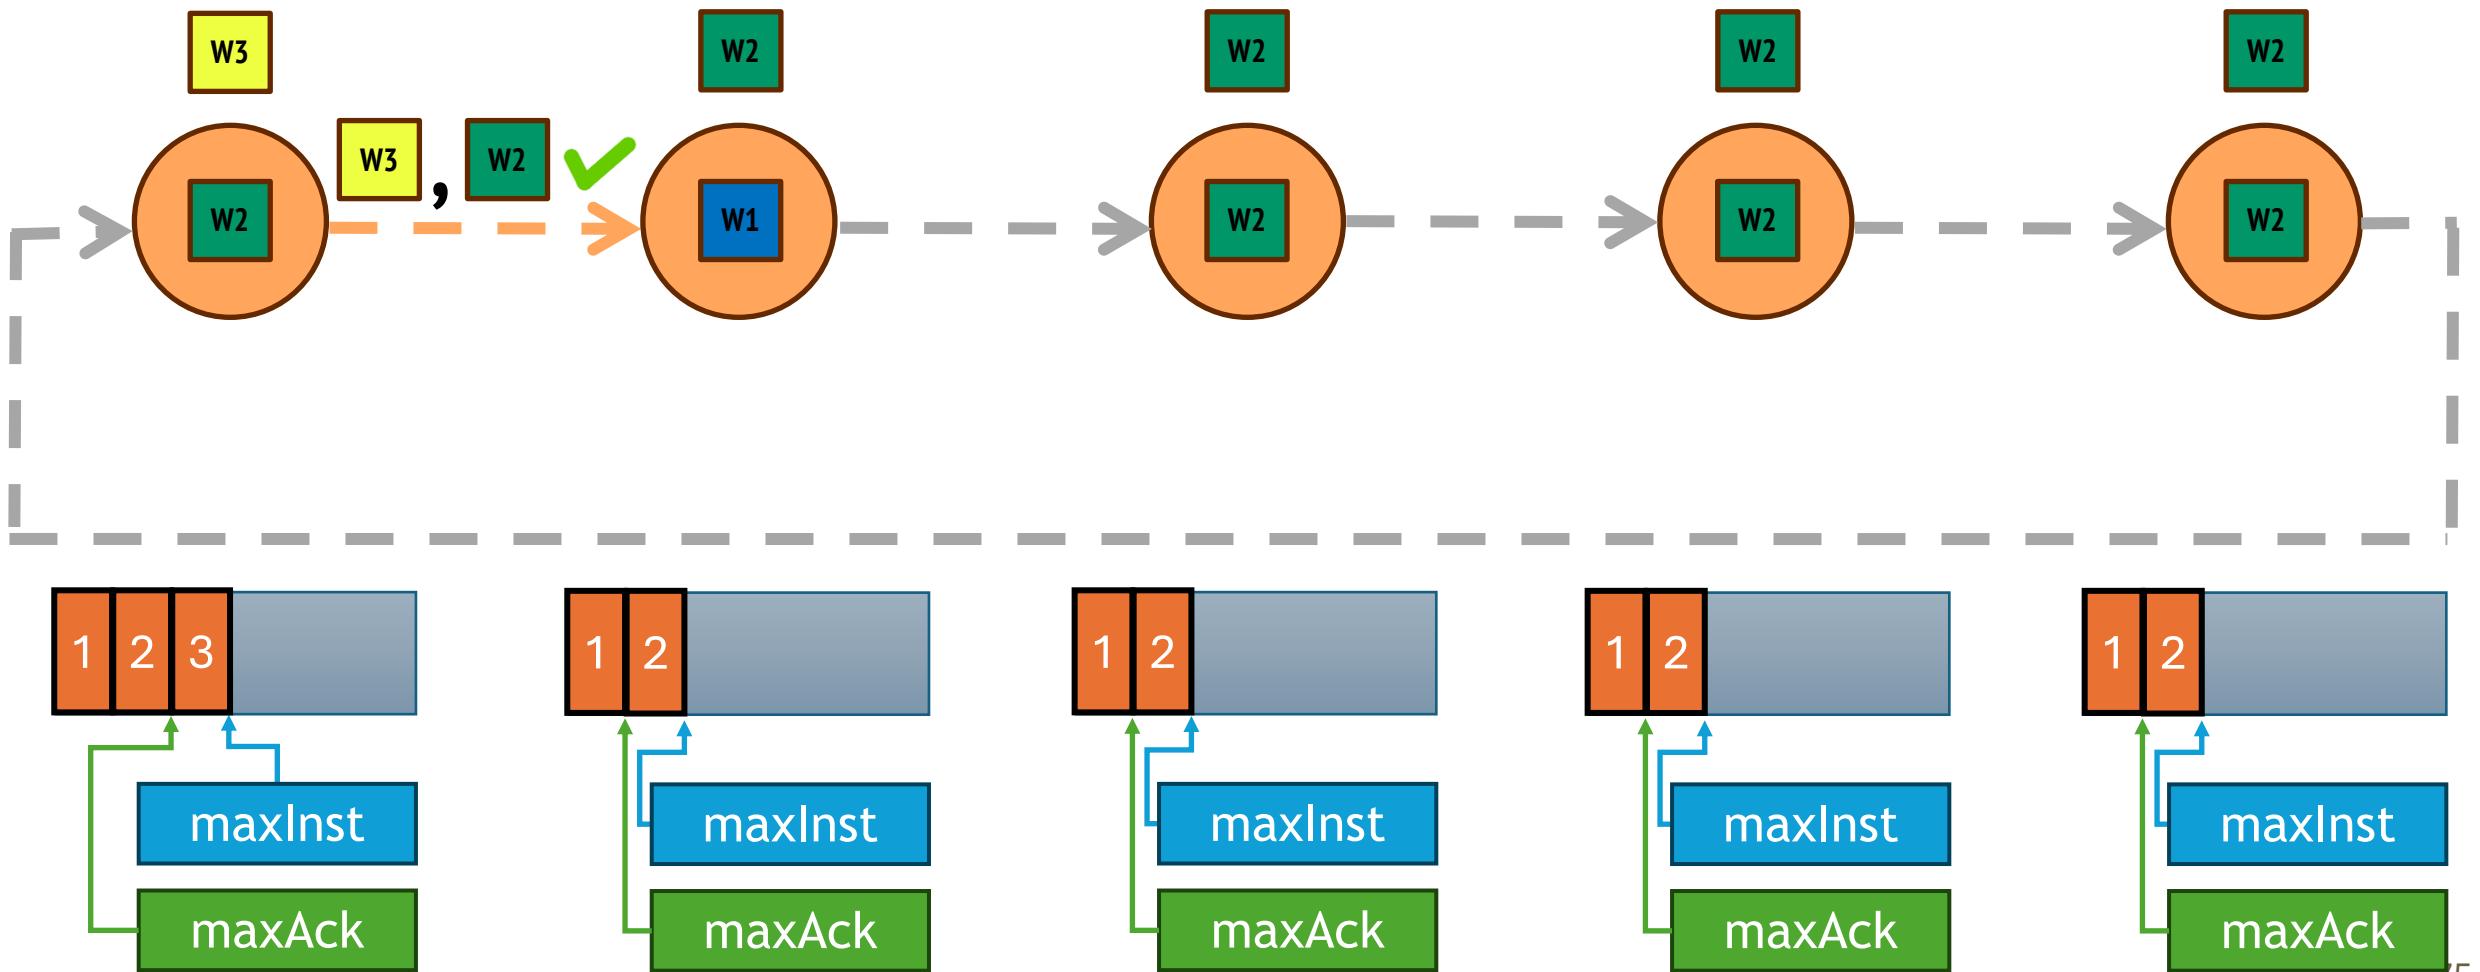


In Depth: Why not ack of max seen instance?

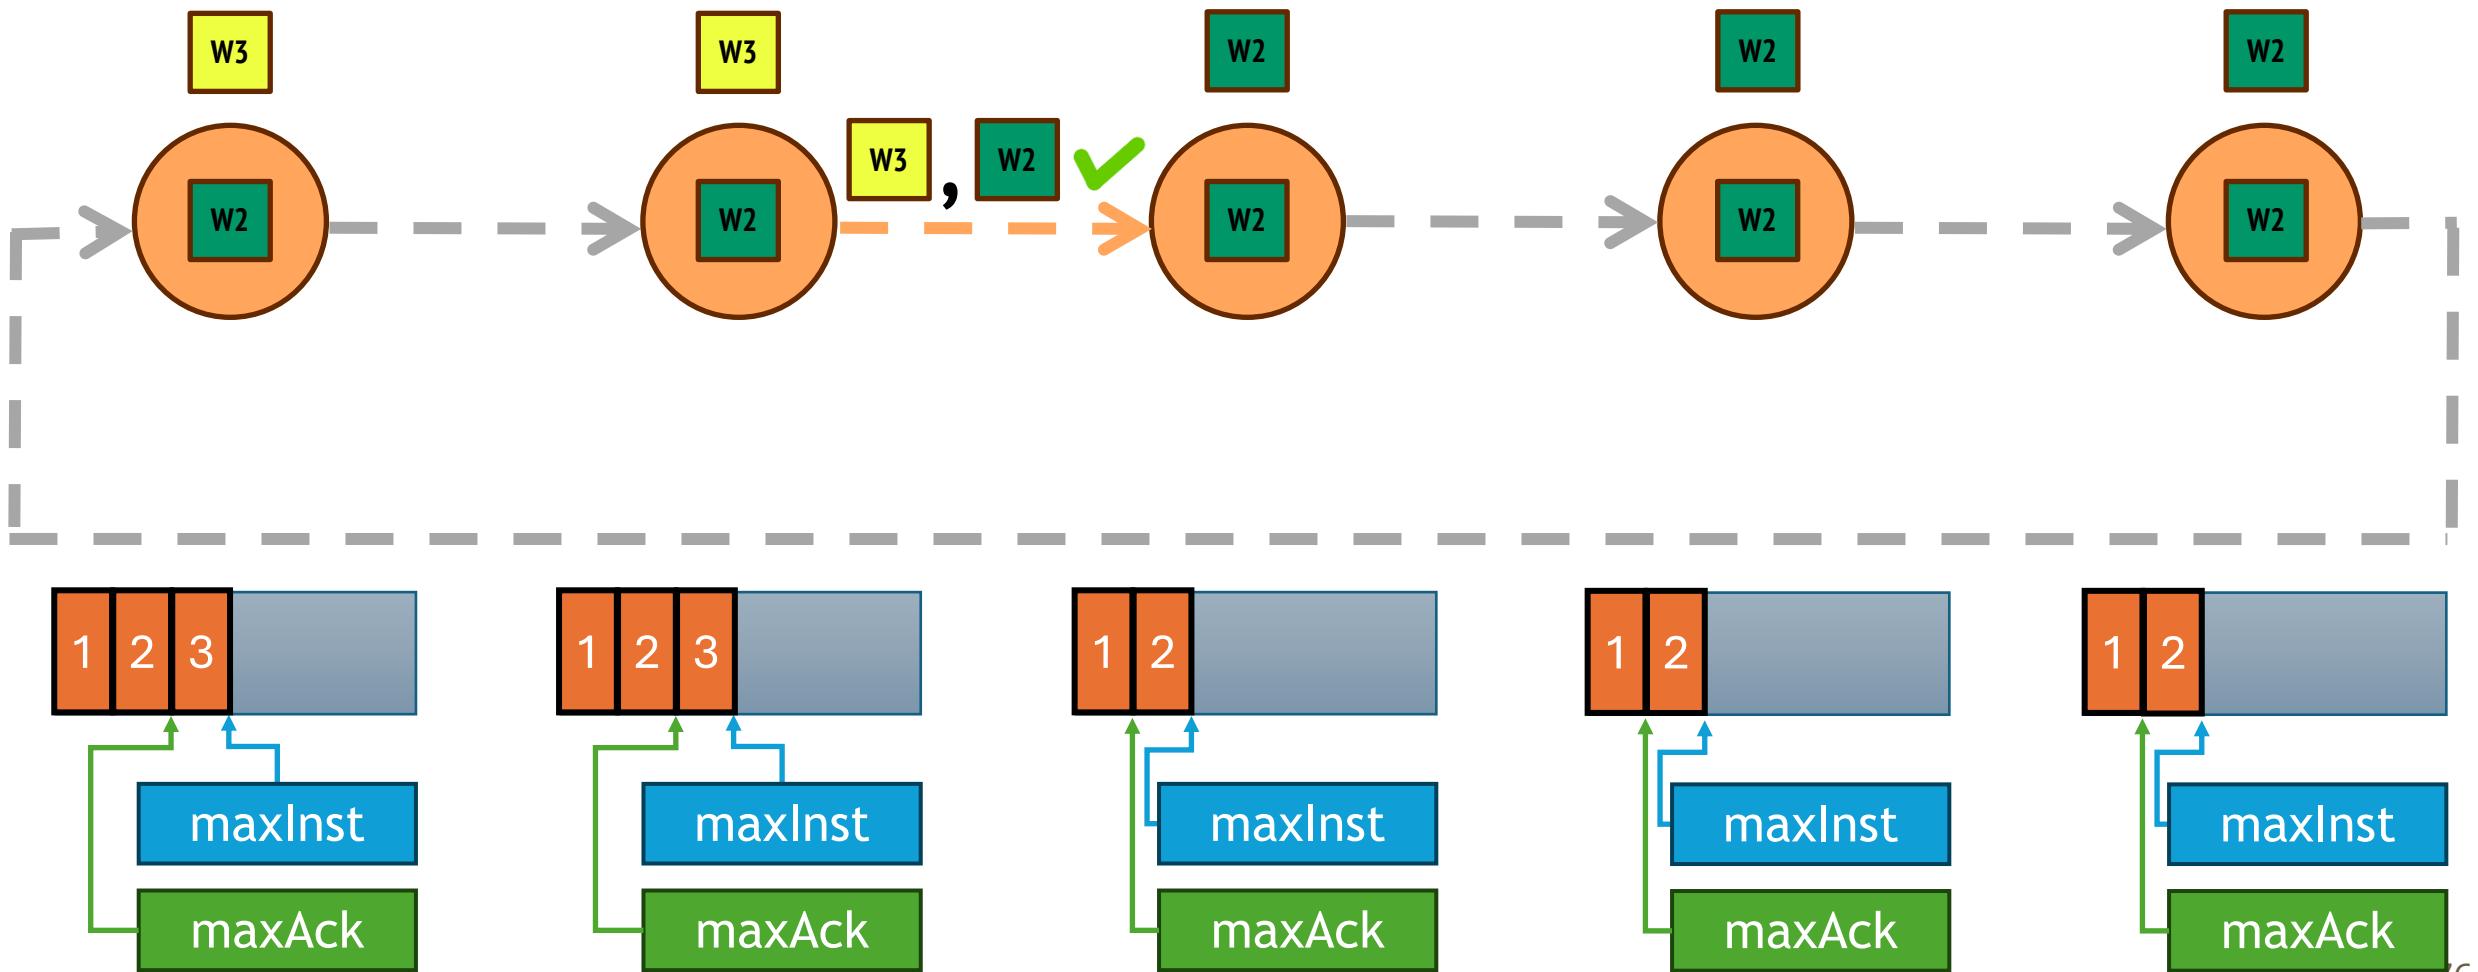
*Ack is "piggybacked" along
with next write*



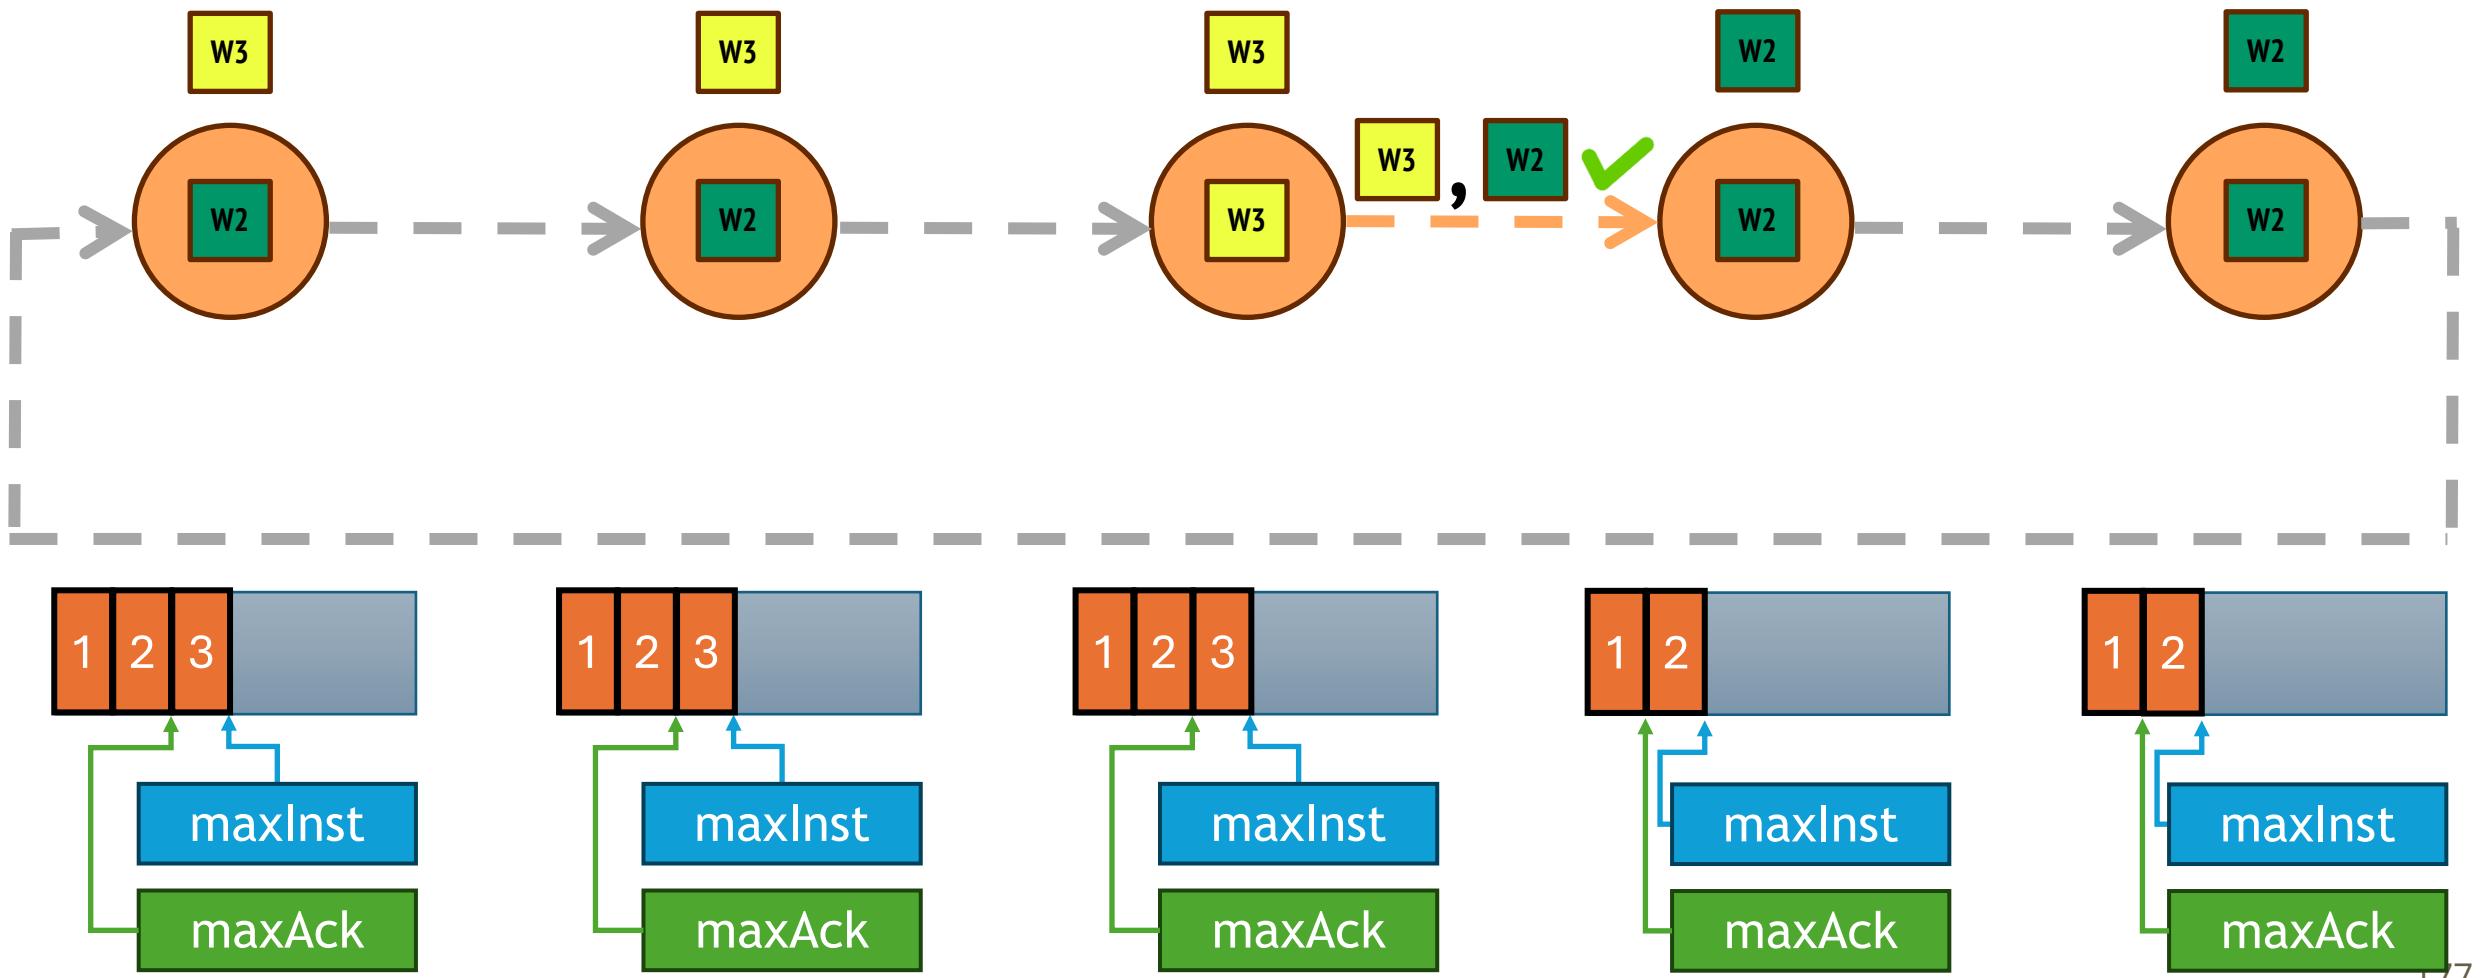
In Depth: Why not ack of max seen instance?



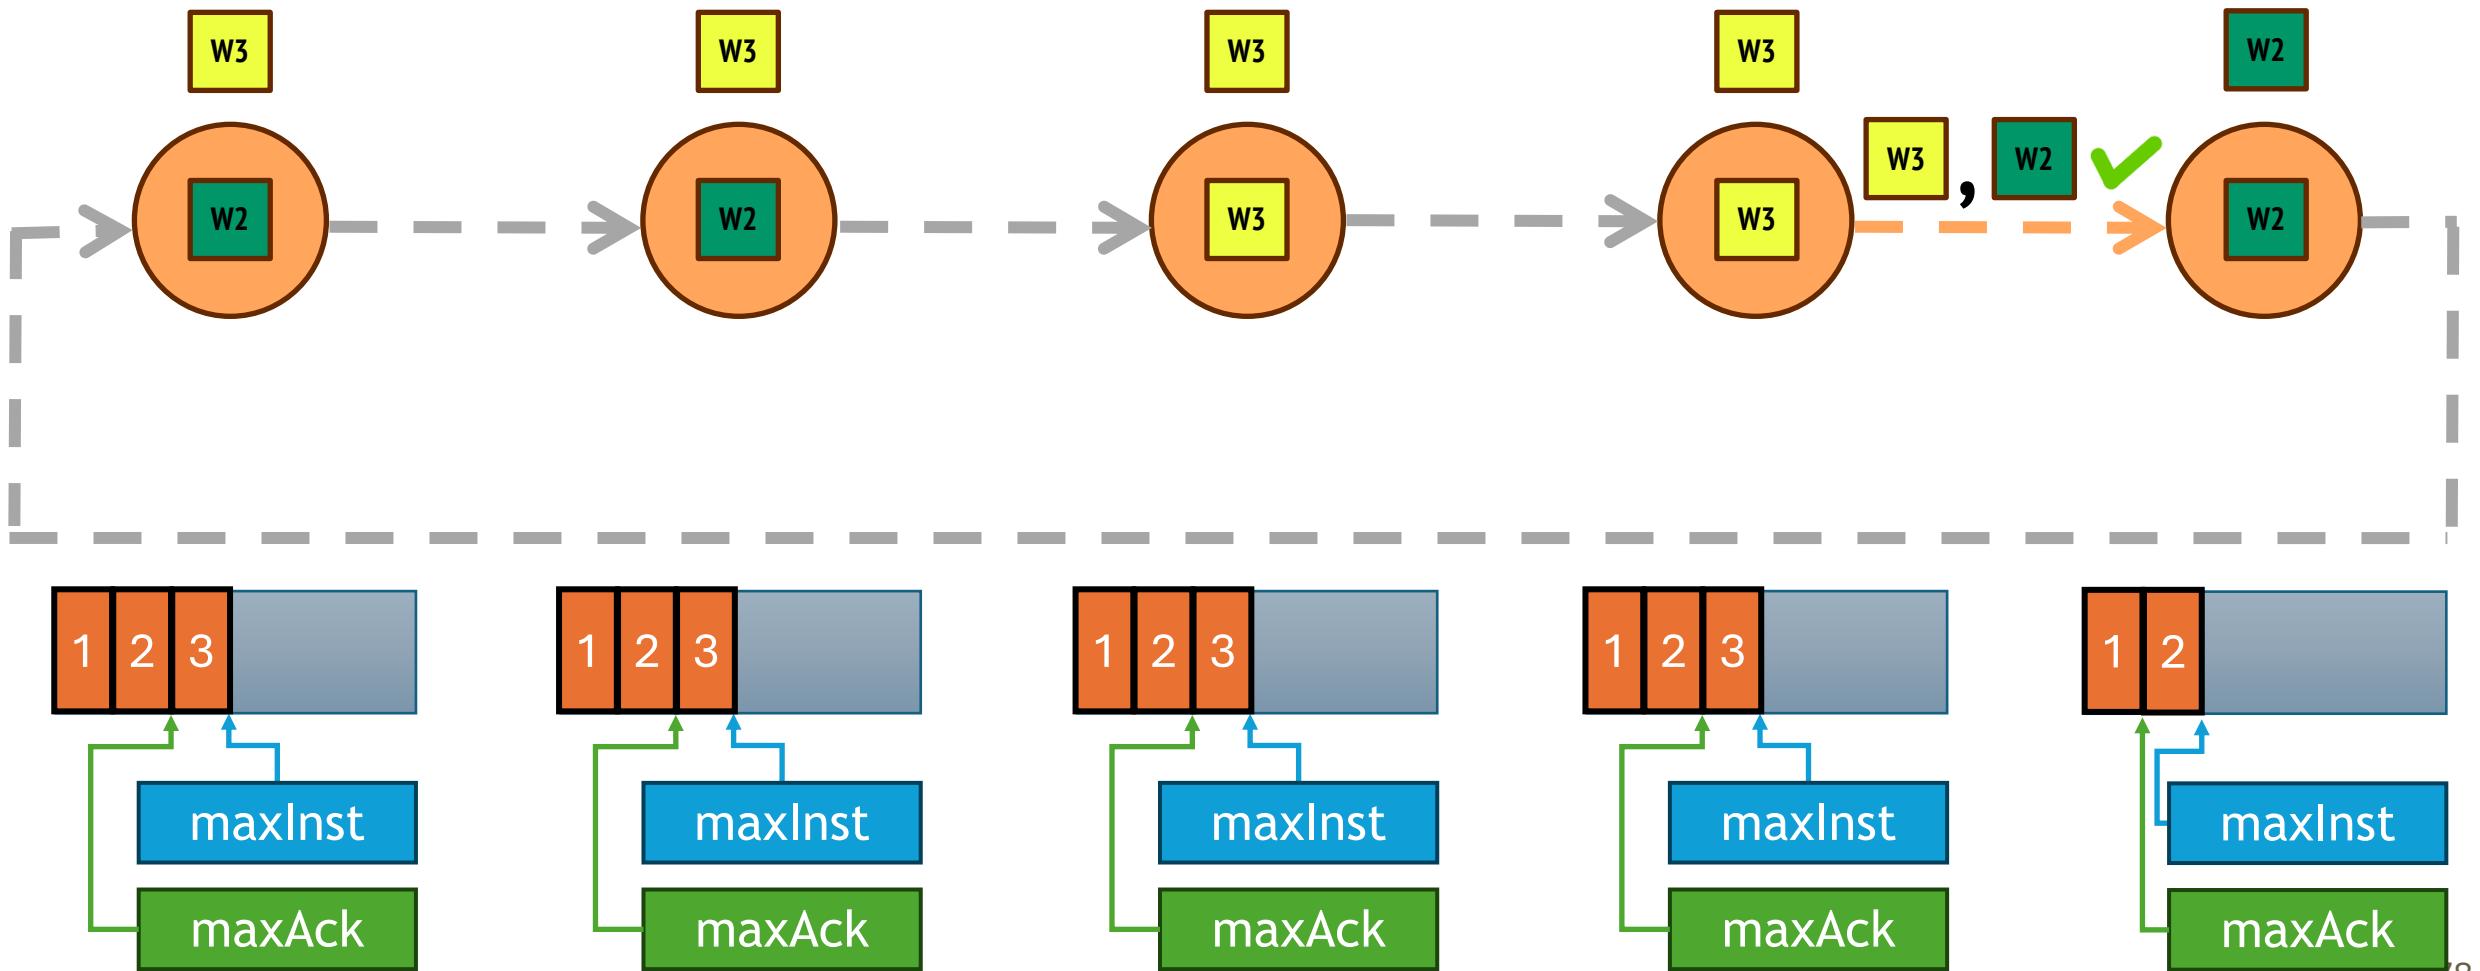
In Depth: Why not ack of max seen instance?



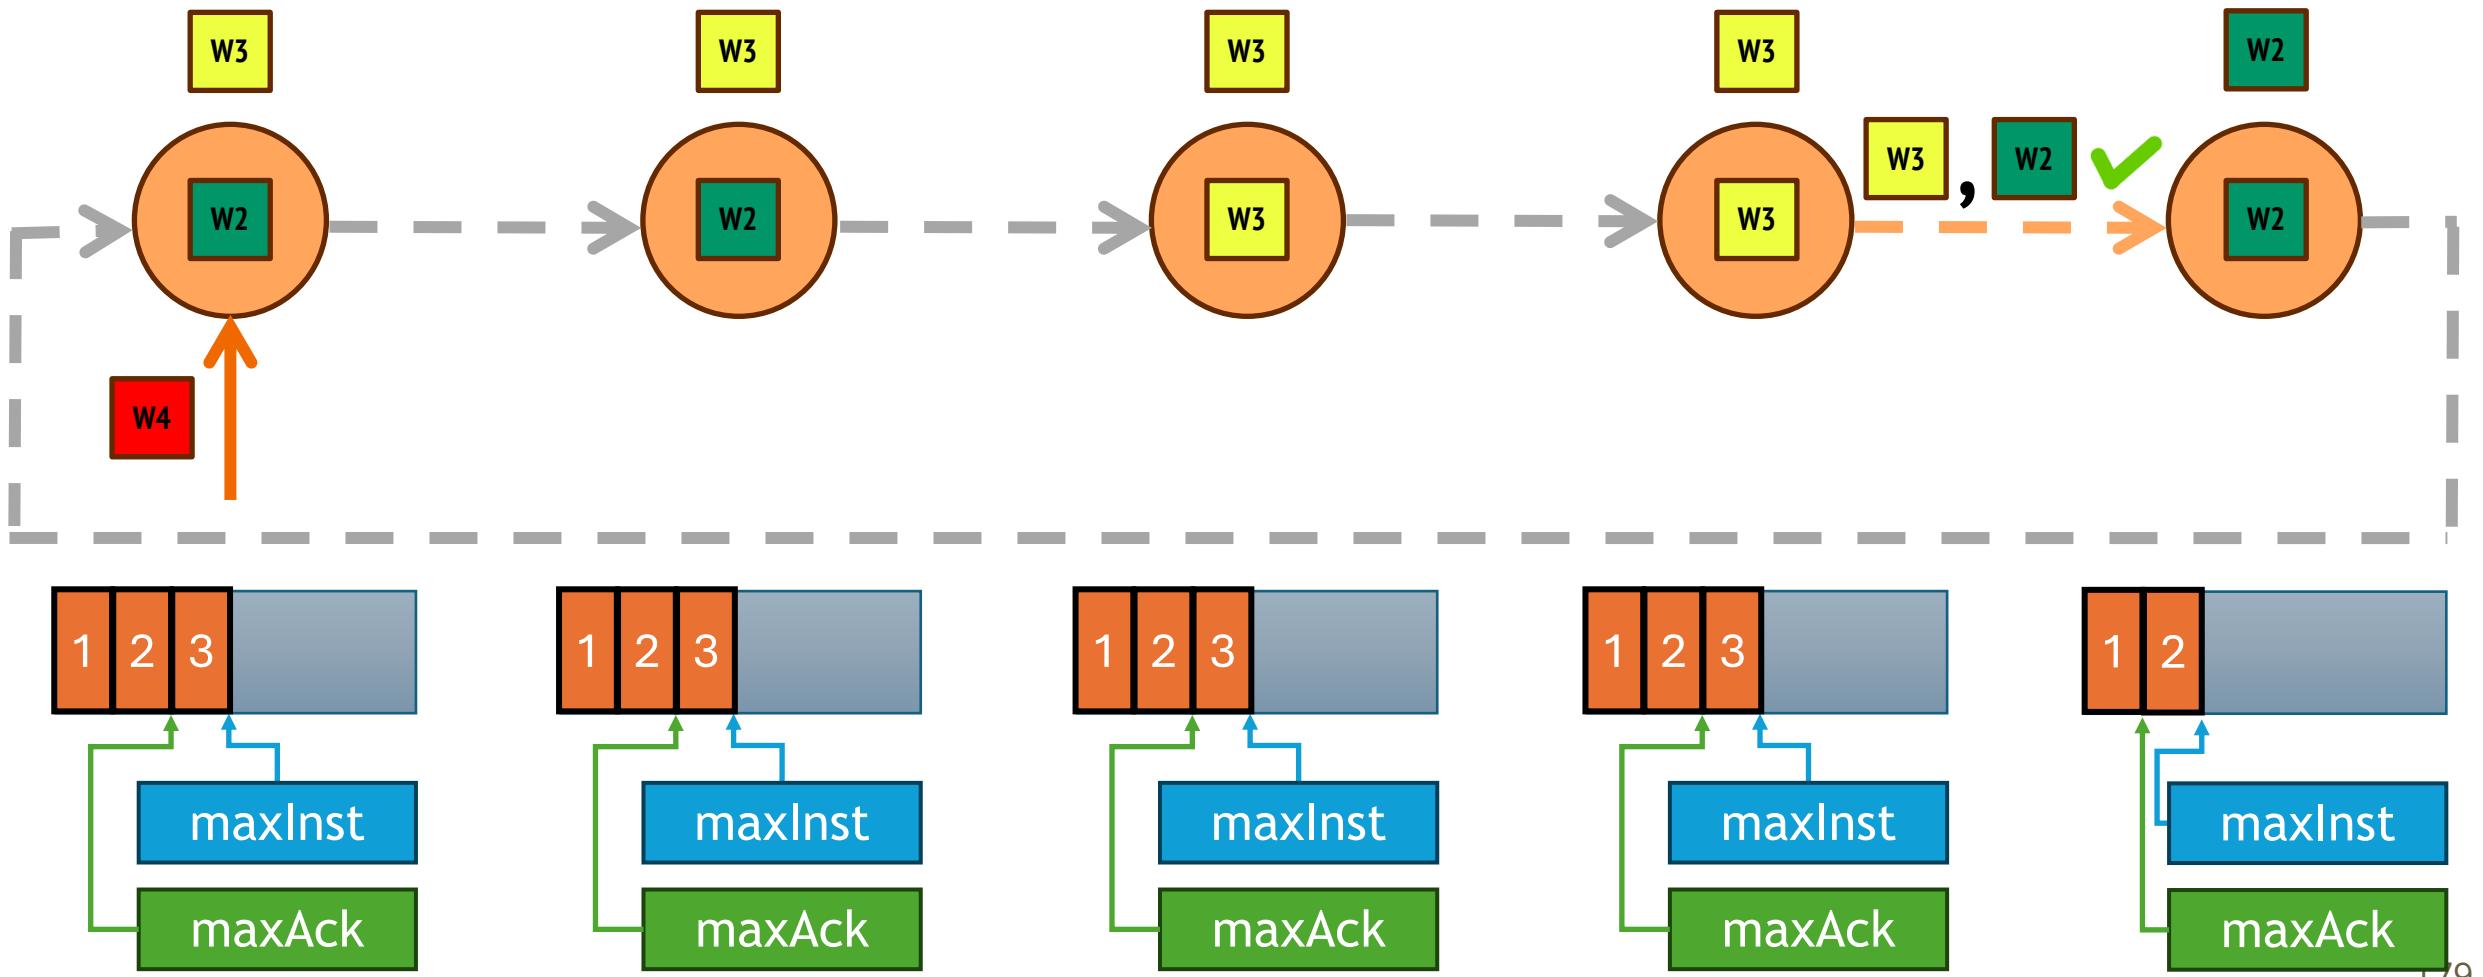
In Depth: Why not ack of max seen instance?



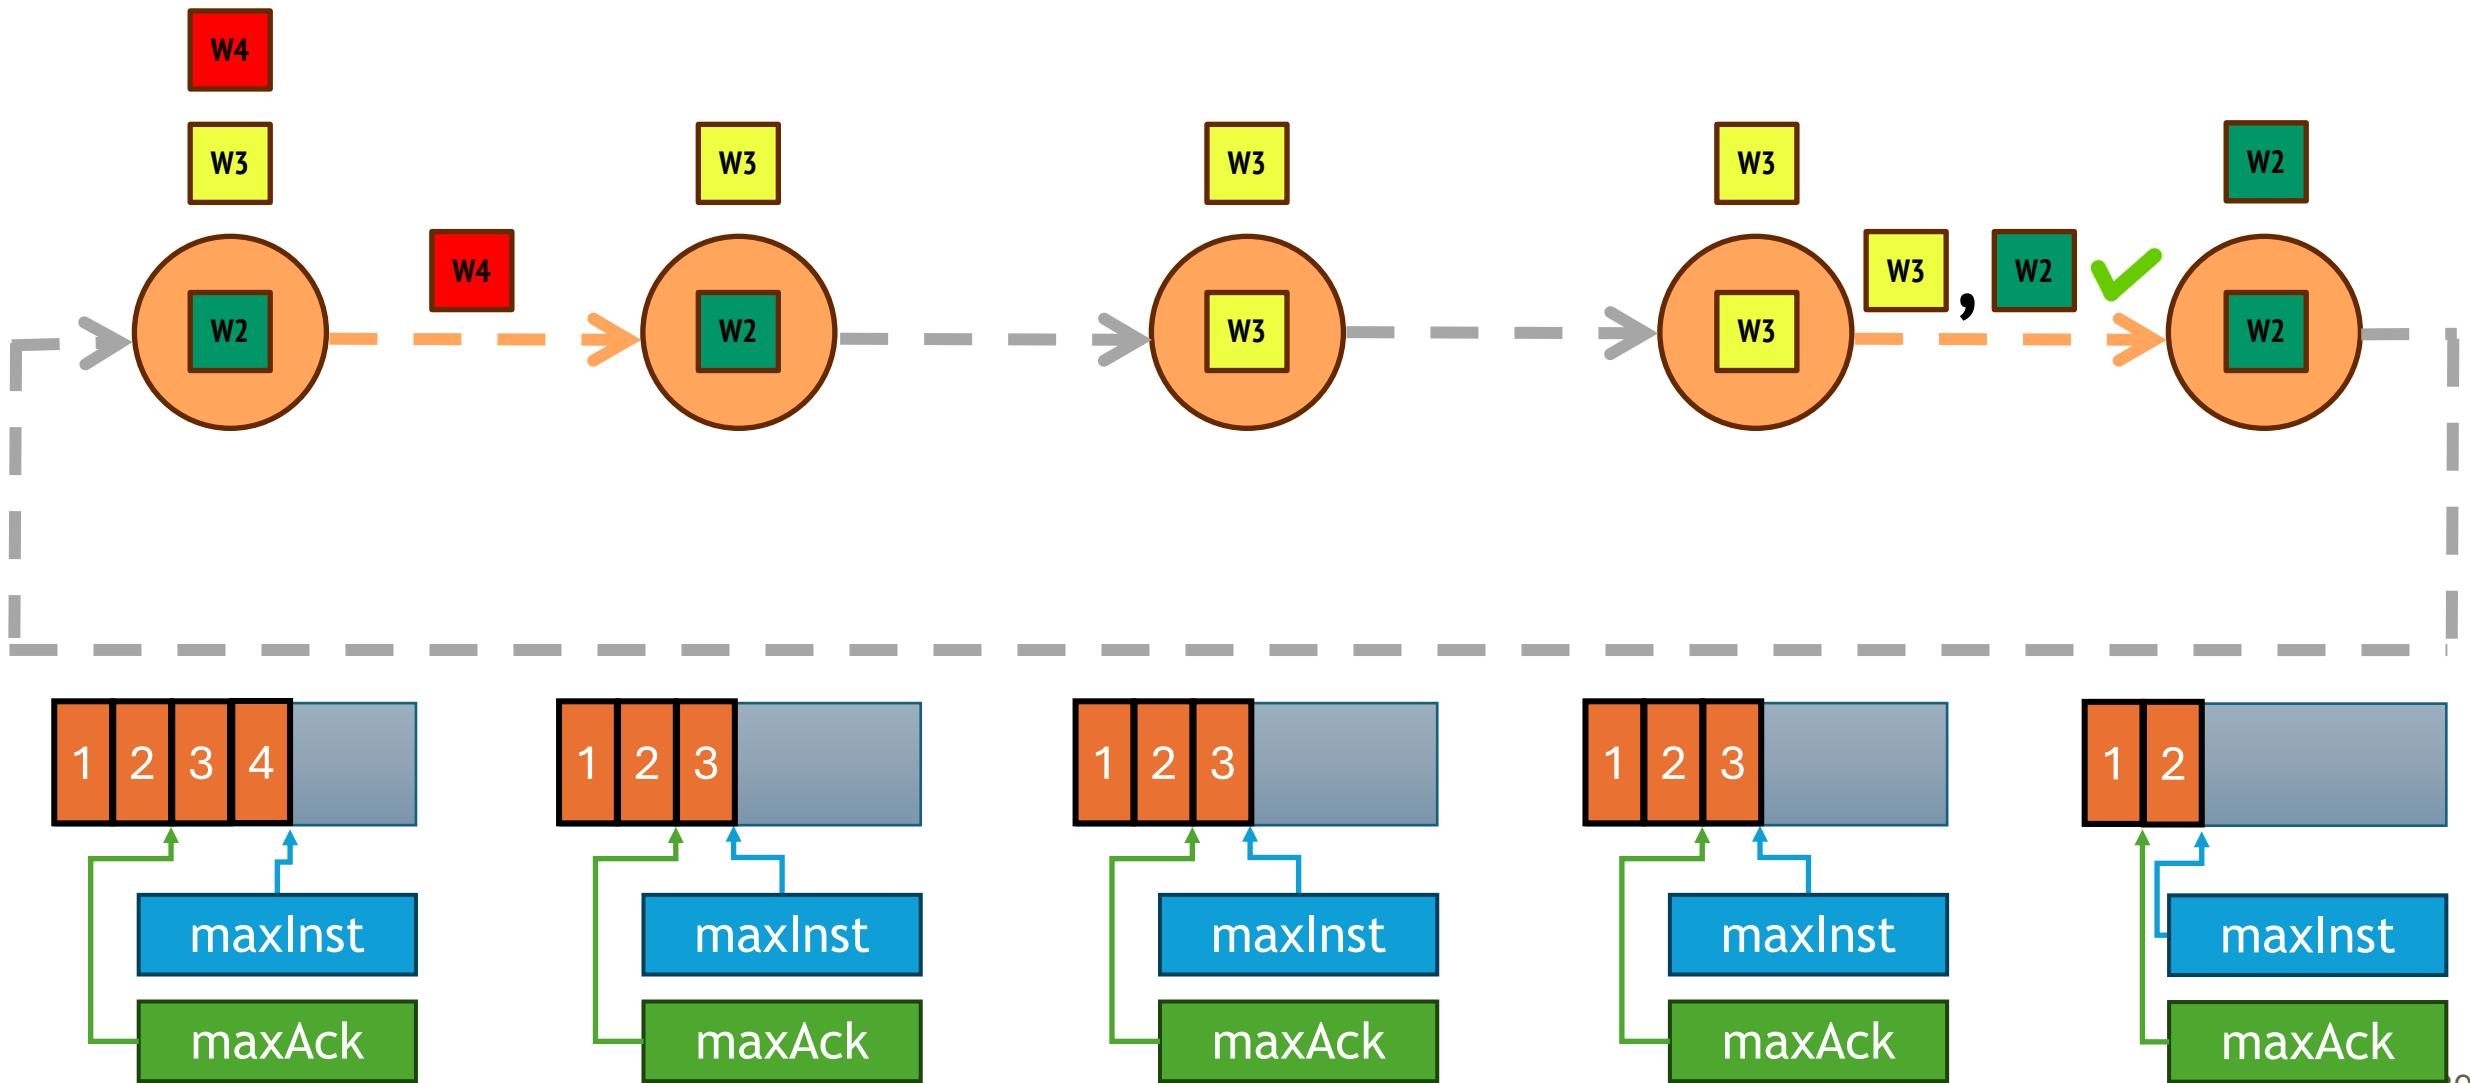
In Depth: Why not ack of max seen instance?



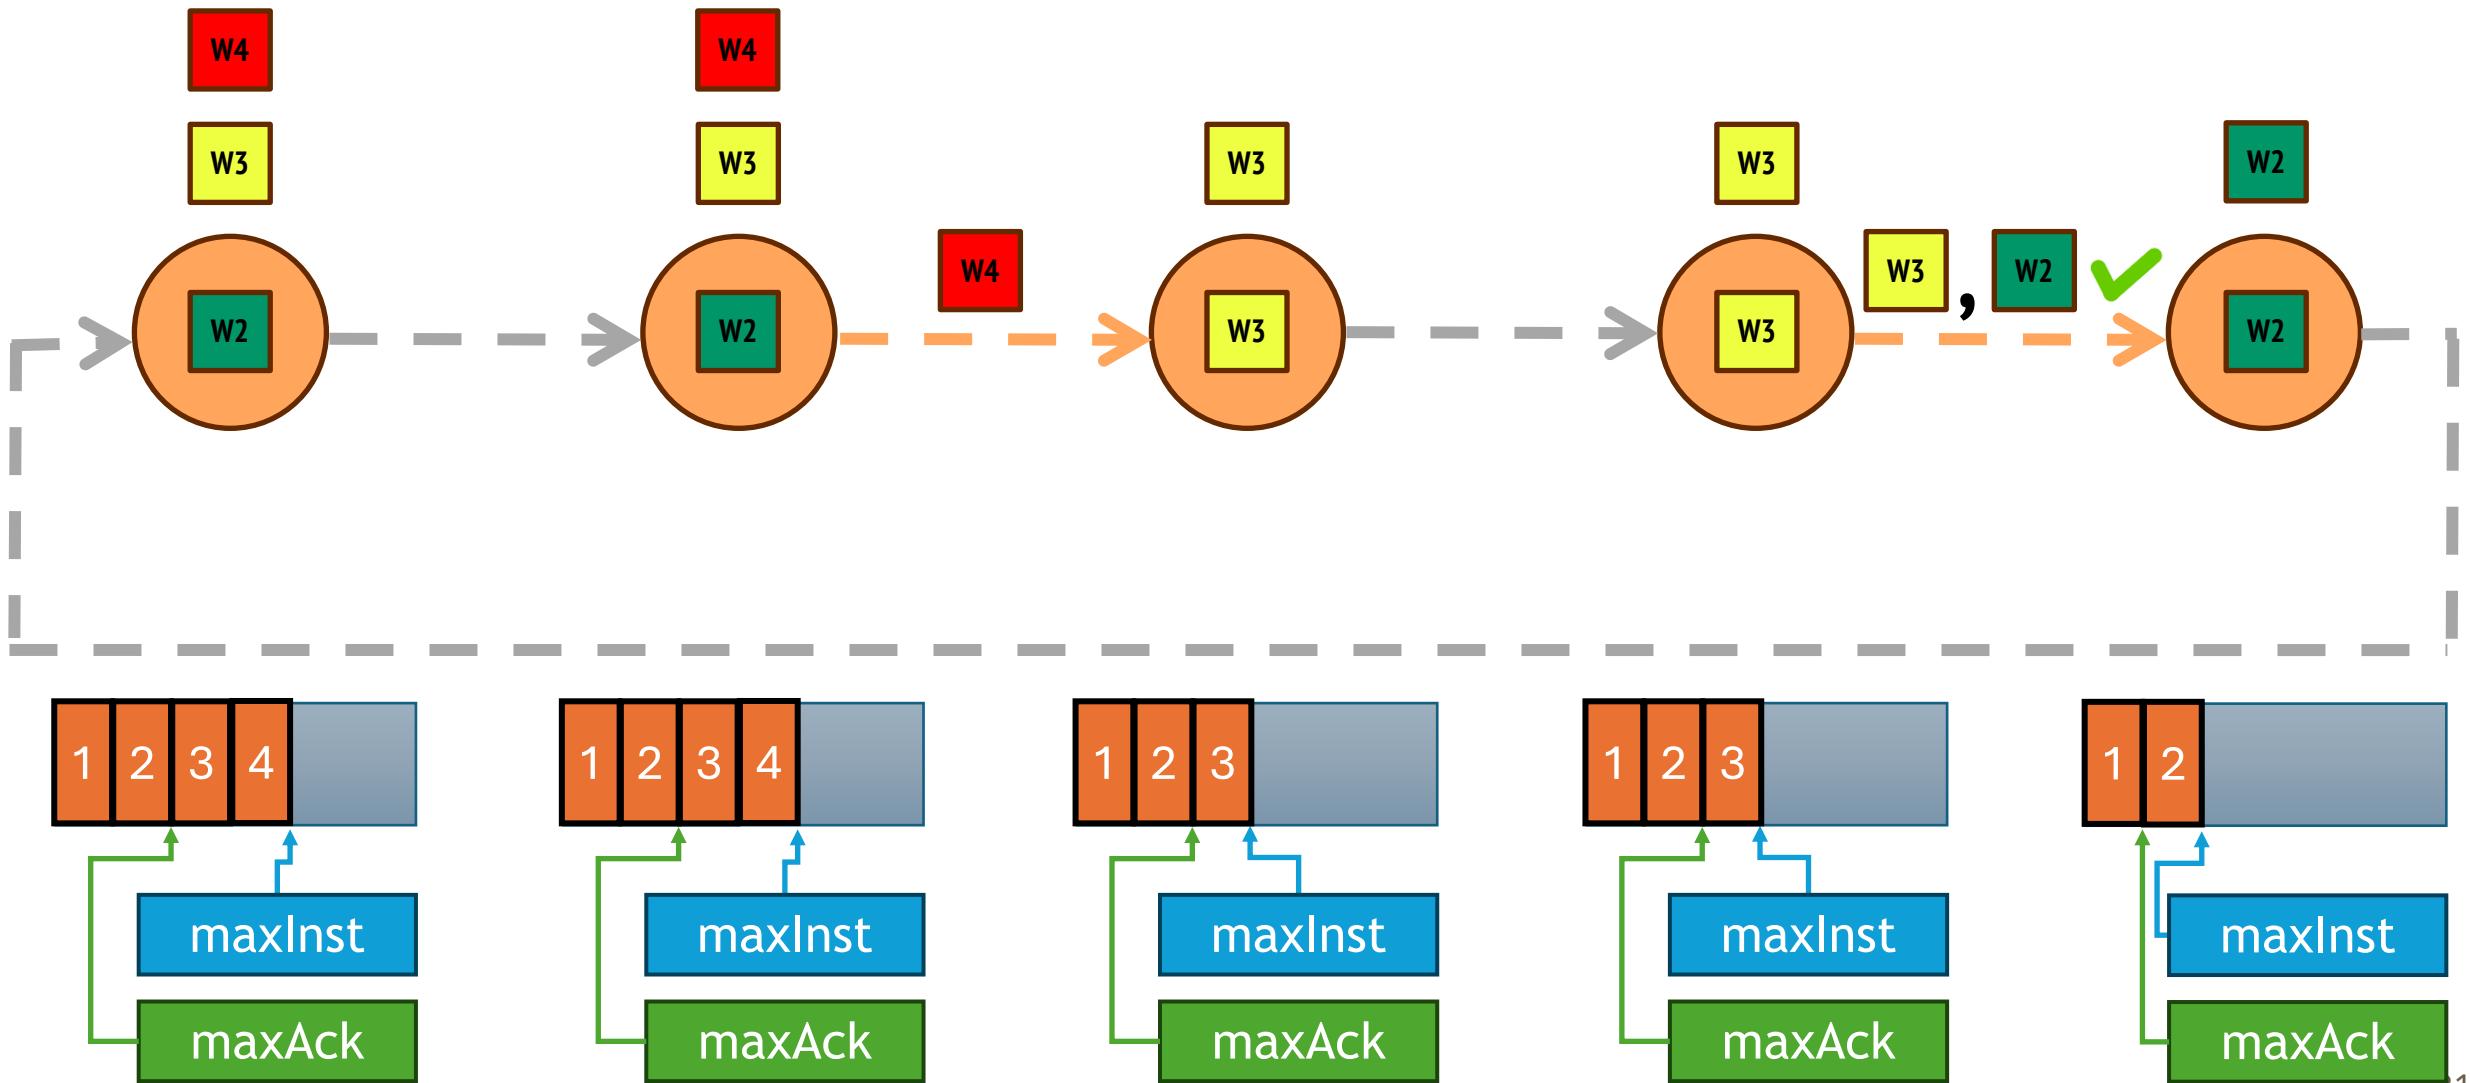
In Depth: Why not ack of max seen instance?



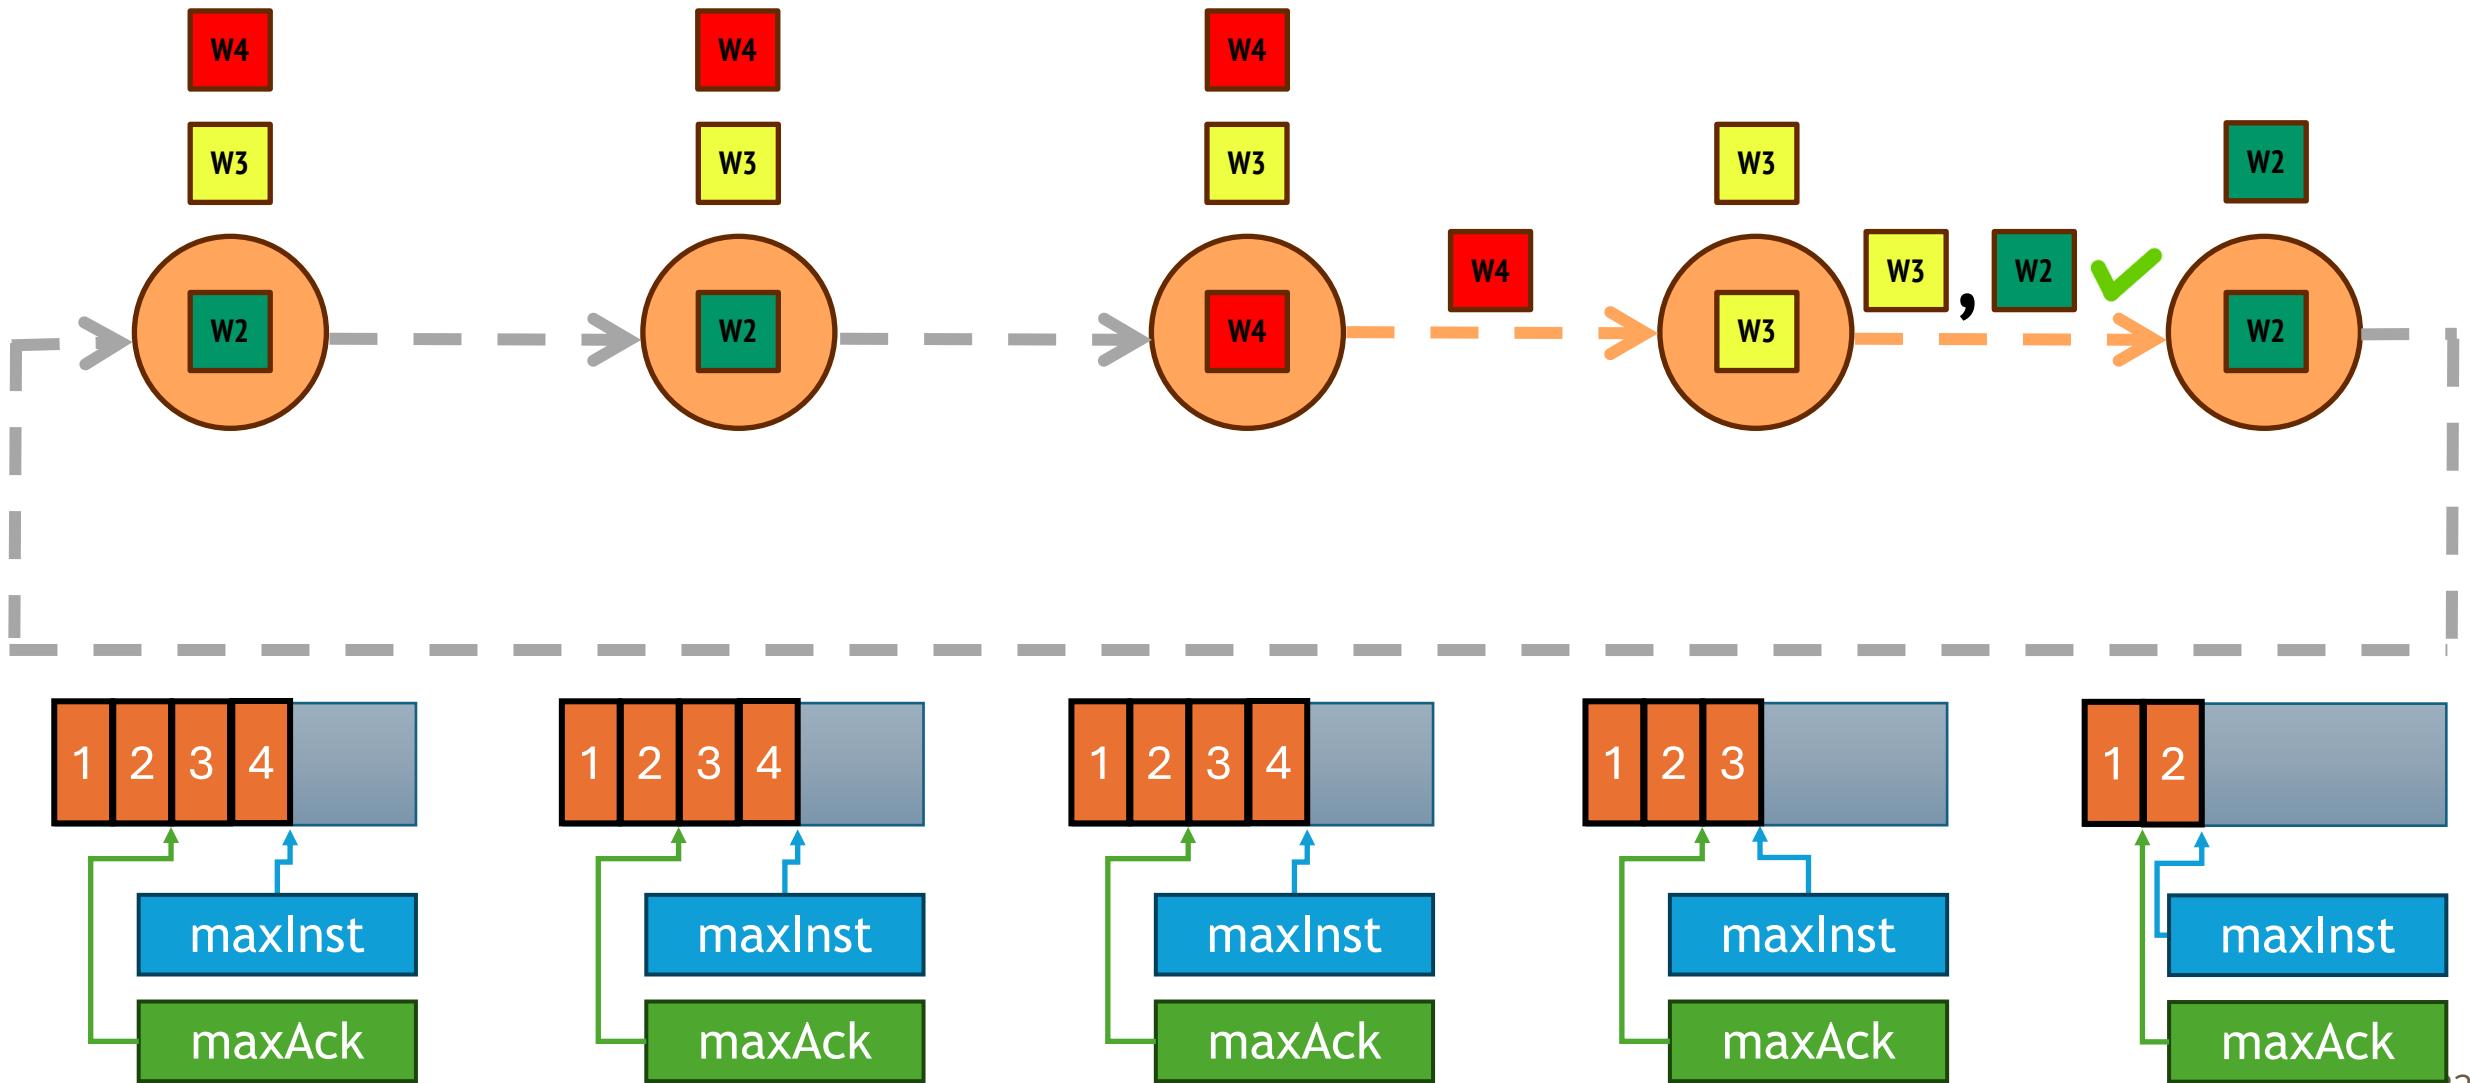
In Depth: Why not ack of max seen instance?



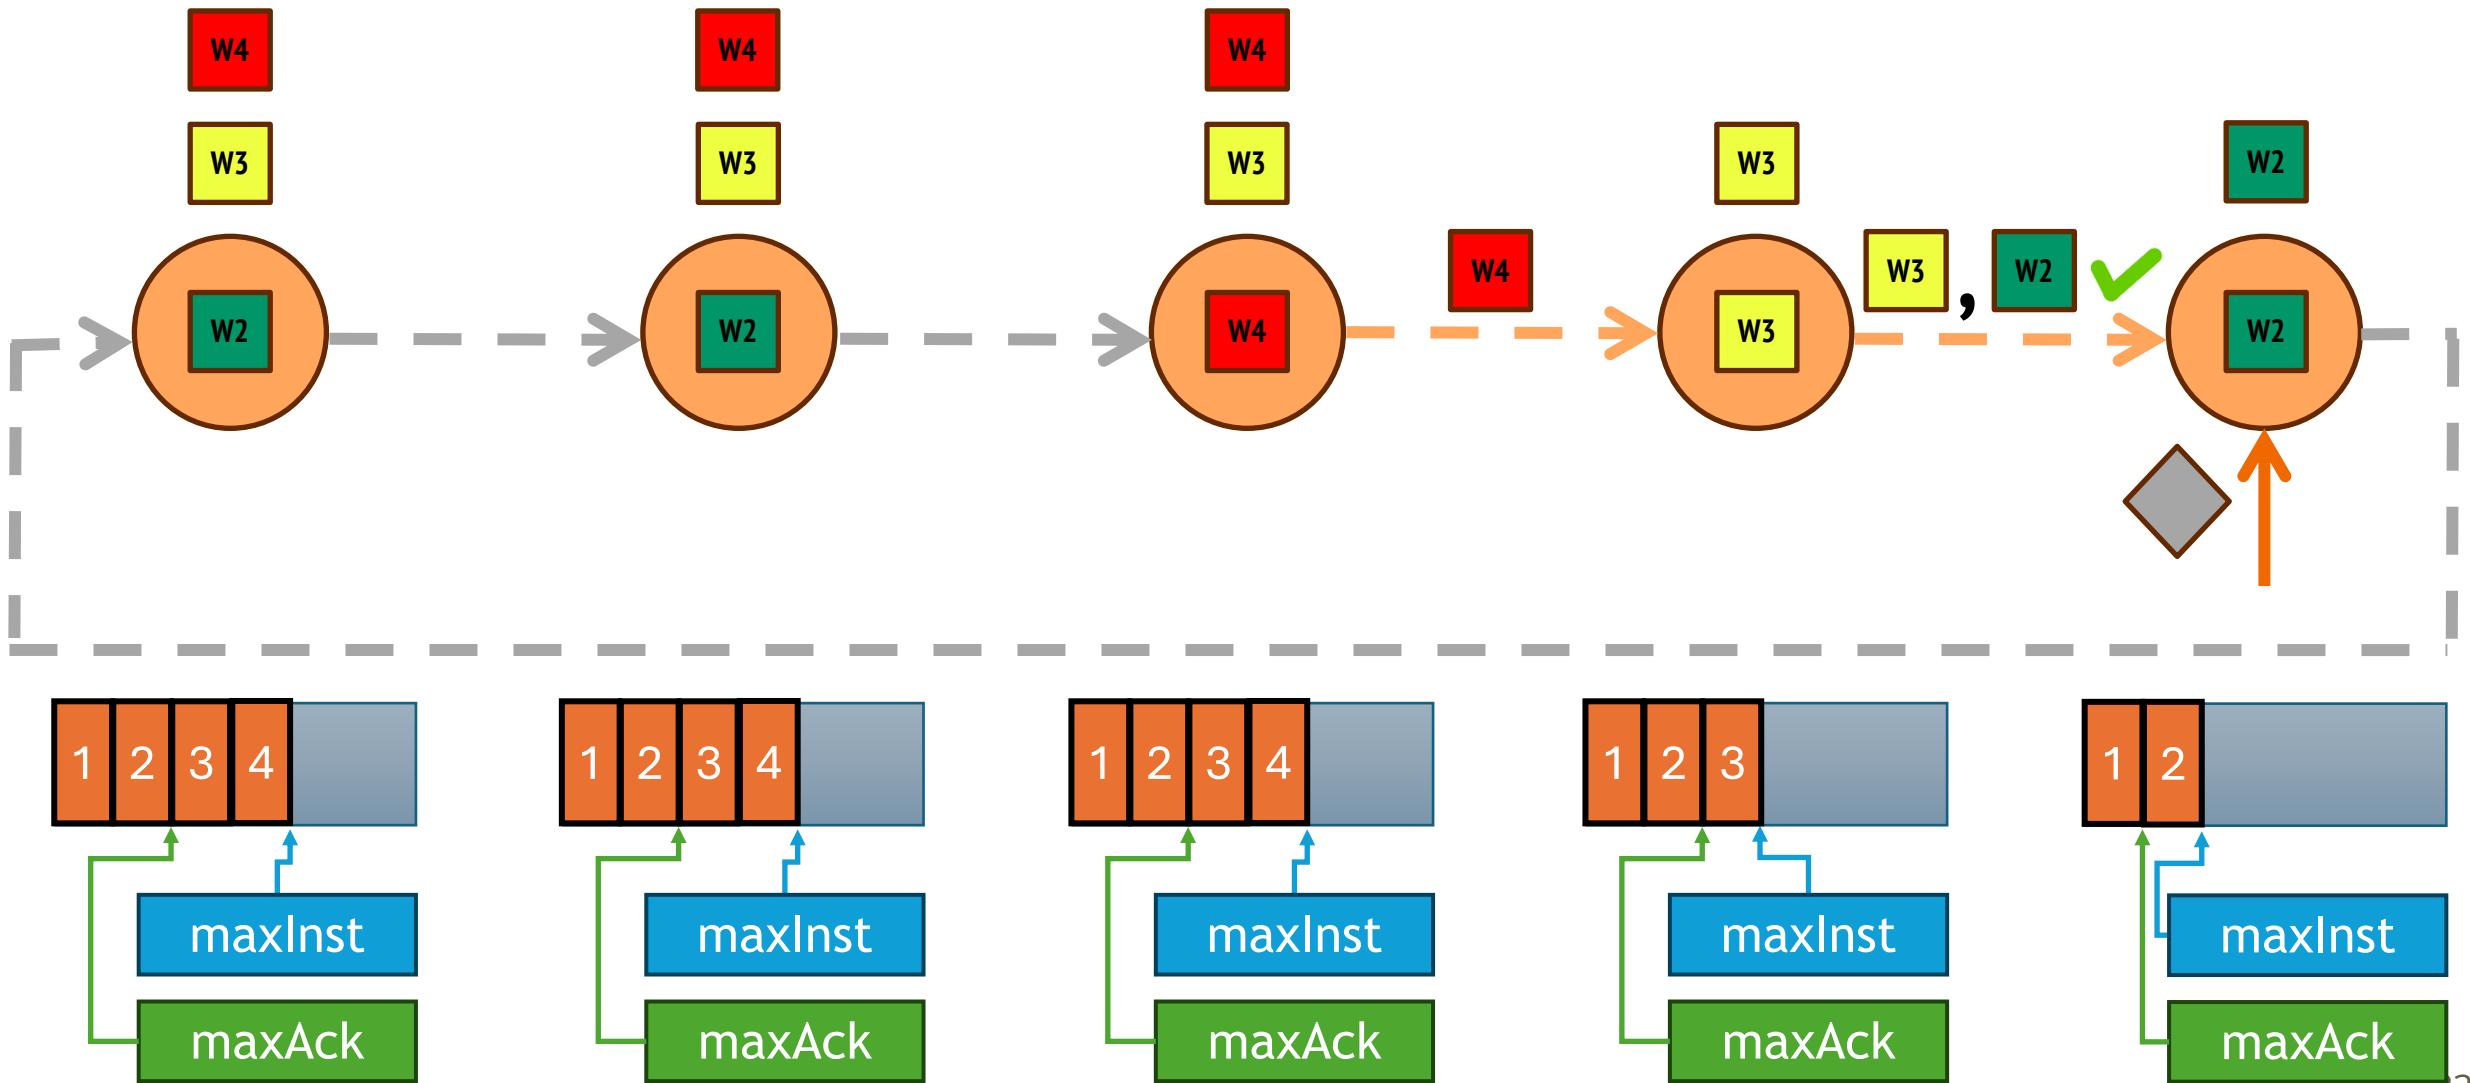
In Depth: Why not ack of max seen instance?



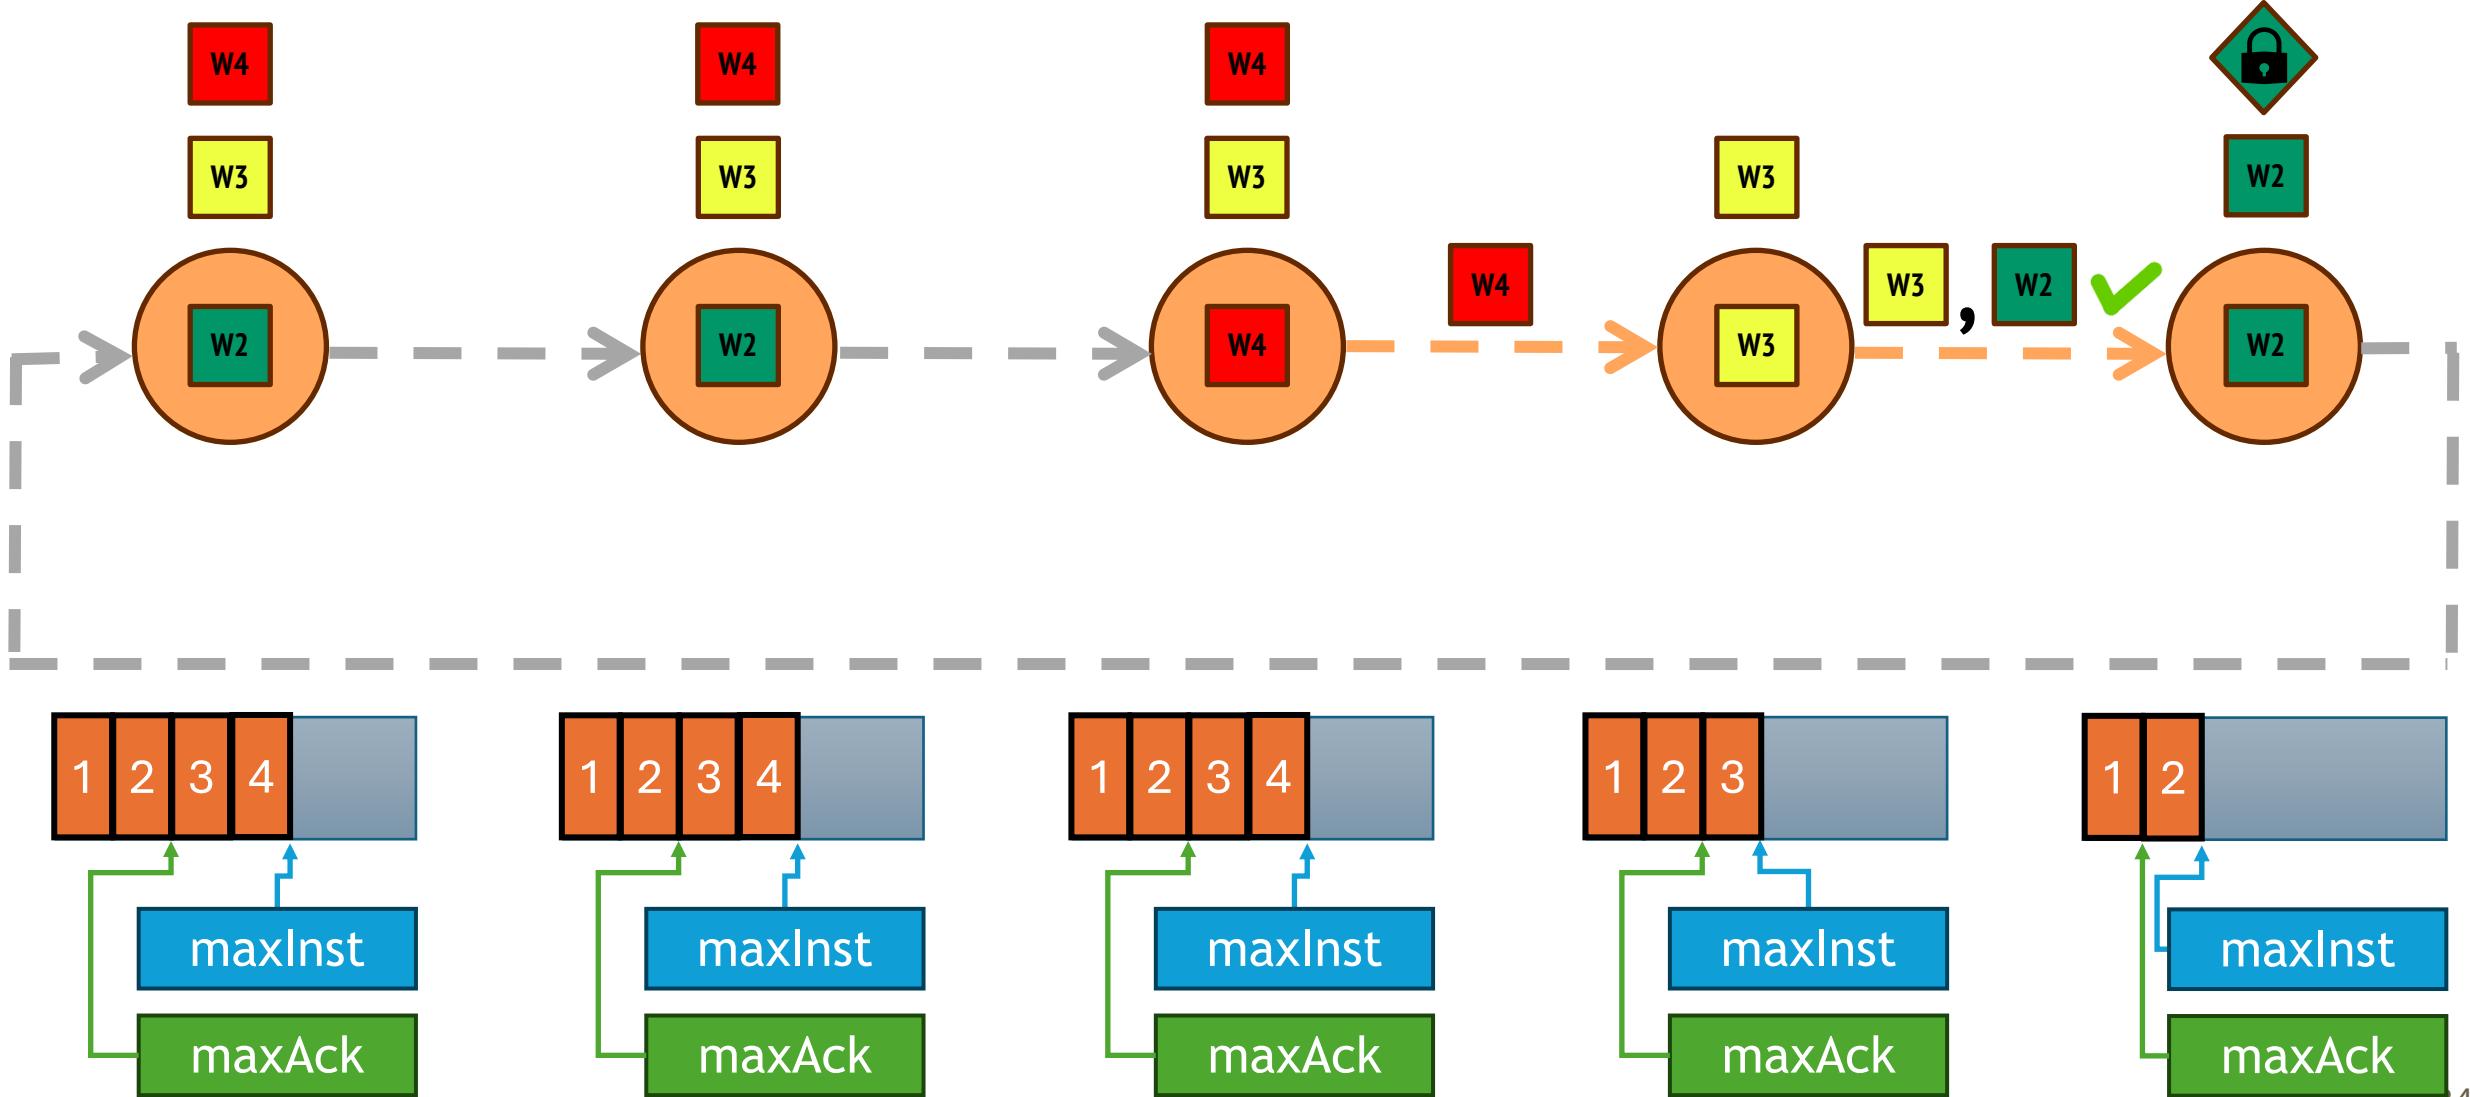
In Depth: Why not ack of max seen instance?



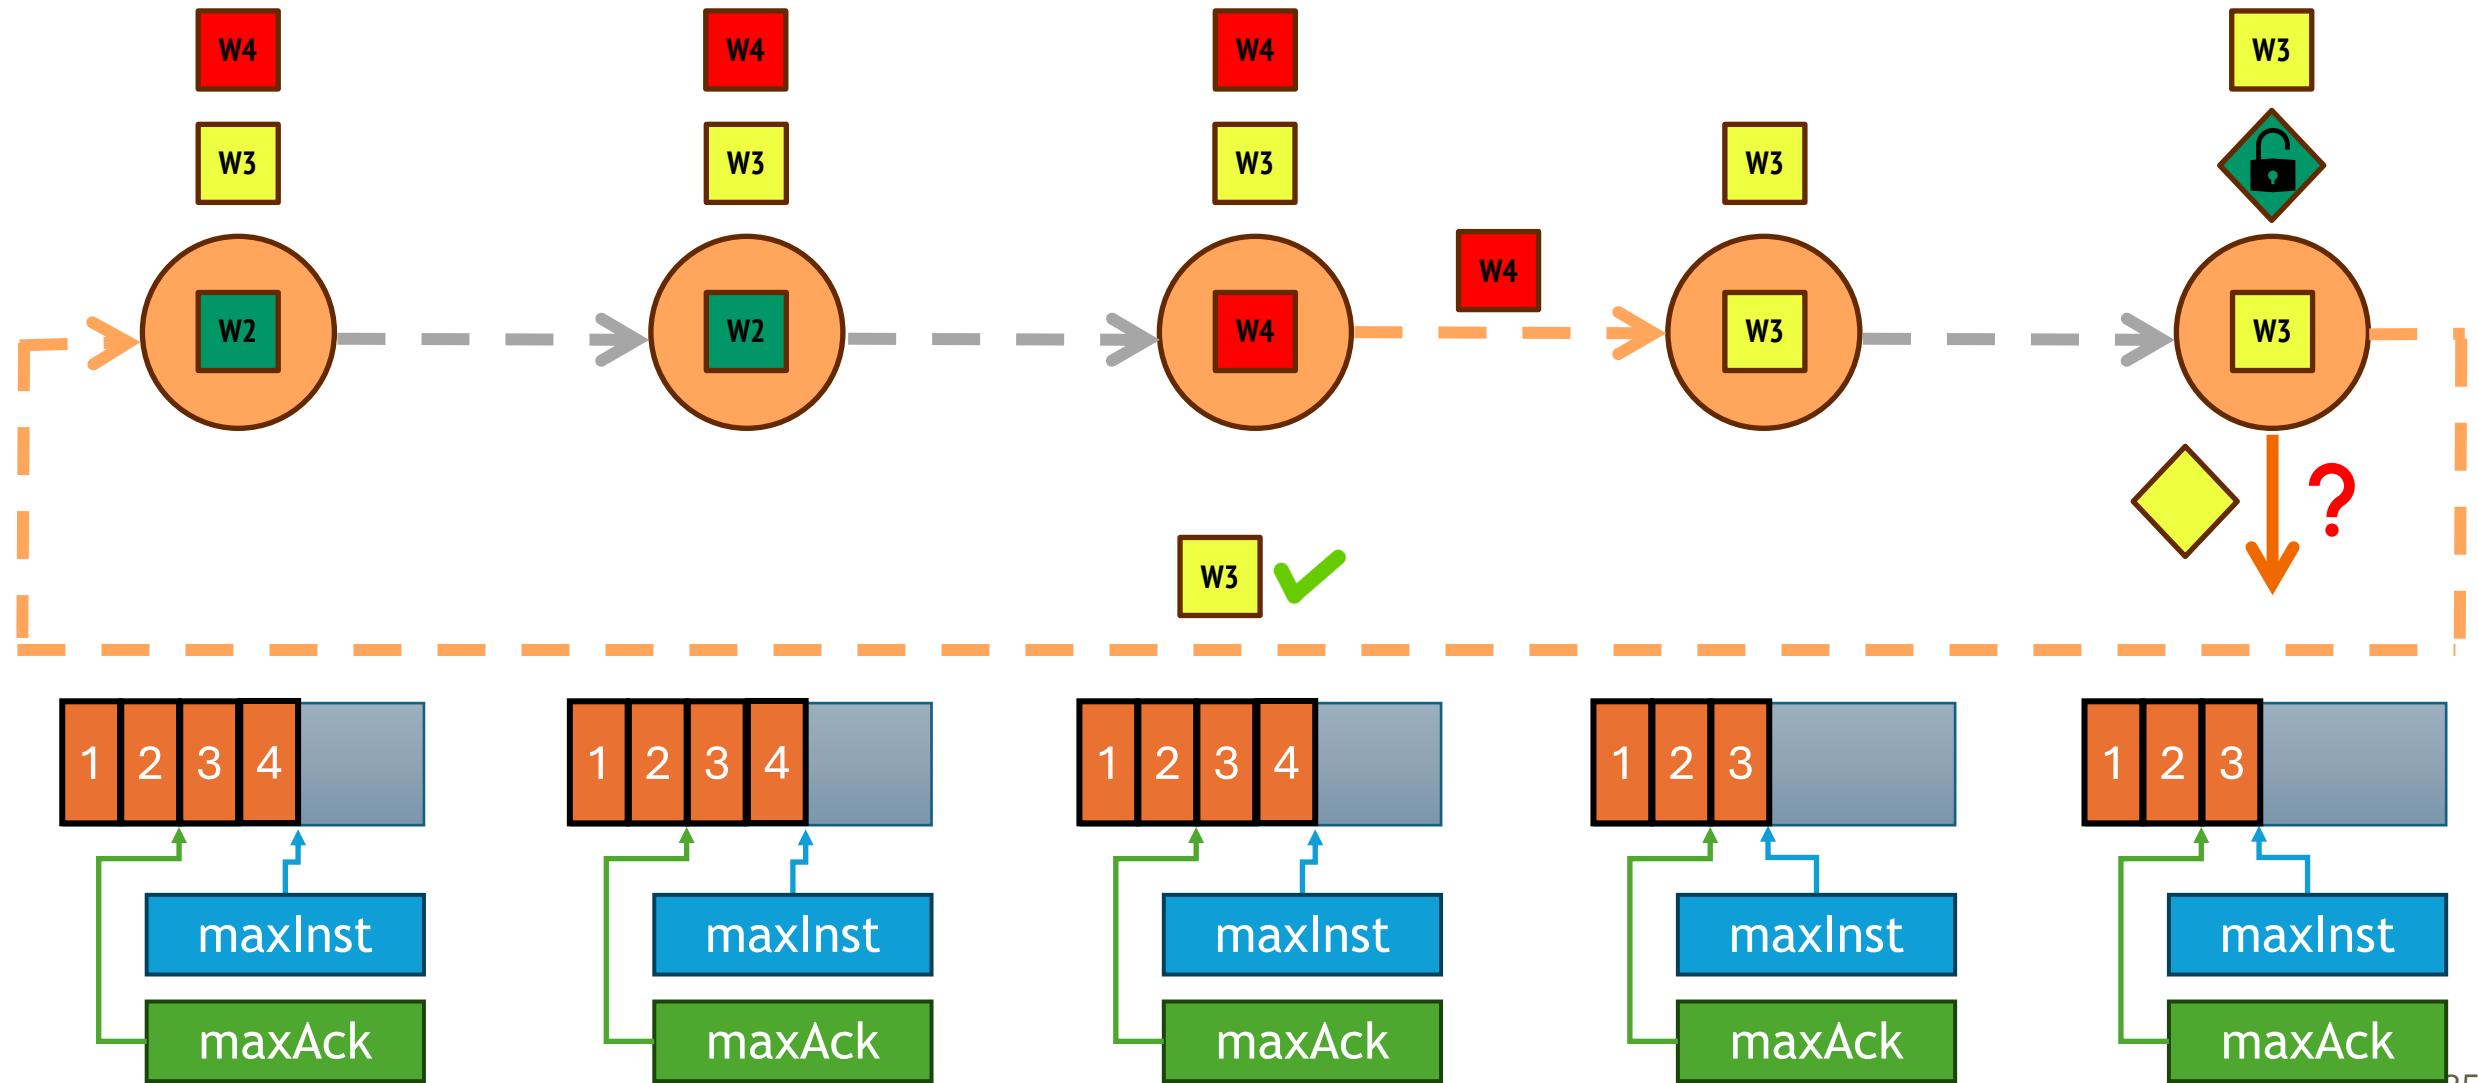
In Depth: Why not ack of max seen instance?



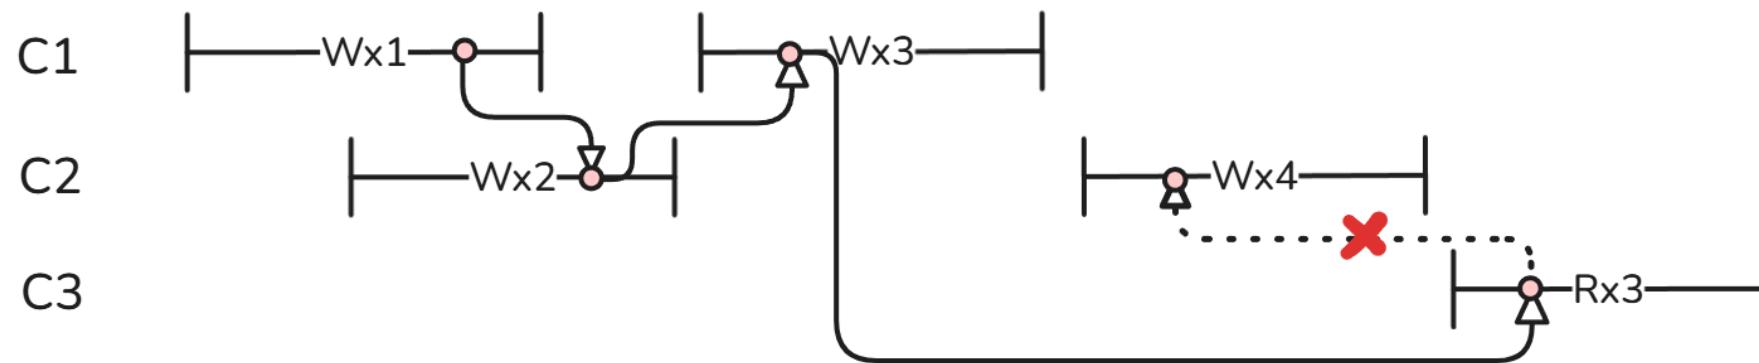
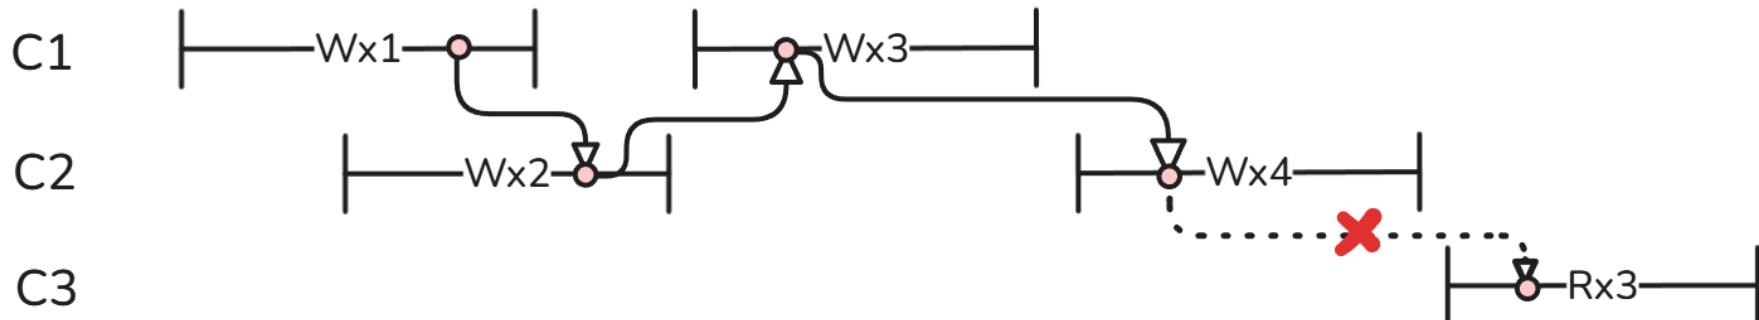
In Depth: Why not ack of max seen instance?



In Depth: Why not ack of max seen instance?



In Depth: Why not ack of max seen instance?



Not Linearizable



In Depth: Reads

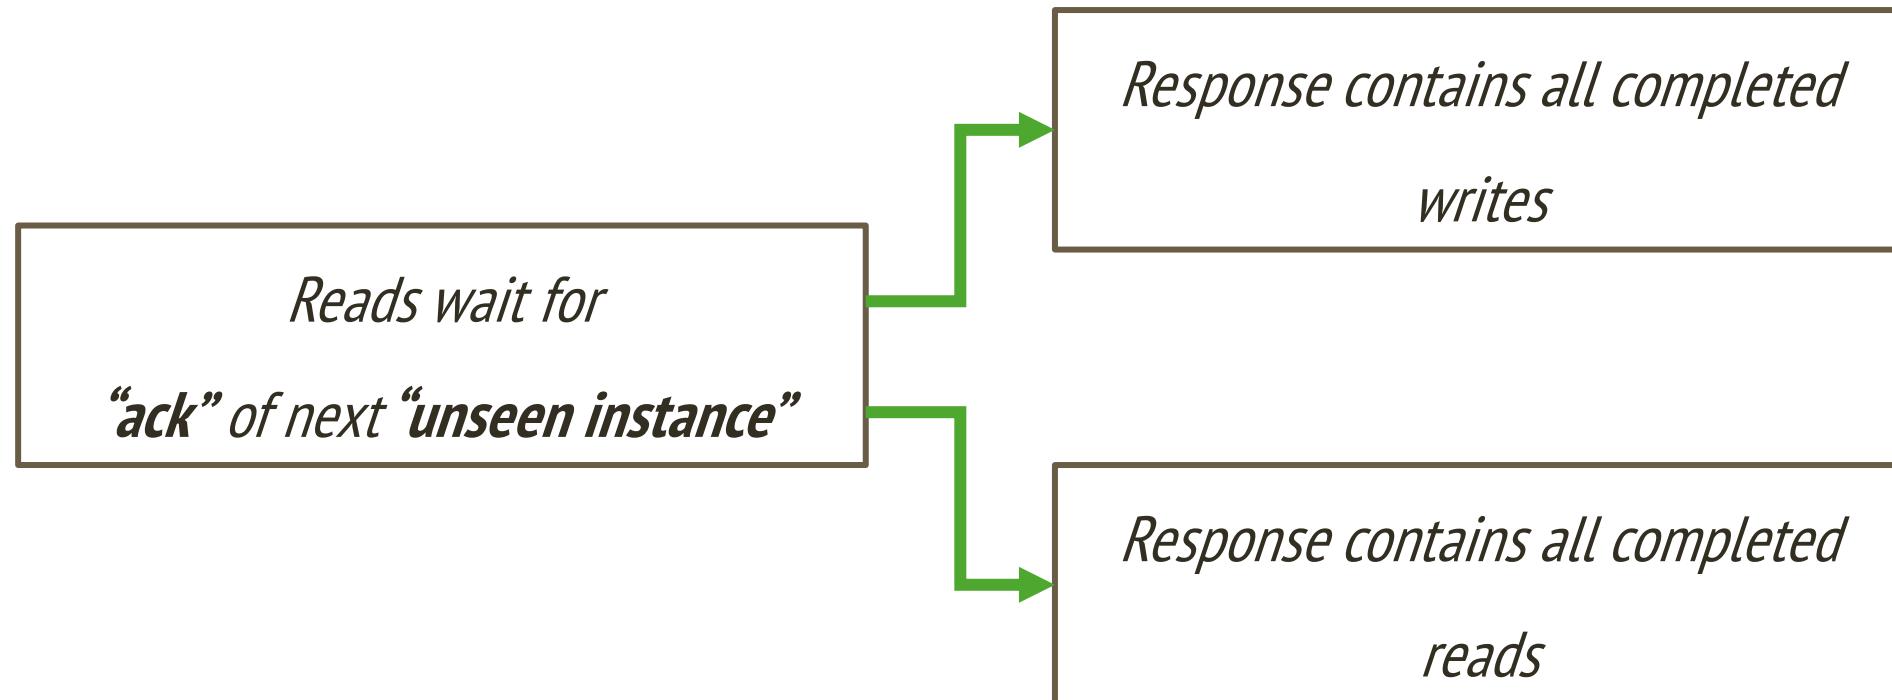
Why acks of next unseen instance?

~~*Why not immediately reply?*~~

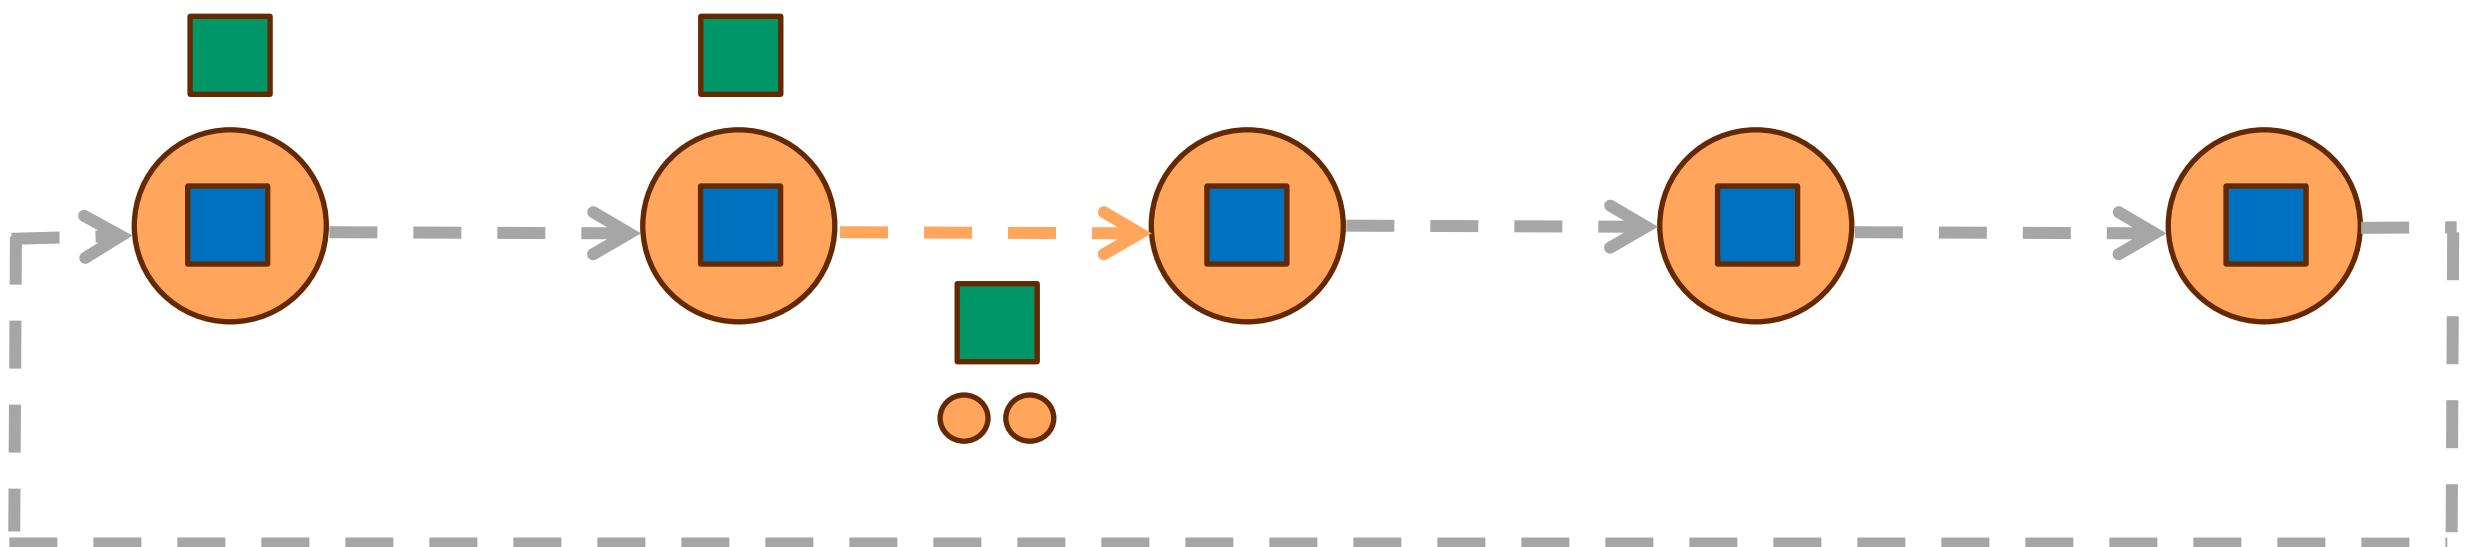
~~*Why not next unseen instance?*~~

~~*Why not ack of max seen instance?*~~

In Depth: Reads

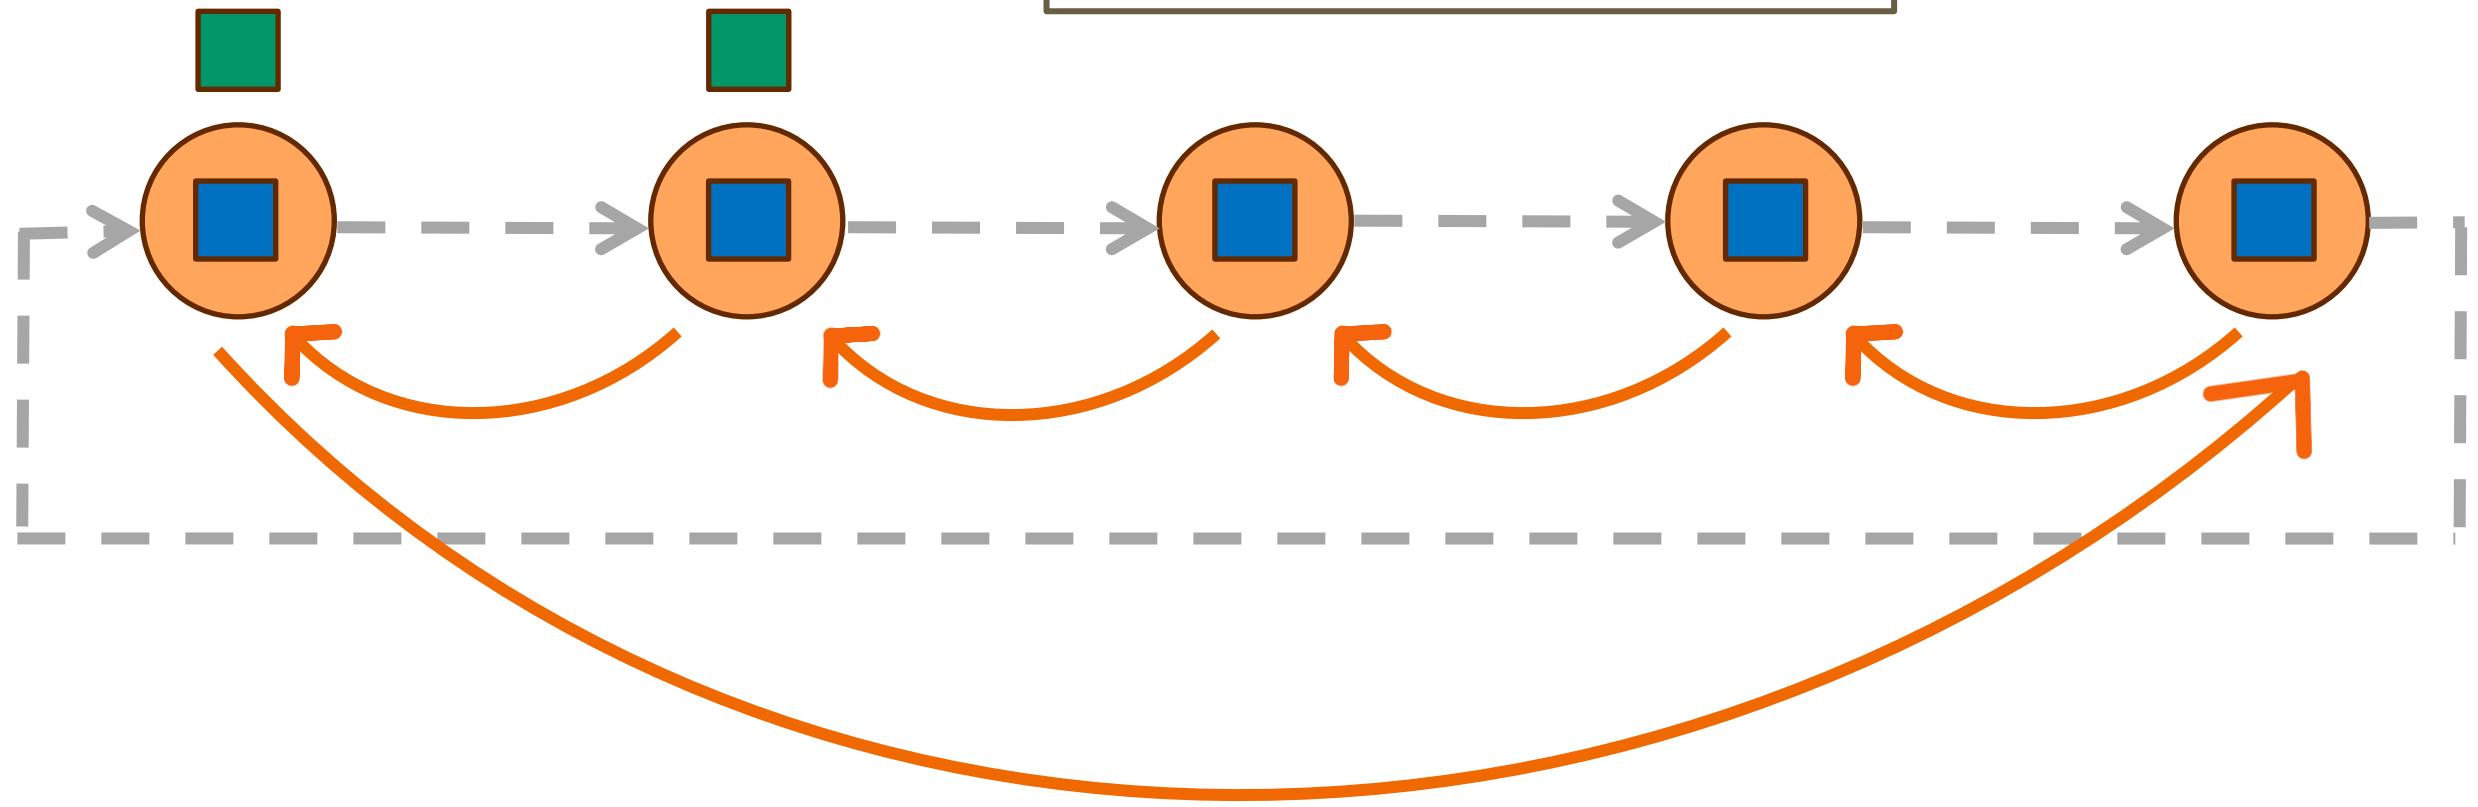


Membership Management: Removal

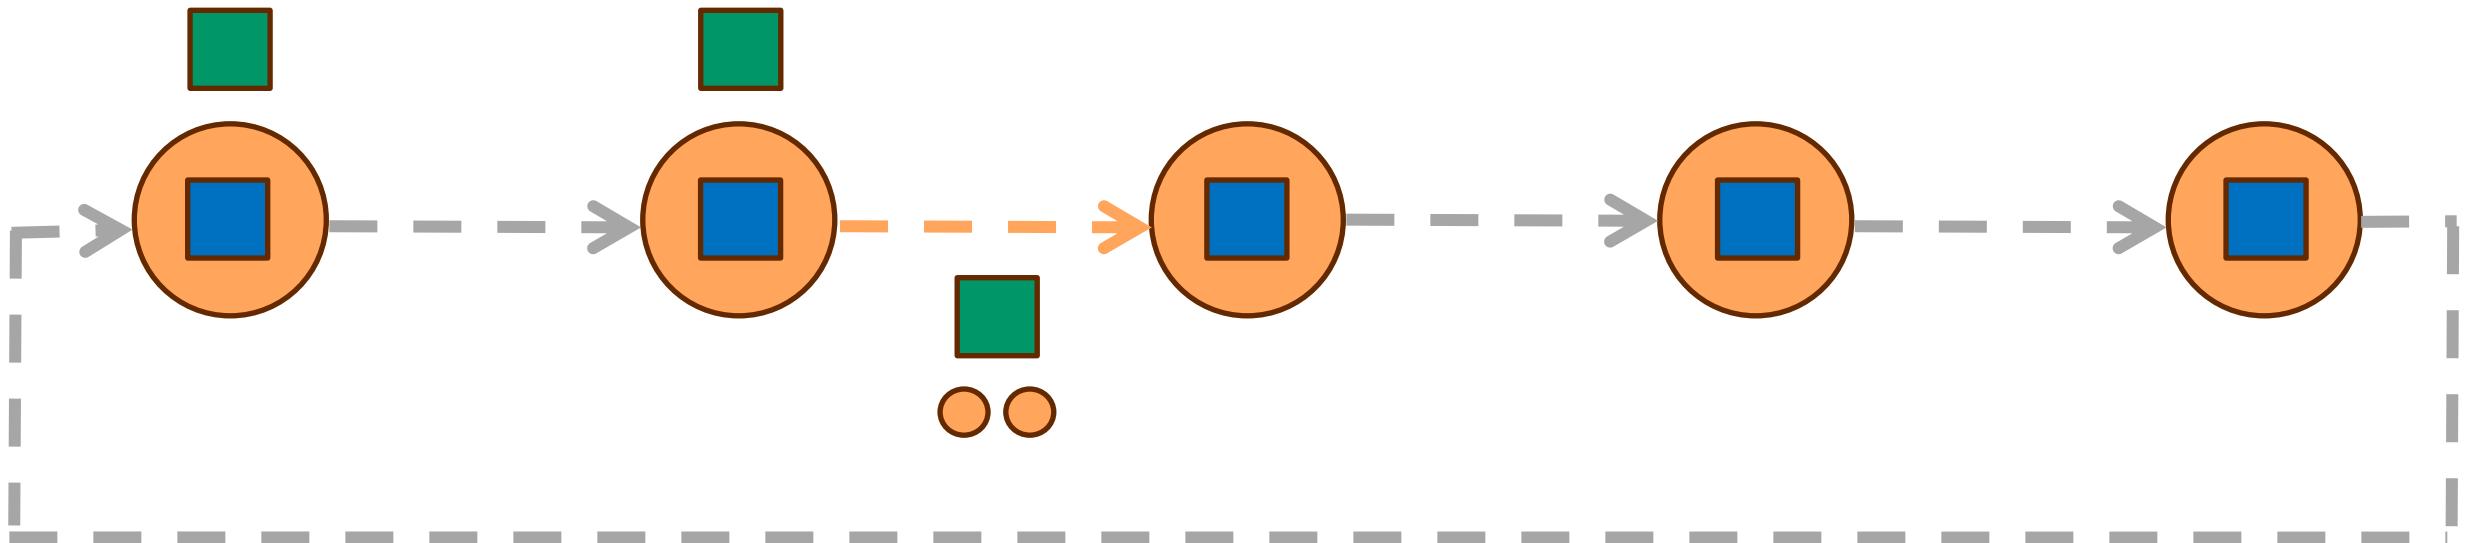


Membership Management: Removal

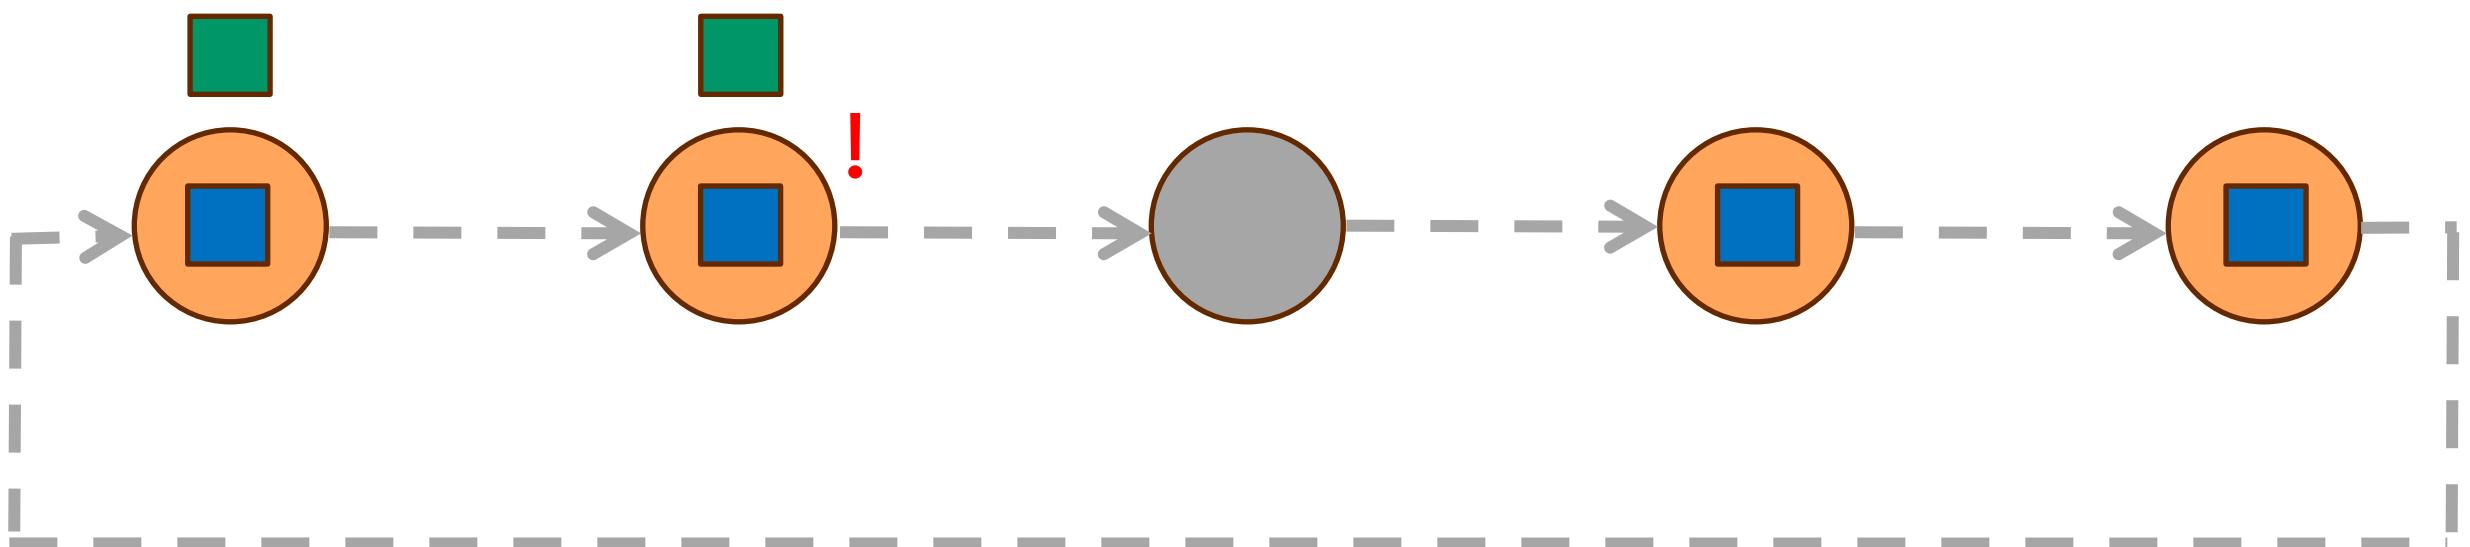
*Replicas send "keep-alive" messages
to previous replica*



Membership Management: Removal

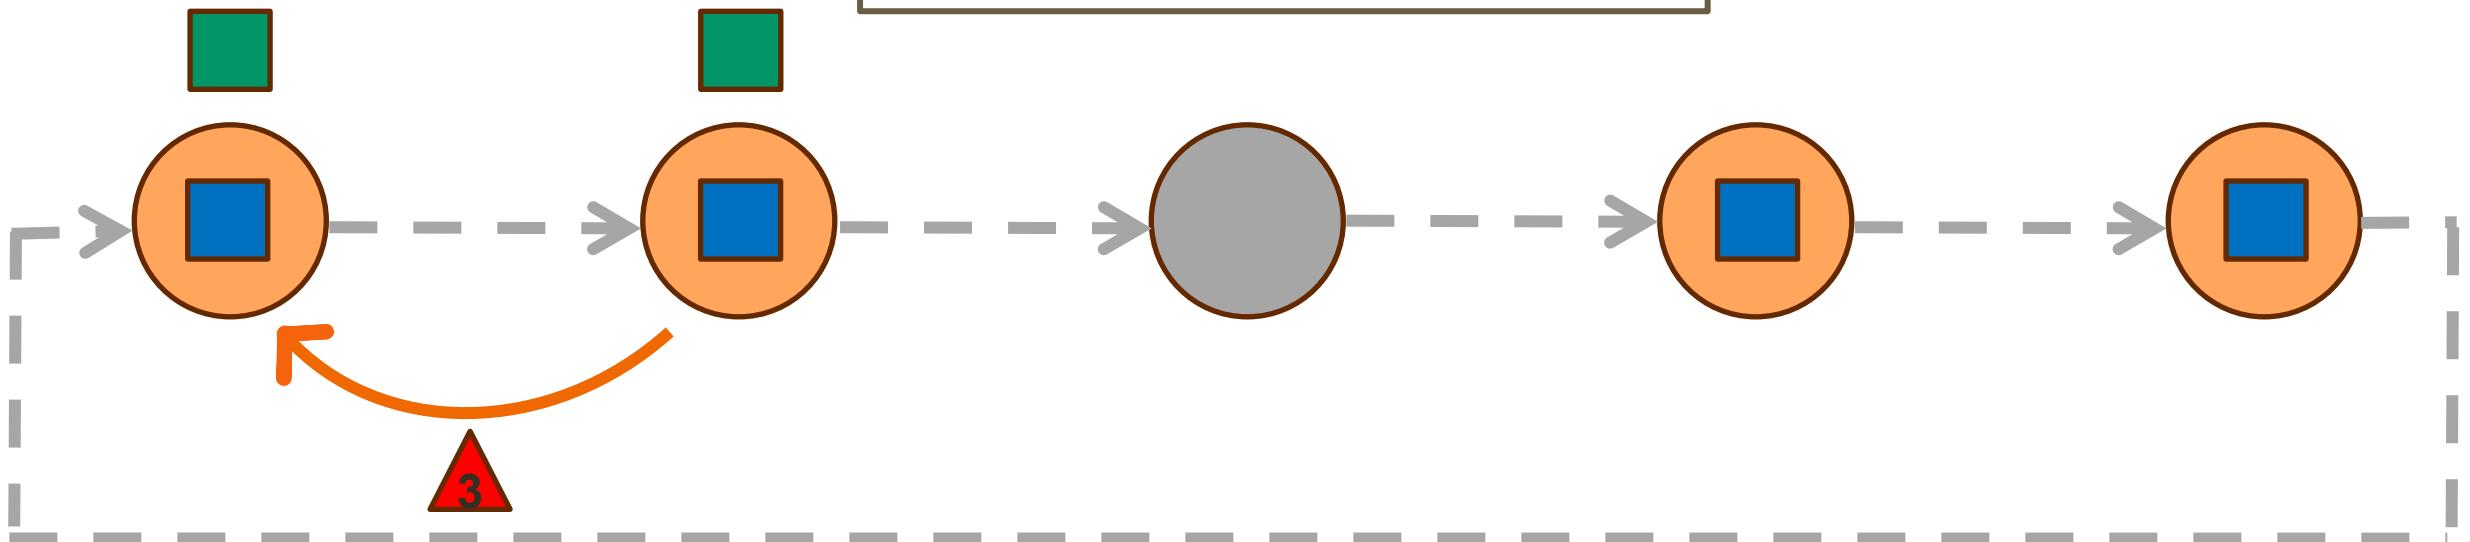


Membership Management: Removal



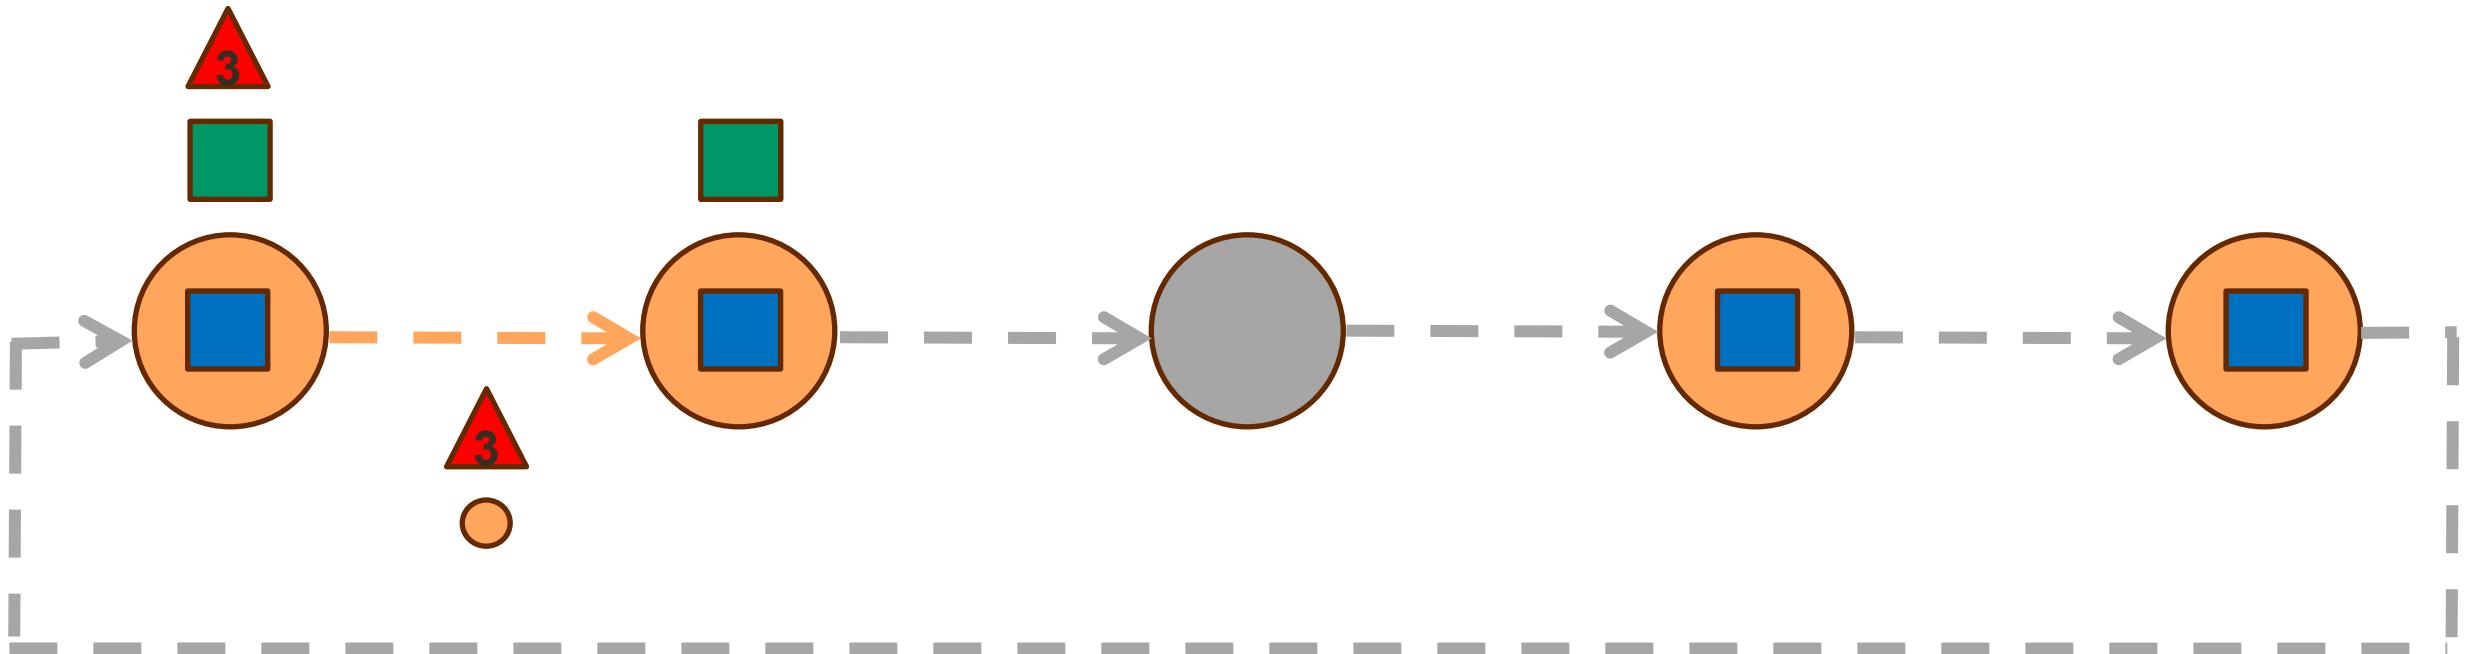
Membership Management: Removal

*Remove requests are handled like
regular write operations*

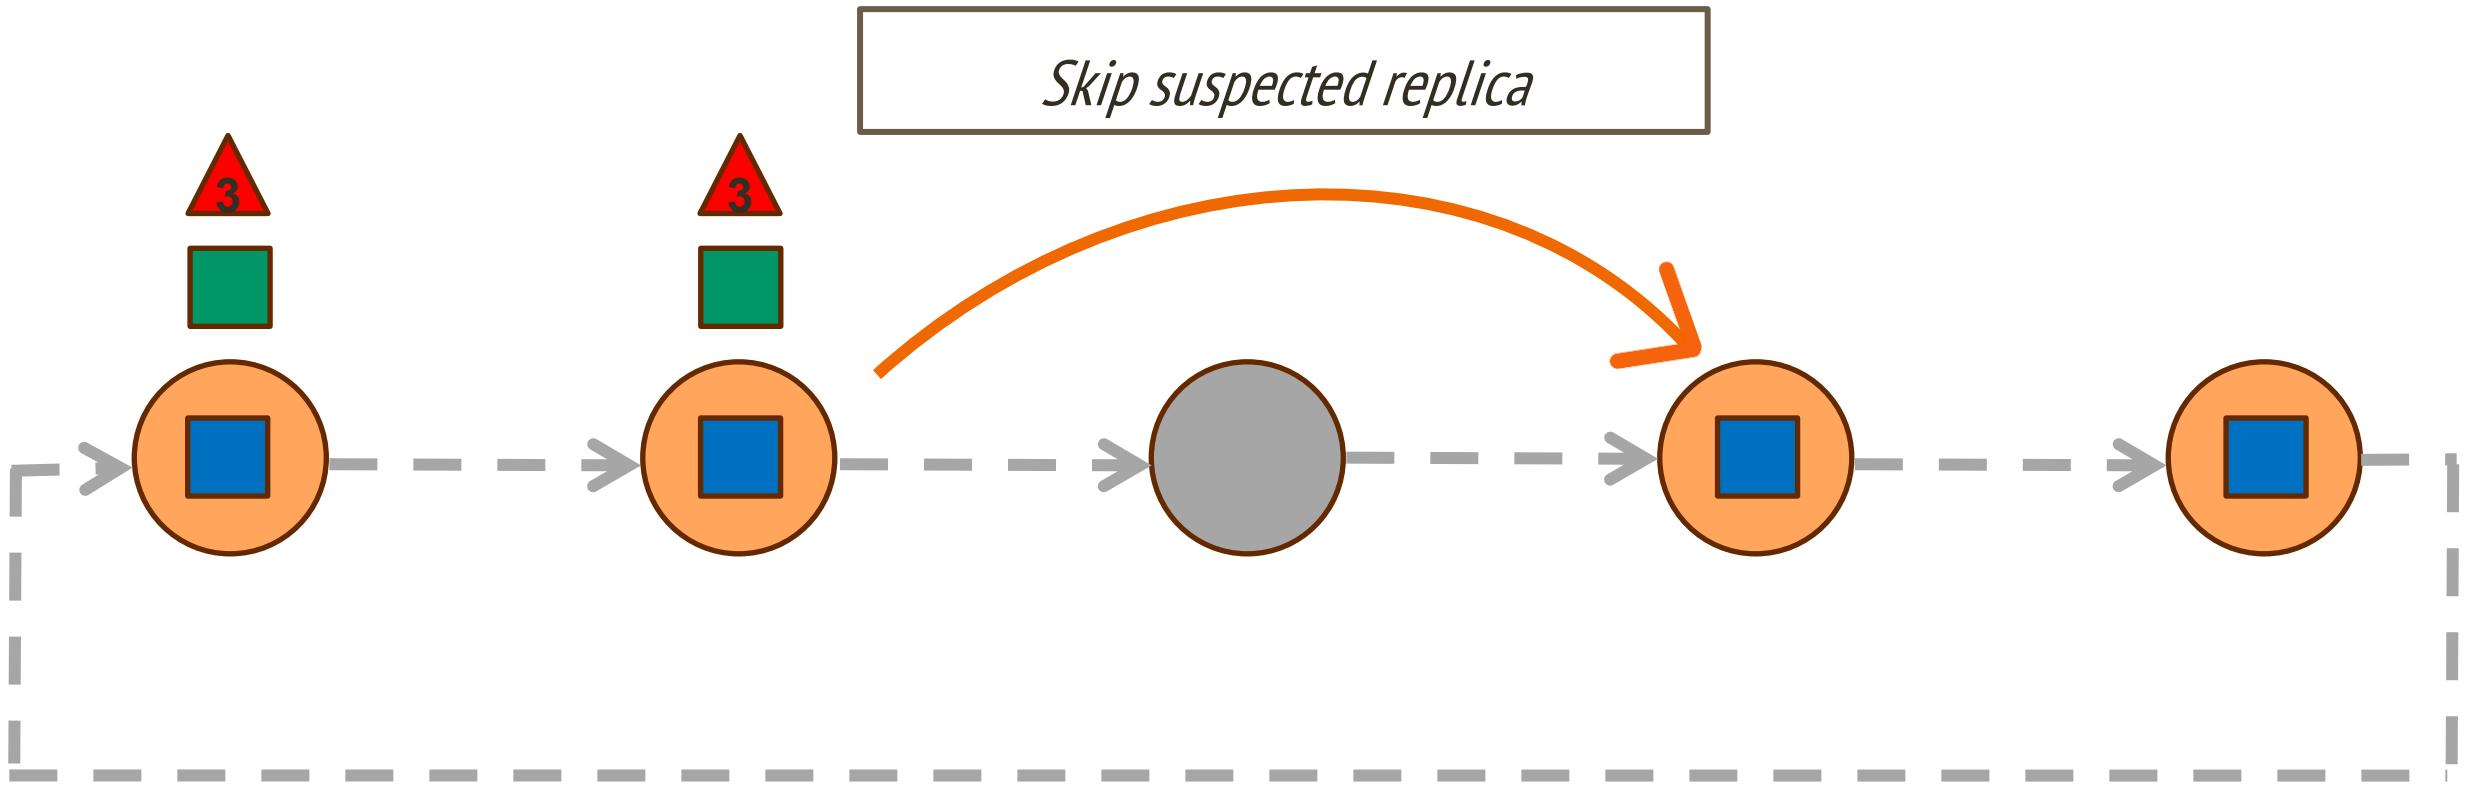


: Remove i'th replica request

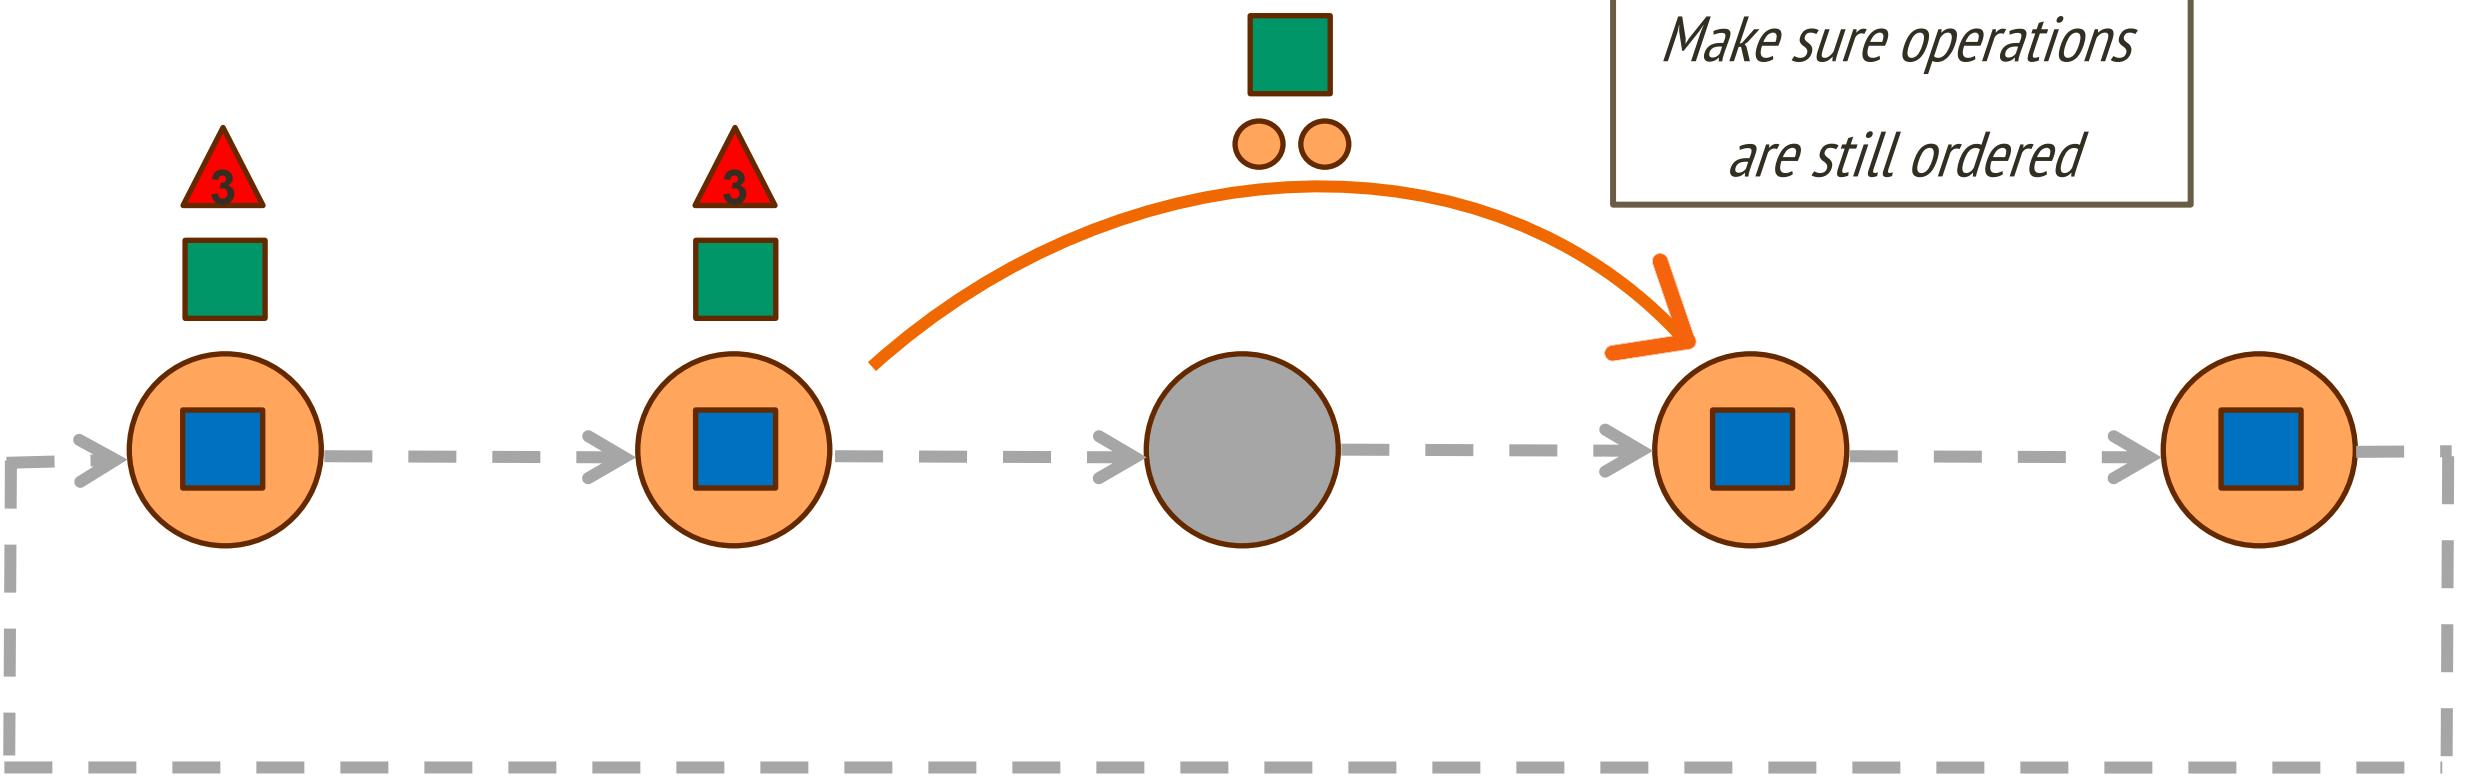
Membership Management: Removal



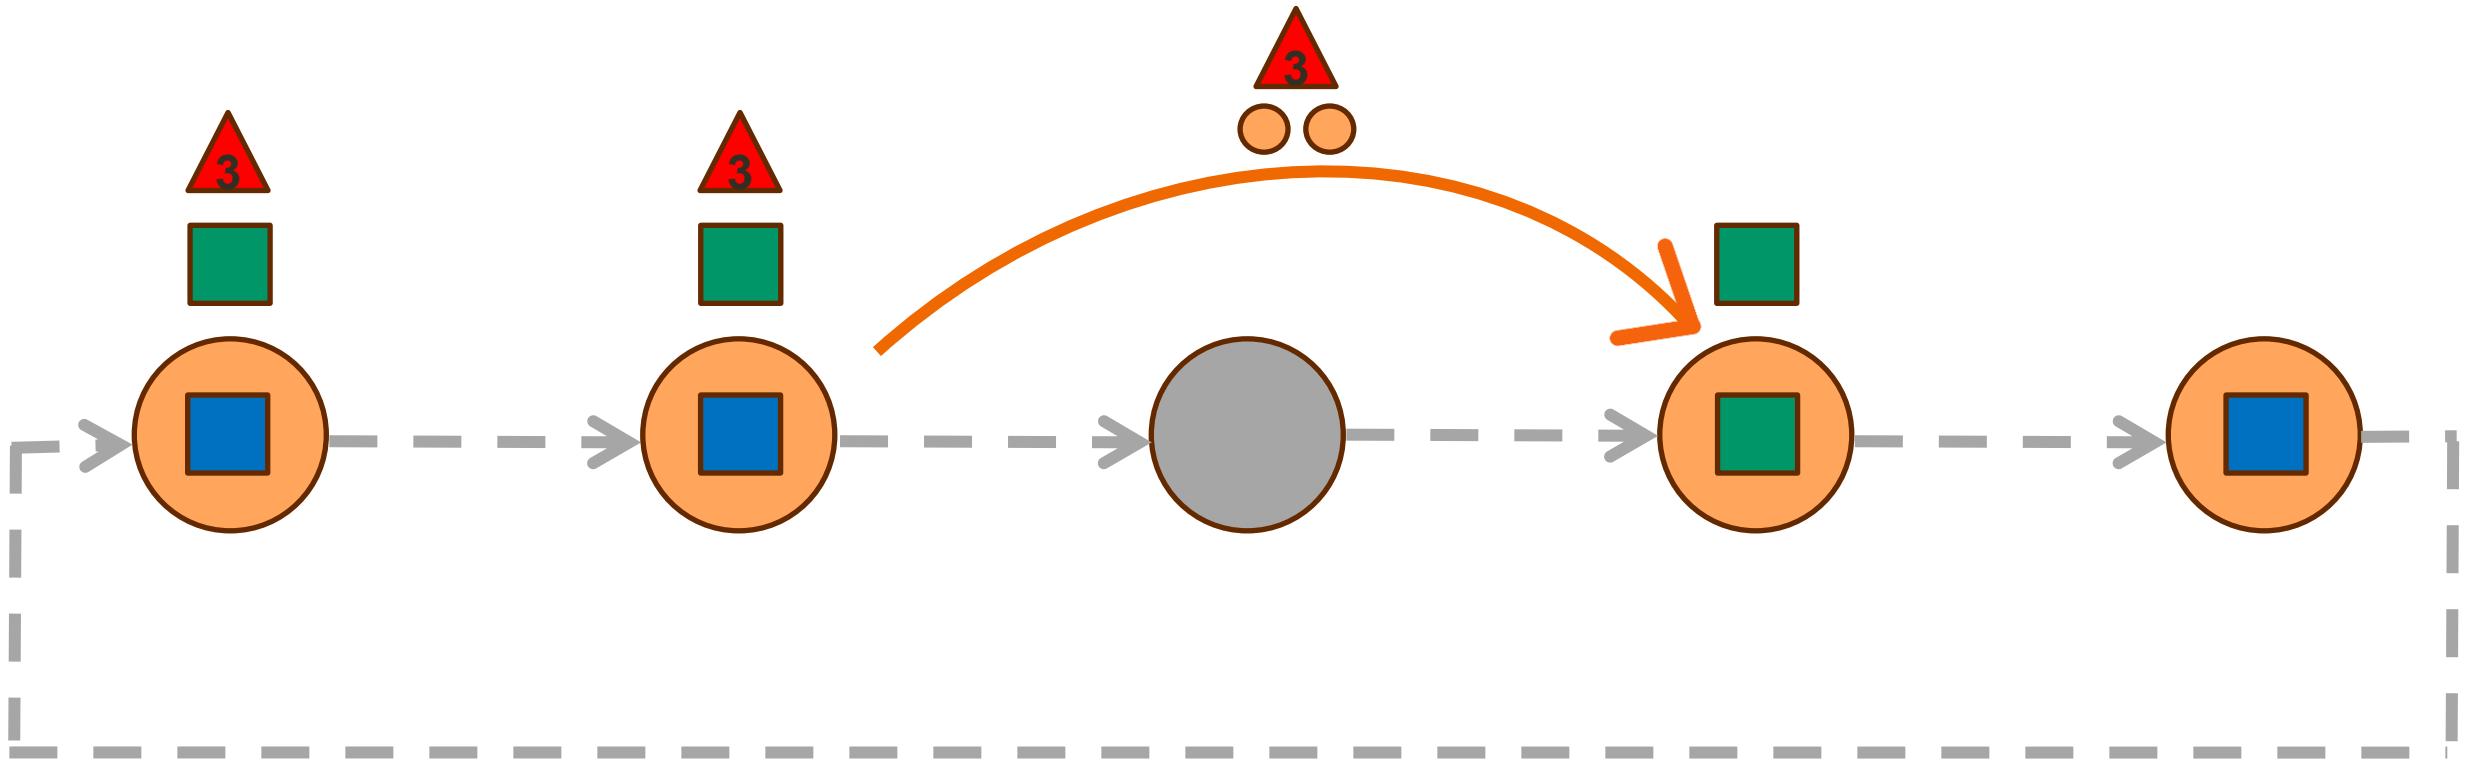
Membership Management: Removal



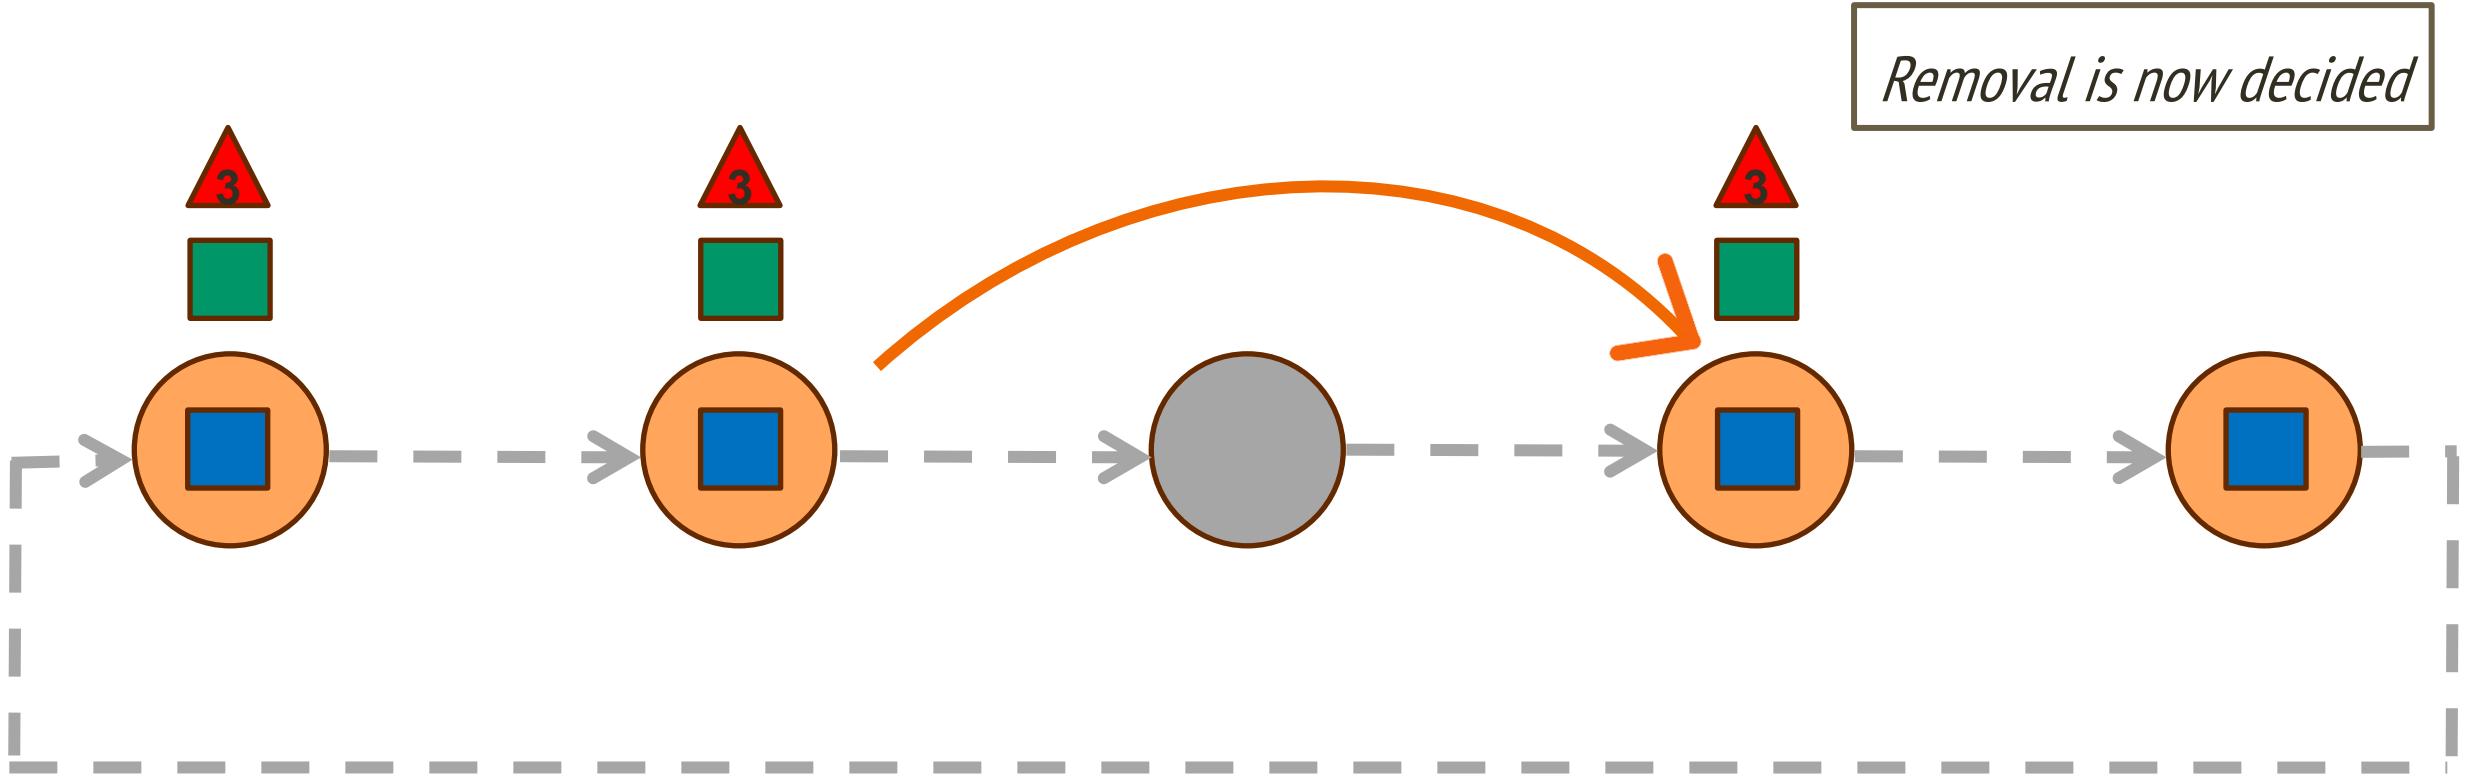
Membership Management: Removal



Membership Management: Removal



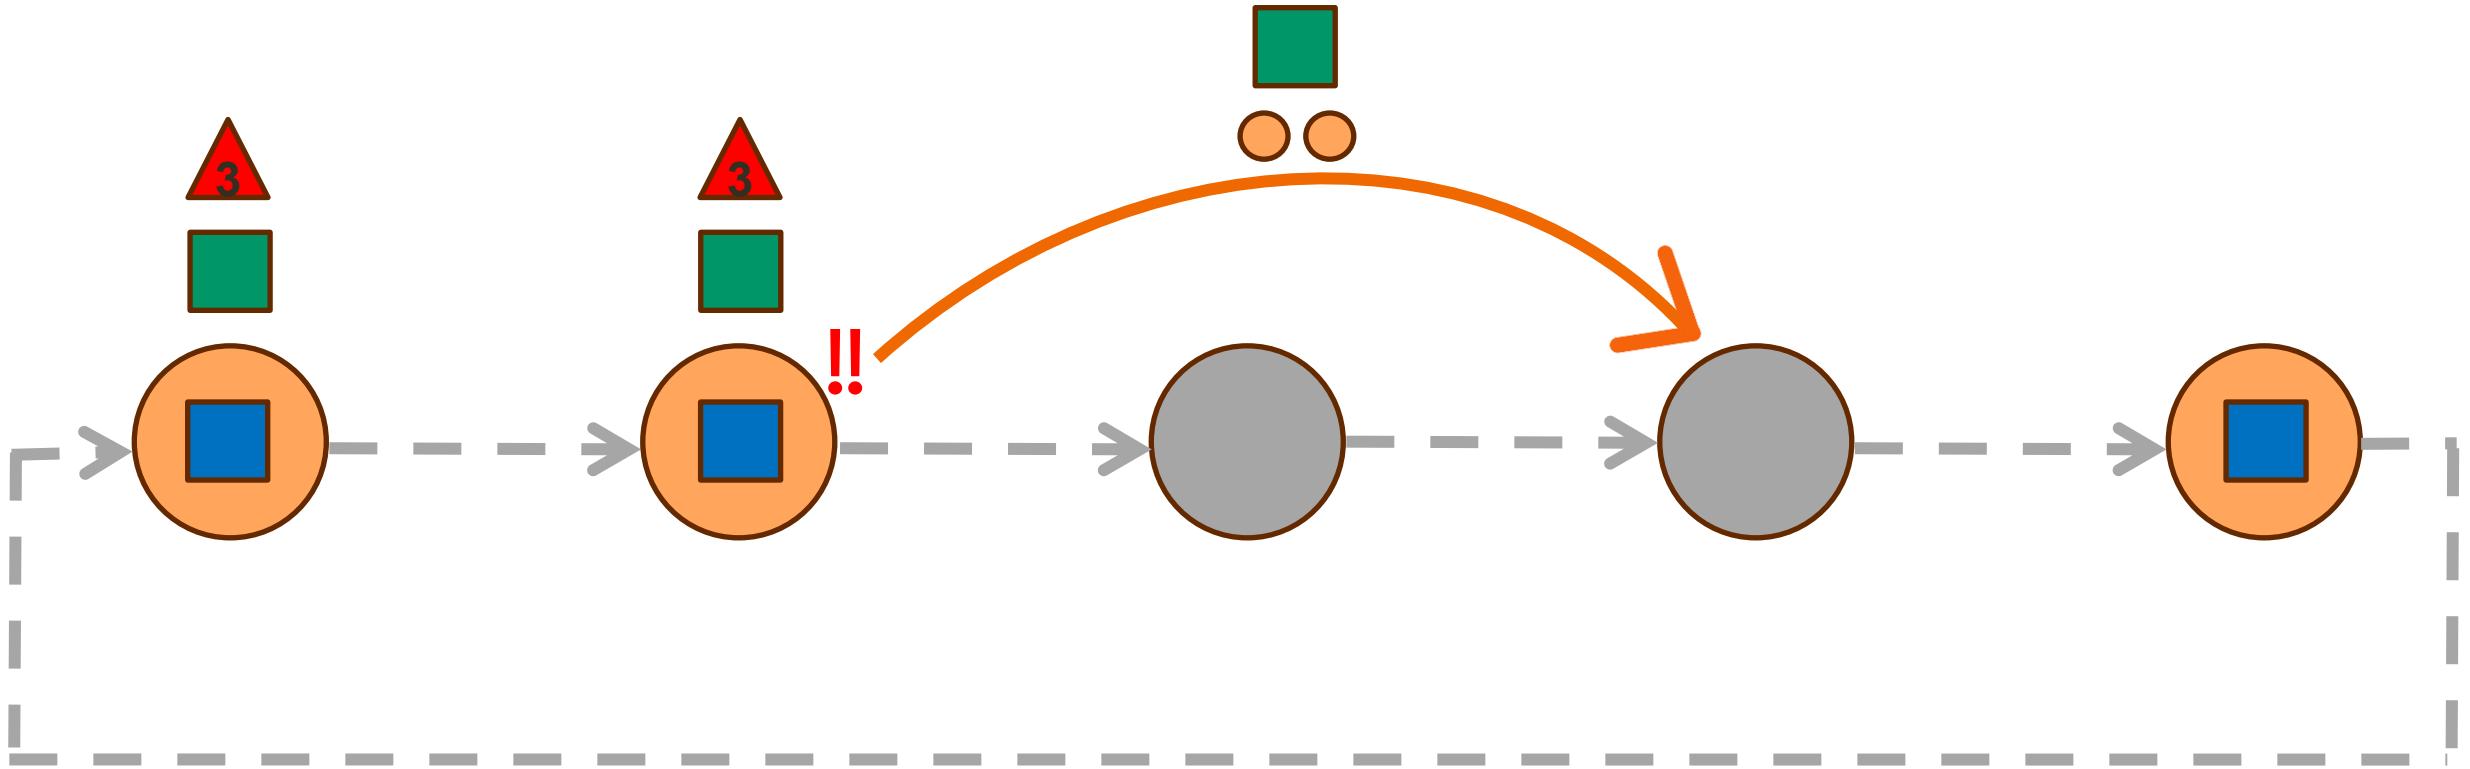
Membership Management: Removal



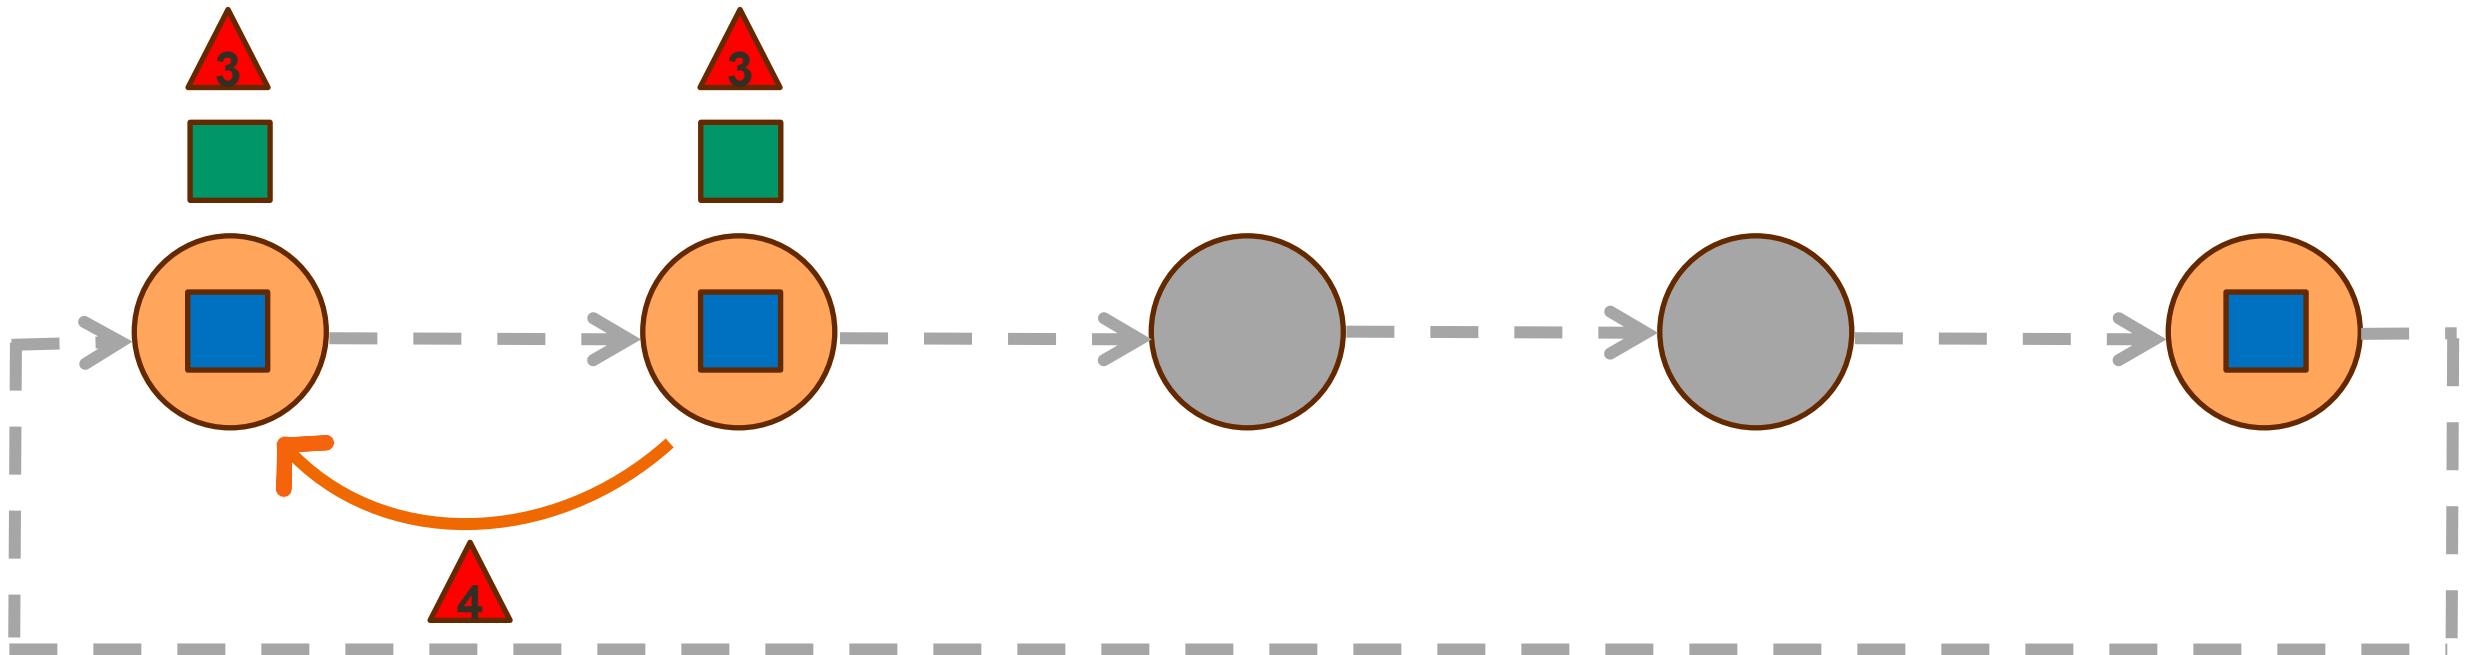
Membership Management: Removal



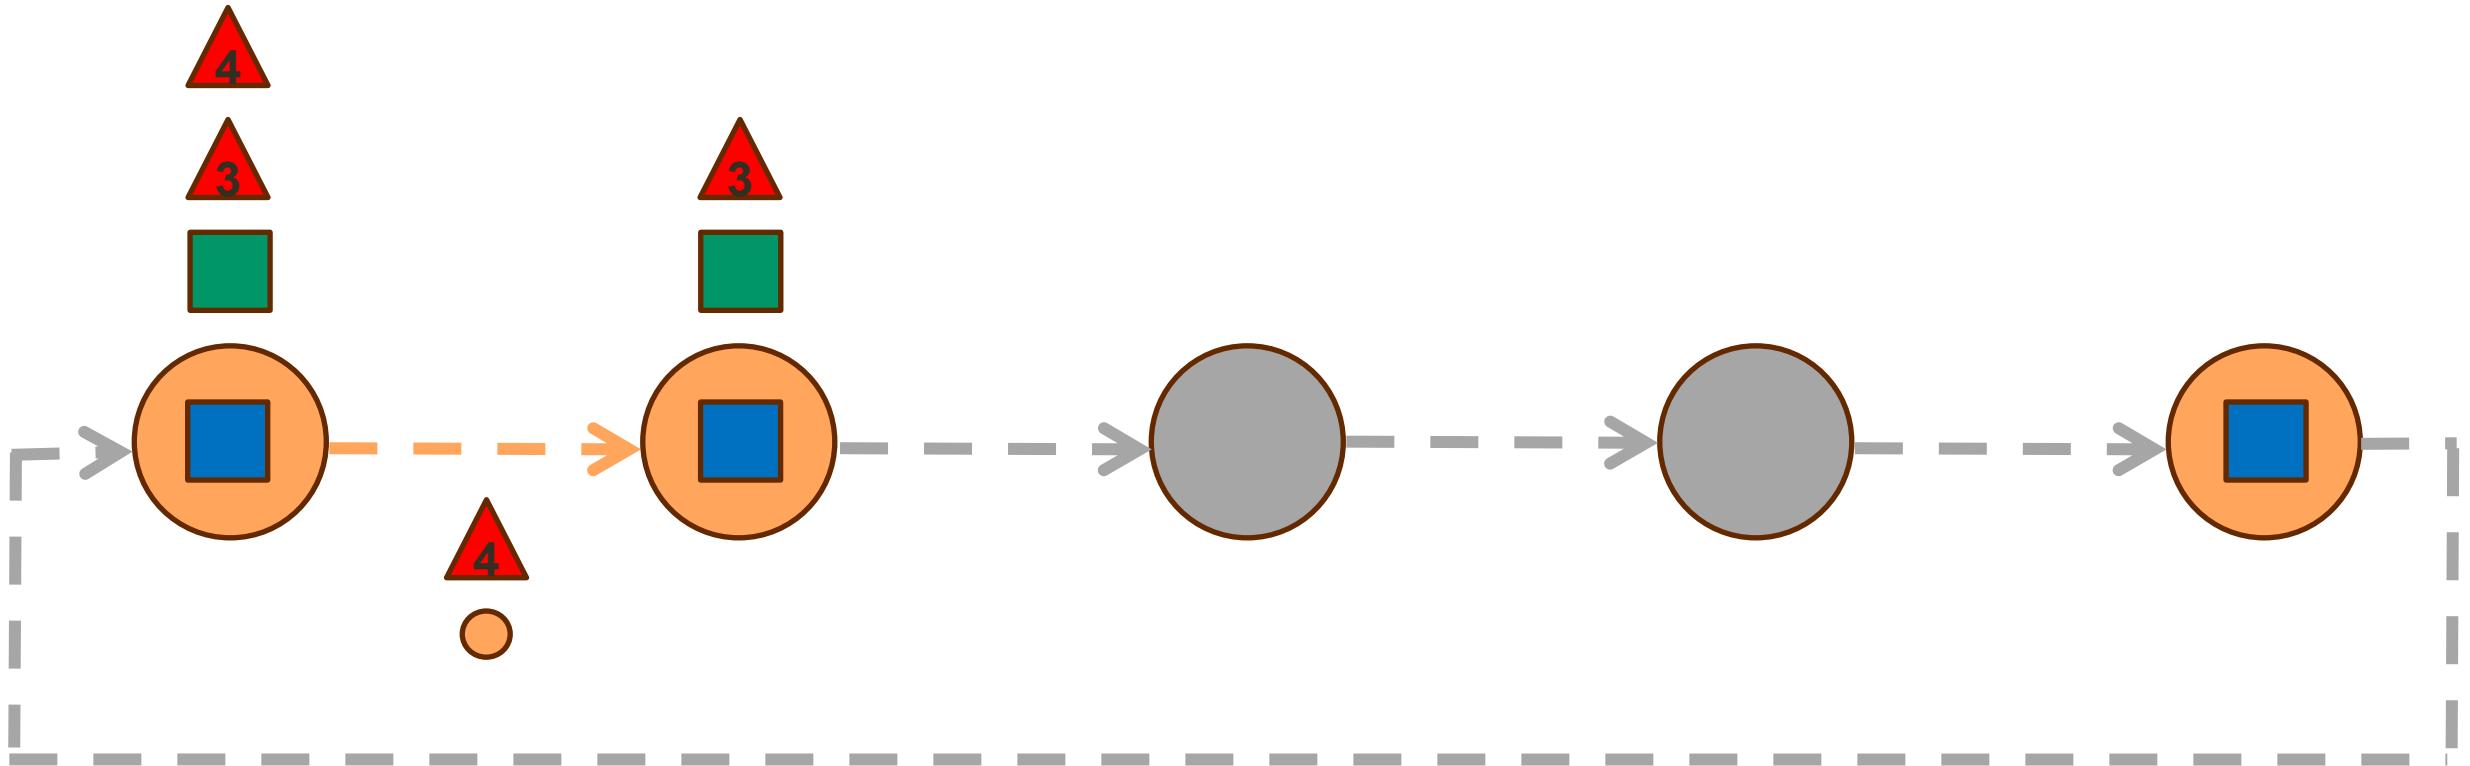
Membership Management: Removal



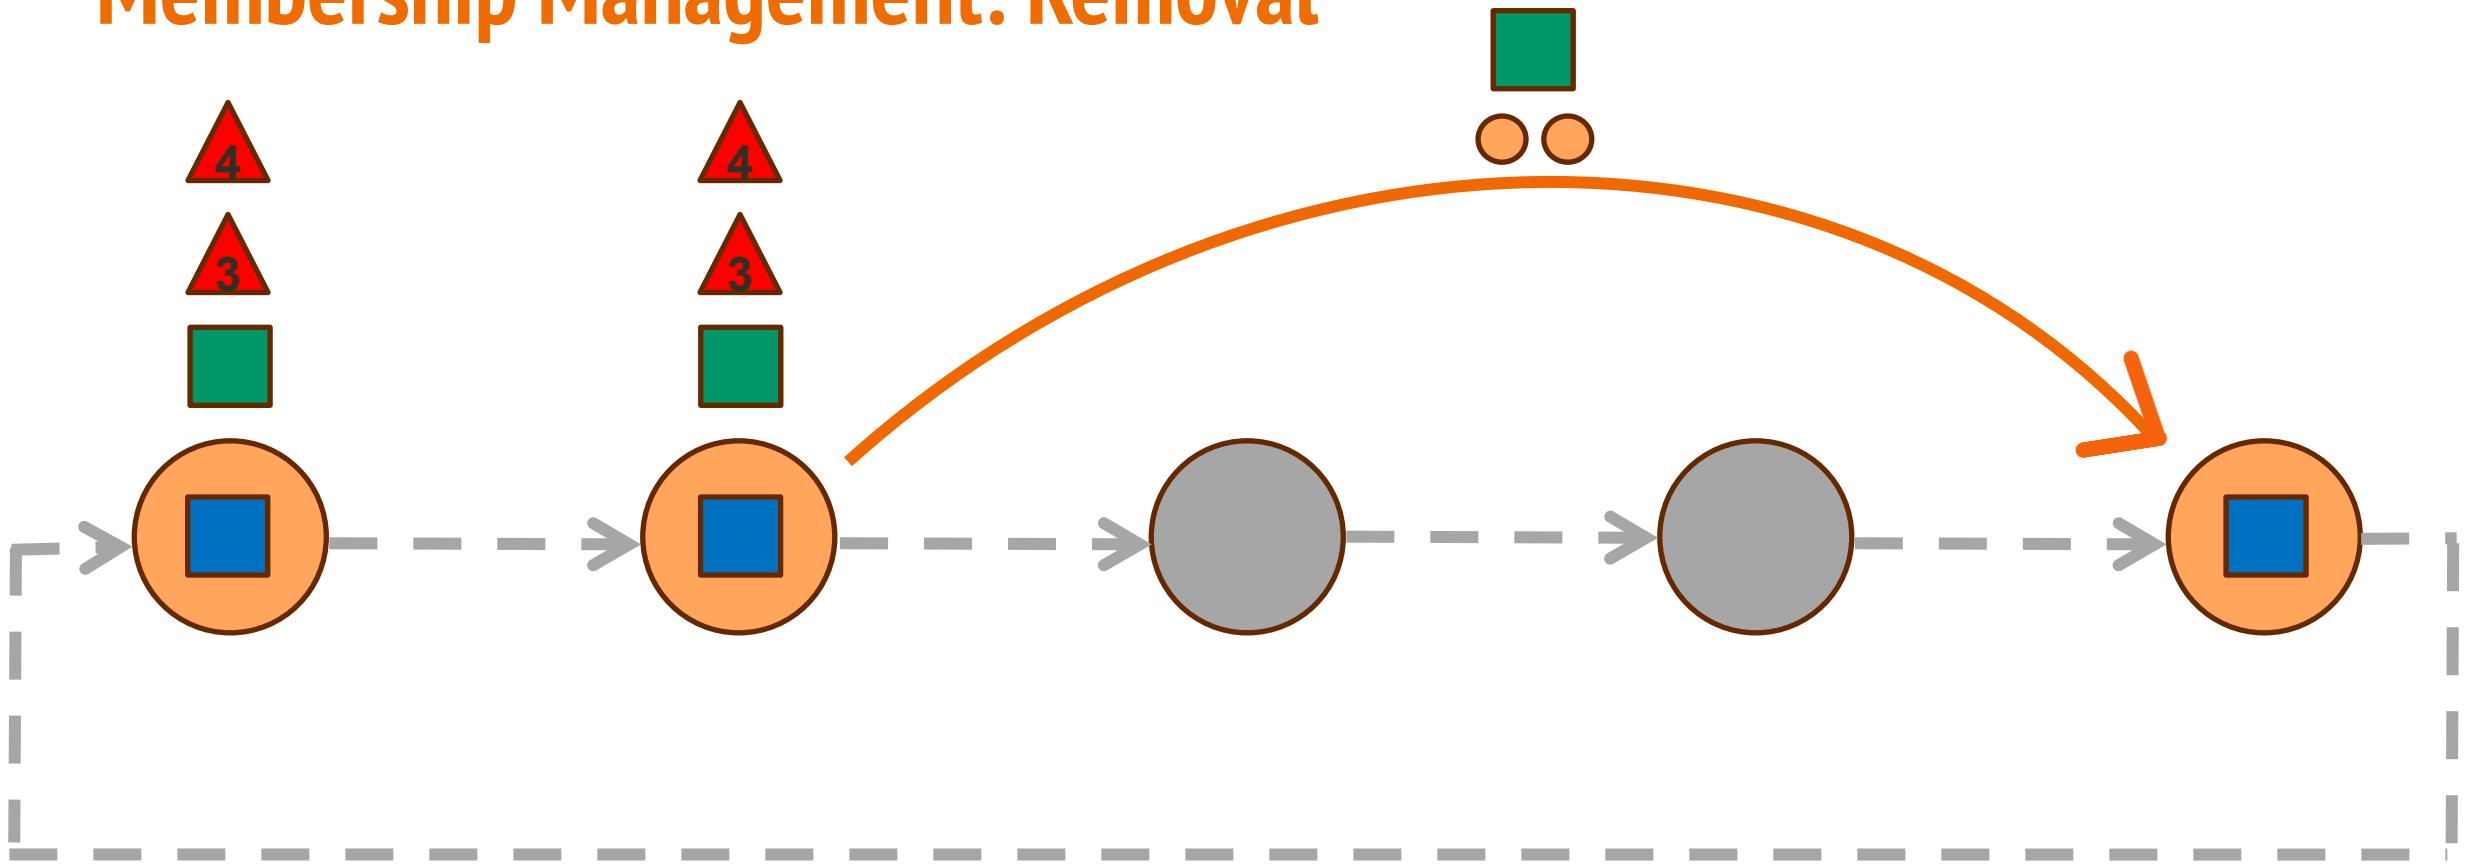
Membership Management: Removal



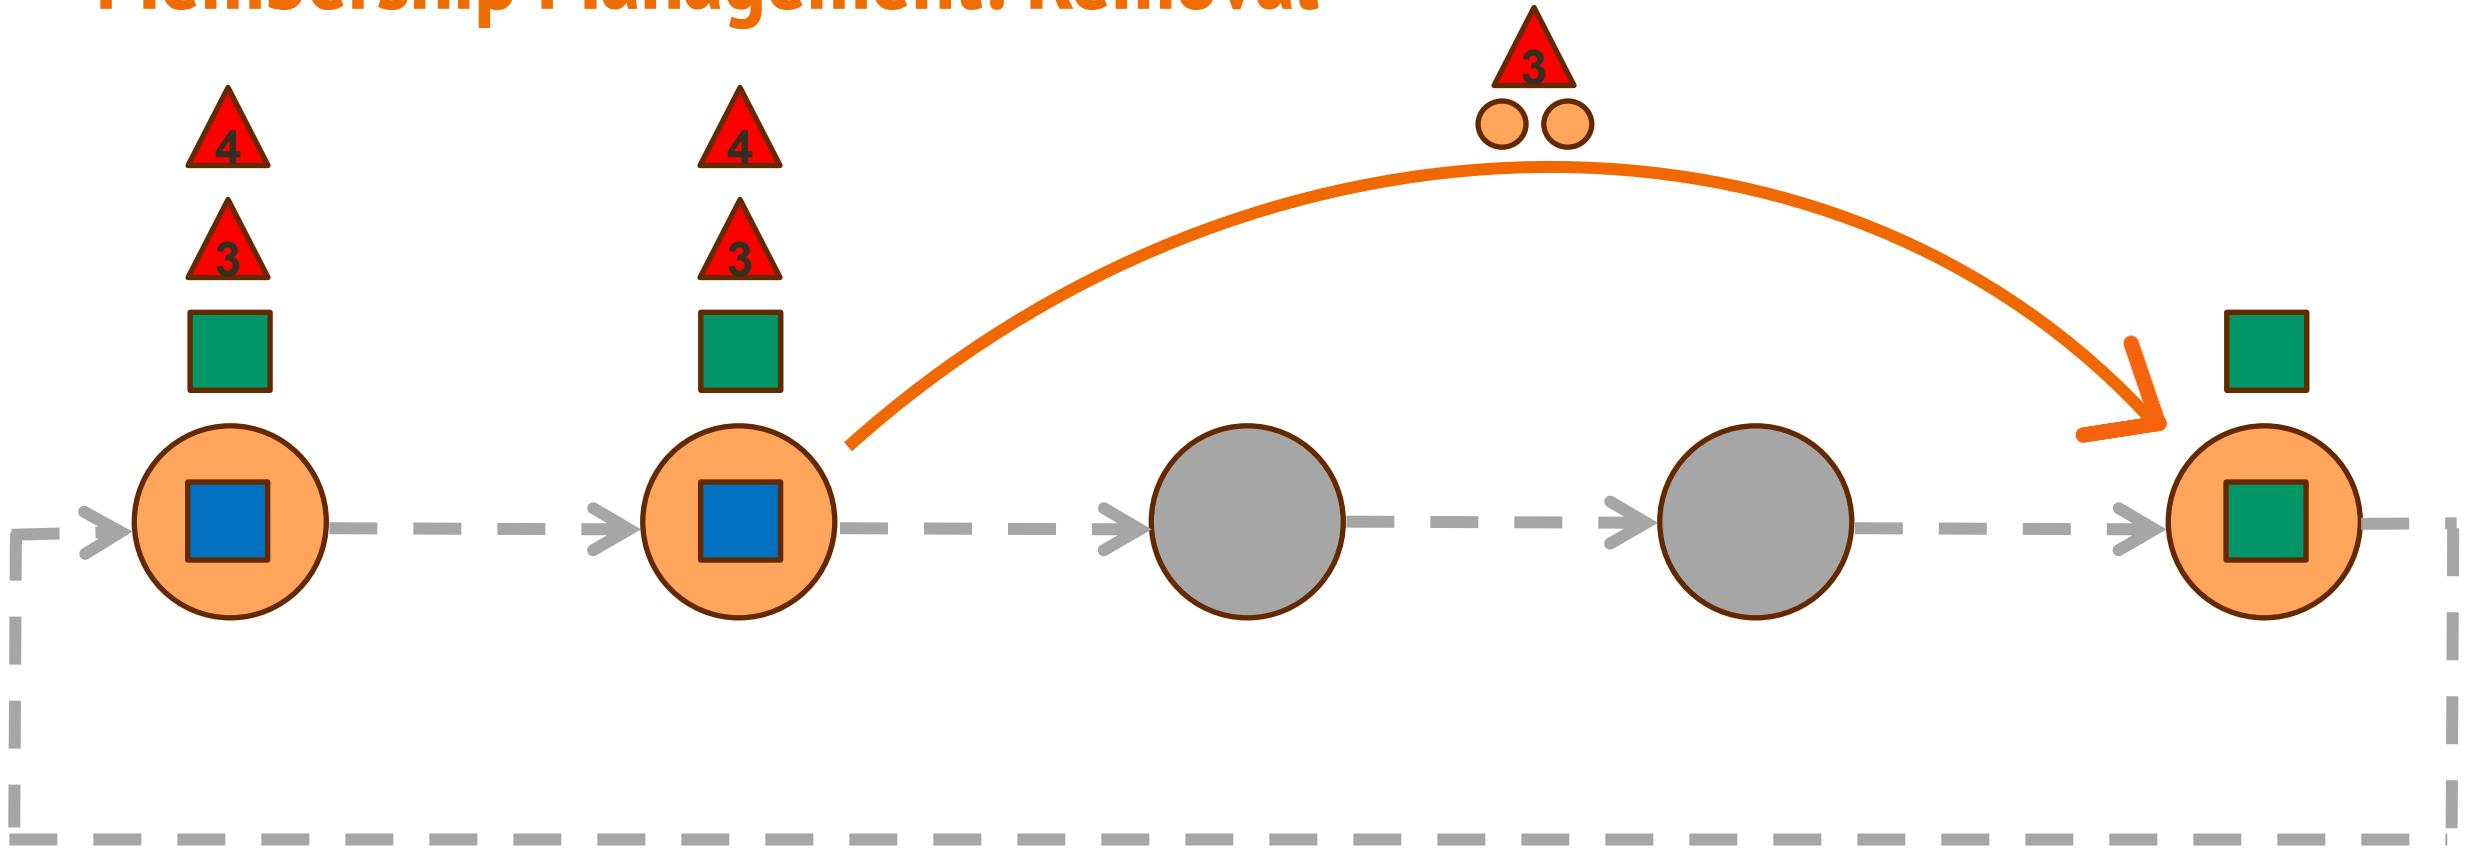
Membership Management: Removal



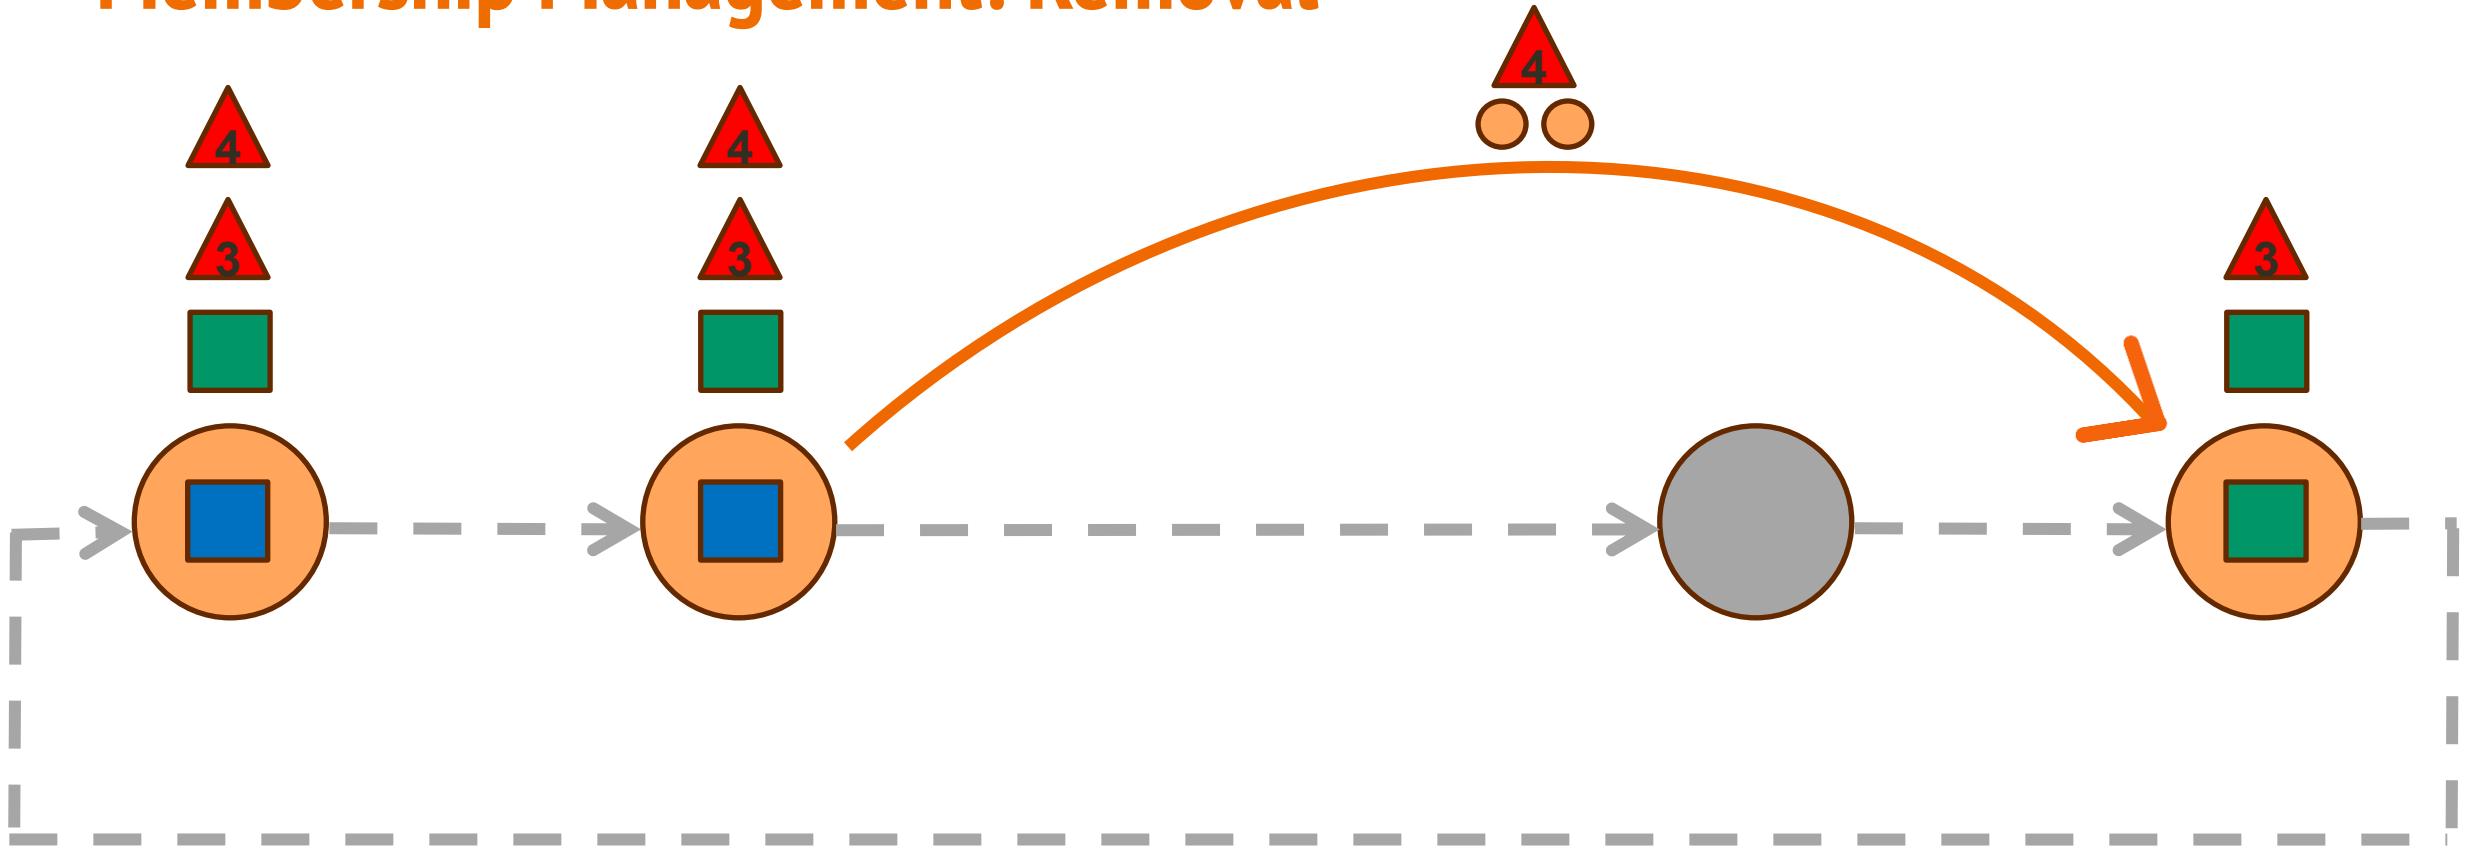
Membership Management: Removal



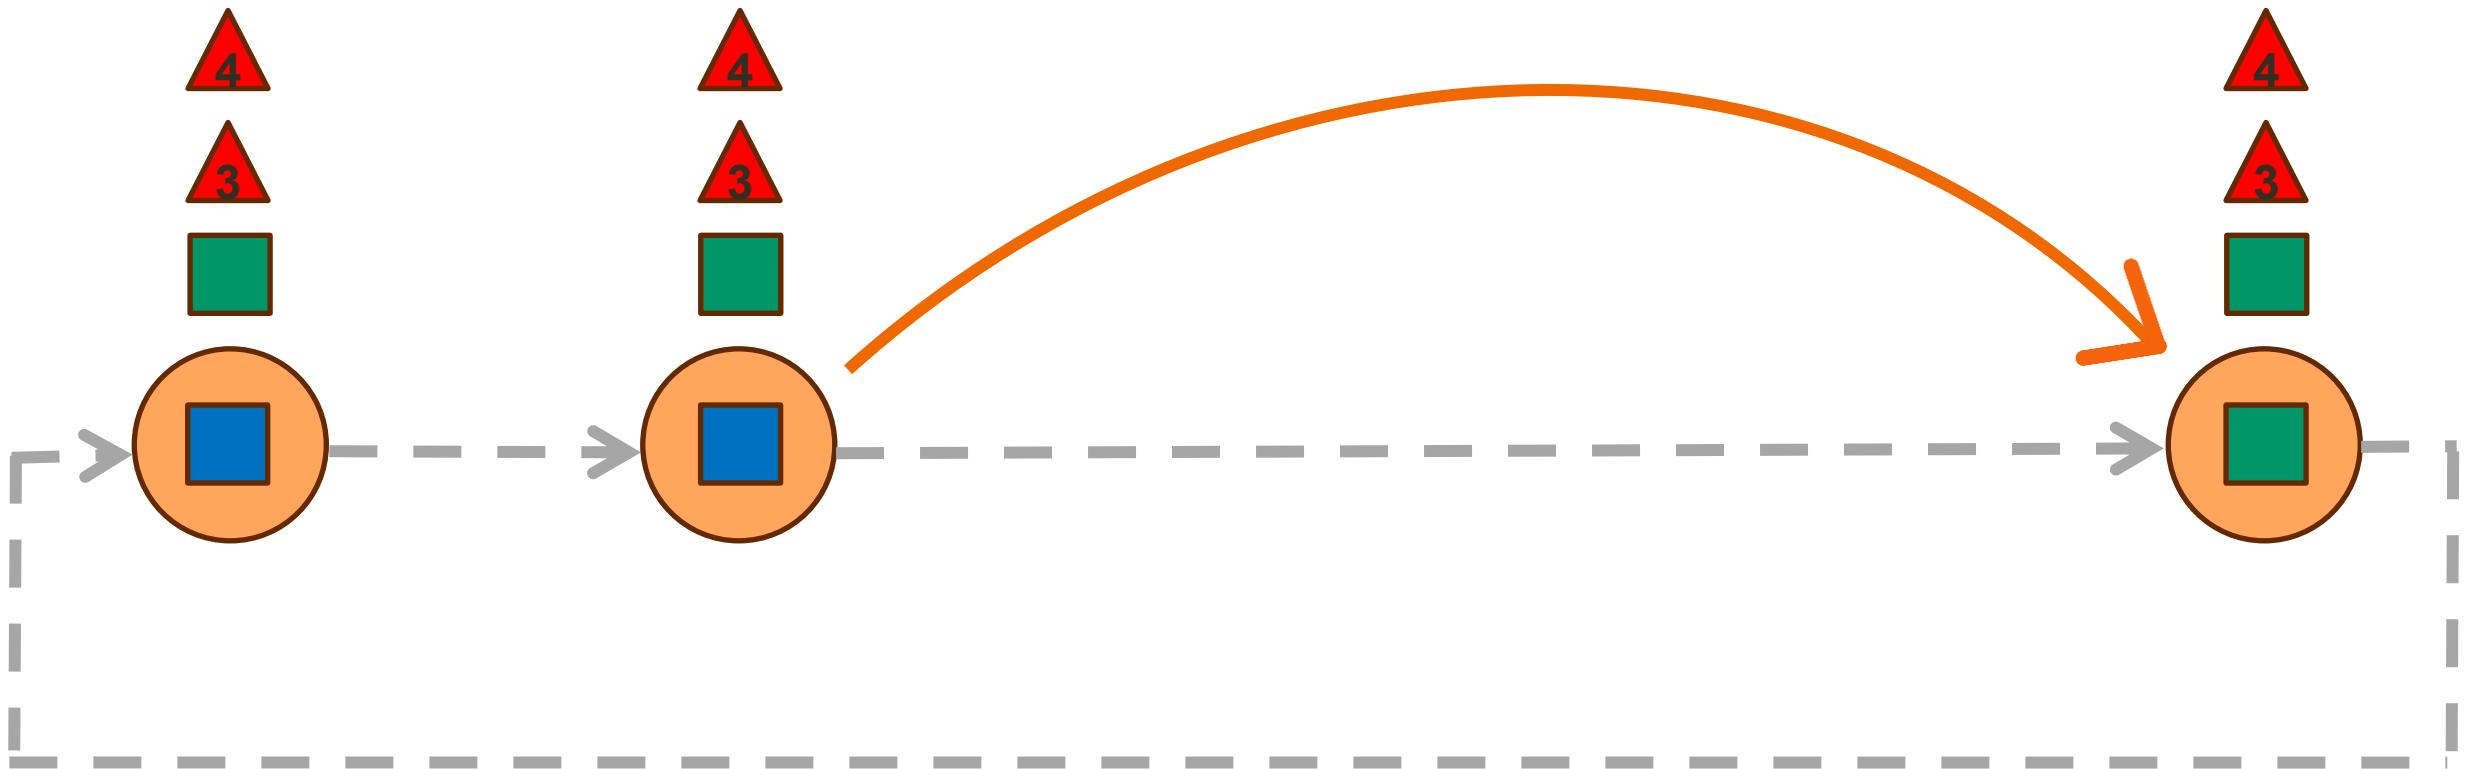
Membership Management: Removal



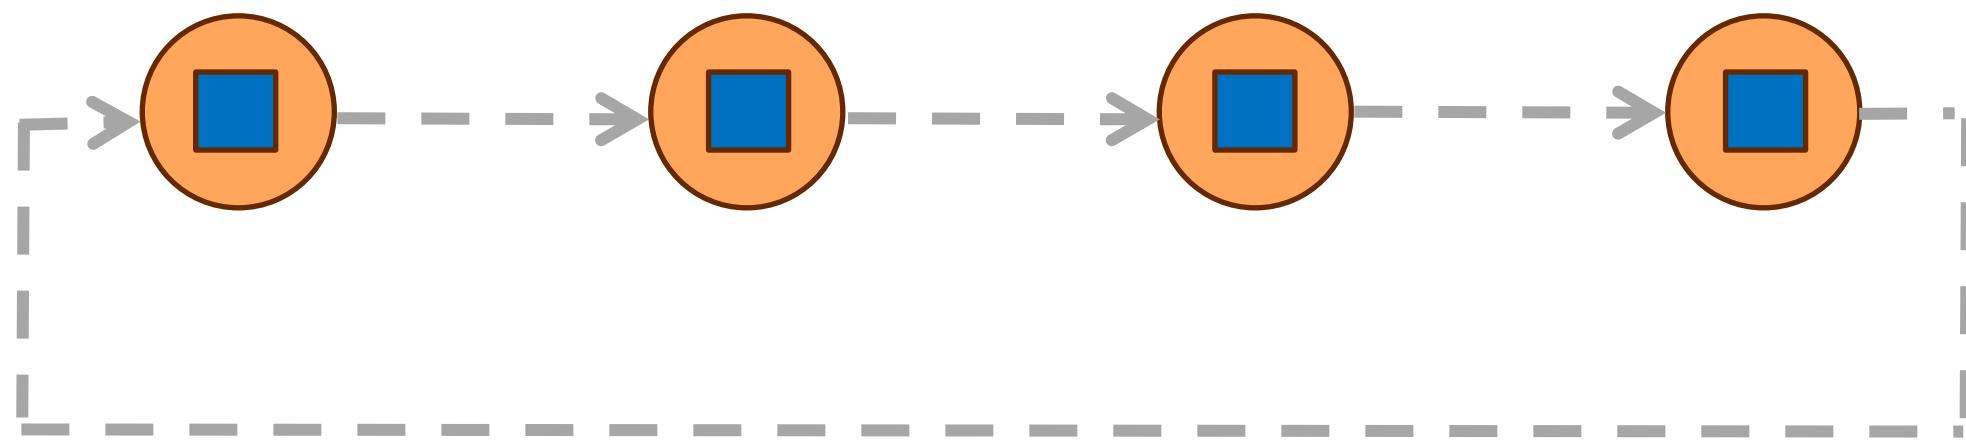
Membership Management: Removal



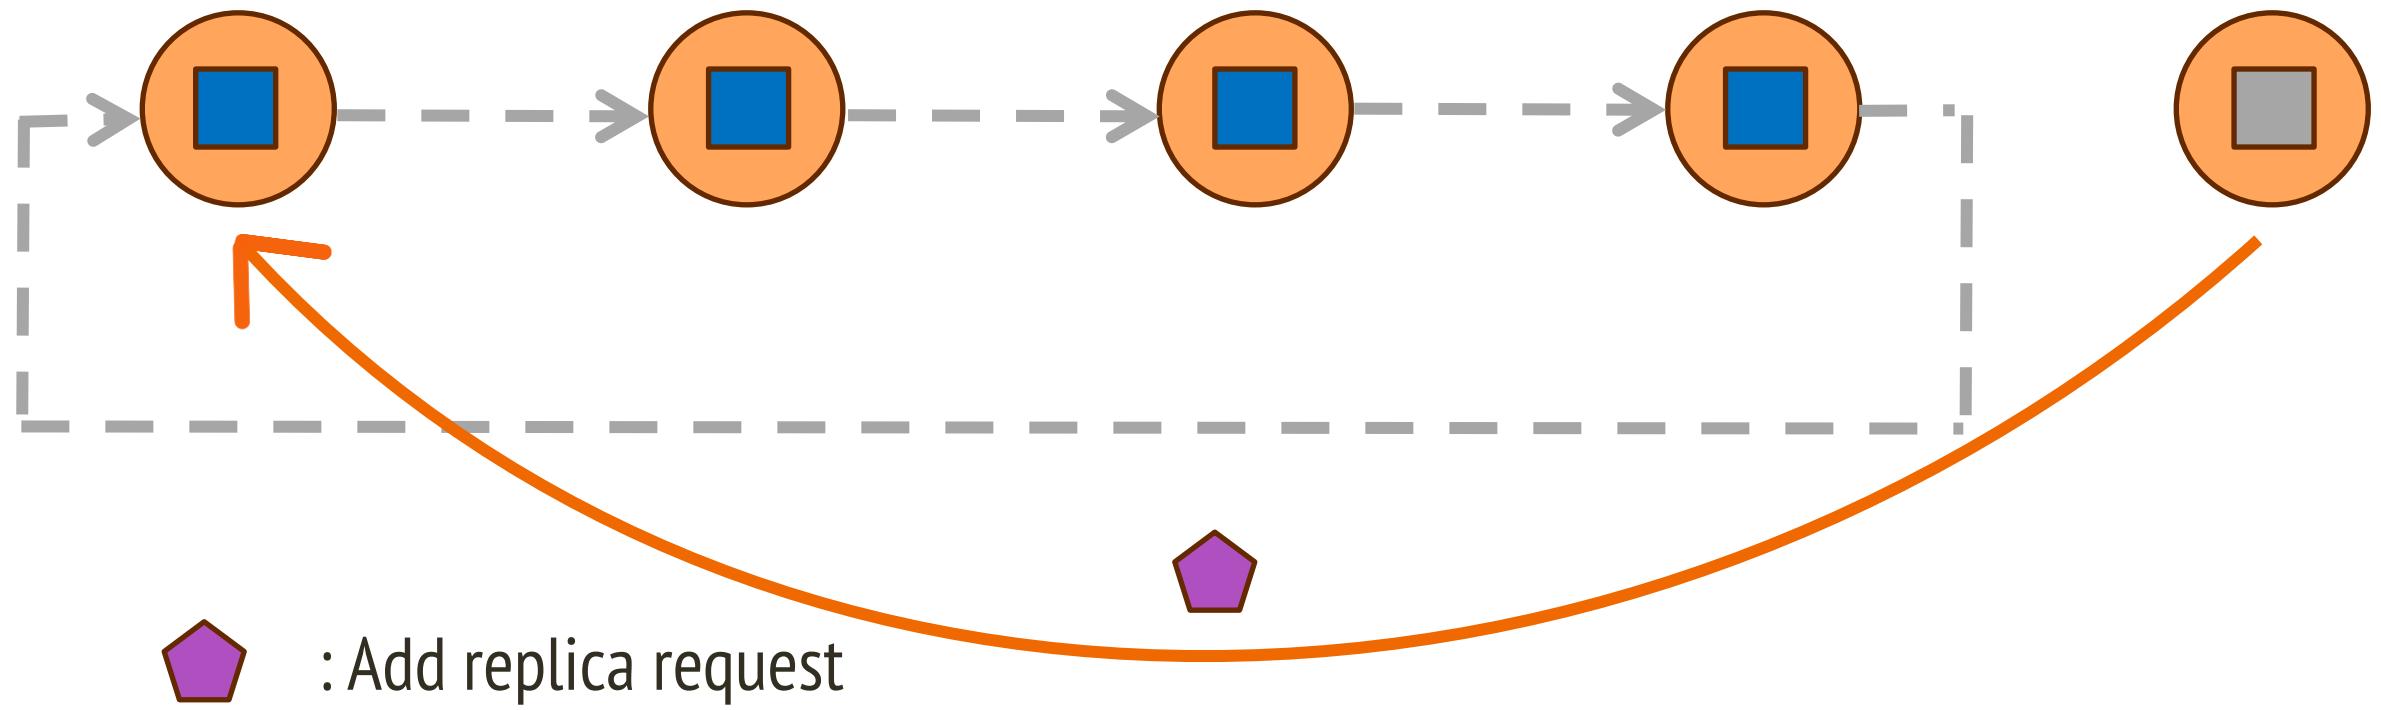
Membership Management: Removal



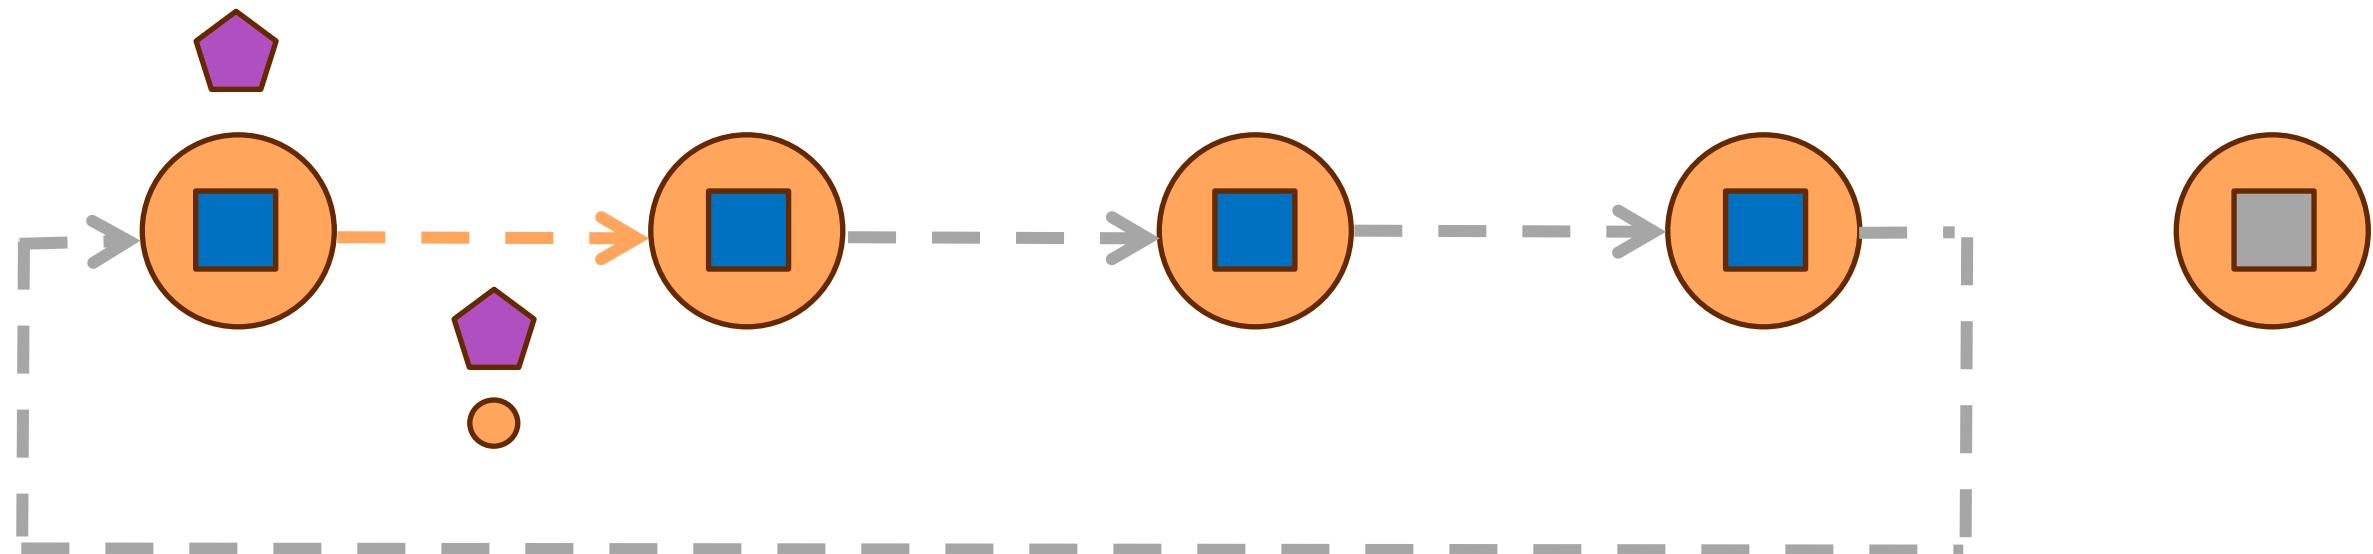
Membership Management: Addition



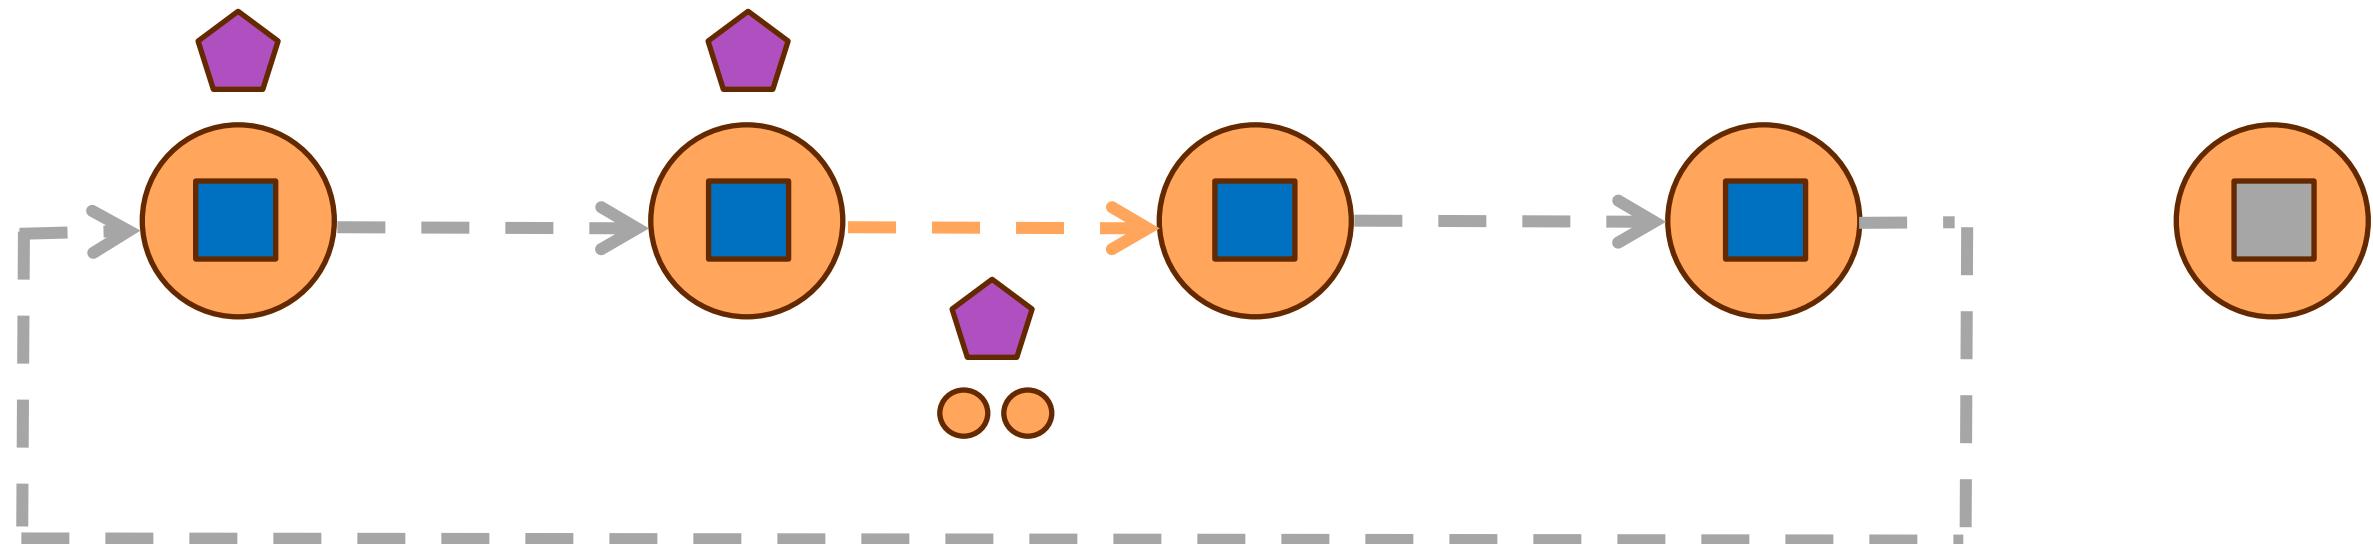
Membership Management: Addition



Membership Management: Addition

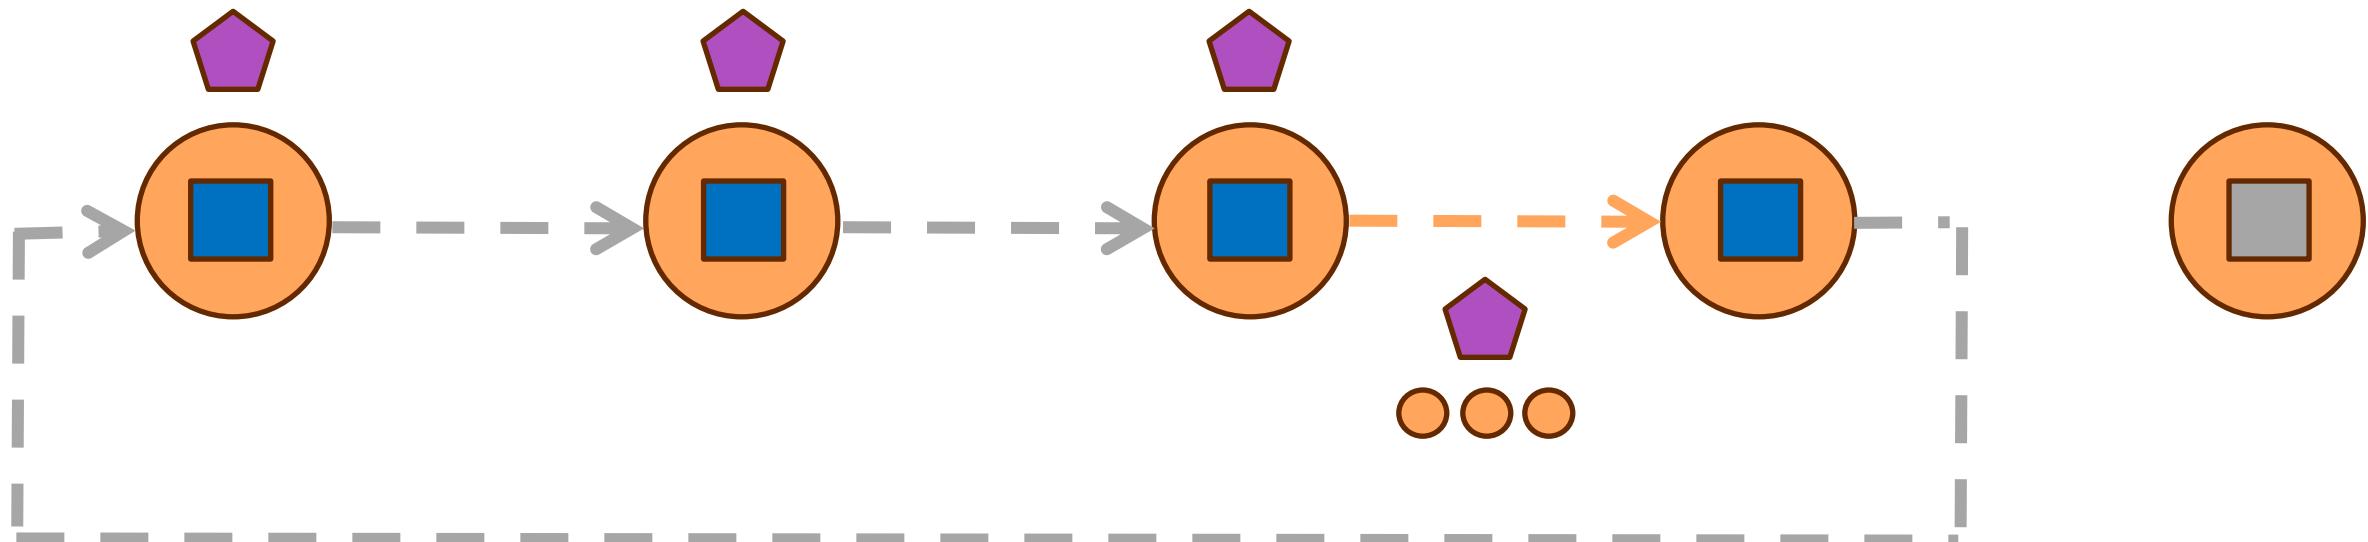


Membership Management: Addition

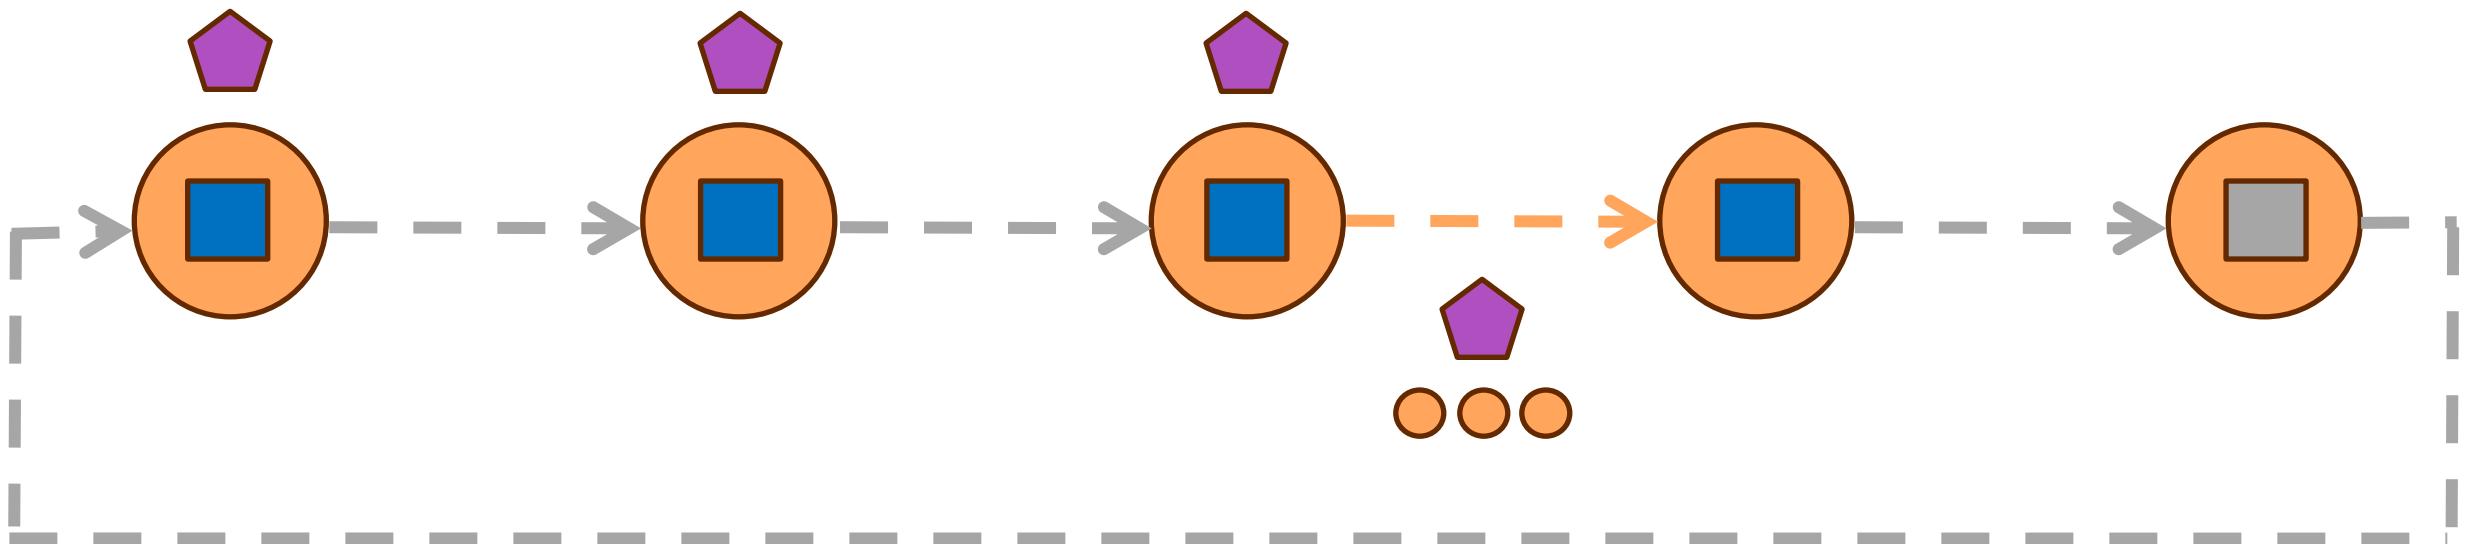


Membership Management: Addition

Addition is now decided

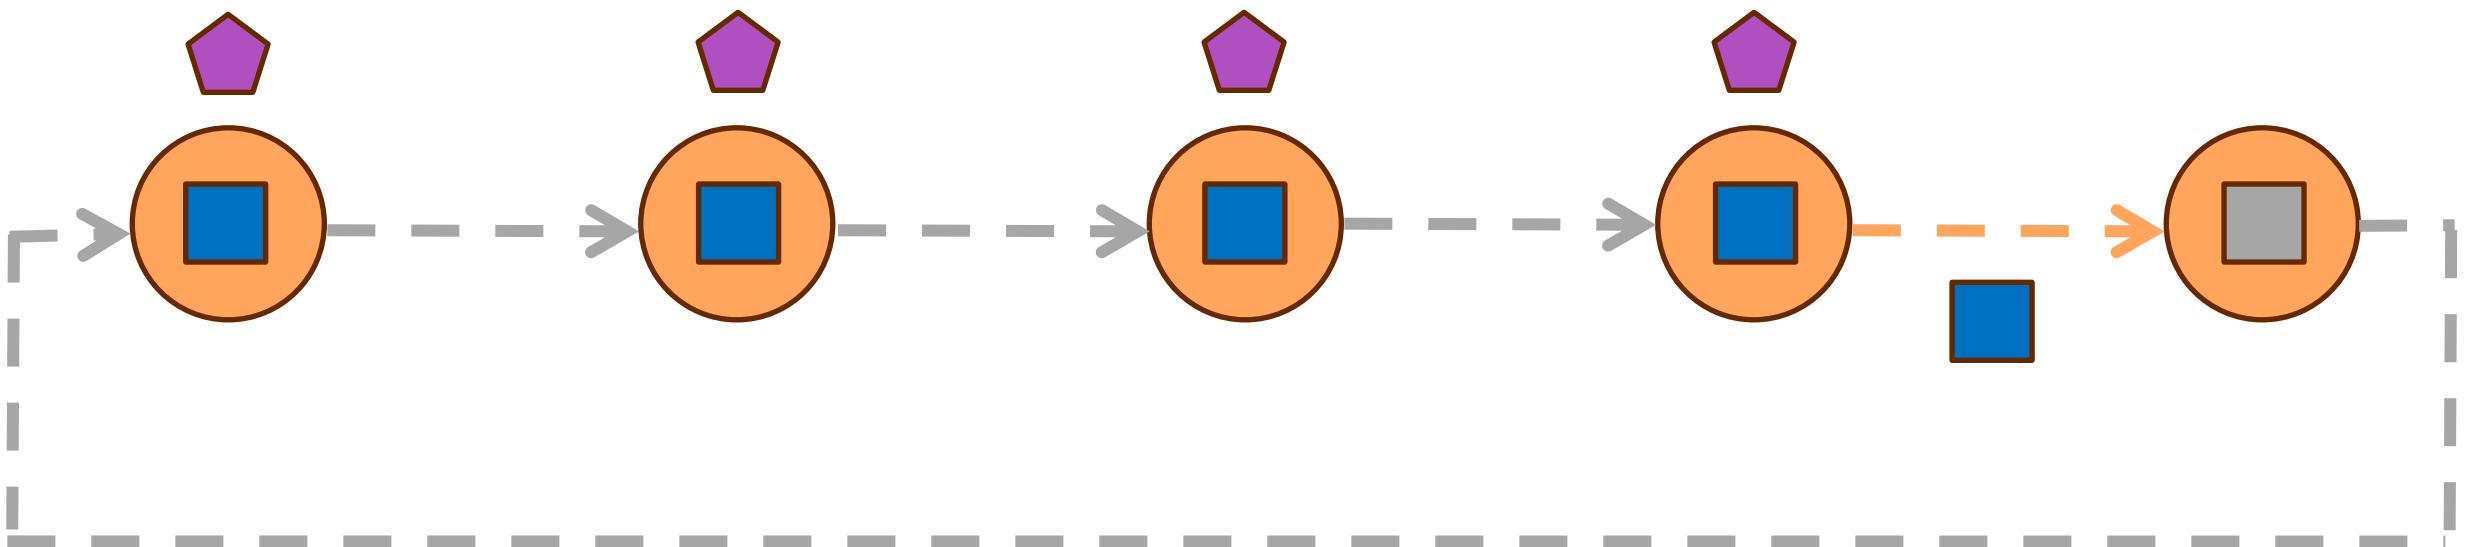


Membership Management: Addition

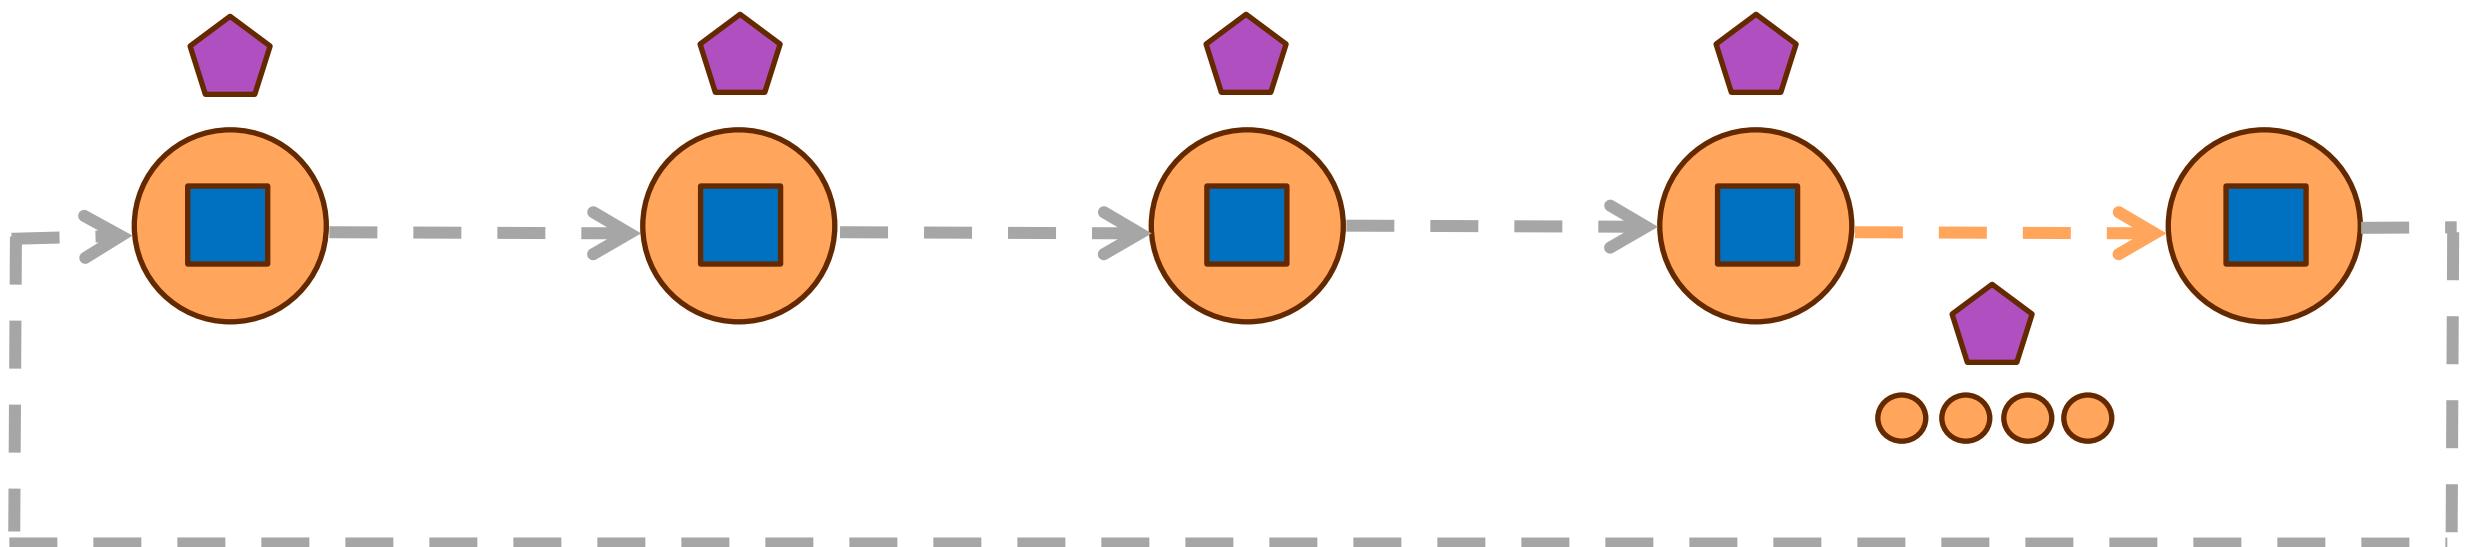


Membership Management: Addition

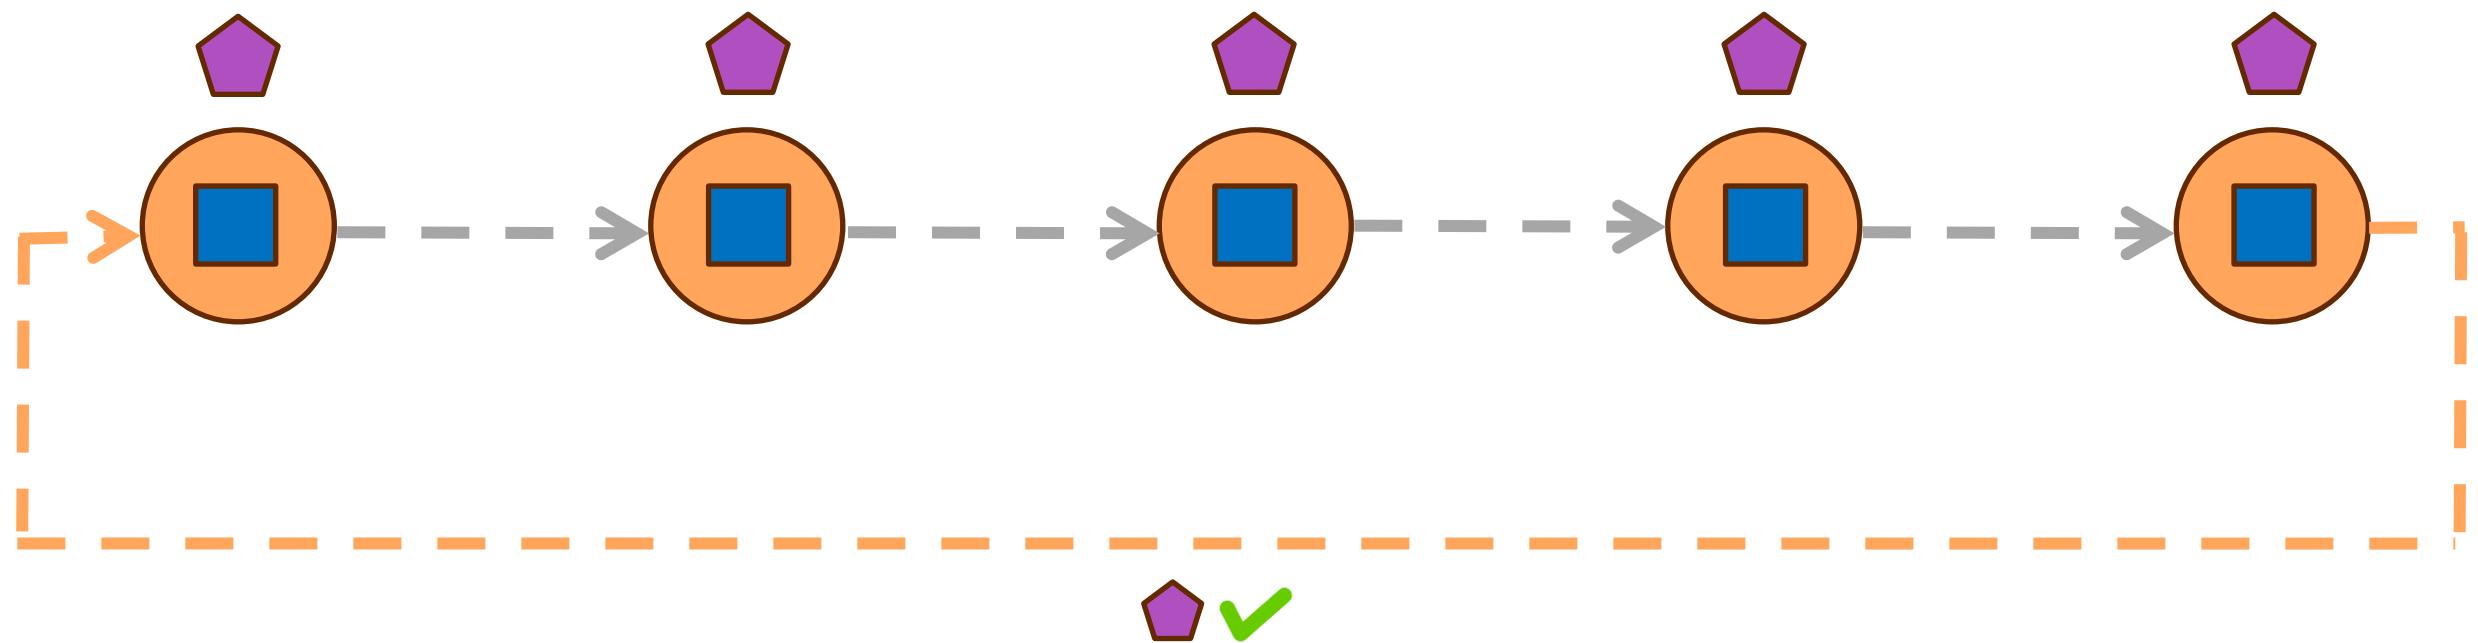
Transfer state



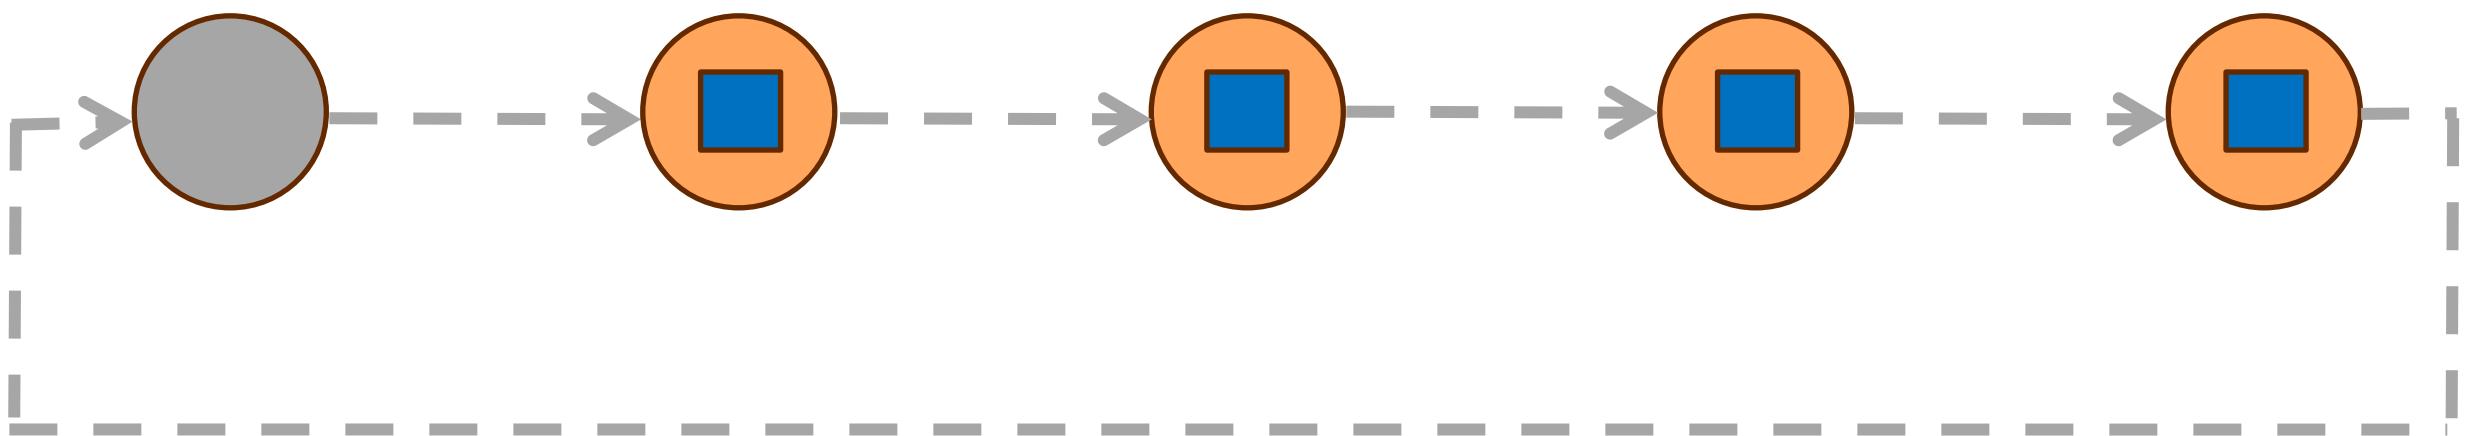
Membership Management: Addition



Membership Management: Addition

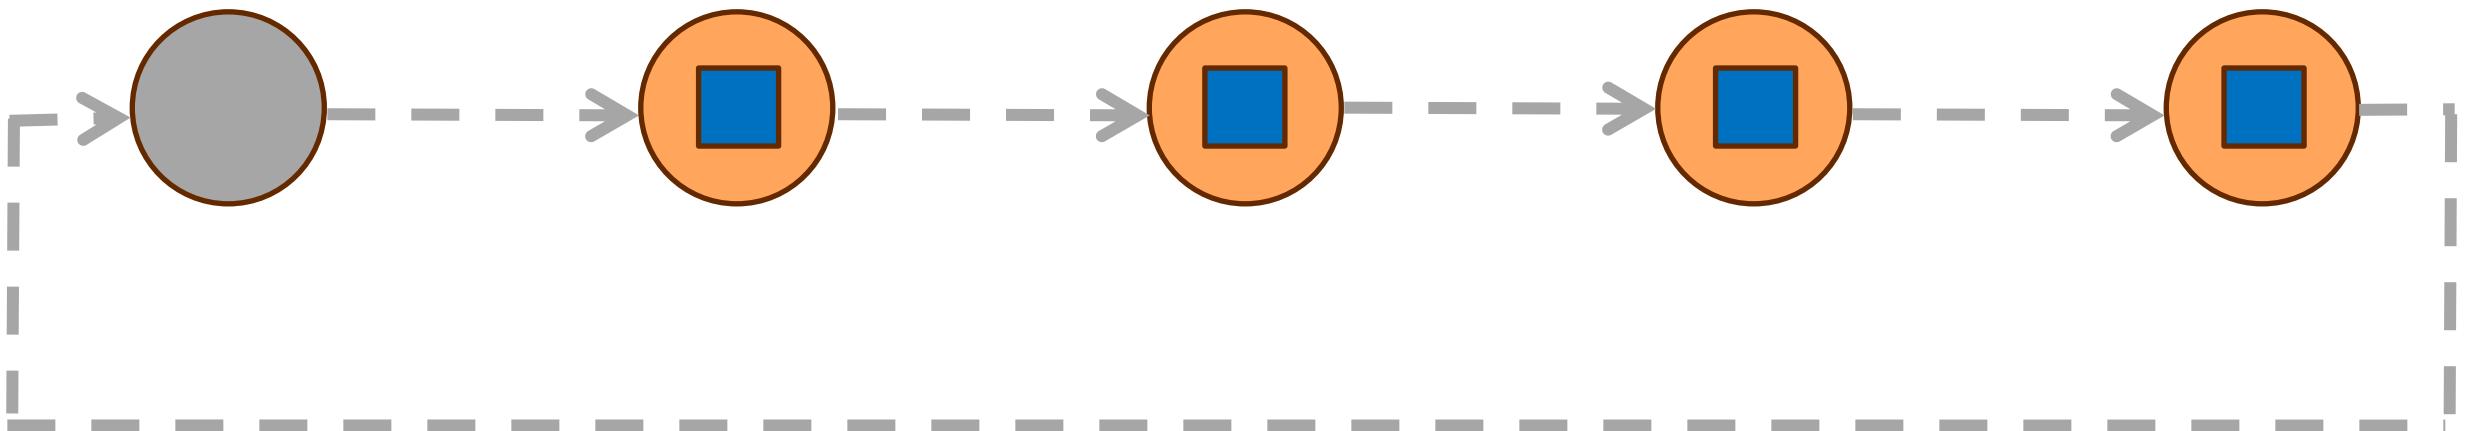


Leader Failure

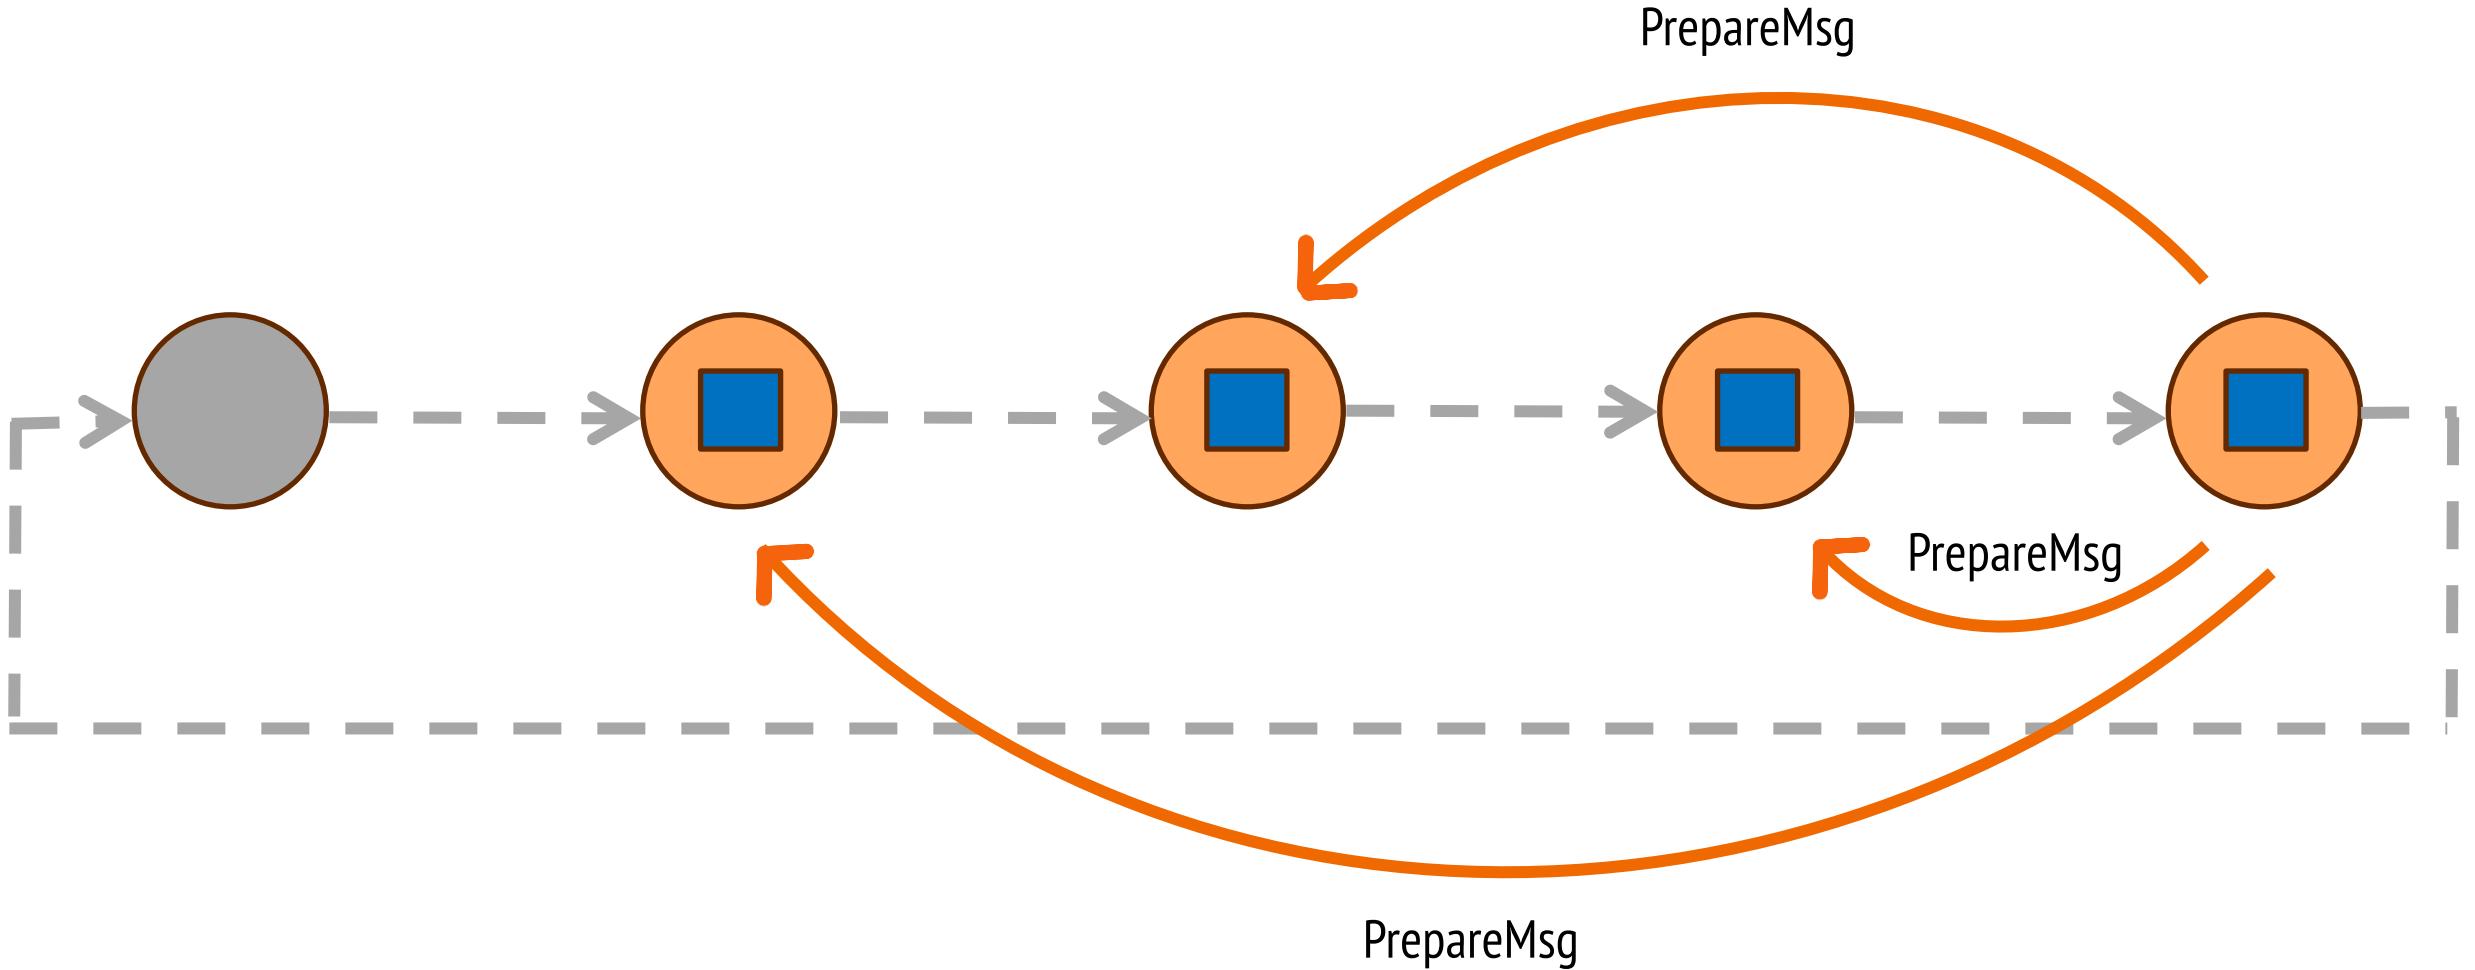


Leader Failure

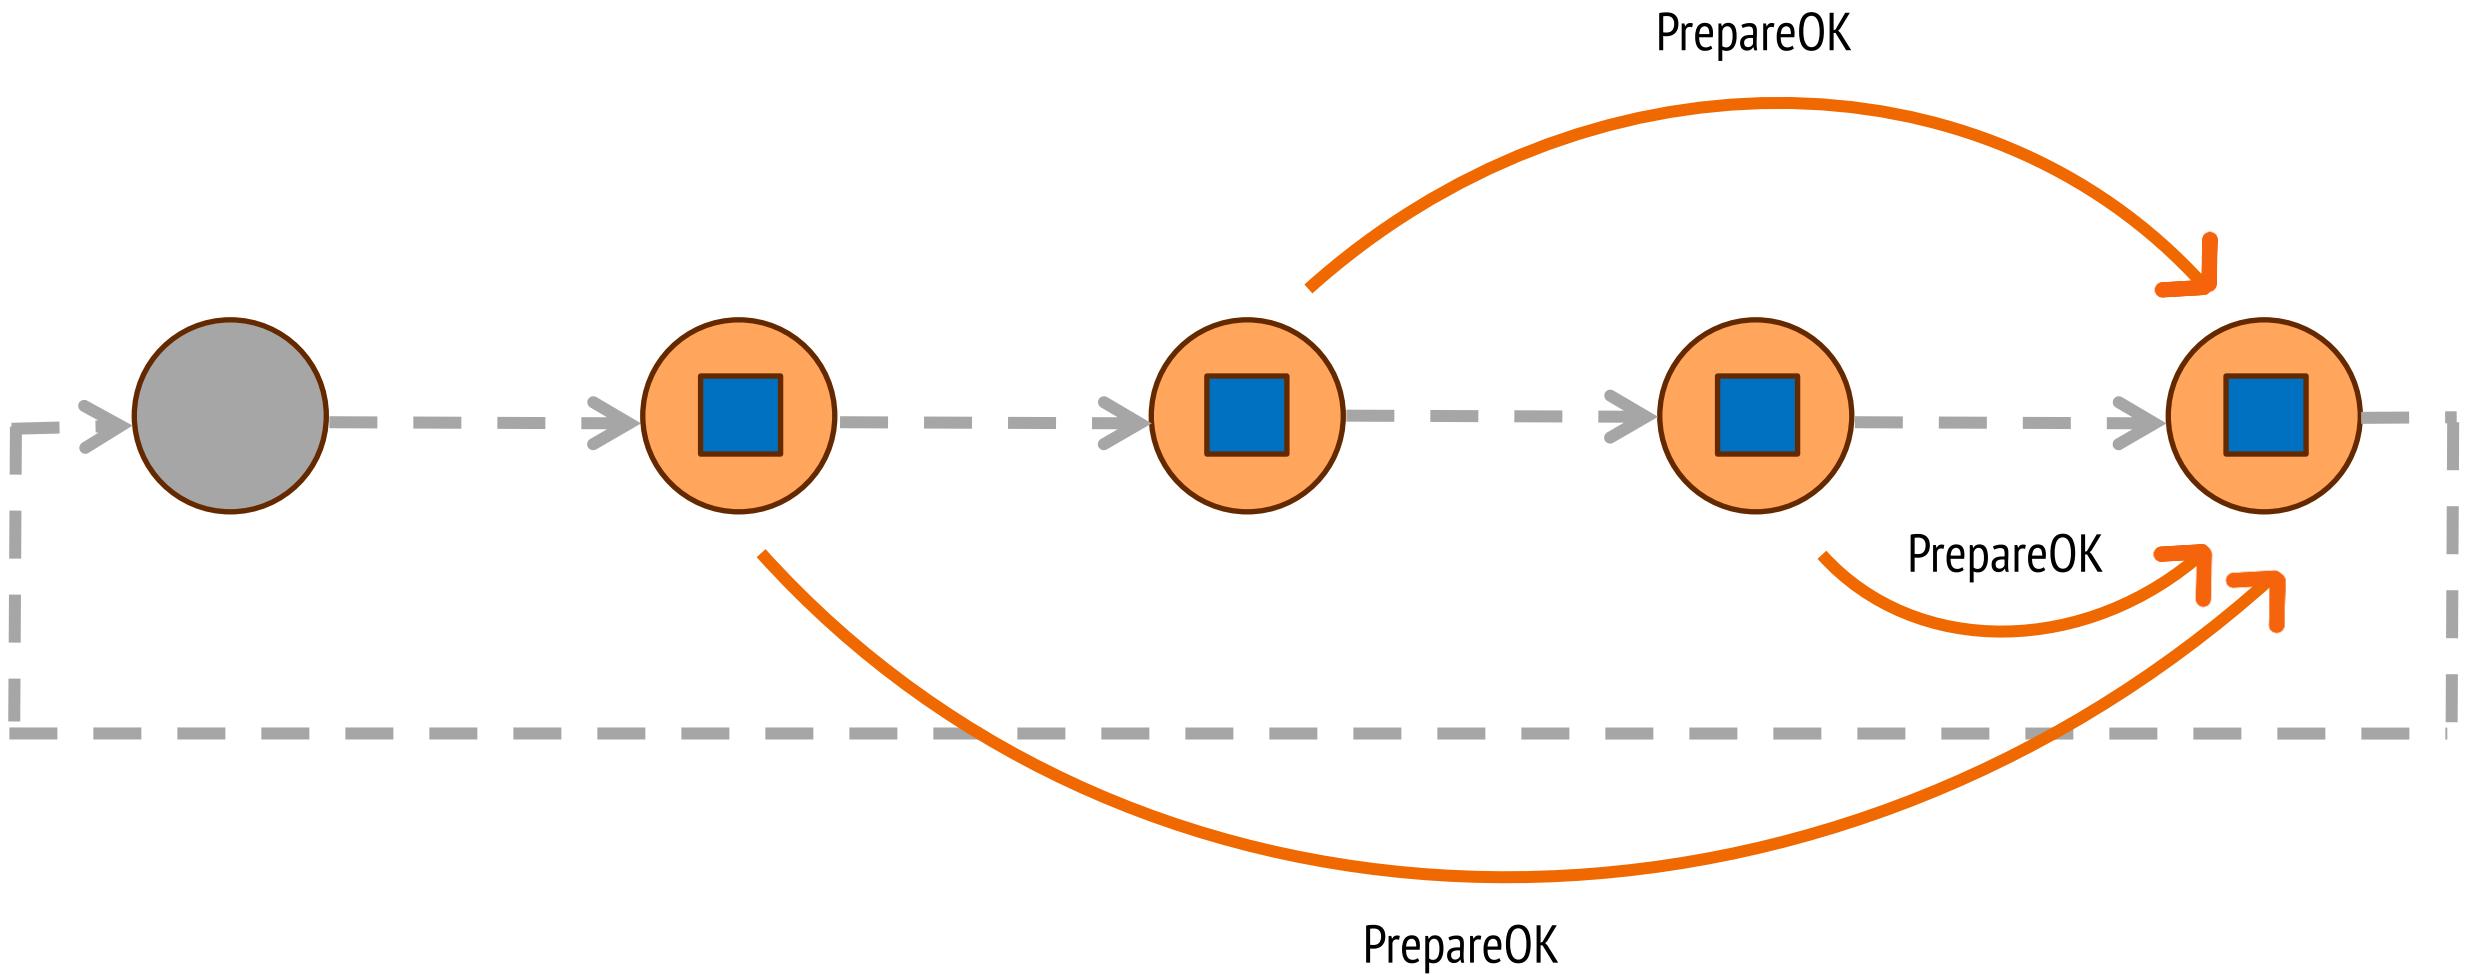
*Run Paxos Leader
Election*



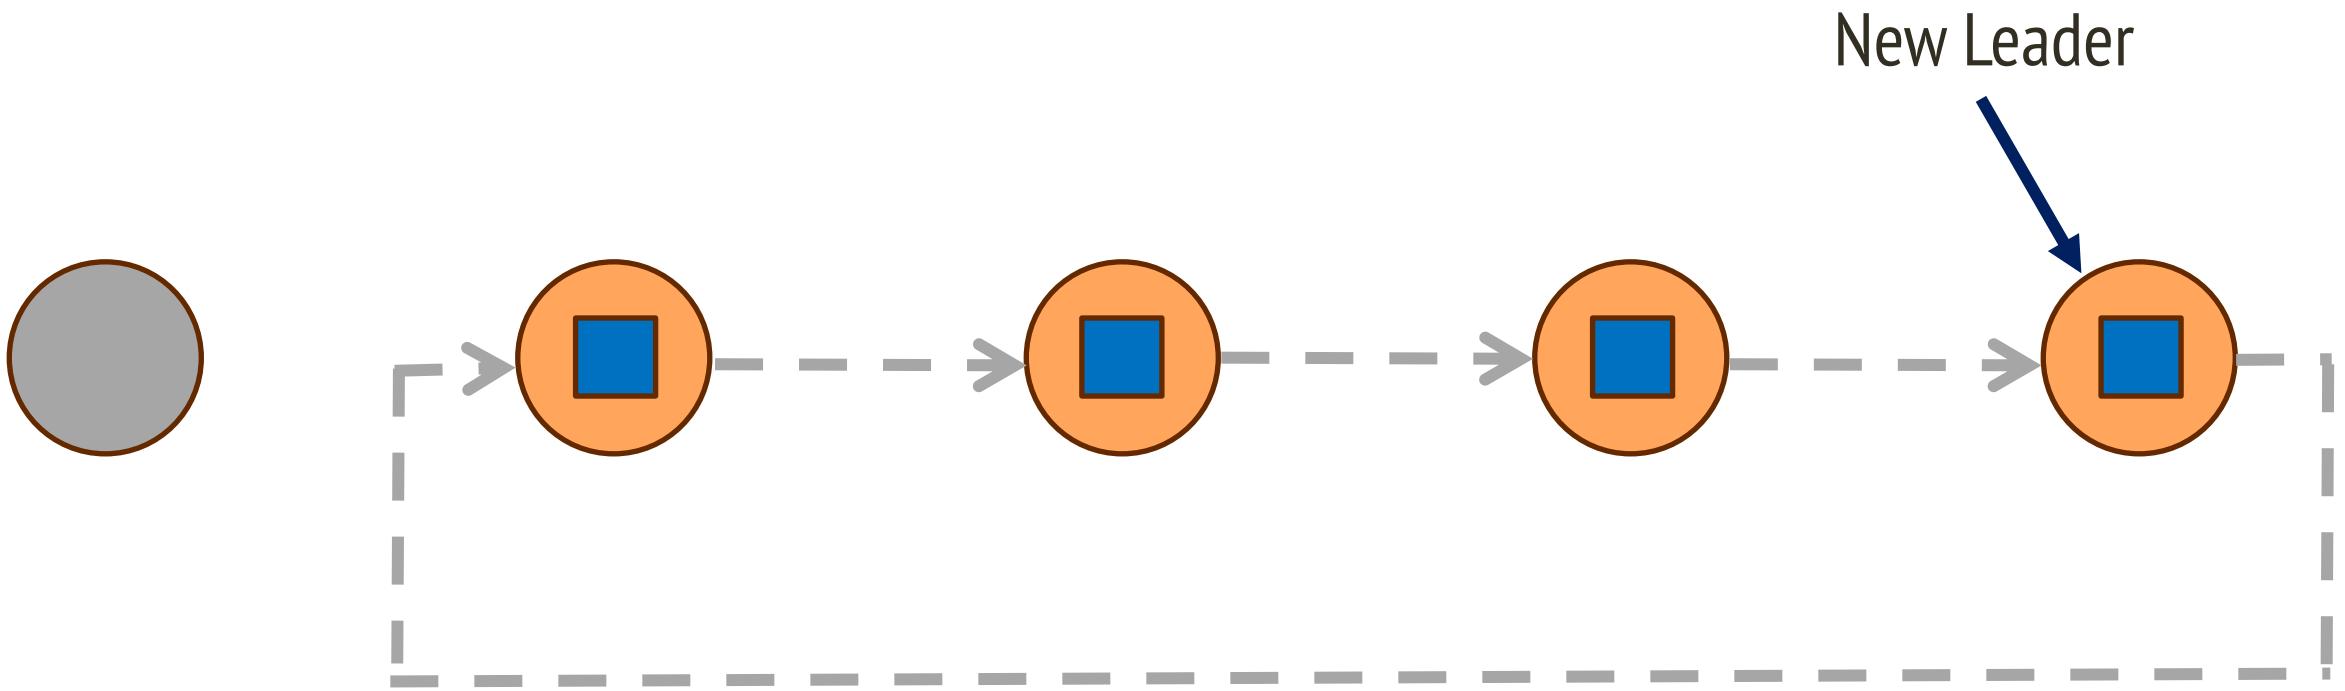
Leader Failure



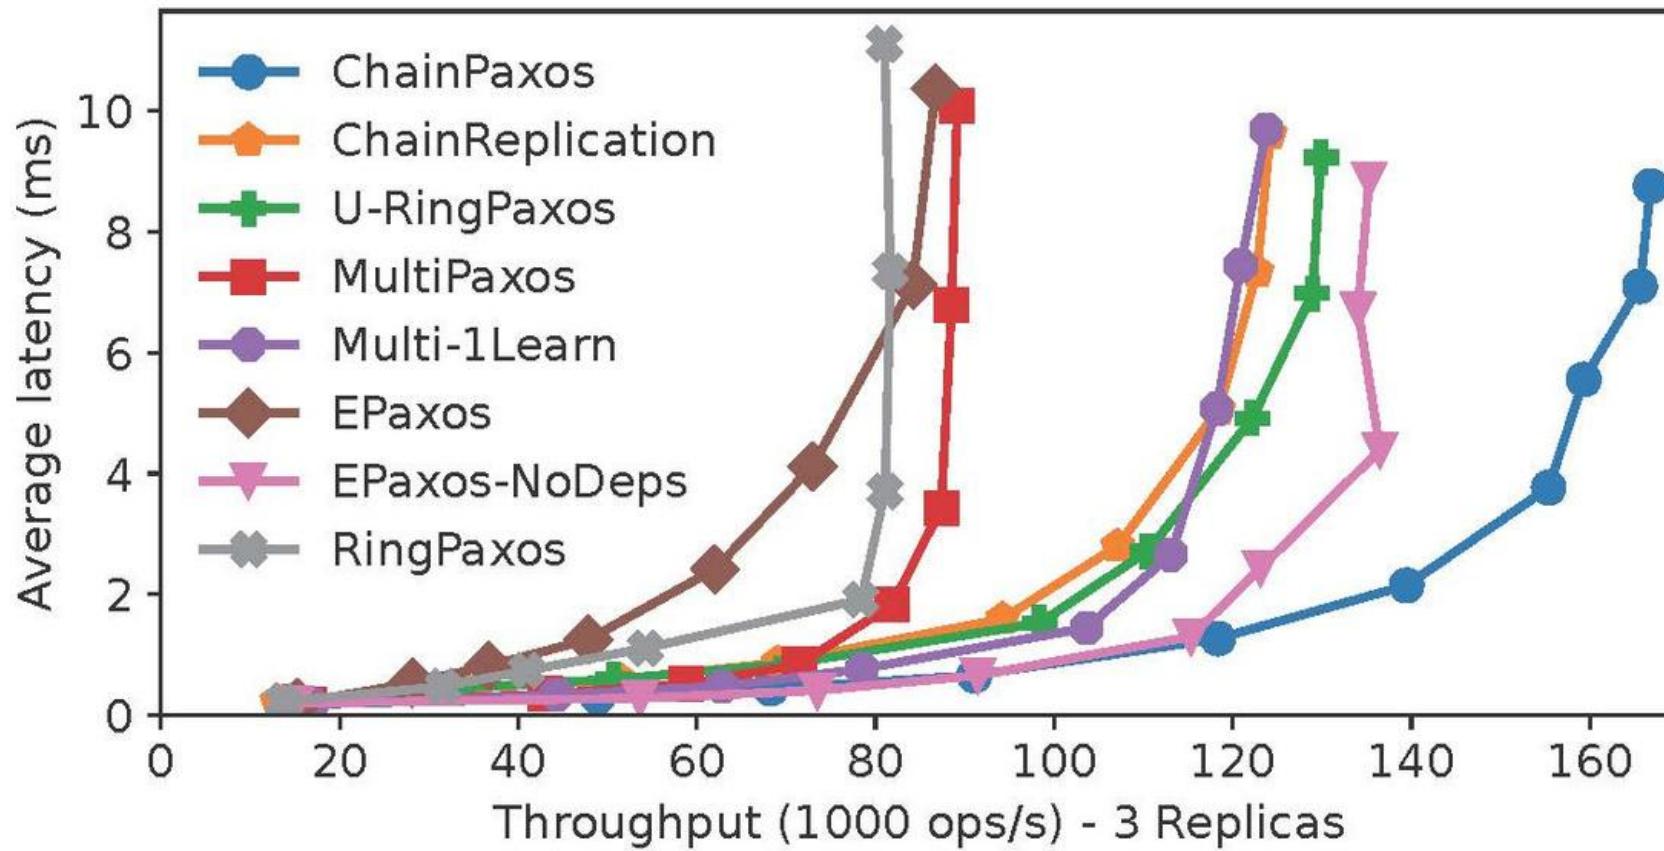
Leader Failure



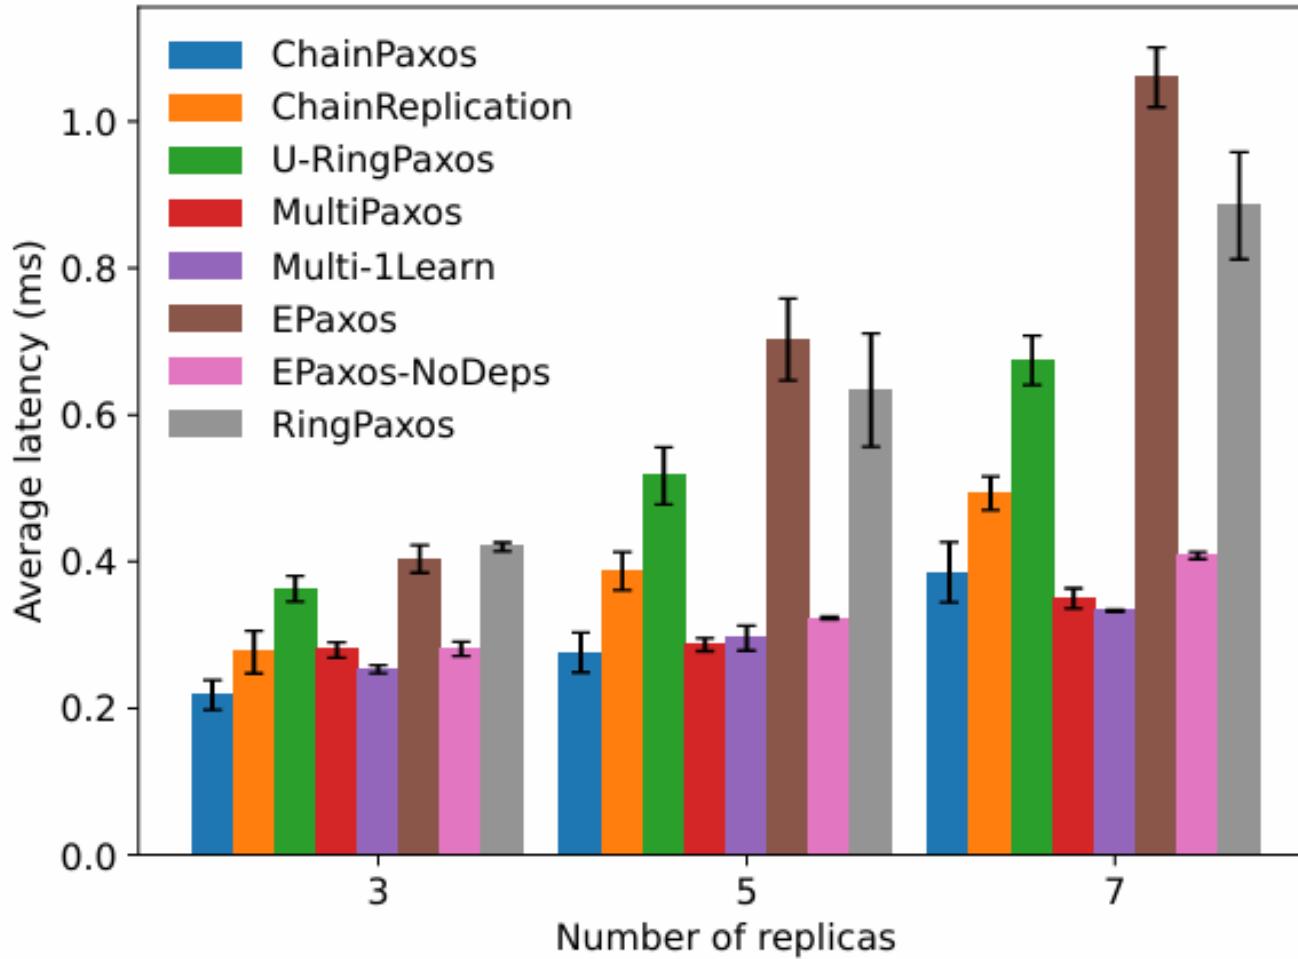
Leader Failure



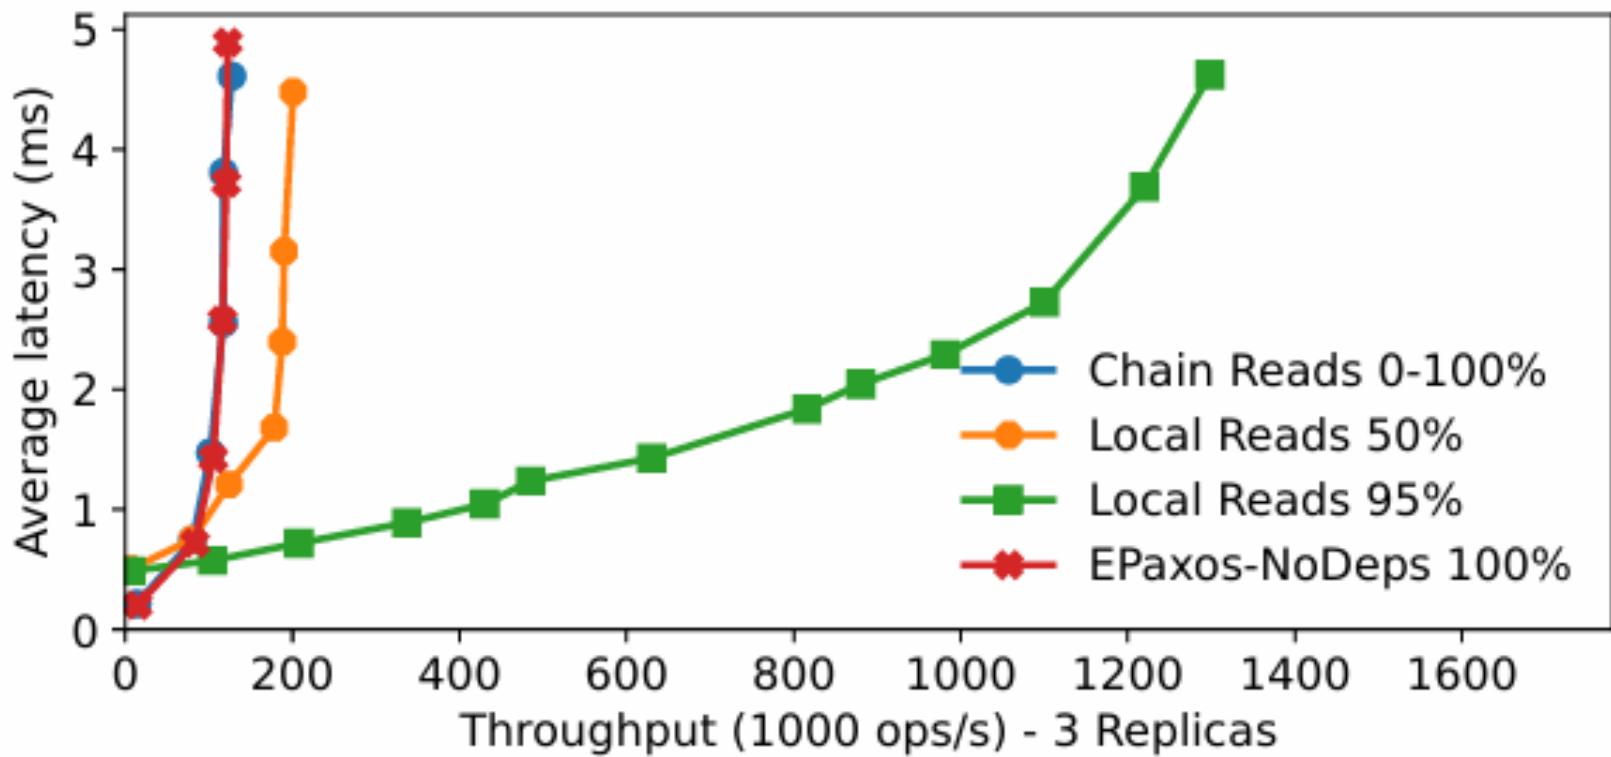
Experiment Results: CPU Bottleneck



Experiment Results: Latency Test



Experiment Results: Local Read Test



References

- Fouto, P., Preguiça, N., & Leitão, J. (2022). High Throughput Replication with Integrated Membership Management. USENIX Annual Technical Conference (USENIX ATC 22), 575–592.

Thank you