

LazyLog: A Shared Log Abstraction for Low-Latency Applications

Xuhao Luo, Shreesha G. Bhat, Jiyu Hu, Ram Alagappan, Aishwarya Ganesan

University of Illinois Urbana-Champaign

ANKITA SINGH
2024MCS2891

ABHIYA M JOSE
2024JCS2046

VADDI JHANCY
2024JCS2040

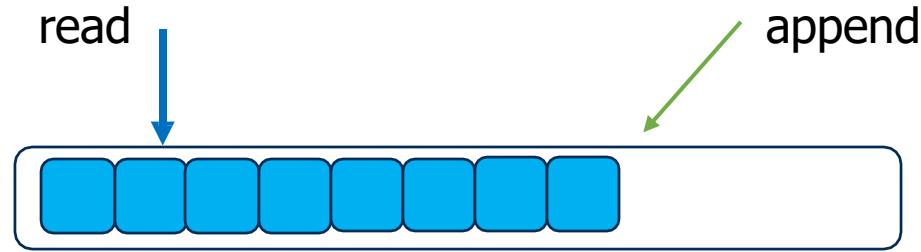
Shared Log: Abstraction and Interface

Abstraction:

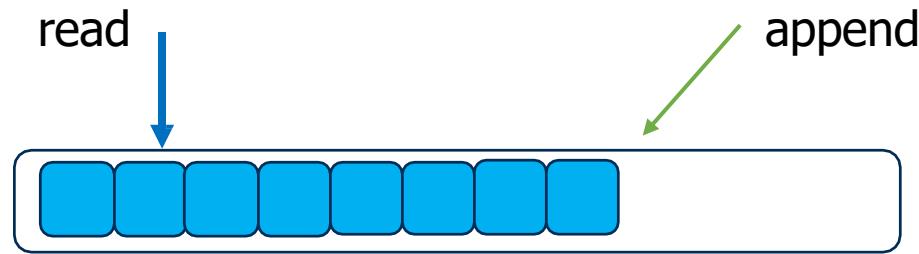
Fault-tolerant, durable,
linearizably ordered sequence of records

Interface:

append
read



Shared Log: Abstraction and Interface



```
// append to log; return log position  
uint64_t append(record r);  
  
// read 'len' records starting at 'from'  
list read(logpos_t from, uint64_t len);
```

The Problem with Current Shared Logs

- high append latencies
 - Append takes multiple RTTs
- Low append latency is critical to applications
- Eager ordering nature of shared logs:

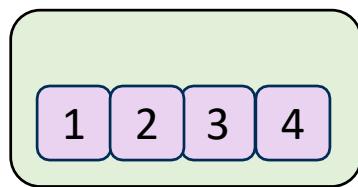
The Problem with Current Shared Logs

- **high ingestion latencies**
 - Append takes multiple RTTs
- Low ingestion latency is critical to applications
- **Eager ordering** nature of shared logs:
 - Order is established eagerly upon appends
 - Position of record is decided by the time append completes

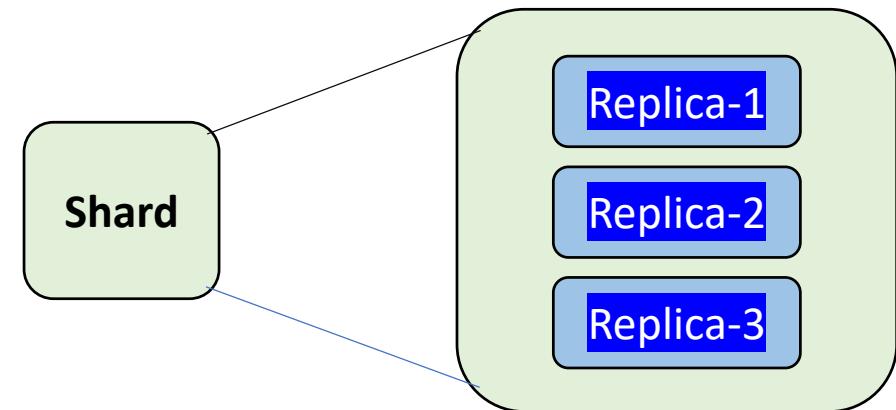
Total Ordering in Shared Logs

Linearizable order:

if append(B) starts after append(A) completes, then B appears after A in the shared log



Shard-1

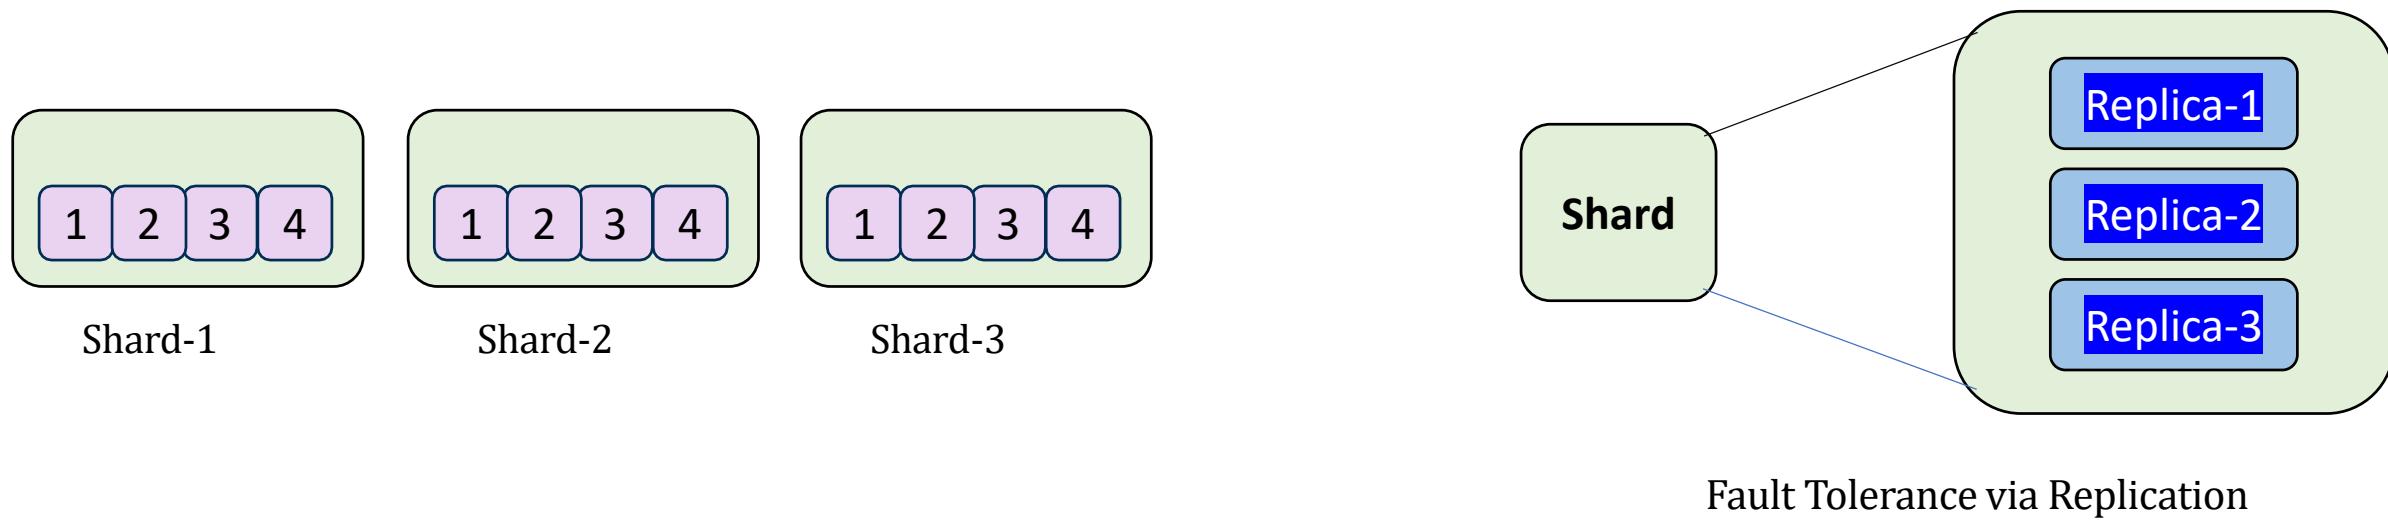


Fault Tolerance via Replication

Total Ordering in Shared Logs

Linearizable order:

if append(B) starts after append(A) completes, then B appears after A in the shared log

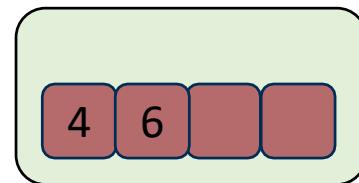
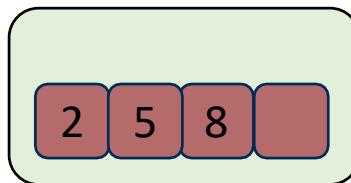
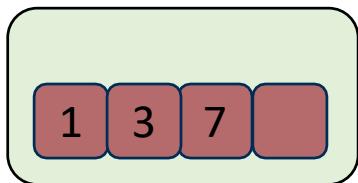
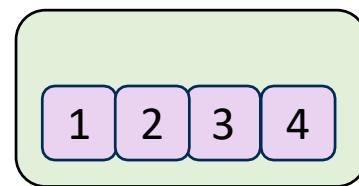
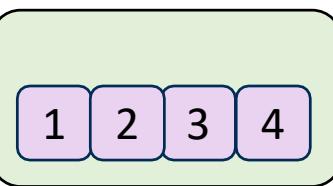
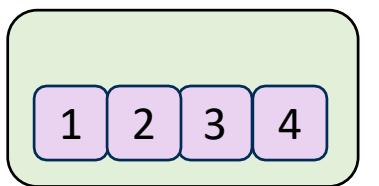


Requires Total Ordering across the shards

Total Ordering in Shared Logs

Linearizable order:

if append(B) starts after append(A) completes, then B appears after A in the shared log

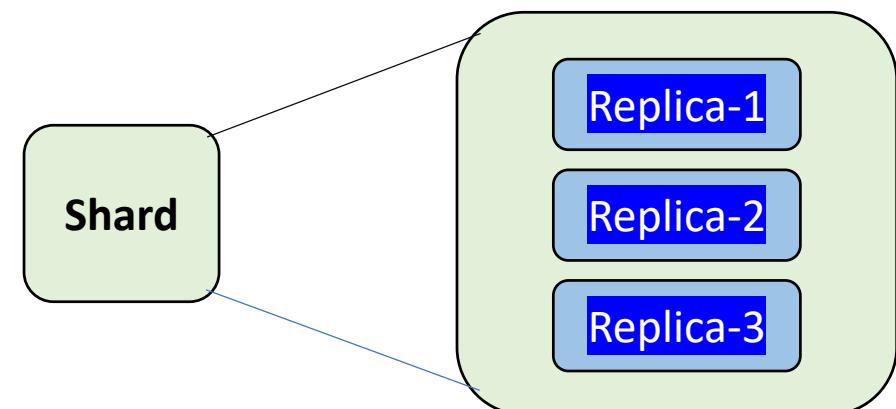


Shard-1

Shard-2

Shard-3

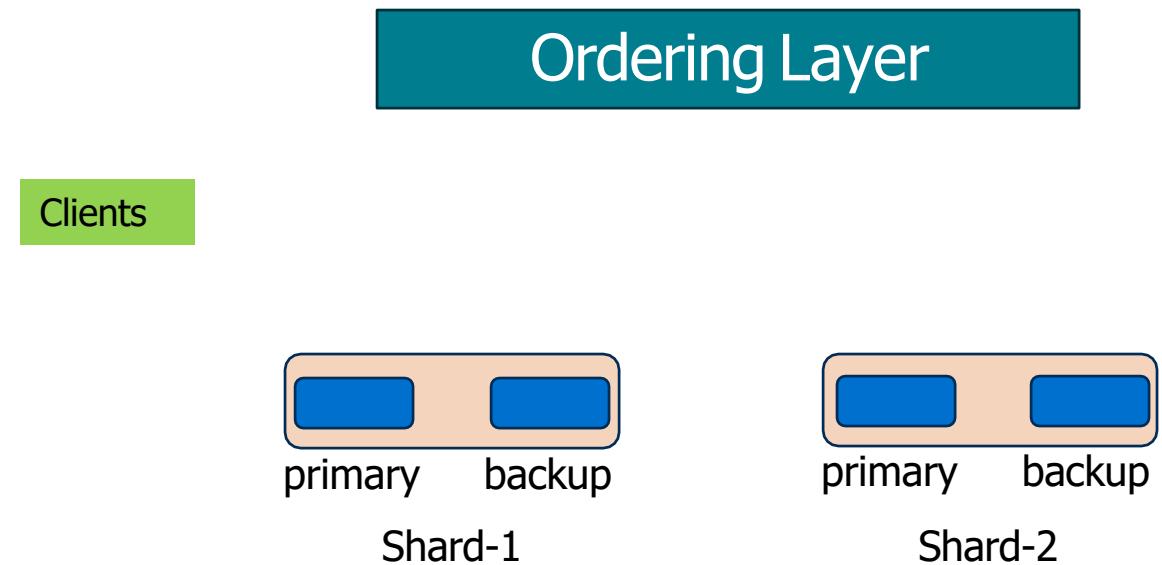
Total Order across the shards



Fault Tolerance via Replication

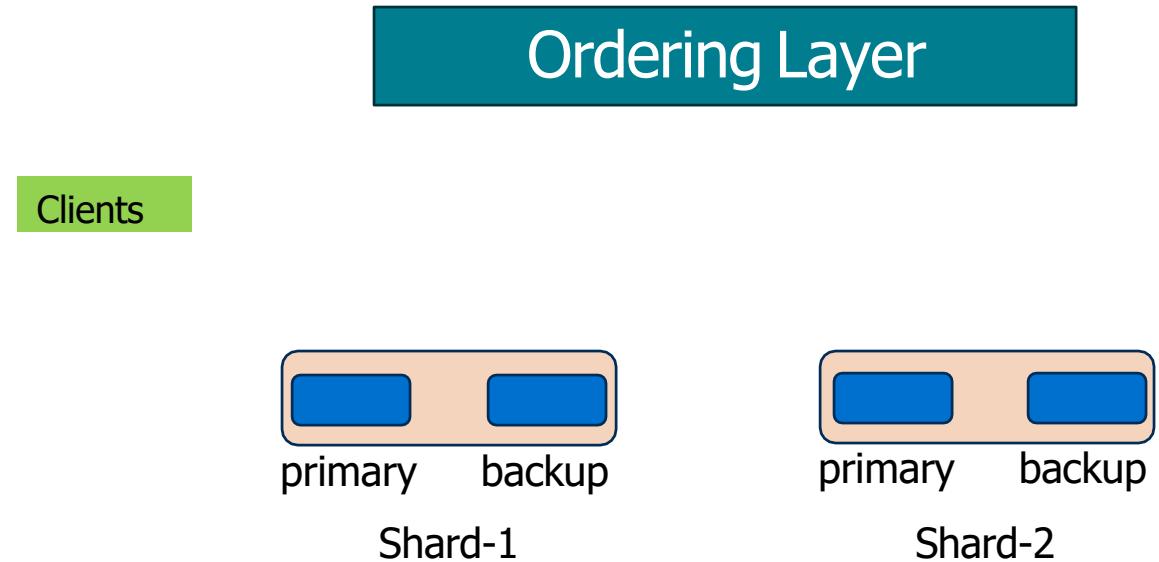
Scalability with total order: Scalog & Corfu

Eager Ordering → High Latency



Eager Ordering → High Latency

- Incur high append latency

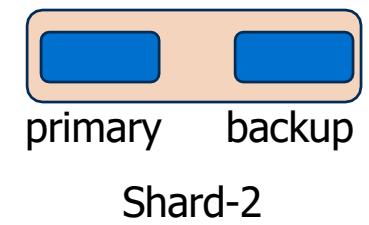
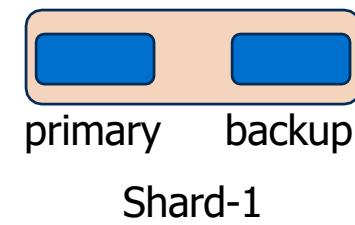


Eager Ordering → High Latency

- Incur high append latency
- Rooted in **eager ordering**

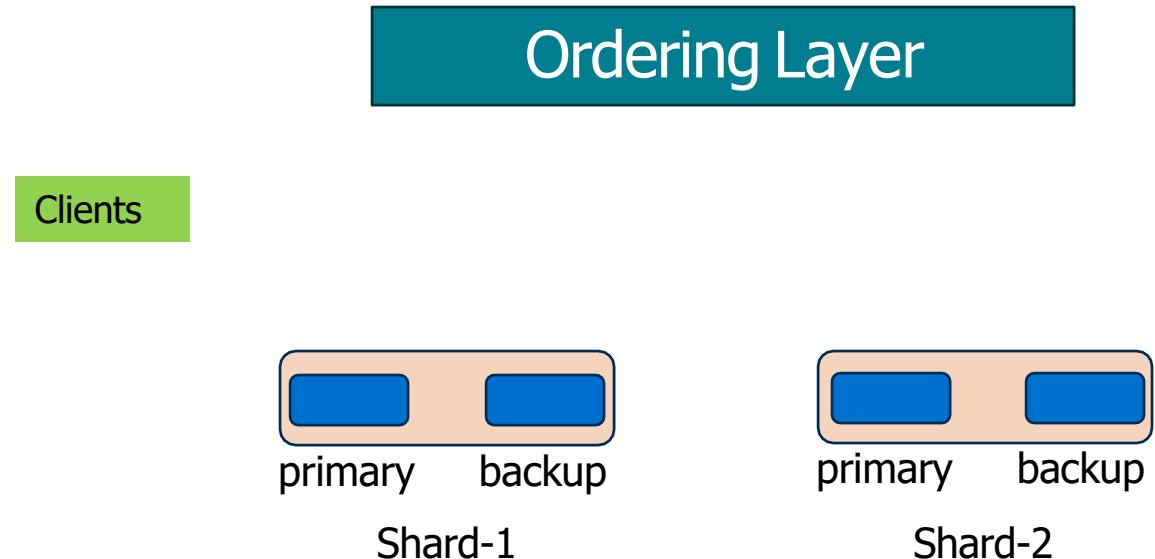
Ordering Layer

Clients



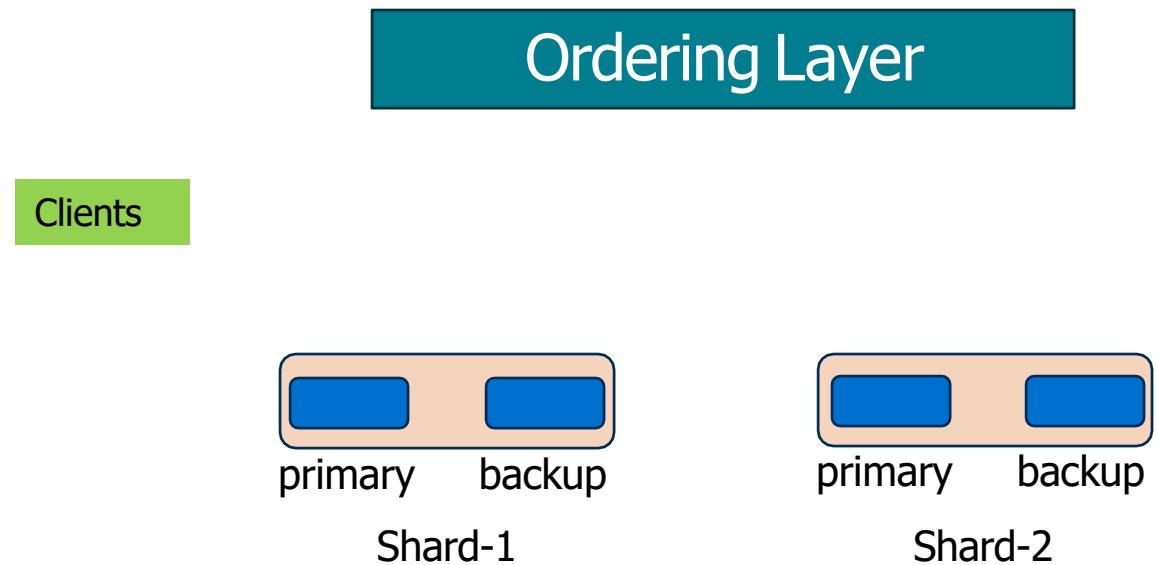
Eager Ordering → High Latency

- Incur high append latency
- Rooted in **eager ordering**
- Both durability and global ordering are completed before getting back to clients



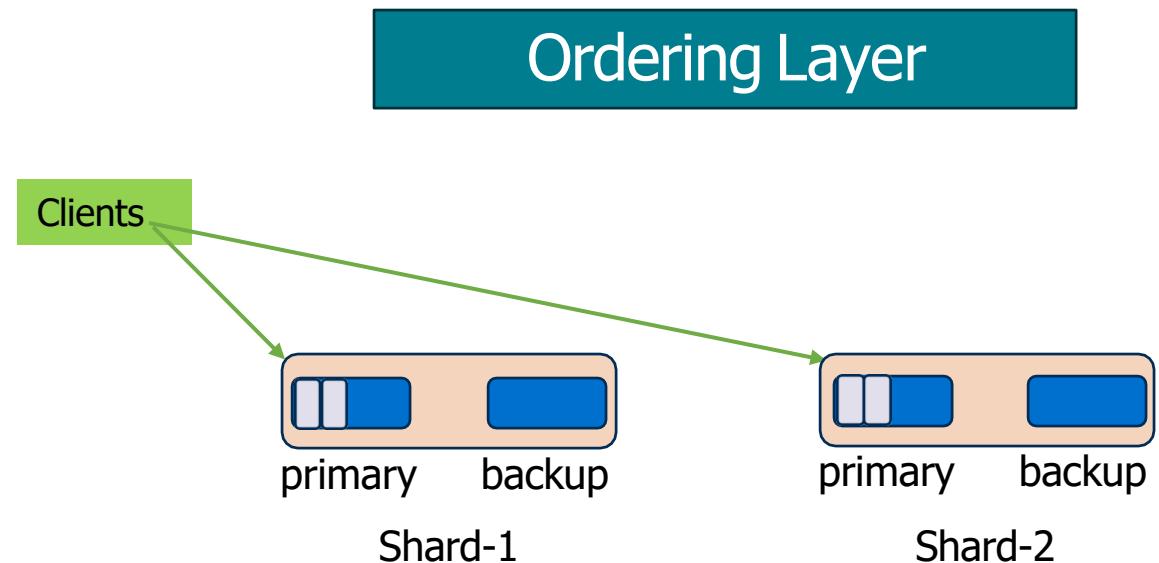
Eager Ordering → High Latency

- **Scalog**



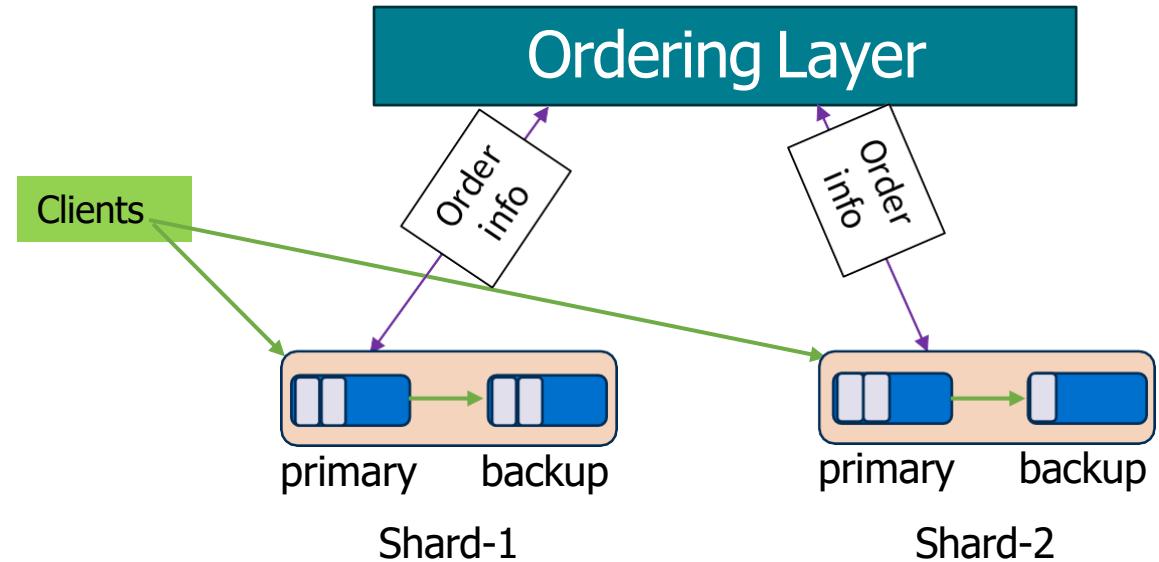
Eager Ordering → High Latency

- Scalog
 - durability first



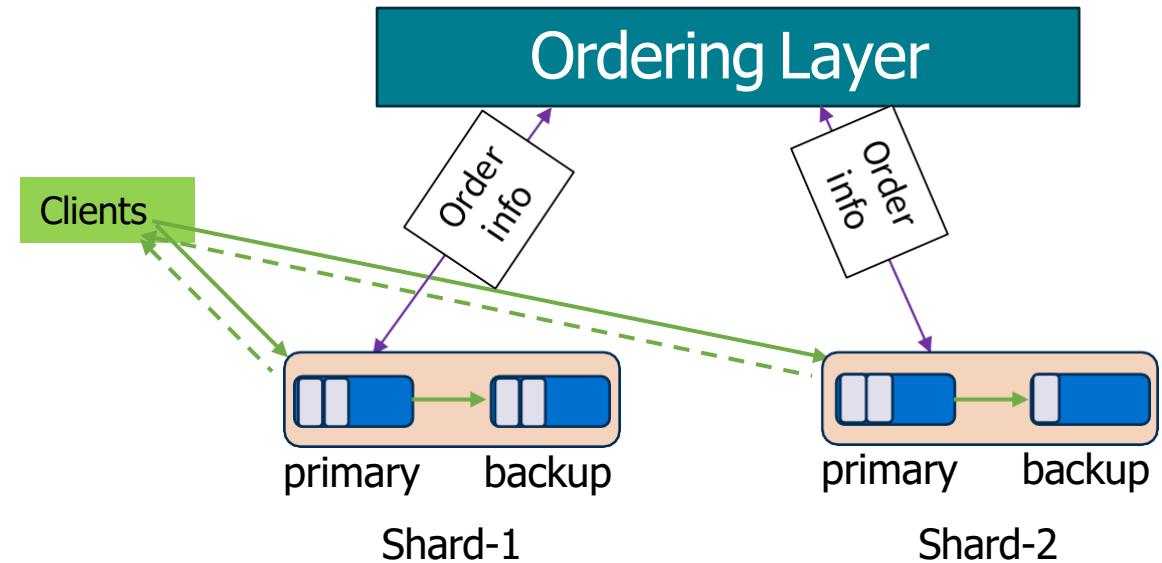
Eager Ordering → High Latency

- **Scalog**
 - durability first
 - then global ordering



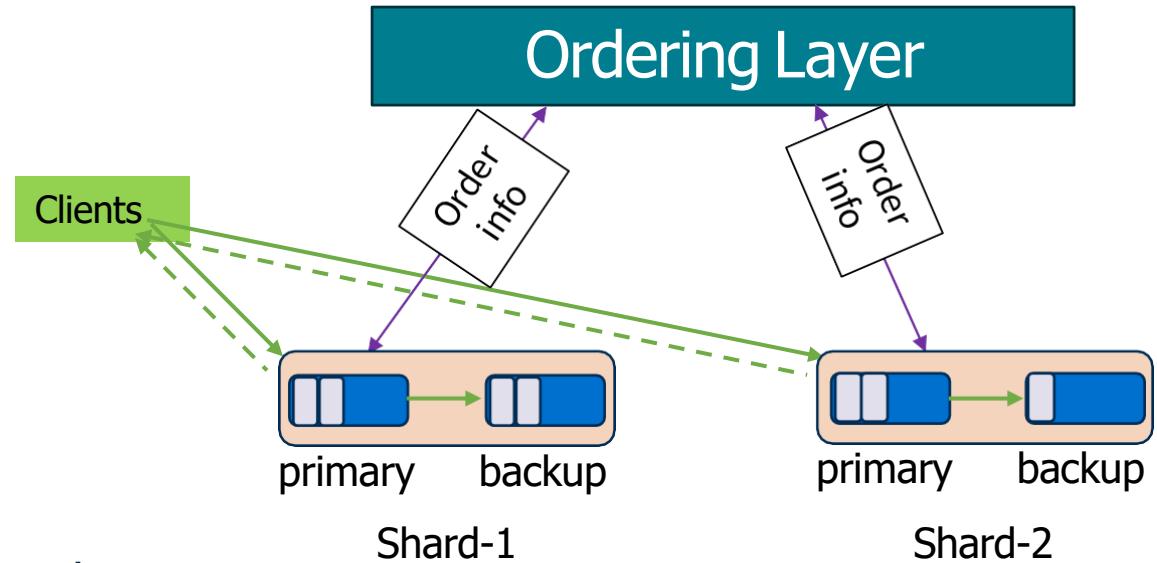
Eager Ordering → High Latency

- **Scalog**
 - durability first
 - then global ordering



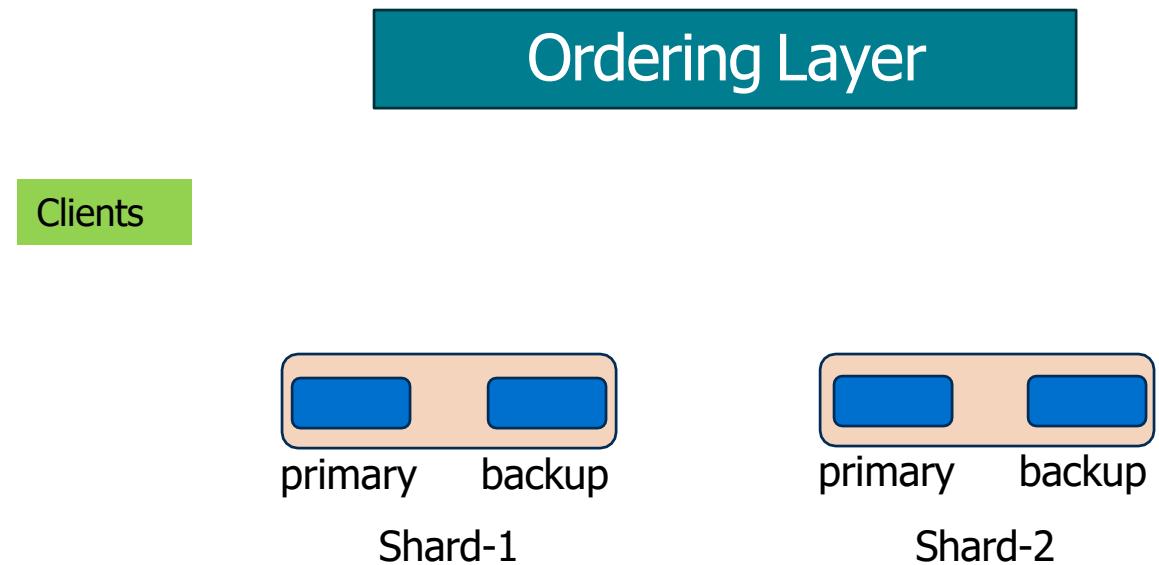
Eager Ordering → High Latency

- **Scalog**
 - durability first
 - then global ordering
 - multiple RTT +batch interval
 - Results in high ingestion latency for applications



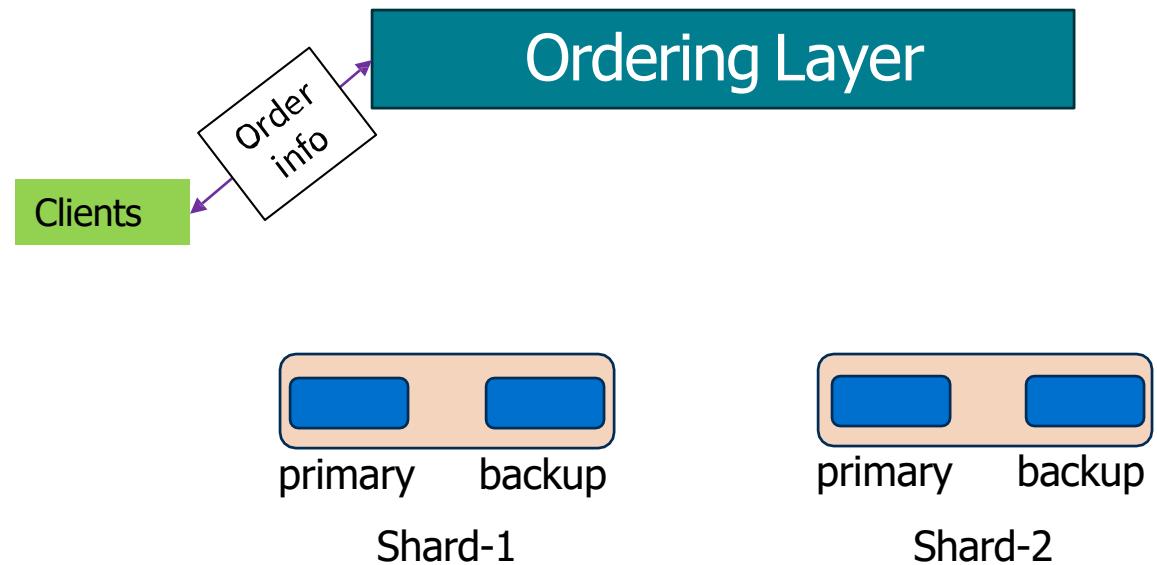
Eager Ordering → High Latency

- Corfu



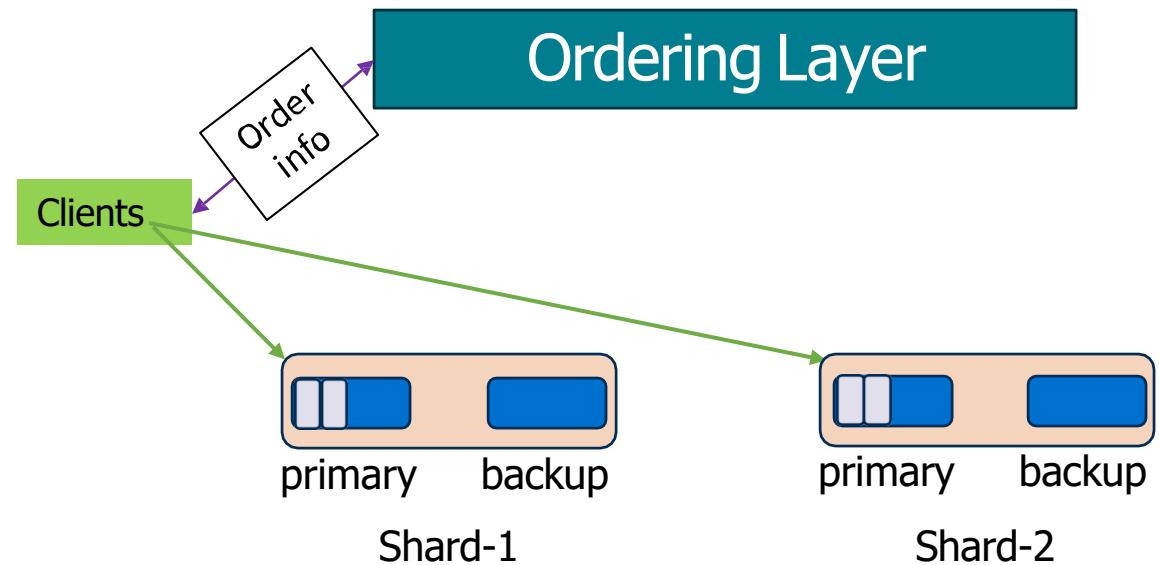
Eager Ordering → High Latency

- **Corfu**
 - global ordering first



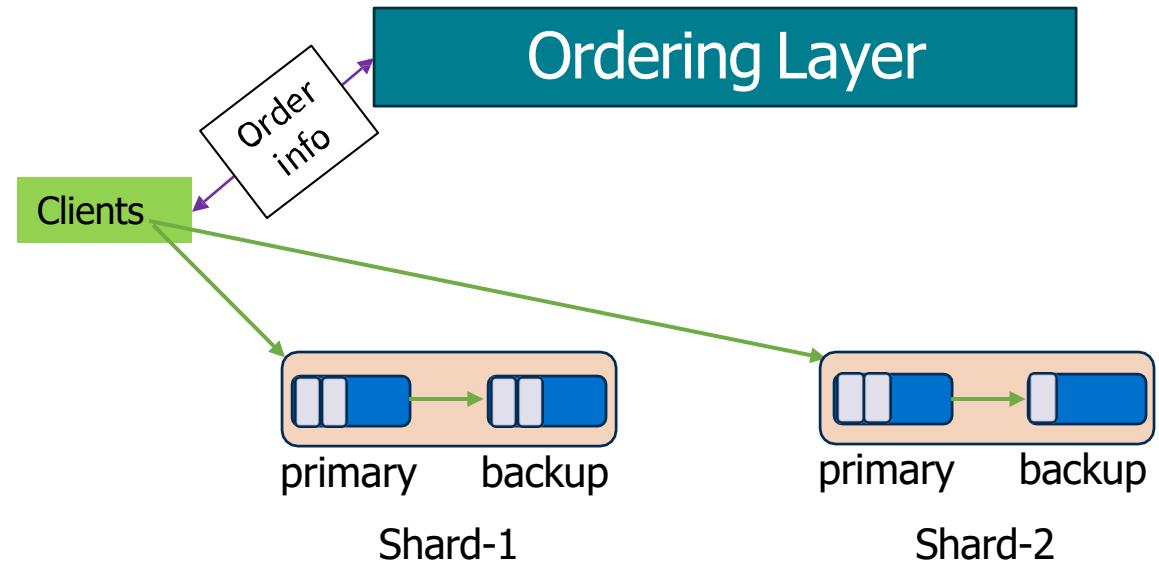
Eager Ordering → High Latency

- **Corfu**
 - global ordering first
 - then durability



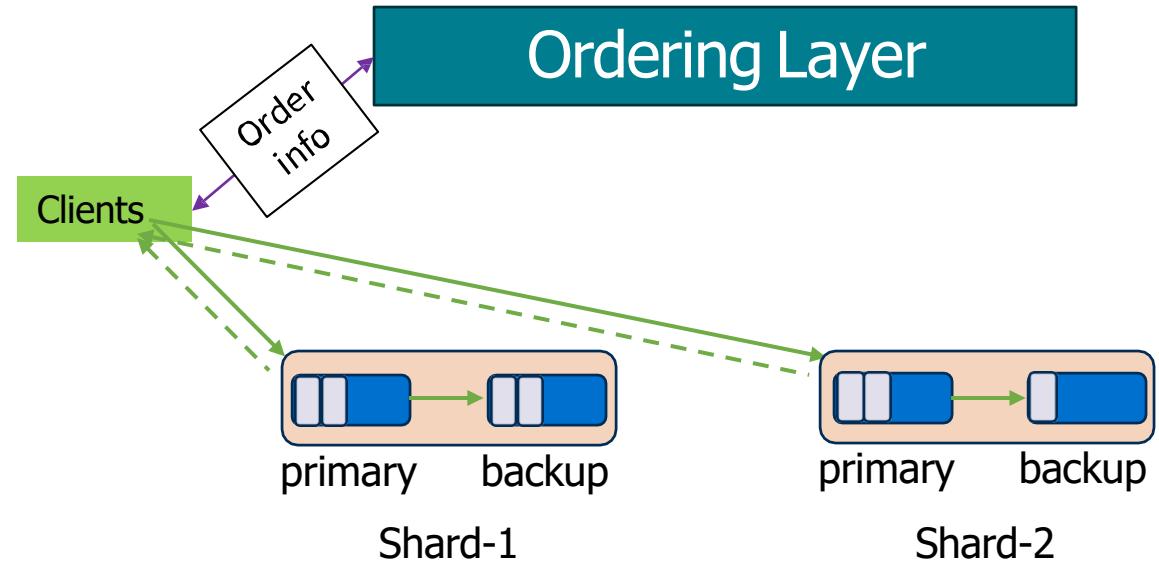
Eager Ordering → High Latency

- Corfu
 - global ordering first
 - then durability



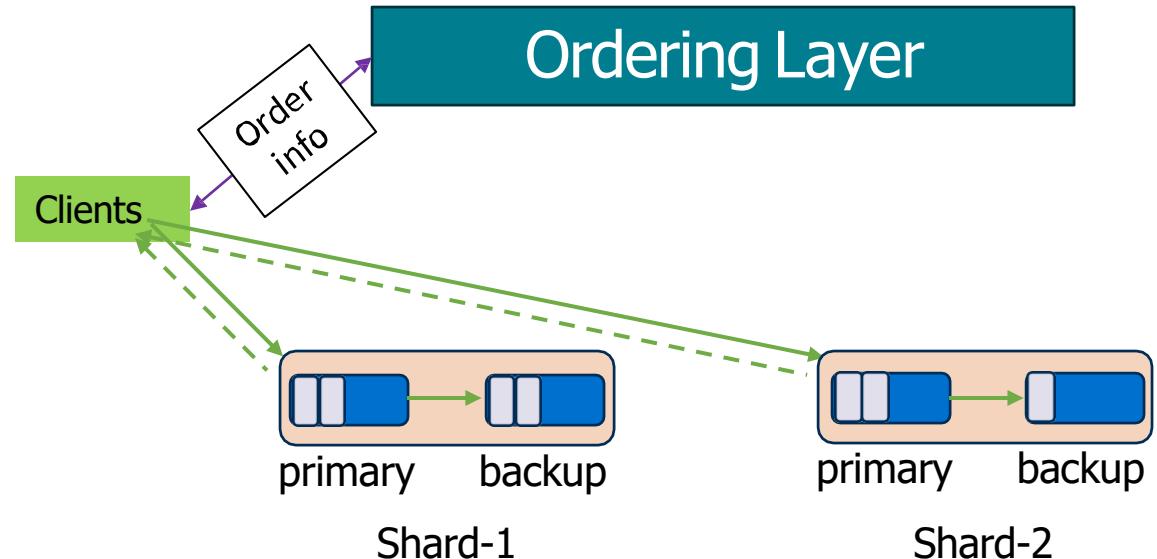
Eager Ordering → High Latency

- **Corfu**
 - global ordering first
 - then durability



Eager Ordering → High Latency

- **Corfu**
 - global ordering first
 - then durability
 - multiple RTT
 - Results in high ingestion latency for applications



Can a shared log **avoid eager ordering**, yet also **preserve the ordering guarantees** of conventional shared logs?

LazyLog: Idea and Abstraction

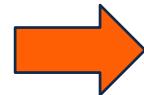
Idea



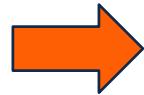
- No need for **eagerly bind** upon an append
 - But only make it durable

LazyLog: Idea and Abstraction

Idea



- No need for **eagerly bind** upon an append
 - But only make it durable



- Bind the records lazily, it must **enforce** ordering before positions can be read

LazyLog: Idea and Abstraction

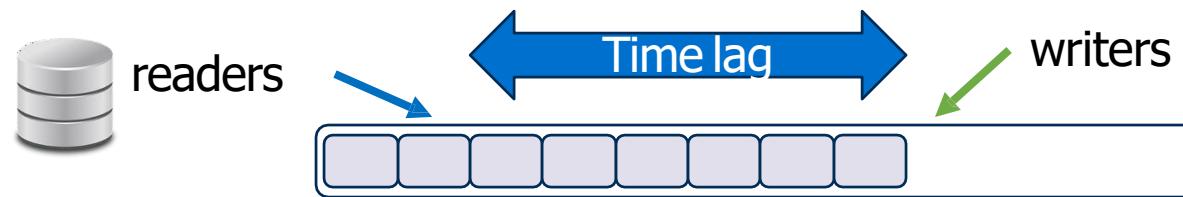
Idea

- No need for **eagerly bind** upon an append
 - But only make it durable
- ○ Bind the records lazily, it must **enforce** ordering before positions can be read
- ○ Shared log can do the ordering comfortably **in the background**

Insight

A shared log can **defer ordering** upon appends But
establish it before reads arrive

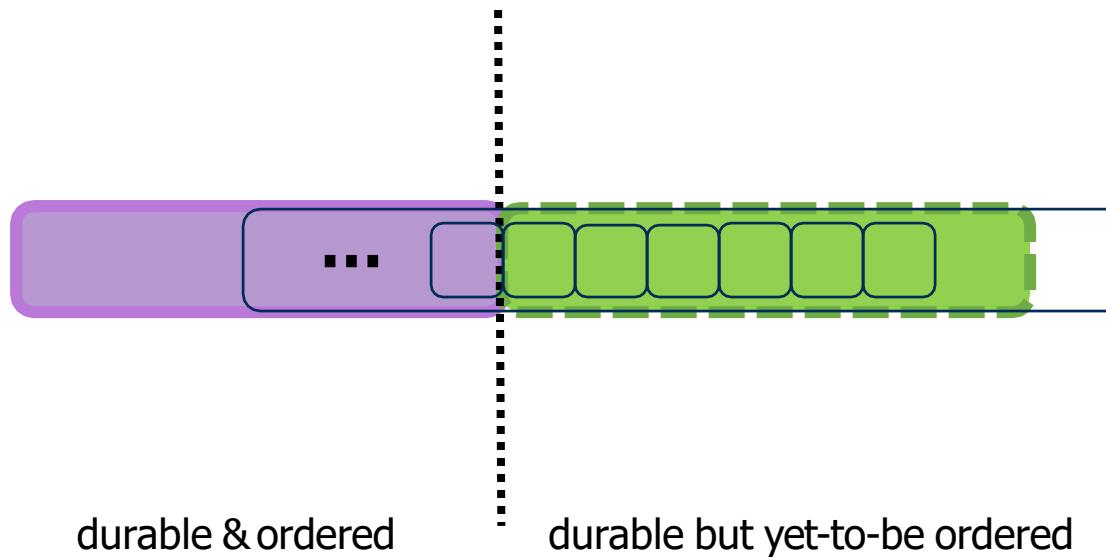
Holds for Many Apps



- Readers and writers are **time-decoupled**: readers typically **lag behind** writers

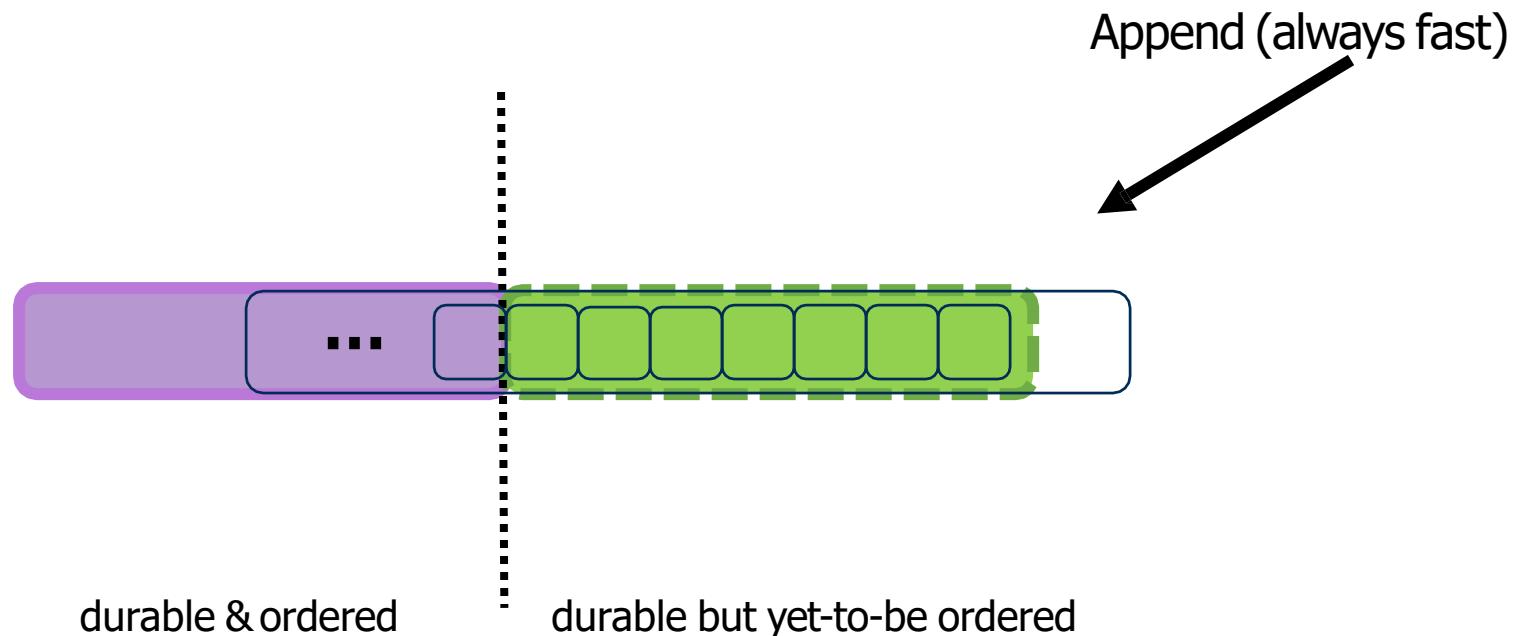
Performance Property

- Cannot be too lazy – keep ordering in the background



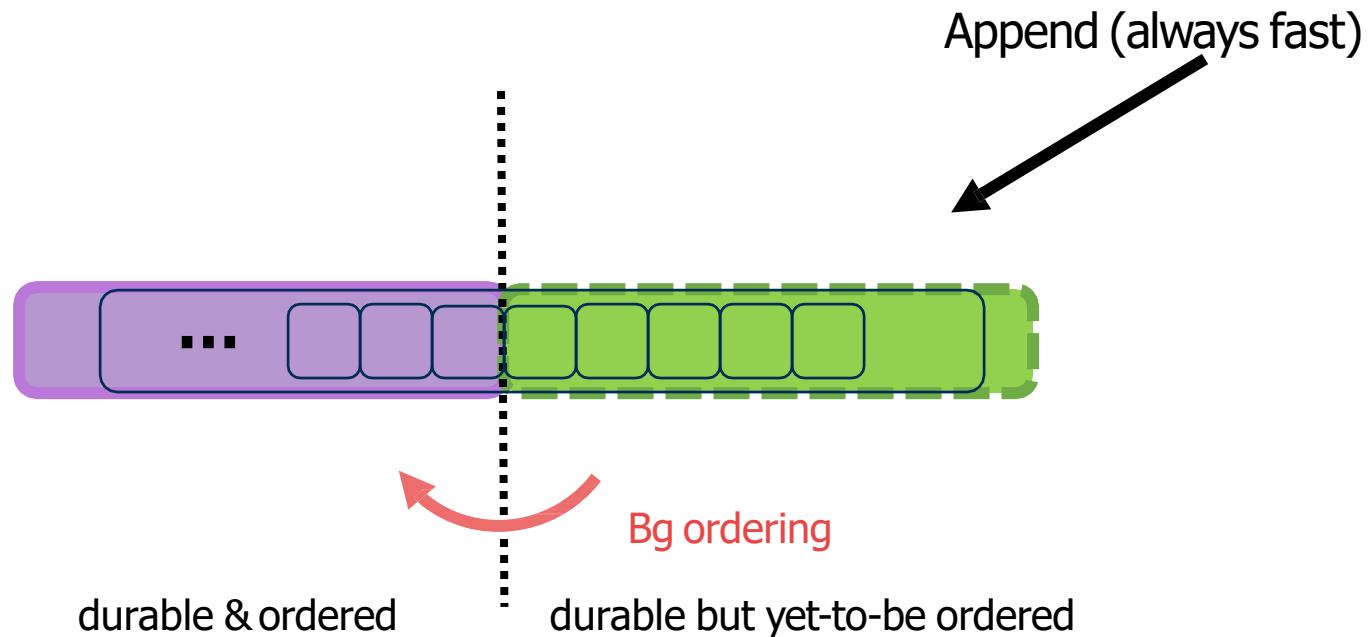
Performance Property

- Cannot be too lazy – keep ordering in the background



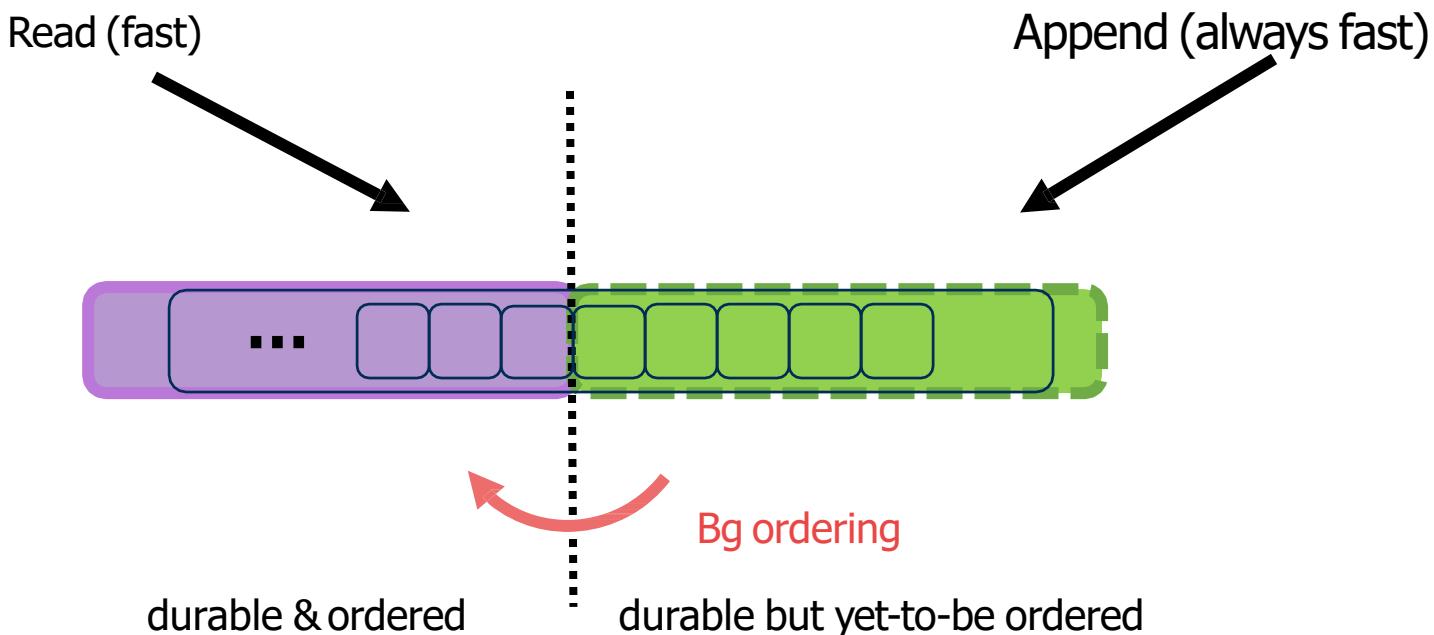
Performance Property

- Cannot be too lazy – keep ordering in the background



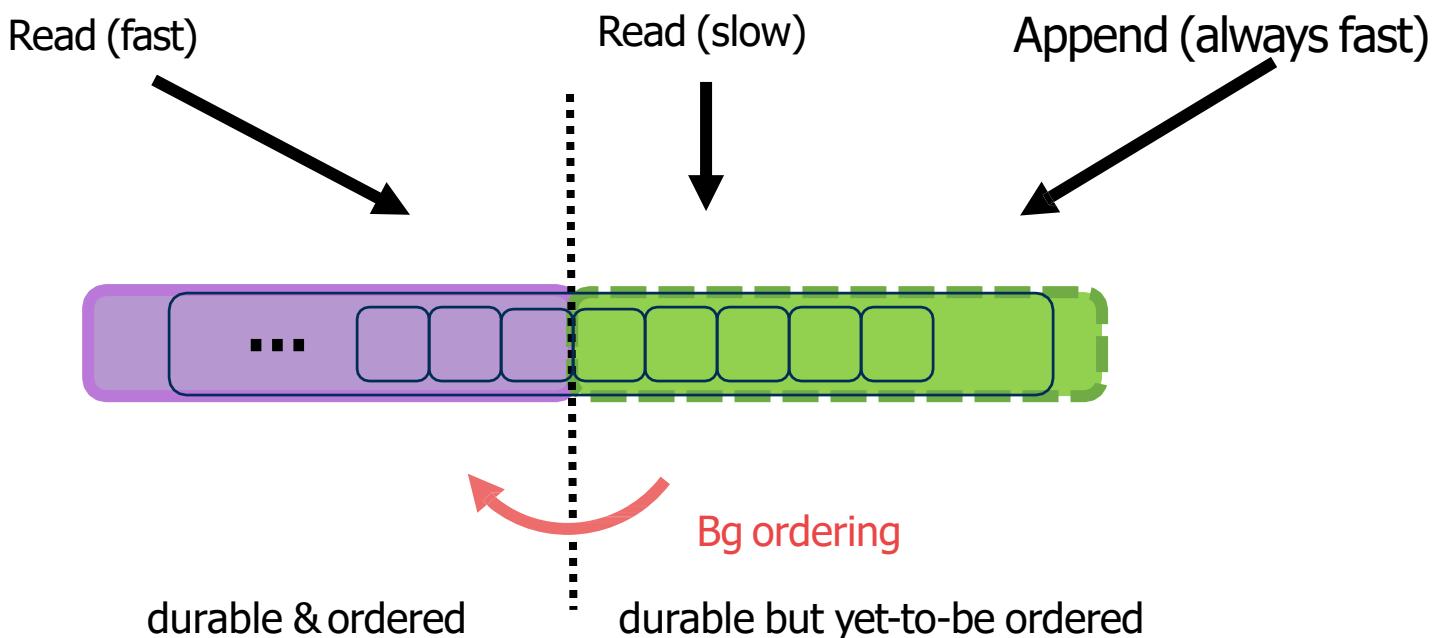
Performance Property

- Cannot be too lazy – keep ordering in the background



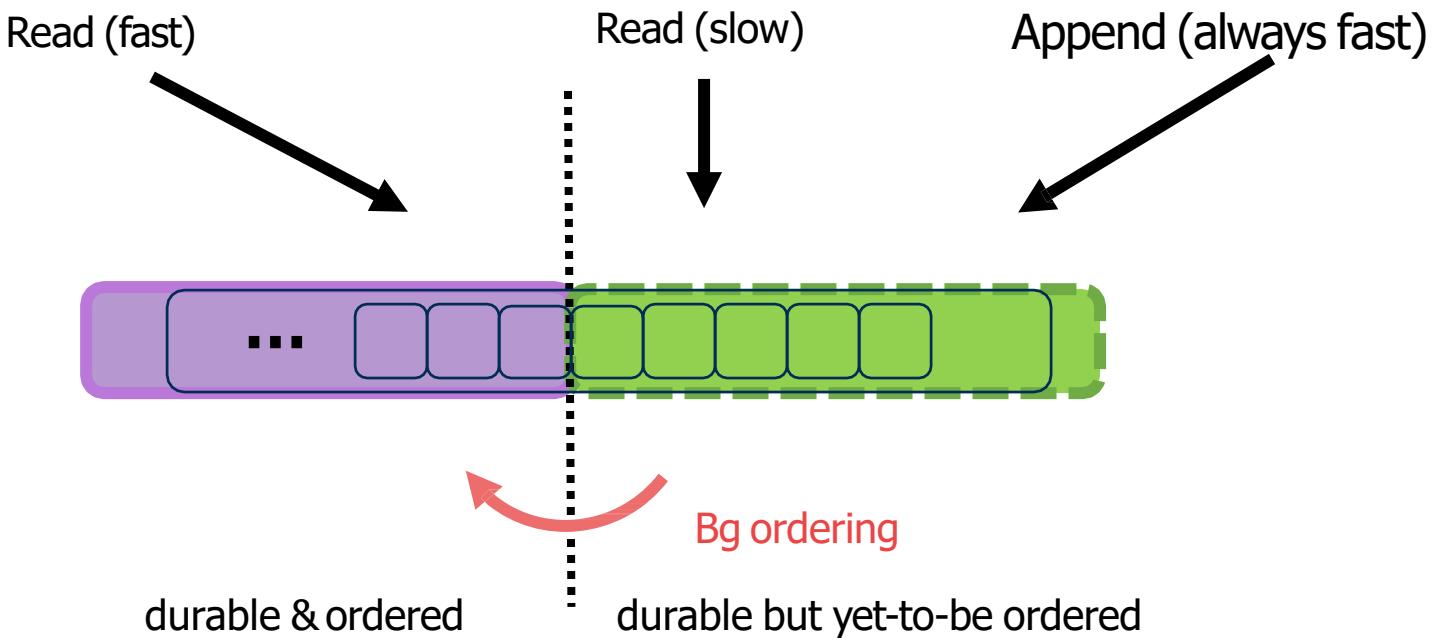
Performance Property

- Cannot be too lazy – keep ordering in the background



Performance Property

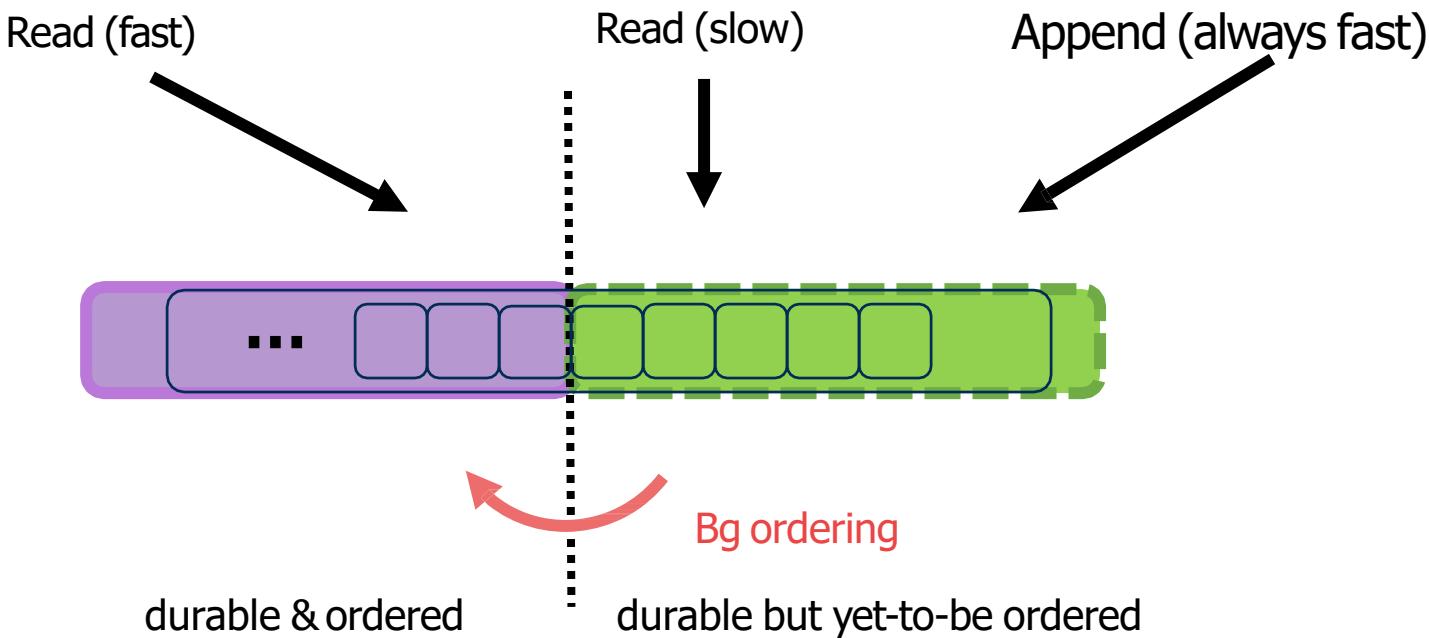
- Cannot be too lazy – keep ordering in the background



- For many apps – reads are always fast

Performance Property

- Cannot be too lazy – keep ordering in the background



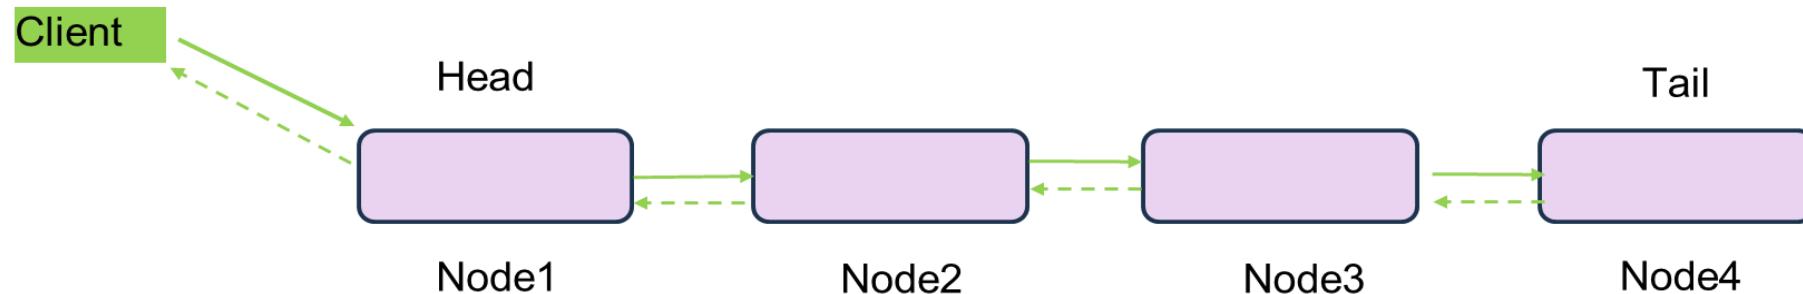
- For many apps – reads are always fast
- For immediate read, LazyLog preserves the performance of eager shared logs
 - never worse than an eager-ordering shared log!

Insight

Why **1-RTT Appends** are hard to achieve in Systems like
Chain Paxos or Raft?

Chain Paxos

Requires multi-step coordination between replicas:



Client → Node1

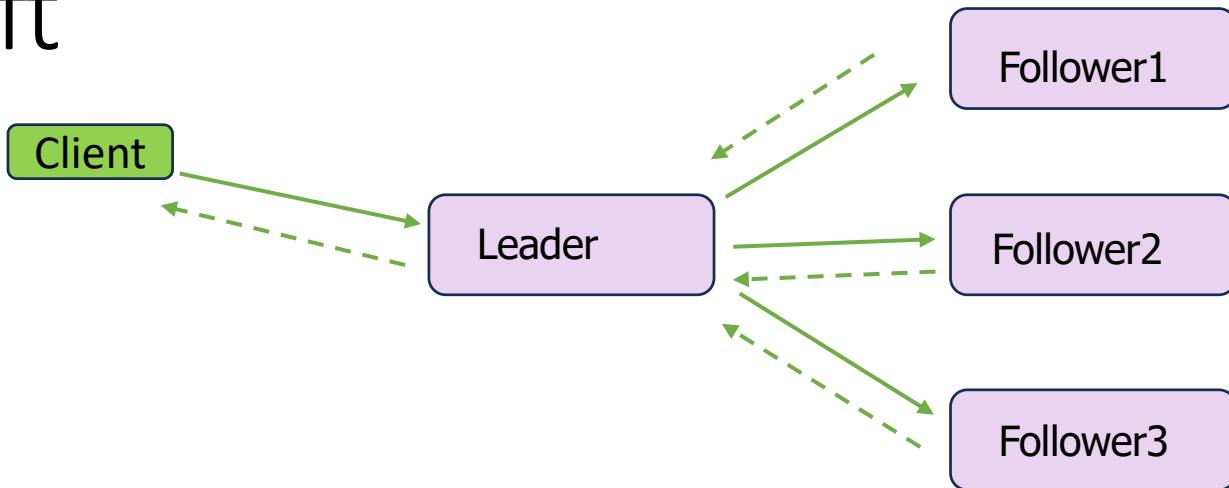
Node1 → Node2

Node2 → Node3

Node3 → Node4

Node4 → Client

Raft



Client → Leader (Send Request)

Leader → Followers (Replicate Entry)

Followers → Leader (Acks)

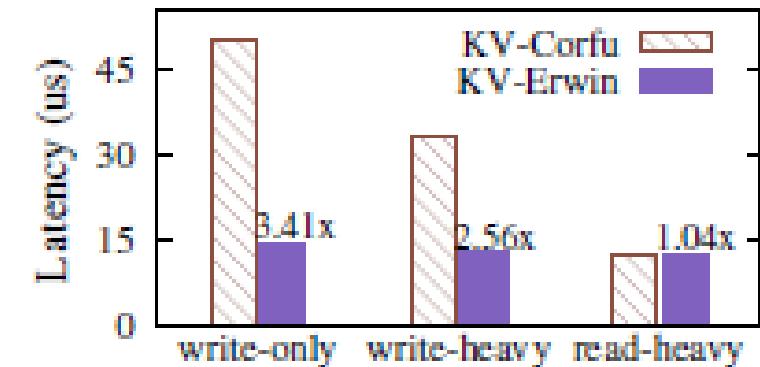
Leader → Client (Reply)

Insight

Application: Lazy-Log

Application: Key-Value Store

- Shared-log based key-value store
- Readers – Writers are decoupled
- Write-processing server- handles append request
- Read server – consume the log
- One writer server, one reader server.
- Shared log with 1 shard (3 replicas).
- Compared Corfu vs Erwin (LazyLog)



(a) KV Store

3 Tests

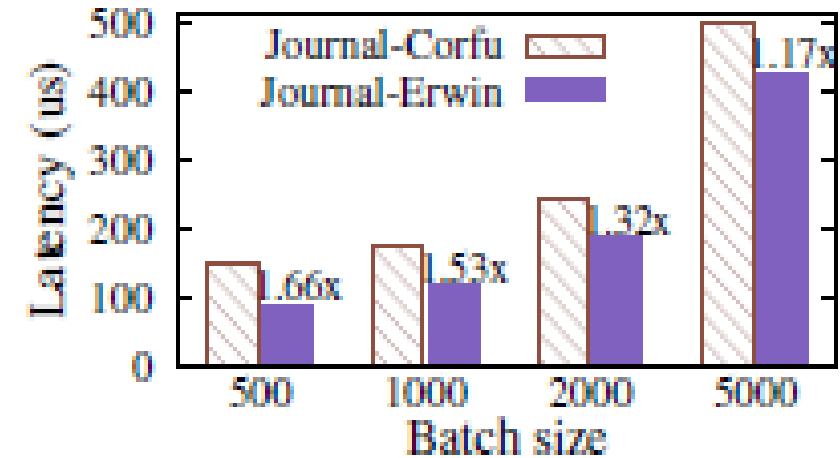
- YCSB 100% writes
- YCSB-A 50% writes & 50% reads
- YCSB-B 5% writes & 95% reads

write only	3.4× lower latency
write heavy	2.5× lower latency
read heavy	No change

Application: Journaling for Stream Processing

- Task Worker-
 - reads the incoming records
 - use shared log to store their states
 - output the result
 - processes in batch

To recover during failure



Latency = Record Reading + Processing + checkpointing + emitting

Compared **Corfu** vs **Erwin**

Tests: stream-processing word count app with 5 workers For Different batches

- 5,000 records more computation **1.17x** faster
- 500 records logging dominates **1.66x** faster

System Design

Designed an implementation of the LazyLog interface: **Erwin**

Goal- Offers linearizable ordering across shards with 1-RTT appends

System Design

Designed an implementation of the LazyLog interface: **Erwin**

Offers linearizable ordering across shards with 1-RTT appends

Treats shards as black boxes and requires them to support:

- append an entry and
- read the entry at a specified index / position

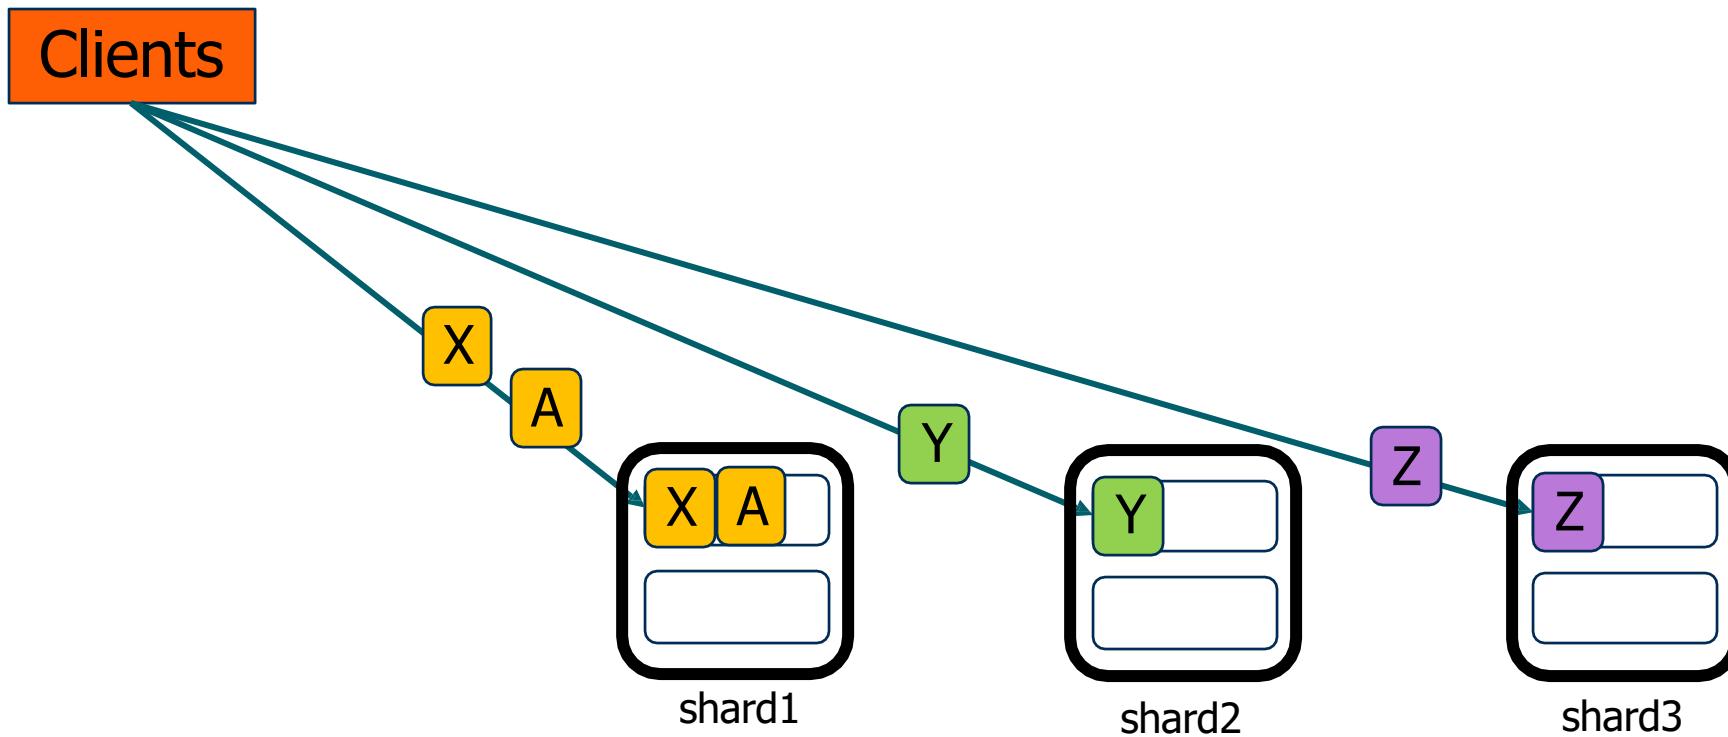
Why Erwin?

Motivation

Motivation



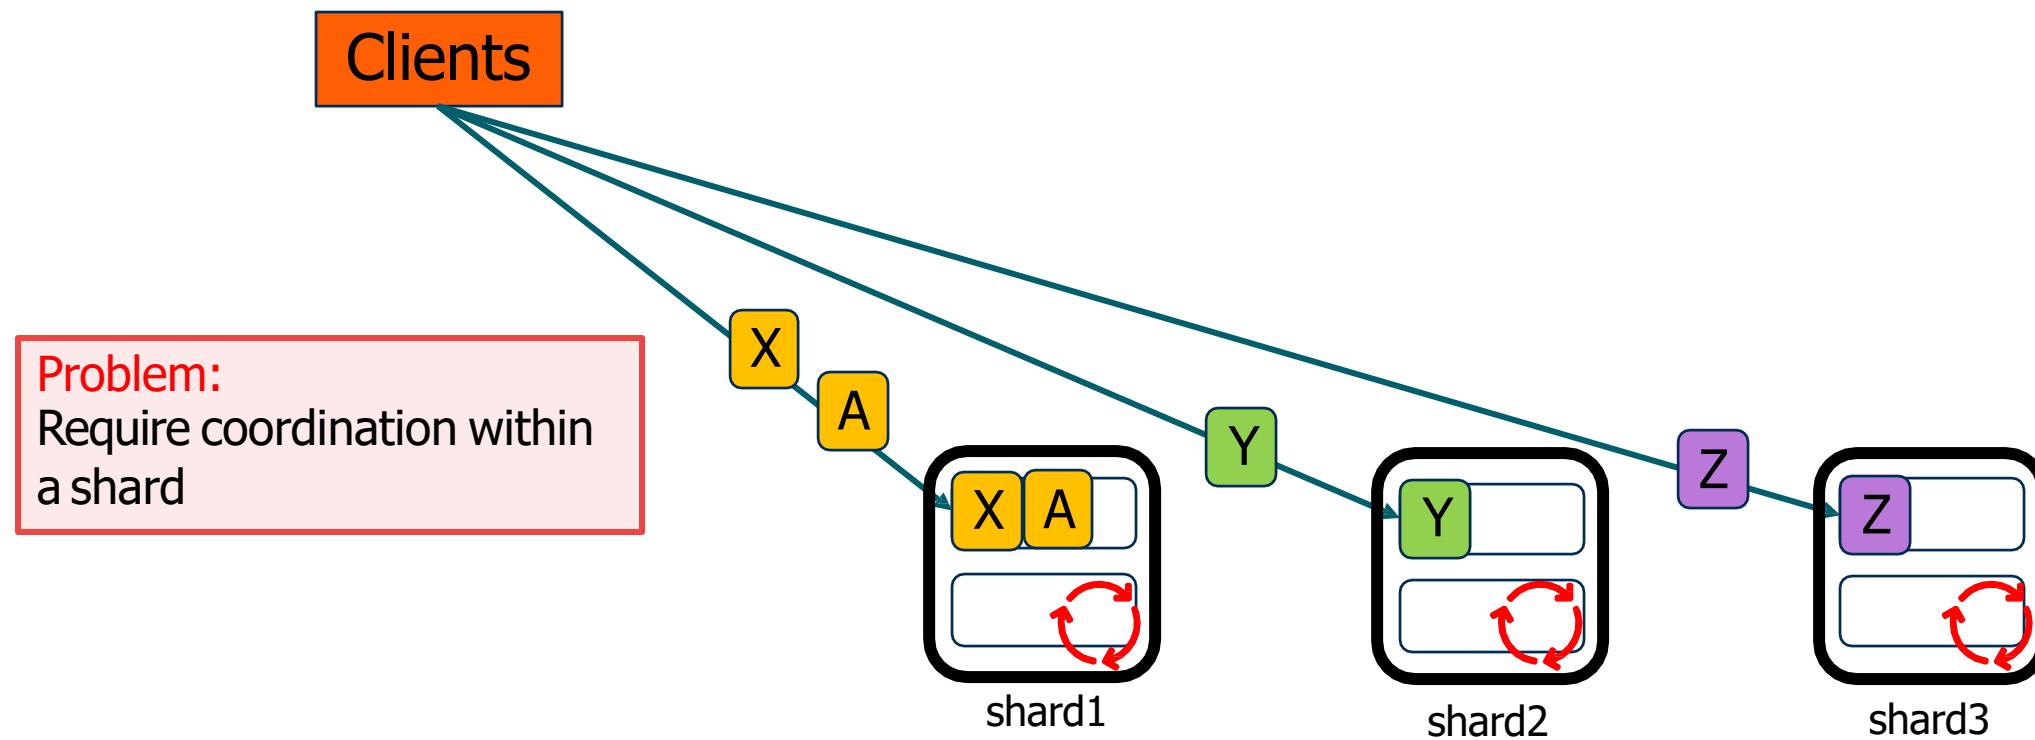
Record



Motivation



Record



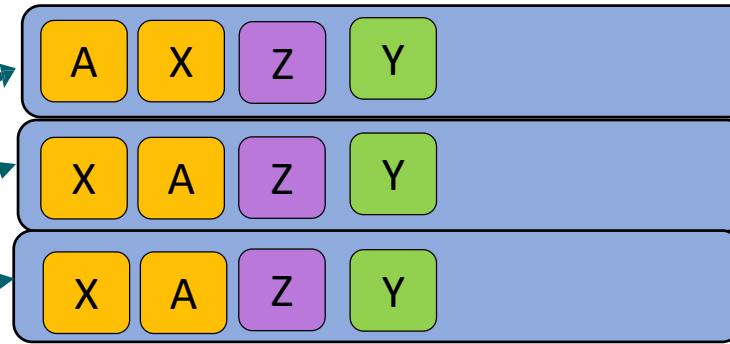
Motivation



Record

Clients

Sequencing layer



Problem:

Require coordination within a shard

Solution:

Use a buffer and send requests to this buffer



shard1



shard2



shard3

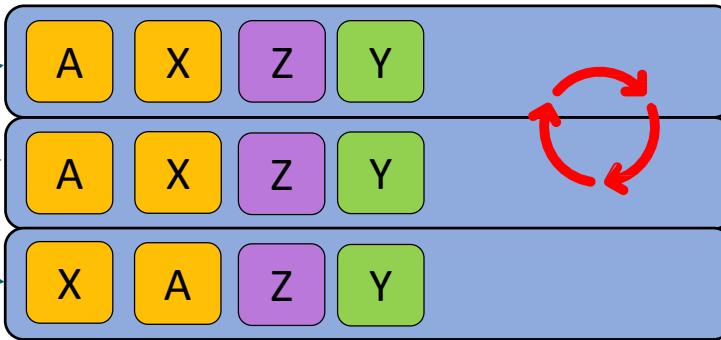
Erwin: 1-RTT Append



Record

Clients

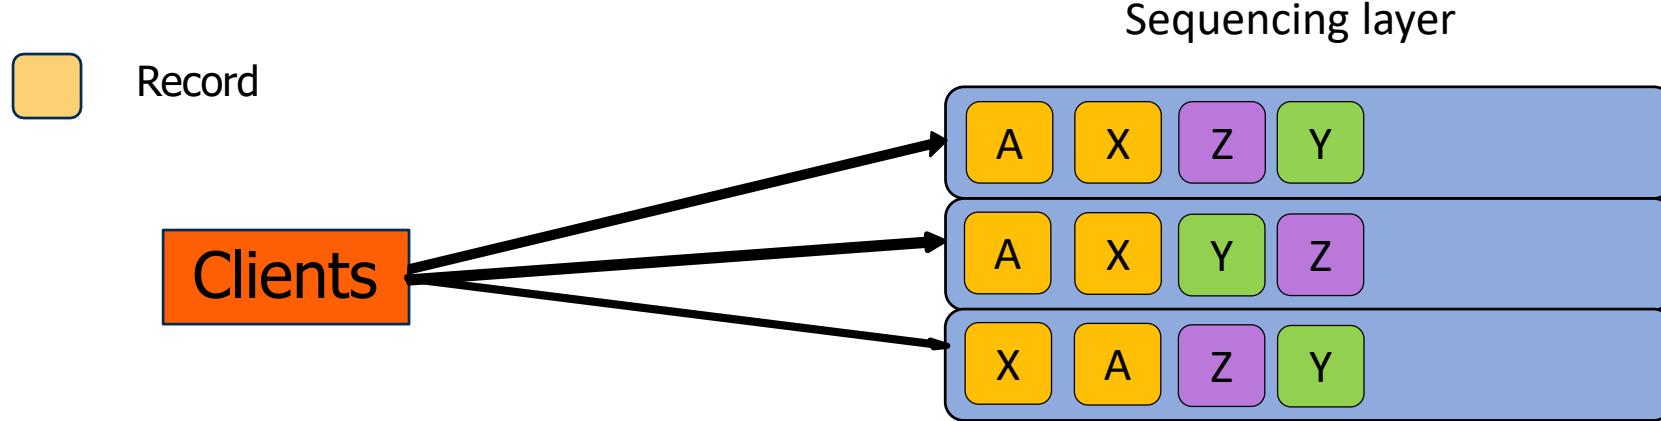
Sequencing layer



Problem:

Sequencing layer must run consensus to make ordering fault-tolerant → Incurs coordination within replicas

Erwin: 1-RTT Append



Problem:
Sequencing layer must run consensus to make ordering fault-tolerant → Incurs coordination within replicas

Solution:
Coordination-free sequencing

- Clients write to shard replicas in 1RTT; in same RTT, write metadata to **all** seq replicas
 - Appends complete in 1 RTT
 - Erwin allows different orders across sequencing replicas
 - but without violating the linearizability

Erwin's Architecture

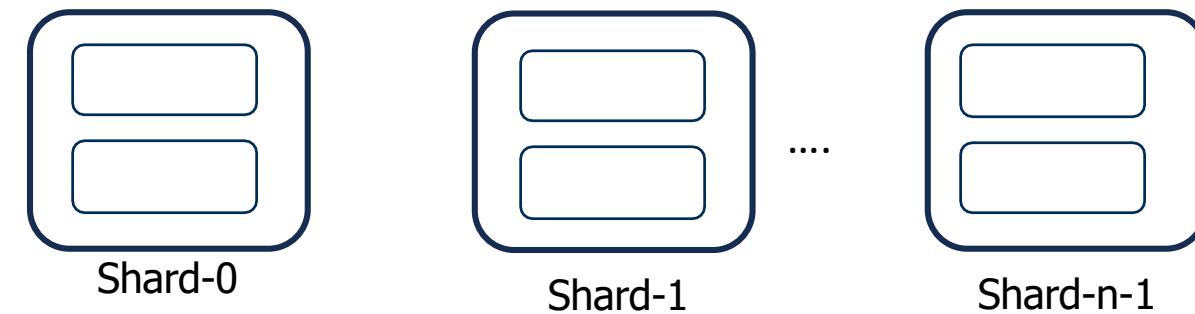
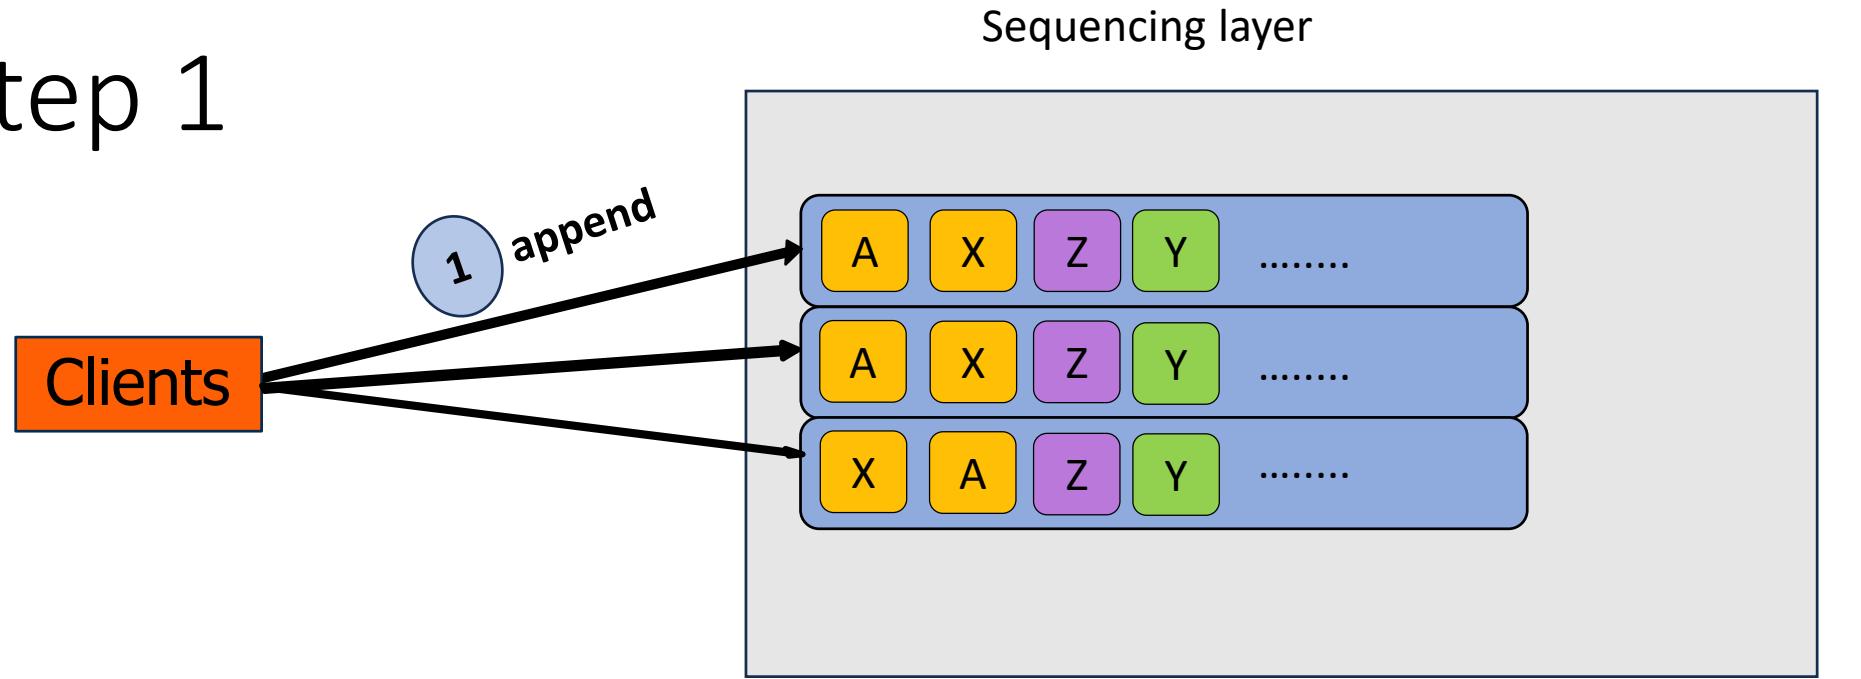
- Has set of unmodified primary-backup shards and **Sequencing layer**
- **Sequencing layer:**
 - Fault-tolerant : $f+1$ replicas to tolerate f failures
 - Coordination-free
 - Provide only short-term durability

How does Erwin Append work?

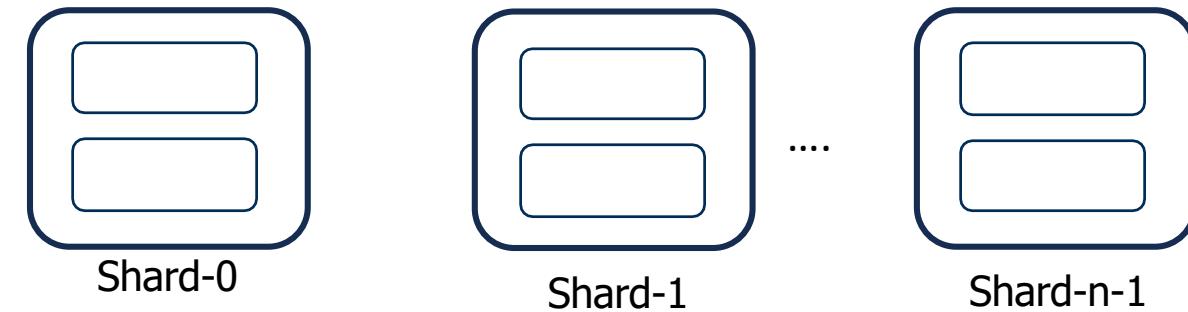
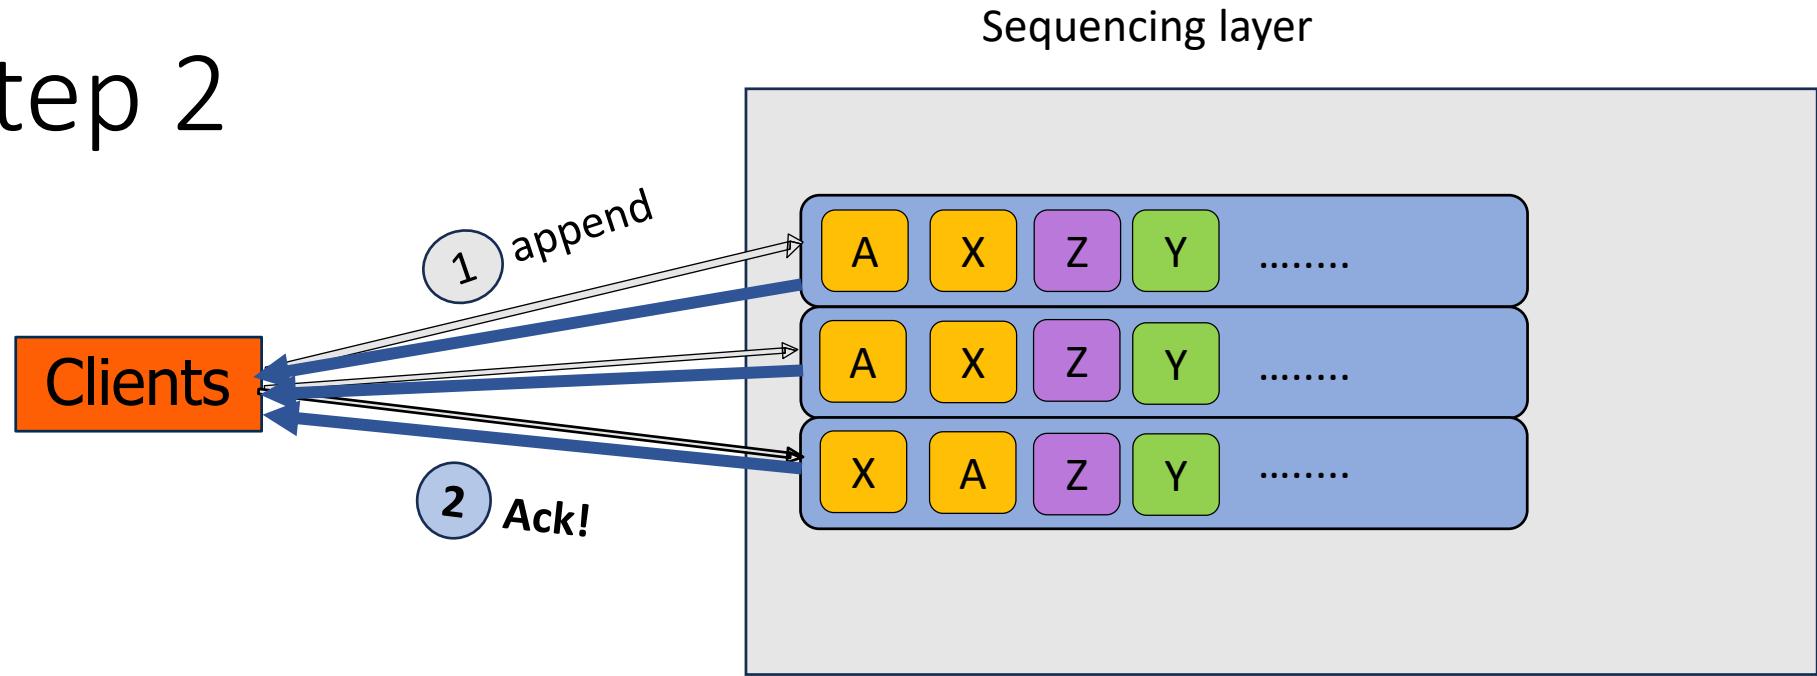
How is Lazy ordering done?

Step by step

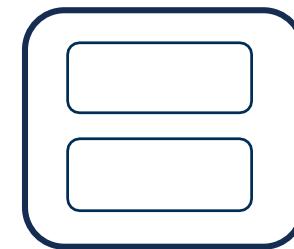
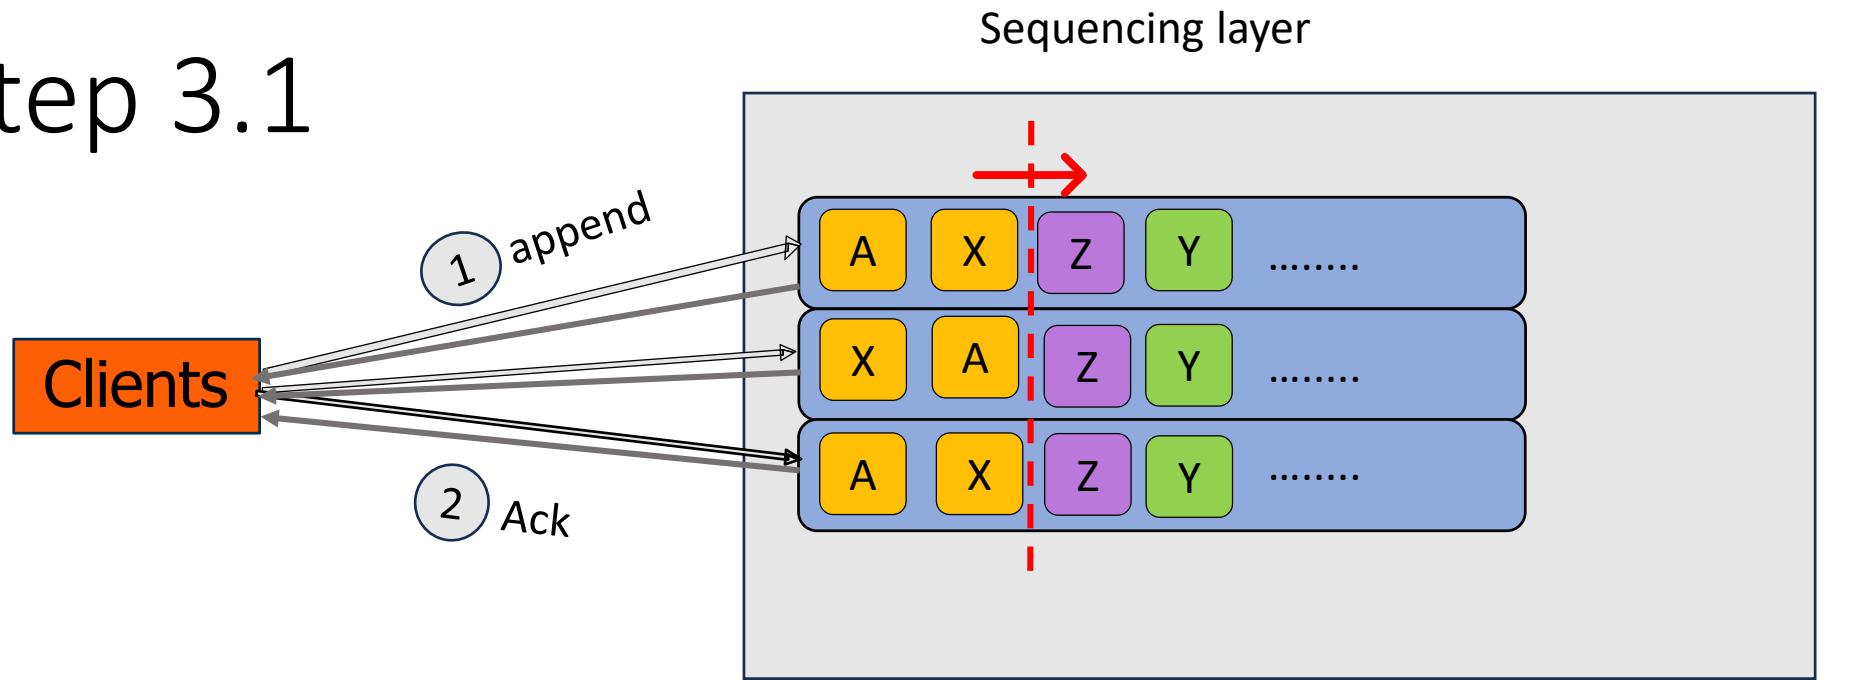
Step 1



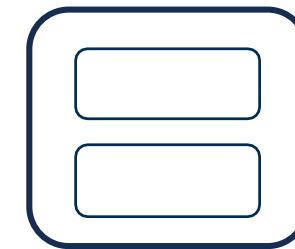
Step 2



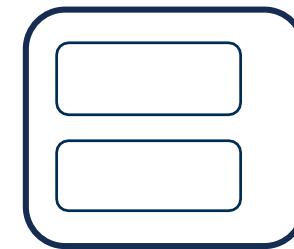
Step 3.1



Shard-0

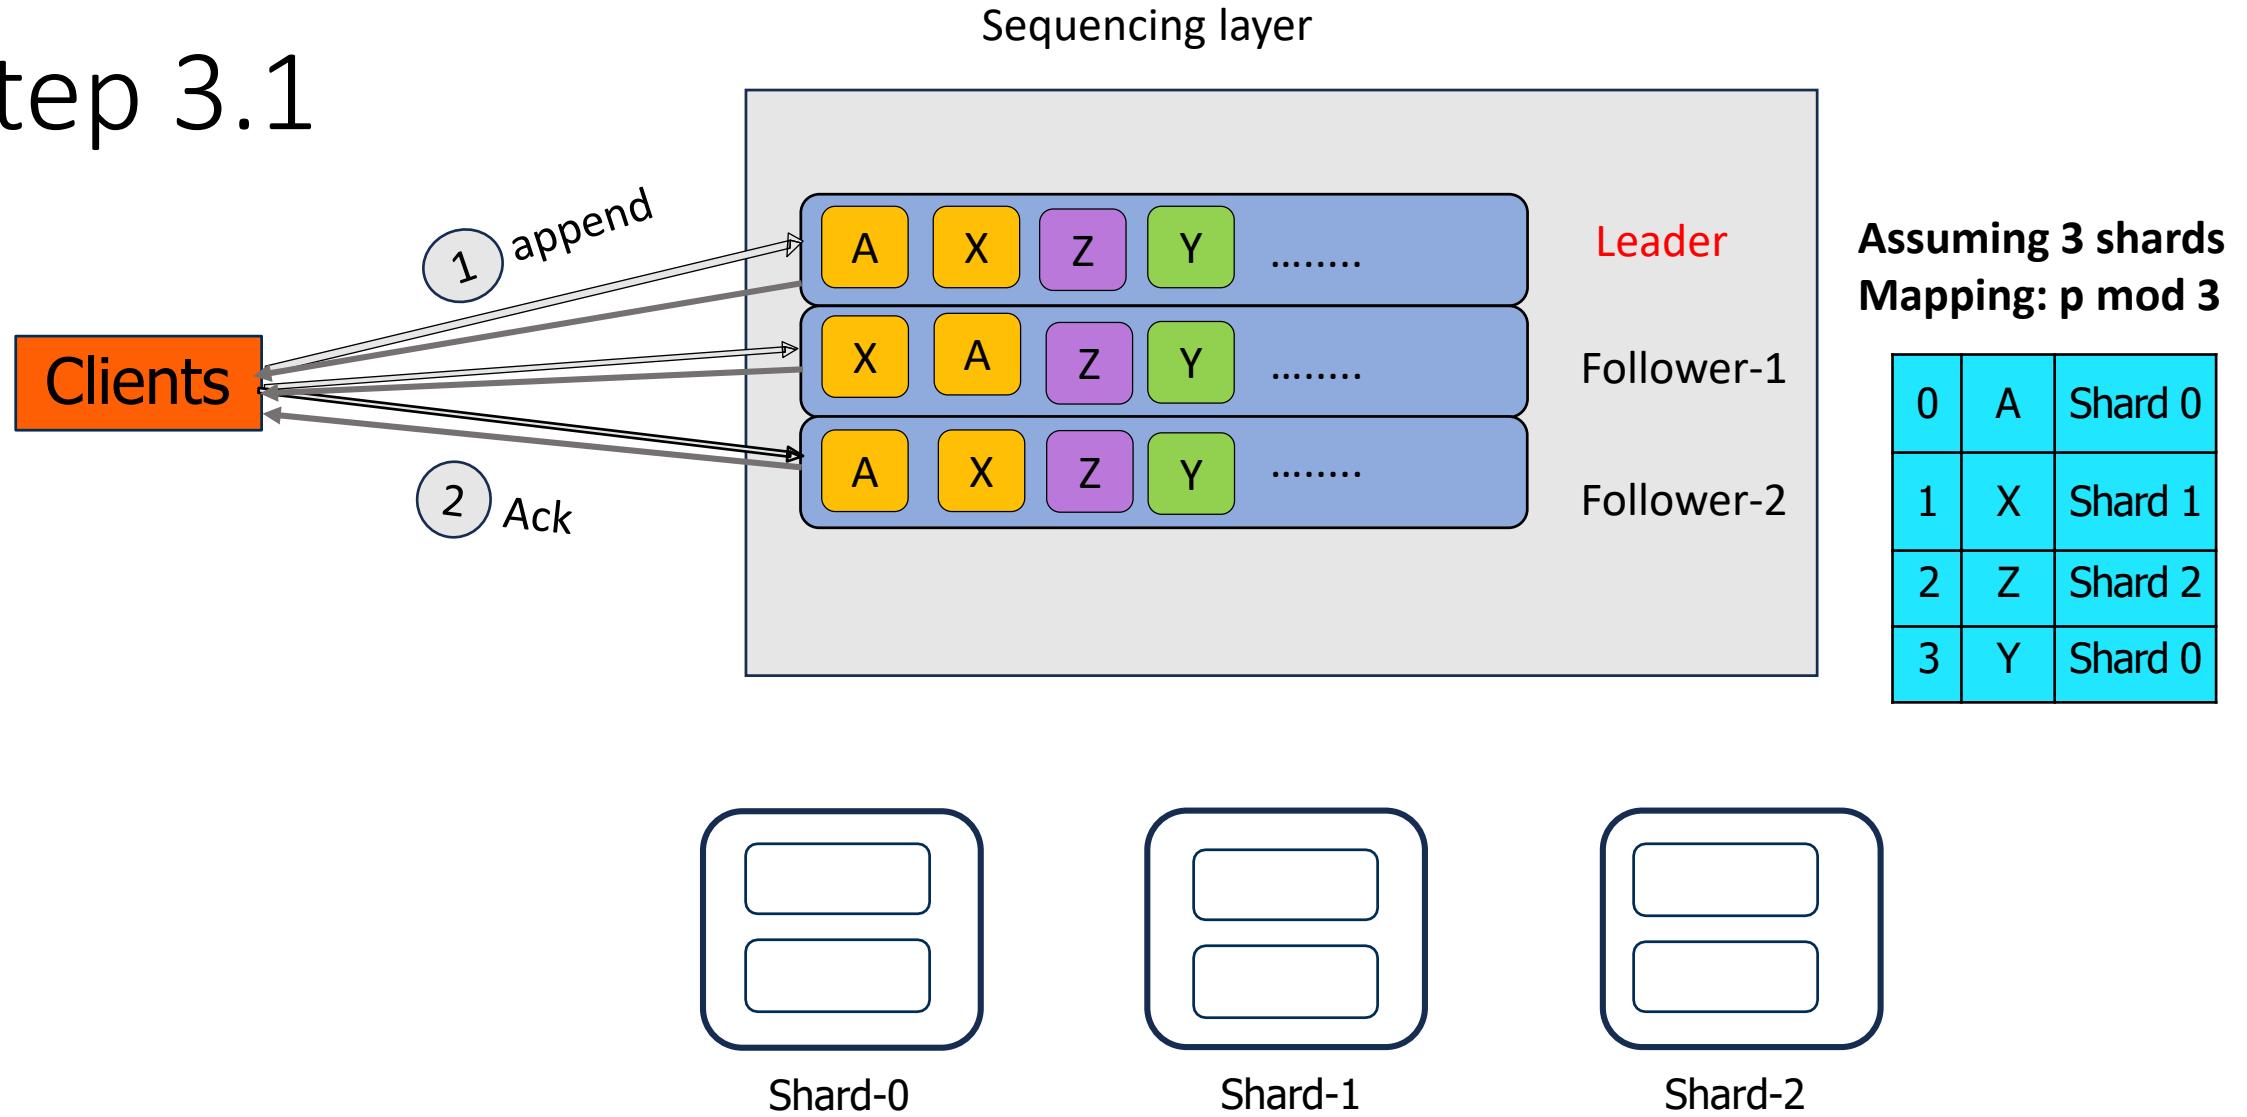


Shard-1

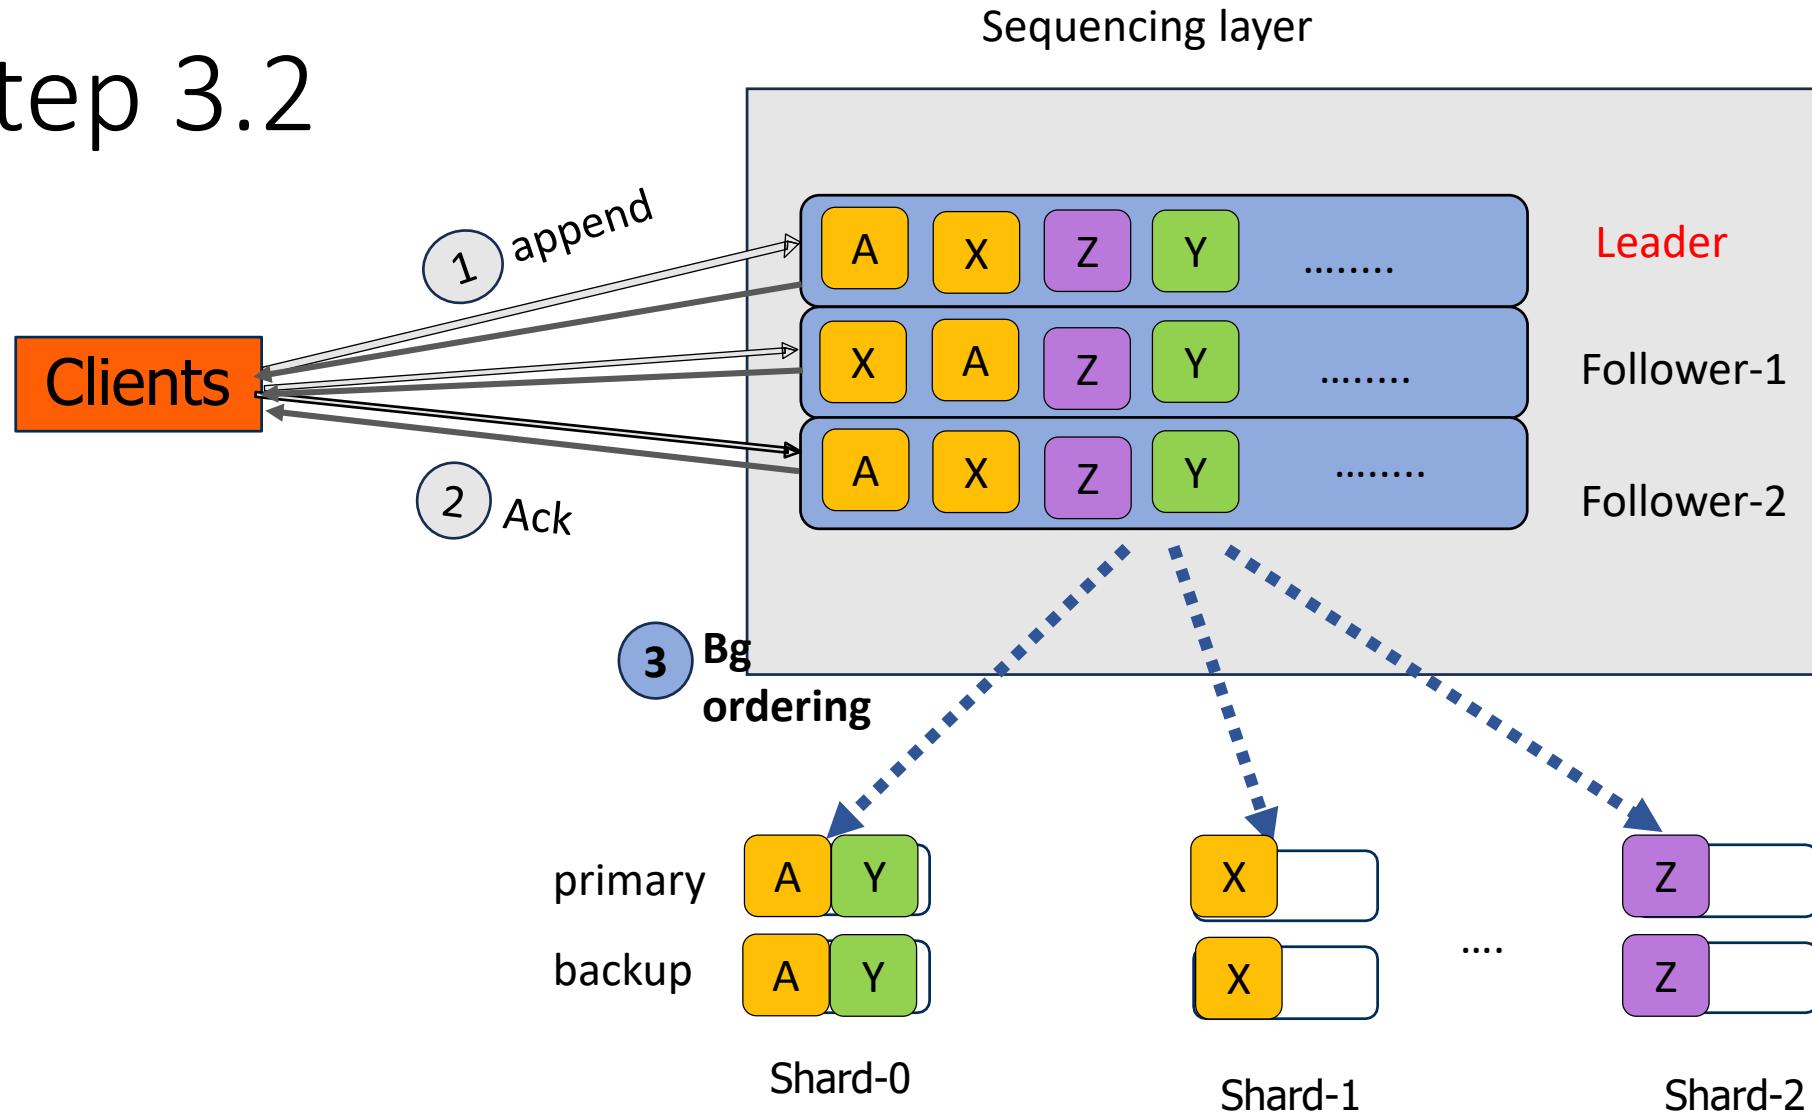


Shard-2

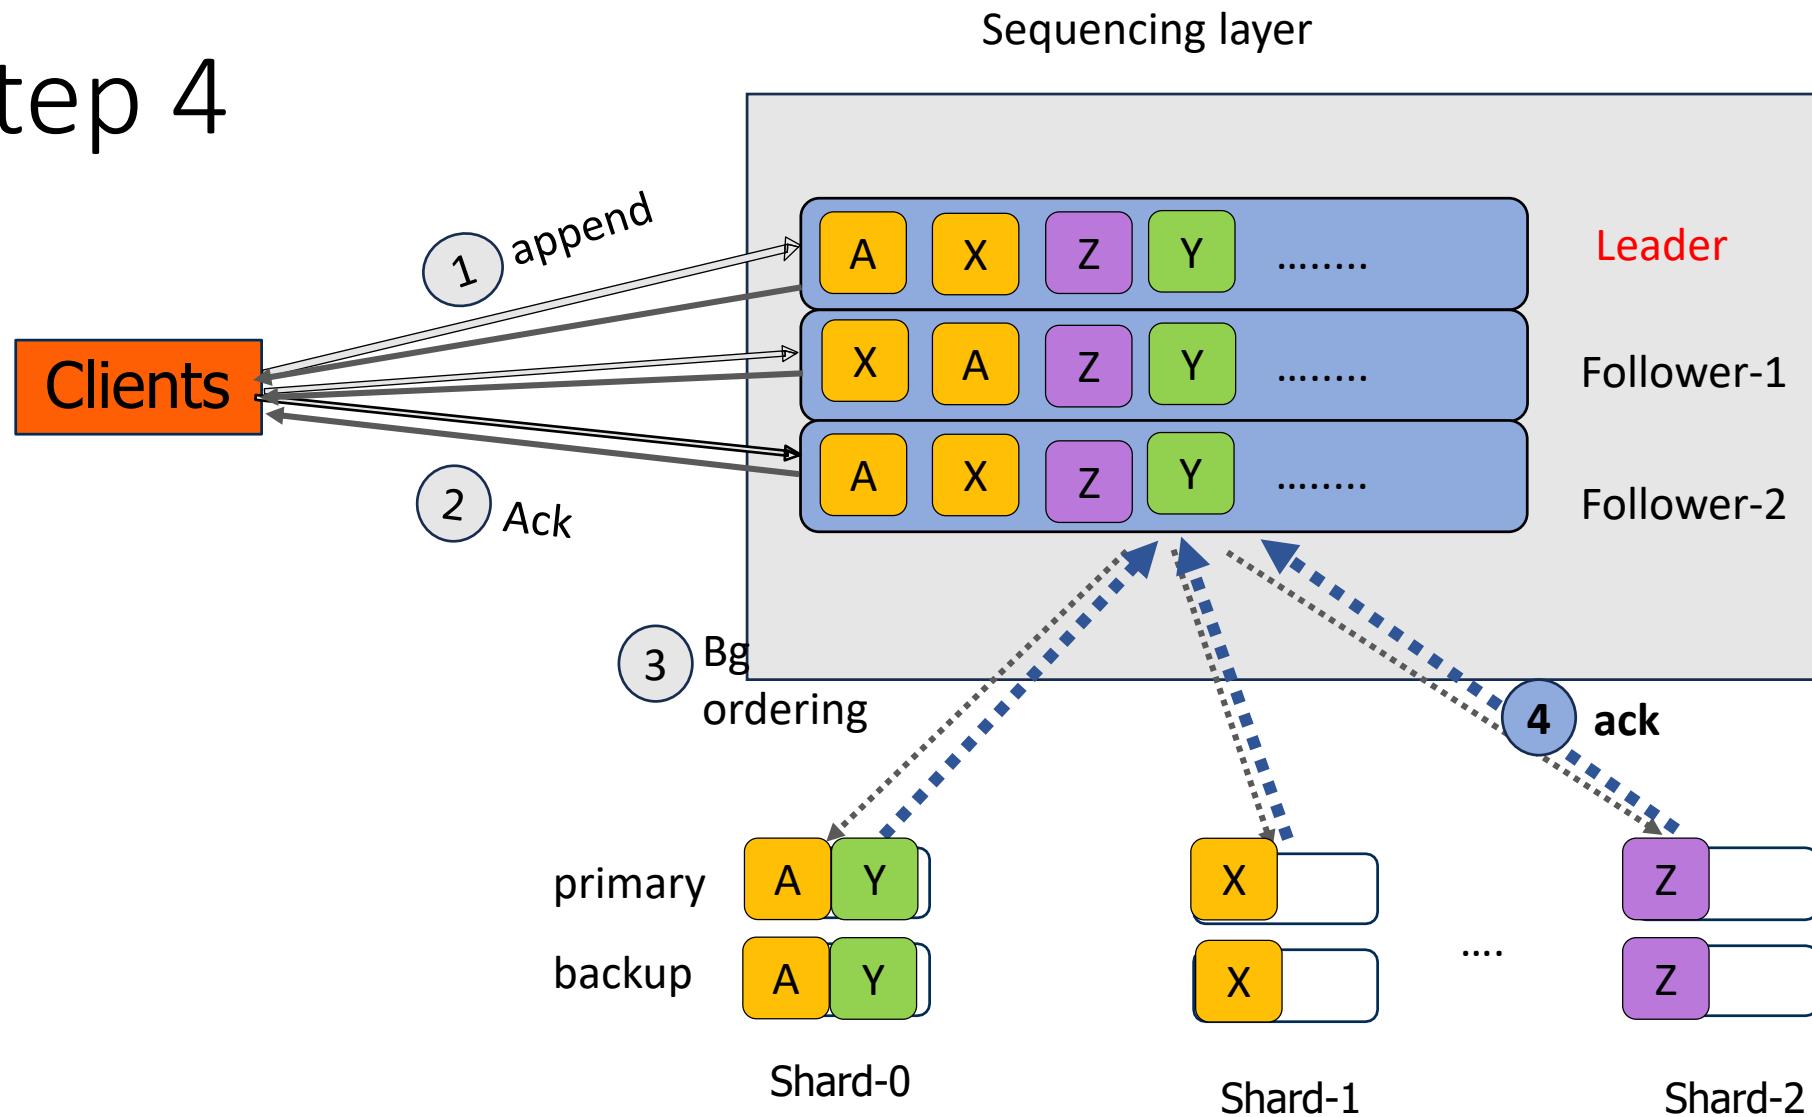
Step 3.1



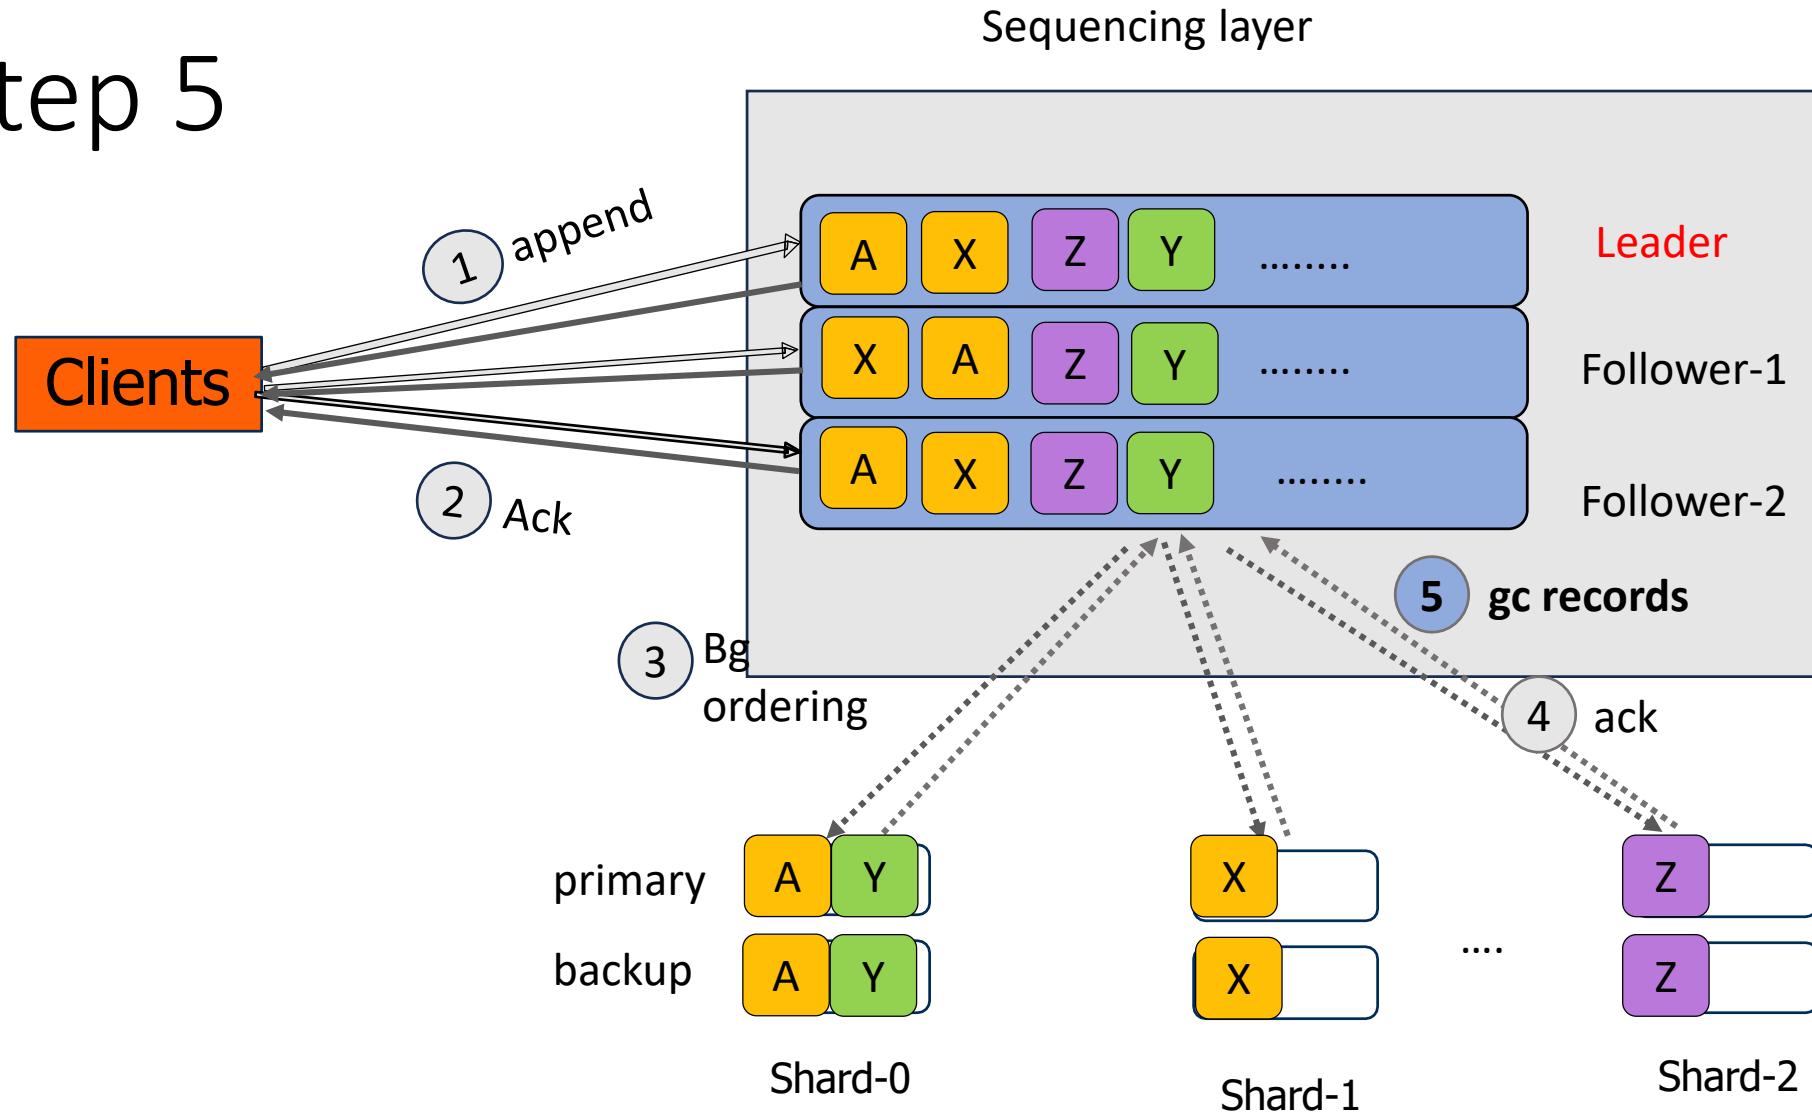
Step 3.2



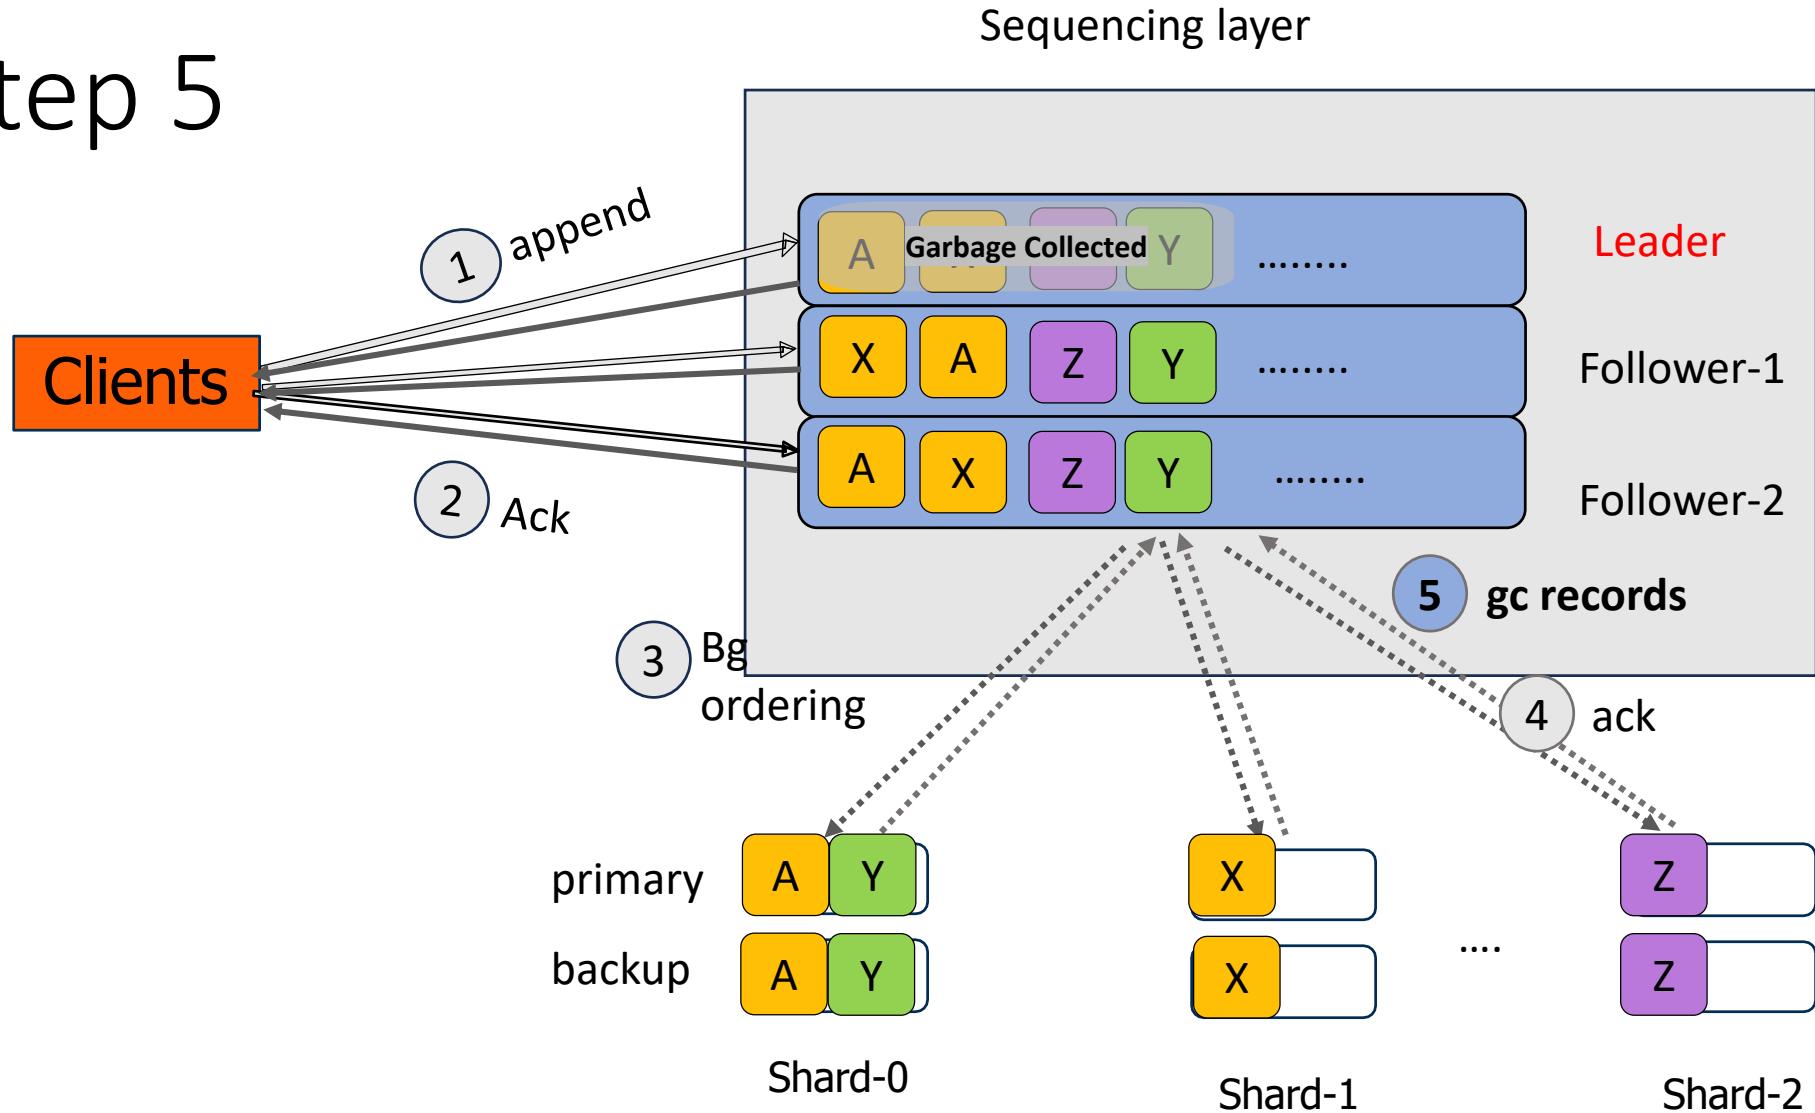
Step 4



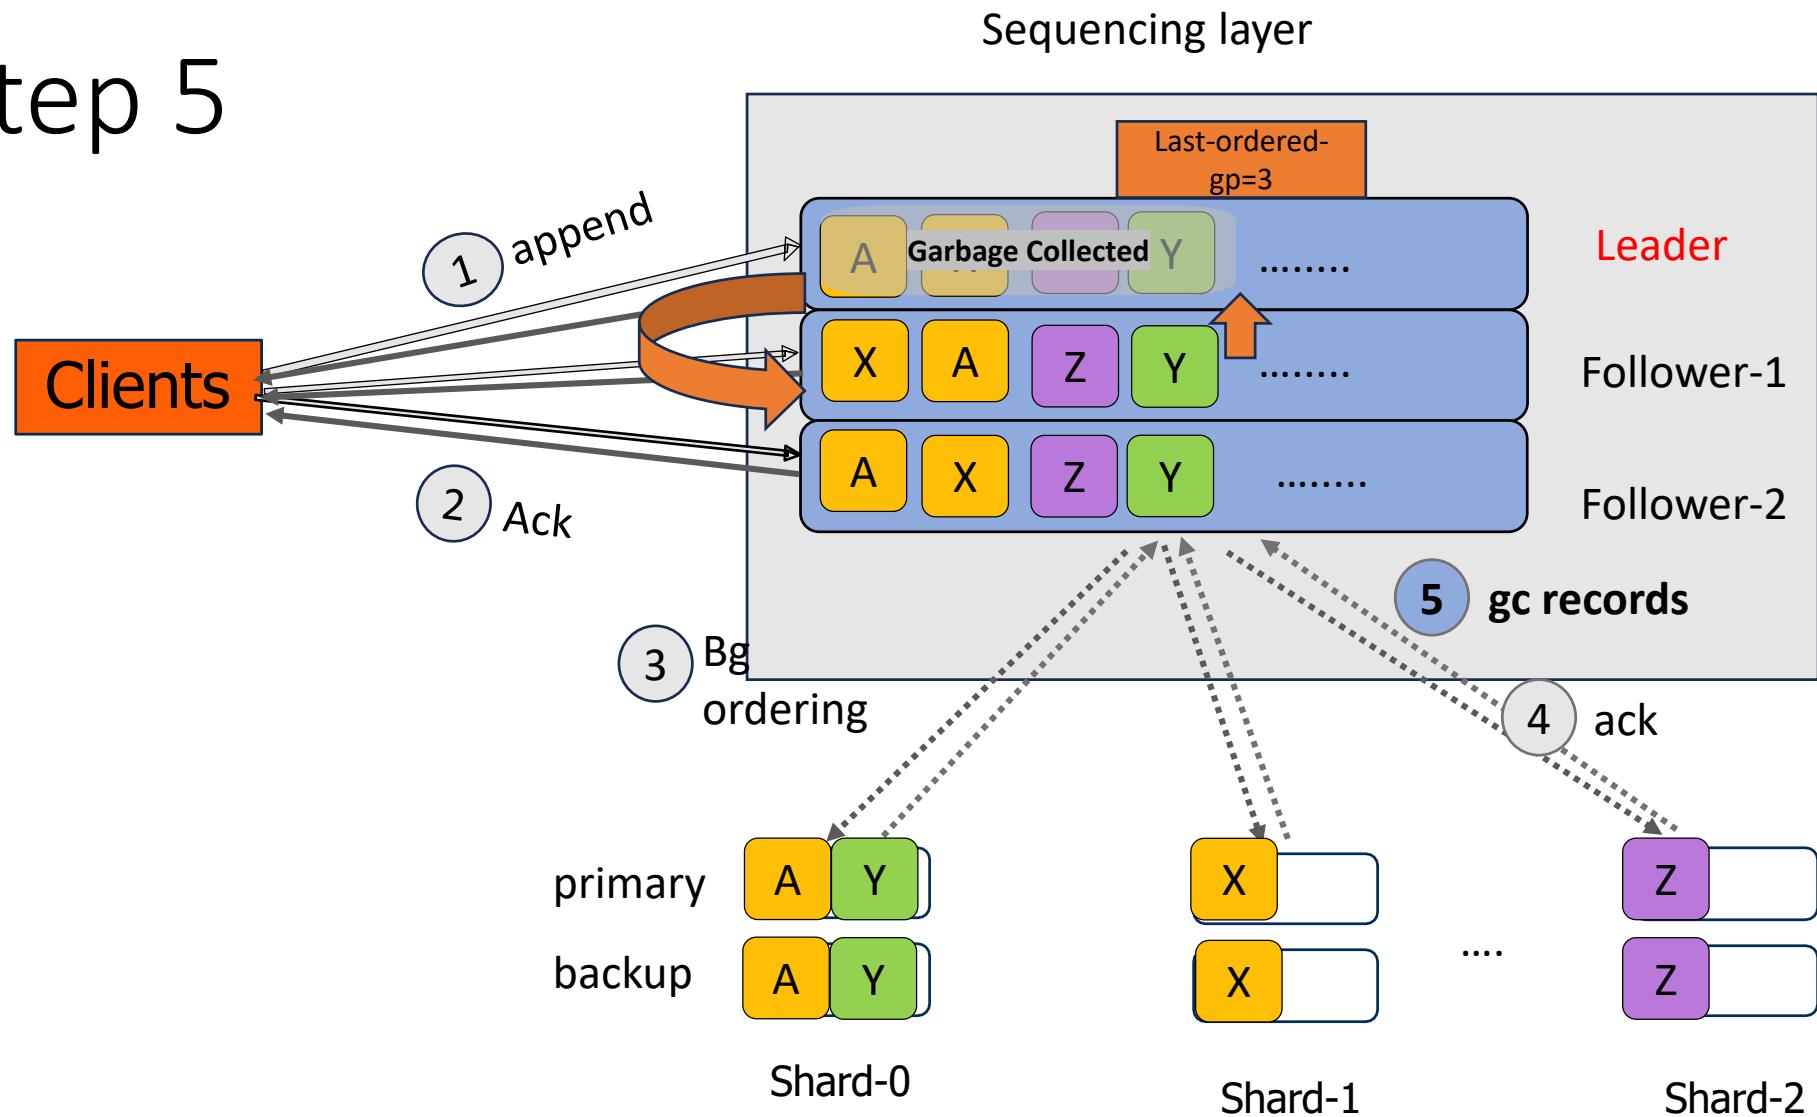
Step 5



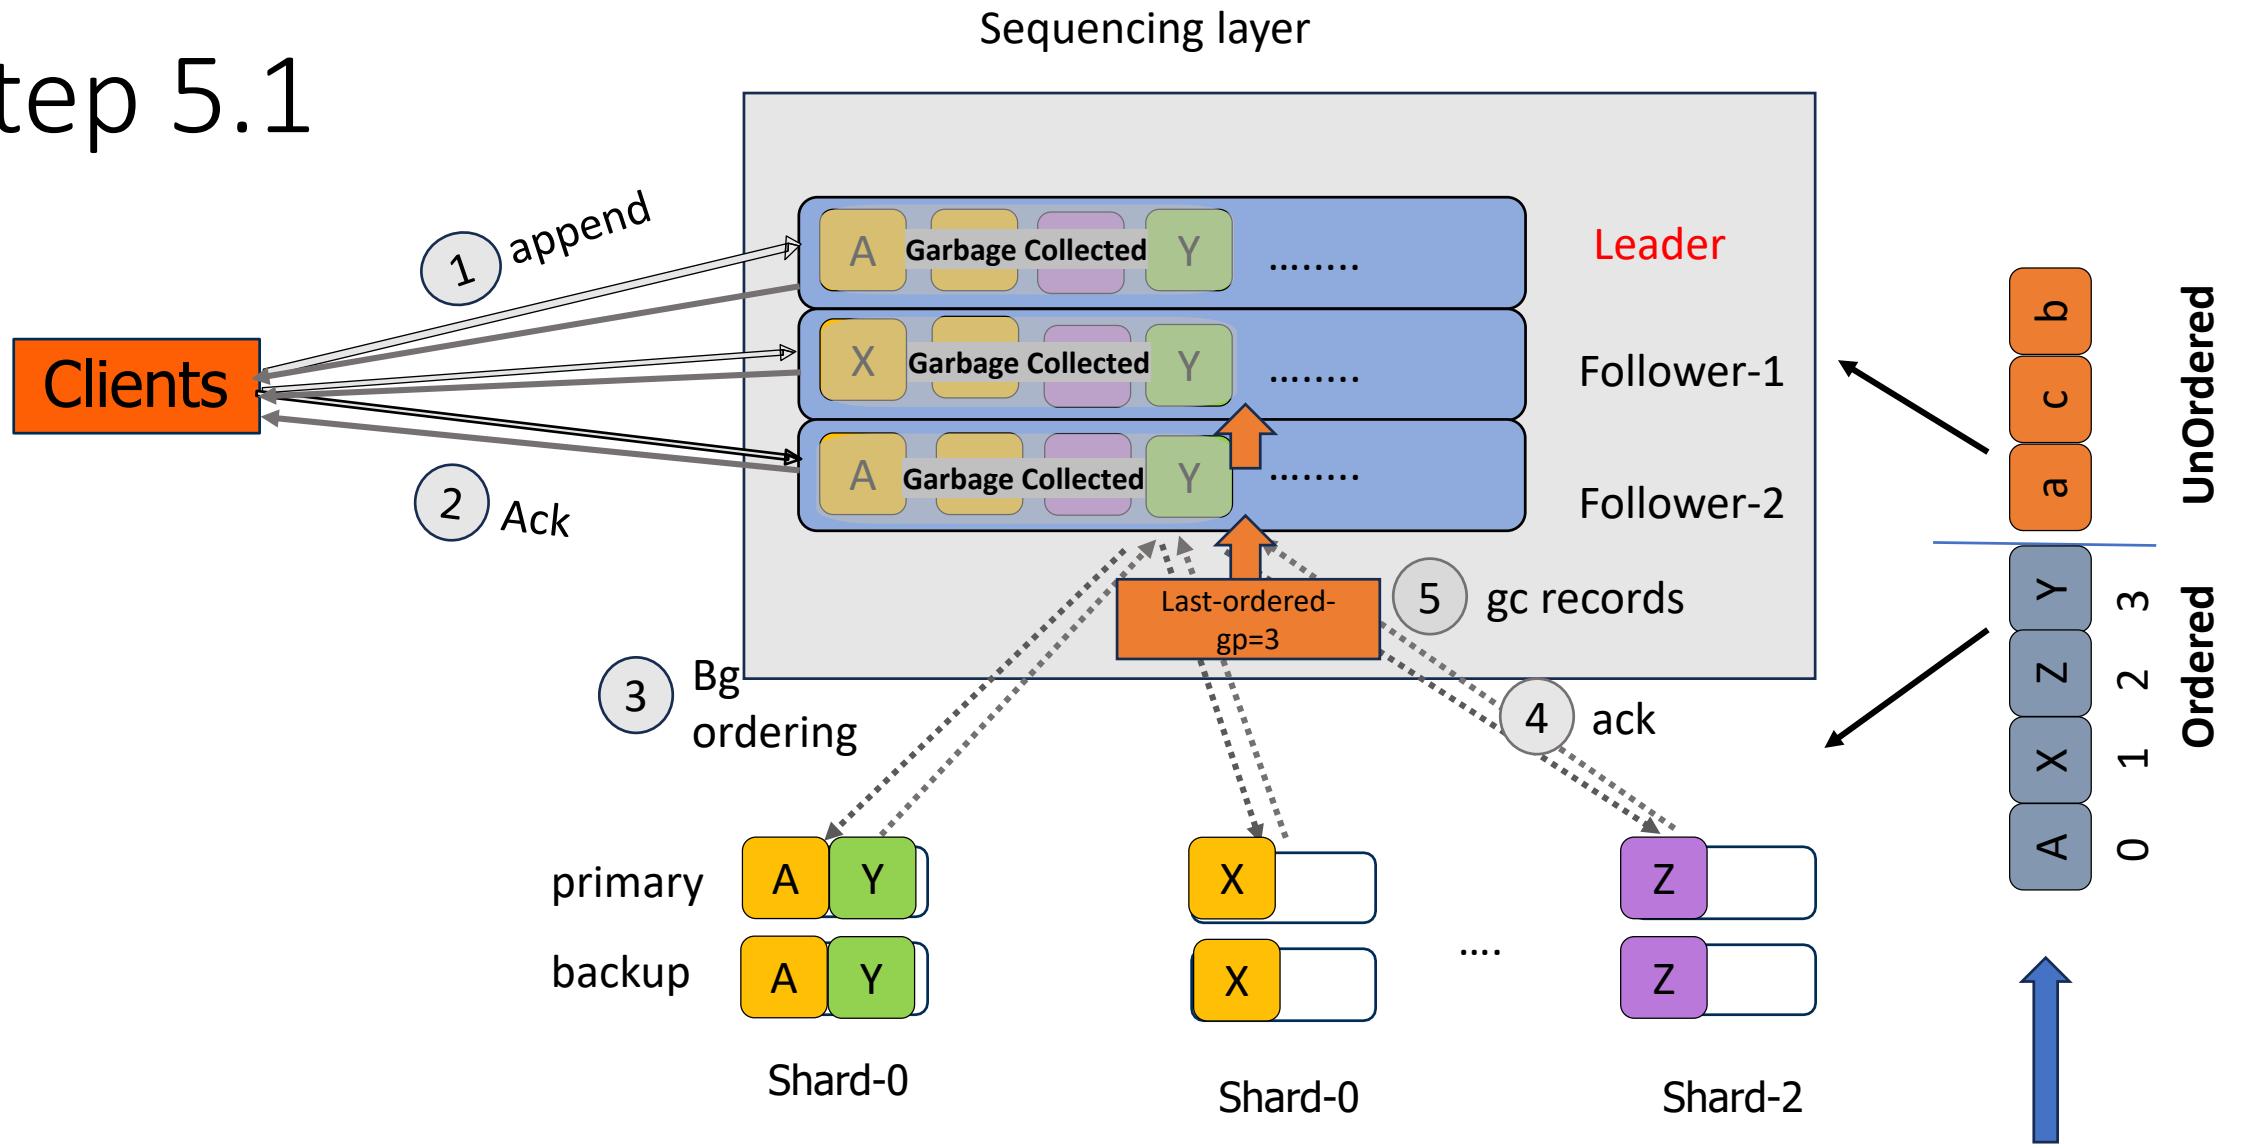
Step 5



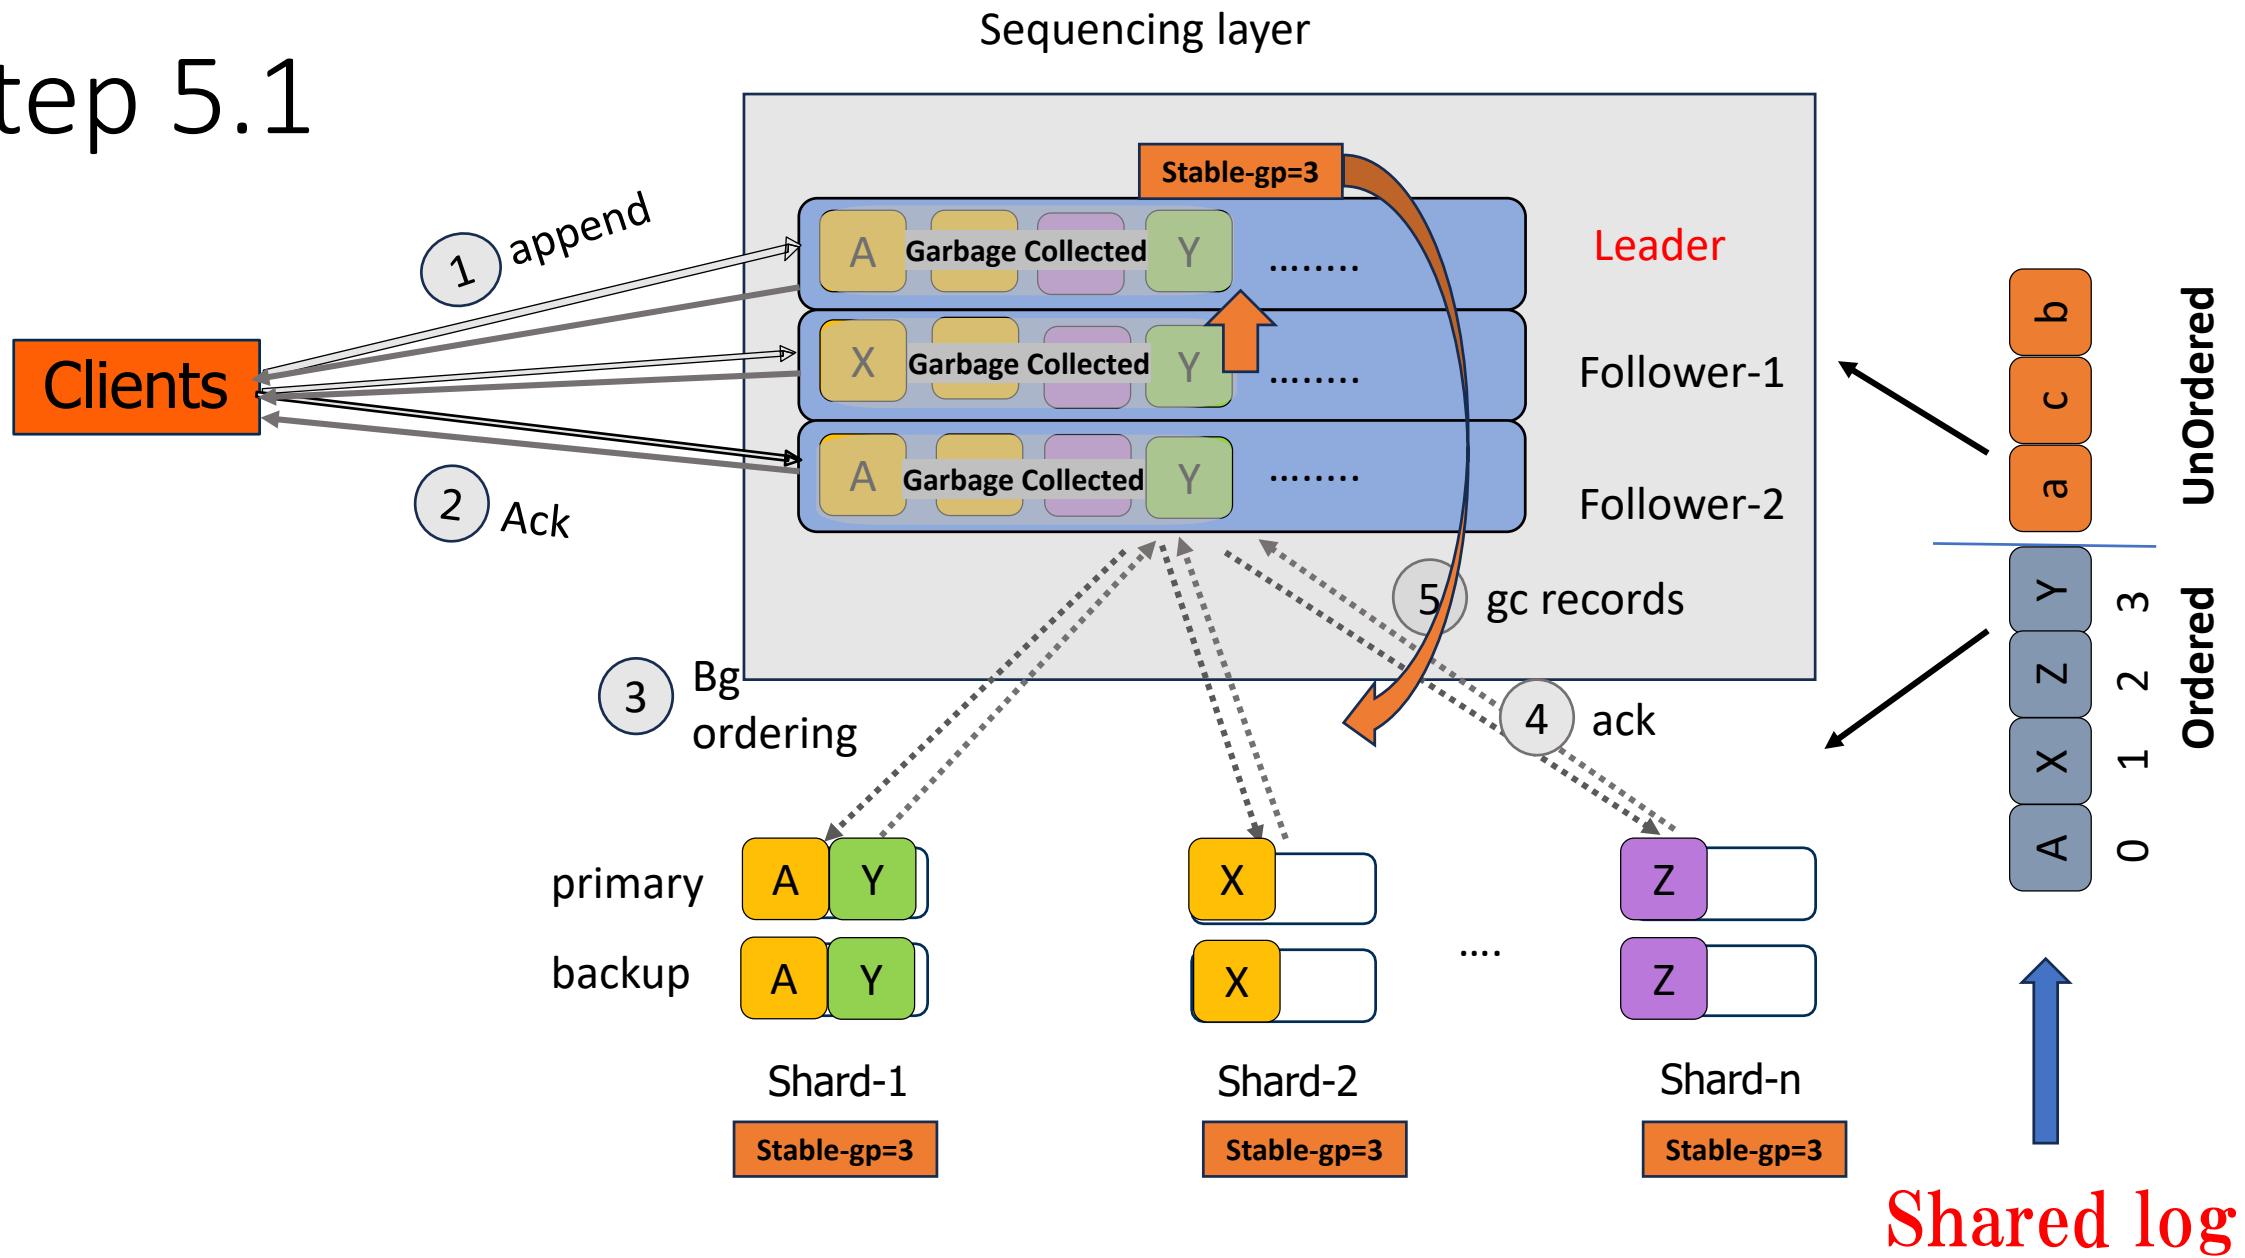
Step 5



Step 5.1



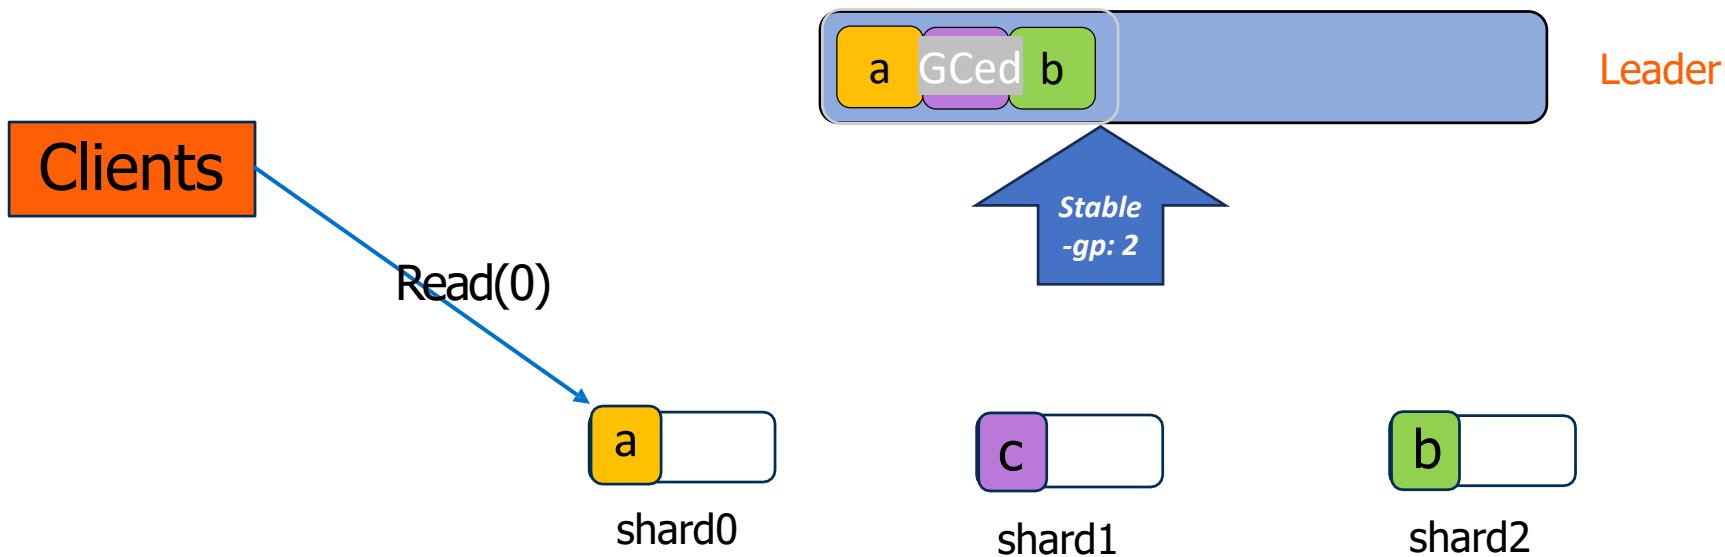
Step 5.1



How does Erwin Read work?

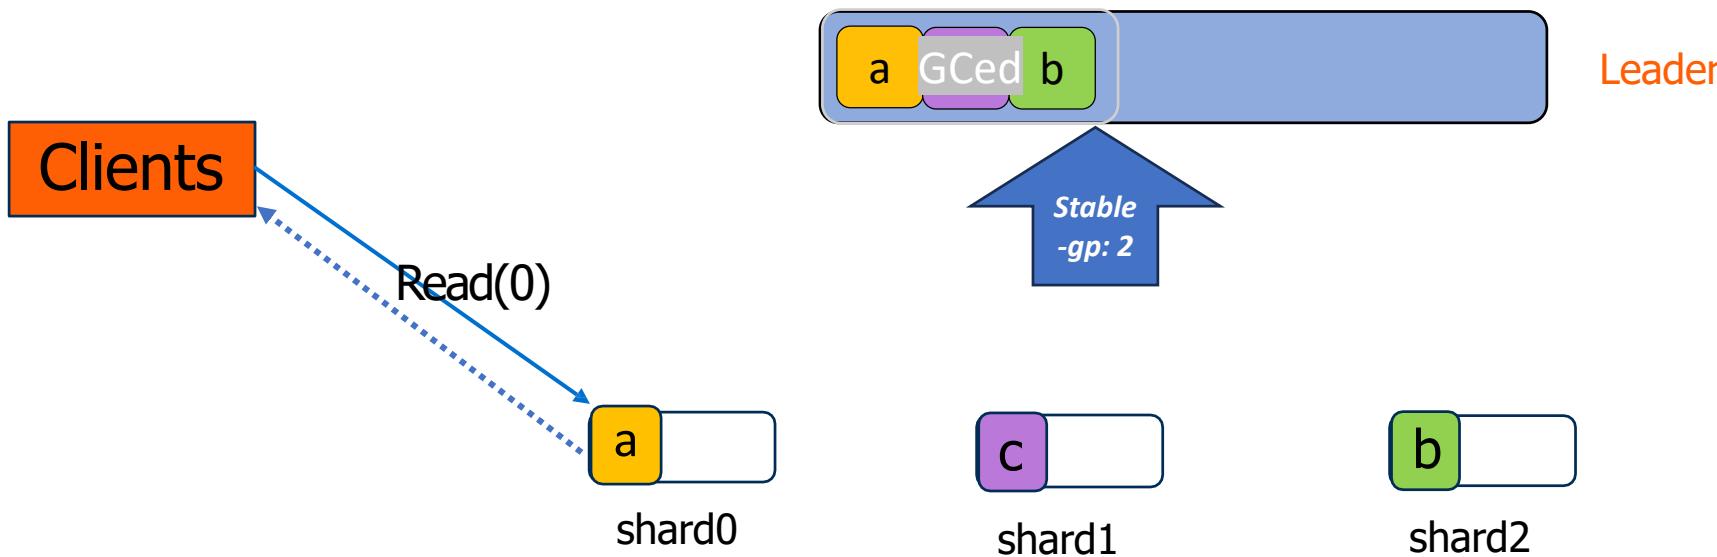
Step by step

Erwin: Read



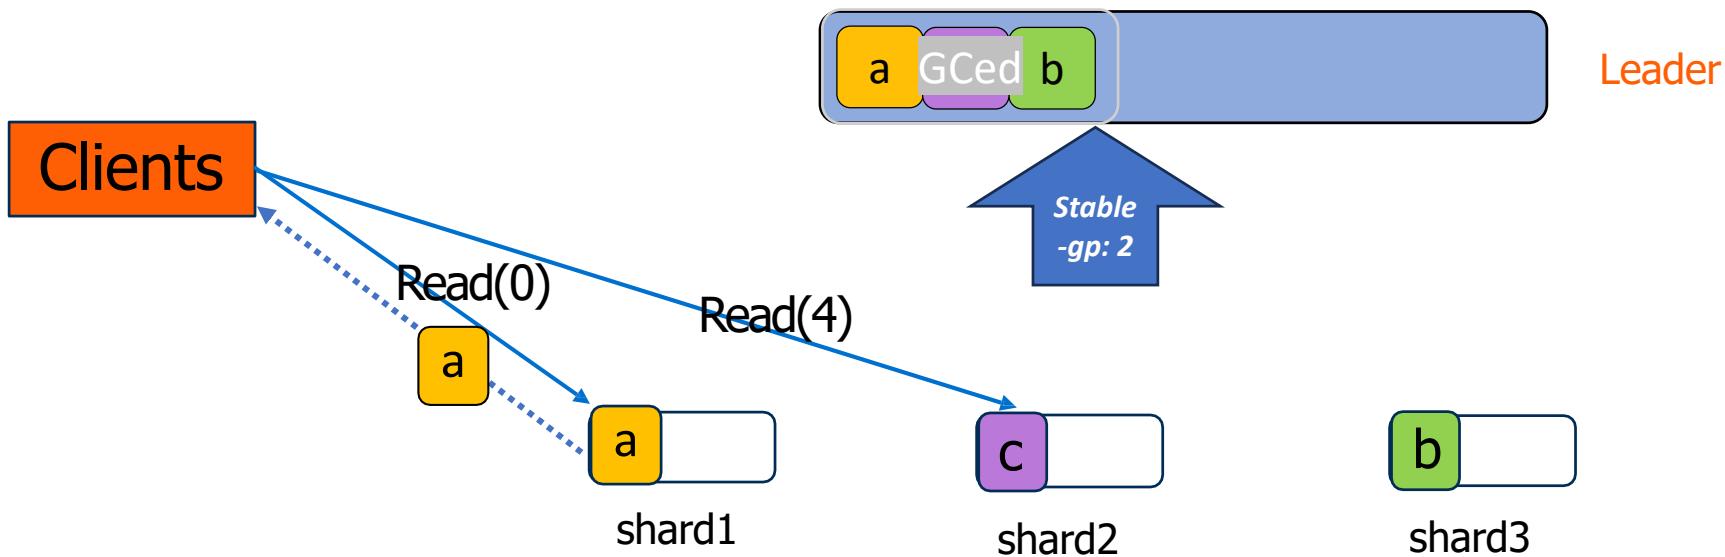
- Reading ordered position (fast read): entry returned directly

Erwin: Read



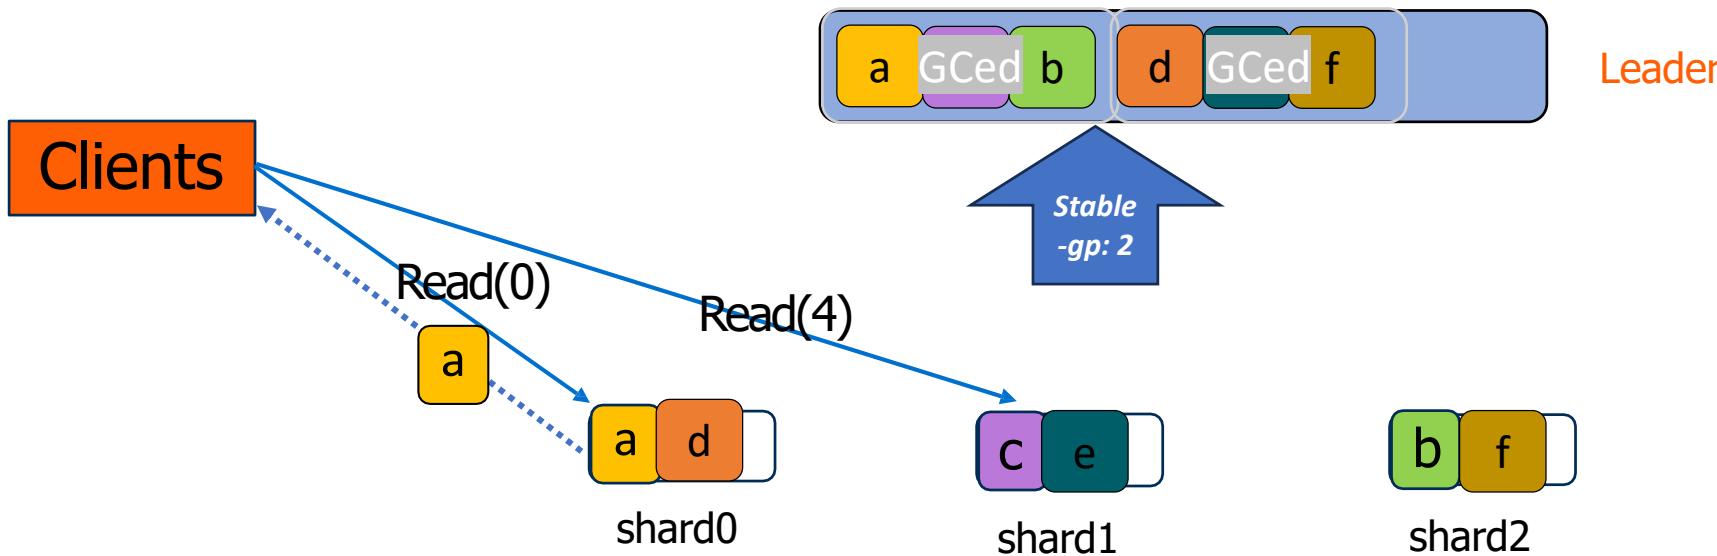
- Reading ordered position (fast read): entry returned directly

Erwin: Read



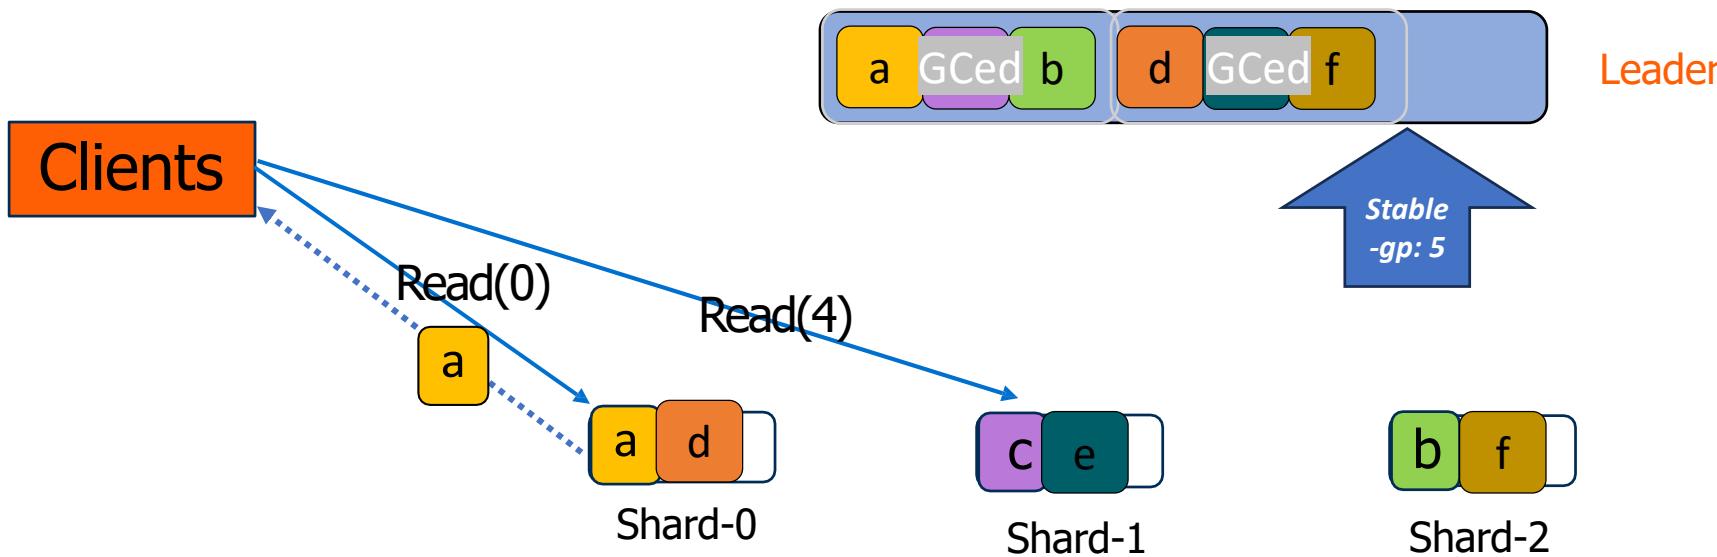
- Reading ordered position (fast read): entry returned directly
- Reading unordered position (slow read): must wait until **stable-gp** is advanced to the read position

Erwin: Read



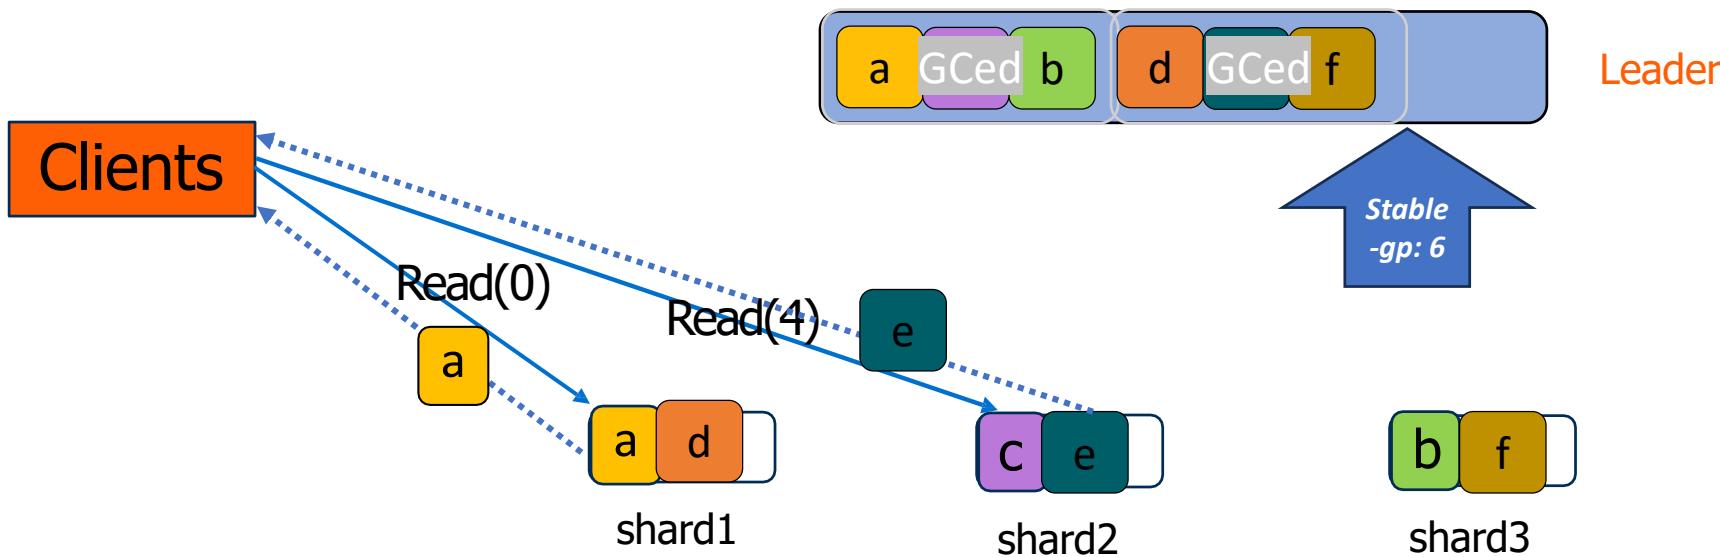
- Reading ordered position (fast read): entry returned directly
- Reading unordered position (slow read): must wait until **stable-gp** is advanced to the read position

Erwin: Read



- Reading ordered position (fast read): entry returned directly
- Reading unordered position (slow read): must wait until **stable-gp** is advanced to the read position

Erwin: Read



- Reading ordered position (fast read): entry returned directly
- Reading unordered position (slow read): must wait until **stable-gp** is advanced to the read position

Problem with Erwin Blackbox

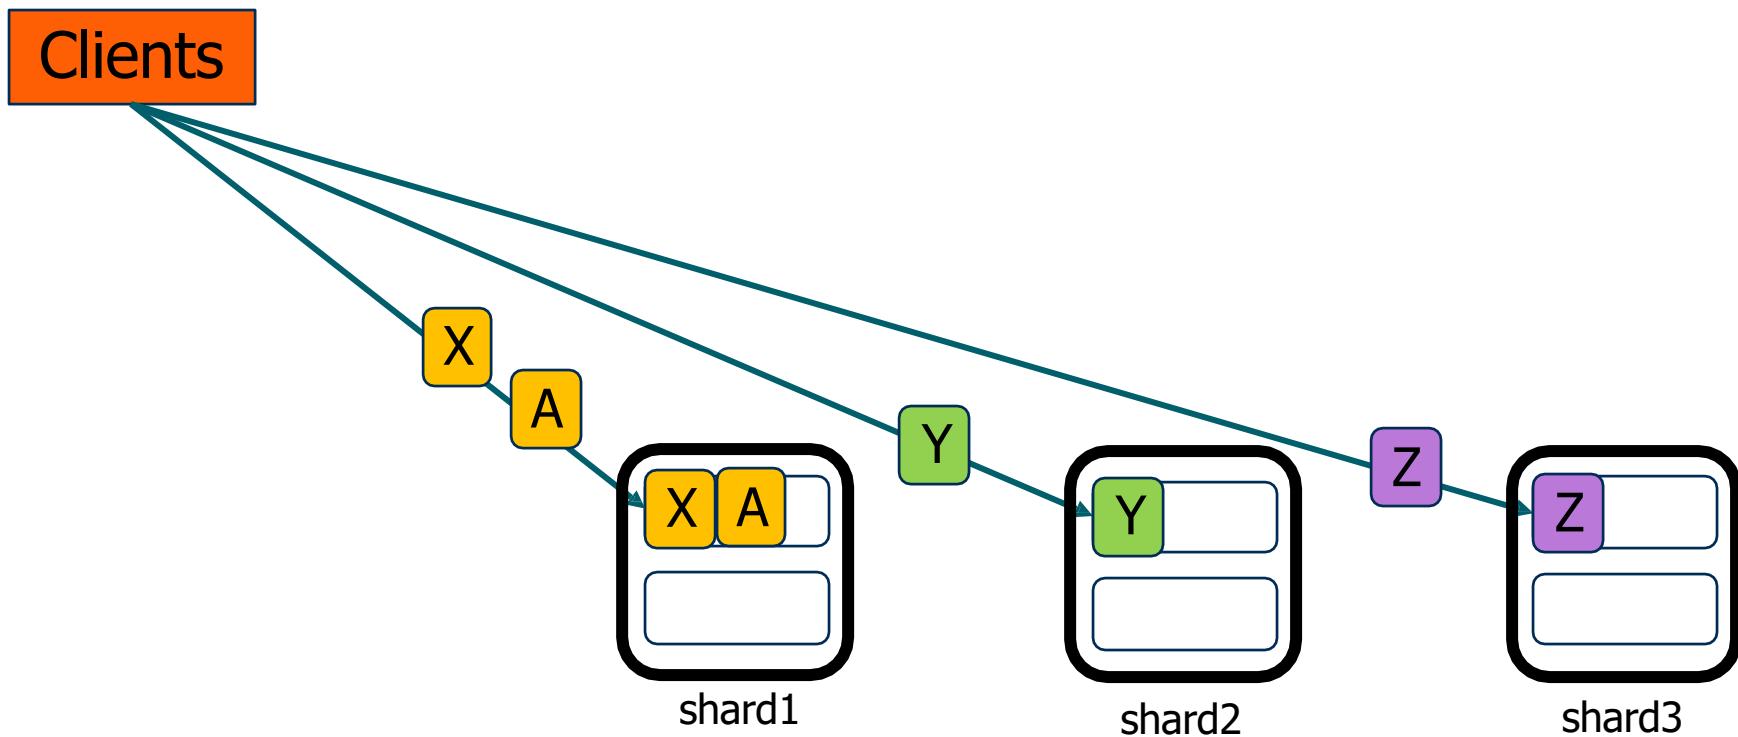
Large Record --> Sequencing layer will be saturated

Solution

Split Record --> into Data and
Metadata (Record_Id, Shard_Id)



Record

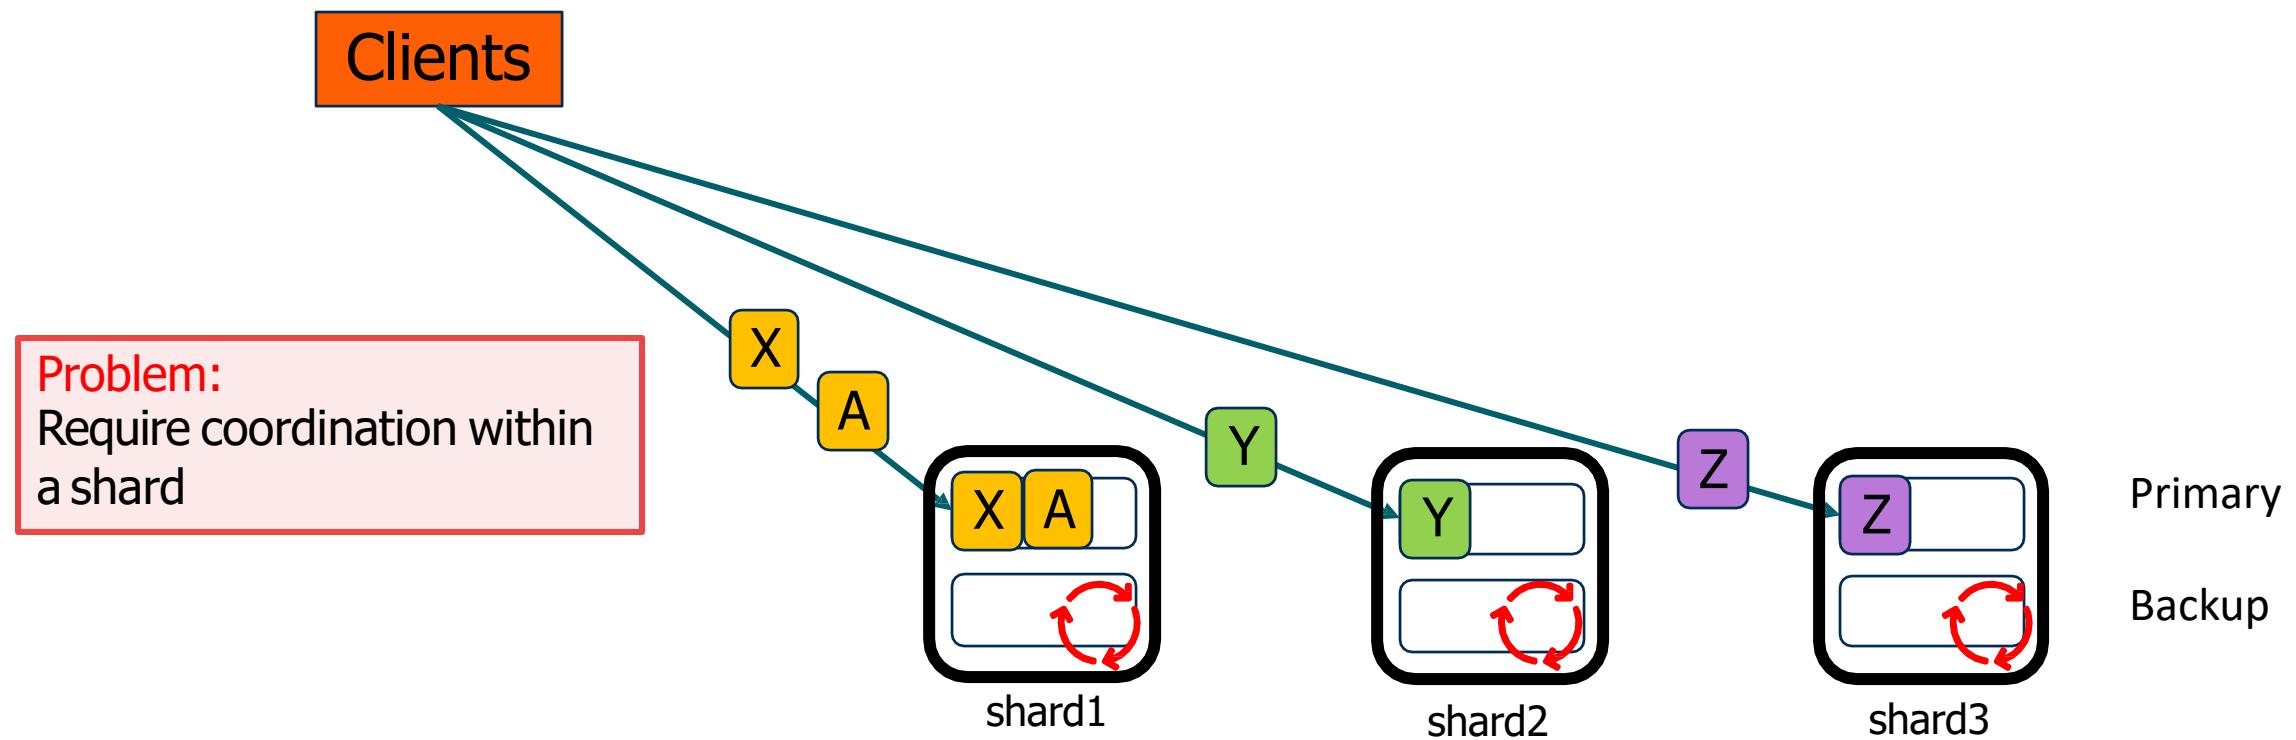


Primary

Backup



Record





Record

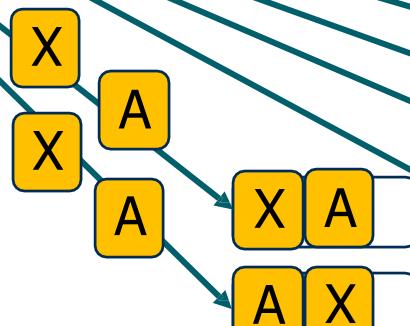
Clients

Problem:

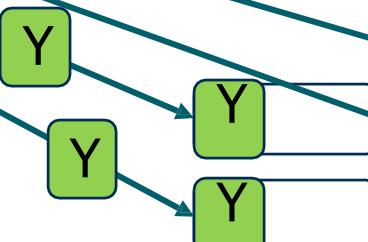
Require coordination within
a shard

Solution:

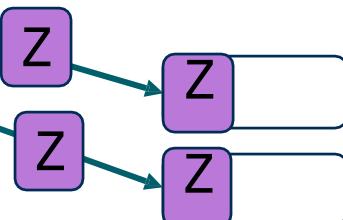
Send record to all replicas in the
shard



shard1



shard2



shard3

Primary

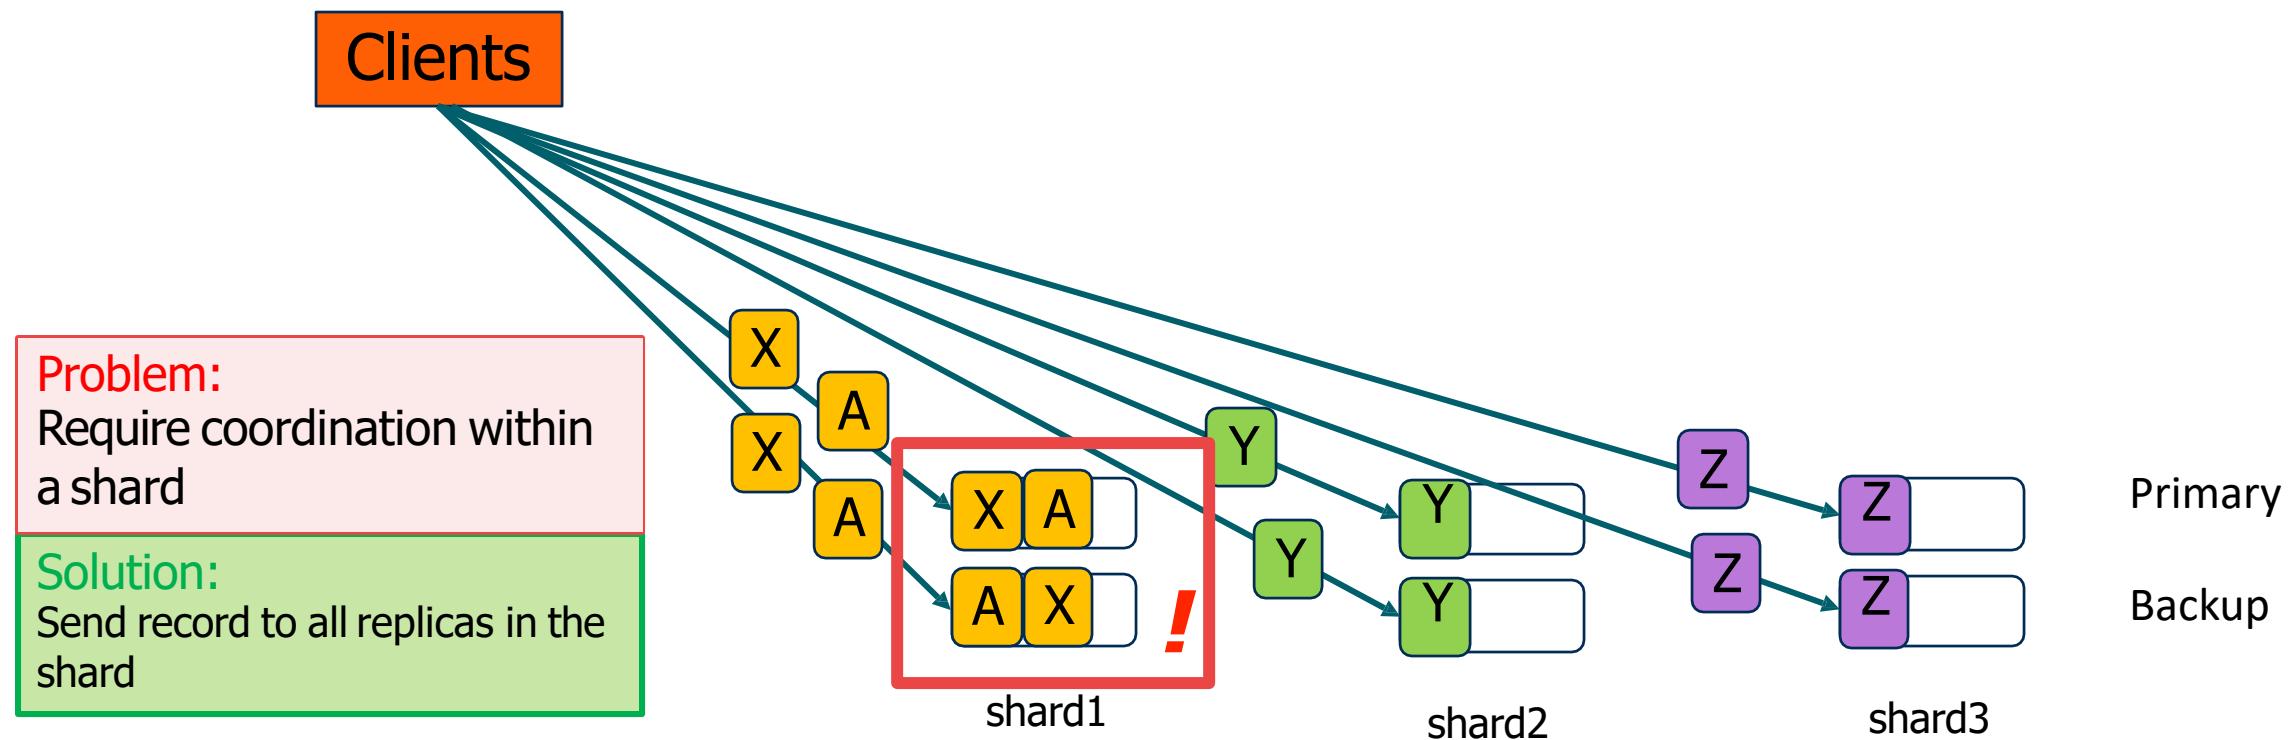
Backup

Motivation



Record

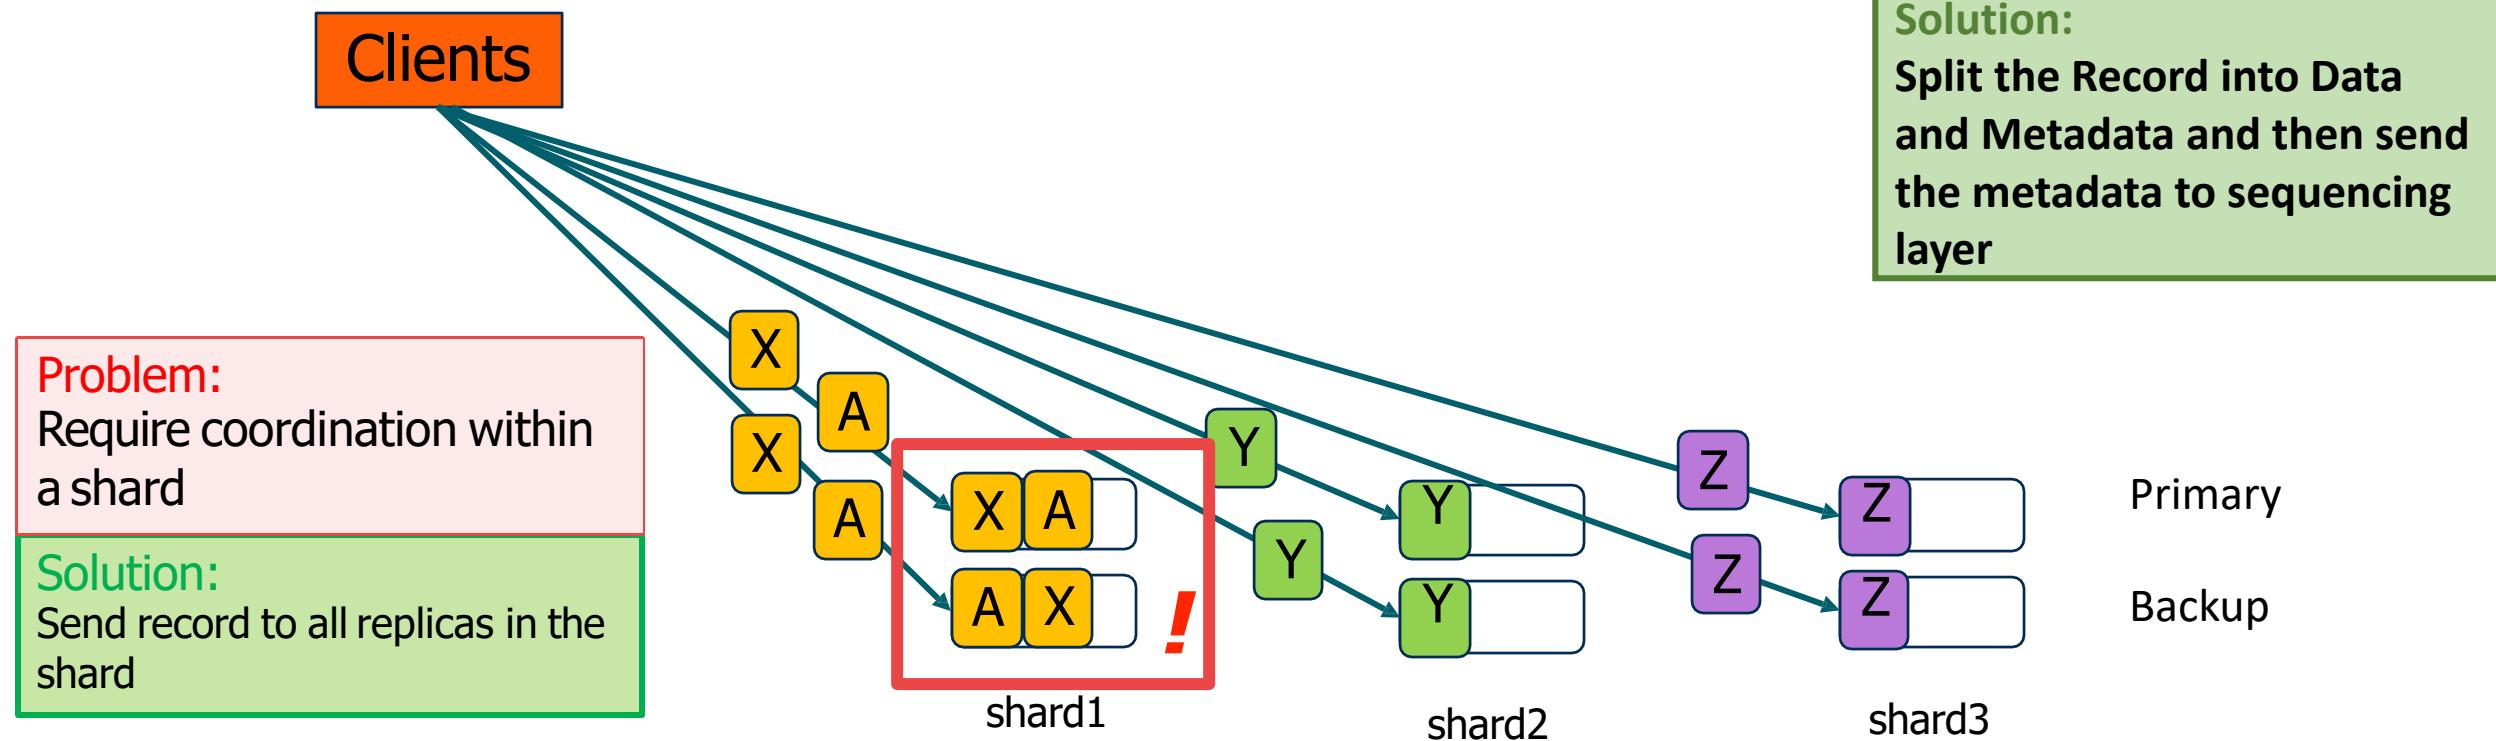
Problem:
No order across and within shards



Motivation

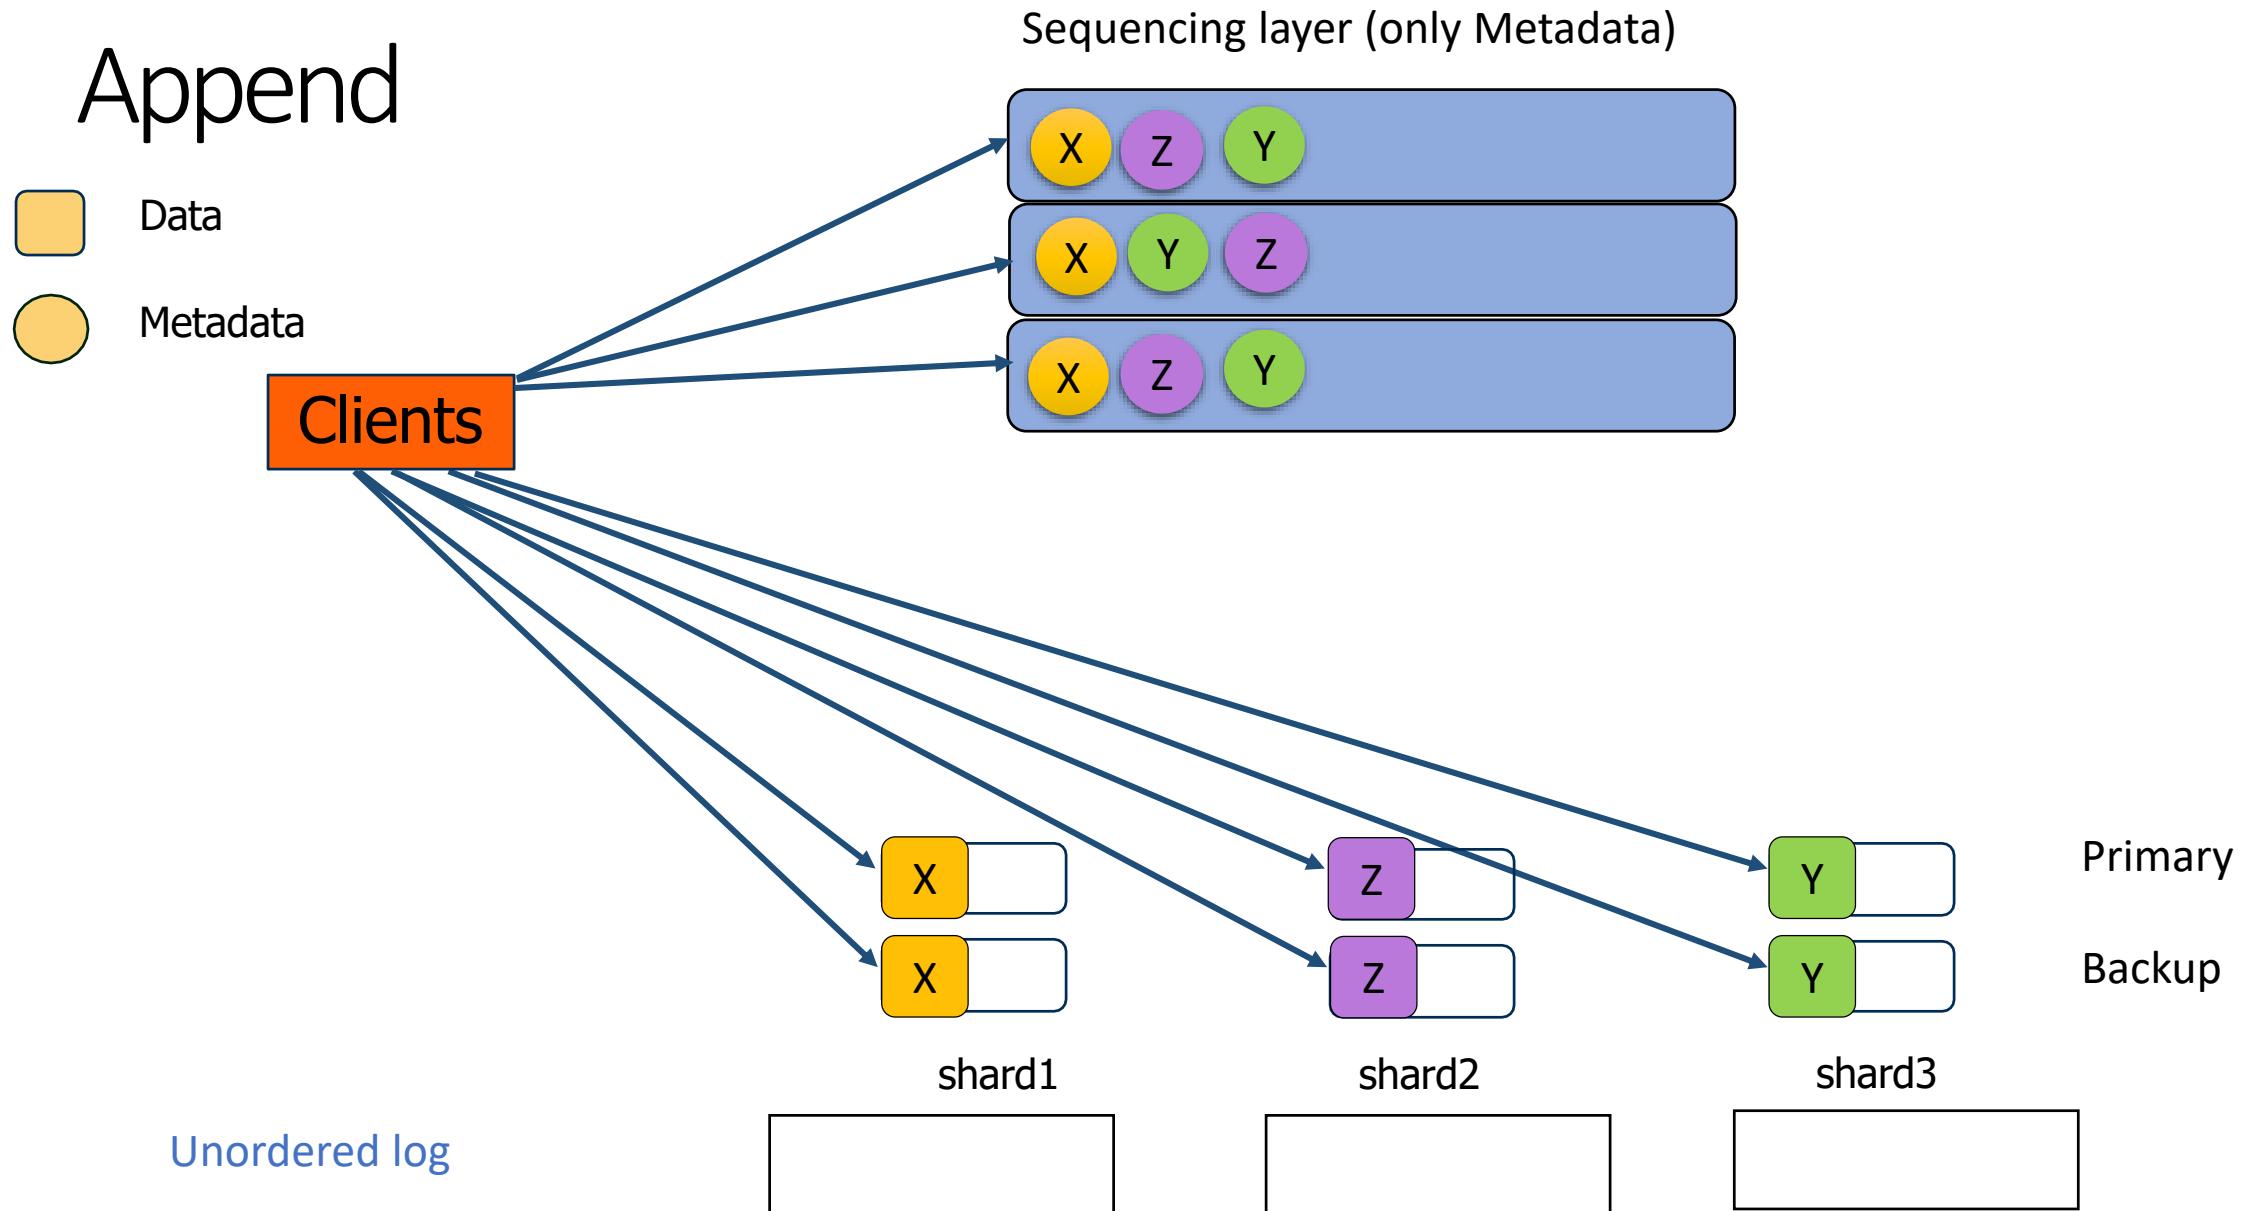


Record

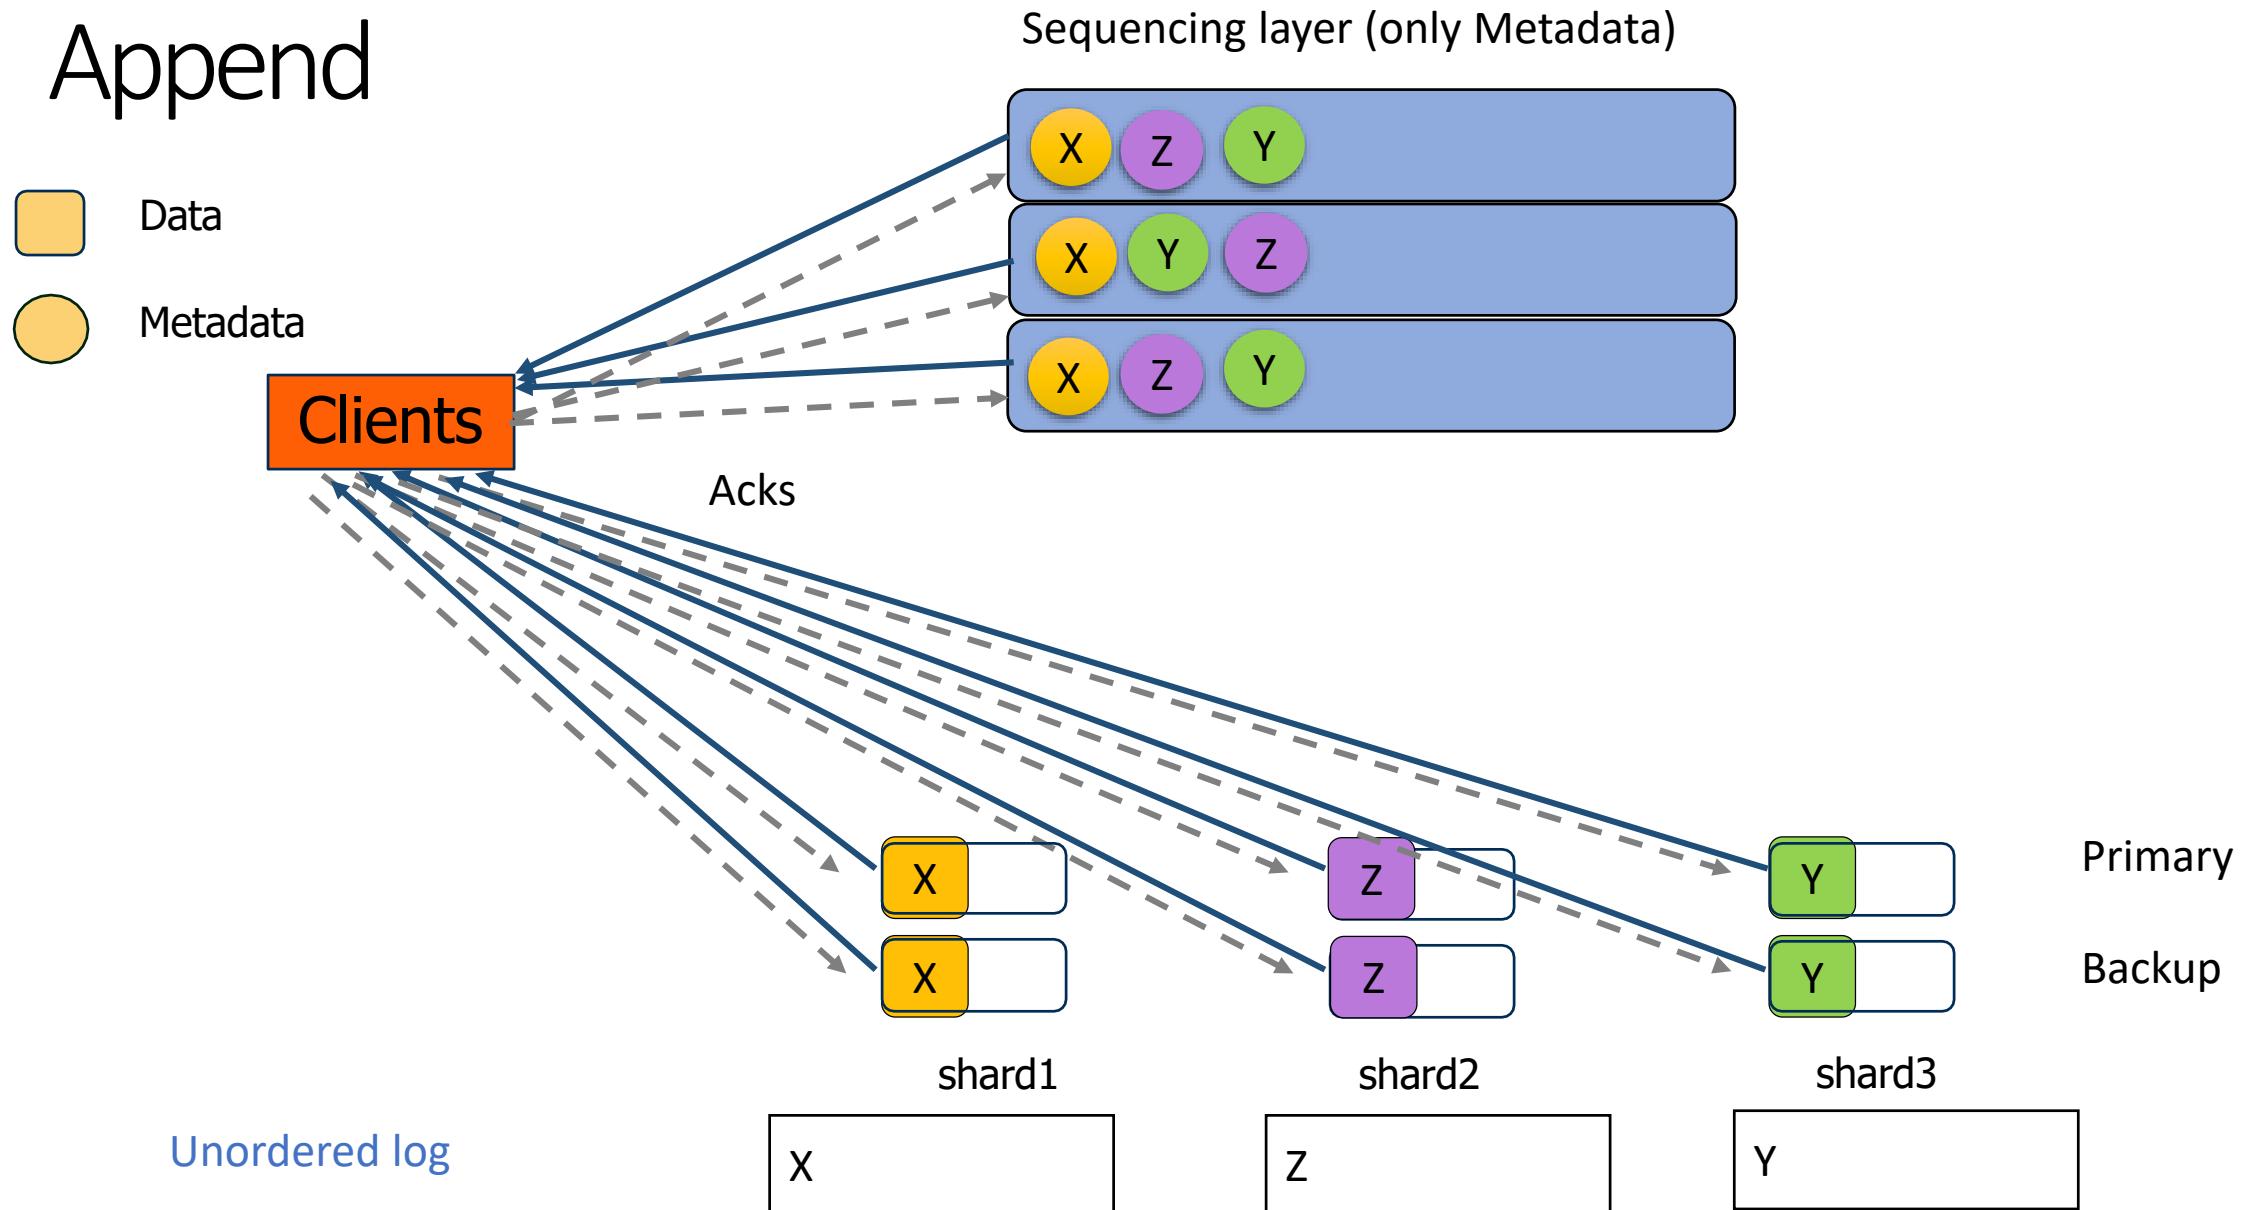


Erwin-st(scalable throughput)

Erwin's Goal: 1-RTT Append

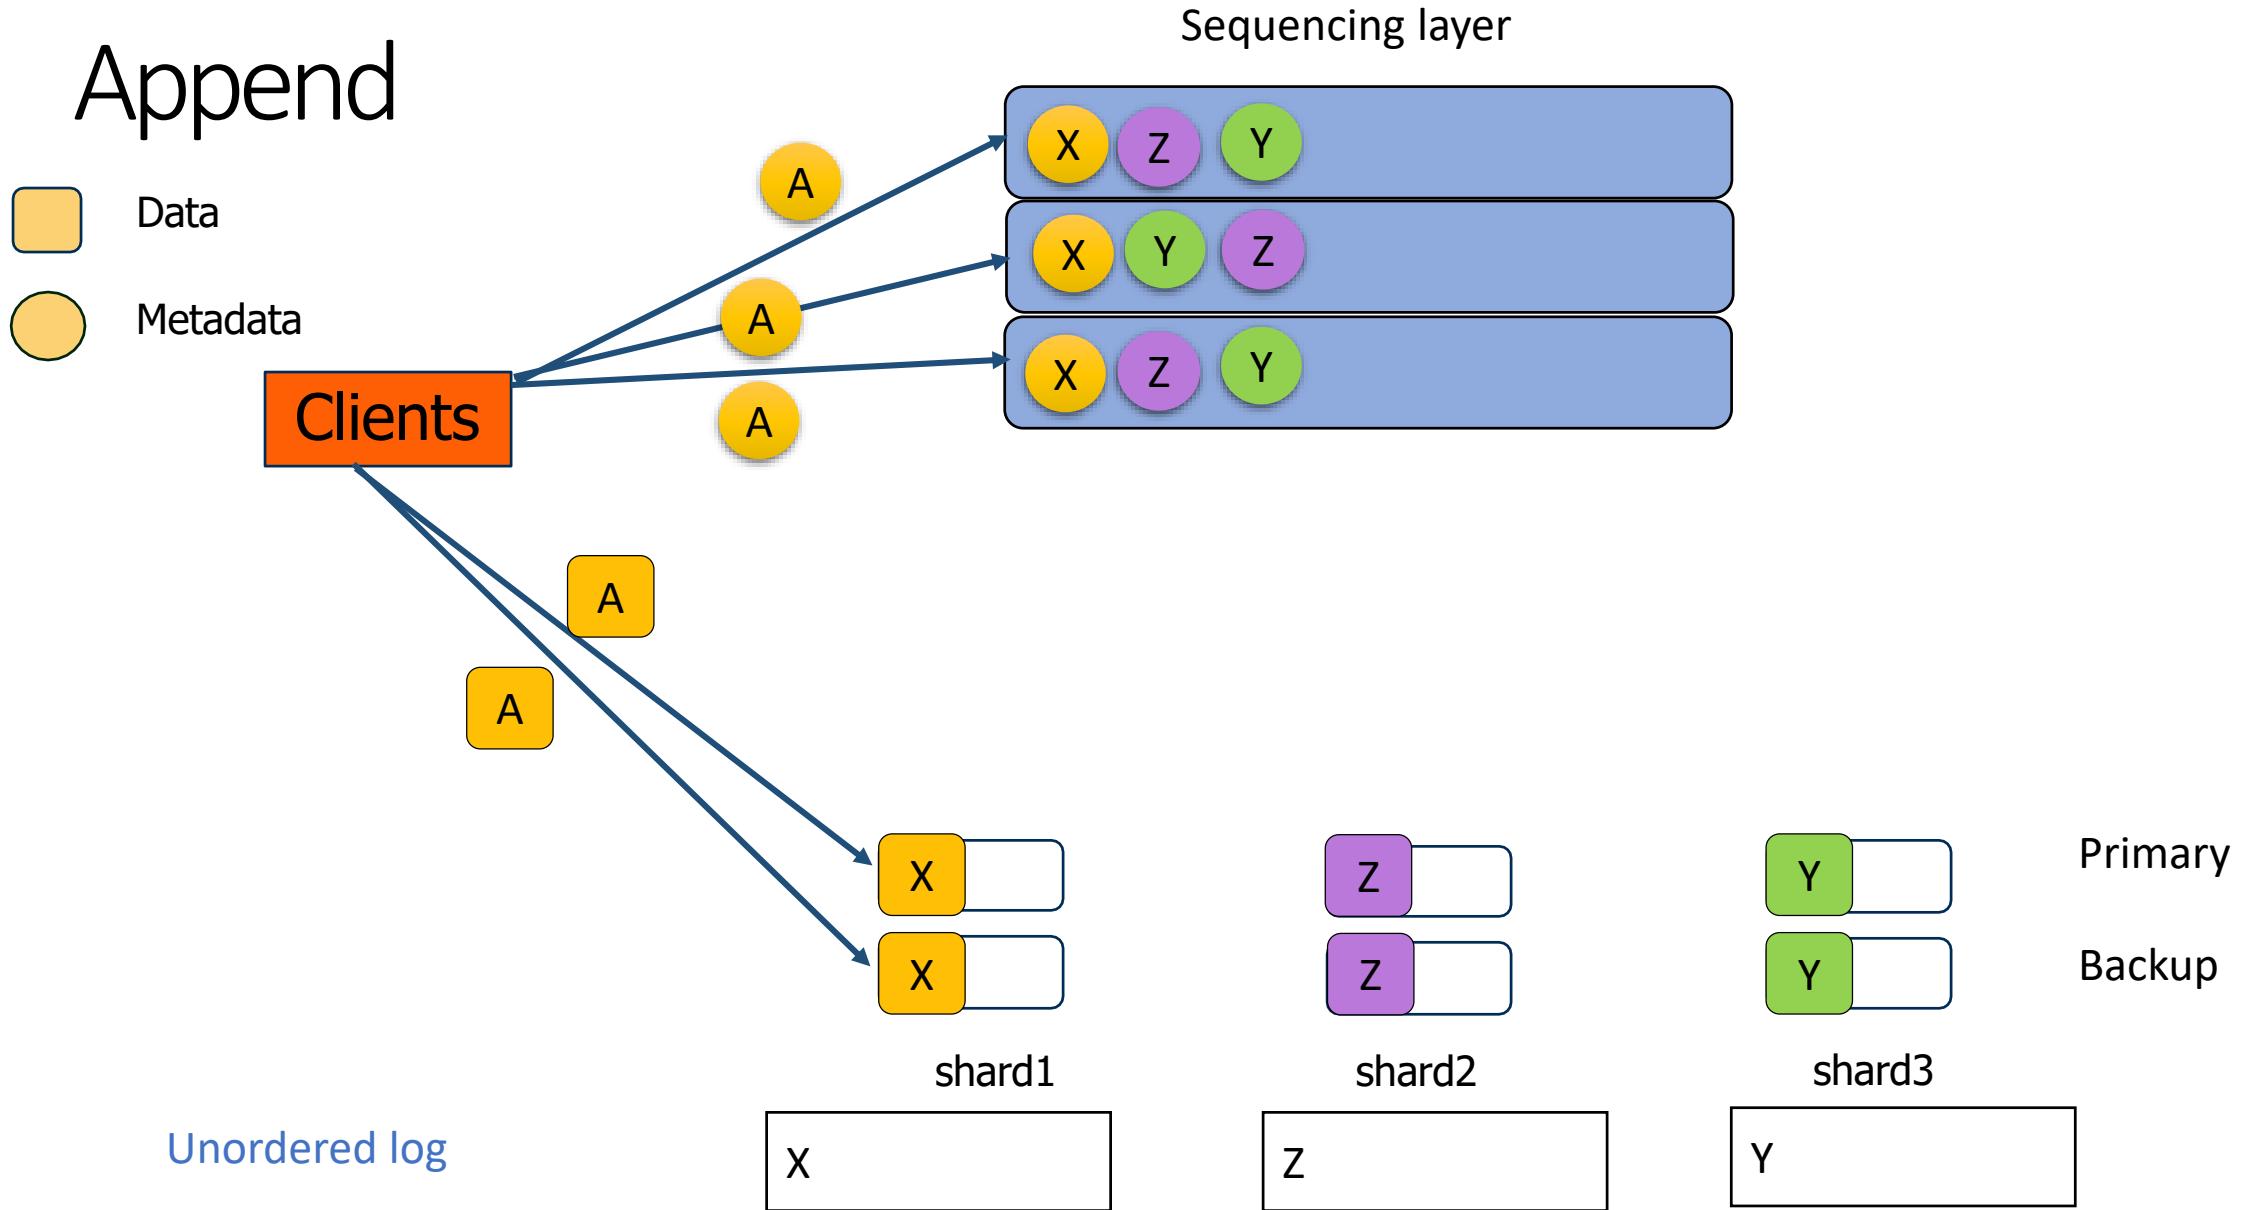


Erwin's Goal: 1-RTT Append



Erwin's Goal: 1-RTT

Append



Erwin's Goal: 1-RTT

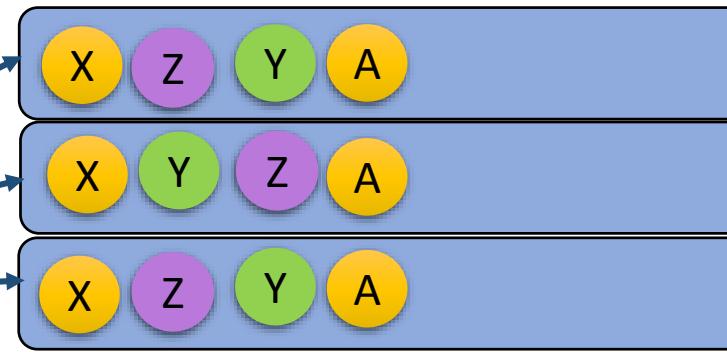
Append

Data

Metadata

Clients

Sequencing layer



X A

X A

shard1

Z

Z

shard2

Y

Y

shard3

Primary

Backup

Unordered log

X, A

Z

Y

Erwin's Goal: 1-RTT

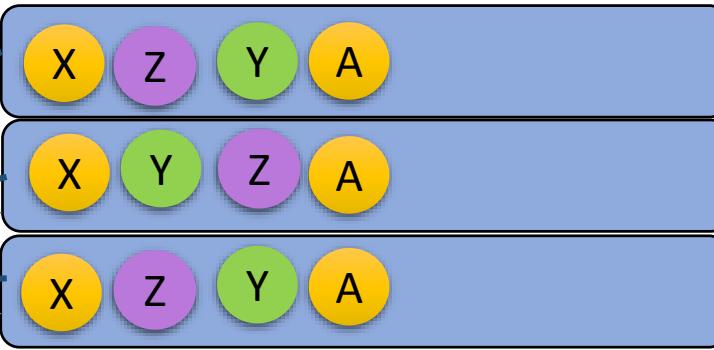
Append

Data

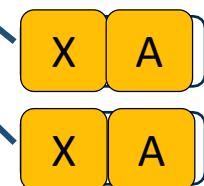
Metadata

Clients

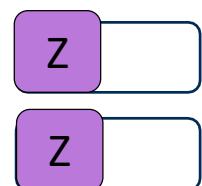
Sequencing layer



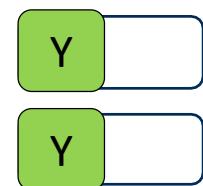
Acks



shard1



shard2



shard3

Unordered log

X , A

Z

Y

Primary

Backup

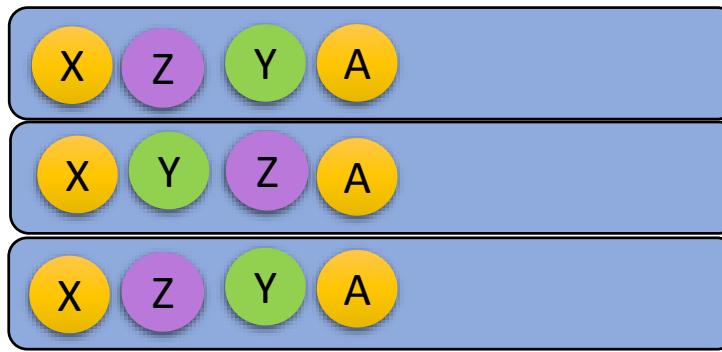
Erwin's Goal: 1-RTT Append

Data

Metadata

Clients

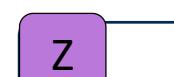
Sequencing layer



Sends position of
the record



shard1



shard2



shard3

Primary

Backup

Unordered log

X , A

Z

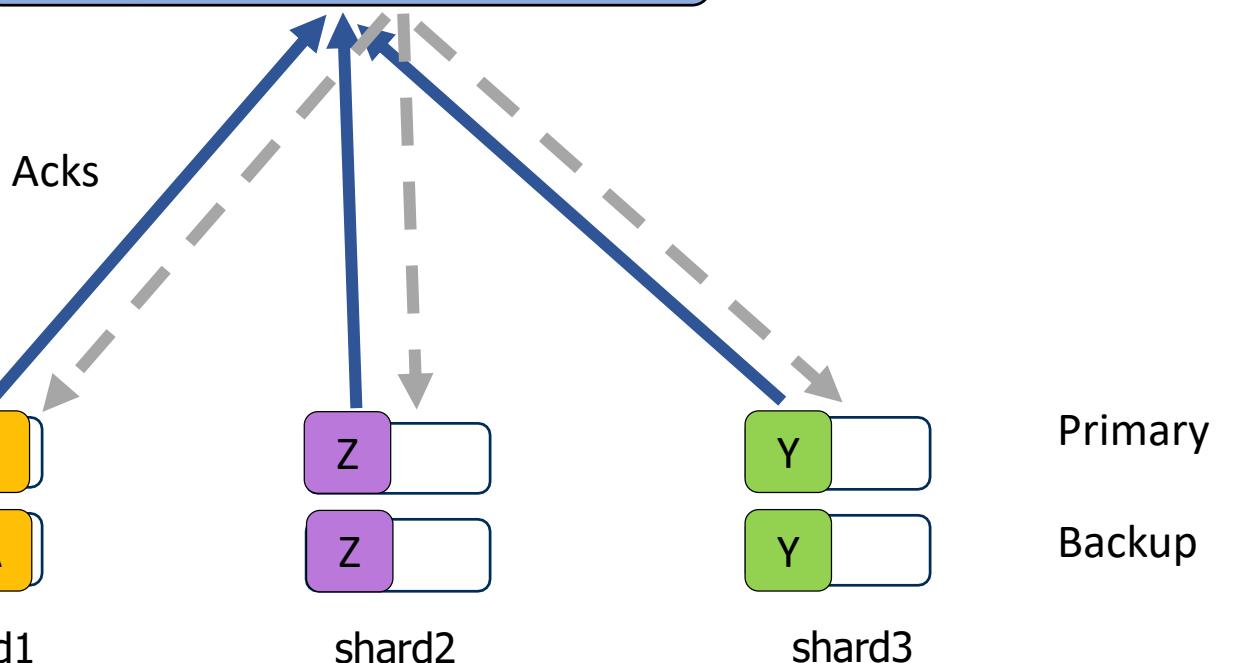
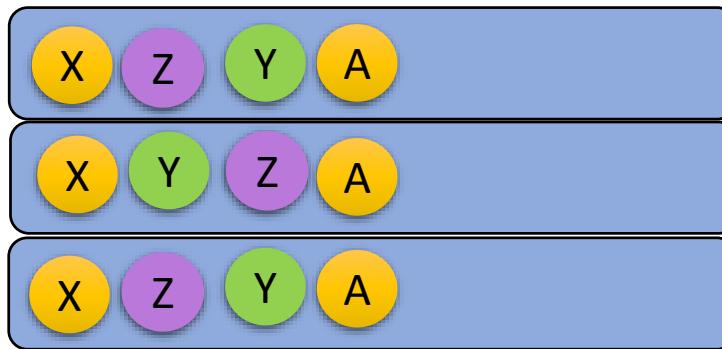
Y

Erwin's Goal: 1-RTT Append

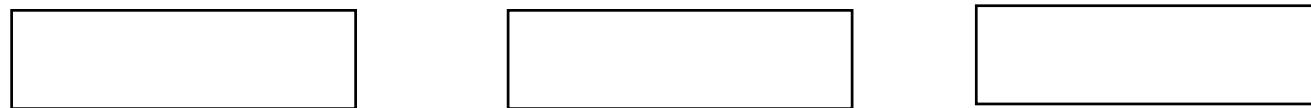
- Data
- Metadata

Clients

Sequencing layer



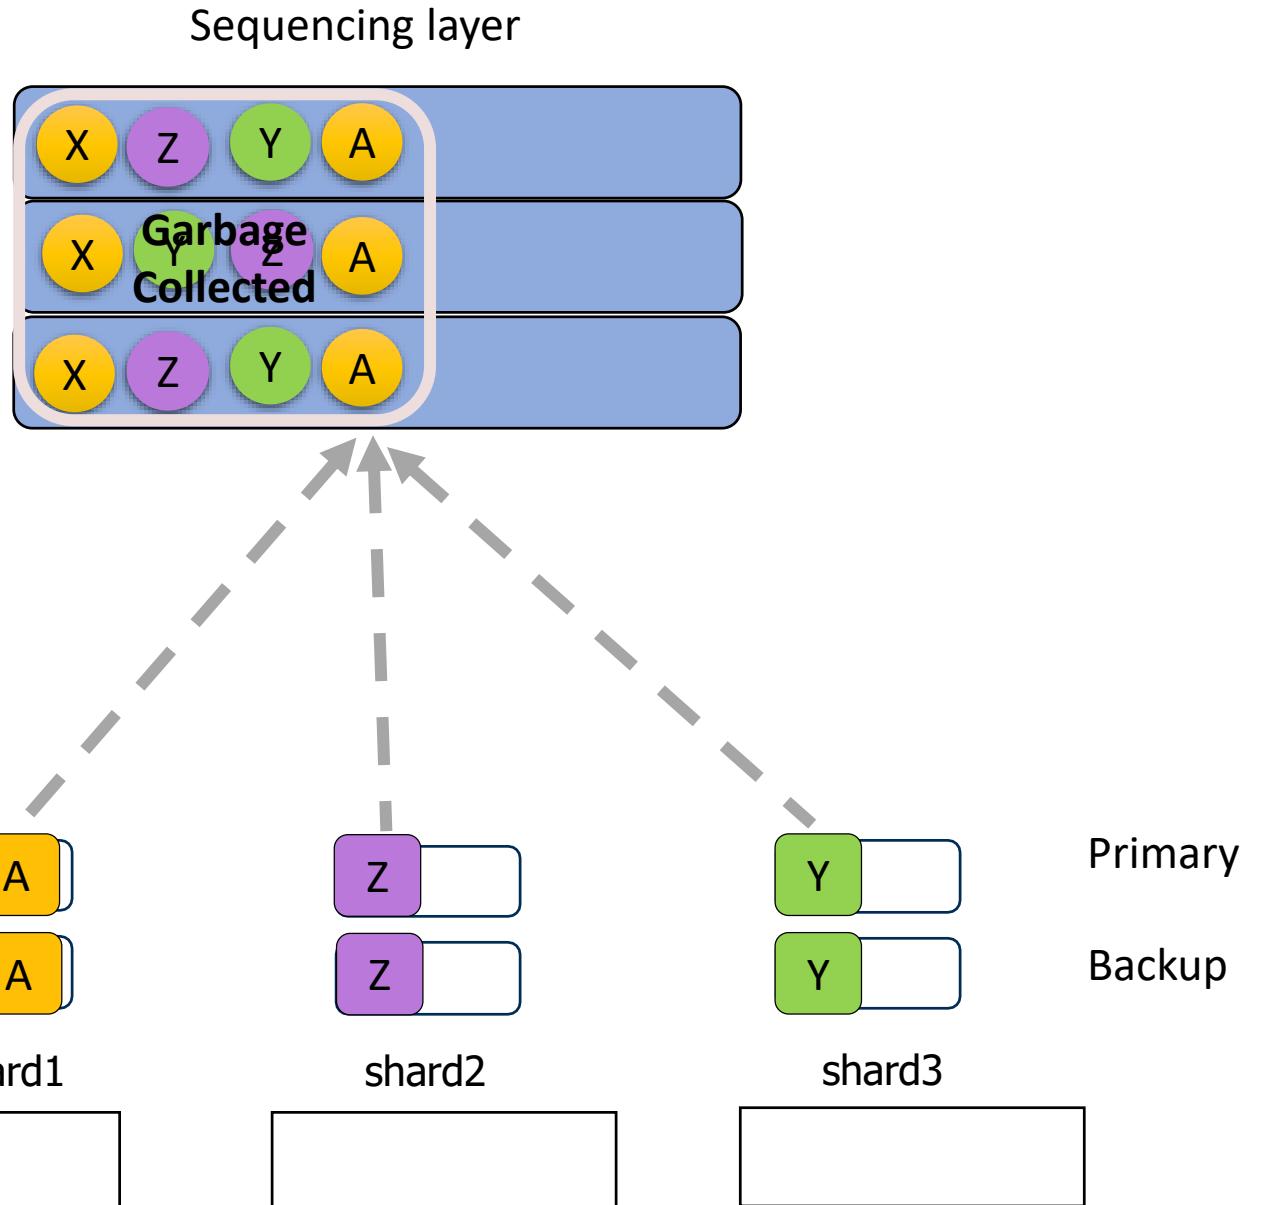
Unordered log



Erwin's Goal: 1-RTT Append

- Data
- Metadata

Clients



Failure Handling

Failure Handling

- Shard failure is handled within the shard- Paxos/Raft
- Network blips – clients can retry their append requests
- Sequencing layer replica failure- handled using **view and reconfiguration**

Step 1: Failure detection

- Using a control plane which has a zookeeper and a controller
- Every replica in the sequencing layer maintains heartbeat with the zookeeper and if it breaks then it informs the controller about the failed replica

Step 2: Start a new view

- Controller now seals the old **view**
- Any record sent with old view won't be appended
- Any of the other replica is assigned as recovery replica
- Flush records to shards starting from last-ordered-gp+1

Failure Handling Erwin St

- Sequencing replica failure
- Shard Failure

Same as Erwin Blackbox

Failure Handling

- Sequencing replica failure
- Shard Failure

Same as Erwin Blackbox

Client Failures?

Failure Handling

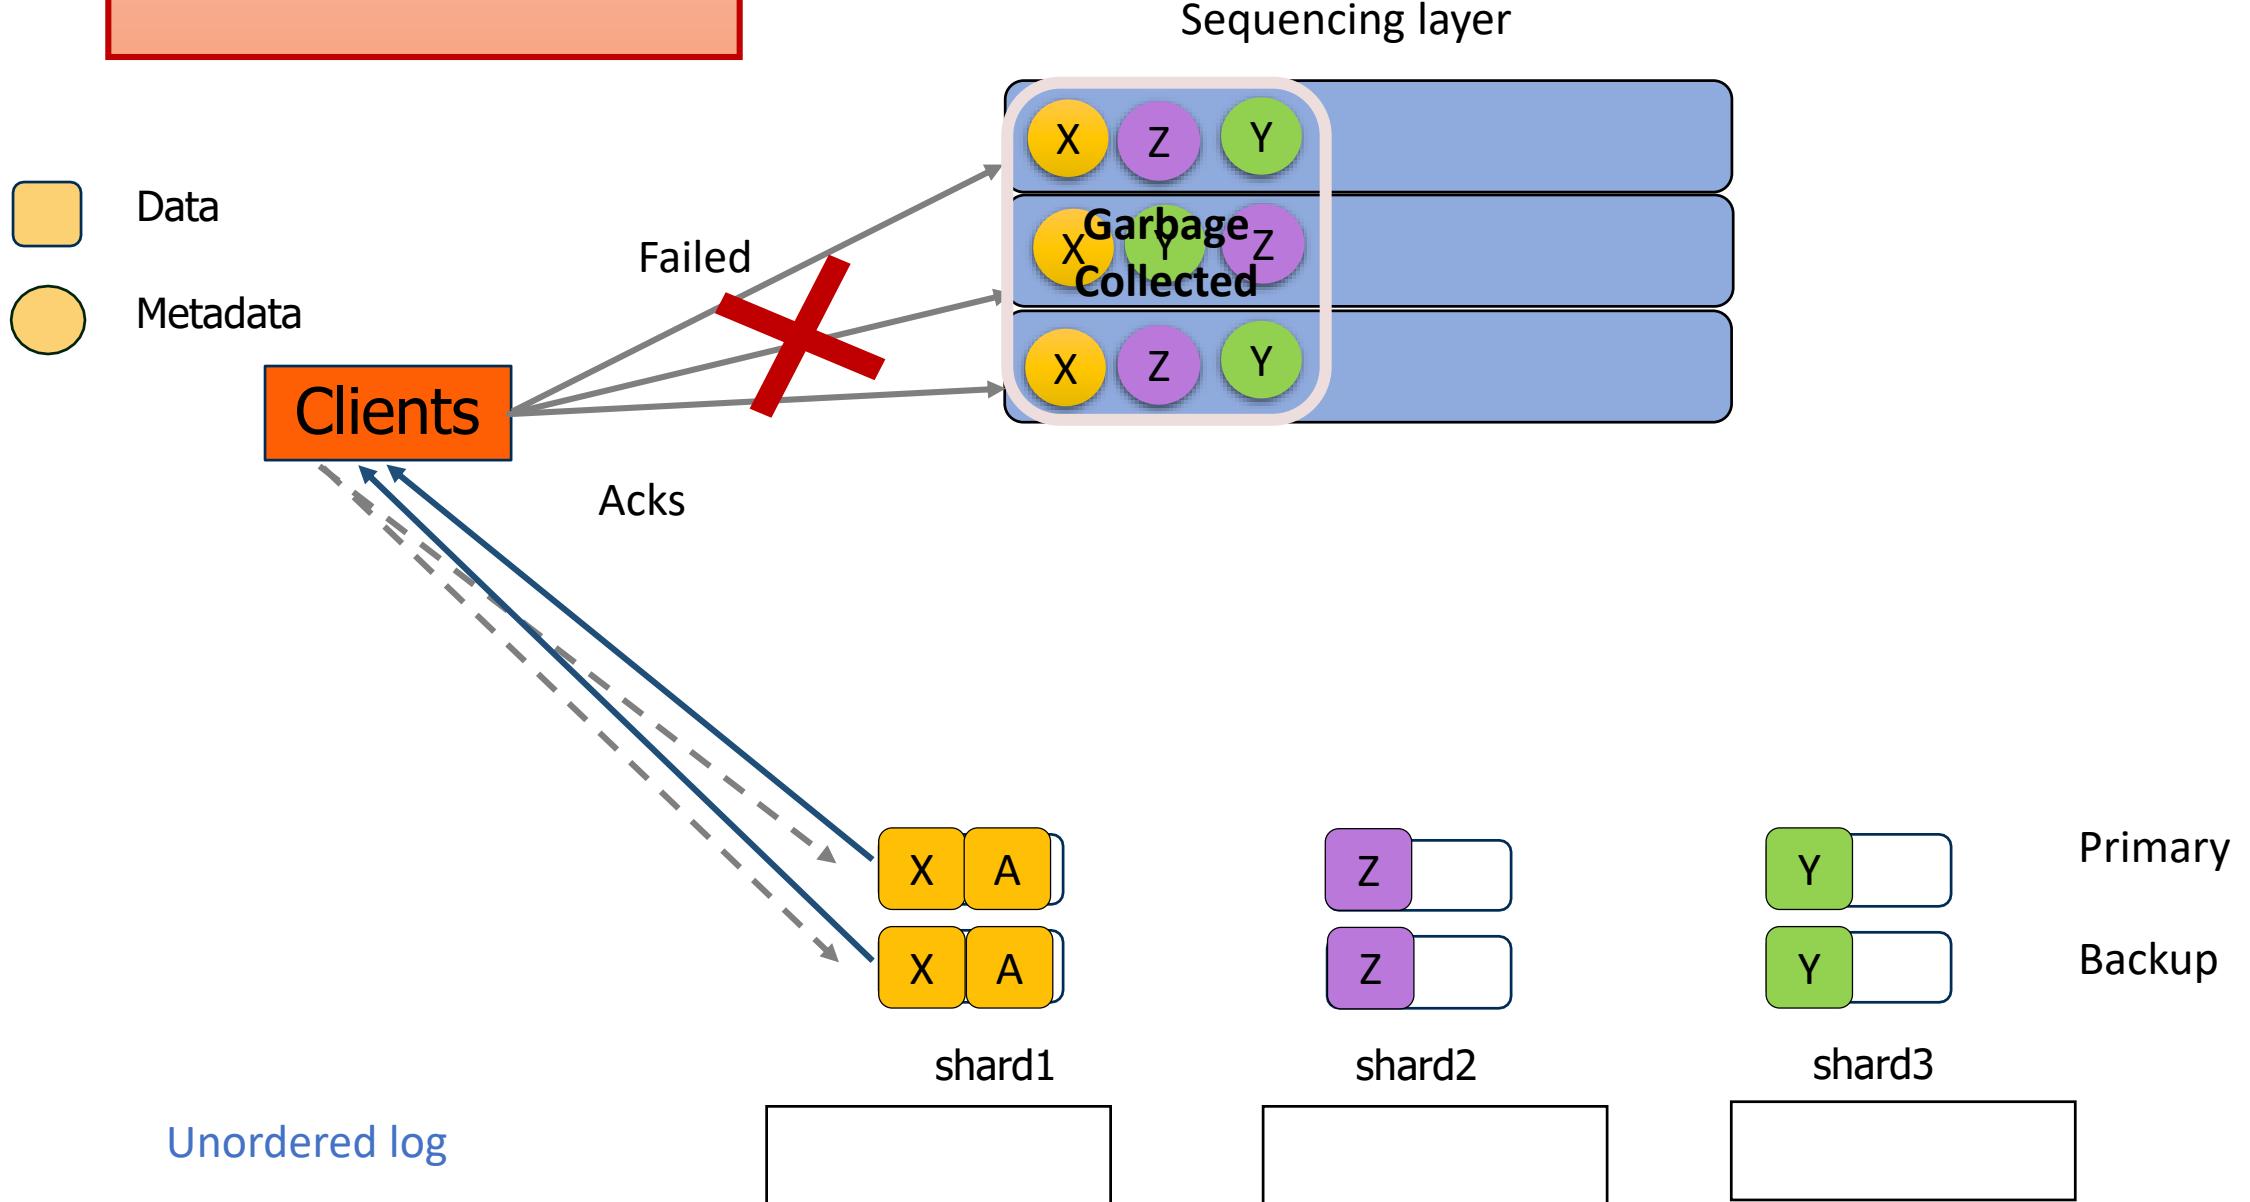
- Sequencing replica failure
- Shard Failure

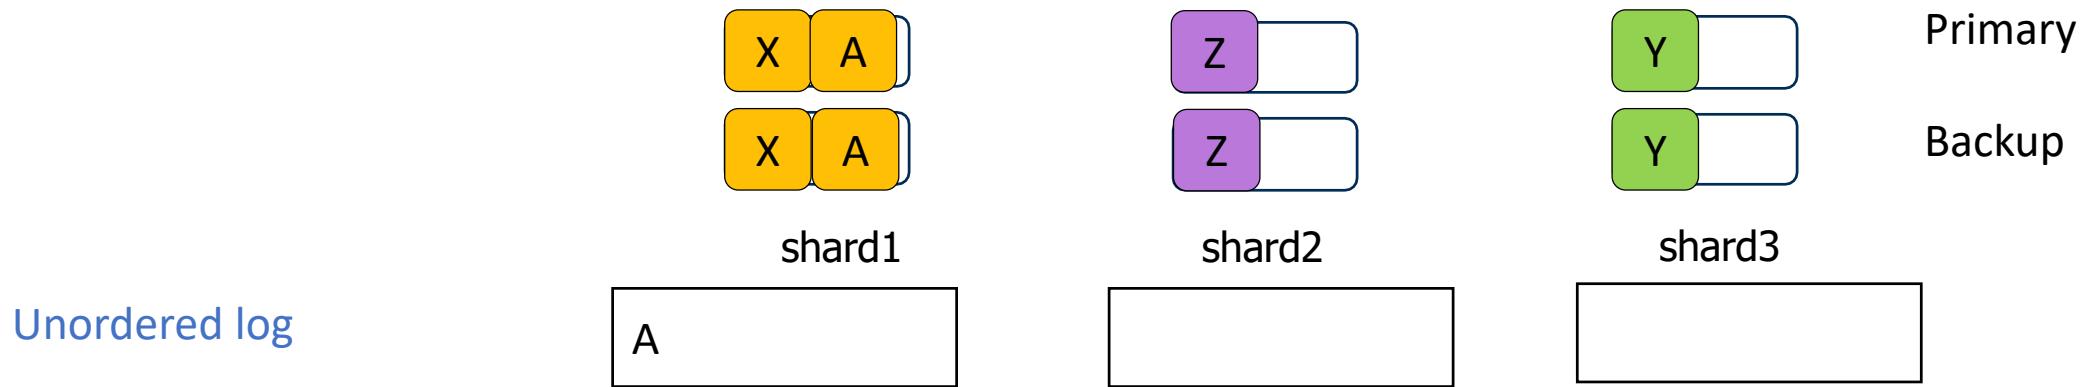
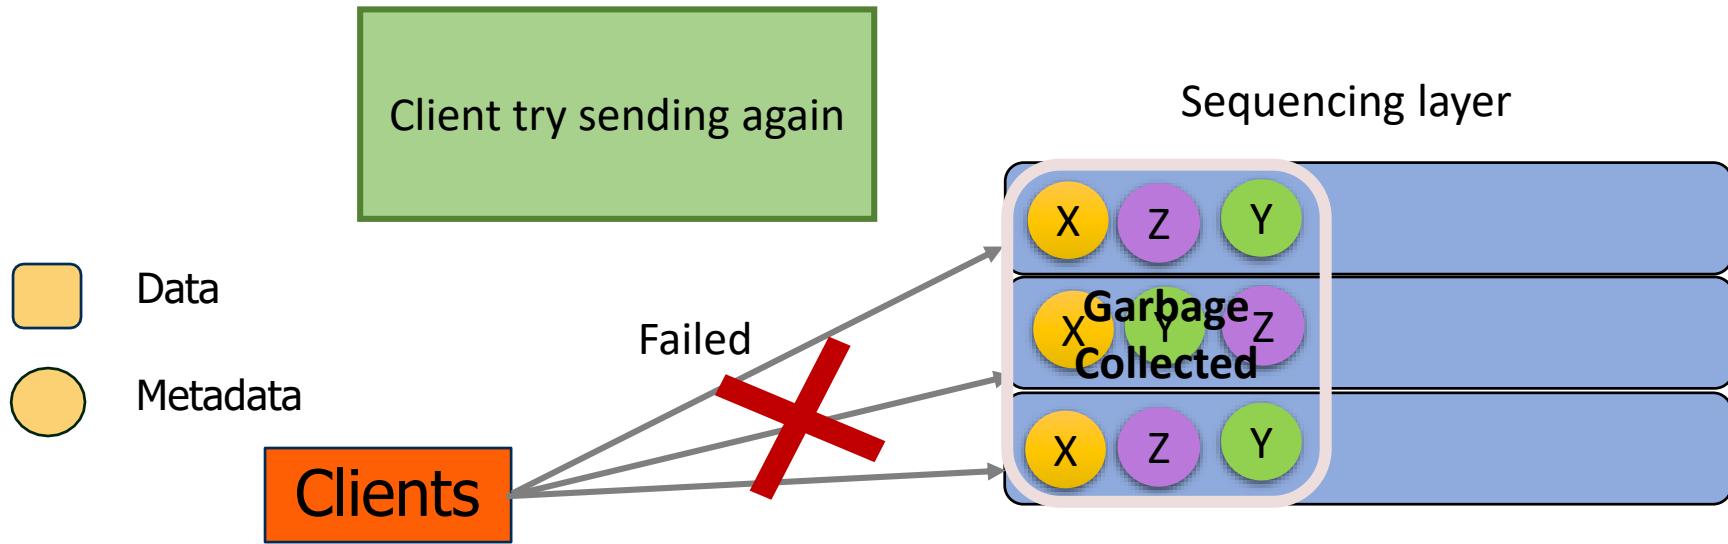
Same as Erwin Blackbox

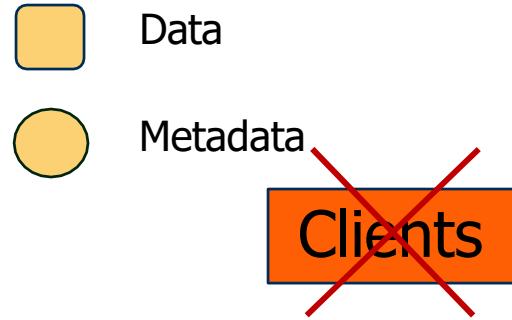
Client
Failures?

Two Problems

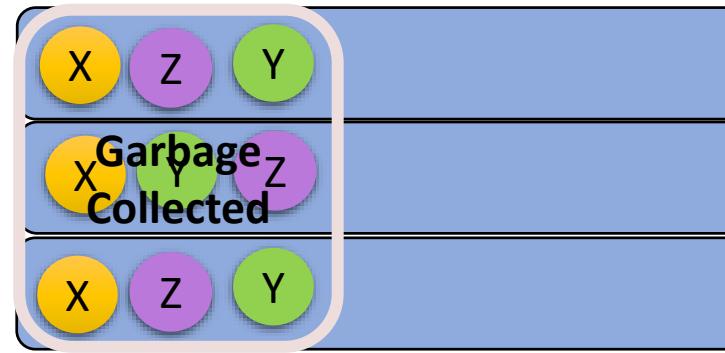
Problem-1



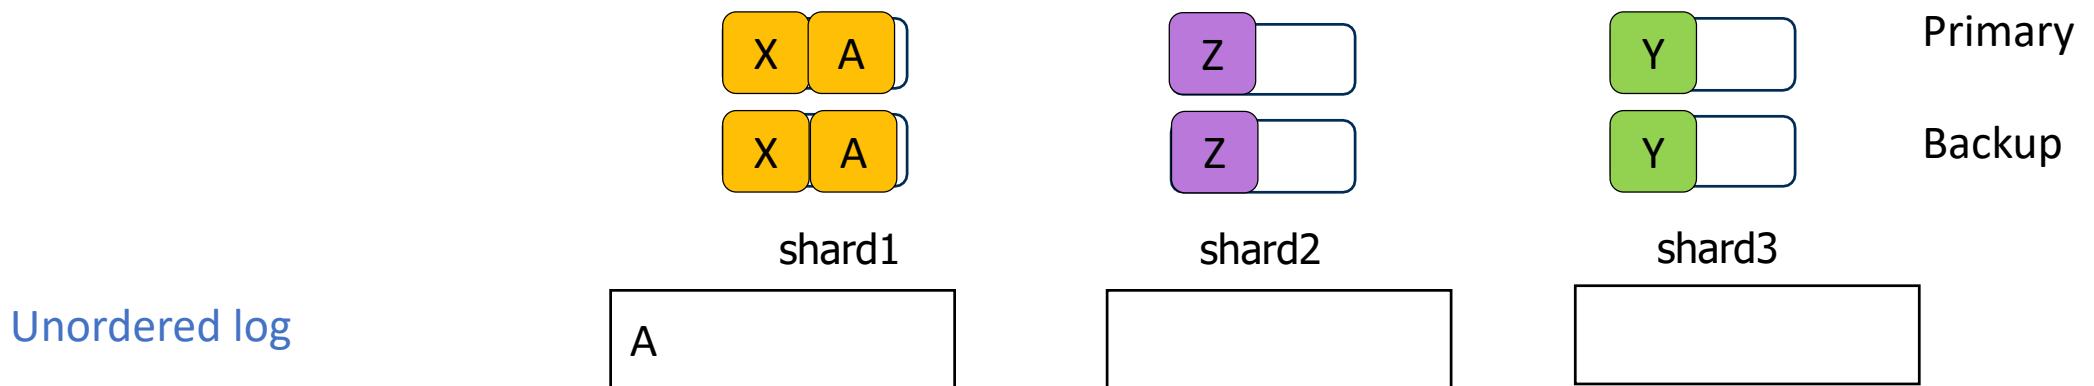


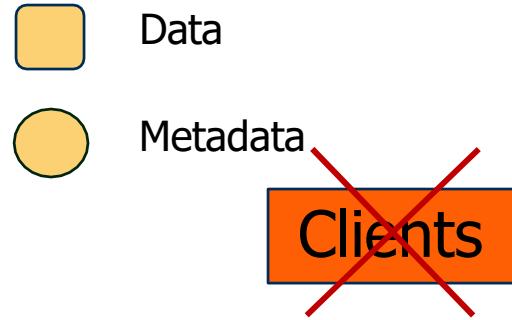


Sequencing layer

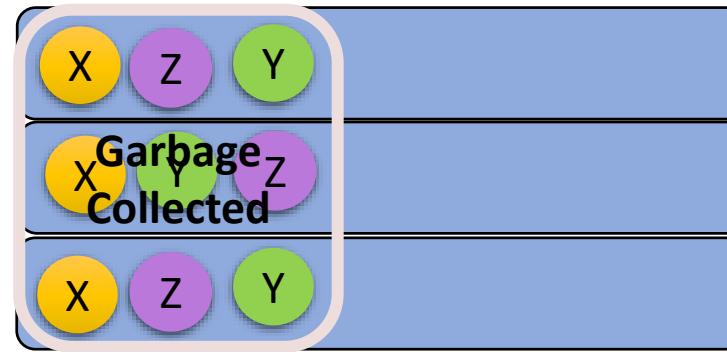


Data is Orphaned

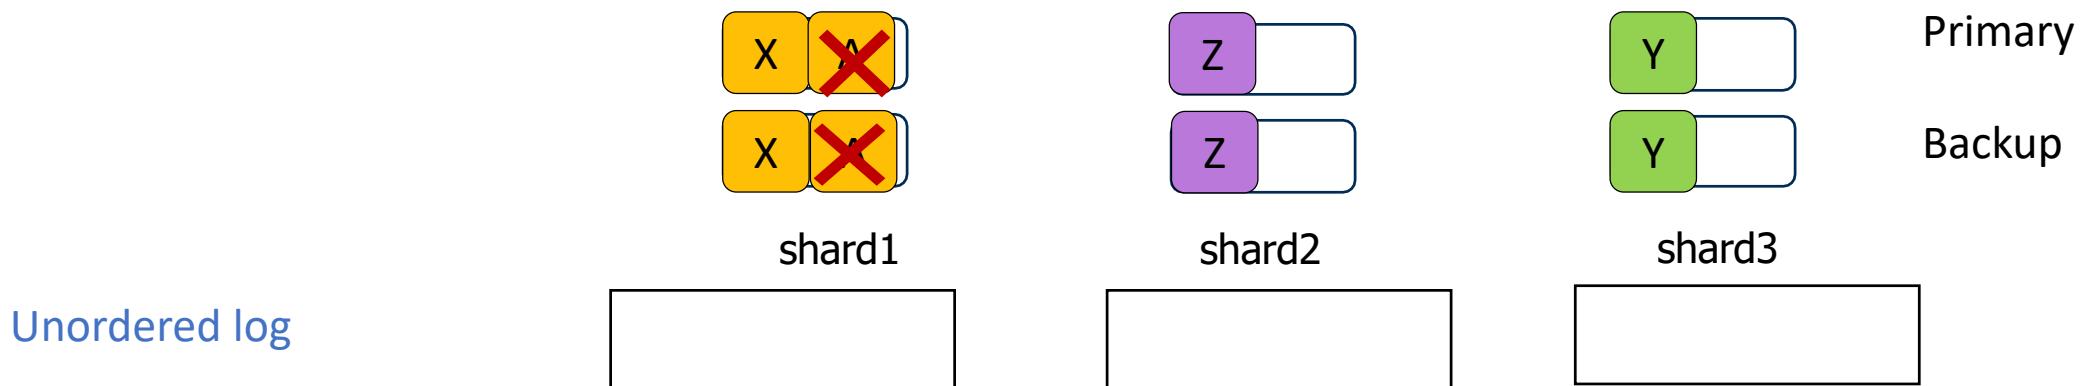




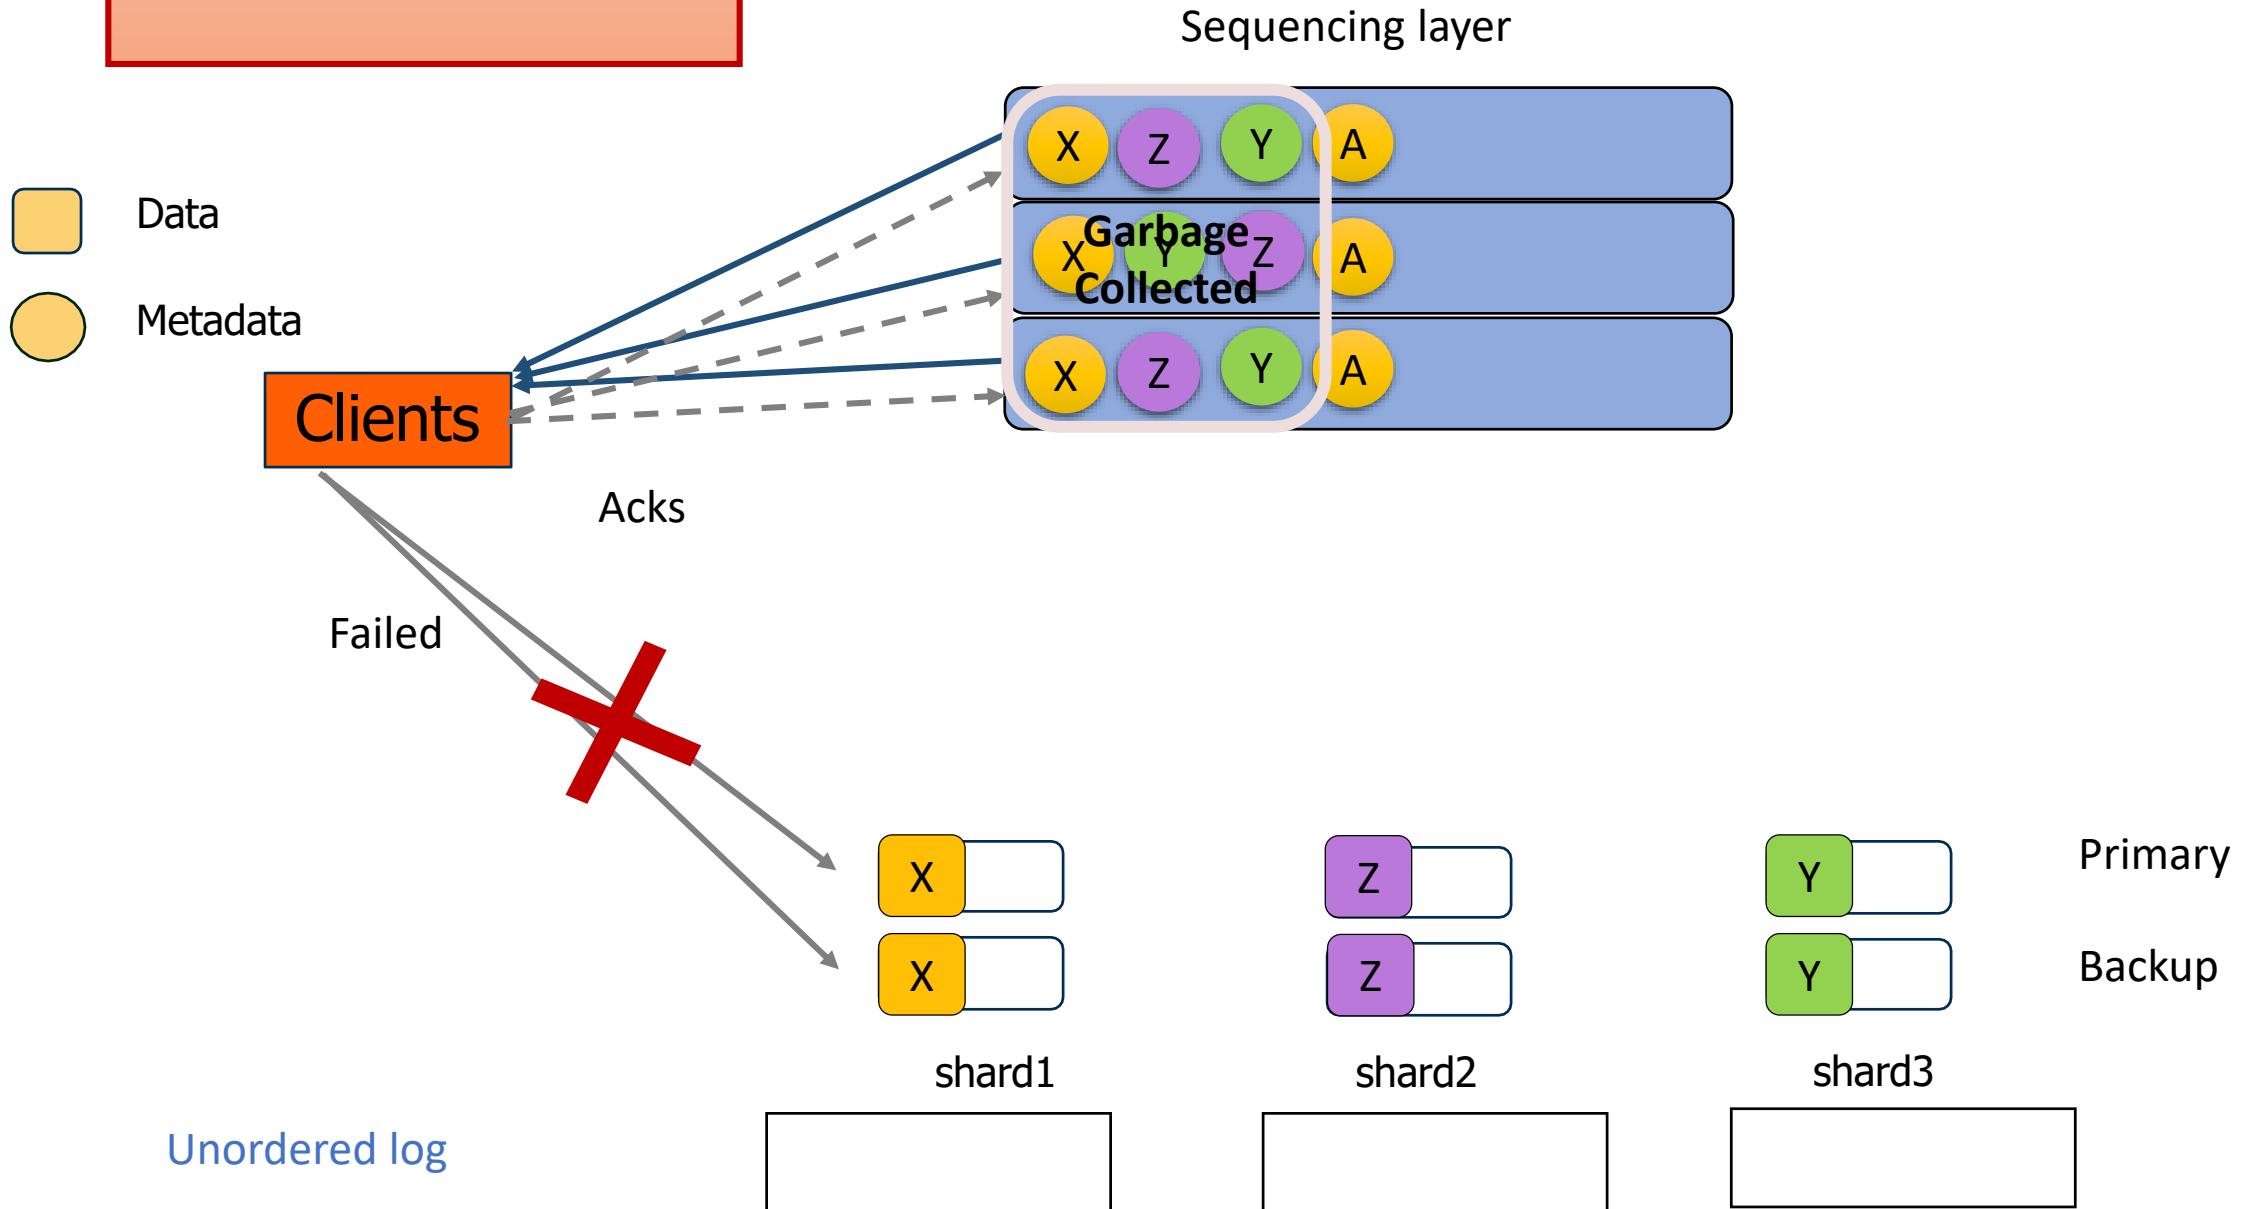
Sequencing layer



Data is Orphaned



Problem-2



Problem-2

 Data

 Metadata

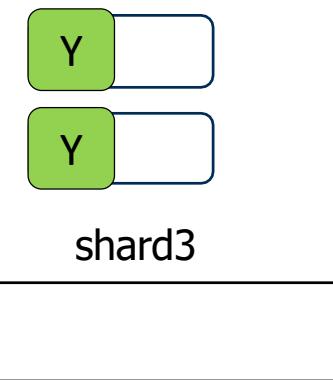
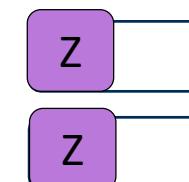
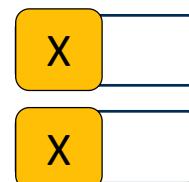
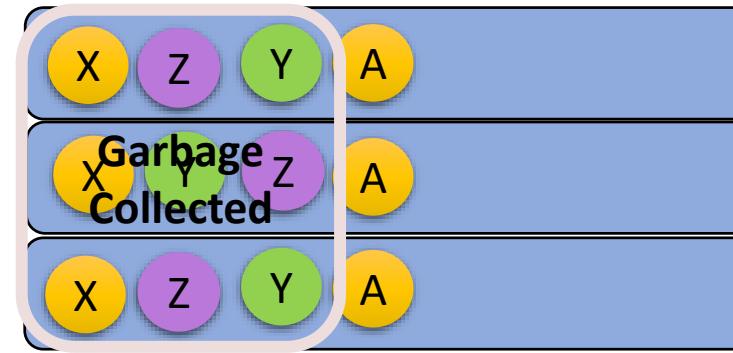
Clients

Failed

Client try sending again

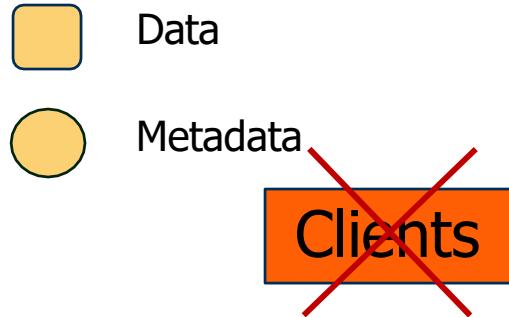
Unordered log

Sequencing layer



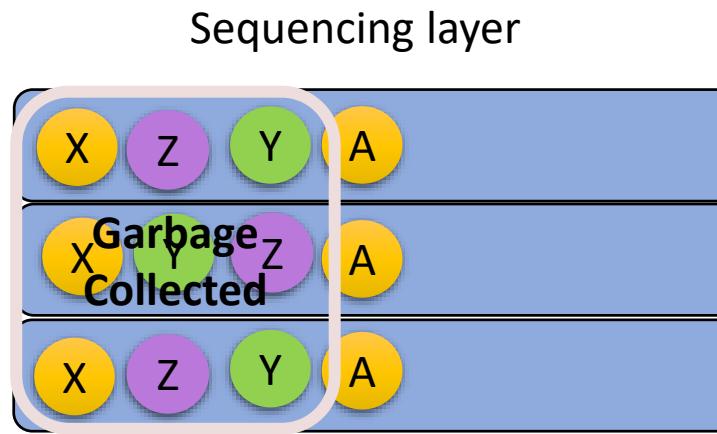
Primary
Backup

Problem-2

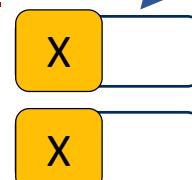


No Data with
record_id

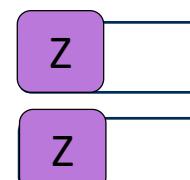
Unordered log



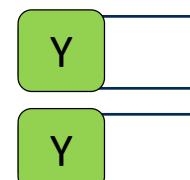
Sends position of
the record



shard1



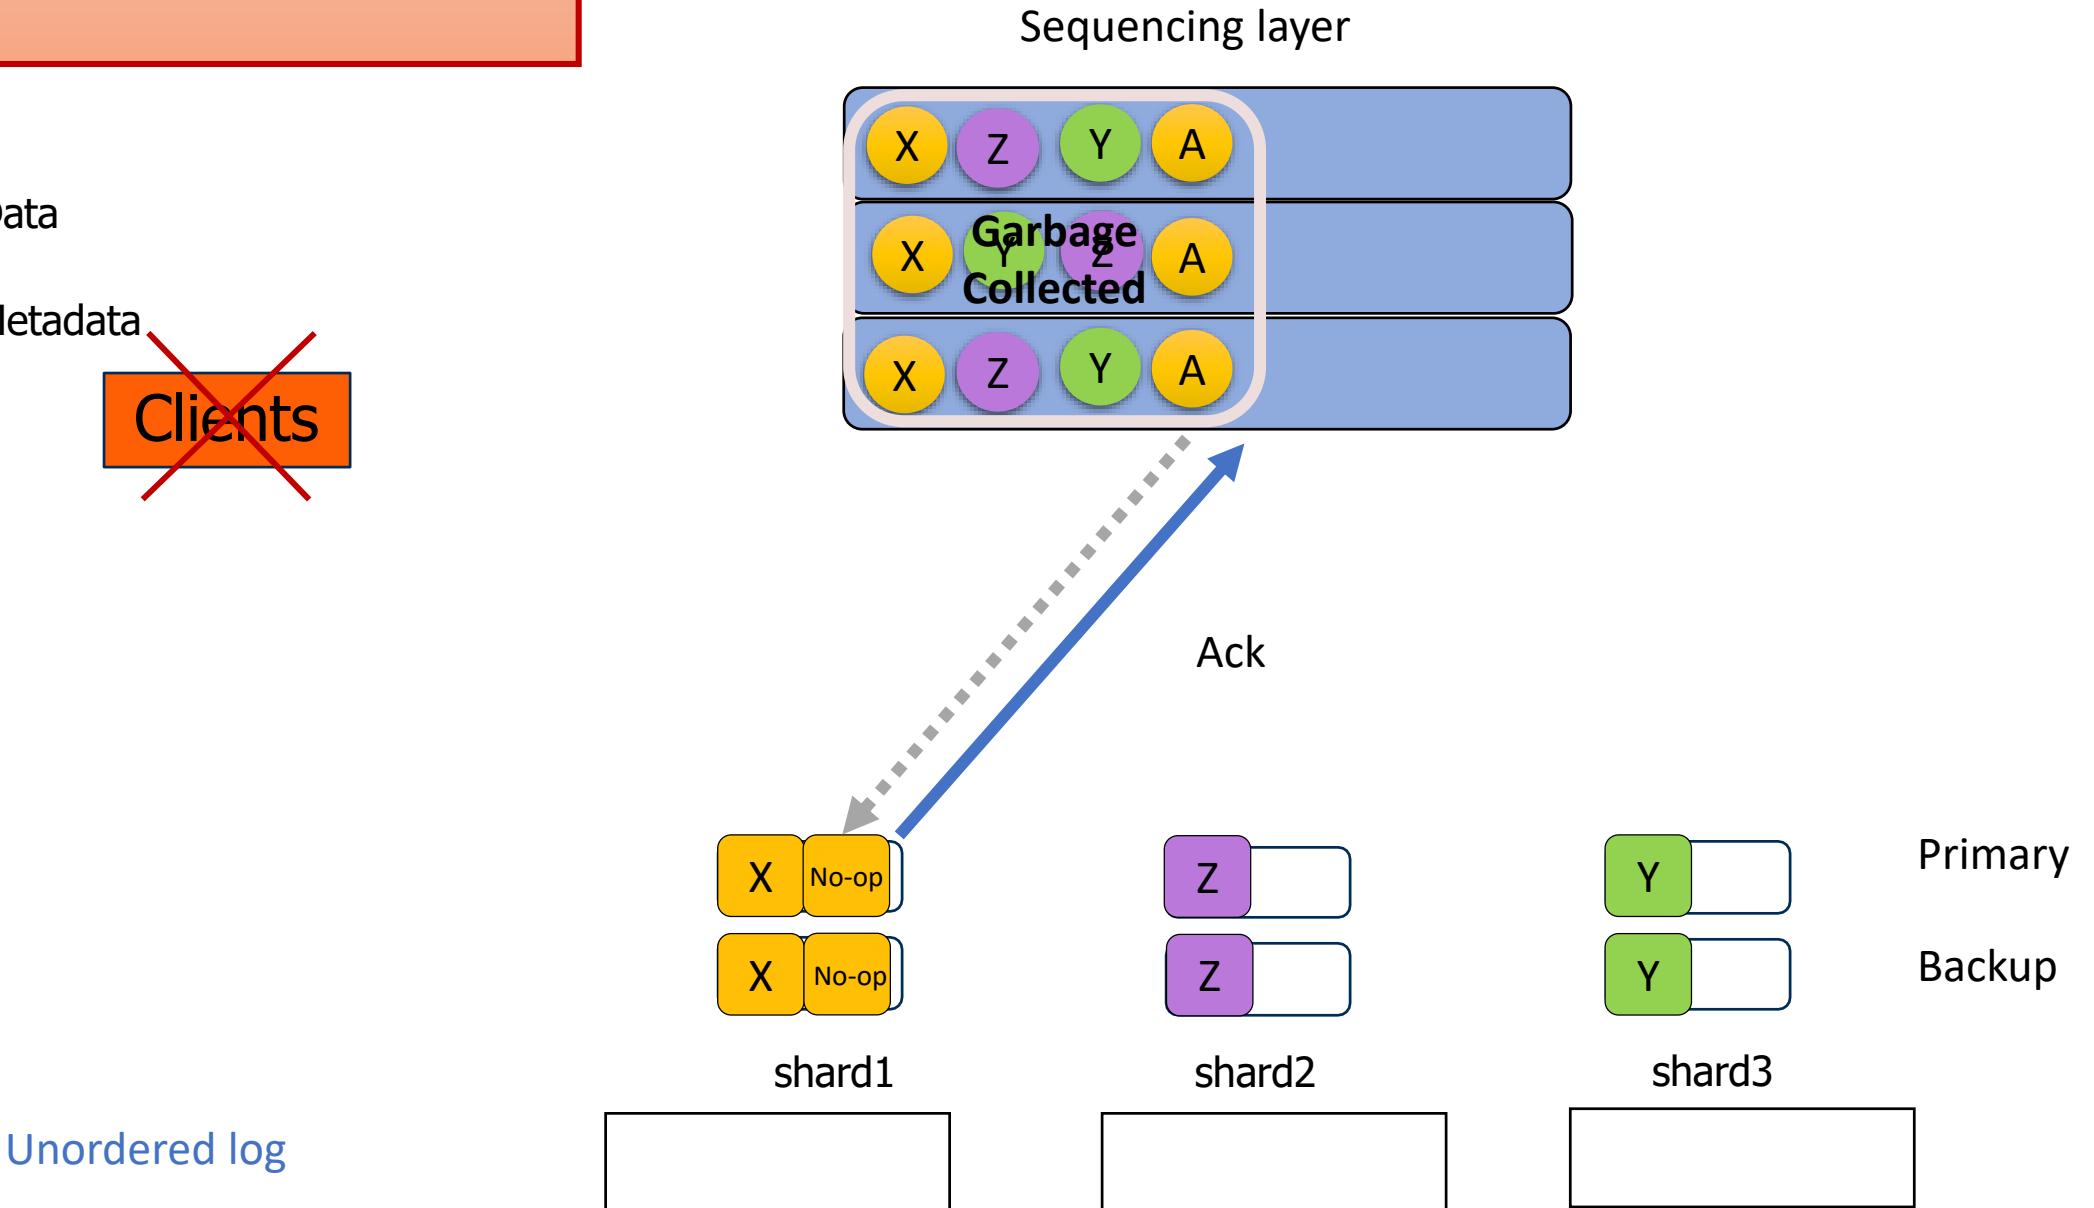
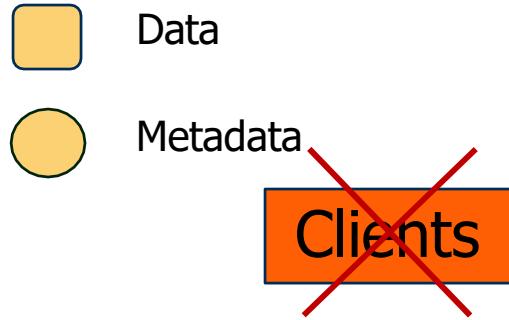
shard2



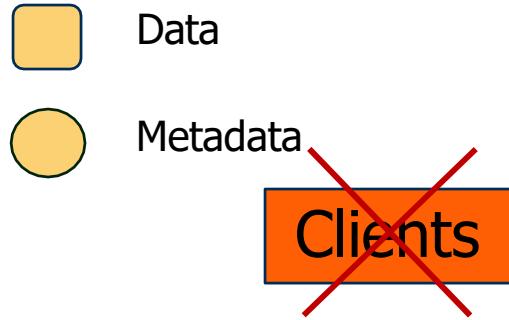
shard3

Primary
Backup

Problem-2

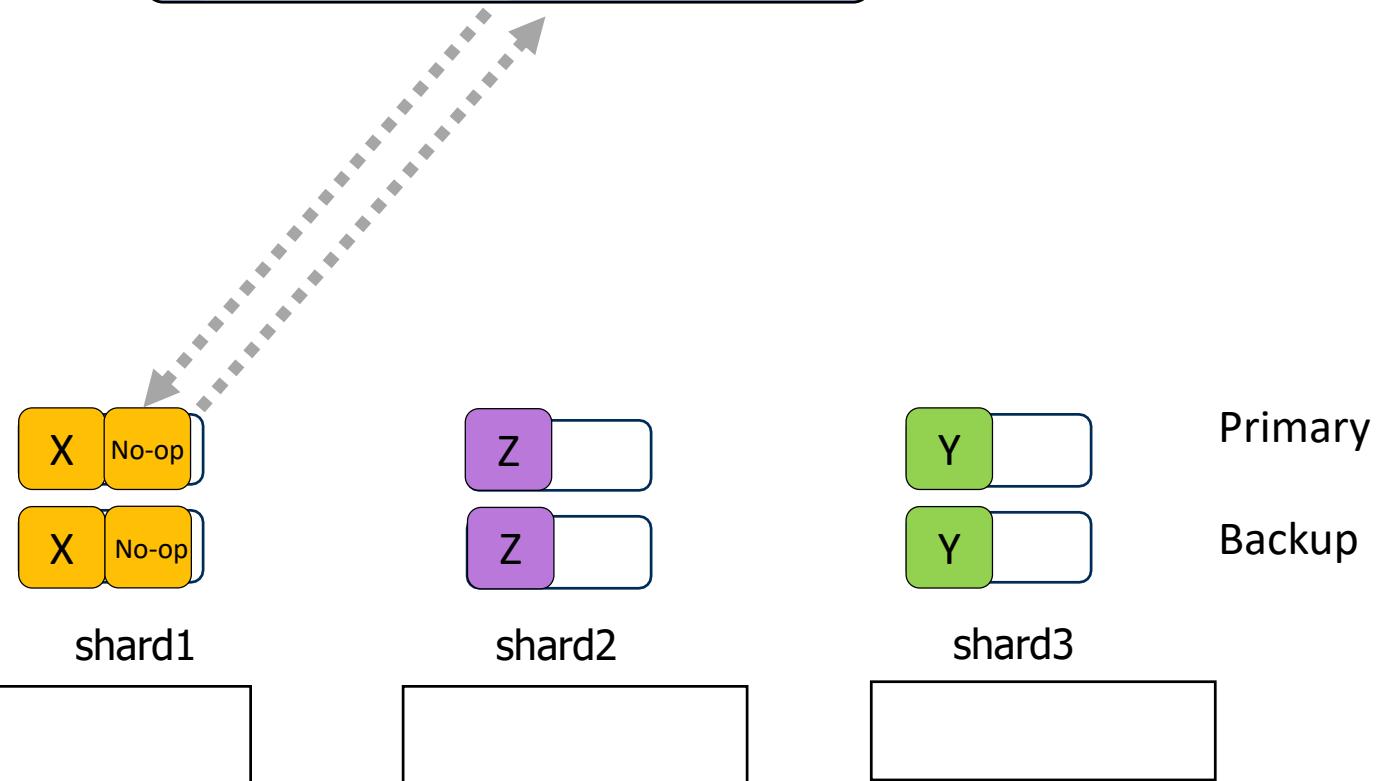
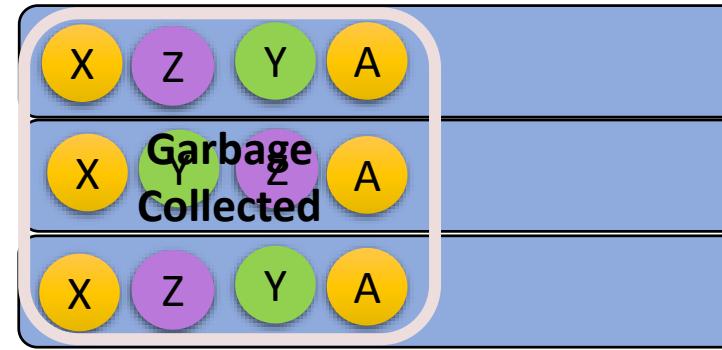


Problem-2



Unordered log

Sequencing layer



Performance Evaluation

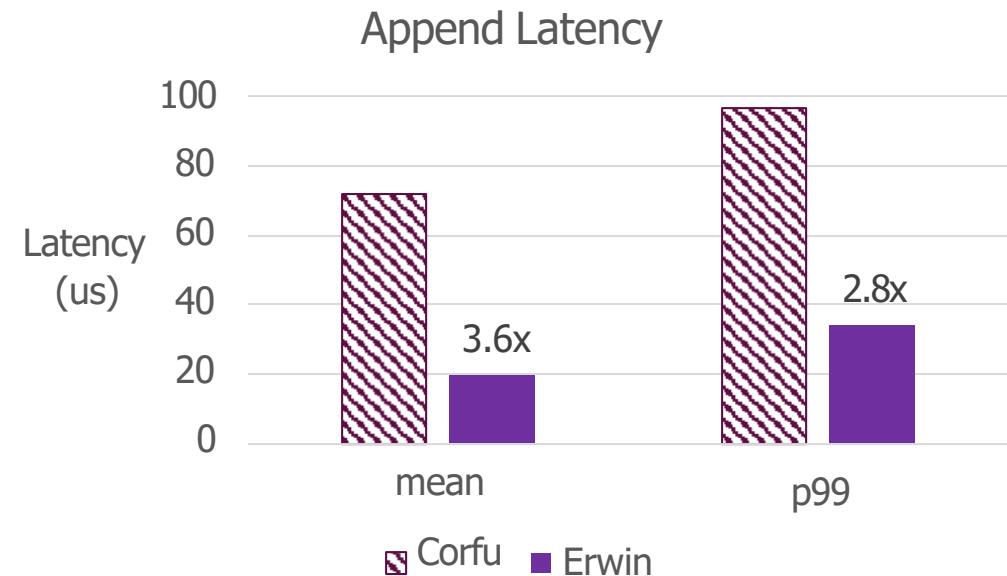
- What's the latency benefit of lazy ordering?
- How do reads perform in LazyLog?
- Do end applications benefit?

What's the Latency Benefit of Lazy Ordering?

Workload: 4KB record append-only
3 replicas per shard with 5 shards

Erwin reduces append latency

- Avg: By 3.6x compared to Corfu
- P99: By 2.8x compared to Corfu



How Do Reads Perform in LazyLog?

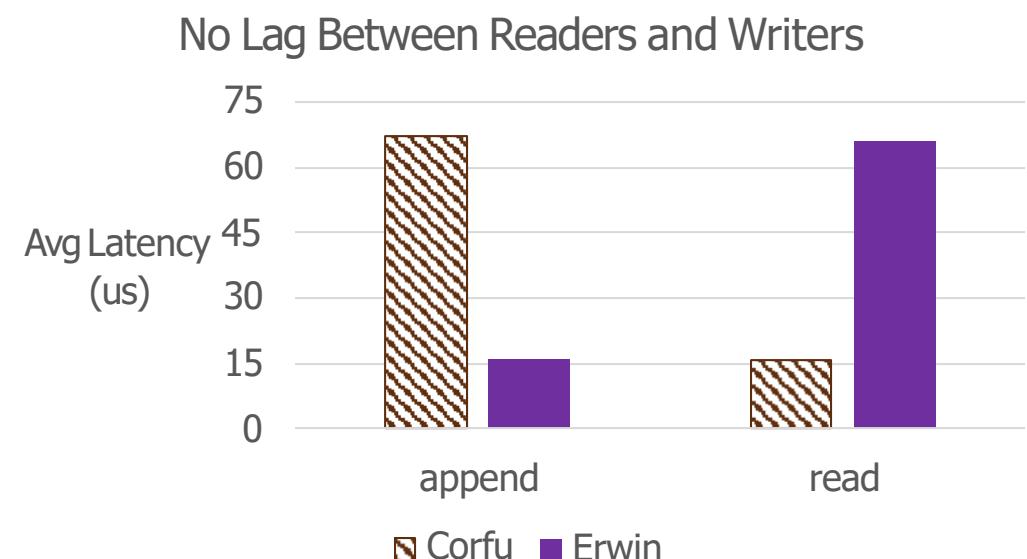
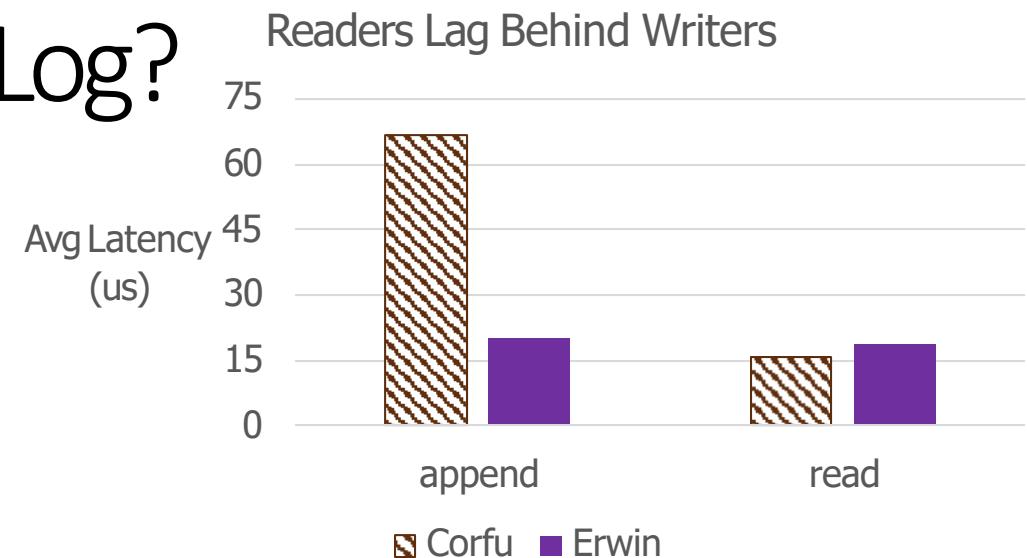
4KB record read after append

For many applications in which reads lag behind writes:

- Erwin achieves low append latency and read latency

In the worst case when there is *no* lag:

- Erwin shifts ordering cost from append to read
- Append +read latency remains the same



Do End Apps Benefit from LazyLog?

Built 3 Apps: *KV Store*, *Audit Log*, and *Journal for stream processing system*

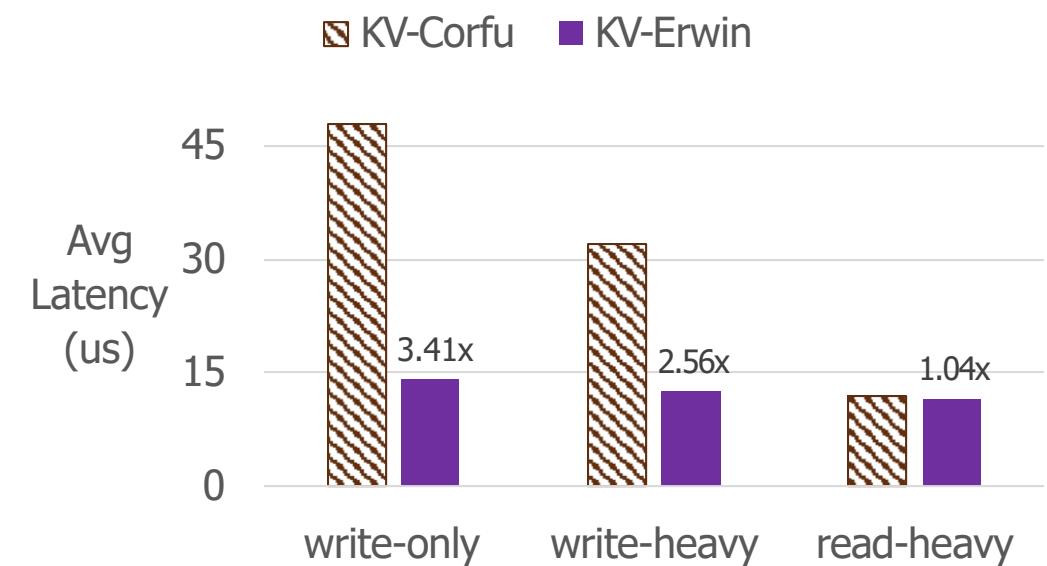
KV Store (decoupled WR-er and RD-er):

Append to log on PUTs

Reader reads log, constructs state, serves
GETs

Erwin benefits applications by reducing
ingestion latency

- Benefit is more pronounced when shared-log interaction takes significant partition of app request execution



Summary

- Eager-ordering shared logs incur high latencies, impacts app performance
- Eager ordering is not needed for many applications and readers are time-decoupled from writers
- LazyLog – a new shared-log abstraction that defers ordering
- Low ingestion latency with little overhead upon reads
- LazyLog systems deliver benefits for applications